

Building Solutions with Progress Agentic RAG

The complete practitioner's guide — start here

This is the master index for the book — seven parts plus a reference appendix. Read the parts in order for a guided build, or jump to the reference to look up any endpoint, data model, or verification status.

Key idea Every request, response, and code sample in the guide chapters was executed against a live Progress Agentic RAG Knowledge Box in the `aws-ap-southeast-2-1` zone. The reference appendix is generated from the official OpenAPI specifications (global, nua, nucliadb, zone — v1) and documents every endpoint, parameter, response, and schema. Appendix G records how each endpoint was verified.

How to read this book

- **New to the platform?** Read Part 1, then do the Chapter 3 quickstart with your own key.
- **Confused by keys and 403s?** Chapter 4 maps every credential to the APIs it can call.
- **Building ingestion?** Part 3. **Building search/answers?** Part 4 — the core of the book.
- **Configuring RAG for a specific use case?** Chapter 13 — a cookbook of tested recipes.
- **Need advanced, multi-source retrieval?** Part 5, the Retrieval Agent.
- **Going to production?** Part 6, especially Chapter 19.
- **Looking up an endpoint or data model?** Appendices A–F. **What's proven?** Appendix G.

Contents

Part 1 — Foundations

- 1. What Is Progress Agentic RAG?
- 2. Core Concepts and Vocabulary
- 3. Your First Solution: A Guided Quickstart

Part 2 — Authentication and Management

- 4. Authentication and Authorization

- 5. Accounts, Zones, and Knowledge Box Management

Part 3 — Data Ingestion

- 6. Resources and Fields: The Data Model
- 7. Uploading and Processing Files
- 8. Data Augmentation: Ingestion Agents

Part 4 — Search, Retrieval, and RAG

- 9. Retrieval Fundamentals: /find and /search
- 10. Generative Answers: The /ask Endpoint
- 11. Catalog, Suggest, and Summarize
- 12. The Knowledge Graph and Graph Search
- 13. Configuring RAG for Your Use Case (cookbook)

Part 5 — The Retrieval Agent

- 14. Agentic Retrieval, Sessions, and Workflows

Part 6 — Building and Operating

- 15. The SDKs: Python, JavaScript/TypeScript, and .NET
- 16. No-Code: The Search Widget and Dashboard
- 17. Quality and Observability
- 18. Backups, Export/Import, and Sync Connectors
- 19. Production Best Practices

Part 7 — Reference Appendix

- A. NucliaDB REST API — full endpoint reference
- B. NUA API — full endpoint reference
- C. Zone API — full endpoint reference
- D. Global / Account API — full endpoint reference
- E. Glossary and Schema Quick Reference
- F. Complete Schema Catalog — all 1,140 data models
- G. API Verification Matrix — how every endpoint was tested

Scope and verification

The guide chapters teach the full pipeline with cURL, Python, and JavaScript/TypeScript examples, culminating in Chapter 13's cookbook of tested RAG configurations for real use cases. Appendices A–D document all **356 operations** — every parameter, request-body field (nested objects expanded),

authorization role, and response. Appendix F documents all **1,140 named schemas** field by field. Appendix G reports live verification: **235 of 356 operations (66%) were exercised against the sandbox**, the remainder skipped only for safety, credential, or agent-provisioning reasons — each marked explicitly.

Building Solutions with Progress Agentic RAG

The complete practitioner's guide — grounded in a live, tested Knowledge Box

Progress Agentic RAG is a Retrieval-Augmented Generation platform delivered as a service. It ingests your unstructured data in any format and any language, indexes it for both keyword and semantic retrieval, and lets you ask questions in natural language and receive trusted, cited answers generated by a large language model of your choice. This book is a hands-on guide to building real solutions on that platform — from the first API call to a production deployment.

About this book

Most API documentation tells you what an endpoint *is*. This book shows you what an endpoint *does* — because every request and response you will read here was executed against a live Progress Agentic RAG Knowledge Box while the book was being written. When you see a JSON response, an answer, a citation, or an error message, it is a real one captured from the platform, not an idealized illustration.

The book is organized as a **guide with a full reference appendix**. Parts 1 through 6 are a narrative you can read front to back: they build a mental model, then walk through ingestion, retrieval, generative answers, the knowledge graph, the Retrieval Agent, the SDKs, and production concerns. Part 7 is an endpoint-by-endpoint reference generated from the platform's official OpenAPI specifications, so you can look up any parameter without leaving the book.

Who this book is for

- **Application developers** integrating search and generative answers into a product, using the REST API or one of the official SDKs (Python, JavaScript/TypeScript, .NET).
- **Solution engineers and architects** designing a RAG pipeline: choosing models, shaping the data model, tuning retrieval, and enforcing security and multi-tenancy.
- **Technical teams evaluating the platform** who want to understand exactly what it can do before committing, with concrete, reproducible examples.

You do not need prior RAG experience. You do need to be comfortable reading HTTP requests and a little Python or JavaScript. Chapter 2 defines every term the platform uses.

How everything in this book was tested

A dedicated, throwaway Knowledge Box was created in the `aws-ap-southeast-2-1` (Australia, AWS) zone and loaded with a small synthetic corpus about home energy — solar panels, battery storage, and electric-vehicle charging. Every capability in the book was exercised against that Knowledge Box: resources were created and read back, files were uploaded, searches were run, generative answers were produced and streamed, and the knowledge graph was queried. The exact responses appear throughout.

Key idea Because the examples are real, they are also honest. Where the platform returns an error, a surprising default, or a limit, this book shows it — those moments are marked as **Gotcha** callouts and are some of the most useful pages here.

Endpoint signatures, parameter names, types, defaults, and enumerations are taken directly from the four official OpenAPI specifications that Progress publishes for the platform:

Specification	Covers	Base URL pattern
<code>global</code> (v1)	Accounts, users, zones	<code>https://rag.progress.cloud/api/v1</code>
<code>nucliadb</code> (v1)	Knowledge Box content, search, RAG	<code>https://<zone>.dp.progress.cloud/api/v1</code>
<code>nua</code> (v1)	Processing, Predict, agents	<code>https://<zone>.dp.progress.cloud/api/v1</code>
<code>zone</code> (v1)	KB management within a zone	<code>https://<zone>.dp.progress.cloud</code>

Conventions used throughout

Code samples appear in three flavors wherever it makes sense: a raw `cURL` call that maps directly to the HTTP request, the **Python SDK** (`nuclia`), and the **JavaScript/TypeScript SDK** (`@nuclia/core`). Reading the `cURL` first is the fastest way to understand what any SDK method does under the hood.

Placeholders are written in uppercase: `ZONE` is your zone slug (for example `aws-ap-southeast-2-1`), `KB` is a Knowledge Box id, `ACCOUNT` is an account id, and `KEY` is whichever token authorizes the call. The chapter on authentication explains which key each call needs.

Callouts flag things worth pausing on:

Note Background or clarifying detail.

Tip A practical shortcut or recommended practice.

Warning Something that can cause data loss, cost, or a hard-to-debug failure.

Gotcha — tested A real, tested surprise — a non-obvious default, limit, or error the platform returned during authoring.

A note on names

The platform is built on the open-source **NucliaDB** engine and the **Nuclia Understanding API (NUA)**. Progress Agentic RAG is the commercial, managed offering. You will see the names *Nuclia* and *NucliaDB* in URLs, headers (**X-NUCLIA-NUAKEY**), SDK package names, and the OpenAPI specifications — they refer to the same platform this book describes.

PART 1 — FOUNDATIONS

Chapter 1

What Is Progress Agentic RAG?

The platform, the problem it solves, and how the pieces fit together

Progress Agentic RAG is a Retrieval-Augmented Generation platform delivered as a managed service. In one sentence: it turns a pile of your unstructured content into a searchable, question-answering knowledge base that a large language model can reason over — without you having to build and operate the retrieval infrastructure yourself.

This chapter builds the mental model the rest of the book depends on. If you already know what RAG is and just want to make calls, skip to Chapter 3's quickstart; but the ten minutes you spend here will make every later chapter easier.

The problem RAG solves

Large language models are fluent but forgetful. They know only what was in their training data, they cannot see your private documents, and when asked about something they do not know they tend to invent a plausible answer. For a consumer chatbot that is an annoyance. For a legal, medical, compliance, or commerce application it is a liability.

Retrieval-Augmented Generation fixes this by changing the order of operations. Instead of asking the model to answer from memory, you first **retrieve** the most relevant passages from your own content, then hand those passages to the model and ask it to answer *using only what it was given*. The model's fluency is preserved; its tendency to hallucinate is constrained by evidence you control. Done well, every answer can point back to the exact source paragraph that supports it.

Key idea RAG is not a model. It is a pipeline: **ingest** → **index** → **retrieve** → **generate**. Progress Agentic RAG is a production-grade implementation of that entire pipeline, exposed through an API.

What "as a service" actually buys you

Building a RAG pipeline by hand means assembling and operating a surprising number of moving parts: document parsers for dozens of file formats, an OCR and transcription stack for images and audio, a text

chunker, an embedding model, a vector database, a keyword index, a reranker, prompt templates, an LLM integration, and the glue that keeps them consistent as data changes. Each piece is a project in its own right.

Progress Agentic RAG provides all of it behind a single API and a dashboard. Specifically, the platform gives you:

- **Ingestion for any data, any language** — upload files, links, or raw text; the platform extracts, cleans, transcribes, and chunks the content automatically.
- **A hybrid index** — every piece of content is indexed for both keyword (BM25) and semantic (vector) retrieval, so you get exact-match precision and meaning-based recall at once.
- **Your choice of models** — the semantic (embedding) model and the generative (LLM) model are configurable per Knowledge Box; you can even bring your own model provider account.
- **A customizable retrieval and RAG pipeline** — you control filtering, ranking, reranking, how much context is assembled, and the exact prompt sent to the LLM.
- **Trusted answers with citations** — generative answers can carry citations back to the source paragraphs, so answers are auditable.
- **An agentic layer** — the Retrieval Agent can plan multi-step retrieval, call tools, fall back to the public internet, and hold stateful conversations.

A first look at the shape of the API

Everything in the platform hangs off one central object: the **Knowledge Box**. A Knowledge Box is an isolated container of content plus its indexes and configuration. You ingest resources into a Knowledge Box, and you search, ask, and generate against that same Knowledge Box. A single account can hold many Knowledge Boxes, each fully isolated from the others.

The two endpoints you will use most illustrate the whole platform in miniature:

Endpoint	What it does	Returns
<code>POST /kb/{kbid}/find</code>	Hybrid retrieval — finds the most relevant passages for a query	Ranked resources and text blocks
<code>POST /kb/{kbid}/ask</code>	Retrieval + generation — finds passages, then asks the LLM to answer from them	A generated answer, its citations, and the retrieval that grounded it

`ask` internally calls `find` to gather evidence, then sends that evidence to the LLM. Because of that, `ask` accepts every parameter `find` does — plus a large set that control answer generation. Chapters 9 and 10 cover both in depth.

Note Throughout the book, `find` and `ask` are shown against a live Knowledge Box loaded with a small corpus about home solar, batteries, and EV charging. You will see the same corpus answer real

questions in Chapter 3.

Where the platform runs: zones, accounts, and the four APIs

Two facts about the platform's topology will save you hours of confusion, so they are worth stating up front.

Zones

Your data lives in a **zone** — a geographic region such as `aws-ap-southeast-2-1` (Australia) or `aws-eu-central-1-1` (Europe). A Knowledge Box belongs to exactly one zone, and content operations for that Knowledge Box are served from that zone's regional host. Choosing a zone is primarily about data residency and latency.

The four APIs

The platform is not a single API but four, each with its own base URL and its own job. Knowing which is which tells you immediately where a given call goes and which credential it needs.

API	Responsibility	Base URL
Global	Accounts, users, zone discovery	<code>https://rag.progress.cloud/api/v1</code>
Zone (Regional Manager)	Managing Knowledge Boxes within a zone — create, configure, delete, keys	<code>https://<zone>.dp.progress.cloud</code>
NucliaDB	Knowledge Box content — resources, fields, search, ask, graph	<code>https://<zone>.dp.progress.cloud/api/v1</code>
NUA	The Understanding API — processing, Predict, agents	<code>https://<zone>.dp.progress.cloud/api/v1</code>

Gotcha — tested These boundaries are enforced strictly. A credential that works on one API can be flatly rejected by another — a NUA key used against the Global API returns `Nuakeys are not valid in the global API`. Chapter 4 maps every credential to the APIs it can actually call. This is the single most common source of early confusion.

How this book is organized around the pipeline

The parts of this book follow the RAG pipeline in order, so you can read them as a build sequence:

- **Foundations (Part 1)** — concepts, vocabulary, and a working quickstart.
- **Authentication and Management (Part 2)** — credentials, and creating and configuring Knowledge Boxes.
- **Data Ingestion (Part 3)** — getting content in, and enriching it as it lands.
- **Search, Retrieval, and RAG (Part 4)** — the heart of the platform: finding, asking, and shaping answers.
- **The Retrieval Agent (Part 5)** — agentic, multi-step, stateful retrieval.
- **Building and Operating (Part 6)** — SDKs, the no-code widget, observability, and production practices.
- **Reference (Part 7)** — every endpoint, generated from the official specifications.

By the end of Part 4 you will be able to build a complete, high-quality question-answering solution. The remaining parts make it agentic, observable, and production-ready.

PART 1 — FOUNDATIONS

Chapter 2

Core Concepts and Vocabulary

The objects and terms the platform is built from

The platform is precise about its vocabulary, and its API uses these terms literally. This chapter defines each one and shows how they nest. Treat it as a reference to return to — the first time you meet `field_type`, `paragraph`, or `vectorset` in a later chapter, this is where the definition lives.

The containment hierarchy

From the outside in, the platform's objects nest like this:

```

TEXT
Account
├── Knowledge Box (belongs to one zone)
│   ├── Resource
│   │   ├── Field (text / file / link / conversation / key-value)
│   │   │   ├── Paragraph (a retrievable text block)
│   │   │   │   └── Vector (an embedding of a paragraph or sentence)

```

Retrieval operates at the **paragraph** level: when you search, the platform finds the most relevant paragraphs and then tells you which resource and field they came from. Understanding this hierarchy is the key to understanding both search results and filtering.

Account and zone

An **account** is the billing and ownership boundary — your organization. It holds users, Knowledge Boxes, and usage. A **zone** is the geographic region a Knowledge Box lives in. One account can own Knowledge Boxes across several zones.

Knowledge Box

A **Knowledge Box** (often abbreviated *KB*, and identified by a `kbid` UUID) is the primary container. It holds resources, the indexes built from them, and a configuration that includes which semantic and generative models to use, security settings, and labels. Knowledge Boxes are fully isolated from one

another — a search never crosses KB boundaries. Most real solutions use one Knowledge Box per tenant, per product area, or per security domain.

Resource

A **resource** is a single item of content you put into a Knowledge Box — a document, a web page, an image, an audio or video file, or a block of text. Each resource has a stable **uuid** and an optional human-friendly **slug**. A resource carries metadata (title, summary, labels, origin, security) and one or more **fields** that hold its actual content.

Tip Set a **slug** when you create a resource. A slug is a stable, human-readable id you choose (for example an external record id), and most content endpoints accept a slug in place of the UUID — which means you can address resources by your own identifiers.

Fields and field types

Content does not sit loose in a resource; it sits in a typed **field**. A single resource can hold several fields of different types — for example a **file** field for an uploaded PDF and a **text** field for an editorial summary. The field types are:

Field type	API key	Holds
Text	texts	Plain text, Markdown, HTML, or RST you supply directly
File	files	An uploaded binary — PDF, Office doc, image, audio, video
Link	links	A URL the platform will fetch and process
Conversation	conversations	A sequence of messages (e.g. a chat or support thread)
Key-value	key_values	Structured key/value data

Every field is addressed by a **field_type** letter and a **field_id** in search results and downloads — **t** for text, **f** for file, **u** for link, **a** for the generic/title field, and so on. You will see identifiers like **ca2a.../t/body/0-301** in results: that is resource **ca2a...**, text field **body**, character range 0–301.

Paragraph and vector

When a field is processed, its text is split into **paragraphs** — the retrievable unit. Each paragraph is indexed two ways: into a **keyword** index (BM25, for exact term matching) and into a **vector** index (an embedding produced by the Knowledge Box's semantic model, for meaning-based matching). A search can use either or both; using both is called *hybrid* retrieval and is the default.

A **vectorset** is a named vector index. A Knowledge Box can hold more than one — for example one per embedding model — and you can select which to search with the **vectorset** parameter. This is what makes it possible to migrate embedding models without rebuilding everything.

Labels, entities, and relations

Beyond raw text, the platform enriches content with structure you can filter and navigate by:

- **Labels** (also called classifications) are tags organized into **labelsets**. You can apply them manually or have an ingestion agent apply them automatically, then filter searches by them.
- **Entities** are named things — people, organizations, places, products — detected by the NER model or defined by you in **entity groups**.
- **Relations** connect entities and resources into a **knowledge graph** you can query directly (Chapter 12).

Processing and the NUA pipeline

Ingestion is asynchronous. When you create or upload a resource, the platform accepts it immediately and then **processes** it in the background through the Nuclia Understanding API (NUA): extracting text, transcribing media, running NER, chunking into paragraphs, and computing vectors. Only after processing completes is the content fully searchable.

Gotcha — tested Because processing is asynchronous, content is not instantly searchable. In testing, a small text resource was searchable within seconds, but a freshly uploaded file returned *"Not enough data to answer this"* from `/ask` until processing finished. Always account for this lag — Chapter 7 shows how to check processing status.

The credentials, in one place

Four kinds of credential appear in the book. Chapter 4 is devoted to them, but here is the one-line version so the code samples make sense sooner:

Credential	Authorizes	Typical use
NUA key	Processing/Predict, and (if enabled) KB management	Backend automation, this book's testing
KB API key (service account)	Content and search on one Knowledge Box	Application backends
Personal Access Token (PAT)	Everything your user can do	Long-running scripts, admin
User token	Everything, briefly (30 min)	Interactive/dashboard use

PART 1 — FOUNDATIONS

Chapter 3

Your First Solution: A Guided Quickstart

From zero to a cited, generated answer in five steps

This chapter walks the entire pipeline end to end: authenticate, create a Knowledge Box, ingest a little content, retrieve it, and ask a question that returns a generated, cited answer. Every command here was run against the live platform while writing — the responses are the real ones. By the end you will have a working solution you can adapt to your own data.

Note The quickstart uses a **NUA key with KB-management enabled** to create the Knowledge Box, then a **Knowledge Box API key** for content and search. Chapter 4 explains why the handoff between the two is necessary; for now, follow the steps.

Before you start

You need three things: your **zone** slug (this walkthrough uses `aws-ap-southeast-2-1`), a **NUA key** whose token encodes `allow_kb_management: true`, and either `curl`, Python 3.9+, or Node 18+. Set two environment variables so the samples are copy-pasteable:

BASH

```
export ZONE="aws-ap-southeast-2-1"
export NUA_KEY="eyJhbGciOiI...your-nua-key..."
```

Step 1 — Create a Knowledge Box

Knowledge Boxes are created through the **Zone API**, under your account. Creating one requires only a **slug** and a **title**. The NUA key is presented in the `X-NUCLIA-NUAKEY` header with the **Bearer** scheme.

BASH

```
curl -X POST "https://$ZONE.dp.progress.cloud/api/v1/account/$ACCOUNT/kbs" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY" \
```

```
-H "Content-Type: application/json" \
-d '{"slug":"quickstart","title":"Quickstart KB"}'
```

The response returns the new Knowledge Box id:

```
JSON
{ "id": "e3bb0d28-3821-44a4-8948-fe511f48aa42" }
```

Save that id as `KB`. In Python and JavaScript the official SDKs handle account resolution and the base URL for you once you authenticate with the NUA key:

```
PYTHON
from nuclia import sdk

# Authenticate once; the SDK extracts the zone and account from the token.
sdk.NucliaAuth().nua(token=NUA_KEY)

kbs = sdk.NucliaAccount()
kb = kbs.create_kb(account=ACCOUNT, zone=ZONE, slug="quickstart",
title="Quickstart KB")
print(kb.id)
```

Gotcha — tested Two limits surfaced immediately in testing. First, the NUA key is bound to one specific account server-side — passing any other account id returns [This NUA key is issued for a different account](#). Second, the Global API cannot be used to discover that account from a NUA key (it rejects NUA keys outright), so the account id must come from your dashboard or the SDK's stored configuration.

Step 2 — Get a Knowledge Box API key

The NUA key created the Knowledge Box, but **content and search calls need a Knowledge Box-scoped credential**. Create a service account on the Knowledge Box and then mint a key for it — both through the Zone API:

```
BASH
# Create a service account with owner rights on the KB
curl -X POST
"https://$ZONE.dp.progress.cloud/api/v1/account/$ACCOUNT/kb/$KB/service_accounts" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY" -H "Content-Type: application/json" \
-d '{"title":"app-backend","role":"SOWNER"}'
# -> { "id": "32af7a29-7692-4806-a973-ec926da2c175" }
```



```
# Mint an API key for that service account (expiry <= 1095 days)
curl -X POST
"https://$ZONE.dp.progress.cloud/api/v1/account/$ACCOUNT/kb/$KB/service_account/
$SA_ID/keys" \
  -H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY" -H "Content-Type: application/json" \
  -d '{"expires": "2028-01-01T00:00:00Z"}'
# -> { "token": "eyJhbGciOi...KB-scoped-key..." }
```

Gotcha — tested The key expiry is capped: requesting an expiry beyond three years returns **Chosen key expiration exceeds the current maximum of 1095 days** (the platform's spelling). Pick a date within 1095 days.

Save the returned **token** as **KEY**. Content calls present it with the standard **Authorization: Bearer** header.

Step 3 — Ingest content

Create a few text resources. Each resource gets a **slug**, a **title**, and one or more fields — here a single **text** field named **body**.

```
BASH
curl -X POST "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/resources" \
  -H "Authorization: Bearer $KEY" -H "Content-Type: application/json" \
  -d '{
    "slug": "battery-storage",
    "title": "Home Battery Storage",
    "texts": { "body": { "body": "Home battery systems store excess solar energy
for use at night. A typical LFP battery offers 10kWh usable capacity. The federal
Cheaper Home Batteries program subsidises installation from July 2025.",
"format": "PLAIN" } }
  }'
```

The platform accepts the resource immediately and returns its **uuid**:

```
JSON
{ "uuid": "86b5f7948b144b008f018cbd47a2bde4", "seqid": null }
```

The same operation in the two SDKs:

```
PYTHON
from nuclia import sdk
kb = sdk.NucliaKB(url=f"https://{ZONE}.dp.progress.cloud/api/v1/kb/{KB}",
api_key=KEY)
```

```
kb.resource.create(
  slug="battery-storage", title="Home Battery Storage",
  texts={"body": {"body": "Home battery systems store excess solar energy..."},
  "format": "PLAIN"}},
)
```

JAVASCRIPT

```
import { Nuclia } from "@nuclia/core";
const nuclia = new Nuclia({ backend: `https://${ZONE}.dp.progress.cloud/api`,
zone: ZONE,
                           knowledgeBox: KB, apiKey: KEY });
nuclia.db.getKnowledgeBox().subscribe(kb =>
  kb.createResource({ slug: "battery-storage", title: "Home Battery Storage",
    texts: { body: { body: "Home battery systems store excess solar energy...",
format: "PLAIN" } } })
    .subscribe(res => console.log(res.uuid)));
```

Warning Ingestion is asynchronous. The resource exists the moment you create it, but it is not searchable until background processing finishes. For small text this is seconds; for large files it can be longer. Do not run Step 4 in the same breath and expect results — poll the resource's status first (Chapter 7).

Step 4 — Retrieve with /find

Once processing completes, **find** performs hybrid (keyword + semantic) retrieval. Note the query is a natural-language phrase, not keywords — semantic matching does the rest.

BASH

```
curl -X POST "https://${ZONE}.dp.progress.cloud/api/v1/kb/$KB/find" \
-H "Authorization: Bearer $KEY" -H "Content-Type: application/json" \
-d '{"query":"how long is the payback period for solar","features":
["keyword","semantic"],"top_k":5}'
```

The response ranks matching resources and returns the specific paragraphs that matched, keyed by a **resource/field/char-range** identifier in **best_matches**:

JSON

```
{
  "resources": { "86b5f794...": { "slug": "battery-storage", "title": "Home
Battery Storage", ... } },
  "best_matches": [
    "ca2a8f3b.../t/body/0-301",
    "86b5f794.../t/body/0-336"
```

```
],
  "total": 4
}
```

Step 5 — Ask a question

`ask` runs the same retrieval, then sends the best passages to the generative model with your question. Request citations so the answer can be traced back to its sources. The `X-Synchronous: true` header waits for the full answer instead of streaming it.

BASH

```
curl -X POST "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/ask" \
  -H "Authorization: Bearer $KEY" -H "Content-Type: application/json" \
  -H "X-Synchronous: true" \
  -d '{"query":"How much does a home battery store and when is it
subsidised?","citations":true}'
```

The real answer the platform generated from the ingested corpus:

JSON

```
{
  "answer": "A home battery stores 10kWh of usable capacity and is subsidised
starting from July 2025 under the federal Cheaper Home Batteries program.",
  "status": "success",
  "citations": { "86b5f794.../t/body/0-336": [[0, 218]] },
  "retrieval_results": { "resources": { "86b5f794...": { ... } } }
}
```

Notice three things. The answer is correct and drawn entirely from the ingested text. The `citations` object maps a source paragraph to the character span of the answer it supports — auditable grounding. And `retrieval_results` echoes exactly which passages were used, so nothing about how the answer was produced is hidden.

Key idea That is the whole platform in five calls: create a Knowledge Box, get a key, ingest, find, ask. Everything else in this book makes each of these steps richer — better ingestion, sharper retrieval, more controllable answers — but the shape never changes.

What to read next

- Confused about which key does what? **Chapter 4 — Authentication.**

- Want to configure models and security on the Knowledge Box? **Chapter 5 — Management.**
- Ready to ingest real documents and files? **Part 3 — Data Ingestion.**
- Want to master **find** and **ask**? **Part 4 — Search, Retrieval, and RAG.**

PART 2 — AUTHENTICATION AND MANAGEMENT

Chapter 4

Authentication and Authorization

Which key opens which door — the credential model, tested end to end

More early failures come from authentication than from any other part of the platform, because there is not one credential but several, and each is accepted by some APIs and flatly rejected by others. This chapter maps the whole system. Every rejection quoted here is a real response captured during testing, so you can recognize them instantly when they happen to you.

The mental model: APIs are doors, credentials are keys

Recall the four APIs from Chapter 1 — Global, Zone, NucliaDB, and NUA. Think of each as a separate door. A credential is a key cut for specific doors. Presenting the right key at the wrong door does not degrade gracefully; it returns 403 with a message naming the mismatch.

Credential	Global	Zone (management)	NucliaDB (content)	NUA (processing)
User token / PAT	Yes	Yes	Yes	Yes
NUA key	No	Yes*	No	Yes
KB API key (service account)	No	No	Yes	Limited

*A NUA key can manage Knowledge Boxes only if its token was minted with the `allow_kb_management` claim, and only within the single account it is bound to.

Gotcha — tested Two rejections you will certainly meet. A NUA key against the Global API returns `Nuakeys are not valid in the global API`. A NUA key against Knowledge Box content returns `NuaKeyUser cannot access context of type NucliaDBKnowledgeBox`. Neither is a bug — they are the door telling you the key is cut for a different lock.

The four credentials in detail

User token and Personal Access Token (PAT)

A **user token** represents an interactively authenticated human. It grants everything that user can do across every API, but it is short-lived — about 30 minutes — and obtaining one requires a browser login flow. It is ideal for the dashboard and for interactive CLI sessions, and unsuitable for unattended scripts.

A **Personal Access Token (PAT)** is the long-lived equivalent: same broad access, but you set its expiry (90 days by default) and can revoke it at any time. Use a PAT for administrative automation — anything that must touch the Global API, manage accounts, or read activity logs.

BASH

```
# Create a PAT with the CLI (opens a browser once to authenticate you)
nuclia auth login
nuclia auth create_personal_token --description="ci-admin" --days=30
```

NUA key

A **NUA key** authenticates to the Nuclia Understanding API — the processing and Predict services. Its token is a JWT whose payload encodes the issuing zone, the account it belongs to, and its capabilities. The key used throughout this book carries `"allow_kb_management": true`, which extends it to Knowledge Box management (list, get, create, delete) — but only inside its own account.

NUA keys are presented in a dedicated header:

BASH

```
# NUA key uses X-NUCLIA-NUAKEY, not Authorization
curl "https://$ZONE.dp.progress.cloud/api/v1/predict/tokens?
text=Tesla+is+in+Austin" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Decoding the token payload shows exactly what a NUA key knows about itself:

JSON

```
{
  "iss": "https://aws-ap-southeast-2-1.dp.progress.cloud/",
  "sub": "9c6108f9-...",           // the NUA client id
  "key": "8c6c99ce-...",           // the key id
  "exp": 253370764800,
  "allow_kb_management": true       // extends the key to KB management
}
```

Gotcha — tested The account id is **not** in the token payload — only the client and key ids are. And you

cannot look the account up through the Global API with a NUA key, because that API rejects NUA keys. In practice the account id comes from the dashboard or from the SDK's stored configuration after you run `nuclia auth nua`. During testing, the SDK's config file was the reliable source of the bound account id.

Knowledge Box API key (service account)

Content and search on a specific Knowledge Box are authorized by a **Knowledge Box API key**. You do not create these directly; you create a **service account** on the Knowledge Box, give it a role, and then mint one or more **keys** for it. This indirection is deliberate: it lets you rotate keys, scope them, and revoke a single leaked key without disturbing the others.

The key is a bearer token used with the standard header, and it is the credential your application backend should carry:

BASH

```
curl -X POST "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/find" \
-H "Authorization: Bearer $KB_API_KEY" \
-H "Content-Type: application/json" \
-d '{"query":"..."}'
```

The full bootstrap, as tested

Here is the exact sequence used to stand up this book's sandbox — the canonical way to go from a management credential to a working content key.

1. **Create the Knowledge Box** with the NUA key on the Zone API: `POST /api/v1/account/{account}/kbs` -> returns the `kbid`.
2. **Create a service account** on that KB: `POST /api/v1/account/{account}/kb/{kbid}/service_accounts` with `{"title": "...", "role": "SOWNER"}` -> returns the service account `id`.
3. **Mint an API key** for it: `POST /.../service_account/{sa_id}/keys` with an `expires` date within 1095 days -> returns the `token`.
4. **Use that token** with `Authorization: Bearer` for all content and search calls on the KB.

Gotcha — tested The key-expiry ceiling is exactly 1095 days (three years). An `expires` further out returns `Chosen key expiration exceeds the current maxium of 1095 days` — note the platform's own spelling of the message, which is how you will see it.

Service account roles

A service account's role determines what its keys can do on the Knowledge Box. Grant the narrowest role that works — a search-only frontend should never hold a writer key.

Role	Grants
<code>SREADER</code>	Read and search only — resources, find, ask, catalog
<code>SWRITER</code>	Reader plus create/modify/delete resources and fields
<code>SCONTRIBUTOR</code>	Write content but not manage KB-level configuration
<code>SOWNER</code>	Full control of the Knowledge Box, including configuration and keys

Tip Issue two keys per Knowledge Box in most applications: an `SREADER` key embedded in or proxied for the frontend, and an `SWRITER/SOWNER` key held only by your ingestion backend. If the frontend key leaks, no one can alter your data with it.

Security controls beyond the key

Authorization does not stop at the token. A Knowledge Box can be locked down further:

- **Allowed origins** — restrict which web origins may call the KB, so an embedded widget key can only be used from your own domains.
- **Allowed IP addresses** — restrict calls to a set of source IPs for backend-only Knowledge Boxes.
- **Security groups on resources** — mark resources with `access_groups` so that a query only returns results the caller's groups are permitted to see (row-level security). Enable `enforce_security` on the KB to make this mandatory.

These are configured at Knowledge Box creation or via management (Chapter 5) and enforced at query time (Chapter 9). For multi-tenant applications, resource security groups are how you keep one tenant's content out of another tenant's answers within a shared Knowledge Box.

PART 2 — AUTHENTICATION AND MANAGEMENT**Chapter 5**

Accounts, Zones, and Knowledge Box Management

Creating, configuring, and retiring Knowledge Boxes

This chapter covers the administrative surface: how accounts and zones frame everything, how to create and configure Knowledge Boxes, how to choose models, and how to delete cleanly. These operations live on the **Zone API** (<https://<zone>.dp.progress.cloud>) and, for account-wide concerns, the **Global API**.

Accounts and zones

An **account** is your organization's boundary for billing, users, and Knowledge Boxes. Each account has a stable **id** (UUID) and a human-friendly **slug**. Most management URLs are scoped by account id: `/api/v1/account/{account_id}/...`

A **zone** is a region your data can live in. Knowledge Boxes are created within a specific zone and never move between zones. The zones available during testing included Australia (`aws-ap-southeast-2-1`), Europe (`aws-eu-central-1-1` and `europa-1`), the USA (`aws-us-east-2-1`), and Israel (`aws-il-central-1-1`). Choose based on data-residency requirements and proximity to your users.

Creating a Knowledge Box

Creation requires only **slug** and **title**, but several optional fields let you configure the Knowledge Box up front rather than patching it later.

Field	Type	Purpose
<code>slug</code>	string (required)	Human-readable id, unique within the account
<code>title</code>	string (required)	Display name
<code>description</code>	string	Free-text description
<code>mode</code>	enum	Knowledge Box mode; immutable

		after creation
<code>learning_configuration</code>	object	Model choices — semantic, generative, NER, etc.
<code>allowed_origins</code>	array<string>	Web origins permitted to call the KB
<code>allowed_ip_addresses</code>	array<string>	Source IPs permitted to call the KB
<code>external_index_provider</code>	object	Use an external vector index instead of the built-in one
<code>enforce_security</code>	boolean	Require resource security groups on every request

BASH

```
curl -X POST "https://$ZONE.dp.progress.cloud/api/v1/account/$ACCOUNT/kbs" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY" -H "Content-Type: application/json" \
-d '{
  "slug": "support-kb",
  "title": "Customer Support KB",
  "description": "Public help-centre content",
  "learning_configuration": { "generative_model": "chatgpt-azure-4o" }
}'
# -> 201 { "id": "e3bb0d28-3821-44a4-8948-fe511f48aa42" }
```

Inspecting and changing configuration

Two different views of a Knowledge Box matter. The **management view** (Zone API) shows account context, security, and top-level settings. The **learning configuration** (NucliaDB API, [GET /kb/{kbid}/configuration](#)) shows the models actually in force. Here is the real learning configuration of the sandbox as returned by the platform:

JSON

```
{
  "semantic_model": "multilingual-2024-05-06",
  "semantic_vector_size": 1024,
  "semantic_vector_similarity": "DOT",
  "semantic_threshold": 0.4,
  "generative_model": "chatgpt-azure-4o",
  "ner_model": "multilingual",
  "relation_model": "base",
  "anonymization_model": "disabled",
  "visual_labeling": "disabled",
  "summary_model": "chatgpt-azure-4o",
  "summary": "simple"
}
```

Change configuration with a **PATCH**. Model changes take effect for subsequent operations; changing the semantic model is where vectorsets (Chapter 2) matter, because existing vectors were produced by the old model.

BASH

```
# Change the generative model (Zone API, management)
curl -X PATCH "https://$ZONE.dp.progress.cloud/api/v1/account/$ACCOUNT/kb/$KB" \
  -H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY" -H "Content-Type: application/json" \
  -d '{"learning_configuration":{"generative_model":"chatgpt-azure-4o"}}'

# Or with the CLI against the learning configuration directly
nuclia kb set_configuration --generative_model=chatgpt-azure-4o
```

Choosing models

Two model choices dominate answer quality: the **semantic model** (which embedding model indexes and matches meaning) and the **generative model** (which LLM writes answers).

Semantic models

The platform offers English-optimized and multilingual embedding models. The sandbox defaulted to **multilingual-2024-05-06** (1024-dimensional, dot-product similarity). Pick a multilingual model if your content or your users span languages; pick an English model for best quality on English-only corpora. You can query the available options from the Predict service.

Generative models

The generative model is any LLM the platform supports — during authoring the default was **chatgpt-azure-4o**, and the platform's model lineup also includes Anthropic Claude models, Llama, and Azure Mistral, among others. You can also **bring your own** provider account (Anthropic, AWS Bedrock, Google Gemini) so that generation runs under your own contract and keys — configured through the learning configuration's **user_keys**.

Tip Set the generative model per Knowledge Box, but override it per request when needed: **/ask** accepts a **generative_model** parameter, so you can route hard questions to a stronger model and routine ones to a cheaper model without reconfiguring the KB.

Listing, and deleting cleanly

List every Knowledge Box in an account with `GET /api/v1/account/{account}/kbs`. Deletion is permanent and removes all content and indexes, so it is the one management call to guard carefully.

```
BASH

# List
curl "https://$ZONE.dp.progress.cloud/api/v1/account/$ACCOUNT/kbs" \
  -H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"

# Delete (irreversible)
curl -X DELETE "https://$ZONE.dp.progress.cloud/api/v1/account/$ACCOUNT/kb/$KB" \
  -H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Warning Deletion cannot be undone, and it takes the KB's service accounts and API keys with it. For anything you might need to restore, take a backup first (Chapter 17) — enterprise accounts can create backup/restore Knowledge Boxes.

Management operations at a glance

The Zone API exposes the full administrative surface for a Knowledge Box. The most useful:

Operation	Endpoint
Create KB	<code>POST /api/v1/account/{account}/kbs</code>
Get / modify / delete KB	<code>GET PATCH DELETE /api/v1/account/{account}/kb/{kb}</code>
List KBs	<code>GET /api/v1/account/{account}/kbs</code>
Service accounts	<code>GET POST /.../kb/{kb}/service_accounts</code>
API keys	<code>POST /.../service_account/{sa}/keys,</code> <code>DELETE .../key/{id}</code>
Invite users to a KB	<code>POST /.../kb/{kb}/invite</code>
KB permissions / users	<code>GET /.../kb/{kb}/permissions, GET </code> <code>PATCH /.../kb/{kb}/users</code>
Activity logs	<code>POST /api/v1/kb/{kb}/activity/{event_type}</code> <code>/query</code> (needs user/PAT)

Gotcha — tested Activity-log endpoints reject NUA keys with `NuaKeyUser` cannot access context of type `ActivityLogsAPI`. Reading activity logs requires a user token or PAT, even though the same NUA key can create and delete the Knowledge Box itself. Chapter 16 covers observability with the right credential.

PART 3 — DATA INGESTION

Chapter 6

Resources and Fields: The Data Model

How content is shaped, addressed, and enriched before it is ever searched

Retrieval quality is decided at ingestion time. How you shape a resource — how you split content into fields, what metadata you attach, how you identify it — determines what you can filter by, how results are grouped, and how precisely answers can cite their sources. This chapter is the data model in practice.

The resource, revisited

A **resource** is one item of content plus its metadata. It has a system-assigned **uuid**, an optional **slug** you choose, and a set of **fields** that hold the actual content. A resource can be as simple as one uploaded PDF, or a composite of several fields — the platform's recipe example bundles an image, an explanatory text, a preparation video, and a comments conversation into a single resource, so a search for the dish returns one coherent result rather than four fragments.

Key idea Model one resource per *thing a user would want to retrieve as a unit*. If two pieces of content should always surface together, make them fields of one resource; if they should rank independently, make them separate resources.

Creating resources

Create a resource with **POST /kb/{kbid}/resources**. The body carries the resource's metadata and dictionaries of fields keyed by field id. Field ids must match `^[a-zA-Z0-9:_-]+$`.

BASH

```
curl -X POST "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/resources" \
-H "Authorization: Bearer $KEY" -H "Content-Type: application/json" \
-d '{
  "slug": "solar-basics",
  "title": "Solar Panel Basics",
  "texts": { "body": { "body": "A photovoltaic solar panel converts sunlight
into DC electricity...", "format": "PLAIN" } },
  "usermetadata": { "classifications": [ { "labelset": "topic", "label":
"energy" } ] }
}'
```

```
# -> 201 { "uuid": "ca2a8f3bf186491facde30cdb2fbc2b6", "seqid": null }
```

Note Content calls accept either `Authorization: Bearer <key>` or `X-NUCLIA-SERVICEACCOUNT: Bearer <key>`; both were verified to work with a KB API key. This book uses `Authorization` for brevity. The older `X-NUCLIA-SERVICEACCOUNT` form appears in some dashboard-generated snippets — they are equivalent.

The field types

Each field type is a key in the create/modify body and behaves differently at processing time.

Text fields (`texts`)

Text you supply directly. The `format` selects how it is interpreted: `PLAIN`, `HTML`, `MARKDOWN` (converted to plain text — use `KEEP_MARKDOWN` to preserve the markup), `RST`, `JSON`, or `JSONL`. Structured JSON is understood: index a JSON array of records and you can then ask questions whose answers depend on the structure.

JSON

```
"texts": {
  "body": { "body": "Plain prose here.", "format": "PLAIN" },
  "specs": { "body": "[{\"book\": \"Dune\", \"author\": \"Herbert\"}]", "format":
"JSON" }
}
```

File fields (`files`)

Uploaded binaries — PDFs, Office documents, images, audio, video, archives. Files can be added inline as base64, uploaded to a field endpoint, or supplied as a URL. The platform runs OCR on images and speech-to-text on audio and video automatically; the extracted text is what gets indexed. Chapter 7 is devoted to file upload.

Link fields (`links`)

A URL the platform fetches and processes. Ideal for indexing web pages and sitemaps without downloading them yourself.

JSON

```
"links": { "home": { "uri": "https://www.progress.com/agent-ic-rag" } }
```

Conversation fields (`conversations`)

An ordered list of messages — a chat log, a support thread, an email chain. Each message carries an author, text, and optional attachments. Conversations can be appended to over time with a dedicated messages endpoint, so a live thread stays a single resource.

Key-value fields (`key\_values`)

Structured key/value data attached to a resource, useful for records whose fields you want available alongside prose.

Metadata: the four buckets

Metadata is what turns a pile of text into a filterable, governable knowledge base. The platform separates it into distinct buckets with different rules.

Bucket	Field	Who sets it	Queryable?
User metadata	<code>usermetadata</code>	You	Yes — labels and relations drive filtering
Origin	<code>origin</code>	You	Yes — source, dates, tags, collaborators, path
System metadata	<code>metadata</code>	Platform (language, etc.)	Yes
Extra	<code>extra</code>	You	No — stored but not indexed for retrieval

User metadata holds `classifications` (labels) and `relations`. Labels are the primary mechanism for faceting and filtering — a `topic=energy` label on a resource lets a query restrict to energy content.

Origin captures provenance: `source_id`, `url`, `created/modified` timestamps, `tags`, `collaborators`, and a `path` (useful for reproducing a folder hierarchy). **Extra** is a free-form JSON store for information you want to keep with the resource but explicitly do *not* want influencing retrieval.

Gotcha — tested Put anything you might filter or facet by into `usermetadata/origin`, not `extra`. The `extra` bucket is deliberately excluded from the query path — data placed there is invisible to `find` filters. This trips people who store a category in `extra` and then wonder why they cannot filter on it.

Reading resources back

GET `/kb/{kbid}/resource/{rid}` returns a resource. The `show` and `extracted` query parameters control how much detail you get — from just the basics to the full extracted text and computed metadata the platform produced during processing.

<code>`show` value</code>	<code>Includes</code>
<code>basic</code>	Title, summary, icon, labels, status (the default)
<code>values</code>	The field values you supplied
<code>extracted</code>	Processed output: extracted text, paragraphs, computed metadata
<code>errors</code>	Any processing errors on the resource
<code>origin / security / relations</code>	The corresponding metadata blocks

BASH

```
curl "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/resource/$RID?
show=basic&show=extracted&extracted=text&extracted=metadata" \
-H "Authorization: Bearer $KEY"
```

Reading back the sandbox's `solar-basics` resource confirmed its status was `PROCESSED`, its detected `language` was `en`, and its text field carried both `text` and `metadata` under `extracted` — the paragraphs and computed metadata the indexes were built from.

Addressing resources by slug

Every by-id endpoint has a by-slug twin under `/kb/{kbid}/slug/{rslug}`. Because the slug is yours to choose, you can address resources by your own external identifiers without ever storing the platform's UUIDs.

TEXT

```
# These two are equivalent when the resource's slug is "solar-basics"
GET /kb/{kbid}/resource/ca2a8f3bf186491facde30cdb2fbc2b6
GET /kb/{kbid}/slug/solar-basics
```

Modifying, hiding, and deleting

PATCH a resource to change metadata or add fields; unset fields are left untouched, so patches are additive by default. Resources can be **hidden** (excluded from search without deletion) by patching `hidden: true` — useful for staging content or soft-retiring it. **DELETE** removes a resource and its fields permanently.

BASH

```
# Add a label and hide the resource in one patch
curl -X PATCH "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/resource/$RID" \
  -H "Authorization: Bearer $KEY" -H "Content-Type: application/json" \
  -d '{"hidden": true, "usermetadata": {"classifications":
[{"labelset": "status", "label": "archived"}]}'
```

PART 3 — DATA INGESTION

Chapter 7

Uploading and Processing Files

Binaries in, extracted knowledge out — upload paths, processing, and strategies

Files are where RAG earns its keep: a folder of PDFs, scanned images, and recorded calls becomes searchable knowledge. This chapter covers the two upload paths, what processing does to a file, how to control chunking and extraction, and how to know when a file is ready.

What the platform accepts

The platform ingests a wide range of formats and extracts text from all of them, applying OCR and speech-to-text automatically:

Category	Examples
Documents	.txt, .html, .docx, .pdf, .json
Spreadsheets	.xlsx, .csv
Presentations	.pptx
Images (OCR)	.jpg, .png, .tiff
Video (speech-to-text)	.mp4, .avi, .mpeg
Audio (speech-to-text)	.mp3, .wav
Web	.html, sitemaps
Archives	.zip, .gzip, .rar

Gotcha — tested An archive is indexed as **one** resource — the text of every file inside is extracted but merged together. If you need per-file results, unpack the archive before ingesting. This is a deliberate design choice, not a limitation you can toggle.

Two ways to upload

Simple upload

For files up to a moderate size, `POST /kb/{kbid}/upload` sends the whole binary in one request. Set `X-Filename` (and optionally `X-Language`, `X-Password` for protected documents, `X-Extract-Strategy`, `X-Split-Strategy`). This was the path verified in testing.

```
BASH
curl -X POST "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/upload" \
  -H "Authorization: Bearer $KEY" \
  -H "X-Filename: wind.txt" \
  -H "Content-Type: text/plain" \
  --data-binary @wind.txt
# -> 201 { "seqid": null, "uuid": "1009f3c7...", "field_id": "1d7c1a94..." }
```

The response returns the new resource `uuid` and the `field_id` of the file field it created — so a single call both makes the resource and attaches the file.

Resumable (TUS) upload

For large files, unreliable networks, or resumable transfers, the platform implements the **TUS** protocol. You create an upload, then send the file in chunks (`PATCH` with an `Upload-Offset`), and can resume from the last acknowledged offset after an interruption. Use TUS for anything large enough that re-sending from scratch would hurt — long videos, big PDFs. The SDKs wrap TUS so you rarely hand-roll the offset bookkeeping.

What processing does

Once a file lands, the Nuclia Understanding API processes it through a pipeline: detect type, extract text (OCR/STT as needed), detect language, split into paragraphs, run NER, and compute vectors with the KB's semantic model. Each step's output is stored under the field's `extracted` data, which you can read back (Chapter 6).

Warning Processing is asynchronous and not instant. In testing, a freshly uploaded file was accepted with 201 but `/ask` still answered *"Not enough data to answer this"* until processing completed, because its vectors did not yet exist. Never assume a file is queryable the moment upload returns.

Knowing when a file is ready

Read the resource and check its processing status. `GET /kb/{kbid}/resource/{rid}?show=basic` returns `metadata.status`, which moves from `PENDING/PROCESSING` to `PROCESSED` (or `ERROR`). Poll it, or — better — have processing notify you.

PYTHON

```
import time, requests
def wait_processed(zone, kb, rid, key, timeout=300):
    url = f"https://{zone}.dp.progress.cloud/api/v1/kb/{kb}/resource/{rid}"
    hdr = {"Authorization": f"Bearer {key}"}
    deadline = time.time() + timeout
    while time.time() < deadline:
        status = requests.get(url, headers=hdr, params={"show": "basic"}).json()
        ["metadata"]["status"]
        if status in ("PROCESSED", "ERROR"):
            return status
        time.sleep(3)
    raise TimeoutError("still processing")
```

Tip For production ingestion at volume, do not poll. Configure a **webhook** (the augmentation agents in Chapter 8 can trigger one) or watch the processing status endpoint in batches. Polling one resource at a time does not scale.

Controlling extraction and chunking

Two kinds of strategy let you override the defaults for how a file is turned into paragraphs.

Extract strategies

An **extract strategy** controls how text is pulled from a document — for example, how tables or layout are handled. Strategies are defined on the Knowledge Box and referenced by name at upload time via the `X-Extract-Strategy` header (or the `extract_strategy` field on a text/link field). Listing them on a fresh KB returns an empty set — none are defined until you create them.

BASH

```
# List extract strategies (empty {} on a new KB)
curl "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/extract_strategies" \
  -H "Authorization: Bearer $KEY"
```

Split strategies

A **split strategy** controls how extracted text is divided into paragraphs — the chunking step. The default splitting works well for prose, but a custom split strategy helps when your documents have unusual structure (dense tables, code, transcripts). Like extract strategies, they are defined on the KB and referenced by name with `X-Split-Strategy`.

Key idea Chunking is the highest-leverage ingestion tuning knob. If answers are missing context or citing fragments, revisit your split strategy before touching the models. Paragraphs that are too small lose context; too large, and retrieval precision drops.

Reprocessing and reindexing

Two maintenance operations handle change. **Reprocess** re-runs the full understanding pipeline on a resource — use it after changing an extract/split strategy or the semantic model. **Reindex** rebuilds the indexes from already-extracted data — cheaper, for when only indexing needs refreshing.

BASH

```
# Reprocess (full pipeline) vs reindex (indexes only)
curl -X POST ".../kb/$KB/resource/$RID/reprocess" -H "Authorization: Bearer $KEY"
curl -X POST ".../kb/$KB/resource/$RID/reindex" -H "Authorization: Bearer $KEY"
```

PART 3 — DATA INGESTION

Chapter 8

Data Augmentation: Ingestion Agents

Automatically enriching content as it lands — labels, structure, graphs, and safety

Ingestion does not have to stop at extraction. **Data augmentation agents** (also called ingestion agents) run over your resources and create new derived data: labels, summaries, structured JSON, question/answer pairs, knowledge-graph relations, and safety flags. They can run once over existing content, continuously over new content, or only over resources matching criteria — and they can fire a webhook when done.

The agent catalog

Querying the platform's task catalog on a live Knowledge Box (`GET /kb/{kbid}/tasks`) returned the following agents. Each comes with its own configuration schema.

Agent	What it produces
labeler	Applies labels to resources or text blocks from label descriptions and examples
llm-graph	Extracts a knowledge graph — entities and relations — over resources
synthetic-questions	Generates question/answer pairs (a resource-specific FAQ)
ask	Generates summaries, JSON, or extracted structured content
llm-align	Aligns/validates generated content against source
memory	Analyses conversation fields and extracts memory facts incrementally
llama-guard	Flags text blocks or resources as unsafe / inappropriate
prompt-guard	Flags jailbreak / prompt-injection content

Note This catalog was read live from the platform during authoring. The dashboard groups them under friendlier names — Labeler, Generator, Q&A Generator, Graph Extraction, Content Safety & LLM Security — but the API task names above are what you configure against.

Labelers — automatic classification

A labeler applies a labelset to content based on a short description of each label. The canonical example: a support Knowledge Box with a **Severity** labelset — **Low**, **Moderate**, **High** — each with a precise description. The labeler reads each incoming ticket and assigns the right severity, which you can then filter and facet on at query time (Chapter 9). Because the classification happens at ingestion, search stays fast — no per-query LLM classification.

Tip Label descriptions are prompts. Vague descriptions produce vague labels. Write each one as if instructing a new analyst: state what qualifies, what does not, and give an example. This single practice does more for label quality than any model choice.

Generators — summaries, JSON, and Q&A

Generator-style agents (the **ask** and **synthetic-questions** tasks) create new text from existing resources. Use them to pre-compute per-resource summaries, extract structured JSON (for example pulling parties, dates, and amounts out of a contract into fields you can filter on), or build a synthetic FAQ so common questions have crisp, pre-written answers to retrieve against.

Graph extraction — building the knowledge graph

The **llm-graph** agent reads resources and extracts named entities and the relations between them, populating the knowledge graph that Chapter 12 shows how to query. On the sandbox, graph extraction over an ingested web page surfaced entities such as *Progress*, *Private Knowledge Box*, and *Enterprise Knowledge Management*, connected by typed relations — all queryable through **/graph**.

Content safety and LLM security

Two guard agents protect the pipeline. **llama-guard** flags unsafe or inappropriate content; **prompt-guard** flags jailbreak and prompt-injection attempts. Running these as ingestion agents means unsafe content is labelled before it can ever be retrieved into an answer, and you can exclude flagged resources at query time. For any application exposed to user-generated or untrusted content, enable both.

Triggering: one-shot, continuous, and conditional

Every agent can run in three modes, and the mode is what makes agents operationally useful:

- **One-shot** — apply to all existing resources now (backfilling a new labelset, for instance).
- **Continuous** — apply automatically to every new resource as it is ingested.
- **Conditional** — apply only to resources matching criteria: resource type, field type, or keywords. The platform also supports triggering agents based on labels, so one agent's output can gate another's — for example, only run the summarizer on resources the labeler tagged **High**.

An agent can also call a **webhook** on completion, which is the clean way to notify downstream systems that a resource is enriched and ready — far better than polling.

Managing agents through the API

Agents are tasks on the Knowledge Box. List the catalog and inspect each agent's configuration schema with `GET /kb/{kbid}/tasks`; the NUA API exposes the endpoints to start, stop, and inspect agent runs. The configuration schema returned for each task tells you exactly which fields it accepts — including which LLM to use and, for the guards, which provider keys are required.

BASH

```
# Read the full agent catalog and each agent's config schema
curl "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/tasks" \
-H "Authorization: Bearer $KEY"
```

PART 4 — SEARCH, RETRIEVAL, AND RAG

Chapter 9

Retrieval Fundamentals: /find and /search

Finding the right passages — hybrid retrieval, filtering, ranking, and reranking

Generation is only as good as the passages it is given. This chapter is about getting that step right: the two retrieval endpoints, how hybrid retrieval works, how to filter precisely, and how ranking and reranking decide the final order. Master `find` and `ask` (Chapter 10) becomes mostly a matter of what to do with the passages it returns.

Two endpoints, two shapes of result

The platform offers two retrieval endpoints. Both accept `GET` (query in the URL) and `POST` (query in a JSON body); use `POST` for anything non-trivial.

Endpoint	Returns	Use when
<code>/find</code>	Results grouped by resource, with the best-matching paragraphs and a <code>best_matches</code> ranking	Building RAG context or a search UI — this is the modern default
<code>/search</code>	Flat lists — separate <code>paragraphs</code> , <code>resources</code> , <code>sentences</code> blocks with scores	You need the raw, ungrouped result sets or legacy-style output

`find` is what `ask` calls internally, and its resource-grouped shape maps naturally onto search interfaces, so this chapter leads with it.

Hybrid retrieval: keyword + semantic

The `features` parameter selects which indexes to search. The default, `["keyword", "semantic"]`, runs both — *hybrid* retrieval — and fuses the results.

Feature	Index	Good at
---------	-------	---------

keyword	BM25 lexical	Exact terms, names, codes, acronyms
semantic	Vector	Meaning, paraphrase, synonyms, cross-lingual matches
relations	Knowledge graph	Entities related to the query
graph	Graph paths	Structured graph traversal (see Chapter 12)

A live `find` for *"how long is the payback period for solar"* — a phrase that does not appear verbatim in the corpus — correctly surfaced the solar and battery resources via semantic matching, returning paragraph-level `best_matches` like `ca2a8f3b.../t/body/0-301`. That identifier is the retrieval currency of the platform: `resource-id / field-type / field-id / char-range`.

BASH

```
curl -X POST "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/find" \
  -H "Authorization: Bearer $KEY" -H "Content-Type: application/json" \
  -d '{"query":"how long is the payback period for solar","features":
["keyword","semantic"],"top_k":5}'
```

JSON

```
{
  "resources": { "86b5f794...": { "slug": "battery-storage", "title": "Home
Battery Storage", ... } },
  "best_matches": [ "ca2a8f3b.../t/body/0-301", "86b5f794.../t/body/0-336",
"ca2a8f3b.../a/title/0-18"],
  "total": 4
}
```

The complete parameter set, by job

`find` accepts 33 parameters. Rather than an alphabetical dump (that is Appendix A's job), here they are grouped by what they *do* — the same three-question framework the cookbook in Chapter 13 builds on: what is eligible, how much comes back, and how it is serialized.

Query and eligibility — what can be retrieved

Parameter	Type / default	What it controls
query	string	The natural-language query
features	array — [keyword, semantic]	Which indexes to search: keyword, semantic, relations, graph

<code>fields</code>	array — <code>[]</code>	Restrict the search to specific fields, e.g. <code>a/title</code>
<code>filter_expression</code>	object	Composable AND/OR/NOT filter over labels, slugs, dates, origin (see below)
<code>security</code>	object	The caller's <code>groups</code> , for row-level (per-resource) security
<code>vectorset</code>	string	Which vector index to search, on a multi-model KB
<code>with_synonyms</code>	bool — <code>false</code>	Expand the query using the KB's custom synonyms
<code>show_hidden</code>	bool — <code>false</code>	Include resources marked hidden
<code>range_creation_start/end</code>	datetime	Only resources created within this window
<code>range_modification_start/end</code>	datetime	Only resources modified within this window
<code>resource_filters</code>	array	Restrict to specific resource ids (legacy; prefer <code>filter_expression</code>)
<code>filters / keyword_filters</code>	array	Legacy flat filters, superseded by <code>filter_expression</code>

Context and ranking — how much comes back, in what order

Parameter	Type / default	What it controls
<code>top_k</code>	integer — <code>20</code>	How many results to retrieve (max 200)
<code>min_score</code>	number or object	Drop results below this score; settable per index (bm25/vector)
<code>rank_fusion</code>	<code>rrf</code> (default)	How keyword and semantic result lists are merged
<code>reranker</code>	<code>predict</code> (default)	Second-pass model reordering of the fused results
<code>vector</code>	array<number>	Supply your own query vector instead of embedding <code>query</code>
<code>with_duplicates</code>	bool — <code>false</code>	Return duplicate paragraphs instead of collapsing them

Query understanding and output

Parameter	Type / default	What it controls
<code>rephrase</code>	bool — <code>false</code>	Rewrite the query before retrieval for better matching
<code>rephrase_prompt</code>	string	Custom instructions for the rephrasing step
<code>query_image</code>	object	An image to search with (visual models)
<code>highlight</code>	bool — <code>false</code>	Wrap matched terms in the returned text
<code>show</code>	array — <code>[basic]</code>	Which resource metadata blocks to serialize
<code>field_type_filter</code>	array	Which field types appear in results
<code>search_configuration</code>	string	Load these parameters from a stored named configuration
<code>generative_model</code>	string	(when relations/graph features drive follow-on generation)
<code>debug / audit_metadata</code>	bool / object	Debugging output and audit tagging

Key idea You will not set most of these on any given call — the defaults are strong. This table is a map, not a checklist. Chapter 13 shows which handful to reach for in each real use case.

Filtering with `filter_expression`

Filtering is how you turn a whole-KB search into a scoped one. The modern mechanism is `filter_expression`, a composable object that supersedes the older flat `filters`, `fields`, `range_*`, `resource_filters`, and `keyword_filters` parameters. It lets you build AND/OR/NOT expressions over resource and field properties — labels, slugs, origin metadata, creation and modification dates, and more.

A verified example restricting a search to a single resource by slug:

```
BASH
curl -X POST "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/find" \
```

```
-H "Authorization: Bearer $KEY" -H "Content-Type: application/json" \
-d '{
  "query": "energy",
  "filter_expression": { "field": { "prop": "resource", "slug": "battery-
storage" } },
  "top_k": 5
}'
```

Filtering by label is the workhorse pattern — combine it with the labels an ingestion agent applied (Chapter 8) to scope answers to a category, tenant, or status:

```
JSON
{
  "query": "outages",
  "filter_expression": {
    "field": { "prop": "label", "labelset": "severity", "label": "High" }
  }
}
```

Tip Prefer `filter_expression` over the legacy flat parameters for anything new. It is strictly more expressive — you can nest AND/OR/NOT — and the platform documents the older parameters as things it *replaces*. Mixing both in one request invites confusion.

Ranking and reranking

Two stages decide the final order. First, **rank fusion** merges the keyword and semantic result lists — the default and currently only strategy is Reciprocal Rank Fusion (**rrf**), which rewards documents that rank well in *both* lists. Second, a **reranker** takes the fused top results and reorders them with a more expensive, more accurate model; the default is **predict**, the platform's own reranking model.

```
BASH
# Explicit RRF fusion + Predict reranking (both verified against the live KB)
curl -X POST "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/find" \
-H "Authorization: Bearer $KEY" -H "Content-Type: application/json" \
-d '{"query":"battery","rank_fusion":"rrf","reranker":"predict","top_k":5}'
```

Key idea The defaults — hybrid features, **rrf** fusion, **predict** reranking — are strong and the right starting point. Tune **top_k** and **min_score** first; only revisit fusion and reranking once you have evidence the default ordering is failing you.

Scoping, security, and hidden content

Three parameters shape *what is eligible* to be returned:

- ``security`` — pass the caller's `groups`, and the platform returns only resources whose `access_groups` permit them. This is row-level security enforced at query time; with `enforce_security` on the KB it is mandatory.
- ``show_hidden`` — by default, resources marked `hidden` are excluded; set this to include them (for admin or staging views).
- ``with_synonyms`` — expand the query using the KB's custom synonyms, so `pv` also matches `photovoltaic` and `solar` (the sandbox had exactly those synonyms configured).

Autocomplete and suggestions

For a search-as-you-type experience, `/suggest` returns fast partial-match paragraphs and entities for a prefix — much cheaper than a full `find` on every keystroke. A live `/suggest?query=batt` returned the EV and battery paragraphs immediately. Chapter 11 covers it alongside catalog and summarize.

PART 4 — SEARCH, RETRIEVAL, AND RAG

Chapter 10

Generative Answers: The /ask Endpoint

From retrieved passages to a trusted, cited, controllable answer

`/ask` is the platform's centerpiece. It performs retrieval (everything in Chapter 9), assembles the best passages into a context, sends that context plus your question to a generative model, and returns the answer — optionally with citations, structured output, and a full account of what it retrieved. It accepts every `find` parameter *plus* a large family of generation controls. This chapter walks the ones that matter.

The simplest ask, and its anatomy

At minimum, `ask` needs a `query`. Requesting `citations` and sending `X-Synchronous: true` (wait for the whole answer rather than streaming) gives the cleanest first result:

BASH

```
curl -X POST "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/ask" \
  -H "Authorization: Bearer $KEY" -H "Content-Type: application/json" \
  -H "X-Synchronous: true" \
  -d '{"query":"How much does a home battery store and when is it
subsidised?","citations":true}'
```

JSON

```
{
  "answer": "A home battery stores 10kWh of usable capacity and is subsidised
starting from July 2025 under the federal Cheaper Home Batteries program.",
  "status": "success",
  "citations": { "86b5f794.../t/body/0-336": [[0, 218]] },
  "retrieval_results": { "resources": { "86b5f794...": { ... } } }
}
```

Three parts of the response deserve attention every time. `answer` is the generated text. `citations` maps each supporting source paragraph to the character spans of the answer it grounds — this is how you show users *why* an answer is trustworthy. `retrieval_results` is the exact `find` output that fed the model, so you can inspect or display the sources.

Gotcha — tested When retrieval finds nothing relevant, the platform does not invent an answer — it returns `status: "no_context"` and the text *"Not enough data to answer this."* This is the anti-hallucination guardrail doing its job. Treat `no_context` as a first-class outcome in your code, not an

error.

The generation parameters, grouped

Beyond the retrieval parameters it inherits, `ask` adds these generation controls. They are grouped here by what they do.

Controlling the answer text

Parameter	Type / default	Effect
<code>prompt</code>	string or object	Your own prompt template, with <code>{context}</code> and <code>{question}</code> placeholders
<code>generative_model</code>	string	Override the KB's LLM for this request
<code>max_tokens</code>	int or object	Cap the answer length (and optionally context tokens)
<code>prefer_markdown</code>	bool — <code>false</code>	Ask the model to format the answer as Markdown
<code>generative_model_seed</code>	int	Seed for more reproducible generations
<code>generate_answer</code>	bool — <code>true</code>	If false, retrieve and assemble context but skip generation
<code>reasoning</code>	bool or object — <code>false</code>	Enable model reasoning for harder questions

A custom `prompt` is the highest-leverage control. A verified request that constrains tone and length:

```
JAVASCRIPT
{
  "query": "payback period?",
  "prompt": "Answer in one short sentence for a homeowner. Context: {context}
Question: {question}"
}
// -> "The payback period for a residential solar system in Australia is
typically 4 to 7 years, depending on feed-in tariffs."
```

Controlling grounding and citations

Parameter	Effect
<code>citations</code>	Return source-paragraph → answer-span mapping
<code>citation_threshold</code>	Minimum confidence for a citation to be included
<code>answer_json_schema</code>	Force the answer into a JSON structure you define
<code>extra_context</code>	Inject extra passages not in the KB into the context
<code>extra_context_images</code>	Inject images into the context (requires a visual LLM)

Conversational context

Parameter	Effect
<code>chat_history</code>	Prior turns, so the query is rephrased with conversation awareness
<code>chat_history_relevance_threshold</code>	How strongly history influences rephrasing
<code>rephrase</code> / <code>rephrase_prompt</code>	Rewrite the user's query before retrieval

Structured answers with `answer_json_schema`

For programmatic use, you often want data, not prose. Supply an `answer_json_schema` and the model returns a validated object instead of free text. A live example extracting a single number:

BASH

```
curl -X POST "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/ask" \
  -H "Authorization: Bearer $KEY" -H "X-Synchronous: true" -H "Content-Type:
application/json" \
  -d '{
    "query": "list the battery capacity in kWh",
    "answer_json_schema": {
      "name": "battery", "description": "battery facts",
      "parameters": { "type": "object", "properties": { "capacity_kwh": { "type":
"number" } },
      "required": ["capacity_kwh"] }
    }
  }'
```

JSON

```
{ "answer": "", "answer_json": { "capacity_kwh": 10 }, "status": "success", ... }
```

Key idea With `answer_json_schema`, the prose `answer` comes back empty and the structured result lands in `answer_json`. This is the clean way to use `/ask` as a data-extraction engine — parse `answer_json`, ignore `answer`.

RAG strategies: shaping the context

By default, `ask` sends the top matching *paragraphs* to the model. **RAG strategies** change what context is assembled — often the difference between a fragmentary answer and a complete one. Supply one or more in `rag_strategies`. The platform offers:

Strategy	What it sends to the model
<code>full_resource</code>	The entire text of the top resources, not just matched paragraphs
<code>field_extension</code>	Specific additional fields of the matched resources
<code>hierarchy</code>	The matched paragraph plus its surrounding document hierarchy
<code>neighbouring_paragraphs</code>	The matched paragraph plus the ones immediately around it
<code>metadata_extension</code>	Resource metadata added to the context
<code>conversational</code>	Surrounding messages for conversation fields
<code>prequeries</code>	Run predefined queries first and add their results to context
<code>graph</code>	Expand context along knowledge-graph relations

A verified `full_resource` request — useful when answers need the whole document, not a snippet:

```
JSON
{ "query": "summarise EV charging options",
  "rag_strategies": [ { "name": "full_resource", "count": 1 } ] }
```

Tip `neighbouring_paragraphs` is the cheapest quality win. RAG answers often fail because the matched paragraph is correct but incomplete — the sentence before or after held the missing detail. Adding a neighbour or two fixes this without the cost of `full_resource`.

Streaming the answer

Omit `X-Synchronous` and `ask` streams its response as newline-delimited JSON items — ideal for a chat UI that shows the answer as it is written. Capturing a live stream revealed the full item sequence:

<code>`item.type`</code>	Payload
<code>answer</code>	One token/chunk of the answer text (many of these)
<code>retrieval</code>	The retrieval results that grounded the answer
<code>citations</code>	The source → span citation map
<code>augmented_context</code>	Any extra context that was assembled
<code>status</code>	Final status, e.g. <code>{code:0, status:success}</code>
<code>metadata</code>	Token counts and timings — input/output tokens, first-chunk latency

PYTHON

```
import requests, json
r = requests.post(f"https://{ZONE}.dp.progress.cloud/api/v1/kb/{KB}/ask",
                  headers={"Authorization": f"Bearer {KEY}"},
                  json={"query": "What is the battery subsidy?", "citations":
True}, stream=True)
for line in r.iter_lines():
    if not line: continue
    item = json.loads(line)["item"]
    if item["type"] == "answer":
        print(item["text"], end="", flush=True)    # render tokens as they arrive
    elif item["type"] == "citations":
        citations = item["citations"]              # show sources when they arrive
```

The `metadata` item is quietly valuable in production: it reports input and output token counts (both raw and in Nuclia's billing units) and timings such as time-to-first-chunk. Log it to track cost and latency per answer.

Asking against a single resource

`POST /kb/{kbid}/resource/{rid}/ask` (and its by-slug twin) restricts generation to one resource — perfect for a "chat with this document" feature. It takes the same generation parameters, but retrieval is confined to that resource's fields.

Reusing configurations

Rather than repeating a long parameter set on every call, define a **search configuration** in the dashboard and reference it by name with `search_configuration`. The stored configuration supplies the defaults; anything you pass in the request overrides it. This keeps application code small and lets non-developers tune retrieval and prompts without a deploy.

Key idea You now know every `find` and `ask` parameter. The next chapter — **Chapter 13, Configuring RAG for Your Use Case** — is the payoff: a cookbook of tested recipes that map real use cases (precise Q&A, high-recall research, structured extraction, chatbots, multi-tenant, low-cost) to concrete parameter settings you can copy and adapt.

PART 4 — SEARCH, RETRIEVAL, AND RAG

Chapter 11

Catalog, Suggest, and Summarize

The supporting retrieval endpoints that round out a search experience

Three endpoints sit alongside `find` and `ask` and handle jobs those two are not built for: browsing the catalog of resources, autocompleting as the user types, and summarizing content on demand. Each is small, fast, and worth knowing.

Catalog — browse and facet the resource index

`/catalog` lists and filters resources by their *stored* attributes — title, slug, labels, origin — without touching the semantic index. It is the endpoint behind an admin table or a faceted browse experience, and it is available the instant a resource is created, before processing finishes.

BASH

```
# List everything
curl "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/catalog" -H "Authorization: Bearer $KEY"

# Facet by a labelset to get counts per label
curl -X POST "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/catalog" \
  -H "Authorization: Bearer $KEY" -H "Content-Type: application/json" \
  -d '{"query":"","faceted":["/classification.labels/topic"]}'
```

Key idea `catalog` reads the title/metadata index, not the vector index, so it works before a resource is processed and never incurs embedding cost. Use it for browse, filter, and faceted-count UIs; use `find` when you need relevance.

Suggest — autocomplete as you type

`/suggest` returns fast partial matches for a prefix — paragraphs and entities — so a search box can show suggestions on each keystroke. A live `?query=batt` returned the battery and EV paragraphs in milliseconds. Keep it on `GET` for simplicity and cache aggressively.

BASH

```
curl "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/suggest?query=batt" \
-H "Authorization: Bearer $KEY"
```

Summarize — condense resources on demand

`/summarize` takes a list of resource ids and returns a summary of each, plus an optional combined summary. It uses the KB's summary model. A live call over the solar and battery resources returned tidy per-resource summaries. Use it for digest views, previews, or pre-computing summaries you will later retrieve against.

BASH

```
curl -X POST "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/summarize" \
-H "Authorization: Bearer $KEY" -H "Content-Type: application/json" \
-d '{"resources":["ca2a8f3b...", "86b5f794..."]}'
```

JSON

```
{ "resources": { "ca2a8f3b...": { "summary": "Photovoltaic solar panels convert
sunlight into DC electricity, which is then converted to AC..." } } }
```

Tip For summaries you will show repeatedly, generate them once with a Generator ingestion agent (Chapter 8) and store them as a field, rather than calling `/summarize` on every page view. Summarize on demand only for ad-hoc selections.

PART 4 — SEARCH, RETRIEVAL, AND RAG

Chapter 12

The Knowledge Graph and Graph Search

Retrieval that follows relationships, not just similarity

Vector and keyword search find passages that *resemble* a query. Some questions instead need passages connected by *relationships* — who reports to whom, which component depends on which, how two entities are linked. For those, the platform builds a **knowledge graph** of entities and relations and lets you query it directly.

Where the graph comes from

Two sources populate the graph. The NER model detects entities (people, organizations, places, dates) during processing automatically. The `llm-graph` ingestion agent (Chapter 8) goes further, using an LLM to extract richer, typed relations between entities across your resources. On the sandbox, graph extraction over an ingested web page produced entities such as *Progress*, *Private Knowledge Box*, and *Enterprise Knowledge Management*, connected by typed relations.

Three graph endpoints

Endpoint	Returns
POST /kb/{kbid}/graph	Paths — source → relation → destination triples
POST /kb/{kbid}/graph/nodes	Nodes — the entities themselves, ranked
POST /kb/{kbid}/graph/relations	Relations — the edge types in the graph

Gotcha — tested All three require a `query` object in the body. Sending only `{"top_k": 10}` returns 422 `Field required` for `query`. The query object selects what part of the graph to traverse — for example `{"prop": "node"}` for nodes or `{"prop": "path"}` for paths.

```
BASH
```



```
# Retrieve nodes (entities) from the graph – verified live
curl -X POST "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/graph/nodes" \
  -H "Authorization: Bearer $KEY" -H "Content-Type: application/json" \
  -d '{"query":{"prop":"node"},"top_k":5}'
```

JSON

```
{
  "nodes": [
    { "value": "Private Knowledge Box", "type": "entity", "group": "ORG",
      "score": 1.037 },
    { "value": "Progress", "type": "entity", "group": "ORG", "score": 1.037 }
  ]
}
```

A path query returns triples. A live [/graph](#) call returned paths linking a resource to a topic via an [ABOUT](#) relation — the shape you traverse to answer relationship questions:

JSON

```
{
  "paths": [
    { "source": { "value": "ca2a8f3b...", "type": "resource" },
      "relation": { "type": "ABOUT", "label": "" },
      "destination": { "value": "topic/energy", "type": "entity" } }
  ]
}
```

Graph in retrieval: features and rag_strategies

You do not always query the graph directly. Two integration points bring graph knowledge into ordinary retrieval and generation:

- **features: ["relations"] or ["graph"] on find/ask** — include graph-related entities and paths in retrieval results.
- **The graph RAG strategy on ask** (Chapter 10) — expand the generation context along graph relations, so the model sees not just the matched paragraph but the entities and facts connected to it.

Key idea Reach for the graph when the *connections* are the answer. "What depends on service X?" or "Which policies mention both A and B?" are graph questions; "What is the refund window?" is a vector question. Many real applications use both, with the graph strategy enriching an otherwise standard [ask](#).

Defining your own entities

The platform auto-detects standard entity groups — `ORG`, `GPE` (places), `DATE`, and others were present on the sandbox by default. For domain-specific entities (product SKUs, internal system names), define your own **entity groups** so the graph and NER recognise them. Read the current groups with `GET /kb/{kbid}/entitiesgroups`; manage domain entity groups from the dashboard or the management API, then reprocess so existing content is re-scanned.

PART 4 — SEARCH, RETRIEVAL, AND RAG

Chapter 13

Configuring RAG for Your Use Case

A cookbook — turning the /find and /ask parameters into concrete recipes

Chapters 9 and 10 introduced every parameter `find` and `ask` accept. This chapter answers the question that actually matters when you build something: *for my use case, how do I set them?* It gives a framework for reasoning about the parameters, a decision table, and a set of copy-and-adapt recipes — each one a complete request, tested against the live sandbox, with the reasoning for every setting.

Key idea There is no single 'best' RAG configuration. A support bot that must never guess and a research tool that must never miss want almost opposite settings. Configuration is where you encode what *your* application values — precision, recall, completeness, cost, or latency.

A framework: three questions

Every RAG request answers three questions in order. Group the parameters by which question they answer and configuration stops being a wall of forty fields.

Question	You are controlling...	Key parameters
1. What is <i>eligible</i> to be retrieved?	The candidate set before ranking	<code>features</code> , <code>filter_expression</code> , <code>security</code> , <code>fields</code> , <code>vectorset</code> , <code>show_hidden</code> , range filters
2. How much <i>context</i> reaches the model?	What gets assembled and sent	<code>top_k</code> , <code>min_score</code> , <code>rank_fusion</code> , <code>reranker</code> , <code>rag_strategies</code>
3. How should it <i>answer</i> ?	The generation itself (ask only)	<code>prompt</code> , <code>generative_model</code> , <code>max_tokens</code> , <code>citations</code> , <code>answer_json_schema</code> , <code>reasoning</code> , <code>prefer_markdown</code> , <code>generate_answer</code>

Work top to bottom. If answers are wrong, the fault is almost always in question 1 or 2 — the model can only be as good as the context it is handed. Tune eligibility and context before you touch the prompt.

The decision table

A fast lookup from what you care about to the parameters that move the needle.

If you need...	Set
Maximum accuracy, no guessing	Low <code>top_k</code> (5–8) · hybrid <code>features</code> · <code>reranker: predict</code> · a strict <code>prompt</code> · <code>citations: true</code>
Maximum recall (miss nothing)	High <code>top_k</code> (30–50) · add <code>relations</code> to <code>features</code> · no <code>min_score</code> · consider <code>full_resource</code>
Complete, non-fragmented answers	<code>rag_strategies: neighbouring_paragraphs</code> or <code>hierarchy</code>
Whole-document reasoning	<code>rag_strategies: full_resource</code> (low <code>top_k</code>)
Structured output for a pipeline	<code>answer_json_schema</code> (read <code>answer_json</code> , ignore <code>answer</code>)
A multi-turn chatbot	<code>chat_history</code> · <code>rephrase: true</code> · <code>citations: true</code>
Exact-match lookup (codes, SKUs)	<code>features: [keyword]</code> · target <code>fields</code> · synonyms off
Lowest latency / cost	Small <code>top_k</code> · cap <code>max_tokens</code> · cheaper <code>generative_model</code> · stream
Bring-your-own LLM	<code>generate_answer: false</code> — take the retrieved context and generate elsewhere
Tenant isolation in a shared KB	<code>security.groups</code> + <code>filter_expression</code> on a tenant label

Recipe 1 — Precise factual Q&A (support, compliance, policy)

The goal is a trustworthy answer or an honest 'I don't know' — never a plausible guess. Keep the candidate set small and high-quality, rerank it, constrain the model with a strict prompt, and demand citations so every answer is auditable.

JSON

```
POST /kb/{kbid}/ask
{
  "query": "How much does a home battery store and when is it subsidised?",
  "features": ["keyword", "semantic"],
  "top_k": 6,
  "reranker": "predict",
  "citations": true,
  "max_tokens": 300,
  "prompt": "Answer ONLY from the context. If the context does not contain the
answer, say you don't know. Be concise. Context: {context} Question: {question}"
}
```

Tested against the sandbox, this returned: *"A home battery stores 10kWh of usable capacity and is subsidised starting from July 2025 under the federal Cheaper Home Batteries program,"* with a citation pointing at the exact source paragraph. **Why each setting:** low `top_k` keeps noise out of the context; `reranker: predict` puts the single best passage first; the `prompt` is the guardrail against invention; `citations` makes the answer defensible.

Gotcha — tested Even without a strict prompt, the platform will not invent an answer when retrieval finds nothing — it returns `status: "no_context"`. The strict prompt makes that behaviour tighter and extends it to *partially* relevant context. Always handle `no_context` in your UI.

Recipe 2 — High-recall research (discovery, literature review, e-discovery)

Here the cost of missing a relevant passage is high and a human reads the results, so favour recall over precision. Cast a wide net, add graph relations, and drop `min_score` so nothing is filtered out prematurely.

JSON

```
POST /kb/{kbid}/find
{
  "query": "risks associated with battery storage",
  "features": ["keyword", "semantic", "relations"],
  "top_k": 40,
  "rank_fusion": "rrf",
  "highlight": true,
  "show": ["basic", "values"]
}
```

Use `find` (not `ask`) for research surfaces — you want the ranked evidence, not a single synthesised answer. `highlight` marks the matched terms for a reviewer; `relations` pulls in graph-connected entities the pure vector search would rank lower.

Recipe 3 — Whole-document answers (contracts, standards, long policies)

When the answer depends on a whole document rather than one paragraph — 'summarise this policy', 'what are the termination clauses' — send the full resource, not snippets.

```
JSON
POST /kb/{kbid}/ask
{
  "query": "Summarise the EV charging options and their power ratings.",
  "top_k": 2,
  "rag_strategies": [ { "name": "full_resource", "count": 1 } ]
}
```

Tested, this produced a complete summary covering both the 7kW single-phase and 22kW three-phase options — detail a paragraph-only retrieval would have split across fragments. Keep `top_k` low with `full_resource`; whole documents are large, and two of them can fill the model's context window on their own.

Recipe 4 — Complete, non-fragmented snippets (the cheapest quality win)

The most common RAG failure is a correct-but-incomplete answer: the matched paragraph was right, but the detail lived in the sentence before or after it. `neighbouring_paragraphs` fixes this for a fraction of the cost of `full_resource`.

```
JSON
POST /kb/{kbid}/ask
{
  "query": "what capacity is a home battery",
  "rag_strategies": [ { "name": "neighbouring_paragraphs", "before": 1, "after":
1 } ]
}
```

Reach for this first whenever answers feel truncated. In testing it returned a clean *"A typical home battery has a usable capacity of 10kWh"* by giving the model the matched paragraph plus its immediate neighbours for context.

Recipe 5 — Structured extraction (feed a pipeline, not a person)

When the consumer is code, not a human, don't parse prose — demand JSON. Supply an `answer_json_schema` and the model returns a validated object.

```
JSON
POST /kb/{kbid}/ask
{
  "query": "list the battery capacity in kWh",
  "answer_json_schema": {
    "name": "battery",
    "description": "battery facts",
    "parameters": {
      "type": "object",
      "properties": { "capacity_kwh": { "type": "number" } },
      "required": ["capacity_kwh"]
    }
  }
}
```

Tested, this returned `{"answer_json": {"capacity_kwh": 10}}` with an empty prose `answer`. Parse `answer_json`; ignore `answer`. This is how you turn `/ask` into a reliable field-extraction engine over documents — parties and dates from contracts, specs from datasheets, entities from reports.

Recipe 6 — Conversational assistant (multi-turn chatbot)

A chatbot must understand follow-ups like 'and for batteries?' that only make sense given the previous turn. Pass `chat_history` so the platform rephrases the query with conversational awareness before retrieving.

```
JSON
POST /kb/{kbid}/ask
{
  "query": "and for batteries?",
  "rephrase": true,
  "citations": true,
  "chat_history": [
    { "author": "USER", "text": "What is the payback period for solar?" },
    { "author": "NUCLIA", "text": "About 4 to 7 years." }
  ]
}
```

Gotcha — tested History changes retrieval. In testing, the follow-up 'and for batteries?' with this history rephrased to a battery query — but if the rephrased query drifts from your corpus you can get

`no_context` where a direct question would have answered. Tune `chat_history_relevance_threshold` if history is pulling retrieval off-topic.

Recipe 7 — Exact-match lookup (SKUs, error codes, case numbers)

Semantic search is the wrong tool for an exact token like `ERR-4021` or a part number — you want lexical precision, not 'similar meaning'. Turn semantics off and target the field.

```
JSON
POST /kb/{kbid}/find
{
  "query": "LFP",
  "features": ["keyword"],
  "fields": ["a/title", "t/body"],
  "top_k": 10
}
```

`features: [keyword]` uses only the BM25 index; `fields` restricts the search to named fields so a code in the body doesn't get lost among unrelated semantic neighbours. This is also the cheapest query type — no embedding of the query is required.

Recipe 8 — Multilingual knowledge base

With a multilingual semantic model (the sandbox defaulted to `multilingual-2024-05-06`), vector search matches *meaning across languages*: a French question can retrieve an English passage. Lead with semantic features and let rephrasing normalise the query.

```
JSON
POST /kb/{kbid}/ask
{
  "query": "Quel est le délai de rentabilité du solaire ?",
  "features": ["semantic", "keyword"],
  "rephrase": true
}
```

Keep `keyword` in the mix for names and codes (which are language-agnostic), but semantic does the cross-lingual work. If your users and content span languages, confirm the KB's semantic model is a multilingual one (Chapter 5).

Recipe 9 — Lowest latency and cost

For high-volume or interactive surfaces, every token and millisecond counts. Shrink the context, cap the output, pick a cheaper model, and stream so the user sees words immediately.

```
JSON
POST /kb/{kbid}/ask          (send WITHOUT X-Synchronous to stream)
{
  "query": "what is the battery subsidy?",
  "top_k": 3,
  "max_tokens": 150,
  "generative_model": "<a-cheaper-model>",
  "prefer_markdown": true
}
```

Streaming's `metadata` item reports input/output tokens and time-to-first-chunk per answer (Chapter 10) — log it and watch the trend. `prefer_markdown` was verified to return formatted output (headings, bold, lists) when you are rendering into a rich UI.

Tip Route by difficulty: set a cheap `generative_model` as the KB default and override `generative_model` per request only for the hard questions. Most traffic is easy; pay for the strong model only when it earns its keep.

Recipe 10 — Retrieval-only: bring your own LLM

Sometimes you want the platform's world-class retrieval but your own model for generation — a fine-tuned model, an on-prem LLM, or one under your own contract. Set `generate_answer: false` and `ask` returns the assembled context without generating.

```
JSON
POST /kb/{kbid}/ask
{
  "query": "battery subsidy",
  "generate_answer": false
}
```

Tested, this returned `status: "success"` with an empty `answer` and the full `retrieval_results`. Take those passages, build your own prompt, and call whatever model you like. You get the platform's hybrid retrieval, reranking, and RAG strategies as a service, decoupled from generation.

Recipe 11 — Secure, multi-tenant retrieval

In a shared Knowledge Box, every query must see only its tenant's data. Combine a `filter_expression` on a tenant label with the caller's `security.groups`, and enable `enforce_security` on the KB so it can never be forgotten (Chapters 4 and 18).

JSON

```
POST /kb/{kbid}/find
{
  "query": "outstanding invoices",
  "security": { "groups": ["tenant:acme"] },
  "filter_expression": { "field": { "prop": "label", "labelset": "tenant",
    "label": "acme" } },
  "top_k": 10
}
```

Warning Test tenant isolation adversarially before shipping — a missing `security` block or a wrong label here is a cross-tenant data leak, not a cosmetic bug. Prefer a Knowledge Box per tenant when isolation matters more than overhead (Chapter 18).

Putting it together

Real applications rarely use one recipe; they compose. A production support assistant might use Recipe 1's strict prompt and citations, Recipe 4's neighbouring paragraphs for completeness, Recipe 6's chat history for follow-ups, and Recipe 11's security scoping — all in one request. Start from the recipe closest to your goal, then borrow settings from the others.

Key idea Store the configuration you settle on as a named **search configuration** (Chapter 10) and call it by name with `search_configuration`. Your application code stays a one-liner, and you can retune retrieval and prompts — swap a recipe — without shipping a new build.

PART 5 — THE RETRIEVAL AGENT

Chapter 14

Agentic Retrieval, Sessions, and Workflows

Beyond fixed pipelines — multi-source, multi-step, reasoning retrieval

Everything so far has used a Knowledge Box's *fixed* RAG pipeline: one query, one retrieval, one answer. That is the right tool for most questions. But complex questions — ones that need several sources, intermediate reasoning, or a decision about *where* to look — call for something more. This is the **Retrieval Agent** (also called the Retrieval-Augmented Orchestrator, or RAO): a configurable agent that plans, routes, and iterates. This is the "Agentic" in Progress Agentic RAG.

Why an agent, not just a pipeline

A Knowledge Box answers "what is the refund window?" perfectly: search, retrieve, generate. It struggles with "compare our refund policy to the one in this SQL database of competitor terms, and if neither covers digital goods, check the web." That question needs multiple sources, a sub-question breakdown, conditional routing, and iteration until the information is sufficient. The Retrieval Agent is built for exactly this.

A Retrieval Agent can:

- Draw on **several sources at once** — multiple Knowledge Boxes, SQL and Snowflake databases, pandas dataframes, internet search services, and MCP servers.
- **Analyse and split** a user question into sub-questions and decide dynamically which sources each needs.
- Apply **conditional logic** to route from one step to the next — including fallbacks.
- **Reason and iterate** autonomously until it has enough to answer.

The two building blocks: drivers and a workflow

A Retrieval Agent is defined by two things.

Drivers — the sources

A **driver** makes one information source available to the agent. Its name and description matter, because the agent uses them to decide whether a source is relevant to the question in hand. Supported drivers include:

Driver	Source it exposes
NucliaDB	One of your Knowledge Boxes (with a chosen search configuration)
Perplexity / Google	Internet search services
SQL	A relational database, queried in natural language (beta)
Snowflake	A Snowflake database, with dialect-aware SQL (beta)
MCP	Tools exposed by Model Context Protocol servers

Tip Write driver descriptions as carefully as you would write a job description. The agent routes queries to sources based on these descriptions; a vague one means the agent sends the wrong questions to the wrong source. For a NucliaDB driver in the same region, the platform will auto-create an API key for you — you only set its expiry.

Workflow — how questions are handled

The **workflow** is the plan. It has four stages, each hosting specific agents:

Stage	Purpose	Available agents
Preprocess	Improve the question before retrieval	Rephrase
Retrieval context	Gather the information to answer	Ask, Perplexity/Google, MCP, SQL, Snowflake, Pandas, Smart Agent
Generation	Produce the answer from what was retrieved	Summarize
Postprocess	Assess quality or trigger actions	Validation (REMi), External Call

Preprocess → **Rephrase** rewrites the user's question for better retrieval, optionally consulting a Knowledge Box for context and following rules you specify. **Retrieval context** is where the work happens: an *Ask* agent queries a Knowledge Box driver; a *Perplexity* agent searches the web; *SQL/Snowflake/Pandas* agents query structured data in natural language. **Generation** → **Summarize** turns the gathered material into an answer. **Postprocess** can run a *Validation* agent that uses the REMi model to score answer quality, or an *External Call* agent that makes an HTTP request to trigger a downstream action.

The Smart Agent: planning and reactive modes

The most powerful retrieval agent is the **Smart Agent**, which orchestrates other agents. Instead of a fixed route, it picks the appropriate sources, splits the question into sub-questions, evaluates relevancy, and iterates until the information is sufficient. It runs in one of two modes:

- **Planning mode** — generate a full plan up front, then execute it.
- **Reactive mode** — select and execute functions iteratively, deciding each next step from the last result.

You register other agents (Knowledge Boxes, MCP tools, databases) with the Smart Agent as *registered agents*, each with a description of what it does and the function it exposes. Only registered agents that expose functions can be orchestrated. You can use different LLMs for the planning and execution phases to balance cost and quality, and enable *feedback support* so the agent can pause to ask the user to validate an intermediate result or supply a missing detail.

Key idea The Smart Agent is where 'agentic' becomes literal: it decides, acts, observes, and decides again. Start with a simple fixed workflow (Ask → Summarize) and graduate to a Smart Agent only when questions genuinely need multi-source planning — the fixed workflow is cheaper, faster, and easier to reason about.

A worked pattern: internet fallback

A common, high-value workflow answers from your Knowledge Box when it can and falls back to the web when it cannot — so users always get an answer, clearly sourced. The wiring:

1. Create a **NucliaDB driver** (your Knowledge Box) and a **Perplexity driver** (web search).
2. In **Retrieval context**, add an **Ask** agent using the Knowledge Box as its source.
3. On that Ask agent's **Fallback** output, add a **Perplexity** agent — it runs only when the Knowledge Box has no relevant match.
4. In **Generation**, add a **Summarize** agent to phrase the final answer from whichever source produced results.

This mirrors, at the agent level, the `no_context` behavior of `/ask` from Chapter 10 — but instead of returning "not enough data," the agent reaches for another source.

Talking to an agent: sessions and the WebSocket API

Retrieval Agents are conversational and stateful, so they are driven over a **WebSocket** rather than a single request/response. This lets the agent stream intermediate steps, stream the answer, and — crucially — pause to ask the user for more information (a missing parameter, or to trigger an OAuth flow for a source).

Opening a connection

Open the socket with an **ephemeral token** obtained from the agent's `/ephemeral_token` endpoint (short-lived, safe to hand to a browser). The URL:

TEXT

```
wss://<zone>.dp.progress.cloud/api/v1/agent/<agent-id>/session/ephemeral/ws?eph-token=<token>&workflow_id=<workflow-id>
```

`workflow_id` is optional; omit it to run the agent's default workflow. Send the user's question as the first message with `operation: 0`; send `operation: 1` to end the chat.

JSON

```
{
  "question": "Compare our refund window to the competitor terms.",
  "operation": 0,
  "chat_history": [ { "question": "...", "answer": "..." } ]
}
```

The agent streams back messages: intermediate step results, citations, answer chunks, and occasionally an agent-to-user request for input. Handle that last case — respond on the same socket — and the agent resumes. The JavaScript SDK's `RetrievalAgent` class wraps this protocol so you work with observables instead of raw frames.

Gotcha — tested Ephemeral tokens are exactly that — ephemeral. Mint one per session from your backend (where the real credential lives) and pass only the ephemeral token to the browser. Never embed a long-lived key in client code that opens the socket directly.

Deploying a Retrieval Agent

Three deployment paths, in increasing order of control:

- **Ready-to-use widget** — create a widget in the dashboard, customize it, and paste the snippet into your site. Fastest path to a working chat UI (Chapter 15).
- **Custom frontend** — implement the WebSocket protocol yourself, or use the JavaScript SDK's `RetrievalAgent` class for full control of the experience.
- **MCP** — expose the agent to a broader AI system (for example Claude or another MCP client) so the agent becomes a tool other assistants can call. The platform exposes MCP handler endpoints for exactly this.

Whichever you choose, the agent's configuration — drivers and workflow — stays server-side and unchanged; deployment only decides how users reach it.

PART 6 — BUILDING AND OPERATING**Chapter 15**

The SDKs: Python, JavaScript/TypeScript, and .NET

Working with the platform in your language of choice

Every capability in this book is reachable over plain HTTP, and the cURL samples map directly onto that. But for real applications the official SDKs remove boilerplate — authentication, base-URL resolution, retries, streaming, and typed models. The SDKs are wrappers over the same API, so anything you have learned about endpoints and parameters carries straight across. As the platform notes, the API always has the newest features first; the SDKs follow closely.

Python SDK (nuclia)

The Python SDK is the most complete, and it ships with a CLI. It is the right choice for ingestion pipelines, data scripts, backends, and — as this book demonstrated — testing.

BASH

```
pip install nuclia
```

Authenticate once; the SDK reads the zone and account out of the token. It supports NUA keys, Knowledge Box keys, user login, and personal access tokens.

PYTHON

```
from nuclia import sdk

# NUA key (management + processing). The SDK extracts zone/account from the
token.
sdk.NucliaAuth().nua(token=NUA_KEY)

# Or work against one Knowledge Box with a KB API key
kb = sdk.NucliaKB(url=f"https://{ZONE}.dp.progress.cloud/api/v1/kb/{KB}",
api_key=KB_KEY)
kb.resource.create(slug="doc-1", title="Doc 1",
                    texts={"body": {"body": "Hello world", "format": "PLAIN"}})

search = sdk.NucliaSearch(url=..., api_key=KB_KEY)
answer = search.ask(query="What does the document say?")
print(answer.answer)
```

The Python SDK is organized into modules that mirror the API surface — a clean map of what is available:

Module	Covers
<code>auth</code>	NUA, user, and personal-access-token authentication
<code>kb / manage</code>	Knowledge Box configuration and labels
<code>upload</code>	Resources and file upload (including resumable)
<code>search</code>	find, search, ask, catalog, summarize
<code>nua</code>	Predict services — embeddings, NER, LLM, query analysis
<code>agents (ai/da)</code>	Retrieval Agents and data-augmentation agents
<code>kb-backup / import-export</code>	Backups and data portability
<code>activity-log</code>	Reading activity logs
<code>rao</code>	Retrieval-Augmented Orchestrator (Retrieval Agent) sessions

Tip The CLI shares the Python SDK's engine, so `nuclia kb ...`, `nuclia auth ...`, and friends are perfect for one-off administration and for scripting in CI without writing Python. `nuclia auth nua REGION NUA_KEY` stores the key for reuse.

JavaScript / TypeScript SDK (@nuclia/core)

The JS/TS SDK targets the frontend and Node backends and is observable-first (RxJS). It is the engine behind the search widget, and the way to build a custom search or agent UI. It is large — the `RetrievalAgent` class alone exposes well over a hundred members — but the entry point is small.

BASH

```
npm install @nuclia/core
```

JAVASCRIPT

```
import { Nuclia } from "@nuclia/core";

const nuclia = new Nuclia({
  backend: `https://${ZONE}.dp.progress.cloud/api`,
  zone: ZONE,
  knowledgeBox: KB,
```



```

    apiKey: KB_KEY,
  });

  // Ask, streaming results as they arrive
  nuclia.knowledgeBox
    .ask("How much does a home battery store?")
    .subscribe((answer) => console.log(answer.text));

```

Key classes you will meet: **Nuclia** (the client), **Db** (account/KB access), **KnowledgeBox** and **WritableKnowledgeBox** (read vs. read-write operations), **Resource/ReadableResource** (a single resource and its fields), and **RetrievalAgent** with **Session/ReadableSession** for agent conversations over WebSocket. The read/write split is deliberate — hand a **KnowledgeBox** (read-only) to frontend code and reserve **WritableKnowledgeBox** for trusted backends.

.NET SDK

The .NET SDK exposes the platform through typed service interfaces, a natural fit for C# backends. The services mirror the API's shape:

Interface	Responsibility	Representative methods
IKnowledgeBoxService	KB-level operations	configuration, labels
IResourceService	Resource CRUD	create, get, modify, delete
IResourceFieldsService	Fields on a resource	add/get/download fields
ISearchService	Retrieval and generation	AskAsync , AskStreamAsync , FindAsync , CatalogAsync , SuggestAsync , SummarizeAsync , GraphSearchAsync

CSHARP

```

// C# – ask a question and read the answer
var answer = await searchService.AskAsync(new AskRequest {
    Query = "How much does a home battery store?",
    Citations = true
});
Console.WriteLine(answer.Answer);

// Stream tokens as they arrive
await foreach (var item in searchService.AskStreamAsync(new AskRequest { Query =
    "..."}))
    Console.Write(item.Text);

```

Which SDK, when

Situation	Reach for
Ingestion pipelines, data/admin scripts, CI	Python SDK or CLI
Frontend search or chat UI	JavaScript/TypeScript SDK (or the widget)
Node backend	JavaScript/TypeScript SDK
C#/.NET backend	.NET SDK
A language with no SDK	The REST API directly — every cURL sample in this book

Key idea The SDKs are conveniences, not gatekeepers. When an SDK lags a new API feature, drop to the REST call for that one operation and keep using the SDK for the rest. Because they share the same endpoints, mixing them is safe.

PART 6 — BUILDING AND OPERATING

Chapter 16

No-Code: The Search Widget and Dashboard

A production search experience without writing a backend

Not every solution needs custom code. The platform ships a dashboard for managing everything through a GUI and an embeddable **widget** that puts a full search-and-answer experience on a web page with a two-line snippet. For many internal tools and public help centres, this is the entire frontend.

The dashboard

The dashboard at rag.progress.cloud covers the whole lifecycle without code: sign up and pick a zone, upload data (files in PDF/Word/Excel/PowerPoint/text, or links to web pages), watch processing, then use the search page to try queries and generative answers. Its most useful trick for developers is the **Get code** button: configure a search visually, then copy the exact API call, Python, or widget snippet it produces. Treat the dashboard as an interactive query builder, not just an admin console.

The embeddable widget

The widget is a set of Web Components loaded from the platform's CDN. Two custom elements — a search bar and a results area — plus your Knowledge Box id and zone give a working experience:

HTML

```
<script src="https://cdn.rag.progress.cloud/nuclia-widget.umd.js"></script>

<nuclia-search-bar
  knowledgebox="YOUR-KB"
  zone="aws-ap-southeast-2-1"
  features="answers,filter">
</nuclia-search-bar>
<nuclia-search-results></nuclia-search-results>
```

The **features** attribute toggles behavior — **answers** enables generative answers, **filter** enables faceted filtering, and so on. Build and preview the snippet in the dashboard's **Widgets** section, then click **Embed widget** to generate it.

The three widget types

Type	Experience
Embedded in page	Search input on the page; results appear below; 'Ask more' opens full chat
Popup	A search button that opens the experience in an overlay
Chat	A conversational chat interface, suited to a Retrieval Agent

Using the widget with a private Knowledge Box

A public Knowledge Box needs no key in the widget. A **private** Knowledge Box does — and you must not paste a long-lived key into a public page. The supported patterns are to restrict the key's **allowed origins** so it only works from your domains, and to serve short-lived credentials from your backend for the widget to use.

Warning Restricting allowed origins is a genuine control, but it is not a secret store. For truly sensitive corpora, keep the Knowledge Box private, proxy search through your backend (which holds an **SREADER** key), and have the widget talk to your proxy — not directly to the platform.

Customizing appearance and behavior

The widget exposes attributes and a JavaScript API for styling, prompts, result rendering, and event hooks, so it can match your brand and wire into your analytics. When you outgrow it, the JavaScript SDK (Chapter 14) picks up exactly where the widget leaves off — the widget is built on it — so migrating from widget to custom UI is an incremental step, not a rewrite.

PART 6 — BUILDING AND OPERATING

Chapter 17

Quality and Observability

Measuring whether your RAG is actually good, and watching it in production

A RAG system can return fluent answers that are subtly wrong, and you will not know unless you measure. This chapter covers the three tools the platform gives you to keep quality honest: REMi evaluation metrics, user feedback capture, and activity logs.

REMi: evaluating the RAG triad

REMi (RAG Evaluation Metrics) is a model built by the platform to score the quality of a RAG interaction along the classic *RAG triad*:

Metric	Question it answers
Answer Relevance	Is the generated answer relevant to the user's query?
Context Relevance	Was the retrieved context relevant to the query?
Groundedness	Is the answer actually supported by the retrieved context?

These three localize failure. Low *context relevance* means retrieval is the problem — revisit chunking, filters, or the semantic model. High context relevance but low *groundedness* means the model is drifting from its evidence — tighten the prompt or lower `max_tokens`. Low *answer relevance* with good grounding means the question was misread — look at rephrasing. REMi is available through the NUA API and, per the platform, its logs are downloadable for quality monitoring.

Tip Sample, score, and track over time rather than scoring every call. Run REMi on a rolling sample of production interactions and watch the three metrics as a dashboard; a sudden drop in groundedness after a content or model change is exactly the signal you want to catch before users do.

BASH

```
# REMi is exposed via the NUA API; the Python SDK wraps it.
# It needs both the KB endpoint/key and a NUA key, since evaluation runs in NUA.
pip install nuclia
```

Capturing user feedback

The `/feedback` endpoint records a user's verdict on an answer — thumbs up/down and optional text — tied to the interaction. A live call succeeded with a simple body:

BASH

```
curl -X POST "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/feedback" \
  -H "Authorization: Bearer $KEY" -H "Content-Type: application/json" \
  -d '{"ident": "<answer-ident>", "good": true, "task": "CHAT", "feedback": "exactly right"}'
```

Feed this back into your evaluation loop: the interactions users flag as bad are the highest-value inputs to a REMi review and to tuning your prompts and chunking. Wire a feedback control into every answer surface — it is the cheapest quality signal you will ever get.

Activity logs

Activity logs are the audit trail of a Knowledge Box: resources processed, created, and edited; questions and generative answers; queries; and user feedback. They are how you monitor usage, troubleshoot, and satisfy compliance. Query them per event type:

BASH

```
curl -X POST
"https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/activity/visited/query" \
  -H "Authorization: Bearer $PAT" -H "Content-Type: application/json" \
  -d '{"year_month": "2026-07"}'
```

Gotcha — tested Activity-log endpoints reject NUA keys — a live call returned `NuaKeyUser cannot access context of type ActivityLogsAPI` (and the metrics variant, `ActivityLogsMetricsAPI`). Read activity logs with a **user token or PAT**, not the NUA key or the KB service-account key. This surprises people because the same NUA key can create the very Knowledge Box whose logs it cannot read.

The platform also surfaces activity logs in the dashboard UI, and — as noted in its own change log — REMi logs are downloadable, so quality and usage data can flow into your own analytics stack.

PART 6 — BUILDING AND OPERATING

Chapter 18

Backups, Export/Import, and Sync Connectors

Protecting data, moving it between Knowledge Boxes, and keeping it fresh

Content has a lifecycle beyond ingestion: it must be protected against mistakes, moved between environments, and kept in sync with the systems of record it came from. The platform provides three mechanisms for this — backups, export/import, and sync connectors.

Backup and restore

The **backup and restore** feature captures the full state of a Knowledge Box — its configuration *and* its resources — as a managed, point-in-time snapshot. Both creating and restoring a backup are asynchronous: you make a request, then poll for completion. Restoring produces a Knowledge Box in the captured state, which makes backups useful not only for disaster recovery but for cloning an environment.

Operation	Effect
Create backup	Snapshot KB configuration + resources (async)
List / get backup	Inspect available snapshots and their status
Restore backup	Recreate a Knowledge Box from a snapshot (async)
Delete backup	Remove a snapshot

Note Backup/restore is an enterprise-account capability. For smaller Knowledge Boxes, export/import (below) may be all you need; reach for backup/restore when you need a complete state capture or are operating at scale.

Export and import

Export/import serializes a Knowledge Box's data into a portable form you can download and re-import elsewhere. Use it to move content between accounts or zones, to seed a staging Knowledge Box from

production, or to keep an offline copy. For small-to-medium Knowledge Boxes it is simpler than backup/restore; for very large ones or full-state capture, prefer backups.

PYTHON

```
# Conceptual flow (Python SDK)
from nuclia import sdk
export = sdk.NucliaExportImport()
export.export_kb(...) # produces a portable export
export.import_kb(...) # re-creates content in a target KB
```

Keeping content fresh: the Sync Agent and cloud connectors

Most content lives in a system of record — a shared drive, a website, a cloud store — and changes there. The **Sync Agent** keeps a Knowledge Box current with such a source automatically.

The Sync Agent

The Sync Agent is a small application (Windows, macOS, Linux) that synchronizes a local folder, a sitemap, or an RSS feed into your Knowledge Box. It runs on your laptop or a server; you point the dashboard at its endpoint (for a local install, <http://localhost:8090>) to configure it. New and changed items in the source flow into the Knowledge Box without manual uploads.

Cloud storage sync

The platform also connects directly to cloud sources through **external connections** and **sync configurations** (visible in the Zone API as `external_connections` and `sync_config` endpoints). These let a Knowledge Box pull from cloud storage and SaaS sources on a schedule, with OAuth authorization handled through the platform, so content stays current without running the local Sync Agent.

Tip Choose by where the content lives and who controls the schedule. Local folders, sitemaps, and RSS -> the Sync Agent. Cloud storage and SaaS sources -> external connections and sync configurations, which run in the platform and need nothing installed.

PART 6 — BUILDING AND OPERATING

Chapter 19

Production Best Practices

Security, multi-tenancy, cost, and performance for real deployments

The gap between a working prototype and a production system is made of the concerns in this chapter: keeping data isolated and secure, controlling cost, staying fast, and failing safely. Everything here builds on mechanisms introduced earlier — this chapter is about applying them deliberately.

Security and access control

Layer your defenses; do not rely on any single one.

- **Least-privilege keys.** Give each component the narrowest service-account role that works — `SREADER` for search frontends, `SWRITER/SOWNER` only for trusted ingestion backends (Chapter 4). Never ship a writer key to a browser.
- **Rotate keys.** Because keys are minted per service account and expire within 1095 days, rotate them on a schedule and revoke individual keys the moment one is suspected leaked, without disturbing others.
- **Constrain the surface.** Set `allowed_origins` and, for backend-only KBs, `allowed_ip_addresses` so a key is useless outside your infrastructure.
- **Row-level security.** Tag resources with `access_groups` and pass the caller's `security` groups on every query; enable `enforce_security` to make it mandatory.
- **Guard the pipeline.** Run the `llama-guard` and `prompt-guard` ingestion agents (Chapter 8) on any untrusted content so unsafe or injection-laden material is flagged before it can reach an answer.

Multi-tenancy

Two models, chosen by how strong your isolation must be:

Model	How	Use when
Knowledge Box per tenant	Each tenant gets its own KB and keys	Strong isolation, per-tenant config/models, easy per-tenant deletion and export
Shared KB with security groups	One KB; resources tagged with tenant <code>access_groups</code> ; queries scoped by <code>security</code>	Many small tenants where a KB each is too much overhead

Key idea Default to a Knowledge Box per tenant when tenants are few or sensitive — isolation is structural, and deleting a tenant is deleting a KB. Use the shared-KB-with-security model only for large numbers of low-risk tenants, and test the `security` filter relentlessly, because there a bug is a cross-tenant data leak.

Cost and token management

Generative answers cost tokens, and retrieval quality is what keeps that cost down — a tight context is cheaper *and* better. Practical levers:

- **Watch the numbers.** The streaming `metadata` item on `/ask` reports input and output tokens per answer (Chapter 10). Log it and aggregate — it is your per-answer cost meter.
- **Right-size context.** Keep `top_k` as low as quality allows, and prefer `neighbouring_paragraphs` over `full_resource` unless whole documents are truly needed. More context is not free and often does not help.
- **Route models by difficulty.** Override `generative_model` per request — a cheaper model for routine questions, a stronger one only when needed.
- **Set budgets.** Account-level consumption controls let you cap and monitor spend; wire alerts before you hit limits.
- **Bring your own provider.** For high volume, connecting your own Anthropic, Bedrock, or Gemini account can change the cost equation and keep generation under your own contract.

Performance and latency

- **Stream answers.** Use streaming `/ask` so users see tokens immediately; the `metadata` item reports time-to-first-chunk so you can track it.
- **Cache the cacheable.** `catalog`, `suggest`, and popular answers are cache-friendly; do not re-embed a query you just served.
- **Pick the near zone.** Latency is dominated by distance — put the Knowledge Box in the zone closest to your users, within your data-residency rules.
- **Ingest asynchronously and in batches.** Do not block user requests on ingestion; let processing run and use webhooks (Chapter 8) rather than polling to learn when content is ready.

Reliability and correctness

- **Treat ``no_context`` as normal.** When retrieval finds nothing, `/ask` says so instead of hallucinating (Chapter 10). Handle that status in your UI — offer to broaden the search or fall back to a Retrieval Agent with internet fallback (Chapter 13).
- **Respect processing lag.** Content is not searchable the instant it is created. Poll status or use webhooks before promising users that new content is live.
- **Use the review cycle.** Enable hidden resources so new content is staged and reviewed before it appears in answers — indexing without exposing until approved.
- **Back up before destructive change.** Deletion is irreversible and takes keys with it; snapshot first (Chapter 17).
- **Monitor quality continuously.** Sample REMi metrics and watch user feedback (Chapter 16); regressions from a content or model change show up there first.

Key idea A production RAG system is judged on trust. Every practice in this chapter — least privilege, tenant isolation, honest `no_context`, staged review, continuous evaluation — exists to make answers something users can rely on. Speed and cost matter, but trust is the product.

PART 7 — REFERENCE APPENDIX

Appendix A — NucliaDB REST API

Reference

Knowledge Box content, search, and RAG — every endpoint, fully expanded

The content and retrieval surface of a Knowledge Box.

115 operations in **11 groups**, generated from the official `nucliadb` v1 OpenAPI specification. Served from `https://<zone>.dp.progress.cloud/api/v1`; authorized with a Knowledge Box API key (`Authorization: Bearer` or `X-NUCLIA-SERVICEACCOUNT`). Every request-body field, parameter, and response is listed; object-typed fields reference named schemas documented in full in Appendix F (Schema Catalog).

Knowledge Boxes (10)

GET /kb/s/{slug}

Get Knowledge Box (by slug) · roles: MANAGER, READER

Parameters

Name	In	Type	Req	Description
<code>slug</code>	path	string	yes	

Example — a representative request:

```
BASH
curl -X GET "https://$ZONE.dp.progress.cloud/api/v1/kb/s/{slug}" \
  -H "Authorization: Bearer $KB_API_KEY"
```

Response codes

Status	Schema	Description
200	<code>KnowledgeBoxObj</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

GET /kb/{kbid}

Get Knowledge Box · roles: MANAGER, READER

Parameters

Name	In	Type	Req	Description
<code>kbid</code>	path	string	yes	

Example — a real call captured against the live sandbox:

BASH

```
curl -X GET "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB" \
-H "Authorization: Bearer $KB_API_KEY"
```

Response:

JSON

```
{
  "slug": "arag-book-sandbox",
  "uuid": "$KB",
  "config": {
    "uuid": null,
    "slug": "$ACCOUNT:arag-book-sandbox",
    "title": "ARAG Book Sandbox (safe to delete)",
    "description": "Throwaway KB for live-testing while writing the book.",
    "learning_configuration": null,
    "external_index_provider": null,
    "configured_external_index_provider": {
      "type": "unset"
    },
    "similarity": null,
    "hidden_resources_enabled": false,
    "hidden_resources_hide_on_creation": false,
    "enforce_security": false
  },
  "model": null
}
```

Response codes

Status	Schema	Description
200	<code>KnowledgeBoxObj</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

POST /kb/{kbid}/configuration

Create Knowledge Box models configuration · roles: MANAGER, OWNER

Create configuration of models assigned to a Knowledge Box

Parameters

Name	In	Type	Req	Description
<code>kbid</code>	path	string	yes	

Example — a representative request:

BASH

```
curl -X POST "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/configuration" \
-H "Authorization: Bearer $KB_API_KEY"
```

Response codes

Status	Schema	Description
204		Successful Response
422	<code>HTTPValidationError</code>	Validation Error

PATCH /kb/{kbid}/configuration

Update Knowledge Box models configuration · roles: MANAGER, OWNER

Update current configuration of models assigned to a Knowledge Box

Parameters

Name	In	Type	Req	Description
<code>kbid</code>	path	string	yes	

Example — a representative request:

BASH

```
curl -X PATCH "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/configuration" \
-H "Authorization: Bearer $KB_API_KEY"
```

Response codes

Status	Schema	Description
204		Successful Response
422	HTTPValidationError	Validation Error

GET /kb/{kbid}/counters

Knowledgebox Counters · roles: READER, MANAGER

Summary of amount of different things inside a knowledgebox

Parameters

Name	In	Type	Req	Description
kbid	path	string	yes	
debug	query	boolean		If set, the response will include some extra metadata for debugging purposes, like the list of queried nodes.

Example — a representative request:

```
BASH
curl -X GET "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/counters" \
  -H "Authorization: Bearer $KB_API_KEY"
```

Response codes

Status	Schema	Description
200	KnowledgeboxCounters	Successful Response
422	HTTPValidationError	Validation Error

GET /kb/{kbid}/export/{export_id}

Download a Knowledge Box export · roles: MANAGER, READER

Parameters

Name	In	Type	Req	Description
<code>kbid</code>	path	string	yes	
<code>export_id</code>	path	string	yes	

Example — a representative request:

```
BASH
curl -X GET "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/export/{export_id}" \
-H "Authorization: Bearer $KB_API_KEY"
```

Response codes

Status	Schema	Description
200		Successful Response
422	<code>HTTPValidationError</code>	Validation Error

GET /kb/{kbid}/export/{export_id}/status

Get the status of a Knowledge Box Export · roles: MANAGER, READER

Parameters

Name	In	Type	Req	Description
<code>kbid</code>	path	string	yes	
<code>export_id</code>	path	string	yes	

Example — a representative request:

```
BASH
curl -X GET "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/export/{export_id}/status" \
-H "Authorization: Bearer $KB_API_KEY"
```

Response codes

Status	Schema	Description
200	<code>StatusResponse</code>	Successful Response

422	HTTPValidationError	Validation Error
-----	-------------------------------------	------------------

GET /kb/{kbid}/import/{import_id}/status

Get the status of a Knowledge Box Import · roles: MANAGER, READER

Parameters

Name	In	Type	Req	Description
kbid	path	string	yes	
import_id	path	string	yes	

Example — a representative request:

```
BASH
curl -X GET "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/import/{import_id}/status" \
-H "Authorization: Bearer $KB_API_KEY"
```

Response codes

Status	Schema	Description
200	StatusResponse	Successful Response
422	HTTPValidationError	Validation Error

POST /kb/{kbid}/upload

Upload binary file on a Knowledge Box · roles: WRITER

Upload a file onto a Knowledge Box, field id will be file and rid will be autogenerated.

Parameters

Name	In	Type	Req	Description
kbid	path	string	yes	
x-filename	header	string null		Name of the file being uploaded.
x-password	header	string null		If the file is password protected, the password must be

				provided here.
x-language	header	string null		
x-md5	header	string null		MD5 hash of the file being uploaded. This is used to check if the file has been uploaded before.
x-extract-strategy	header	string null		Extract strategy to use when uploading a file. If not provided, the default strategy will be used.
x-split-strategy	header	string null		Split strategy to use when uploading a file. If not provided, the default strategy will be used.

Example — a real call captured against the live sandbox:

BASH

```
curl -X POST "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/upload" \
-H "Authorization: Bearer $KB_API_KEY"
```

Response:

JSON

```
{
  "seqid": null,
  "uuid": "1009f3c71b934964aecb012f5d18e390",
  "field_id": "1d7c1a94d9ca4a4ab6ad9a12a4342e8a"
}
```

Response codes

Status	Schema	Description
201	ResourceFileUploaded	Successful Response
422	HTTPValidationError	Validation Error

GET /learning/configuration/schema

Learning Configuration Schema

Get jsonschema definition for learningconfiguration field of knowledgebox creation payload

Example — a representative request:

```
BASH
curl -X GET
"https://$ZONE.dp.progress.cloud/api/v1/learning/configuration/schema" \
-H "Authorization: Bearer $KB_API_KEY"
```

Response codes

Status	Schema	Description
200	<code>object</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

Search (19)

POST /kb/{kbid}/ask

Ask Knowledge Box · roles: `READER`

Ask questions on a Knowledge Box

Parameters

Name	In	Type	Req	Description
<code>kbid</code>	path	string	yes	
<code>x-ndb-client</code>	header	NucliaDBClientType		
<code>x-show-consumption</code>	header	boolean		
<code>x-nucliadb-user</code>	header	string		
<code>x-forwarded-for</code>	header	string		
<code>x-synchronous</code>	header	boolean		When set to true, outputs response as JSON in a non-streaming way. This is slower and requires waiting for entire

				answer to be ready.
--	--	--	--	---------------------

Request body — `application/json` (required)

Field	Type	Req	Default	Description
<code>audit_metadata</code>	<code>map<string, string> null</code>			A dictionary containing optional audit-specific metadata, such as <code>userid</code> , <code>environment</code> , or other contextual information. This metadata can be leveraged for filtering and analyzing activity logs in futu...
<code>query</code>	<code>string</code>	yes		The query to get a generative answer for
<code>top_k</code>	<code>integer</code>		20	The top most relevant results to fetch at the retrieval step. The maximum number of results allowed is 200.
<code>filter_expression</code>	<code>FilterExpression null</code>			Returns only documents that match this filter expression. Filtering examples can be found here: (see docs) This allows building complex filtering expressions and replaces the following parameters: <code>field...</code>
<code>fields</code>	<code>array<string></code>		[]	The list of fields to search in. For instance: <code>a/title</code> to search only on title field. For more details on filtering by field, see: (see docs)

<code>filters</code>	<code>array<string> array<Filter></code>		<code>[]</code>	The list of filters to apply. Filtering examples can be found here: (see docs)
<code>keyword_filters</code>	<code>array<string> array<Filter></code>		<code>[]</code>	List of keyword filter expressions to apply to the retrieval step. The text block search will only be performed on the documents that contain the specified keywords. The filters are case-insensitive, ...
<code>vectorset</code>	<code>string null</code>			Vectors index to perform the search in. If not provided, NucliaDB will use the default one
<code>min_score</code>	<code>number MinScore null</code>			Minimum score to filter search results. Results with a lower score will be ignored. Accepts either a float or a dictionary with the minimum scores for the bm25 and vector indexes. If a float is provid...
<code>features</code>	<code>array<ChatOptions></code>		<code>['semantic', 'keyword']</code>	Features enabled for the chat endpoint. Semantic search is done if semantic is included. If keyword is included, the results will include matching paragraphs from the bm25 index. If relations is inclu...
<code>range_creation_start</code>	<code>string null</code>			Resources created before this date will be filtered out of search results. Datetime are

				represented as a str in ISO 8601 format, like: 2008-09-15T15:53:00+05:00.
range_creation_end	string null			Resources created after this date will be filtered out of search results. Datetime are represented as a str in ISO 8601 format, like: 2008-09-15T15:53:00+05:00.
range_modification_start	string null			Resources modified before this date will be filtered out of search results. Datetime are represented as a str in ISO 8601 format, like: 2008-09-15T15:53:00+05:00.
range_modification_end	string null			Resources modified after this date will be filtered out of search results. Datetime are represented as a str in ISO 8601 format, like: 2008-09-15T15:53:00+05:00.
show	array<ResourceProperties>		['basic']	Controls which types of metadata are serialized on resources of search results
field_type_filter	array<FieldTypeName>		['text', 'file', 'link', 'conversation', 'generic', 'key_value']	Define which field types are serialized on resources of search results
extracted	array<ExtractedData TypeName>		[]	[Deprecated] Please use GET resource endpoint instead to get extracted

				metadata
context	array<ChatContextMessage> null			DEPRECATED! Please, use chathistory instead.
chat_history	array<ChatContextMessage> null			Use to rephrase the new LLM query by taking into account the chat conversation history. This will be passed to the LLM so that it is aware of the previous conversation.
extra_context	array<string> null			Additional context that is added to the retrieval context sent to the LLM. It allows extending the chat feature with content that may not be in the Knowledge Box.
extra_context_images	array<Image> null			Additional images added to the retrieval context sent to the LLM." It allows extending the chat feature with content that may not be in the Knowledge Box.
query_image	Image null			Image that will be used together with the query text for retrieval and then sent to the LLM as part of the context. If a query image is provided, the extracontextimages and ragimagesstrategies will be...
highlight	boolean		False	If set to true, the query terms will be highlighted in the results between

				<mark>...</mark> tags
resource_filters	array<string>		[]	List of resource ids to filter search results for. Only paragraphs from the specified resources will be returned.
prompt	string CustomPrompt null			Use to customize the prompts given to the generative model. Both system and user prompts can be customized. If a string is provided, it is interpreted as the user prompt.
rank_fusion	RankFusionName ReciprocalRankFusion		rrf	Rank fusion algorithm to use to merge results from multiple retrievers (keyword, semantic)
reranker	RerankerName PredictReranker		predict	Reranker let you specify which method you want to use to rerank your results at the end of retrieval
citations	boolean CitationsType null			Whether to include citations in the response. If set to None or False, no citations will be computed. If set to True or 'default', citations will be computed after answer generation and send as a sepa...
citation_threshold	number null			If citations is set to True or 'default', this will be the similarity threshold. Value between 0 and 1, lower values will produce more

				citations. If not set, it will be set to the optimized threshold ...
security	RequestSecurity null			Security metadata for the request. Please refer to the documentation for more details on how security works: (see docs)
show_hidden	boolean		False	If set to false (default), excludes hidden resources from search
rag_strategies	array<FieldExtensionStrategy FullResourceStrategy HierarchyResourceStrategy NeighbouringParagraphsStrategy MetadataExtensionStrategy ConversationalStrategy PreQueriesStrategy GraphStrategy>		[]	Options for tweaking how the context for the LLM model is crafted: - fullresource will add the full text of the matching resources to the context. This strategy cannot be combined with hierarchy, neig...
rag_images_strategies	array<PageImageStrategy ParagraphImageStrategy TableImageStrategy>		[]	Options for tweaking how the image based context for the LLM model is crafted: - pageimage will add the full page image of the matching resources to the context. - tables will send the table images fo...
debug	boolean		False	If set, the response will include some extra metadata for debugging purposes, like the list of queried

				nodes.
<code>generative_model</code>	string null			The generative model to use for the chat endpoint. If not provided, the model configured for the Knowledge Box is used.
<code>generative_model_seed</code>	integer null			The seed to use for the generative model for deterministic generation. Only supported by some models.
<code>max_tokens</code>	integer MaxTokens null			Use to limit the amount of tokens used in the LLM context and/or for generating the answer. If not provided, the default maximum tokens of the generative model will be used. If an integer is provided,...
<code>rephrase</code>	boolean		<code>False</code>	Rephrase the query for a more efficient retrieval. This will consume LLM tokens and make the request slower.
<code>chat_history_relevance_threshold</code>	number null			Threshold to determine if the past chat history is relevant to rephrase the user's question. 0 - Always treat previous messages as relevant (always rephrase). 1 - Always treat previous messages as irre...
<code>prefer_markdown</code>	boolean		<code>False</code>	If set to true, the response will be in markdown format

<code>answer_json_schema</code>	object (free-form map) null			Desired JSON schema for the LLM answer. This schema is passed to the LLM so that it answers in a structured format following the schema. If not provided, textual response is returned. Note that when u...
<code>generate_answer</code>	boolean		<code>True</code>	Whether to generate an answer using the generative model. If set to false, the response will only contain the retrieval results.
<code>search_configuration</code>	string null			Load ask parameters from this configuration. Parameters in the request override parameters from the configuration.
<code>reasoning</code>	Reasoning boolean		<code>False</code>	Reasoning options for the generative model. Set to True to enable default reasoning, False to disable, or provide a Reasoning object for custom options.

Example — a real call captured against the live sandbox:

```
BASH
curl -X POST "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/ask" \
  -H "Authorization: Bearer $KB_API_KEY" \
  -H "Content-Type: application/json" \
  -d '{
    "query": "How much does a home battery store and when is it subsidised?",
    "top_k": 5,
    "features": [
      "keyword",
      "semantic"
    ],
  },
```

```
"citations": true
}'
```

Response:

```
JSON
{
  "answer": "A home battery stores 10kWh of usable capacity and is subsidised
starting from July 2025 under the federal Cheaper Home Batteries program.",
  "status": "success",
  "retrieval_results": {
    "resources": {
      "86b5f7948b144b008f018cbd47a2bde4": {
        "id": "86b5f7948b144b008f018cbd47a2bde4",
        "slug": "battery-storage",
        "title": "Home Battery Storage",
        "summary": "",
        "icon": "text/plain",
        "thumbnail": "",
        "metadata": {
          "metadata": {},
          "language": "en",
          "languages": [
            "en"
          ],
        },
        "status": "PROCESSED"
      },
    },
    "usermetadata": {
      "classifications": [
        {
          "labelset": "topic",
          "label": "energy",
          "cancelled_by_user": false
        }
      ],
      "relations": []
    },
  },
  "fi
... (truncated)
```

Response codes

Status	Schema	Description
200	SyncAskResponse	Successful Response
422	HTTPValidationError	Validation Error

GET /kb/{kbid}/catalog

List resources of a Knowledge Box · roles: READER

Parameters

Name	In	Type	Req	Description
kbid	path	string	yes	
query	query	string		The query to search for
filter_expression	query	string null		Returns only documents that match this filter expression. Filtering examples can be found here: (see docs) This allows building complex filtering expressions
filters	query	array<string>		The list of filters to apply. Filtering examples can be found here: (see docs)
faceted	query	array<string>		The list of facets to calculate. The facets follow the same syntax as filters: (see docs)
sort_field	query	SortField		Field to sort results with (Score not supported in catalog)
sort_order	query	SortOrder		Order to sort results with
page_number	query	integer		The page number of the results to return
page_size	query	integer		The number of results to return per page. The maximum number of results per page allowed is 200.
with_status	query	ResourceProcessingStatus null		Filter results by resource processing

				status
<code>range_creation_start</code>	query	string null		Resources created before this date will be filtered out of search results. Datetime are represented as a str in ISO 8601 format, like: 2008-09-15T15:5...
<code>range_creation_end</code>	query	string null		Resources created after this date will be filtered out of search results. Datetime are represented as a str in ISO 8601 format, like: 2008-09-15T15:53...
<code>range_modification_start</code>	query	string null		Resources modified before this date will be filtered out of search results. Datetime are represented as a str in ISO 8601 format, like: 2008-09-15T15:...
<code>range_modification_end</code>	query	string null		Resources modified after this date will be filtered out of search results. Datetime are represented as a str in ISO 8601 format, like: 2008-09-15T15:5...
<code>hidden</code>	query	boolean null		Set to filter only hidden or only non-hidden resources. Default is to return everything
<code>show</code>	query	array<ResourceProperties>		Controls which types of metadata are serialized on resources of search results

Example — a real call captured against the live sandbox:

BASH

```
curl -X GET "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/catalog" \
-H "Authorization: Bearer $KB_API_KEY"
```

Response:

JSON

```
{
  "resources": {
    "df6191306a97408db56025102074aac0": {
      "id": "df6191306a97408db56025102074aac0",
      "slug": "progress-site",
      "title": "Progress site",
      "summary": "",
      "icon": "application/generic",
      "thumbnail": "",
      "metadata": {
        "metadata": {},
        "language": "",
        "languages": [],
        "status": "PENDING"
      },
      "usermetadata": {
        "classifications": []
      },
      "fieldmetadata": [],
      "computedmetadata": {
        "field_classifications": []
      },
      "created": "2026-07-27T05:24:19.555446",
      "modified": "2026-07-27T05:24:19.555461",
      "last_seqid": 0,
      "last_account_seq": null,
      "queue": "private",
      "hidden": false,
      "data": {
        "links": {
          "home": {
            "status": "PENDING"
          }
        },
        "generics": {
          "title": {
            ... (truncated)
          }
        }
      }
    }
  }
}
```

Response codes

Status	Schema	Description
--------	--------	-------------

200	<code>KnowledgeboxSearchResults</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

POST /kb/{kbid}/catalog

List resources of a Knowledge Box · roles: READER

Parameters

Name	In	Type	Req	Description
<code>kbid</code>	path	string	yes	

Request body — `application/json` (required)

Field	Type	Req	Default	Description
<code>query</code>	string <code>CatalogQuery</code>			The query to search for
<code>filter_expression</code>	<code>CatalogFilterExpression</code> null			Returns only documents that match this filter expression. Filtering examples can be found here: (see docs) This allows building complex filtering expressions and replaces the following parameters: <code>filters:filter...</code>
<code>faceted</code>	array<string>		<code>[]</code>	The list of facets to calculate. The facets follow the same syntax as filters: (see docs)
<code>sort</code>	<code>SortOptions</code> null			Options for results sorting
<code>page_number</code>	integer		<code>0</code>	The page number of the results to return

<code>page_size</code>	integer		<code>20</code>	The number of results to return per page. The maximum number of results per page allowed is 200.
<code>hidden</code>	boolean null			Set to filter only hidden or only non-hidden resources. Default is to return everything
<code>show</code>	array<ResourceProperties>		<code>['basic', 'errors']</code>	Controls which types of metadata are serialized on resources of search results
<code>filters</code>	array<string> array<Filter>		<code>[]</code>	The list of filters to apply. Filtering examples can be found here: (see docs)
<code>with_status</code>	ResourceProcessingStatus null			Filter results by resource processing status
<code>range_creation_start</code>	string null			Resources created before this date will be filtered out of search results. Datetime are represented as a str in ISO 8601 format, like: 2008-09-15T15:53:00+05:00.
<code>range_creation_end</code>	string null			Resources created after this date will be filtered out of search results. Datetime are represented as a str in ISO 8601 format, like: 2008-09-15T15:53:00+05:00.
<code>range_modification_start</code>	string null			Resources modified before this date will be filtered out of search results. Datetime are

				represented as a str in ISO 8601 format, like: 2008-09-15T15:53:00+05:00.
<code>range_modification_end</code>	string null			Resources modified after this date will be filtered out of search results. Datetime are represented as a str in ISO 8601 format, like: 2008-09-15T15:53:00+05:00.

Example — a real call captured against the live sandbox:

BASH

```
curl -X POST "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/catalog" \
-H "Authorization: Bearer $KB_API_KEY" \
-H "Content-Type: application/json" \
-d '{
  "query": "",
  "faceted": [
    "/classification.labels/topic"
  ]
}'
```

Response:

JSON

```
{
  "resources": {
    "df6191306a97408db56025102074aac0": {
      "id": "df6191306a97408db56025102074aac0",
      "slug": "progress-site",
      "title": "Progress site",
      "summary": "",
      "icon": "application/stf-link",
      "thumbnail":
"/kb/$KB/resource/df6191306a97408db56025102074aac0/link/home/download/extracted/
link_thumbnail",
      "metadata": {
        "metadata": {},
        "language": "en",
        "languages": [
          "en",
          "fr"
        ],
        "status": "PROCESSED"
      },
    },
  },
}
```

```

"usermetadata": {
  "classifications": []
},
"fieldmetadata": [],
"computedmetadata": {
  "field_classifications": []
},
"created": "2026-07-27T05:24:19.555446",
"modified": "2026-07-27T05:24:19.555461",
"last_seqid": 0,
"last_account_seq": null,
"queue": "private",
"hidden": false,
"data": {
... (truncated)

```

Response codes

Status	Schema	Description
200	<code>KnowledgeboxSearchResults</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

POST /kb/{kbid}/feedback

Send Feedback · roles: READER

Send feedback for a search operation in a Knowledge Box

Parameters

Name	In	Type	Req	Description
<code>kbid</code>	path	string	yes	
<code>x-ndb-client</code>	header	NucliaDBClientType		
<code>x-nucliadb-user</code>	header	string		
<code>x-forwarded-for</code>	header	string		

Request body — `application/json` (required)

Field	Type	Req	Default	Description
-------	------	-----	---------	-------------

<code>ident</code>	string	yes		Id of the request to provide feedback for. This id is returned in the response header Nuclia-Learning-Id of the chat endpoint.
<code>good</code>	boolean	yes		Whether the result was good or not
<code>task</code>	FeedbackTasks	yes		The task the feedback is for. For now, only CHAT task is available
<code>feedback</code>	string null			Feedback text
<code>text_block_id</code>	string null			Text block id

Example — a real call captured against the live sandbox:

BASH

```
curl -X POST "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/feedback" \
-H "Authorization: Bearer $KB_API_KEY" \
-H "Content-Type: application/json" \
-d '{
  "ident": "test",
  "good": true,
  "task": "CHAT",
  "feedback": "great"
}'
```

Response:

JSON

None

Response codes

Status	Schema	Description
200	<code>object</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

GET /kb/{kbid}/find

Find Knowledge Box · roles: READER

Find on a Knowledge Box

Parameters

Name	In	Type	Req	Description
<code>kbid</code>	path	string	yes	
<code>query</code>	query	string		The query to search for
<code>filter_expression</code>	query	string null		Returns only documents that match this filter expression. Filtering examples can be found here: (see docs) This allows building complex filtering expressions...
<code>fields</code>	query	array<string>		The list of fields to search in. For instance: a/title to search only on title field. For more details on filtering by field, see: (see docs)
<code>filters</code>	query	array<string>		The list of filters to apply. Filtering examples can be found here: (see docs)
<code>top_k</code>	query	integer null		The number of results search should return. The maximum number of results allowed is 200.
<code>min_score</code>	query	number null		Minimum similarity score to filter vector index results. If not specified, the default minimum score of the semantic model associated to the Knowledge...
<code>min_score_semantic</code>	query	number null		Minimum semantic similarity score to

				filter vector index results. If not specified, the default minimum score of the semantic model associated to the ...
<code>min_score_bm25</code>	query	number		Minimum bm25 score to filter paragraph and document index results
<code>vectorset</code>	query	string null		Vectors index to perform the search in. If not provided, NucliaDB will use the default one
<code>range_creation_start</code>	query	string null		Resources created before this date will be filtered out of search results. Datetime are represented as a str in ISO 8601 format, like: 2008-09-15T15:5...
<code>range_creation_end</code>	query	string null		Resources created after this date will be filtered out of search results. Datetime are represented as a str in ISO 8601 format, like: 2008-09-15T15:53...
<code>range_modification_start</code>	query	string null		Resources modified before this date will be filtered out of search results. Datetime are represented as a str in ISO 8601 format, like: 2008-09-15T15:...
<code>range_modification_end</code>	query	string null		Resources modified after this date will be filtered out of search

				results. Datetime are represented as a str in ISO 8601 format, like: 2008-09-15T15:5...
features	query	array<FindOptions>		List of search features to use. Each value corresponds to a lookup into one of the different indexes
debug	query	boolean		If set, the response will include some extra metadata for debugging purposes, like the list of queried nodes.
highlight	query	boolean		If set to true, the query terms will be highlighted in the results between <mark>...</mark> tags
show	query	array<ResourceProperties>		Controls which types of metadata are serialized on resources of search results
field_type	query	array<FieldTypeName>		Define which field types are serialized on resources of search results
extracted	query	array<ExtractedDataTypeName>		[Deprecated] Please use GET resource endpoint instead to get extracted metadata
with_duplicates	query	boolean		Whether to return duplicate paragraphs on the same document
with_synonyms	query	boolean		Whether to return matches for custom knowledge box synonyms of the

				query terms. Note: only supported for keyword and fulltext search options.
<code>security_groups</code>	query	array<string> null		List of security groups to filter search results for. Only resources matching the query and containing the specified security groups will be returned....
<code>show_hidden</code>	query	boolean		If set to false (default), excludes hidden resources from search
<code>rank_fusion</code>	query	RankFusionName		Rank fusion algorithm to use to merge results from multiple retrievers (keyword, semantic)
<code>reranker</code>	query	RerankerName		Reranker let you specify which method you want to use to rerank your results at the end of retrieval
<code>search_configuration</code>	query	string null		Load find parameters from this configuration. Parameters in the request override parameters from the configuration.
<code>x-ndb-client</code>	header	NucliaDBClientType		
<code>x-nucliadb-user</code>	header	string		
<code>x-forwarded-for</code>	header	string		

Example — a representative request:

```
BASH
```



```
curl -X GET "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/find" \
-H "Authorization: Bearer $KB_API_KEY"
```

Response codes

Status	Schema	Description
200	<code>KnowledgeboxFindResults</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

POST /kb/{kbid}/find

Find Knowledge Box · roles: READER

Find on a Knowledge Box

Parameters

Name	In	Type	Req	Description
<code>kbid</code>	path	string	yes	
<code>x-ndb-client</code>	header	NucliaDBClientType		
<code>x-nucliadb-user</code>	header	string		
<code>x-forwarded-for</code>	header	string		

Request body — `application/json` (required)

Field	Type	Req	Default	Description
<code>audit_metadata</code>	<code>map<string, string> null</code>			A dictionary containing optional audit-specific metadata, such as userid, environment, or other contextual information. This metadata can be leveraged for filtering and analyzing activity logs in futu...

query	string			The query to search for
filter_expression	FilterExpression null			Returns only documents that match this filter expression. Filtering examples can be found here: (see docs) This allows building complex filtering expressions and replaces the following parameters: field...
fields	array<string>		[]	The list of fields to search in. For instance: a/title to search only on title field. For more details on filtering by field, see: (see docs)
filters	array<string> array<Filter>		[]	The list of filters to apply. Filtering examples can be found here: (see docs)
top_k	integer		20	The number of results search should return. The maximum number of results allowed is 200.
min_score	number MinScore null			Minimum score to filter search results. Results with a lower score will be ignored. Accepts either a float or a dictionary with the minimum scores for the bm25 and vector indexes. If a float is provid...
range_creation_start	string null			Resources created before this date will be filtered out of search results.

				Datetime are represented as a str in ISO 8601 format, like: 2008-09-15T15:53:00+05:00.
<code>range_creation_end</code>	string null			Resources created after this date will be filtered out of search results. Datetime are represented as a str in ISO 8601 format, like: 2008-09-15T15:53:00+05:00.
<code>range_modification_start</code>	string null			Resources modified before this date will be filtered out of search results. Datetime are represented as a str in ISO 8601 format, like: 2008-09-15T15:53:00+05:00.
<code>range_modification_end</code>	string null			Resources modified after this date will be filtered out of search results. Datetime are represented as a str in ISO 8601 format, like: 2008-09-15T15:53:00+05:00.
<code>debug</code>	boolean		<code>False</code>	If set, the response will include some extra metadata for debugging purposes, like the list of queried nodes.
<code>highlight</code>	boolean		<code>False</code>	If set to true, the query terms will be highlighted in the results between <code><mark>...</mark></code> tags
<code>show</code>	array<ResourceProperties>		<code>['basic']</code>	Controls which types of metadata are serialized on

				resources of search results
<code>field_type_filter</code>	<code>array<FieldTypeName></code>		<code>['text', 'file', 'link', 'conversation', 'generic', 'key_value']</code>	Define which field types are serialized on resources of search results
<code>extracted</code>	<code>array<ExtractedData TypeName></code>		<code>[]</code>	[Deprecated] Please use GET resource endpoint instead to get extracted metadata
<code>vector</code>	<code>array<number> null</code>			The vector to perform the search with. If not provided, NucliaDB will use Nuclia Predict API to create the vector off from the query.
<code>vectorset</code>	<code>string null</code>			Vectors index to perform the search in. If not provided, NucliaDB will use the default one
<code>with_duplicates</code>	<code>boolean</code>		<code>False</code>	Whether to return duplicate paragraphs on the same document
<code>with_synonyms</code>	<code>boolean</code>		<code>False</code>	Whether to return matches for custom knowledge box synonyms of the query terms. Note: only supported for keyword and fulltext search options.
<code>resource_filters</code>	<code>array<string></code>		<code>[]</code>	List of resource ids to filter search results for. Only paragraphs from the specified resources will be returned.

security	RequestSecurity null			Security metadata for the request. Please refer to the documentation for more details on how security works: (see docs)
show_hidden	boolean		False	If set to false (default), excludes hidden resources from search
rephrase	boolean		False	Rephrase the query for a more efficient retrieval. This will consume LLM tokens and make the request slower.
rephrase_prompt	string null			Rephrase prompt given to the generative model responsible for rephrasing the query for a more effective retrieval step. This is only used if the rephrase flag is set to true in the request. If not spe...
query_image	Image null			Image that will be used together with the query text for retrieval.
graph_query	And_GraphPathQuery_ Or_GraphPathQuery_ Not_GraphPathQuery_ GraphPath-Input SourceNode DestinationNode AnyNode Relation-Input nucliadb_models__graph__requests__Generated null			Query for the knowledge graph. Paths (node-relation-node) extracted from a paragraphid will be used to extend the results

features	array<FindOptions>		['keyword', 'semantic']	List of search features to use. Each value corresponds to a lookup into one of the different indexes
rank_fusion	RankFusionName ReciprocalRankFusion		rrf	Rank fusion algorithm to use to merge results from multiple retrievers (keyword, semantic)
reranker	RerankerName PredictReranker		predict	Reranker let you specify which method you want to use to rerank your results at the end of retrieval
keyword_filters	array<string> array<Filter>		[]	List of keyword filter expressions to apply to the retrieval step. The text block search will only be performed on the documents that contain the specified keywords. The filters are case-insensitive, ...
search_configuration	string null			Load find parameters from this configuration. Parameters in the request override parameters from the configuration.
generative_model	string null			The generative model used to rephrase the query. If not provided, the model configured for the Knowledge Box is used.

Example — a real call captured against the live sandbox:

```
BASH
```

```
curl -X POST "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/find" \
-H "Authorization: Bearer $KB_API_KEY" \
-H "Content-Type: application/json" \
-d '{
  "query": "how long is the payback period for solar",
  "features": [
    "keyword",
    "semantic"
  ],
  "top_k": 5
}'
```

Response:

```
JSON
{
  "resources": {
    "86b5f7948b144b008f018cbd47a2bde4": {
      "id": "86b5f7948b144b008f018cbd47a2bde4",
      "slug": "battery-storage",
      "title": "Home Battery Storage",
      "summary": "",
      "icon": "text/plain",
      "thumbnail": "",
      "metadata": {
        "metadata": {},
        "language": "en",
        "languages": [
          "en"
        ],
        "status": "PROCESSED"
      },
      "usermetadata": {
        "classifications": [
          {
            "labelset": "topic",
            "label": "energy",
            "cancelled_by_user": false
          }
        ]
      },
      "fieldmetadata": [],
      "computedmetadata": {
        "field_classifications": []
      },
      "created": "2026-07-27T05:24:18.542195",
      "modified": "2026-07-27T05:24:18.542214",
      "last_seqid": 0,
      "last_account_seq": null,
      "queue": "private",
      "hidden": false,
      ... (truncated)
    }
  }
}
```

Response codes

Status	Schema	Description
200	<code>KnowledgeboxFindResults</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

POST /kb/{kbid}/graph

Search Knowledge Box graph · roles: `READER`

Search on the Knowledge Box graph and retrieve triplets of vertex-edge-vertex

Parameters

Name	In	Type	Req	Description
<code>kbid</code>	path	string	yes	
<code>x-ndb-client</code>	header	NucliaDBClientType		
<code>x-nucliadb-user</code>	header	string		
<code>x-forwarded-for</code>	header	string		

Request body — `application/json` (required)

Field	Type	Req	Default	Description
<code>top_k</code>	integer		<code>50</code>	
<code>filter_expression</code>	GraphFilterExpression null			Returns only relations from documents that match this filter expression. Filtering examples can be found here: (see docs)
<code>security</code>	RequestSecurity null			Security metadata for the request. Please refer to the documentation for more details on how security works: (see docs)

show_hidden	boolean		False	If set to false (default), excludes hidden resources from search
query	And_GraphPathQuery Or_GraphPathQuery Not_GraphPathQuery GraphPath-Input SourceNode DestinationNode AnyNode Relation-Input nucliadb_models__graph__requests__Generated	yes		

Example — a real call captured against the live sandbox:

BASH

```
curl -X POST "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/graph" \
-H "Authorization: Bearer $KB_API_KEY" \
-H "Content-Type: application/json" \
-d '{
  "query": {
    "prop": "path"
  },
  "top_k": 10
}'
```

Response:

JSON

```
{
  "paths": [
    {
      "source": {
        "value": "ca2a8f3bf186491facde30cdb2fbc2b6",
        "type": "resource",
        "group": ""
      },
      "relation": {
        "label": "",
        "type": "ABOUT"
      },
      "destination": {
        "value": "topic/energy",
        "type": "label",

```

```

    "group": ""
  },
  "metadata": {
    "field_id": "ca2a8f3bf186491facde30cdb2fbc2b6/a/metadata",
    "paragraph_id": null
  },
  "score": 1.0
},
{
  "source": {
    "value": "ca2a8f3bf186491facde30cdb2fbc2b6",
    "type": "resource",
    "group": ""
  },
  "relation": {
    "label": "",
    "type": "ENTITY"
  },
  "destination": {
    "value": "Australia",
    "type": "entity",
    "group": "GPE"
  },
  "metadata": {
    "field_id": "ca2a8f3bf186491facde30cdb2fbc2b6/t/body",
    "para
... (truncated)

```

Response codes

Status	Schema	Description
200	GraphSearchResponse	Successful Response
422	HTTPValidationError	Validation Error

POST /kb/{kbid}/graph/nodes

Search Knowledge Box graph nodes · roles: READER

Search on the Knowledge Box graph and retrieve nodes (vertices)

Parameters

Name	In	Type	Req	Description
kbid	path	string	yes	
x-ndb-client	header	NucliaDBClientType		
x-nucliadb-	header	string		

user				
x-forwarded-for	header	string		

Request body — `application/json` (required)

Field	Type	Req	Default	Description
top_k	integer		50	
filter_expression	GraphFilterExpression null			Returns only relations from documents that match this filter expression. Filtering examples can be found here: (see docs)
security	RequestSecurity null			Security metadata for the request. Please refer to the documentation for more details on how security works: (see docs)
show_hidden	boolean		False	If set to false (default), excludes hidden resources from search
query	And_GraphNodesQuery_ Or_GraphNodesQuery_ Not_GraphNodesQuery_ AnyNode nucliadb_models__graph__requests__Generated	yes		

Example — a real call captured against the live sandbox:

BASH

```
curl -X POST "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/graph/nodes" \
-H "Authorization: Bearer $KB_API_KEY" \
-H "Content-Type: application/json" \
-d '{
```

```
"query": {
  "prop": "node"
},
"top_k": 5
}'
```

Response:

```
JSON
{
  "nodes": [
    {
      "value": "Private Knowledge Box",
      "type": "entity",
      "group": "ORG",
      "score": 1.037416696548462
    },
    {
      "value": "Enterprise Knowledge Management",
      "type": "entity",
      "group": "ORG",
      "score": 1.037416696548462
    },
    {
      "value": "Progress",
      "type": "entity",
      "group": "ORG",
      "score": 1.037416696548462
    },
    {
      "value": "90-day",
      "type": "entity",
      "group": "DATE",
      "score": 1.037416696548462
    },
    {
      "value": "2026",
      "type": "entity",
      "group": "DATE",
      "score": 1.037416696548462
    }
  ]
}
```

Response codes

Status	Schema	Description
200	GraphNodesSearchResponse	Successful Response
422	HTTPValidationError	Validation Error

POST /kb/{kbid}/graph/relationsSearch Knowledge Box graph relations · roles: **READER**

Search on the Knowledge Box graph and retrieve relations (edges)

Parameters

Name	In	Type	Req	Description
kbid	path	string	yes	
x-ndb-client	header	NucliaDBClientType		
x-nucliadb-user	header	string		
x-forwarded-for	header	string		

Request body — `application/json` (required)

Field	Type	Req	Default	Description
top_k	integer		50	
filter_expression	GraphFilterExpression null			Returns only relations from documents that match this filter expression. Filtering examples can be found here: (see docs)
security	RequestSecurity null			Security metadata for the request. Please refer to the documentation for more details on how security works: (see docs)
show_hidden	boolean		False	If set to false (default), excludes hidden resources from search
query	And_GraphRelationsQuery_	yes		

	Or_GraphRelationsQuery_ Not_GraphRelationsQuery_ Relation-Input nucliadb_models__graph__requests__Generated			
--	---	--	--	--

Example — a representative request:

```
BASH
curl -X POST "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/graph/relations" \
-H "Authorization: Bearer $KB_API_KEY" \
-H "Content-Type: application/json" \
-d '{
  "query": {
    "operands": [
      {}
    ]
  },
  "top_k": 50,
  "show_hidden": false
}'
```

Response codes

Status	Schema	Description
200	GraphRelationsSearchResponse	Successful Response
422	HTTPValidationError	Validation Error

GET /kb/{kbid}/predict/{endpoint}

Predict API Proxy · roles: **READER**

Convenience endpoint that proxies requests to the Predict API. It adds the Knowledge Box configuration settings as headers to the predict API request. Refer to the Predict API documentation for more details about the request and response models: (see docs)

Parameters

Name	In	Type	Req	Description
kbid	path	string	yes	

endpoint	path	PredictProxiedEndpoints	yes	
x-nucliadb-user	header	string		
x-ndb-client	header	NucliaDBClientType		
x-forwarded-for	header	string		

Example — a representative request:

```
BASH
curl -X GET "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/predict/{endpoint}" \
-H "Authorization: Bearer $KB_API_KEY"
```

Response codes

Status	Schema	Description
200	object	Successful Response
422	HTTPValidationError	Validation Error

POST /kb/{kbid}/predict/{endpoint}

Predict API Proxy · roles: READER

Convenience endpoint that proxies requests to the Predict API. It adds the Knowledge Box configuration settings as headers to the predict API request. Refer to the Predict API documentation for more details about the request and response models: (see docs)

Parameters

Name	In	Type	Req	Description
kbid	path	string	yes	
endpoint	path	PredictProxiedEndpoints	yes	
x-nucliadb-user	header	string		
x-ndb-client	header	NucliaDBClientType		

x-forwarded-for	header	string		
-----------------	--------	--------	--	--

Example — a representative request:

BASH

```
curl -X POST "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/predict/{endpoint}" \
-H "Authorization: Bearer $KB_API_KEY"
```

Response codes

Status	Schema	Description
200	object	Successful Response
422	HTTPValidationError	Validation Error

POST /kb/{kbid}/resource/{rid}/ask

Ask a resource (by id) · roles: READER

Ask questions to a resource

Parameters

Name	In	Type	Req	Description
kbid	path	string	yes	
rid	path	string	yes	
x-show-consumption	header	boolean		
x-ndb-client	header	NucliaDBClientType		
x-nucliadb-user	header	string		
x-forwarded-for	header	string		
x-synchronous	header	boolean		When set to true, outputs response as JSON in a non-streaming way. This is slower and requires waiting for entire

				answer to be ready.
--	--	--	--	---------------------

Request body — `application/json` (required)

Field	Type	Req	Default	Description
<code>audit_metadata</code>	<code>map<string, string> null</code>			A dictionary containing optional audit-specific metadata, such as <code>userid</code> , <code>environment</code> , or other contextual information. This metadata can be leveraged for filtering and analyzing activity logs in futu...
<code>query</code>	<code>string</code>	yes		The query to get a generative answer for
<code>top_k</code>	<code>integer</code>		20	The top most relevant results to fetch at the retrieval step. The maximum number of results allowed is 200.
<code>filter_expression</code>	<code>FilterExpression null</code>			Returns only documents that match this filter expression. Filtering examples can be found here: (see docs) This allows building complex filtering expressions and replaces the following parameters: <code>field...</code>
<code>fields</code>	<code>array<string></code>		[]	The list of fields to search in. For instance: <code>a/title</code> to search only on title field. For more details on filtering by field, see: (see docs)

<code>filters</code>	<code>array<string> array<Filter></code>		<code>[]</code>	The list of filters to apply. Filtering examples can be found here: (see docs)
<code>keyword_filters</code>	<code>array<string> array<Filter></code>		<code>[]</code>	List of keyword filter expressions to apply to the retrieval step. The text block search will only be performed on the documents that contain the specified keywords. The filters are case-insensitive, ...
<code>vectorset</code>	<code>string null</code>			Vectors index to perform the search in. If not provided, NucliaDB will use the default one
<code>min_score</code>	<code>number MinScore null</code>			Minimum score to filter search results. Results with a lower score will be ignored. Accepts either a float or a dictionary with the minimum scores for the bm25 and vector indexes. If a float is provid...
<code>features</code>	<code>array<ChatOptions></code>		<code>['semantic', 'keyword']</code>	Features enabled for the chat endpoint. Semantic search is done if semantic is included. If keyword is included, the results will include matching paragraphs from the bm25 index. If relations is inclu...
<code>range_creation_start</code>	<code>string null</code>			Resources created before this date will be filtered out of search results. Datetime are

				represented as a str in ISO 8601 format, like: 2008-09-15T15:53:00+05:00.
<code>range_creation_end</code>	string null			Resources created after this date will be filtered out of search results. Datetime are represented as a str in ISO 8601 format, like: 2008-09-15T15:53:00+05:00.
<code>range_modification_start</code>	string null			Resources modified before this date will be filtered out of search results. Datetime are represented as a str in ISO 8601 format, like: 2008-09-15T15:53:00+05:00.
<code>range_modification_end</code>	string null			Resources modified after this date will be filtered out of search results. Datetime are represented as a str in ISO 8601 format, like: 2008-09-15T15:53:00+05:00.
<code>show</code>	array<ResourceProperties>		<code>['basic']</code>	Controls which types of metadata are serialized on resources of search results
<code>field_type_filter</code>	array<FieldTypeName>		<code>['text', 'file', 'link', 'conversation', 'generic', 'key_value']</code>	Define which field types are serialized on resources of search results
<code>extracted</code>	array<ExtractedData TypeName>		<code>[]</code>	[Deprecated] Please use GET resource endpoint instead to get extracted

				metadata
context	array<ChatContextMessage> null			DEPRECATED! Please, use chathistory instead.
chat_history	array<ChatContextMessage> null			Use to rephrase the new LLM query by taking into account the chat conversation history. This will be passed to the LLM so that it is aware of the previous conversation.
extra_context	array<string> null			Additional context that is added to the retrieval context sent to the LLM. It allows extending the chat feature with content that may not be in the Knowledge Box.
extra_context_images	array<Image> null			Additional images added to the retrieval context sent to the LLM." It allows extending the chat feature with content that may not be in the Knowledge Box.
query_image	Image null			Image that will be used together with the query text for retrieval and then sent to the LLM as part of the context. If a query image is provided, the extracontextimages and ragimagesstrategies will be...
highlight	boolean		False	If set to true, the query terms will be highlighted in the results between

				<mark>...</mark> tags
resource_filters	array<string>		[]	List of resource ids to filter search results for. Only paragraphs from the specified resources will be returned.
prompt	string CustomPrompt null			Use to customize the prompts given to the generative model. Both system and user prompts can be customized. If a string is provided, it is interpreted as the user prompt.
rank_fusion	RankFusionName ReciprocalRankFusion		rrf	Rank fusion algorithm to use to merge results from multiple retrievers (keyword, semantic)
reranker	RerankerName PredictReranker		predict	Reranker let you specify which method you want to use to rerank your results at the end of retrieval
citations	boolean CitationsType null			Whether to include citations in the response. If set to None or False, no citations will be computed. If set to True or 'default', citations will be computed after answer generation and send as a sepa...
citation_threshold	number null			If citations is set to True or 'default', this will be the similarity threshold. Value between 0 and 1, lower values will produce more

				citations. If not set, it will be set to the optimized threshold ...
security	RequestSecurity null			Security metadata for the request. Please refer to the documentation for more details on how security works: (see docs)
show_hidden	boolean		False	If set to false (default), excludes hidden resources from search
rag_strategies	array<FieldExtensionStrategy FullResourceStrategy HierarchyResourceStrategy NeighbouringParagraphsStrategy MetadataExtensionStrategy ConversationalStrategy PreQueriesStrategy GraphStrategy>		[]	Options for tweaking how the context for the LLM model is crafted: - fullresource will add the full text of the matching resources to the context. This strategy cannot be combined with hierarchy, neig...
rag_images_strategies	array<PageImageStrategy ParagraphImageStrategy TableImageStrategy>		[]	Options for tweaking how the image based context for the LLM model is crafted: - pageimage will add the full page image of the matching resources to the context. - tables will send the table images fo...
debug	boolean		False	If set, the response will include some extra metadata for debugging purposes, like the list of queried

				nodes.
<code>generative_model</code>	string null			The generative model to use for the chat endpoint. If not provided, the model configured for the Knowledge Box is used.
<code>generative_model_seed</code>	integer null			The seed to use for the generative model for deterministic generation. Only supported by some models.
<code>max_tokens</code>	integer MaxTokens null			Use to limit the amount of tokens used in the LLM context and/or for generating the answer. If not provided, the default maximum tokens of the generative model will be used. If an integer is provided,...
<code>rephrase</code>	boolean		<code>False</code>	Rephrase the query for a more efficient retrieval. This will consume LLM tokens and make the request slower.
<code>chat_history_relevance_threshold</code>	number null			Threshold to determine if the past chat history is relevant to rephrase the user's question. 0 - Always treat previous messages as relevant (always rephrase). 1 - Always treat previous messages as irre...
<code>prefer_markdown</code>	boolean		<code>False</code>	If set to true, the response will be in markdown format

<code>answer_json_schema</code>	object (free-form map) null			Desired JSON schema for the LLM answer. This schema is passed to the LLM so that it answers in a structured format following the schema. If not provided, textual response is returned. Note that when u...
<code>generate_answer</code>	boolean		<code>True</code>	Whether to generate an answer using the generative model. If set to false, the response will only contain the retrieval results.
<code>search_configuration</code>	string null			Load ask parameters from this configuration. Parameters in the request override parameters from the configuration.
<code>reasoning</code>	Reasoning boolean		<code>False</code>	Reasoning options for the generative model. Set to True to enable default reasoning, False to disable, or provide a Reasoning object for custom options.

Example — a representative request:

```
BASH
curl -X POST "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/resource/$RID/ask" \
-H "Authorization: Bearer $KB_API_KEY" \
-H "Content-Type: application/json" \
-d '{
  "query": "your question here",
  "top_k": 20,
  "fields": [],
  "features": [
    "semantic",
    "keyword"
  ]
}
```



```

],
"show": [
  "basic"
],
"field_type_filter": [
  "text",
  "file",
  "link",
  "conversation",
  "generic",
  "key_value"
]
}'

```

Response codes

Status	Schema	Description
200	SyncAskResponse	Successful Response
422	HTTPValidationError	Validation Error

GET /kb/{kbid}/resource/{rid}/search

Search on Resource · roles: **READER**

Search on a single resource

Parameters

Name	In	Type	Req	Description
kbid	path	string	yes	
rid	path	string	yes	
query	query	string	yes	
filter_expression	query	string null		Returns only documents that match this filter expression. Filtering examples can be found here: (see docs) This allows building complex filtering expre...
fields	query	array<string>		The list of fields to search in. For

				instance: a/title to search only on title field. For more details on filtering by field, see: (see docs)
<code>filters</code>	query	array<string>		The list of filters to apply. Filtering examples can be found here: (see docs)
<code>faceted</code>	query	array<string>		The list of facets to calculate. The facets follow the same syntax as filters: (see docs)
<code>sort_field</code>	query	SortField null		Field to sort results with (Score not supported in catalog)
<code>sort_order</code>	query	SortOrder		Order to sort results with
<code>top_k</code>	query	integer null		The number of results search should return. The maximum number of results allowed is 200.
<code>range_creation_start</code>	query	string null		Resources created before this date will be filtered out of search results. Datetime are represented as a str in ISO 8601 format, like: 2008-09-15T15:5...
<code>range_creation_end</code>	query	string null		Resources created after this date will be filtered out of search results. Datetime are represented as a str in ISO 8601 format, like: 2008-09-15T15:53...
<code>range_modification_start</code>	query	string null		Resources modified before this date will

				be filtered out of search results. Datetime are represented as a str in ISO 8601 format, like: 2008-09-15T15:...
<code>range_modification_end</code>	query	string null		Resources modified after this date will be filtered out of search results. Datetime are represented as a str in ISO 8601 format, like: 2008-09-15T15:5...
<code>highlight</code>	query	boolean		If set to true, the query terms will be highlighted in the results between <code><mark>...</mark></code> tags
<code>debug</code>	query	boolean		If set, the response will include some extra metadata for debugging purposes, like the list of queried nodes.
<code>x-ndb-client</code>	header	NucliaDBClientType		

Example — a representative request:

```
BASH
curl -X GET "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/resource/$RID/search" \
  -H "Authorization: Bearer $KB_API_KEY"
```

Response codes

Status	Schema	Description
200	<code>ResourceSearchResults</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

GET /kb/{kbid}/search

Search Knowledge Box · roles: READER

Search on a Knowledge Box and retrieve separate results for documents, paragraphs, and sentences.

Usually, it is better to use find

Parameters

Name	In	Type	Req	Description
kbid	path	string	yes	
query	query	string		The query to search for
filter_expression	query	string null		Returns only documents that match this filter expression. Filtering examples can be found here: (see docs) This allows building complex filtering expressions.
fields	query	array<string>		The list of fields to search in. For instance: a/title to search only on title field. For more details on filtering by field, see: (see docs)
filters	query	array<string>		The list of filters to apply. Filtering examples can be found here: (see docs)
faceted	query	array<string>		The list of facets to calculate. The facets follow the same syntax as filters: (see docs)
sort_field	query	SortField		Field to sort results with (Score not supported in catalog)
sort_order	query	SortOrder		Order to sort results with

<code>top_k</code>	query	integer		The number of results search should return. The maximum number of results allowed is 200.
<code>offset</code>	query	integer		The number of results to skip, starting from the beginning in sort order. Used for pagination. It can only be used with the keyword and fulltext index...
<code>min_score</code>	query	number null		Minimum similarity score to filter vector index results. If not specified, the default minimum score of the semantic model associated to the Knowledge...
<code>min_score_semantic</code>	query	number null		Minimum semantic similarity score to filter vector index results. If not specified, the default minimum score of the semantic model associated to the ...
<code>min_score_bm25</code>	query	number		Minimum bm25 score to filter paragraph and document index results
<code>vectorset</code>	query	string null		Vectors index to perform the search in. If not provided, NucliaDB will use the default one
<code>range_creation_start</code>	query	string null		Resources created before this date will be filtered out of search results. Datetime are represented as a str in

				ISO 8601 format, like: 2008-09-15T15:5...
<code>range_creation_end</code>	query	string null		Resources created after this date will be filtered out of search results. Datetime are represented as a str in ISO 8601 format, like: 2008-09-15T15:53...
<code>range_modification_start</code>	query	string null		Resources modified before this date will be filtered out of search results. Datetime are represented as a str in ISO 8601 format, like: 2008-09-15T15:...
<code>range_modification_end</code>	query	string null		Resources modified after this date will be filtered out of search results. Datetime are represented as a str in ISO 8601 format, like: 2008-09-15T15:5...
<code>features</code>	query	array<SearchOptions>		List of search features to use. Each value corresponds to a lookup into one of the different indexes
<code>debug</code>	query	boolean		If set, the response will include some extra metadata for debugging purposes, like the list of queried nodes.
<code>highlight</code>	query	boolean		If set to true, the query terms will be highlighted in the results between <mark>...</mark> tags

show	query	array<ResourceProperties>		Controls which types of metadata are serialized on resources of search results
field_type	query	array<FieldTypeName>		Define which field types are serialized on resources of search results
extracted	query	array<ExtractedData TypeName>		[Deprecated] Please use GET resource endpoint instead to get extracted metadata
with_duplicates	query	boolean		Whether to return duplicate paragraphs on the same document
with_synonyms	query	boolean		Whether to return matches for custom knowledge box synonyms of the query terms. Note: only supported for keyword and fulltext search options.
security_groups	query	array<string> null		List of security groups to filter search results for. Only resources matching the query and containing the specified security groups will be returned....
show_hidden	query	boolean		If set to false (default), excludes hidden resources from search
x-ndb-client	header	NucliaDBClientType		
x-nucladb-user	header	string		
x-forwarded-	header	string		

for				
-----	--	--	--	--

Example — a representative request:

```
BASH
curl -X GET "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/search" \
-H "Authorization: Bearer $KB_API_KEY"
```

Response codes

Status	Schema	Description
200	KnowledgeboxSearchResults	Successful Response
422	HTTPValidationError	Validation Error

POST /kb/{kbid}/search

Search Knowledge Box · roles: READER

Search on a Knowledge Box and retrieve separate results for documents, paragraphs, and sentences. Usually, it is better to use find

Parameters

Name	In	Type	Req	Description
kbid	path	string	yes	
x-ndb-client	header	NucliaDBClientType		
x-nucliadb-user	header	string		
x-forwarded-for	header	string		

Request body — `application/json` (required)

Field	Type	Req	Default	Description
audit_metadata	map<string, string> null			A dictionary containing optional audit-specific metadata, such as

				userid, environment, or other contextual information. This metadata can be leveraged for filtering and analyzing activity logs in futu...
query	string			The query to search for
filter_expression	FilterExpression null			Returns only documents that match this filter expression. Filtering examples can be found here: (see docs) This allows building complex filtering expressions and replaces the following parameters:field...
fields	array<string>		[]	The list of fields to search in. For instance: a/title to search only on title field. For more details on filtering by field, see: (see docs)
filters	array<string> array<Filter>		[]	The list of filters to apply. Filtering examples can be found here: (see docs)
top_k	integer		20	The number of results search should return. The maximum number of results allowed is 200.
min_score	number MinScore null			Minimum score to filter search results. Results with a lower score will be ignored. Accepts either a float or a dictionary with the minimum scores

				for the bm25 and vector indexes. If a float is provid...
<code>range_creation_start</code>	string null			Resources created before this date will be filtered out of search results. Datetime are represented as a str in ISO 8601 format, like: 2008-09-15T15:53:00+05:00.
<code>range_creation_end</code>	string null			Resources created after this date will be filtered out of search results. Datetime are represented as a str in ISO 8601 format, like: 2008-09-15T15:53:00+05:00.
<code>range_modification_start</code>	string null			Resources modified before this date will be filtered out of search results. Datetime are represented as a str in ISO 8601 format, like: 2008-09-15T15:53:00+05:00.
<code>range_modification_end</code>	string null			Resources modified after this date will be filtered out of search results. Datetime are represented as a str in ISO 8601 format, like: 2008-09-15T15:53:00+05:00.
<code>debug</code>	boolean		<code>False</code>	If set, the response will include some extra metadata for debugging purposes, like the list of queried nodes.
<code>highlight</code>	boolean		<code>False</code>	If set to true, the query terms will be

				highlighted in the results between <mark>...</mark> tags
show	array<ResourceProperties>		['basic']	Controls which types of metadata are serialized on resources of search results
field_type_filter	array<FieldTypeName>		['text', 'file', 'link', 'conversation', 'generic', 'key_value']	Define which field types are serialized on resources of search results
extracted	array<ExtractedDataTypeName>		[]	[Deprecated] Please use GET resource endpoint instead to get extracted metadata
vector	array<number> null			The vector to perform the search with. If not provided, NucliaDB will use Nuclia Predict API to create the vector off from the query.
vectorset	string null			Vectors index to perform the search in. If not provided, NucliaDB will use the default one
with_duplicates	boolean		False	Whether to return duplicate paragraphs on the same document
with_synonyms	boolean		False	Whether to return matches for custom knowledge box synonyms of the query terms. Note: only supported for keyword and fulltext search options.

<code>resource_filters</code>	<code>array<string></code>		<code>[]</code>	List of resource ids to filter search results for. Only paragraphs from the specified resources will be returned.
<code>security</code>	<code>RequestSecurity null</code>			Security metadata for the request. Please refer to the documentation for more details on how security works: (see docs)
<code>show_hidden</code>	<code>boolean</code>		<code>False</code>	If set to false (default), excludes hidden resources from search
<code>rephrase</code>	<code>boolean</code>		<code>False</code>	Rephrase the query for a more efficient retrieval. This will consume LLM tokens and make the request slower.
<code>rephrase_prompt</code>	<code>string null</code>			Rephrase prompt given to the generative model responsible for rephrasing the query for a more effective retrieval step. This is only used if the rephrase flag is set to true in the request. If not spe...
<code>query_image</code>	<code>Image null</code>			Image that will be used together with the query text for retrieval.
<code>features</code>	<code>array<SearchOptions></code>		<code>['keyword', 'fulltext', 'semantic']</code>	List of search features to use. Each value corresponds to a lookup into one of the different indexes
<code>faceted</code>	<code>array<string></code>		<code>[]</code>	The list of facets to calculate. The facets

				follow the same syntax as filters: (see docs)
<code>sort</code>	SortOptions null			Options for results sorting
<code>offset</code>	integer		<code>0</code>	The number of results to skip, starting from the beginning in sort order. Used for pagination. It can only be used with the keyword and fulltext indexes.

Example — a real call captured against the live sandbox:

BASH

```
curl -X POST "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/search" \
-H "Authorization: Bearer $KB_API_KEY" \
-H "Content-Type: application/json" \
-d '{
  "query": "battery capacity",
  "features": [
    "keyword",
    "semantic"
  ],
  "top_k": 5
}'
```

Response:

JSON

```
{
  "resources": {
    "f08c5ae15670407a9256059d4ea3b5a4": {
      "id": "f08c5ae15670407a9256059d4ea3b5a4",
      "slug": "ev-charging",
      "title": "EV Home Charging",
      "summary": "",
      "icon": "text/plain",
      "thumbnail": "",
      "metadata": {
        "metadata": {},
        "language": "en",
        "languages": [
          "en"
        ]
      },
      "status": "PROCESSED"
    }
  }
}
```

```

    },
    "usermetadata": {
      "classifications": [
        {
          "labelset": "topic",
          "label": "transport",
          "cancelled_by_user": false
        }
      ]
    },
    "fieldmetadata": [],
    "computedmetadata": {
      "field_classifications": []
    },
    "created": "2026-07-27T05:24:19.052088",
    "modified": "2026-07-27T05:24:19.052103",
    "last_seqid": 0,
    "last_account_seq": null,
    "queue": "private",
    "hidden": false
  ... (truncated)

```

Response codes

Status	Schema	Description
200	<code>KnowledgeboxSearchResults</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

POST /kb/{kbid}/slug/{slug}/ask

Ask a resource (by slug) · roles: READER

Ask questions to a resource

Parameters

Name	In	Type	Req	Description
<code>kbid</code>	path	string	yes	
<code>slug</code>	path	string	yes	
<code>x-show-consumption</code>	header	boolean		
<code>x-ndb-client</code>	header	NucliaDBClientType		
<code>x-nucliadb-</code>	header	string		

<code>user</code>				
<code>x-forwarded-for</code>	header	string		
<code>x-synchronous</code>	header	boolean		When set to true, outputs response as JSON in a non-streaming way. This is slower and requires waiting for entire answer to be ready.

Request body — `application/json` (required)

Field	Type	Req	Default	Description
<code>audit_metadata</code>	<code>map<string, string> null</code>			A dictionary containing optional audit-specific metadata, such as userid, environment, or other contextual information. This metadata can be leveraged for filtering and analyzing activity logs in futu...
<code>query</code>	string	yes		The query to get a generative answer for
<code>top_k</code>	integer		20	The top most relevant results to fetch at the retrieval step. The maximum number of results allowed is 200.
<code>filter_expression</code>	<code>FilterExpression null</code>			Returns only documents that match this filter expression. Filtering examples can be found here: (see docs) This allows building complex filtering expressions and replaces the following

				parameters:field...
fields	array<string>		[]	The list of fields to search in. For instance: a/title to search only on title field. For more details on filtering by field, see: (see docs)
filters	array<string> array<Filter>		[]	The list of filters to apply. Filtering examples can be found here: (see docs)
keyword_filters	array<string> array<Filter>		[]	List of keyword filter expressions to apply to the retrieval step. The text block search will only be performed on the documents that contain the specified keywords. The filters are case-insensitive, ...
vectorset	string null			Vectors index to perform the search in. If not provided, NucliaDB will use the default one
min_score	number MinScore null			Minimum score to filter search results. Results with a lower score will be ignored. Accepts either a float or a dictionary with the minimum scores for the bm25 and vector indexes. If a float is provid...
features	array<ChatOptions>		['semantic', 'keyword']	Features enabled for the chat endpoint. Semantic search is done if semantic is included. If keyword is included, the

				results will include matching paragraphs from the bm25 index. If relations is inclu...
<code>range_creation_start</code>	string null			Resources created before this date will be filtered out of search results. Datetime are represented as a str in ISO 8601 format, like: 2008-09-15T15:53:00+05:00.
<code>range_creation_end</code>	string null			Resources created after this date will be filtered out of search results. Datetime are represented as a str in ISO 8601 format, like: 2008-09-15T15:53:00+05:00.
<code>range_modification_start</code>	string null			Resources modified before this date will be filtered out of search results. Datetime are represented as a str in ISO 8601 format, like: 2008-09-15T15:53:00+05:00.
<code>range_modification_end</code>	string null			Resources modified after this date will be filtered out of search results. Datetime are represented as a str in ISO 8601 format, like: 2008-09-15T15:53:00+05:00.
<code>show</code>	array<ResourceProperties>		<code>['basic']</code>	Controls which types of metadata are serialized on resources of search results
<code>field_type_filter</code>	array<FieldTypeName>		<code>['text', 'file',</code>	Define which field types are serialized

			'link', 'conversation', 'generic', 'key_value']	on resources of search results
extracted	array<ExtractedData TypeName>		[]	[Deprecated] Please use GET resource endpoint instead to get extracted metadata
context	array<ChatContextMessage> null			DEPRECATED! Please, use chathistory instead.
chat_history	array<ChatContextMessage> null			Use to rephrase the new LLM query by taking into account the chat conversation history. This will be passed to the LLM so that it is aware of the previous conversation.
extra_context	array<string> null			Additional context that is added to the retrieval context sent to the LLM. It allows extending the chat feature with content that may not be in the Knowledge Box.
extra_context_images	array<Image> null			Additional images added to the retrieval context sent to the LLM." It allows extending the chat feature with content that may not be in the Knowledge Box.
query_image	Image null			Image that will be used together with the query text for retrieval and then sent to the LLM as part of the context. If a query image is provided, the

				extracontextimages and ragimagesstrategies will be...
<code>highlight</code>	boolean		<code>False</code>	If set to true, the query terms will be highlighted in the results between <code><mark>...</mark></code> tags
<code>resource_filters</code>	array<string>		<code>[]</code>	List of resource ids to filter search results for. Only paragraphs from the specified resources will be returned.
<code>prompt</code>	string CustomPrompt null			Use to customize the prompts given to the generative model. Both system and user prompts can be customized. If a string is provided, it is interpreted as the user prompt.
<code>rank_fusion</code>	RankFusionName ReciprocalRankFusion		<code>rrf</code>	Rank fusion algorithm to use to merge results from multiple retrievers (keyword, semantic)
<code>reranker</code>	RerankerName PredictReranker		<code>predict</code>	Reranker let you specify which method you want to use to rerank your results at the end of retrieval
<code>citations</code>	boolean CitationsType null			Whether to include citations in the response. If set to None or False, no citations will be computed. If set to True or 'default', citations will be computed after answer generation

				and send as a sepa...
<code>citation_threshold</code>	number null			If citations is set to True or 'default', this will be the similarity threshold. Value between 0 and 1, lower values will produce more citations. If not set, it will be set to the optimized threshold ...
<code>security</code>	RequestSecurity null			Security metadata for the request. Please refer to the documentation for more details on how security works: (see docs)
<code>show_hidden</code>	boolean		<code>False</code>	If set to false (default), excludes hidden resources from search
<code>rag_strategies</code>	array<FieldExtensionStrategy FullResourceStrategy HierarchyResourceStrategy NeighbouringParagraphsStrategy MetadataExtensionStrategy ConversationalStrategy PreQueriesStrategy GraphStrategy>		<code>[]</code>	Options for tweaking how the context for the LLM model is crafted: - fullresource will add the full text of the matching resources to the context. This strategy cannot be combined with hierarchy, neig...
<code>rag_images_strategies</code>	array<PageImageStrategy ParagraphImageStrategy TableImageStrategy>		<code>[]</code>	Options for tweaking how the image based context for the LLM model is crafted: - pageimage will add the full page image of the matching resources to the context. - tables will

				send the table images fo...
<code>debug</code>	boolean		<code>False</code>	If set, the response will include some extra metadata for debugging purposes, like the list of queried nodes.
<code>generative_model</code>	string null			The generative model to use for the chat endpoint. If not provided, the model configured for the Knowledge Box is used.
<code>generative_model_seed</code>	integer null			The seed to use for the generative model for deterministic generation. Only supported by some models.
<code>max_tokens</code>	integer MaxTokens null			Use to limit the amount of tokens used in the LLM context and/or for generating the answer. If not provided, the default maximum tokens of the generative model will be used. If an integer is provided,...
<code>rephrase</code>	boolean		<code>False</code>	Rephrase the query for a more efficient retrieval. This will consume LLM tokens and make the request slower.
<code>chat_history_relevance_threshold</code>	number null			Threshold to determine if the past chat history is relevant to rephrase the user's question. 0 - Always treat previous messages as

				relevant (always rephrase).1 - Always treat previous messages as irre...
<code>prefer_markdown</code>	boolean		<code>False</code>	If set to true, the response will be in markdown format
<code>answer_json_schema</code>	object (free-form map) null			Desired JSON schema for the LLM answer. This schema is passed to the LLM so that it answers in a structured format following the schema. If not provided, textual response is returned. Note that when u...
<code>generate_answer</code>	boolean		<code>True</code>	Whether to generate an answer using the generative model. If set to false, the response will only contain the retrieval results.
<code>search_configuration</code>	string null			Load ask parameters from this configuration. Parameters in the request override parameters from the configuration.
<code>reasoning</code>	Reasoning boolean		<code>False</code>	Reasoning options for the generative model. Set to True to enable default reasoning, False to disable, or provide a Reasoning object for custom options.

Example — a representative request:

```
BASH
```

```
curl -X POST "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/slug/{slug}/ask" \
-H "Authorization: Bearer $KB_API_KEY" \
-H "Content-Type: application/json" \
-d '{
  "query": "your question here",
  "top_k": 20,
  "fields": [],
  "features": [
    "semantic",
    "keyword"
  ],
  "show": [
    "basic"
  ],
  "field_type_filter": [
    "text",
    "file",
    "link",
    "conversation",
    "generic",
    "key_value"
  ]
}'
```

Response codes

Status	Schema	Description
200	SyncAskResponse	Successful Response
422	HTTPValidationError	Validation Error

GET /kb/{kbid}/suggest

Suggest on a knowledge box · roles: **READER**

Suggestions on a knowledge box

Parameters

Name	In	Type	Req	Description
kbid	path	string	yes	
query	query	string	yes	The query to get suggestions for
fields	query	array<string>		The list of fields to search in. For instance: a/title to search only on title

				field. For more details on filtering by field, see: (see docs)
<code>filters</code>	query	array<string>		The list of filters to apply. Filtering examples can be found here: (see docs)
<code>range_creation_start</code>	query	string null		Resources created before this date will be filtered out of search results. Datetime are represented as a str in ISO 8601 format, like: 2008-09-15T15:5...
<code>range_creation_end</code>	query	string null		Resources created after this date will be filtered out of search results. Datetime are represented as a str in ISO 8601 format, like: 2008-09-15T15:53...
<code>range_modification_start</code>	query	string null		Resources modified before this date will be filtered out of search results. Datetime are represented as a str in ISO 8601 format, like: 2008-09-15T15:...
<code>range_modification_end</code>	query	string null		Resources modified after this date will be filtered out of search results. Datetime are represented as a str in ISO 8601 format, like: 2008-09-15T15:5...
<code>features</code>	query	array<SuggestOptions>		Features enabled for the suggest endpoint.

<code>debug</code>	query	boolean		If set, the response will include some extra metadata for debugging purposes, like the list of queried nodes.
<code>highlight</code>	query	boolean		If set to true, the query terms will be highlighted in the results between <code><mark>...</mark></code> tags
<code>show_hidden</code>	query	boolean		If set to false (default), excludes hidden resources from search
<code>security_groups</code>	query	array<string> null		List of security groups to filter search results for. Only resources matching the query and containing the specified security groups will be returned....

Example — a real call captured against the live sandbox:

BASH

```
curl -X GET "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/suggest" \
-H "Authorization: Bearer $KB_API_KEY"
```

Response:

JSON

```
{
  "paragraphs": {
    "results": [
      {
        "score": 0.5418849587440491,
        "rid": "f08c5ae15670407a9256059d4ea3b5a4",
        "field_type": "t",
        "field": "body",
        "text": "An electric vehicle can be charged at home using a 7kW single-phase wallbox or a 22kW three-phase charger. Charging overnight from rooftop solar and battery reduces running costs dramatically. Smart chargers support
```

```

scheduled charging aligned to cheap tariff windows and solar surplus.",
  "labels": [
    "topic/transport"
  ],
  "start_seconds": null,
  "end_seconds": null,
  "position": {
    "page_number": 0,
    "index": 0,
    "start": 0,
    "end": 285
  }
},
{
  "score": 0.5418849587440491,
  "rid": "86b5f7948b144b008f018cbd47a2bde4",
  "field_type": "t",
  "field": "body",
  "
... (truncated)

```

Response codes

Status	Schema	Description
200	<code>KnowledgeboxSuggestResults</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

POST /kb/{kbid}/suggest

Suggest on a knowledge box · roles: `READER`

Suggestions on a knowledge box

Parameters

Name	In	Type	Req	Description
<code>kbid</code>	path	string	yes	

Request body — `application/json` (required)

Field	Type	Req	Default	Description
<code>query</code>	string	yes		The query to get suggestions for

features	array<SuggestOptions>		['paragraph', 'entities']	Features enabled for the suggest endpoint.
filter_expression	FilterExpression null			Returns only documents that match this filter expression. Filtering examples can be found here: (see docs) This allows building complex filtering expressions and replaces the following parameters:field...
security	RequestSecurity null			Security metadata for the request. Please refer to the documentation for more details on how security works: (see docs)
show_hidden	boolean		False	If set to false (default), excludes hidden resources from search
highlight	boolean		False	If set to true, the query terms will be highlighted in the results between <mark>...</mark> tags

Example — a representative request:

BASH

```
curl -X POST "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/suggest" \
-H "Authorization: Bearer $KB_API_KEY" \
-H "Content-Type: application/json" \
-d '{
  "query": "your question here",
  "features": [
    "paragraph",
    "entities"
  ],
  "show_hidden": false,
  "highlight": false
}
```

```
}'
```

Response codes

Status	Schema	Description
200	KnowledgeboxSuggestResults	Successful Response
422	HTTPValidationError	Validation Error

POST /kb/{kbid}/summarize

Summarize your documents · roles: READER

Parameters

Name	In	Type	Req	Description
kbid	path	string	yes	
x-show-consumption	header	boolean		

Request body — application/json (required)

Field	Type	Req	Default	Description
generative_model	string null			The generative model to use for the summarization. If not provided, the model configured for the Knowledge Box is used.
user_prompt	string null			Optional custom prompt input by the user
resources	array<string>	yes		Uids or slugs of the resources to summarize. If the resources are not found, they will be ignored.

<code>summary_kind</code>	SummaryKind		<code>simple</code>	Option to customize how the summary will be
---------------------------	-------------	--	---------------------	---

Example — a real call captured against the live sandbox:

BASH

```
curl -X POST "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/summarize" \
-H "Authorization: Bearer $KB_API_KEY" \
-H "Content-Type: application/json" \
-d '{
  "resources": [
    "ca2a8f3bf186491facde30cdb2fbc2b6",
    "86b5f7948b144b008f018cbd47a2bde4"
  ]
}'
```

Response:

JSON

```
{
  "resources": {
    "ca2a8f3bf186491facde30cdb2fbc2b6": {
      "summary": "Photovoltaic solar panels convert sunlight into direct current (DC) electricity, which is then converted to alternating current (AC) by inverters. In Australia, residential rooftop systems typically range from 5kW to 10kW. Certified installers by the Clean Energy Council ensure proper installation. Payback periods range from 4 to 7 years, influenced by feed-in tariffs.",
      "tokens": 104
    },
    "86b5f7948b144b008f018cbd47a2bde4": {
      "summary": "Home battery systems, like lithium-iron-phosphate (LFP) batteries, store excess solar energy for nighttime use, offering 10kWh capacity and about 6000 cycles. They can provide blackout protection with a compatible hybrid inverter. Starting July 2025, the federal Cheaper Home Batteries program will subsidize installations.",
      "tokens": 109
    }
  },
  "
  ... (truncated)
```

Response codes

Status	Schema	Description
200	<code>SummarizedResponse</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

--	--	--

Models (6)

GET /kb/{kbid}/configuration

Get Knowledge Box models configuration · roles: READER, MANAGER

Current configuration of models assigned to a Knowledge Box

Parameters

Name	In	Type	Req	Description
kbid	path	string	yes	

Example — a real call captured against the live sandbox:

```
BASH
curl -X GET "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/configuration" \
-H "Authorization: Bearer $KB_API_KEY"
```

Response:

```
JSON
{
  "resource_labelers_models": null,
  "paragraph_labelers_models": null,
  "intent_models": null,
  "visual_labeling": "disabled",
  "ner_model": "multilingual",
  "relation_model": "base",
  "anonymization_model": "disabled",
  "semantic_model": "multilingual-2024-05-06",
  "semantic_models": [
    "multilingual-2024-05-06"
  ],
  "default_semantic_model": "multilingual-2024-05-06",
  "semantic_graph_node_models": [],
  "default_semantic_graph_node_model": null,
  "semantic_graph_edge_models": [],
  "default_semantic_graph_edge_model": null,
  "semantic_vector_similarity": "DOT",
  "semantic_vector_size": 1024,
  "semantic_matryoshka_dims": [],
  "semantic_threshold": 0.4,
```

```
"generative_model": "chatgpt-azure-4o",
"user_keys": null,
"user_prompts": null,
"summary": "simple",
"summary_model": "chatgpt-azure-4o",
"summary_prompt": null,
"prefer_markdown_generative_response":
... (truncated)
```

Response codes

Status	Schema	Description
200	<code>object</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

GET /kb/{kbid}/generative_providers

Available models for a knowledge box · roles: READER, MANAGER

Get all available models for a knowledge box grouped by provider

Parameters

Name	In	Type	Req	Description
<code>kbid</code>	path	string	yes	

Example — a representative request:

```
BASH
curl -X GET "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/generative_providers" \
-H "Authorization: Bearer $KB_API_KEY"
```

Response codes

Status	Schema	Description
200	<code>object</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

GET /kb/{kbid}/model/{model_id}

Get model metadata · roles: READER, MANAGER

Get metadata for a particular model

Parameters

Name	In	Type	Req	Description
kbid	path	string	yes	
model_id	path	string	yes	

Example — a representative request:

```
BASH
curl -X GET "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/model/{model_id}" \
-H "Authorization: Bearer $KB_API_KEY"
```

Response codes

Status	Schema	Description
200	object	Successful Response
422	HTTPValidationError	Validation Error

GET /kb/{kbid}/models

Get available models · roles: READER, MANAGER

Parameters

Name	In	Type	Req	Description
kbid	path	string	yes	

Example — a real call captured against the live sandbox:

```
BASH
curl -X GET "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/models" \
-H "Authorization: Bearer $KB_API_KEY"
```

Response:

JSON

```
{
  "available": {},
  "trained": {}
}
```

Response codes

Status	Schema	Description
200	<code>object</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

GET /kb/{kbid}/models/{model_id}/{filename}

Download the Knowledge Box model · roles: READER, MANAGER

Download the trained model or any other generated file as a result of a training task on a Knowledge Box.

Parameters

Name	In	Type	Req	Description
<code>kbid</code>	path	string	yes	
<code>model_id</code>	path	string	yes	
<code>filename</code>	path	string	yes	

Example — a representative request:

BASH

```
curl -X GET
"https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/models/{model_id}/{filename}" \
-H "Authorization: Bearer $KB_API_KEY"
```

Response codes

Status	Schema	Description
200	<code>object</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

GET /kb/{kbid}/schema

Learning configuration schema · roles: READER, MANAGER

Get jsonschema definition to update the learningconfiguration of your Knowledge Box

Parameters

Name	In	Type	Req	Description
kbid	path	string	yes	

Example — a representative request:

```
BASH
curl -X GET "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/schema" \
-H "Authorization: Bearer $KB_API_KEY"
```

Response codes

Status	Schema	Description
200	object	Successful Response
422	HTTPValidationError	Validation Error

Knowledge Box Services (21)

GET /kb/{kbid}/custom-synonyms

Get Knowledge Box Custom Synonyms · roles: READER

Parameters

Name	In	Type	Req	Description
kbid	path	string	yes	

Example — a real call captured against the live sandbox:

```
BASH
curl -X GET "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/custom-synonyms" \
-H "Authorization: Bearer $KB_API_KEY"
```

Response:

```
JSON
{
  "synonyms": {
    "ev": [
      "electric vehicle"
    ],
    "pv": [
      "photovoltaic",
      "solar"
    ]
  }
}
```

Response codes

Status	Schema	Description
200	KnowledgeBoxSynonyms	Successful Response
422	HTTPValidationError	Validation Error

PUT /kb/{kbid}/custom-synonyms

Set Knowledge Box Custom Synonyms · roles: WRITER

Parameters

Name	In	Type	Req	Description
kbid	path	string	yes	

Request body — `application/json` (required)

Field	Type	Req	Default	Description
synonyms	map<string, array<string>>	yes		

Example — a real call captured against the live sandbox:

```
BASH
curl -X PUT "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/custom-synonyms" \
-H "Authorization: Bearer $KB_API_KEY" \
-H "Content-Type: application/json" \
-d '{
```

```

    "synonyms": {
      "pv": [
        "photovoltaic",
        "solar"
      ],
      "ev": [
        "electric vehicle"
      ]
    }
  }
}'

```

Response codes

Status	Schema	Description
204		Successful Response
422	HTTPValidationError	Validation Error

DELETE /kb/{kbid}/custom-synonyms

Delete Knowledge Box Custom Synonyms · roles: WRITER

Parameters

Name	In	Type	Req	Description
kbid	path	string	yes	

Example — a representative request:

```

BASH
curl -X DELETE "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/custom-synonyms" \
  -H "Authorization: Bearer $KB_API_KEY"

```

Response codes

Status	Schema	Description
204		Successful Response
422	HTTPValidationError	Validation Error

GET /kb/{kbid}/entitiesgroup/{group}

Get a Knowledge Box Entities Group · roles: READER

Parameters

Name	In	Type	Req	Description
<code>kbid</code>	path	string	yes	
<code>group</code>	path	string	yes	

Example — a representative request:

```
BASH
curl -X GET "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/entitiesgroup/{group}" \
  -H "Authorization: Bearer $KB_API_KEY"
```

Response codes

Status	Schema	Description
200	<code>EntitiesGroup</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

GET /kb/{kbid}/entitiesgroups

Get Knowledge Box Entities · roles: READER

Parameters

Name	In	Type	Req	Description
<code>kbid</code>	path	string	yes	
<code>show_entities</code>	query	boolean		

Example — a real call captured against the live sandbox:

```
BASH
curl -X GET "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/entitiesgroups" \
  -H "Authorization: Bearer $KB_API_KEY"
```

Response:

```
JSON
{
```

```

"uuid": "$KB",
"groups": {
  "ORG": {
    "title": "",
    "color": "",
    "custom": false,
    "entities": {}
  },
  "GPE": {
    "title": "",
    "color": "",
    "custom": false,
    "entities": {}
  },
  "DATE": {
    "title": "",
    "color": "",
    "custom": false,
    "entities": {}
  },
  "TIME": {
    "title": "",
    "color": "",
    "custom": false,
    "entities": {}
  }
}
}

```

Response codes

Status	Schema	Description
200	KnowledgeBoxEntities	Successful Response
422	HTTPValidationError	Validation Error

GET /kb/{kbid}/kv-schemas

List KV schemas · roles: **READER**

Parameters

Name	In	Type	Req	Description
kbid	path	string	yes	

Example — a representative request:

BASH

```
curl -X GET "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/kv-schemas" \
-H "Authorization: Bearer $KB_API_KEY"
```

Response codes

Status	Schema	Description
200	KBKVSSchemas	Successful Response
422	HTTPValidationError	Validation Error

POST /kb/{kbid}/kv-schemas

Create KV schema · roles: WRITER

Parameters

Name	In	Type	Req	Description
kbid	path	string	yes	

Request body — [application/json](#) (required)

Field	Type	Req	Default	Description
id	string	yes		
description	string			
fields	array<KVSchemaField>			

Example — a representative request:

BASH

```
curl -X POST "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/kv-schemas" \
-H "Authorization: Bearer $KB_API_KEY" \
-H "Content-Type: application/json" \
-d '{
  "id": "...",
  "description": ""
}'
```

Response codes

Status	Schema	Description
201	KVSchema	Successful Response
422	HTTPValidationError	Validation Error

GET /kb/{kbid}/kv-schemas/{schema_id}

Get a KV schema · roles: READER

Parameters

Name	In	Type	Req	Description
kbid	path	string	yes	
schema_id	path	string	yes	

Example — a representative request:

```
BASH
curl -X GET
"https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/kv-schemas/{schema_id}" \
-H "Authorization: Bearer $KB_API_KEY"
```

Response codes

Status	Schema	Description
200	KVSchema	Successful Response
422	HTTPValidationError	Validation Error

PUT /kb/{kbid}/kv-schemas/{schema_id}

Update KV schema · roles: WRITER

Parameters

Name	In	Type	Req	Description
kbid	path	string	yes	
schema_id	path	string	yes	

Request body — `application/json` (required)

Field	Type	Req	Default	Description
<code>description</code>	string null			
<code>fields</code>	array<KVSchemaField> null			

Example — a representative request:

```
BASH
curl -X PUT
"https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/kv-schemas/{schema_id}" \
-H "Authorization: Bearer $KB_API_KEY" \
-H "Content-Type: application/json" \
-d '{}'
```

Response codes

Status	Schema	Description
200	<code>KVSchema</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

DELETE /kb/{kbid}/kv-schemas/{schema_id}

Delete KV schema · roles: WRITER

Parameters

Name	In	Type	Req	Description
<code>kbid</code>	path	string	yes	
<code>schema_id</code>	path	string	yes	

Example — a representative request:

```
BASH
curl -X DELETE
"https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/kv-schemas/{schema_id}" \
-H "Authorization: Bearer $KB_API_KEY"
```

Response codes

Status	Schema	Description
--------	--------	-------------

204		Successful Response
422	HTTPValidationError	Validation Error

GET /kb/{kbid}/labelset/{labelset}

Get a Knowledge Box Label Set · roles: READER

Parameters

Name	In	Type	Req	Description
kbid	path	string	yes	
labelset	path	string	yes	

Example — a representative request:

```
BASH
curl -X GET "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/labelset/{labelset}" \
-H "Authorization: Bearer $KB_API_KEY"
```

Response codes

Status	Schema	Description
200	LabelSet	Successful Response
404		Knowledge Box or Label Set not found
422	HTTPValidationError	Validation Error

POST /kb/{kbid}/labelset/{labelset}

Set Knowledge Box Labels · roles: WRITER

Parameters

Name	In	Type	Req	Description
kbid	path	string	yes	
labelset	path	string	yes	

Request body — `application/json` (required)

Field	Type	Req	Default	Description
<code>title</code>	string null			Title of the labelset. It is a prettier display name for the labelset shown in the UI but it is not intended to be used for searching.
<code>color</code>	string null		<code>blue</code>	
<code>multiple</code>	boolean		<code>True</code>	
<code>kind</code>	array<LabelSetKind>		<code>[]</code>	
<code>labels</code>	array<nucliadb_model_labels__Label>			List of labels in the labelset. The titles of the labels must be unique within the labelset.

Example — a representative request:

```
BASH
curl -X POST "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/labelset/{labelset}" \
  -H "Authorization: Bearer $KB_API_KEY" \
  -H "Content-Type: application/json" \
  -d '{
    "multiple": true,
    "kind": []
  }'
```

Response codes

Status	Schema	Description
200	<code>object</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

DELETE `/kb/{kbid}/labelset/{labelset}`

Delete Knowledge Box Label · roles: WRITER

Parameters

Name	In	Type	Req	Description
<code>kbid</code>	path	string	yes	
<code>labelset</code>	path	string	yes	

Example — a representative request:

```
BASH
curl -X DELETE
"https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/labelset/{labelset}" \
-H "Authorization: Bearer $KB_API_KEY"
```

Response codes

Status	Schema	Description
200	<code>object</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

GET /kb/{kbid}/labelsets

Get Knowledge Box Label Sets · roles: READER

Parameters

Name	In	Type	Req	Description
<code>kbid</code>	path	string	yes	

Example — a real call captured against the live sandbox:

```
BASH
curl -X GET "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/labelsets" \
-H "Authorization: Bearer $KB_API_KEY"
```

Response:

```
JSON
{
  "uuid": "$KB",
  "labelsets": {
    "topic": {
      "title": "Topic",
```

```

    "color": "#ff0000",
    "multiple": true,
    "kind": [
      "RESOURCES"
    ],
    "labels": [
      {
        "title": "energy",
        "related": "",
        "text": "",
        "uri": ""
      },
      {
        "title": "transport",
        "related": "",
        "text": "",
        "uri": ""
      }
    ]
  }
}
}

```

Response codes

Status	Schema	Description
200	KnowledgeBoxLabels	Successful Response
422	HTTPValidationError	Validation Error

GET /kb/{kbid}/notifications

Knowledge Box Notifications Stream · roles: **READER**

Provides a stream of activity notifications for the given Knowledge Box. The stream will be automatically closed after 2 minutes.

Parameters

Name	In	Type	Req	Description
kbid	path	string	yes	

Example — a representative request:

```

BASH
curl -X GET "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/notifications" \

```

```
-H "Authorization: Bearer $KB_API_KEY"
```

Response codes

Status	Schema	Description
200	<code>object</code>	Each line of the response is a Base64-encoded JSON object representing a notification. Refer to [the internal documentat...
404		Knowledge Box not found
422	<code>HTTPValidationError</code>	Validation Error

GET /kb/{kbid}/processing-status

Knowledge Box Processing Status · roles: READER

Provides the status of the processing of the given Knowledge Box.

Parameters

Name	In	Type	Req	Description
<code>kbid</code>	path	string	yes	
<code>cursor</code>	query	string null		
<code>scheduled</code>	query	boolean null		
<code>limit</code>	query	integer		

Example — a representative request:

```
BASH
curl -X GET "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/processing-status" \
-H "Authorization: Bearer $KB_API_KEY"
```

Response codes

Status	Schema	Description
200	<code>RequestsResults</code>	Successful Response
404		Knowledge Box not found

422	HTTPValidationError	Validation Error
-----	---------------------	------------------

GET /kb/{kbid}/search_configurations

List search configurations · roles: READER

Parameters

Name	In	Type	Req	Description
kbid	path	string	yes	

Example — a representative request:

```
BASH
curl -X GET "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/search_configurations" \
  -H "Authorization: Bearer $KB_API_KEY"
```

Response codes

Status	Schema	Description
200	map<string, FindSearchConfiguration AskSearchConfiguration>	Successful Response
422	HTTPValidationError	Validation Error

GET /kb/{kbid}/search_configurations/{config_name}

Get search configuration · roles: READER

Parameters

Name	In	Type	Req	Description
kbid	path	string	yes	
config_name	path	string	yes	

Example — a representative request:

```
BASH
```

```
curl -X GET
"https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/search_configurations/
{config_name}" \
-H "Authorization: Bearer $KB_API_KEY"
```

Response codes

Status	Schema	Description
200	FindSearchConfiguration AskSearchConfiguration	Successful Response
422	HTTPValidationError	Validation Error

POST /kb/{kbid}/search_configurations/{config_name}

Create search configuration · roles: OWNER

Parameters

Name	In	Type	Req	Description
kbid	path	string	yes	
config_name	path	string	yes	

Example — a representative request:

```
BASH
curl -X POST
"https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/search_configurations/
{config_name}" \
-H "Authorization: Bearer $KB_API_KEY" \
-H "Content-Type: application/json" \
-d '{
  "kind": "string",
  "config": {
    "query": "",
    "fields": [],
    "top_k": 20,
    "debug": false,
    "highlight": false,
    "show": [
      "basic"
    ]
  }
}'
```


Response codes

Status	Schema	Description
201	<code>object</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

PATCH /kb/{kbid}/search_configurations/{config_name}

Update search configuration · roles: OWNER

Parameters

Name	In	Type	Req	Description
<code>kbid</code>	path	string	yes	
<code>config_name</code>	path	string	yes	

Example — a representative request:

```
BASH
curl -X PATCH
"https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/search_configurations/
{config_name}" \
-H "Authorization: Bearer $KB_API_KEY" \
-H "Content-Type: application/json" \
-d '{
  "kind": "string",
  "config": {
    "query": "",
    "fields": [],
    "top_k": 20,
    "debug": false,
    "highlight": false,
    "show": [
      "basic"
    ]
  }
}'
```

Response codes

Status	Schema	Description
200	<code>object</code>	Successful Response

422	HTTPValidationError	Validation Error
-----	---------------------	------------------

DELETE /kb/{kbid}/search_configurations/{config_name}

Delete search configuration · roles: OWNER

Parameters

Name	In	Type	Req	Description
kbid	path	string	yes	
config_name	path	string	yes	

Example — a representative request:

```
BASH
curl -X DELETE
"https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/search_configurations/
{config_name}" \
-H "Authorization: Bearer $KB_API_KEY"
```

Response codes

Status	Schema	Description
204		Successful Response
422	HTTPValidationError	Validation Error

Extract Strategies (4)

GET /kb/{kbid}/extract_strategies

Learning extract strategies · roles: READER, MANAGER

Get available extract strategies

Parameters

Name	In	Type	Req	Description
kbid	path	string	yes	

Example — a real call captured against the live sandbox:

BASH

```
curl -X GET "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/extract_strategies" \
-H "Authorization: Bearer $KB_API_KEY"
```

Response:

JSON

```
{}
```

Response codes

Status	Schema	Description
200	<code>map<string, ExtractConfig></code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

POST /kb/{kbid}/extract_strategies

Add a extract strategy to a KB · roles: MANAGER, OWNER

Parameters

Name	In	Type	Req	Description
<code>kbid</code>	path	string	yes	

Request body — `application/json` (required)

Field	Type	Req	Default	Description
<code>name</code>	string			
<code>vllm_config</code>	VLLMExtractionConfig null			
<code>ai_tables</code>	AITables null			
<code>split</code>	SplitConfig null			

Example — a representative request:

BASH

```
curl -X POST "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/extract_strategies" \
-H "Authorization: Bearer $KB_API_KEY" \
-H "Content-Type: application/json" \
-d '{
  "name": ""
}'
```

Response codes

Status	Schema	Description
200	<code>string</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

GET /kb/{kbid}/extract_strategies/strategy/{strategy_id}

Extract strategy configuration · roles: READER, MANAGER

Get extract strategy for a given id

Parameters

Name	In	Type	Req	Description
<code>kbid</code>	path	string	yes	
<code>strategy_id</code>	path	string	yes	

Example — a representative request:

BASH

```
curl -X GET "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/extract_strategies/
strategy/{strategy_id}" \
-H "Authorization: Bearer $KB_API_KEY"
```

Response codes

Status	Schema	Description
200	<code>object</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

DELETE /kb/{kbid}/extract_strategies/strategy/{strategy_id}

Remove a extract strategy from a KB · roles: MANAGER, OWNER

Removes a extract strategy from a KB

Parameters

Name	In	Type	Req	Description
kbid	path	string	yes	
strategy_id	path	string	yes	

Example — a representative request:

```
BASH
curl -X DELETE
"https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/extract_strategies/strategy/
{strategy_id}" \
-H "Authorization: Bearer $KB_API_KEY"
```

Response codes

Status	Schema	Description
204		Successful Response
422	HTTPValidationError	Validation Error

Resource field TUS uploads (6)

POST /kb/{kbid}/resource/{path_rid}/file/{field}/tusupload

Create new upload on a Resource (by id) · roles: WRITER

Parameters

Name	In	Type	Req	Description
kbid	path	string	yes	
path_rid	path	string	yes	
field	path	string	yes	

x-extract-strategy	header	string null		Extract strategy to use when uploading a file. If not provided, the default strategy will be used.
x-split-strategy	header	string null		Split strategy to use when uploading a file. If not provided, the default strategy will be used.

Example — a representative request:

```
BASH
curl -X POST "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/resource/{path_rid}/file/{field}/tusupload" \
  -H "Authorization: Bearer $KB_API_KEY" \
  -H "Content-Type: application/json" \
  -d '{}'
```

Response codes

Status	Schema	Description
200	object	Successful Response
422	HTTPValidationError	Validation Error

HEAD /kb/{kbid}/resource/{path_rid}/file/{field}/tusupload/{upload_id}

Upload information · roles: WRITER

Parameters

Name	In	Type	Req	Description
kbid	path	string	yes	
path_rid	path	string	yes	
field	path	string	yes	
upload_id	path	string	yes	

Example — a representative request:

BASH

```
curl -X HEAD "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/resource/{path_rid}/file/{field}/tusupload/{upload_id}" \
-H "Authorization: Bearer $KB_API_KEY"
```

Response codes

Status	Schema	Description
200	<code>object</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

PATCH /kb/{kbid}/resource/{rid}/file/{field}/tusupload/{upload_id}

Upload data on a Resource (by id) · roles: WRITER

Parameters

Name	In	Type	Req	Description
<code>kbid</code>	path	string	yes	
<code>rid</code>	path	string	yes	
<code>field</code>	path	string	yes	
<code>upload_id</code>	path	string	yes	

Example — a representative request:

BASH

```
curl -X PATCH "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/resource/$RID/file/{field}/tusupload/{upload_id}" \
-H "Authorization: Bearer $KB_API_KEY"
```

Response codes

Status	Schema	Description
200	<code>object</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

POST /kb/{kbid}/slug/{rslug}/file/{field}/tusupload

Create new upload on a Resource (by slug) · roles: WRITER

Parameters

Name	In	Type	Req	Description
kbid	path	string	yes	
rslug	path	string	yes	
field	path	string	yes	
x-extract-strategy	header	string null		Extract strategy to use when uploading a file. If not provided, the default strategy will be used.
x-split-strategy	header	string null		Split strategy to use when uploading a file. If not provided, the default strategy will be used.

Example — a representative request:

```
BASH
curl -X POST
"https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/slug/{rslug}/file/{field}/
tusupload" \
-H "Authorization: Bearer $KB_API_KEY" \
-H "Content-Type: application/json" \
-d '{}'
```

Response codes

Status	Schema	Description
200	object	Successful Response
422	HTTPValidationError	Validation Error

PATCH /kb/{kbid}/slug/{rslug}/file/{field}/tusupload/{upload_id}

Upload data on a Resource (by slug) · roles: WRITER

Parameters

Name	In	Type	Req	Description
<code>kbid</code>	path	string	yes	
<code>rslug</code>	path	string	yes	
<code>field</code>	path	string	yes	
<code>upload_id</code>	path	string	yes	

Example — a representative request:

```
BASH
curl -X PATCH
"https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/slug/{rslug}/file/{field}/
tusupload/{upload_id}" \
-H "Authorization: Bearer $KB_API_KEY"
```

Response codes

Status	Schema	Description
200	<code>object</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

HEAD /kb/{kbid}/slug/{rslug}/file/{field}/tusupload/{upload_id}

Upload information · roles: WRITER

Parameters

Name	In	Type	Req	Description
<code>kbid</code>	path	string	yes	
<code>rslug</code>	path	string	yes	
<code>field</code>	path	string	yes	
<code>upload_id</code>	path	string	yes	

Example — a representative request:

```
BASH
curl -X HEAD
"https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/slug/{rslug}/file/{field}/"
```

```
tusupload/{upload_id}" \
-H "Authorization: Bearer $KB_API_KEY"
```

Response codes

Status	Schema	Description
200	<code>object</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

Resource fields (25)

POST /kb/{kbid}/resource/{path_rid}/file/{field}/upload

Upload binary file on a Resource (by id) · roles: WRITER

Upload a file as a field on an existing resource, if the field exists will return a conflict (419)

Parameters

Name	In	Type	Req	Description
<code>kbid</code>	path	string	yes	
<code>path_rid</code>	path	string	yes	
<code>field</code>	path	string	yes	
<code>x-filename</code>	header	string null		Name of the file being uploaded.
<code>x-password</code>	header	string null		If the file is password protected, the password must be provided here.
<code>x-language</code>	header	string null		
<code>x-md5</code>	header	string null		MD5 hash of the file being uploaded. This is used to check if the file has been uploaded before.
<code>x-extract-strategy</code>	header	string null		Extract strategy to use when uploading a file. If not provided,

				the default strategy will be used.
x-split-strategy	header	string null		Split strategy to use when uploading a file. If not provided, the default strategy will be used.

Example — a representative request:

BASH

```
curl -X POST "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/resource/{path_rid}/file/{field}/upload" \
-H "Authorization: Bearer $KB_API_KEY"
```

Response codes

Status	Schema	Description
201	ResourceFileUploaded	Successful Response
422	HTTPValidationError	Validation Error

PUT /kb/{kbid}/resource/{rid}/conversation/{field_id}

Add resource conversation field (by id) · roles: WRITER

Parameters

Name	In	Type	Req	Description
kbid	path	string	yes	
rid	path	string	yes	
field_id	path	string	yes	

Request body — application/json (required)

Field	Type	Req	Default	Description
messages	array<InputMessage>			List of messages in the conversation field. Each message must have a unique ident. A single

				conversation can contain up to 51,200 messages. You can add up to 2,048 messages per request.
<code>extract_strategy</code>	string null			Id of the Nuclia extract strategy used at processing time. If not set, the default strategy was used. Extract strategies are defined at the learning configuration api.
<code>split_strategy</code>	string null			Id of the Nuclia split strategy used at processing time. If not set, the default strategy was used. Split strategies are defined at the learning configuration api.

Example — a representative request:

```
BASH
curl -X PUT
"https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/resource/$RID/conversation/
{field_id}" \
-H "Authorization: Bearer $KB_API_KEY" \
-H "Content-Type: application/json" \
-d '{}'
```

Response codes

Status	Schema	Description
201	<code>ResourceFieldAdded</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

**GET /kb/{kbid}/resource/{rid}/conversation/{field_id}/download/field/{message_id}/
{file_num}**

Download conversation binary field (by id) · roles: READER

Parameters

Name	In	Type	Req	Description
<code>kbid</code>	path	string	yes	
<code>rid</code>	path	string	yes	
<code>field_id</code>	path	string	yes	
<code>message_id</code>	path	string	yes	
<code>file_num</code>	path	integer	yes	
<code>range</code>	header	string null		Standard HTTP Range header that enable multipart requests

Example — a representative request:

```
BASH
curl -X GET
"https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/resource/$RID/conversation/
{field_id}/download/field/{message_id}/{file_num}" \
-H "Authorization: Bearer $KB_API_KEY"
```

Response codes

Status	Schema	Description
200	<code>object</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

PUT /kb/{kbid}/resource/{rid}/conversation/{field_id}/messages

Append messages to conversation field (by id) · roles: WRITER

Parameters

Name	In	Type	Req	Description
<code>kbid</code>	path	string	yes	
<code>rid</code>	path	string	yes	
<code>field_id</code>	path	string	yes	

Request body — `application/json` (required)

Field	Type	Req	Default	Description
(array root)	array<InputMessage>			
<code>timestamp</code>	string null			Time at which the message was sent, in ISO 8601 format.
<code>who</code>	string null			Sender of the message, e.g. 'user' or 'assistant'
<code>to</code>	array<string>			List of recipients of the message, e.g. ['assistant'] or ['user']
<code>content</code>	InputMessageContent	yes		
<code>ident</code>	string	yes		Unique identifier for the message. Must be unique within the conversation.
<code>type</code>	MessageType null			

Example — a representative request:

BASH

```
curl -X PUT
"https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/resource/$RID/conversation/
{field_id}/messages" \
-H "Authorization: Bearer $KB_API_KEY" \
-H "Content-Type: application/json" \
-d '[
{
  "content": {
    "text": "string",
    "format": "PLAIN"
  },
  "ident": "string",
  "to": [
    "string"
  ]
}]'
```

Response codes

Status	Schema	Description
--------	--------	-------------

200	ResourceFieldAdded	Successful Response
422	HTTPValidationError	Validation Error

DELETE /kb/{kbid}/resource/{rid}/conversation/{field_id}/messages/{message_ident}

Delete conversation message (by id) · roles: WRITER

Delete a message from a conversation field by its ident. This is a convenience endpoint intended for occasional deletions — frequent use on the same conversation may degrade read performance over time. Deleted message identifiers are permanently reserved and cannot be reused.

Parameters

Name	In	Type	Req	Description
kbid	path	string	yes	
rid	path	string	yes	
field_id	path	string	yes	
message_ident	path	string	yes	The ident of the message to delete

Example — a representative request:

```
BASH
curl -X DELETE
"https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/resource/$RID/conversation/
{field_id}/messages/{message_ident}" \
-H "Authorization: Bearer $KB_API_KEY"
```

Response codes

Status	Schema	Description
204		Successful Response
422	HTTPValidationError	Validation Error

PUT /kb/{kbid}/resource/{rid}/file/{field_id}

Add resource file field (by id) · roles: WRITER

Parameters

Name	In	Type	Req	Description
<code>kbid</code>	path	string	yes	
<code>rid</code>	path	string	yes	
<code>field_id</code>	path	string	yes	
<code>x-skip-store</code>	header	boolean		If set to true, file fields will not be saved in the blob storage. They will only be sent to process.

Request body — `application/json` (required)

Field	Type	Req	Default	Description
<code>language</code>	string null			
<code>password</code>	string null			
<code>file</code>	File	yes		
<code>extract_strategy</code>	string null			Id of the Nuclia extract strategy to use at processing time. If not set, the default strategy will be used. Extract strategies are defined at the learning configuration api.
<code>split_strategy</code>	string null			Id of the Nuclia split strategy used at processing time. If not set, the default strategy was used. Split strategies are defined at the learning configuration api.

Example — a representative request:

BASH

```
curl -X PUT "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/resource/$RID/file/{field_id}" \
-H "Authorization: Bearer $KB_API_KEY" \
-H "Content-Type: application/json" \
-d '{
  "file": {
    "content_type": "application/octet-stream"
  }
}'
```

Response codes

Status	Schema	Description
201	ResourceFieldAdded	Successful Response
422	HTTPValidationError	Validation Error

GET /kb/{kbid}/resource/{rid}/file/{field_id}/download/field

Download field binary field (by id) · roles: **READER**

Parameters

Name	In	Type	Req	Description
kbid	path	string	yes	
rid	path	string	yes	
field_id	path	string	yes	
inline	query	boolean		
range	header	string null		Standard HTTP Range header that enable multipart requests

Example — a representative request:

BASH

```
curl -X GET "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/resource/$RID/file/{field_id}/download/field" \
-H "Authorization: Bearer $KB_API_KEY"
```

Response codes

Status	Schema	Description
200	<code>object</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

POST /kb/{kbid}/resource/{rid}/file/{field_id}/reprocess

Reprocess file field (by id) · roles: WRITER

Parameters

Name	In	Type	Req	Description
<code>kbid</code>	path	string	yes	
<code>rid</code>	path	string	yes	
<code>field_id</code>	path	string	yes	
<code>reset_title</code>	query	boolean		Reset the title of the resource so that the file or link computed titles are set after processing.
<code>x-nucliadb-user</code>	header	string		
<code>x-file-password</code>	header	string null		If a file is password protected, the password must be provided here for the file to be processed

Example — a representative request:

BASH

```
curl -X POST "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/resource/$RID/file/{field_id}/reprocess" \
-H "Authorization: Bearer $KB_API_KEY"
```

Response codes

Status	Schema	Description
202	<code>ResourceUpdated</code>	Successful Response

422	HTTPValidationError	Validation Error
-----	---------------------	------------------

PUT /kb/{kbid}/resource/{rid}/link/{field_id}

Add resource link field (by id) · roles: WRITER

Parameters

Name	In	Type	Req	Description
kbid	path	string	yes	
rid	path	string	yes	
field_id	path	string	yes	

Request body — application/json (required)

Field	Type	Req	Default	Description
headers	map<string, string> null		{}	
cookies	map<string, string> null		{}	
uri	string	yes		
language	string null			
localstorage	map<string, string> null		{}	
css_selector	string null			
xpath	string null			
extract_strategy	string null			Id of the Nuclia extract strategy to use at processing time. If not set, the default strategy will be used. Extract strategies are defined at the learning configuration api.
split_strategy	string null			Id of the Nuclia split strategy used at processing time. If

				not set, the default strategy was used. Split strategies are defined at the learning configuration api.
--	--	--	--	---

Example — a representative request:

BASH

```
curl -X PUT "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/resource/$RID/link/{field_id}" \
  -H "Authorization: Bearer $KB_API_KEY" \
  -H "Content-Type: application/json" \
  -d '{
    "uri": "https://example.com"
  }'
```

Response codes

Status	Schema	Description
201	ResourceFieldAdded	Successful Response
422	HTTPValidationError	Validation Error

PUT /kb/{kbid}/resource/{rid}/text/{field_id}

Add resource text field (by id) · roles: WRITER

Parameters

Name	In	Type	Req	Description
kbid	path	string	yes	
rid	path	string	yes	
field_id	path	string	yes	

Request body — [application/json](#) (required)

Field	Type	Req	Default	Description
body	string	yes		The text body. The format of the text should be specified in

				the format field. The sum of all text fields in the request may not exceed 2MB. If you need to store more text, consider using a file field...
format	TextFormat		PLAIN	The format of the text.
extract_strategy	string null			Id of the Nuclia extract strategy to use at processing time. If not set, the default strategy will be used. Extract strategies are defined at the learning configuration api.
split_strategy	string null			Id of the Nuclia split strategy used at processing time. If not set, the default strategy was used. Split strategies are defined at the learning configuration api.

Example — a representative request:

BASH

```
curl -X PUT "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/resource/$RID/text/{field_id}" \
  -H "Authorization: Bearer $KB_API_KEY" \
  -H "Content-Type: application/json" \
  -d '{
    "body": "string",
    "format": "PLAIN"
  }'
```

Response codes

Status	Schema	Description
201	ResourceFieldAdded	Successful Response

422	<code>HTTPValidationError</code>	Validation Error
-----	----------------------------------	------------------

GET /kb/{kbid}/resource/{rid}/{field_type}/{field_id}

Get Resource field (by id) · roles: READER

Parameters

Name	In	Type	Req	Description
<code>kbid</code>	path	string	yes	
<code>rid</code>	path	string	yes	
<code>field_type</code>	path	FieldTypeName	yes	
<code>field_id</code>	path	string	yes	
<code>show</code>	query	array<ResourceField Properties>		
<code>extracted</code>	query	array<ExtractedData TypeName>		
<code>page</code>	query	string or integer		

Example — a representative request:

```
BASH
curl -X GET
"https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/resource/$RID/{field_type}/{field_id}" \
-H "Authorization: Bearer $KB_API_KEY"
```

Response codes

Status	Schema	Description
200	<code>ResourceField</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

DELETE /kb/{kbid}/resource/{rid}/{field_type}/{field_id}

Delete Resource field (by id) · roles: WRITER

Parameters

Name	In	Type	Req	Description
<code>kbid</code>	path	string	yes	
<code>rid</code>	path	string	yes	
<code>field_type</code>	path	FieldTypeName	yes	
<code>field_id</code>	path	string	yes	

Example — a representative request:

```
BASH
curl -X DELETE
"https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/resource/$RID/{field_type}/
{field_id}" \
-H "Authorization: Bearer $KB_API_KEY"
```

Response codes

Status	Schema	Description
204		Successful Response
422	<code>HTTPValidationError</code>	Validation Error

GET /kb/{kbid}/resource/{rid}/{field_type}/{field_id}/download/extracted/{download_field}

Download extracted binary file (by id) · roles: READER

Parameters

Name	In	Type	Req	Description
<code>kbid</code>	path	string	yes	
<code>rid</code>	path	string	yes	
<code>field_type</code>	path	FieldTypeName	yes	
<code>field_id</code>	path	string	yes	
<code>download_field</code>	path	string	yes	
<code>range</code>	header	string null		Standard HTTP Range header that enable multipart requests

Example — a representative request:

```
BASH
curl -X GET
"https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/resource/$RID/{field_type}/
{field_id}/download/extracted/{download_field}" \
-H "Authorization: Bearer $KB_API_KEY"
```

Response codes

Status	Schema	Description
200	<code>object</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

PUT /kb/{kbid}/slug/{rslug}/conversation/{field_id}

Add resource conversation field (by slug) · roles: WRITER

Parameters

Name	In	Type	Req	Description
<code>kbid</code>	path	string	yes	
<code>rslug</code>	path	string	yes	
<code>field_id</code>	path	string	yes	

Request body — `application/json` (required)

Field	Type	Req	Default	Description
<code>messages</code>	<code>array<InputMessage></code>			List of messages in the conversation field. Each message must have a unique ident. A single conversation can contain up to 51,200 messages. You can add up to 2,048 messages per request.
<code>extract_strat</code>	<code>string null</code>			Id of the Nuclia

egy				extract strategy used at processing time. If not set, the default strategy was used. Extract strategies are defined at the learning configuration api.
split_strategy	string null			Id of the Nuclia split strategy used at processing time. If not set, the default strategy was used. Split strategies are defined at the learning configuration api.

Example — a representative request:

BASH

```
curl -X PUT
"https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/slug/{rslug}/conversation/
{field_id}" \
-H "Authorization: Bearer $KB_API_KEY" \
-H "Content-Type: application/json" \
-d '{}'
```

Response codes

Status	Schema	Description
201	ResourceFieldAdded	Successful Response
422	HTTPValidationError	Validation Error

**GET /kb/{kbid}/slug/{rslug}/conversation/{field_id}/download/field/{message_id}/
{file_num}**

Download conversation binary field (by slug) · roles: READER

Parameters

Name	In	Type	Req	Description
kbid	path	string	yes	

rslug	path	string	yes	
field_id	path	string	yes	
message_id	path	string	yes	
file_num	path	integer	yes	
range	header	string null		Standard HTTP Range header that enable multipart requests

Example — a representative request:

```
BASH
curl -X GET
"https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/slug/{rslug}/conversation/
{field_id}/download/field/{message_id}/{file_num}" \
-H "Authorization: Bearer $KB_API_KEY"
```

Response codes

Status	Schema	Description
200	object	Successful Response
422	HTTPValidationError	Validation Error

PUT /kb/{kbid}/slug/{rslug}/conversation/{field_id}/messages

Append messages to conversation field (by slug) · roles: WRITER

Parameters

Name	In	Type	Req	Description
kbid	path	string	yes	
rslug	path	string	yes	
field_id	path	string	yes	

Request body — application/json (required)

Field	Type	Req	Default	Description
-------	------	-----	---------	-------------

(array root)	array<InputMessage>			
timestamp	string null			Time at which the message was sent, in ISO 8601 format.
who	string null			Sender of the message, e.g. 'user' or 'assistant'
to	array<string>			List of recipients of the message, e.g. ['assistant'] or ['user']
content	InputMessageContent	yes		
ident	string	yes		Unique identifier for the message. Must be unique within the conversation.
type	MessageType null			

Example — a representative request:

```
BASH
curl -X PUT
"https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/slug/{rslug}/conversation/
{field_id}/messages" \
-H "Authorization: Bearer $KB_API_KEY" \
-H "Content-Type: application/json" \
-d '[
{
  "content": {
    "text": "string",
    "format": "PLAIN"
  },
  "ident": "string",
  "to": [
    "string"
  ]
}
]'
```

Response codes

Status	Schema	Description
200	ResourceFieldAdded	Successful Response
422	HTTPValidationError	Validation Error

DELETE /kb/{kbid}/slug/{rslug}/conversation/{field_id}/messages/{message_ident}

Delete conversation message (by slug) · roles: WRITER

Delete a message from a conversation field by its ident. This is a convenience endpoint intended for occasional deletions — frequent use on the same conversation may degrade read performance over time. Deleted message identifiers are permanently reserved and cannot be reused.

Parameters

Name	In	Type	Req	Description
kbid	path	string	yes	
rslug	path	string	yes	
field_id	path	string	yes	
message_ident	path	string	yes	The ident of the message to delete

Example — a representative request:

```
BASH
curl -X DELETE
"https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/slug/{rslug}/conversation/
{field_id}/messages/{message_ident}" \
-H "Authorization: Bearer $KB_API_KEY"
```

Response codes

Status	Schema	Description
204		Successful Response
422	HTTPValidationError	Validation Error

PUT /kb/{kbid}/slug/{rslug}/file/{field_id}

Add resource file field (by slug) · roles: WRITER

Parameters

Name	In	Type	Req	Description
kbid	path	string	yes	

rslug	path	string	yes	
field_id	path	string	yes	
x-skip-store	header	boolean		If set to true, file fields will not be saved in the blob storage. They will only be sent to process.

Request body — `application/json` (required)

Field	Type	Req	Default	Description
language	string null			
password	string null			
file	File	yes		
extract_strategy	string null			Id of the Nuclia extract strategy to use at processing time. If not set, the default strategy will be used. Extract strategies are defined at the learning configuration api.
split_strategy	string null			Id of the Nuclia split strategy used at processing time. If not set, the default strategy was used. Split strategies are defined at the learning configuration api.

Example — a representative request:

BASH

```
curl -X PUT
"https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/slug/{rslug}/file/{field_id}" \
-H "Authorization: Bearer $KB_API_KEY" \
-H "Content-Type: application/json" \
-d '{
```

```
"file": {
  "content_type": "application/octet-stream"
}
```

Response codes

Status	Schema	Description
201	ResourceFieldAdded	Successful Response
422	HTTPValidationError	Validation Error

GET /kb/{kbid}/slug/{rslug}/file/{field_id}/download/field

Download field binary field (by slug) · roles: READER

Parameters

Name	In	Type	Req	Description
kbid	path	string	yes	
rslug	path	string	yes	
field_id	path	string	yes	
inline	query	boolean		
range	header	string null		Standard HTTP Range header that enable multipart requests

Example — a representative request:

```
BASH
curl -X GET
"https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/slug/{rslug}/file/{field_id}/
download/field" \
-H "Authorization: Bearer $KB_API_KEY"
```

Response codes

Status	Schema	Description
200	object	Successful Response

422	HTTPValidationError	Validation Error
-----	---------------------	------------------

POST /kb/{kbid}/slug/{rslug}/file/{field}/upload

Upload binary file on a Resource (by slug) · roles: WRITER

Upload a file as a field on an existing resource, if the field exists will return a conflict (419)

Parameters

Name	In	Type	Req	Description
kbid	path	string	yes	
rslug	path	string	yes	
field	path	string	yes	
x-filename	header	string null		Name of the file being uploaded.
x-password	header	string null		If the file is password protected, the password must be provided here.
x-language	header	string null		
x-md5	header	string null		MD5 hash of the file being uploaded. This is used to check if the file has been uploaded before.
x-extract-strategy	header	string null		Extract strategy to use when uploading a file. If not provided, the default strategy will be used.
x-split-strategy	header	string null		Split strategy to use when uploading a file. If not provided, the default strategy will be used.

Example — a representative request:

```
BASH
curl -X POST
```

```
"https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/slug/{rslug}/file/{field}/upload"
\
-H "Authorization: Bearer $KB_API_KEY"
```

Response codes

Status	Schema	Description
201	ResourceFileUploaded	Successful Response
422	HTTPValidationError	Validation Error

PUT /kb/{kbid}/slug/{rslug}/link/{field_id}

Add resource link field (by slug) · roles: WRITER

Parameters

Name	In	Type	Req	Description
kbid	path	string	yes	
rslug	path	string	yes	
field_id	path	string	yes	

Request body — application/json (required)

Field	Type	Req	Default	Description
headers	map<string, string> null		{}	
cookies	map<string, string> null		{}	
uri	string	yes		
language	string null			
localStorage	map<string, string> null		{}	
css_selector	string null			
xpath	string null			
extract_strategy	string null			Id of the Nuclia extract strategy to use at processing time. If

				not set, the default strategy will be used. Extract strategies are defined at the learning configuration api.
<code>split_strategy</code>	string null			Id of the Nuclia split strategy used at processing time. If not set, the default strategy was used. Split strategies are defined at the learning configuration api.

Example — a representative request:

```
BASH
curl -X PUT
"https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/slug/{rslug}/link/{field_id}" \
-H "Authorization: Bearer $KB_API_KEY" \
-H "Content-Type: application/json" \
-d '{
  "uri": "https://example.com"
}'
```

Response codes

Status	Schema	Description
201	<code>ResourceFieldAdded</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

PUT /kb/{kbid}/slug/{rslug}/text/{field_id}

Add resource text field (by slug) · roles: WRITER

Parameters

Name	In	Type	Req	Description
<code>kbid</code>	path	string	yes	
<code>rslug</code>	path	string	yes	

field_id	path	string	yes	
----------	------	--------	-----	--

Request body — `application/json` (required)

Field	Type	Req	Default	Description
body	string	yes		The text body. The format of the text should be specified in the format field. The sum of all text fields in the request may not exceed 2MB. If you need to store more text, consider using a file field...
format	TextFormat		PLAIN	The format of the text.
extract_strategy	string null			Id of the Nuclia extract strategy to use at processing time. If not set, the default strategy will be used. Extract strategies are defined at the learning configuration api.
split_strategy	string null			Id of the Nuclia split strategy used at processing time. If not set, the default strategy was used. Split strategies are defined at the learning configuration api.

Example — a representative request:

BASH

```
curl -X PUT
"https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/slug/{rslug}/text/{field_id}" \
-H "Authorization: Bearer $KB_API_KEY" \
-H "Content-Type: application/json" \
-d '{
```

```
"body": "string",
"format": "PLAIN"
}'
```

Response codes

Status	Schema	Description
201	ResourceFieldAdded	Successful Response
422	HTTPValidationError	Validation Error

GET /kb/{kbid}/slug/{rslug}/{field_type}/{field_id}

Get Resource field (by slug) · roles: READER

Parameters

Name	In	Type	Req	Description
kbid	path	string	yes	
rslug	path	string	yes	
field_type	path	FieldTypeName	yes	
field_id	path	string	yes	
show	query	array<ResourceField Properties>		
extracted	query	array<ExtractedData TypeName>		
page	query	string or integer		

Example — a representative request:

```
BASH
curl -X GET
"https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/slug/{rslug}/{field_type}/{field_id}" \
-H "Authorization: Bearer $KB_API_KEY"
```

Response codes

Status	Schema	Description
--------	--------	-------------

200	ResourceField	Successful Response
422	HTTPValidationError	Validation Error

DELETE /kb/{kbid}/slug/{rslug}/{field_type}/{field_id}

Delete Resource field (by slug) · roles: WRITER

Parameters

Name	In	Type	Req	Description
kbid	path	string	yes	
rslug	path	string	yes	
field_type	path	FieldTypeName	yes	
field_id	path	string	yes	

Example — a representative request:

```
BASH
curl -X DELETE
"https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/slug/{rslug}/{field_type}/{field_id}" \
-H "Authorization: Bearer $KB_API_KEY"
```

Response codes

Status	Schema	Description
204		Successful Response
422	HTTPValidationError	Validation Error

GET /kb/{kbid}/slug/{rslug}/{field_type}/{field_id}/download/extracted/{download_field}

Download extracted binary file (by slug) · roles: READER

Parameters

Name	In	Type	Req	Description
kbid	path	string	yes	
rslug	path	string	yes	

<code>field_type</code>	path	FieldName	yes	
<code>field_id</code>	path	string	yes	
<code>download_field</code>	path	string	yes	
<code>range</code>	header	string null		Standard HTTP Range header that enable multipart requests

Example — a representative request:

```
BASH
curl -X GET
"https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/slug/{rslug}/{field_type}/
{field_id}/download/extracted/{download_field}" \
-H "Authorization: Bearer $KB_API_KEY"
```

Response codes

Status	Schema	Description
200	<code>object</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

Resources (14)

GET /kb/{kbid}/resource/{rid}

Get Resource (by id) · roles: READER

Parameters

Name	In	Type	Req	Description
<code>kbid</code>	path	string	yes	
<code>rid</code>	path	string	yes	
<code>show</code>	query	array<ResourceProperties>		
<code>field_type</code>	query	array<FieldName>		

extracted	query	array<ExtractedData TypeName>		
x-nucliadb- user	header	string		
x-forwarded- for	header	string		

Example — a real call captured against the live sandbox:

BASH

```
curl -X GET "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/resource/$RID" \
-H "Authorization: Bearer $KB_API_KEY"
```

Response:

JSON

```
{
  "id": "ca2a8f3bf186491facde30cdb2fbc2b6",
  "slug": "solar-basics",
  "title": "Solar Panel Basics",
  "summary": "",
  "icon": "text/plain",
  "thumbnail": "",
  "metadata": {
    "metadata": {},
    "language": "en",
    "languages": [
      "en"
    ],
    "status": "PROCESSED"
  },
  "usermetadata": {
    "classifications": [
      {
        "labelset": "topic",
        "label": "energy",
        "cancelled_by_user": false
      }
    ]
  },
  "fieldmetadata": [],
  "computedmetadata": {
    "field_classifications": []
  },
  "created": "2026-07-27T05:24:17.119810",
  "modified": "2026-07-27T05:24:17.119827",
  "last_seqid": 0,
  "last_account_seq": null,
  "queue": "private",
}
```

```

"hidden": false,
"data": {
  "texts": {
    "body": {
      "value": {
        "body": "A photovoltaic solar panel converts sunlight into direct
current electricity. A residential rooftop system
... (truncated)

```

Response codes

Status	Schema	Description
200	<code>nucliadb_models__resource__Resource</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

PATCH /kb/{kbid}/resource/{rid}

Modify Resource (by id) · roles: WRITER

Parameters

Name	In	Type	Req	Description
<code>kbid</code>	path	string	yes	
<code>rid</code>	path	string	yes	
<code>x-nucliadb-user</code>	header	string		
<code>x-skip-store</code>	header	boolean		If set to true, file fields will not be saved in the blob storage. They will only be sent to process.

Request body — `application/json` (required)

Field	Type	Req	Default	Description
<code>title</code>	string null			
<code>summary</code>	string null			

slug	string null			The slug is the user-defined id for the resource
thumbnail	string null			
metadata	InputMetadata null			
usermetadata	UserMetadata null			
fieldmetadata	array<UserFieldMetadata> null			
origin	InputOrigin null			
extra	Extra null			Extra metadata for the resource. It can be used to store structured information about the resource that can't be used to query at retrieval time. If not set, the existing extra metadata will not be mo...
files	object		{ }	Dictionary of file fields to be added to the resource. The keys correspond to the field id, and must comply with the regex: <code>^[a-zA-Z0-9:-]+\$</code>
links	object		{ }	Dictionary of link fields to be added to the resource. The keys correspond to the field id, and must comply with the regex: <code>^[a-zA-Z0-9:-]+\$</code>
texts	object		{ }	Dictionary of text fields to be added to the resource. The keys correspond to the field id, and must comply with the regex: <code>^[a-zA-Z0-9:-]+\$</code>

				+\$
<code>conversations</code>	object		<code>{}</code>	Dictionary of conversation fields to be added to the resource. The keys correspond to the field id, and must comply with the regex: <code>^[a-zA-Z0-9:-]+\$</code>
<code>key_values</code>	object		<code>{}</code>	Dictionary of key-value fields to be added to the resource. The key must be the schema name and must comply with the regex: <code>^[a-zA-Z0-9:-]+\$</code>
<code>processing_options</code>	PushProcessingOptions null		<code>{'ml_text': True}</code>	Options for processing the resource. If not set, the default options will be used.
<code>security</code>	ResourceSecurity null			Security metadata for the resource. It can be used to have fine-grained control over who can access the resource.
<code>hidden</code>	boolean null			Modify the hidden status of the resource. If not set, the hidden status will not be modified.

Example — a representative request:

```
BASH
curl -X PATCH "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/resource/$RID" \
  -H "Authorization: Bearer $KB_API_KEY" \
  -H "Content-Type: application/json" \
  -d '{}'
```

Response codes

Status	Schema	Description
200	ResourceUpdated	Successful Response
422	HTTPValidationError	Validation Error

DELETE /kb/{kbid}/resource/{rid}

Delete Resource (by id) · roles: WRITER

Parameters

Name	In	Type	Req	Description
kbid	path	string	yes	
rid	path	string	yes	

Example — a representative request:

```
BASH
curl -X DELETE "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/resource/$RID" \
-H "Authorization: Bearer $KB_API_KEY"
```

Response codes

Status	Schema	Description
204		Successful Response
422	HTTPValidationError	Validation Error

HEAD /kb/{kbid}/resource/{rid}

Head Resource (by id) · roles: READER

Parameters

Name	In	Type	Req	Description
kbid	path	string	yes	
rid	path	string	yes	

Example — a representative request:

BASH

```
curl -X HEAD "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/resource/$RID" \
-H "Authorization: Bearer $KB_API_KEY"
```

Response codes

Status	Schema	Description
200	<code>object</code>	Successful Response
404		Resource does not exist
422	<code>HTTPValidationError</code>	Validation Error

POST /kb/{kbid}/resource/{rid}/reindex

Reindex Resource (by id) · roles: WRITER

Parameters

Name	In	Type	Req	Description
<code>kbid</code>	path	string	yes	
<code>rid</code>	path	string	yes	
<code>reindex_vectors</code>	query	boolean		

Example — a representative request:

BASH

```
curl -X POST
"https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/resource/$RID/reindex" \
-H "Authorization: Bearer $KB_API_KEY"
```

Response codes

Status	Schema	Description
204		Successful Response
422	<code>HTTPValidationError</code>	Validation Error

POST /kb/{kbid}/resource/{rid}/reprocess

Reprocess resource (by id) · roles: WRITER

Parameters

Name	In	Type	Req	Description
kbid	path	string	yes	
rid	path	string	yes	
reset_title	query	boolean		Reset the title of the resource so that the file or link computed titles are set after processing.
x-nucliadb-user	header	string		

Example — a representative request:

```
BASH
curl -X POST
"https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/resource/$RID/reprocess" \
-H "Authorization: Bearer $KB_API_KEY"
```

Response codes

Status	Schema	Description
202	ResourceUpdated	Successful Response
422	HTTPValidationError	Validation Error

GET /kb/{kbid}/resources

List Resources · roles: READER

List of resources of a knowledgebox

Parameters

Name	In	Type	Req	Description
kbid	path	string	yes	

page	query	integer		Requested page number (0-based)
size	query	integer		Page size

Example — a representative request:

```
BASH
curl -X GET "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/resources" \
-H "Authorization: Bearer $KB_API_KEY"
```

Response codes

Status	Schema	Description
200	ResourceList	Successful Response
422	HTTPValidationError	Validation Error

POST /kb/{kbid}/resources

Create Resource · roles: WRITER

Create a new Resource in a Knowledge Box

Parameters

Name	In	Type	Req	Description
kbid	path	string	yes	
x-skip-store	header	boolean		If set to true, file fields will not be saved in the blob storage. They will only be sent to process.
x-nucliadb-user	header	string		

Request body — application/json (required)

Field	Type	Req	Default	Description
title	string null			

summary	string null			
slug	string null			The slug is the user-defined id for the resource
icon	string null			The icon should be a media type string: (see docs)
thumbnail	string null			
metadata	InputMetadata null			Generic metadata for the resource. It can be used to store structured information about the resource that later is serialized on retrieval results, however this metadata can not be used for searching ...
usermetadata	UserMetadata null			
fieldmetadata	array<UserFieldMetadata> null			
origin	InputOrigin null			Origin metadata for the resource. Used to store information about the resource on the origin system. Most of its fields can later be used to filter at search time.
extra	Extra null			Extra metadata for the resource. It can be used to store structured information about the resource that can't be used to query at retrieval time.
hidden	boolean null			Set the hidden status of the resource. If not set, the default value for new resources in the KnowledgeBox will be used.

<code>files</code>	object		<code>{}</code>	Dictionary of file fields to be added to the resource. The keys correspond to the field id, and must comply with the regex: <code>^[a-zA-Z0-9:-]+\$</code>
<code>links</code>	object		<code>{}</code>	Dictionary of link fields to be added to the resource. The keys correspond to the field id, and must comply with the regex: <code>^[a-zA-Z0-9:-]+\$</code>
<code>texts</code>	object		<code>{}</code>	Dictionary of text fields to be added to the resource. The keys correspond to the field id, and must comply with the regex: <code>^[a-zA-Z0-9:-]+\$</code>
<code>conversations</code>	object		<code>{}</code>	Dictionary of conversation fields to be added to the resource. The keys correspond to the field id, and must comply with the regex: <code>^[a-zA-Z0-9:-]+\$</code>
<code>key_values</code>	object		<code>{}</code>	Dictionary of key-value fields to be added to the resource. The key must be the schema name and must comply with the regex: <code>^[a-zA-Z0-9:-]+\$</code>
<code>processing_options</code>	PushProcessingOptions null		<code>{'ml_text': True}</code>	Options for processing the resource. If not set, the default options

				will be used.
security	ResourceSecurity null			Security metadata for the resource. It can be used to have fine-grained control over who can access the resource.

Example — a real call captured against the live sandbox:

BASH

```
curl -X POST "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/resources" \
-H "Authorization: Bearer $KB_API_KEY" \
-H "Content-Type: application/json" \
-d '{
  "slug": "solar-basics",
  "title": "Solar Panel Basics",
  "texts": {
    "body": {
      "body": "A photovoltaic solar panel converts sunlight into direct current electricity. A residential rooftop system in Australia typically ranges from 5kW to 10kW. Inverters convert DC to AC. The Clean Energy Council certifies installers. Payback periods are commonly 4 to 7 years depending on feed-in tariffs.",
      "format": "PLAIN"
    }
  },
  "usermetadata": {
    "classifications": [
      {
        "labelset": "topic",
        "label": "energy"
      }
    ]
  }
}'
```

Response:

JSON

```
{
  "uuid": "ca2a8f3bf186491facde30cdb2fbc2b6",
  "elapsed": 0.23978948593139648,
  "seqid": null
}
```

Response codes

Status	Schema	Description
--------	--------	-------------

201	ResourceCreated	Successful Response
422	HTTPValidationError	Validation Error

GET /kb/{kbid}/slug/{rslug}

Get Resource (by slug) · roles: READER

Parameters

Name	In	Type	Req	Description
kbid	path	string	yes	
rslug	path	string	yes	
show	query	array<ResourceProperties>		
field_type	query	array<FieldTypeName>		
extracted	query	array<ExtractedDataTypeName>		
x-nucliadb-user	header	string		
x-forwarded-for	header	string		

Example — a representative request:

```
BASH
curl -X GET "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/slug/{rslug}" \
-H "Authorization: Bearer $KB_API_KEY"
```

Response codes

Status	Schema	Description
200	nucliadb_models__resource__Resource	Successful Response
422	HTTPValidationError	Validation Error

PATCH /kb/{kbid}/slug/{rslug}

Modify Resource (by slug) · roles: WRITER

Parameters

Name	In	Type	Req	Description
kbid	path	string	yes	
rslug	path	string	yes	
x-skip-store	header	boolean		If set to true, file fields will not be saved in the blob storage. They will only be sent to process.
x-nucliadb-user	header	string		

Request body — `application/json` (required)

Field	Type	Req	Default	Description
title	string null			
summary	string null			
slug	string null			The slug is the user-defined id for the resource
thumbnail	string null			
metadata	InputMetadata null			
usermetadata	UserMetadata null			
fieldmetadata	array<UserFieldMetadata> null			
origin	InputOrigin null			
extra	Extra null			Extra metadata for the resource. It can be used to store structured information about the resource that can't be used to query at

				retrieval time. If not set, the existing extra metadata will not be mo...
files	object		{}	Dictionary of file fields to be added to the resource. The keys correspond to the field id, and must comply with the regex: <code>^[a-zA-Z0-9:-]+\$</code>
links	object		{}	Dictionary of link fields to be added to the resource. The keys correspond to the field id, and must comply with the regex: <code>^[a-zA-Z0-9:-]+\$</code>
texts	object		{}	Dictionary of text fields to be added to the resource. The keys correspond to the field id, and must comply with the regex: <code>^[a-zA-Z0-9:-]+\$</code>
conversations	object		{}	Dictionary of conversation fields to be added to the resource. The keys correspond to the field id, and must comply with the regex: <code>^[a-zA-Z0-9:-]+\$</code>
key_values	object		{}	Dictionary of key-value fields to be added to the resource. The key must be the schema name and must comply with the regex: <code>^[a-zA-Z0-9:-]+\$</code>

<code>processing_options</code>	PushProcessingOptions null		<code>{'ml_text': True}</code>	Options for processing the resource. If not set, the default options will be used.
<code>security</code>	ResourceSecurity null			Security metadata for the resource. It can be used to have fine-grained control over who can access the resource.
<code>hidden</code>	boolean null			Modify the hidden status of the resource. If not set, the hidden status will not be modified.

Example — a representative request:

BASH

```
curl -X PATCH "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/slug/{rslug}" \
-H "Authorization: Bearer $KB_API_KEY" \
-H "Content-Type: application/json" \
-d '{}'
```

Response codes

Status	Schema	Description
200	<code>ResourceUpdated</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

DELETE /kb/{kbid}/slug/{rslug}

Delete Resource (by slug) · roles: WRITER

Parameters

Name	In	Type	Req	Description
<code>kbid</code>	path	string	yes	
<code>rslug</code>	path	string	yes	

Example — a representative request:

```
BASH
curl -X DELETE "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/slug/{rslug}" \
-H "Authorization: Bearer $KB_API_KEY"
```

Response codes

Status	Schema	Description
204		Successful Response
422	HTTPValidationError	Validation Error

HEAD /kb/{kbid}/slug/{rslug}

Head Resource (by slug) · roles: READER

Parameters

Name	In	Type	Req	Description
kbid	path	string	yes	
rslug	path	string	yes	

Example — a representative request:

```
BASH
curl -X HEAD "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/slug/{rslug}" \
-H "Authorization: Bearer $KB_API_KEY"
```

Response codes

Status	Schema	Description
200	object	Successful Response
404		Resource does not exist
422	HTTPValidationError	Validation Error

POST /kb/{kbid}/slug/{rslug}/reindex

Reindex Resource (by slug) · roles: WRITER

Parameters

Name	In	Type	Req	Description
<code>kbid</code>	path	string	yes	
<code>rslug</code>	path	string	yes	
<code>reindex_vectors</code>	query	boolean		

Example — a representative request:

```
BASH
curl -X POST "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/slug/{rslug}/reindex" \
  -H "Authorization: Bearer $KB_API_KEY"
```

Response codes

Status	Schema	Description
204		Successful Response
422	<code>HTTPValidationError</code>	Validation Error

POST /kb/{kbid}/slug/{rslug}/reprocess

Reprocess resource (by slug) · roles: WRITER

Parameters

Name	In	Type	Req	Description
<code>kbid</code>	path	string	yes	
<code>rslug</code>	path	string	yes	
<code>reset_title</code>	query	boolean		Reset the title of the resource so that the file or link computed titles are set after processing.
<code>x-nucliadb-user</code>	header	string		

Example — a representative request:

```
BASH
curl -X POST
"https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/slug/{rslug}/reprocess" \
-H "Authorization: Bearer $KB_API_KEY"
```

Response codes

Status	Schema	Description
202	ResourceUpdated	Successful Response
422	HTTPValidationError	Validation Error

Ingestion Agents (2)

POST /kb/{kbid}/resource/{rid}/run-agents

Run Agents on Resource · roles: READER

Parameters

Name	In	Type	Req	Description
kbid	path	string	yes	
rid	path	string	yes	
x-nucliadb-user	header	string		

Request body — `application/json` (required)

Field	Type	Req	Default	Description
filters	array<AgentsFilter> null			Filters to apply to the agents. If None, all currently configured agents are applied.
agent_ids	array<string> null			

Example — a representative request:

BASH

```
curl -X POST "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/resource/$RID/run-
agents" \
-H "Authorization: Bearer $KB_API_KEY" \
-H "Content-Type: application/json" \
-d '{}'
```

Response codes

Status	Schema	Description
200	ResourceAgentsResponse	Successful Response
422	HTTPValidationError	Validation Error

POST /kb/{kbid}/slug/{slug}/run-agents

Run Agents on Resource (by slug) · roles: **READER**

Parameters

Name	In	Type	Req	Description
kbid	path	string	yes	
slug	path	string	yes	
x-nucliadb-user	header	string		

Request body — [application/json](#) (required)

Field	Type	Req	Default	Description
filters	array<AgentsFilter> null			Filters to apply to the agents. If None, all currently configured agents are applied.
agent_ids	array<string> null			

Example — a representative request:**BASH**

```
curl -X POST "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/slug/{slug}/run-
agents" \
-H "Authorization: Bearer $KB_API_KEY" \
```



```
-H "Content-Type: application/json" \
-d '{}'
```

Response codes

Status	Schema	Description
200	<code>ResourceAgentsResponse</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

Split Strategies (4)

GET /kb/{kbid}/split_strategies

Learning split strategies · roles: READER, MANAGER

Get available split strategies

Parameters

Name	In	Type	Req	Description
<code>kbid</code>	path	string	yes	

Example — a real call captured against the live sandbox:

```
BASH
curl -X GET "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/split_strategies" \
-H "Authorization: Bearer $KB_API_KEY"
```

Response:

```
JSON
{}
```

Response codes

Status	Schema	Description
200	<code>map<string, SplitConfiguration></code>	Successful Response

422	HTTPValidationError	Validation Error
-----	---------------------	------------------

POST /kb/{kbid}/split_strategies

Add a split strategy to a KB · roles: MANAGER, OWNER

Parameters

Name	In	Type	Req	Description
kbid	path	string	yes	

Request body — application/json (required)

Field	Type	Req	Default	Description
name	string			
max_paragraph	integer		0	
custom_split	CustomSplitStrategy null			
llm_split	LLMSplitConfig null			
manual_split	ManualSplitConfig null			

Example — a representative request:

```
BASH
curl -X POST "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/split_strategies" \
-H "Authorization: Bearer $KB_API_KEY" \
-H "Content-Type: application/json" \
-d '{
  "name": "",
  "max_paragraph": 0
}'
```

Response codes

Status	Schema	Description
200	string	Successful Response
422	HTTPValidationError	Validation Error

GET /kb/{kbid}/split_strategies/strategy/{strategy_id}

Extract split configuration · roles: READER, MANAGER

Get split strategy for a given id

Parameters

Name	In	Type	Req	Description
<code>kbid</code>	path	string	yes	
<code>strategy_id</code>	path	string	yes	

Example — a representative request:

```
BASH
curl -X GET
"https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/split_strategies/strategy/
{strategy_id}" \
-H "Authorization: Bearer $KB_API_KEY"
```

Response codes

Status	Schema	Description
200	<code>object</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

DELETE /kb/{kbid}/split_strategies/strategy/{strategy_id}

Remove a split strategy from a KB · roles: MANAGER, OWNER

Removes a split strategy from a KB

Parameters

Name	In	Type	Req	Description
<code>kbid</code>	path	string	yes	
<code>strategy_id</code>	path	string	yes	

Example — a representative request:

BASH

```
curl -X DELETE
"https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/split_strategies/strategy/
{strategy_id}" \
-H "Authorization: Bearer $KB_API_KEY"
```

Response codes

Status	Schema	Description
204		Successful Response
422	HTTPValidationError	Validation Error

Knowledge Box TUS uploads (4)

POST /kb/{kbid}/tusupload

Create new upload on a Knowledge Box · roles: WRITER

Parameters

Name	In	Type	Req	Description
kbid	path	string	yes	
x-extract-strategy	header	string null		Extract strategy to use when uploading a file. If not provided, the default strategy will be used.
x-split-strategy	header	string null		Split strategy to use when uploading a file. If not provided, the default strategy will be used.

Example — a representative request:

BASH

```
curl -X POST "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/tusupload" \
-H "Authorization: Bearer $KB_API_KEY" \
-H "Content-Type: application/json" \
-d '{}'
```

Response codes

Status	Schema	Description
200	<code>object</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

OPTIONS /kb/{kbid}/tusupload

TUS Server information

Parameters

Name	In	Type	Req	Description
<code>kbid</code>	path	string	yes	
<code>rid</code>	query	string null		
<code>rslug</code>	query	string null		
<code>upload_id</code>	query	string null		
<code>field</code>	query	string null		

Example — a representative request:

```
BASH
curl -X OPTIONS "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/tusupload" \
  -H "Authorization: Bearer $KB_API_KEY"
```

Response codes

Status	Schema	Description
200	<code>object</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

PATCH /kb/{kbid}/tusupload/{upload_id}

Upload data on a Knowledge Box · roles: WRITER

Parameters

Name	In	Type	Req	Description
<code>kbid</code>	path	string	yes	
<code>upload_id</code>	path	string	yes	

Example — a representative request:

```
BASH
curl -X PATCH
"https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/tusupload/{upload_id}" \
-H "Authorization: Bearer $KB_API_KEY"
```

Response codes

Status	Schema	Description
200	<code>object</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

HEAD /kb/{kbid}/tusupload/{upload_id}

Upload information · roles: WRITER

Parameters

Name	In	Type	Req	Description
<code>kbid</code>	path	string	yes	
<code>upload_id</code>	path	string	yes	

Example — a representative request:

```
BASH
curl -X HEAD
"https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/tusupload/{upload_id}" \
-H "Authorization: Bearer $KB_API_KEY"
```

Response codes

Status	Schema	Description
200	<code>object</code>	Successful Response

422	<code>HTTPValidationError</code>	Validation Error
-----	----------------------------------	------------------

PART 7 — REFERENCE APPENDIX

Appendix B — NUA API Reference

The Nuclia Understanding API — processing, Predict, and agents, fully expanded

The understanding services: processing pipeline, Predict, agents and tasks.

157 operations in **18 groups**, generated from the official [nua](#) v1 OpenAPI specification. Served from <https://<zone>.dp.progress.cloud/api/v1>; authorized with a NUA key (**X-NUCLIA-NUAKEY: Bearer**). Every request-body field, parameter, and response is listed; object-typed fields reference named schemas documented in full in Appendix F (Schema Catalog).

Models (16)

GET /account/{account_id}/assume_role/bedrock

Account Assume Role For Bedrock Models

Start a new assume role authentication flow for an account

Parameters

Name	In	Type	Req	Description
account_id	path	string	yes	

Example — a representative request:

```
BASH
curl -X GET
"https://$ZONE.dp.progress.cloud/api/v1/account/$ACCOUNT/assume_role/bedrock" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
200	BedrockAssumeRoleStartResponse	Successful Response
422	HTTPValidationError	Validation Error

POST /account/{account_id}/assume_role/bedrock

Account Bedrock Models Authentication

Finish a new bedrock authentication flow for an account

Parameters

Name	In	Type	Req	Description
<code>account_id</code>	path	string	yes	

Request body — `application/json` (required)

Field	Type	Req	Default	Description
<code>role_arn</code>	string	yes		ARN of the role created in AWS IAM

Example — a representative request:

```
BASH
curl -X POST
"https://$ZONE.dp.progress.cloud/api/v1/account/$ACCOUNT/assume_role/bedrock" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY" \
-H "Content-Type: application/json" \
-d '{
  "role_arn": "string"
}'
```

Response codes

Status	Schema	Description
200	<code>object</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

DELETE /account/{account_id}/assume_role/bedrock

Account Bedrock Models Authentication

Delete existing bedrock assume role configuration for an account

Parameters

Name	In	Type	Req	Description
------	----	------	-----	-------------

<code>account_id</code>	path	string	yes	

Example — a representative request:

```
BASH
curl -X DELETE
"https://$ZONE.dp.progress.cloud/api/v1/account/$ACCOUNT/assume_role/bedrock" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
200	<code>object</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

GET /account/{account_id}/assume_role/bedrock/validate

Account Bedrock Models Authentication

Validate the bedrock assume role configuration for an account

Parameters

Name	In	Type	Req	Description
<code>account_id</code>	path	string	yes	

Example — a representative request:

```
BASH
curl -X GET
"https://$ZONE.dp.progress.cloud/api/v1/account/$ACCOUNT/assume_role/bedrock/validate" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
200	<code>object</code>	Successful Response

422	HTTPValidationError	Validation Error
-----	---------------------	------------------

GET /account/{account_id}/dataset/{dataset_id}/models

Learning Configuration Schema

Get available models

Parameters

Name	In	Type	Req	Description
account_id	path	string	yes	
dataset_id	path	string	yes	

Example — a representative request:

```
BASH
curl -X GET
"https://$ZONE.dp.progress.cloud/api/v1/account/$ACCOUNT/dataset/{dataset_id}/models" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
200	object (free-form map)	Successful Response
422	HTTPValidationError	Validation Error

GET /account/{account_id}/default_model/{model_id}

Get Public Model Configuration On Account

Get details of a public model configuration on an account

Parameters

Name	In	Type	Req	Description
account_id	path	string	yes	
model_id	path	string	yes	

Example — a representative request:

```
BASH
curl -X GET
"https://$ZONE.dp.progress.cloud/api/v1/account/$ACCOUNT/default_model/
{model_id}" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
200	DefaultModelConfig	Successful Response
422	HTTPValidationError	Validation Error

DELETE /account/{account_id}/default_model/{model_id}

Public Model Configuration Deletion On Account

Delete specific public model configuration from the account

Parameters

Name	In	Type	Req	Description
account_id	path	string	yes	
model_id	path	string	yes	

Example — a representative request:

```
BASH
curl -X DELETE
"https://$ZONE.dp.progress.cloud/api/v1/account/$ACCOUNT/default_model/
{model_id}" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
204		Successful Response
422	HTTPValidationError	Validation Error

GET /account/{account_id}/default_models

List Public Models Configured On An Account

Parameters

Name	In	Type	Req	Description
<code>account_id</code>	path	string	yes	

Example — a representative request:

```
BASH
curl -X GET
"https://$ZONE.dp.progress.cloud/api/v1/account/$ACCOUNT/default_models" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
200	<code>array<DefaultModelConfigListItem></code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

GET /account/{account_id}/model/{model_id}

Learning Configuration Schema

Get trained models

Parameters

Name	In	Type	Req	Description
<code>account_id</code>	path	string	yes	
<code>model_id</code>	path	string	yes	

Example — a representative request:

```
BASH
curl -X GET
"https://$ZONE.dp.progress.cloud/api/v1/account/$ACCOUNT/model/{model_id}" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
200	<code>CustomAccountModel</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

DELETE /account/{account_id}/model/{model_id}

Model Deletion On Account

Delete specific model from account, destroying action

Parameters

Name	In	Type	Req	Description
<code>account_id</code>	path	string	yes	
<code>model_id</code>	path	string	yes	

Example — a representative request:

```
BASH
curl -X DELETE
"https://$ZONE.dp.progress.cloud/api/v1/account/$ACCOUNT/model/{model_id}" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
204		Successful Response
422	<code>HTTPValidationError</code>	Validation Error

GET /account/{account_id}/models

Learning Configuration Models From Account

Get trained models

Parameters

Name	In	Type	Req	Description
------	----	------	-----	-------------

<code>account_id</code>	path	string	yes	
-------------------------	------	--------	-----	--

Example — a representative request:

```
BASH
curl -X GET "https://$ZONE.dp.progress.cloud/api/v1/account/$ACCOUNT/models" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
200	<code>array<CustomAccountModelListItem></code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

DELETE /account/{account_id}/models/{kbid}/{model_id}

Delete A Model From A Kb

Remove from availability a model from a KB

Parameters

Name	In	Type	Req	Description
<code>kbid</code>	path	string	yes	
<code>model_id</code>	path	string	yes	
<code>account_id</code>	path	string	yes	

Example — a representative request:

```
BASH
curl -X DELETE
"https://$ZONE.dp.progress.cloud/api/v1/account/$ACCOUNT/models/$KB/{model_id}" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
204		Successful Response

422	HTTPValidationError	Validation Error
-----	---------------------	------------------

GET /account/{account_id}/schema

Learning Configuration Schema

Get schema for creating a KB

Parameters

Name	In	Type	Req	Description
account_id	path	string	yes	

Example — a representative request:

```
BASH
curl -X GET "https://$ZONE.dp.progress.cloud/api/v1/account/$ACCOUNT/schema" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
200	ConfigSchema	Successful Response
422	HTTPValidationError	Validation Error

GET /models/{kbid}

Learning Configuration Schema

Get available models onprem

Parameters

Name	In	Type	Req	Description
kbid	path	string	yes	

Example — a representative request:

```
BASH
curl -X GET "https://$ZONE.dp.progress.cloud/api/v1/models/$KB" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```


Response codes

Status	Schema	Description
200	<code>object (free-form map)</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

POST /models/{kbid}

Add A Model To A Kb Onprem

Add a model to a KB to be used

Parameters

Name	In	Type	Req	Description
<code>kbid</code>	path	string	yes	

Request body — `application/json` (required)

Field	Type	Req	Default	Description
<code>id</code>	string	yes		

Example — a representative request:

```
BASH
curl -X POST "https://$ZONE.dp.progress.cloud/api/v1/models/$KB" \
  -H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY" \
  -H "Content-Type: application/json" \
  -d '{
    "id": "..."}'
```

Response codes

Status	Schema	Description
200	<code>object</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

GET /models/{kbid}/model/{model_id}

Model Configuration

Get trained models

Parameters

Name	In	Type	Req	Description
kbid	path	string	yes	
model_id	path	string	yes	

Example — a representative request:

```
BASH
curl -X GET "https://$ZONE.dp.progress.cloud/api/v1/models/$KB/model/{model_id}" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
200	object (free-form map)	Successful Response
422	HTTPValidationError	Validation Error

Datasets (10)

GET /account/{account_id}/dataset/{dataset_id}

Get Dataset

Get a dataset

Parameters

Name	In	Type	Req	Description
dataset_id	path	string	yes	
account_id	path	string	yes	

Example — a representative request:

BASH

```
curl -X GET
"https://$ZONE.dp.progress.cloud/api/v1/account/$ACCOUNT/dataset/{dataset_id}" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
200	DatasetCreated	Successful Response
422	HTTPValidationError	Validation Error

DELETE /account/{account_id}/dataset/{dataset_id}

Delete Dataset

Deletes a dataset

Parameters

Name	In	Type	Req	Description
dataset_id	path	string	yes	
account_id	path	string	yes	

Example — a representative request:

BASH

```
curl -X DELETE
"https://$ZONE.dp.progress.cloud/api/v1/account/$ACCOUNT/dataset/{dataset_id}" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
204		Successful Response
422	HTTPValidationError	Validation Error

PUT /account/{account_id}/dataset/{dataset_id}/partition/{partition_id}

Upload Partition Account

Upload a partition of a dataset

Parameters

Name	In	Type	Req	Description
<code>dataset_id</code>	path	string	yes	
<code>account_id</code>	path	string	yes	
<code>partition_id</code>	path	string	yes	
<code>content-length</code>	header	integer	yes	

Example — a representative request:

```
BASH
curl -X PUT
"https://$ZONE.dp.progress.cloud/api/v1/account/$ACCOUNT/dataset/{dataset_id}/
partition/{partition_id}" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
204		Successful Response
422	<code>HTTPValidationError</code>	Validation Error

DELETE /account/{account_id}/dataset/{dataset_id}/partition/{partition_id}

Delete Dataset Partition

Deletes a partition of a dataset

Parameters

Name	In	Type	Req	Description
<code>dataset_id</code>	path	string	yes	
<code>partition_id</code>	path	string	yes	
<code>account_id</code>	path	string	yes	

Example — a representative request:

BASH

```
curl -X DELETE
"https://$ZONE.dp.progress.cloud/api/v1/account/$ACCOUNT/dataset/{dataset_id}/
partition/{partition_id}" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
204		Successful Response
422	<code>HTTPValidationError</code>	Validation Error

DELETE /account/{account_id}/dataset/{dataset_id}/partitions

Delete Dataset Partitions

Deletes all partitions of a dataset

Parameters

Name	In	Type	Req	Description
<code>dataset_id</code>	path	string	yes	
<code>account_id</code>	path	string	yes	

Example — a representative request:

BASH

```
curl -X DELETE
"https://$ZONE.dp.progress.cloud/api/v1/account/$ACCOUNT/dataset/{dataset_id}/
partitions" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
204		Successful Response
422	<code>HTTPValidationError</code>	Validation Error

GET /account/{account_id}/datasets

List Datasets

Lists all datasets

Parameters

Name	In	Type	Req	Description
<code>account_id</code>	path	string	yes	

Example — a representative request:

```
BASH
curl -X GET "https://$ZONE.dp.progress.cloud/api/v1/account/$ACCOUNT/datasets" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
201	<code>DatasetsList</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

POST /account/{account_id}/datasets

Create Dataset Account

Creates a new dataset

Parameters

Name	In	Type	Req	Description
<code>account_id</code>	path	string	yes	

Request body — `application/json` (required)

Field	Type	Req	Default	Description
<code>name</code>	string	yes		
<code>filter</code>	DatasetFilter	yes		
<code>type</code>	Task-Input	yes		

Example — a representative request:

```
BASH
curl -X POST "https://$ZONE.dp.progress.cloud/api/v1/account/$ACCOUNT/datasets" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY" \
-H "Content-Type: application/json" \
-d '{
  "name": "string",
  "filter": {
    "labels": [
      "string"
    ]
  },
  "type": "PARAGRAPH_CLASSIFICATION"
}'
```

Response codes

Status	Schema	Description
201	DatasetCreated	Successful Response
422	HTTPValidationError	Validation Error

DELETE /dataset/{dataset_id}

Dataset Delete

Deletes a dataset

Parameters

Name	In	Type	Req	Description
dataset_id	path	string	yes	

Example — a representative request:

```
BASH
curl -X DELETE "https://$ZONE.dp.progress.cloud/api/v1/dataset/{dataset_id}" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
--------	--------	-------------

204		Successful Response
422	HTTPValidationError	Validation Error

PUT /dataset/{dataset_id}/partition/{partition_id}

Upload Partition

Upload a partition of a dataset

Parameters

Name	In	Type	Req	Description
dataset_id	path	string	yes	
partition_id	path	string	yes	
content-length	header	integer	yes	

Example — a representative request:

```
BASH
curl -X PUT
"https://$ZONE.dp.progress.cloud/api/v1/dataset/{dataset_id}/partition/{partition_id}" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
204		Successful Response
422	HTTPValidationError	Validation Error

POST /datasets

Create Dataset

Creates a new dataset

Request body — `application/json` (required)

Field	Type	Req	Default	Description
-------	------	-----	---------	-------------

<code>name</code>	string	yes		
<code>filter</code>	DatasetFilter	yes		
<code>type</code>	Task-Input	yes		

Example — a representative request:

```
BASH
curl -X POST "https://$ZONE.dp.progress.cloud/api/v1/datasets" \
  -H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY" \
  -H "Content-Type: application/json" \
  -d '{
    "name": "string",
    "filter": {
      "labels": [
        "string"
      ]
    },
    "type": "PARAGRAPH_CLASSIFICATION"
  }'
```

Response codes

Status	Schema	Description
201	<code>DatasetCreated</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

Downloads (3)

DELETE /agent/{agent_id}/audit/download_request/{request_id}

Delete Download Request

Delete a download request by its ID

Parameters

Name	In	Type	Req	Description
<code>agent_id</code>	path	string	yes	
<code>request_id</code>	path	string	yes	

Example — a representative request:

```
BASH
curl -X DELETE
"https://$ZONE.dp.progress.cloud/api/v1/agent/{agent_id}/audit/download_request/{request_id}" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
200	<code>object</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

GET /agent/{agent_id}/audit/download_request/{request_id}/status

Get Download Request Status

Get status of an audit download request by its ID

Parameters

Name	In	Type	Req	Description
<code>agent_id</code>	path	string	yes	
<code>request_id</code>	path	string	yes	

Example — a representative request:

```
BASH
curl -X GET
"https://$ZONE.dp.progress.cloud/api/v1/agent/{agent_id}/audit/download_request/{request_id}/status" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
200	<code>DownloadStatus</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

GET /agent/{agent_id}/audit/download_requests

List Download Requests

List all download requests for the account

Parameters

Name	In	Type	Req	Description
agent_id	path	string	yes	

Example — a representative request:

```
BASH
curl -X GET
"https://$ZONE.dp.progress.cloud/api/v1/agent/{agent_id}/audit/download_requests"
\
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
200	array<DownloadStatus>	Successful Response
422	HTTPValidationError	Validation Error

Retrieval Agent (28)

POST /agent/{agent_id}/audit/interactions/download

Request Interactions Audit Download

Request a download for the audit log of an agent's interactions

Parameters

Name	In	Type	Req	Description
agent_id	path	string	yes	
accept	header	enum(application/x-ndjson, text/csv)		

Request body — application/json (required)

Field	Type	Req	Default	Description
<code>session_id</code>	string null			Filter by session ID
<code>year</code>	integer null			Filter by year (e.g., 2024). If not specified, defaults to the current year.
<code>month</code>	integer null			Filter by month (1-12). If not specified, defaults to the past month.

Example — a representative request:

```
BASH
curl -X POST
"https://$ZONE.dp.progress.cloud/api/v1/agent/{agent_id}/audit/interactions/
download" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY" \
-H "Content-Type: application/json" \
-d '{}'
```

Response codes

Status	Schema	Description
200	<code>DownloadStatus</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

PATCH /agent/{agent_id}/driver/{driver}

Patch Driver

Set Agent Configuration

Parameters

Name	In	Type	Req	Description
<code>agent_id</code>	path	string	yes	
<code>driver</code>	path	string	yes	

Request body — `application/json` (required)

Field	Type	Req	Default	Description
<code>id</code>	string null			
<code>identifier</code>	string	yes		
<code>name</code>	string	yes		
<code>provider</code>	object	yes		The type of driver, e.g., 'google', 'marklogic', etc.
<code>config</code>	EncryptedPayload	yes		The configuration specific to the driver.

Example — a representative request:

```
BASH
curl -X PATCH
"https://$ZONE.dp.progress.cloud/api/v1/agent/{agent_id}/driver/{driver}" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY" \
-H "Content-Type: application/json" \
-d '{
  "identifier": "string",
  "name": "string",
  "provider": null,
  "config": {}
}'
```

Response codes

Status	Schema	Description
200	<code>object</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

DELETE /agent/{agent_id}/driver/{driver}

Delete Driver

Set Agent Configuration

Parameters

Name	In	Type	Req	Description
<code>agent_id</code>	path	string	yes	

driver	path	string	yes	
--------	------	--------	-----	--

Example — a representative request:

```
BASH
curl -X DELETE
"https://$ZONE.dp.progress.cloud/api/v1/agent/{agent_id}/driver/{driver}" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
200	object	Successful Response
422	HTTPValidationError	Validation Error

GET /agent/{agent_id}/drivers

Get Drivers

Get Agent Configuration

Parameters

Name	In	Type	Req	Description
agent_id	path	string	yes	

Example — a representative request:

```
BASH
curl -X GET "https://$ZONE.dp.progress.cloud/api/v1/agent/{agent_id}/drivers" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
200	array<object (free-form map)>	Successful Response
422	HTTPValidationError	Validation Error

POST /agent/{agent_id}/drivers

Add Driver

Set Agent Configuration

Parameters

Name	In	Type	Req	Description
agent_id	path	string	yes	

Request body — `application/json` (required)

Field	Type	Req	Default	Description
id	string null			
identifier	string	yes		
name	string	yes		
provider	object	yes		The type of driver, e.g., 'google', 'marklogic', etc.
config	EncryptedPayload	yes		The configuration specific to the driver.

Example — a representative request:

```
BASH
curl -X POST "https://$ZONE.dp.progress.cloud/api/v1/agent/{agent_id}/drivers" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY" \
-H "Content-Type: application/json" \
-d '{
  "identifier": "string",
  "name": "string",
  "provider": null,
  "config": {}
}'
```

Response codes

Status	Schema	Description
200	DriverID	Successful Response
422	HTTPValidationError	Validation Error

POST /agent/{agent_id}/export

Request Export Retrieval Agent

Request the export of a Retrieval Agent. This will send an email with a download link to the current user.

Parameters

Name	In	Type	Req	Description
agent_id	path	string	yes	

Request body — `application/json` (required)

Field	Type	Req	Default	Description
passphrase	string	yes		Passphrase to encrypt the exported configuration. Will be required for import.

Example — a representative request:

```
BASH
curl -X POST "https://$ZONE.dp.progress.cloud/api/v1/agent/{agent_id}/export" \
  -H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY" \
  -H "Content-Type: application/json" \
  -d '{
    "passphrase": "string"
  }'
```

Response codes

Status	Schema	Description
200	object	Successful Response
422	HTTPValidationError	Validation Error

GET /agent/{agent_id}/generation

Get Generation

Set Agent Configuration

Parameters

Name	In	Type	Req	Description
<code>agent_id</code>	path	string	yes	

Example — a representative request:

```
BASH
curl -X GET "https://$ZONE.dp.progress.cloud/api/v1/agent/{agent_id}/generation" \
  -H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
200	<code>array<object></code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

POST /agent/{agent_id}/generation

Add Generation

Add Generation Agent Configuration

Parameters

Name	In	Type	Req	Description
<code>agent_id</code>	path	string	yes	

Example — a representative request:

```
BASH
curl -X POST "https://$ZONE.dp.progress.cloud/api/v1/agent/{agent_id}/generation" \
  -H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY" \
  -H "Content-Type: application/json" \
  -d '{}'
```

Response codes

Status	Schema	Description
--------	--------	-------------

200	AgentID	Successful Response
422	HTTPValidationError	Validation Error

PATCH /agent/{agent_id}/generation/{generation}

Patch Generation

Set PreProcess Agent Configuration

Parameters

Name	In	Type	Req	Description
agent_id	path	string	yes	
generation	path	string	yes	

Example — a representative request:

```
BASH
curl -X PATCH
"https://$ZONE.dp.progress.cloud/api/v1/agent/{agent_id}/generation/{generation}" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY" \
-H "Content-Type: application/json" \
-d '{}'
```

Response codes

Status	Schema	Description
200	object	Successful Response
422	HTTPValidationError	Validation Error

DELETE /agent/{agent_id}/generation/{generation}

Delete Generation

Delete Generation Agent Configuration

Parameters

Name	In	Type	Req	Description
------	----	------	-----	-------------

<code>agent_id</code>	path	string	yes	
<code>generation</code>	path	string	yes	

Example — a representative request:

```
BASH
curl -X DELETE
"https://$ZONE.dp.progress.cloud/api/v1/agent/{agent_id}/generation/{generation}"
\
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
200	<code>object</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

POST /agent/{agent_id}/import

Import Retrieval Agent

Import Retrieval Agent Configuration

Parameters

Name	In	Type	Req	Description
<code>agent_id</code>	path	string	yes	

Request body — `multipart/form-data` (required)

Field	Type	Req	Default	Description
<code>file</code>	string	yes		
<code>passphrase</code>	string	yes		Passphrase for the encrypted export file
<code>overwrite</code>	boolean		<code>False</code>	Whether to overwrite existing agent configuration. If false

				and the agent configuration is not empty, the import will fail.
--	--	--	--	---

Example — a representative request:

BASH

```
curl -X POST "https://$ZONE.dp.progress.cloud/api/v1/agent/{agent_id}/import" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
200	<code>object</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

GET /agent/{agent_id}/postprocess

Get Postprocess

Set Agent Configuration

Parameters

Name	In	Type	Req	Description
<code>agent_id</code>	path	string	yes	

Example — a representative request:

BASH

```
curl -X GET "https://$ZONE.dp.progress.cloud/api/v1/agent/{agent_id}/postprocess" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
200	<code>array<object></code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

POST /agent/{agent_id}/postprocess

Add Postprocess

Add PostProcess Agent Configuration

Parameters

Name	In	Type	Req	Description
<code>agent_id</code>	path	string	yes	

Example — a representative request:

```
BASH
curl -X POST
"https://$ZONE.dp.progress.cloud/api/v1/agent/{agent_id}/postprocess" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY" \
-H "Content-Type: application/json" \
-d '{}'
```

Response codes

Status	Schema	Description
200	<code>AgentID</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

PATCH /agent/{agent_id}/postprocess/{postprocess}

Patch Postprocess

Set PostProcess Agent Configuration

Parameters

Name	In	Type	Req	Description
<code>agent_id</code>	path	string	yes	
<code>postprocess</code>	path	string	yes	

Example — a representative request:

BASH

```
curl -X PATCH
"https://$ZONE.dp.progress.cloud/api/v1/agent/{agent_id}/postprocess/
{postprocess}" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY" \
-H "Content-Type: application/json" \
-d '{}'
```

Response codes

Status	Schema	Description
200	<code>object</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

DELETE /agent/{agent_id}/postprocess/{postprocess}

Delete Postprocess

Delete PreProcess Agent Configuration

Parameters

Name	In	Type	Req	Description
<code>agent_id</code>	path	string	yes	
<code>postprocess</code>	path	string	yes	

Example — a representative request:**BASH**

```
curl -X DELETE
"https://$ZONE.dp.progress.cloud/api/v1/agent/{agent_id}/postprocess/
{postprocess}" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
200	<code>object</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

GET /agent/{agent_id}/preprocess

Get Preprocess

Set Agent Configuration

Parameters

Name	In	Type	Req	Description
agent_id	path	string	yes	

Example — a representative request:

```
BASH
curl -X GET "https://$ZONE.dp.progress.cloud/api/v1/agent/{agent_id}/preprocess" \
  -H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
200	array<object>	Successful Response
422	HTTPValidationError	Validation Error

POST /agent/{agent_id}/preprocess

Add Preprocess

Add PreProcess Agent Configuration

Parameters

Name	In	Type	Req	Description
agent_id	path	string	yes	

Example — a representative request:

```
BASH
curl -X POST "https://$ZONE.dp.progress.cloud/api/v1/agent/{agent_id}/preprocess" \
  -H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY" \
  -H "Content-Type: application/json" \
  -d '{}'
```

Response codes

Status	Schema	Description
200	AgentID	Successful Response
422	HTTPValidationError	Validation Error

PATCH /agent/{agent_id}/preprocess/{preprocess}

Patch Preprocess

Set PreProcess Agent Configuration

Parameters

Name	In	Type	Req	Description
agent_id	path	string	yes	
preprocess	path	string	yes	

Example — a representative request:

```
BASH
curl -X PATCH
"https://$ZONE.dp.progress.cloud/api/v1/agent/{agent_id}/preprocess/{preprocess}" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY" \
-H "Content-Type: application/json" \
-d '{}'
```

Response codes

Status	Schema	Description
200	object	Successful Response
422	HTTPValidationError	Validation Error

DELETE /agent/{agent_id}/preprocess/{preprocess}

Delete Preprocess

Delete PreProcess Agent Configuration

Parameters

Name	In	Type	Req	Description
<code>agent_id</code>	path	string	yes	
<code>preprocess</code>	path	string	yes	

Example — a representative request:

```
BASH
curl -X DELETE
"https://$ZONE.dp.progress.cloud/api/v1/agent/{agent_id}/preprocess/{preprocess}" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
200	<code>object</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

GET /agent/{agent_id}/rules

Get Rules

Set Agent Configuration

Parameters

Name	In	Type	Req	Description
<code>agent_id</code>	path	string	yes	

Example — a representative request:

```
BASH
curl -X GET "https://$ZONE.dp.progress.cloud/api/v1/agent/{agent_id}/rules" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
--------	--------	-------------

200	<code>object</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

POST /agent/{agent_id}/rules

Set Rules

Set Agent Configuration

Parameters

Name	In	Type	Req	Description
<code>agent_id</code>	path	string	yes	

Request body — `application/json` (required)

Field	Type	Req	Default	Description
<code>rules</code>	<code>array<Rule string></code>			List of rules that the workflow should follow. Each rule can be a string or a Rule object with a prompt.

Example — a representative request:

BASH

```
curl -X POST "https://$ZONE.dp.progress.cloud/api/v1/agent/{agent_id}/rules" \
  -H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY" \
  -H "Content-Type: application/json" \
  -d '{}'
```

Response codes

Status	Schema	Description
200	<code>object</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

GET /agent/{agent_id}/session/{session}

Get Session

Create session

Parameters

Name	In	Type	Req	Description
agent_id	path	string	yes	
session	path	string	yes	

Example — a representative request:

```
BASH
curl -X GET
"https://$ZONE.dp.progress.cloud/api/v1/agent/{agent_id}/session/{session}" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
200	Resource-Output	Successful Response
422	HTTPValidationError	Validation Error

POST /agent/{agent_id}/session/{session}

Interaction

Interact session

Parameters

Name	In	Type	Req	Description
agent_id	path	string	yes	
session	path	string	yes	
workflow_id	query	string		

Request body — `application/json` (required)

Field	Type	Req	Default	Description
-------	------	-----	---------	-------------

<code>question</code>	string	yes		
<code>headers</code>	map<string, string>		<code>{}</code>	
<code>arguments</code>	map<string, string>		<code>{}</code>	
<code>chat_history</code>	array<HistoryQuestionAnswer> null			Client-managed chat history. When set (even to an empty list), overrides any server-side session history for agents that use previous Q&A context (rephrase, summarize, smart, etc.). Omit the field ent...
<code>operation</code>	InteractionOperation		<code>0</code>	
<code>streaming</code>	boolean		<code>False</code>	

Example — a representative request:

```
BASH
curl -X POST
"https://$ZONE.dp.progress.cloud/api/v1/agent/{agent_id}/session/{session}" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY" \
-H "Content-Type: application/json" \
-d '{
  "question": "your question here",
  "operation": 0,
  "streaming": false
}'
```

Response codes

Status	Schema	Description
200	<code>object</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

PATCH /agent/{agent_id}/session/{session}

Patch Session

Create session

Parameters

Name	In	Type	Req	Description
<code>agent_id</code>	path	string	yes	
<code>session</code>	path	string	yes	

Request body — `application/json` (required)

Field	Type	Req	Default	Description
<code>slug</code>	string	yes		
<code>name</code>	string	yes		
<code>summary</code>	string	yes		
<code>data</code>	string	yes		
<code>format</code>	TextFormat	yes		

Example — a representative request:

```
BASH
curl -X PATCH
"https://$ZONE.dp.progress.cloud/api/v1/agent/{agent_id}/session/{session}" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY" \
-H "Content-Type: application/json" \
-d '{
  "slug": "my-resource",
  "name": "string",
  "summary": "string",
  "data": "string",
  "format": "PLAIN"
}'
```

Response codes

Status	Schema	Description
200	<code>object</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

DELETE /agent/{agent_id}/session/{session}

Delete Session

Parameters

Name	In	Type	Req	Description
agent_id	path	string	yes	
session	path	string	yes	

Example — a representative request:

```
BASH
curl -X DELETE
"https://$ZONE.dp.progress.cloud/api/v1/agent/{agent_id}/session/{session}" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
200	object	Successful Response
422	HTTPValidationError	Validation Error

GET /agent/{agent_id}/sessions

Get Sessions

Create session

Parameters

Name	In	Type	Req	Description
agent_id	path	string	yes	
page	query	integer		Requested page number (0-based)
size	query	integer		Page size

Example — a representative request:

```
BASH
```

```
curl -X GET "https://$ZONE.dp.progress.cloud/api/v1/agent/{agent_id}/sessions" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
200	ResourceList	Successful Response
422	HTTPValidationError	Validation Error

POST /agent/{agent_id}/sessions

Create Session

Parameters

Name	In	Type	Req	Description
agent_id	path	string	yes	

Request body — [application/json](#) (required)

Field	Type	Req	Default	Description
slug	string	yes		
name	string	yes		
summary	string	yes		
data	string	yes		
format	TextFormat	yes		

Example — a representative request:

```
BASH
curl -X POST "https://$ZONE.dp.progress.cloud/api/v1/agent/{agent_id}/sessions" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY" \
-H "Content-Type: application/json" \
-d '{
  "slug": "my-resource",
  "name": "string",
  "summary": "string",
  "data": "string",
  "format": "PLAIN"
}
```

```
}'
```

Response codes

Status	Schema	Description
200	ResourceCreated	Successful Response
422	HTTPValidationError	Validation Error

POST /agent/{agent_id}/workflow/{workflow_id}/session/{session}

Interaction

Interact session

Parameters

Name	In	Type	Req	Description
agent_id	path	string	yes	
session	path	string	yes	
workflow_id	path	string	yes	

Request body — application/json (required)

Field	Type	Req	Default	Description
question	string	yes		
headers	map<string, string>		{}	
arguments	map<string, string>		{}	
chat_history	array<HistoryQuestionAnswer> null			Client-managed chat history. When set (even to an empty list), overrides any server-side session history for agents that use previous Q&A context (rephrase, summarize, smart, etc.). Omit the field ent...

<code>operation</code>	InteractionOperation		<code>0</code>	
<code>streaming</code>	boolean		<code>False</code>	

Example — a representative request:

BASH

```
curl -X POST "https://$ZONE.dp.progress.cloud/api/v1/agent/{agent_id}/workflow/{workflow_id}/session/{session}" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY" \
-H "Content-Type: application/json" \
-d '{
  "question": "your question here",
  "operation": 0,
  "streaming": false
}'
```

Response codes

Status	Schema	Description
200	<code>object</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

Prompt Management (5)

GET /agent/{agent_id}/prompt/{prompt_id}

Get Prompt

Get Agent Configuration

Parameters

Name	In	Type	Req	Description
<code>agent_id</code>	path	string	yes	
<code>prompt_id</code>	path	string	yes	

Example — a representative request:

BASH

```
curl -X GET
```

```
"https://$ZONE.dp.progress.cloud/api/v1/agent/{agent_id}/prompt/{prompt_id}" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
200	PromptConfig	Successful Response
422	HTTPValidationError	Validation Error

PATCH /agent/{agent_id}/prompt/{prompt_id}

Patch Prompt

Update Prompt of Agent

Parameters

Name	In	Type	Req	Description
agent_id	path	string	yes	
prompt_id	path	string	yes	

Request body — application/json (required)

Field	Type	Req	Default	Description
name	string	yes		
description	string	yes		
prompt	string	yes		
arguments	array<PromptArgument> null			
icons	map<string, string> null			
meta	map<string, string> null			
prompt_id	string null			

Example — a representative request:

BASH

```
curl -X PATCH
"https://$ZONE.dp.progress.cloud/api/v1/agent/{agent_id}/prompt/{prompt_id}" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY" \
-H "Content-Type: application/json" \
-d '{
  "name": "string",
  "description": "string",
  "prompt": "string"
}'
```

Response codes

Status	Schema	Description
200	<code>object</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

DELETE /agent/{agent_id}/prompt/{prompt_id}

Delete Prompt

Delete prompt id

Parameters

Name	In	Type	Req	Description
<code>agent_id</code>	path	string	yes	
<code>prompt_id</code>	path	string	yes	

Example — a representative request:**BASH**

```
curl -X DELETE
"https://$ZONE.dp.progress.cloud/api/v1/agent/{agent_id}/prompt/{prompt_id}" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
200	<code>object</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

GET /agent/{agent_id}/prompts

Get Prompts

Get Agent Configuration

Parameters

Name	In	Type	Req	Description
agent_id	path	string	yes	

Example — a representative request:

BASH

```
curl -X GET "https://$ZONE.dp.progress.cloud/api/v1/agent/{agent_id}/prompts" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
200	array<PromptConfig>	Successful Response
422	HTTPValidationError	Validation Error

POST /agent/{agent_id}/prompts

Add Prompt

Add Prompt to Agent

Parameters

Name	In	Type	Req	Description
agent_id	path	string	yes	

Request body — `application/json` (required)

Field	Type	Req	Default	Description
name	string	yes		

description	string	yes		
prompt	string	yes		
arguments	array<PromptArgument> null			
icons	map<string, string> null			
meta	map<string, string> null			
prompt_id	string null			

Example — a representative request:

```
BASH
curl -X POST "https://$ZONE.dp.progress.cloud/api/v1/agent/{agent_id}/prompts" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY" \
-H "Content-Type: application/json" \
-d '{
  "name": "string",
  "description": "string",
  "prompt": "string"
}'
```

Response codes

Status	Schema	Description
200	PromptID	Successful Response
422	HTTPValidationError	Validation Error

General (5)

GET /agent/{agent_id}/session/{session}/mcp

Interaction Mcp Handler

Parameters

Name	In	Type	Req	Description
agent_id	path	string	yes	

<code>session</code>	path	string	yes	
----------------------	------	--------	-----	--

Example — a representative request:

```
BASH
curl -X GET
"https://$ZONE.dp.progress.cloud/api/v1/agent/{agent_id}/session/{session}/mcp" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
200	<code>object</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

POST /agent/{agent_id}/session/{session}/mcp

Interaction Mcp Handler

Parameters

Name	In	Type	Req	Description
<code>agent_id</code>	path	string	yes	
<code>session</code>	path	string	yes	

Example — a representative request:

```
BASH
curl -X POST
"https://$ZONE.dp.progress.cloud/api/v1/agent/{agent_id}/session/{session}/mcp" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
200	<code>object</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

DELETE /agent/{agent_id}/session/{session}/mcp

Mcp Handler Delete

Parameters

Name	In	Type	Req	Description
<code>agent_id</code>	path	string	yes	
<code>session</code>	path	string	yes	

Example — a representative request:

```
BASH
curl -X DELETE
"https://$ZONE.dp.progress.cloud/api/v1/agent/{agent_id}/session/{session}/mcp" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
200	<code>object</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

GET /kb/{kbid}/mcp

Mcp Handler

Parameters

Name	In	Type	Req	Description
<code>kbid</code>	path	string	yes	
<code>x-ndb-client</code>	header	NucliaDBClientType		
<code>x-nucliadb-user</code>	header	string		
<code>x-forwarded-for</code>	header	string		

Example — a representative request:

BASH

```
curl -X GET "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/mcp" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
200	<code>object</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

POST /kb/{kbid}/mcp

Mcp Handler

Parameters

Name	In	Type	Req	Description
<code>kbid</code>	path	string	yes	
<code>x-ndb-client</code>	header	NucliaDBClientType		
<code>x-nucliadb-user</code>	header	string		
<code>x-forwarded-for</code>	header	string		

Example — a representative request:

BASH

```
curl -X POST "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/mcp" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
200	<code>object</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

Workflows (18)

PATCH /agent/{agent_id}/workflow/{workflow_id}

Set Workflow

Set Workflow Configuration

Parameters

Name	In	Type	Req	Description
agent_id	path	string	yes	
workflow_id	path	string	yes	

Request body — `application/json` (required)

Field	Type	Req	Default	Description
name	string	yes		
description	string	yes		
parameters	object (free-form map)	yes		
required	array<string>			
rules	Rules null			

Example — a representative request:

```
BASH
curl -X PATCH "https://$ZONE.dp.progress.cloud/api/v1/agent/{agent_id}/workflow/{workflow_id}" \
  -H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY" \
  -H "Content-Type: application/json" \
  -d '{
    "name": "string",
    "description": "string",
    "parameters": {},
    "required": [
      "string"
    ]
  }'
```

Response codes

Status	Schema	Description
200	<code>object</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

DELETE /agent/{agent_id}/workflow/{workflow_id}

Delete Workflow

Delete Workflow Configuration

Parameters

Name	In	Type	Req	Description
<code>agent_id</code>	path	string	yes	
<code>workflow_id</code>	path	string	yes	

Example — a representative request:

```
BASH
curl -X DELETE
"https://$ZONE.dp.progress.cloud/api/v1/agent/{agent_id}/workflow/{workflow_id}"
\
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
200	<code>object</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

GET /agent/{agent_id}/workflow/{workflow_id}/generation

Get Generation

Get Generation Workflows Configuration

Parameters

Name	In	Type	Req	Description
------	----	------	-----	-------------

<code>agent_id</code>	path	string	yes	
<code>workflow_id</code>	path	string	yes	

Example — a representative request:

```
BASH
curl -X GET "https://$ZONE.dp.progress.cloud/api/v1/agent/{agent_id}/workflow/{workflow_id}/generation" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
200	<code>array<object></code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

POST /agent/{agent_id}/workflow/{workflow_id}/generation

Add Generation

Add Generation Workflows Configuration

Parameters

Name	In	Type	Req	Description
<code>agent_id</code>	path	string	yes	
<code>workflow_id</code>	path	string	yes	

Example — a representative request:

```
BASH
curl -X POST "https://$ZONE.dp.progress.cloud/api/v1/agent/{agent_id}/workflow/{workflow_id}/generation" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY" \
-H "Content-Type: application/json" \
-d '{}'
```

Response codes

Status	Schema	Description
--------	--------	-------------

200	AgentID	Successful Response
422	HTTPValidationError	Validation Error

PATCH /agent/{agent_id}/workflow/{workflow_id}/generation/{generation}

Patch Generation

Set Generation Workflows Configuration

Parameters

Name	In	Type	Req	Description
agent_id	path	string	yes	
workflow_id	path	string	yes	
generation	path	string	yes	

Example — a representative request:

BASH

```
curl -X PATCH "https://$ZONE.dp.progress.cloud/api/v1/agent/{agent_id}/workflow/{workflow_id}/generation/{generation}" \
  -H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY" \
  -H "Content-Type: application/json" \
  -d '{}'
```

Response codes

Status	Schema	Description
200	object	Successful Response
422	HTTPValidationError	Validation Error

DELETE /agent/{agent_id}/workflow/{workflow_id}/generation/{generation}

Delete Generation

Delete Generation Workflows Configuration

Parameters

Name	In	Type	Req	Description
------	----	------	-----	-------------

<code>agent_id</code>	path	string	yes	
<code>workflow_id</code>	path	string	yes	
<code>generation</code>	path	string	yes	

Example — a representative request:

```
BASH
curl -X DELETE
"https://$ZONE.dp.progress.cloud/api/v1/agent/{agent_id}/workflow/{workflow_id}/
generation/{generation}" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
200	<code>object</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

GET /agent/{agent_id}/workflow/{workflow_id}/postprocess

Get Postprocess

Get PostProcess Workflows Configuration

Parameters

Name	In	Type	Req	Description
<code>agent_id</code>	path	string	yes	
<code>workflow_id</code>	path	string	yes	

Example — a representative request:

```
BASH
curl -X GET "https://$ZONE.dp.progress.cloud/api/v1/agent/{agent_id}/workflow/
{workflow_id}/postprocess" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
--------	--------	-------------

200	<code>array<object></code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

POST /agent/{agent_id}/workflow/{workflow_id}/postprocess

Add Postprocess

Add PostProcess Workflows Configuration

Parameters

Name	In	Type	Req	Description
<code>agent_id</code>	path	string	yes	
<code>workflow_id</code>	path	string	yes	

Example — a representative request:

BASH

```
curl -X POST "https://$ZONE.dp.progress.cloud/api/v1/agent/{agent_id}/workflow/{workflow_id}/postprocess" \
  -H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY" \
  -H "Content-Type: application/json" \
  -d '{}'
```

Response codes

Status	Schema	Description
200	<code>AgentID</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

PATCH /agent/{agent_id}/workflow/{workflow_id}/postprocess/{postprocess}

Patch Postprocess

Set PostProcess Workflows Configuration

Parameters

Name	In	Type	Req	Description
------	----	------	-----	-------------

<code>agent_id</code>	path	string	yes	
<code>workflow_id</code>	path	string	yes	
<code>postprocess</code>	path	string	yes	

Example — a representative request:

```
BASH
curl -X PATCH "https://$ZONE.dp.progress.cloud/api/v1/agent/{agent_id}/workflow/{workflow_id}/postprocess/{postprocess}" \
  -H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY" \
  -H "Content-Type: application/json" \
  -d '{}'
```

Response codes

Status	Schema	Description
200	<code>object</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

DELETE /agent/{agent_id}/workflow/{workflow_id}/postprocess/{postprocess}

Delete Postprocess

Delete PostProcess Workflows Configuration

Parameters

Name	In	Type	Req	Description
<code>agent_id</code>	path	string	yes	
<code>workflow_id</code>	path	string	yes	
<code>postprocess</code>	path	string	yes	

Example — a representative request:

```
BASH
curl -X DELETE
"https://$ZONE.dp.progress.cloud/api/v1/agent/{agent_id}/workflow/{workflow_id}/postprocess/{postprocess}" \
  -H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
200	<code>object</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

GET /agent/{agent_id}/workflow/{workflow_id}/preprocess

Get Preprocess

Set Workflow Configuration

Parameters

Name	In	Type	Req	Description
<code>agent_id</code>	path	string	yes	
<code>workflow_id</code>	path	string	yes	

Example — a representative request:

```
BASH
curl -X GET "https://$ZONE.dp.progress.cloud/api/v1/agent/{agent_id}/workflow/{workflow_id}/preprocess" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
200	<code>array<object></code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

POST /agent/{agent_id}/workflow/{workflow_id}/preprocess

Add Preprocess

Add PreProcess Workflows Configuration

Parameters

Name	In	Type	Req	Description
------	----	------	-----	-------------

<code>agent_id</code>	path	string	yes	
<code>workflow_id</code>	path	string	yes	

Example — a representative request:

BASH

```
curl -X POST "https://$ZONE.dp.progress.cloud/api/v1/agent/{agent_id}/workflow/{workflow_id}/preprocess" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY" \
-H "Content-Type: application/json" \
-d '{}'
```

Response codes

Status	Schema	Description
200	<code>AgentID</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

PATCH /agent/{agent_id}/workflow/{workflow_id}/preprocess/{preprocess}

Patch Preprocess

Set PreProcess Workflows Configuration

Parameters

Name	In	Type	Req	Description
<code>agent_id</code>	path	string	yes	
<code>workflow_id</code>	path	string	yes	
<code>preprocess</code>	path	string	yes	

Example — a representative request:

BASH

```
curl -X PATCH "https://$ZONE.dp.progress.cloud/api/v1/agent/{agent_id}/workflow/{workflow_id}/preprocess/{preprocess}" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY" \
-H "Content-Type: application/json" \
-d '{}'
```

Response codes

Status	Schema	Description
200	<code>object</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

DELETE /agent/{agent_id}/workflow/{workflow_id}/preprocess/{preprocess}

Delete Preprocess

Delete PreProcess Workflows Configuration

Parameters

Name	In	Type	Req	Description
<code>agent_id</code>	path	string	yes	
<code>preprocess</code>	path	string	yes	
<code>workflow_id</code>	path	string	yes	

Example — a representative request:

```
BASH
curl -X DELETE
"https://$ZONE.dp.progress.cloud/api/v1/agent/{agent_id}/workflow/{workflow_id}/
preprocess/{preprocess}" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
200	<code>object</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

GET /agent/{agent_id}/workflow/{workflow_id}/rules

Get Rules

Get Workflow Rules Configuration

Parameters

Name	In	Type	Req	Description
------	----	------	-----	-------------

<code>agent_id</code>	path	string	yes	
<code>workflow_id</code>	path	string	yes	

Example — a representative request:

```
BASH
curl -X GET "https://$ZONE.dp.progress.cloud/api/v1/agent/{agent_id}/workflow/{workflow_id}/rules" \
  -H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
200	<code>object</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

POST /agent/{agent_id}/workflow/{workflow_id}/rules

Set Rules

Set Workflow Rules Configuration

Parameters

Name	In	Type	Req	Description
<code>agent_id</code>	path	string	yes	
<code>workflow_id</code>	path	string	yes	

Request body — `application/json` (required)

Field	Type	Req	Default	Description
<code>rules</code>	<code>array<Rule string></code>			List of rules that the workflow should follow. Each rule can be a string or a Rule object with a prompt.

Example — a representative request:

```
BASH
curl -X POST "https://$ZONE.dp.progress.cloud/api/v1/agent/{agent_id}/workflow/{workflow_id}/rules" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY" \
-H "Content-Type: application/json" \
-d '{}'
```

Response codes

Status	Schema	Description
200	<code>object</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

GET /agent/{agent_id}/workflows

Get Workflows

Get Workflow Configuration

Parameters

Name	In	Type	Req	Description
<code>agent_id</code>	path	string	yes	

Example — a representative request:

```
BASH
curl -X GET "https://$ZONE.dp.progress.cloud/api/v1/agent/{agent_id}/workflows" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
200	<code>array<WorkflowData></code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

POST /agent/{agent_id}/workflows

Add Workflow

Add Workflow Configuration

Parameters

Name	In	Type	Req	Description
<code>agent_id</code>	path	string	yes	

Request body — `application/json` (required)

Field	Type	Req	Default	Description
<code>id</code>	string	yes		
<code>name</code>	string	yes		
<code>description</code>	string null			
<code>parameters</code>	object (free-form map) null			
<code>rules</code>	Rules			
<code>required</code>	array<string>			

Example — a representative request:

```
BASH
curl -X POST "https://$ZONE.dp.progress.cloud/api/v1/agent/{agent_id}/workflows" \
  -H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY" \
  -H "Content-Type: application/json" \
  -d '{
    "id": "...",
    "name": "string",
    "required": [
      "string"
    ]
  }'
```

Response codes

Status	Schema	Description
200	<code>object</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

Collector (2)

GET /collect/feedback/{kbid}

Feedback Available Months

List of months within the last year with feedback data

Parameters

Name	In	Type	Req	Description
kbid	path	string	yes	

Example — a representative request:

BASH

```
curl -X GET "https://$ZONE.dp.progress.cloud/api/v1/collect/feedback/$KB" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
200	object	Successful Response
422	HTTPValidationError	Validation Error

GET /collect/feedback/{kbid}/{month}

Get Feedback

Given a month, returns a csv with all the feedback data from it

Parameters

Name	In	Type	Req	Description
kbid	path	string	yes	
month	path	string	yes	

Example — a representative request:

BASH

```
curl -X GET "https://$ZONE.dp.progress.cloud/api/v1/collect/feedback/$KB/{month}" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
200	<code>object</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

Knowledge Boxes (8)

GET /config/{kbid}

Knowledgebox Models Configuration

Current configuration of models assigned to a knowledgebox

Parameters

Name	In	Type	Req	Description
<code>kbid</code>	path	string	yes	

Example — a representative request:

```
BASH
curl -X GET "https://$ZONE.dp.progress.cloud/api/v1/config/$KB" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
200	<code>LearningConfiguration</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

POST /config/{kbid}

Knowledgebox Models Add Configuration

Current configuration of models assigned to a knowledgebox

Parameters

Name	In	Type	Req	Description
<code>kbid</code>	path	string	yes	

Request body — `application/json` (required)

Field	Type	Req	Default	Description
<code>anonymization_model</code>	AnonymizationModel string		<code>disabled</code>	
<code>visual_labeling</code>	VisualLabeling string		<code>disabled</code>	
<code>generative_model</code>	string		<code>chatgpt-azure-4o</code>	
<code>ner_model</code>	NERModel string		<code>multilingual</code>	
<code>relation_model</code>	RelationModel string		<code>base</code>	
<code>user_keys</code>	UserLearningKeys null			
<code>user_prompts</code>	UserPrompts null			
<code>summary</code>	Summary null		<code>simple</code>	
<code>summary_model</code>	string		<code>chatgpt-azure-4o</code>	
<code>summary_prompt</code>	SummaryPrompt null			
<code>prefer_markdown_generative_response</code>	boolean		<code>False</code>	
<code>default_semantic_model</code>	string null			The default vector set / semantic model to use. Must be one of the models defined in the <code>semanticmodels</code> field. If not provided, the first model in the list will be used.

<code>semantic_model</code>	ParagraphEmbeddingModel null			
<code>semantic_models</code>	array<ParagraphEmbeddingModel> null			
<code>semantic_graph_node_models</code>	array<NodeEmbeddingModel> null			
<code>default_semantic_graph_node_model</code>	string null			
<code>semantic_graph_edge_models</code>	array<EdgeEmbeddingModel> null			
<code>default_semantic_graph_edge_model</code>	string null			

Example — a representative request:

```
BASH
curl -X POST "https://$ZONE.dp.progress.cloud/api/v1/config/$KB" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY" \
-H "Content-Type: application/json" \
-d '{
  "generative_model": "chatgpt-azure-4o",
  "summary_model": "chatgpt-azure-4o",
  "prefer_markdown_generative_response": false
}'
```

Response codes

Status	Schema	Description
200	<code>LearningConfiguration</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

PATCH /config/{kbid}

Knowledgebox Models Update Configuration

Current configuration of models assigned to a knowledgebox

Parameters

Name	In	Type	Req	Description
------	----	------	-----	-------------

kbid	path	string	yes	

Request body — `application/json` (required)

Field	Type	Req	Default	Description
anonymization_model	AnonymizationModel string		disabled	
visual_labeling	VisualLabeling string		disabled	
generative_model	string		chatgpt-azure-4o	
ner_model	NERModel string		multilingual	
relation_model	RelationModel string		base	
user_keys	UserLearningKeys null			
user_prompts	UserPrompts null			
summary	Summary null		simple	
summary_model	string		chatgpt-azure-4o	
summary_prompt	SummaryPrompt null			
prefer_markdown_generative_response	boolean		False	
default_semantic_model	string null			The default vector set / semantic model to use. Must be one of the models defined in the semanticmodels field. If not provided, the first model in the list will be used.

Example — a representative request:

BASH

```
curl -X PATCH "https://$ZONE.dp.progress.cloud/api/v1/config/$KB" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY" \
-H "Content-Type: application/json" \
-d '{
  "generative_model": "chatgpt-azure-4o",
  "summary_model": "chatgpt-azure-4o",
  "prefer_markdown_generative_response": false
}'
```

Response codes

Status	Schema	Description
200	LearningConfiguration	Successful Response
422	HTTPValidationError	Validation Error

DELETE /config/{kbid}

Knowledgebox Models Configuration

Delete a KB configuration

Parameters

Name	In	Type	Req	Description
kbid	path	string	yes	

Example — a representative request:

BASH

```
curl -X DELETE "https://$ZONE.dp.progress.cloud/api/v1/config/$KB" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
204		Successful Response
422	HTTPValidationError	Validation Error

GET /generative_providers/{kbid}

Available Models For A Knowledge Box

Get all available models for a knowledge box grouped by provider

Parameters

Name	In	Type	Req	Description
kbid	path	string	yes	

Example — a representative request:

```
BASH
curl -X GET "https://$ZONE.dp.progress.cloud/api/v1/generative_providers/$KB" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
200	ProvidersResponse	Successful Response
422	HTTPValidationError	Validation Error

GET /learning/configuration/schema

Learning Configuration Schema

Get jsonschema definition for learningconfiguration field of knowledgebox creation payload

Example — a representative request:

```
BASH
curl -X GET
"https://$ZONE.dp.progress.cloud/api/v1/learning/configuration/schema" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
200	ConfigSchema	Successful Response
422	HTTPValidationError	Validation Error

GET /schema

Learning Configuration Schema

Get jsonschema definition for learningconfiguration field of knowledgebox creation payload

Example — a representative request:

BASH

```
curl -X GET "https://$ZONE.dp.progress.cloud/api/v1/schema" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
200	ConfigSchema	Successful Response
422	HTTPValidationError	Validation Error

GET /schema/{kbid}

Learning Configuration Schema

Get jsonschema definition for learningconfiguration field of knowledgebox update payload

Parameters

Name	In	Type	Req	Description
kbid	path	string	yes	

Example — a representative request:

BASH

```
curl -X GET "https://$ZONE.dp.progress.cloud/api/v1/schema/$KB" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
200	ConfigSchema	Successful Response
422	HTTPValidationError	Validation Error

Ingestion Agents and Tasks (15)

POST /dataset/{dataset_id}/task/start

Start Dataset Task

Start an Ingestion Agent or a Task on a Dataset

Parameters

Name	In	Type	Req	Description
<code>dataset_id</code>	path	string	yes	

Request body — `application/json` (required)

Field	Type	Req	Default	Description
<code>name</code>	TaskName	yes		
<code>parameters</code>	DataAugmentation-Input SemanticModelMigrationParams null			Parameters to be passed to the task. These must match the validation field for the Task definition class
<code>uuid_task</code>	string null			UUID of an already configured task. This is used to start a task that was already configured

Example — a representative request:

BASH

```
curl -X POST "https://$ZONE.dp.progress.cloud/api/v1/dataset/{dataset_id}/task/start" \
  -H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY" \
  -H "Content-Type: application/json" \
  -d '{
    "name": "dummy"
  }'
```

Response codes

Status	Schema	Description
--------	--------	-------------

200	<code>TaskResponse</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

DELETE /dataset/{dataset_id}/task/{task_id}

Delete Dataset Task

Delete an Ingestion Agent or a Task of a Dataset

Parameters

Name	In	Type	Req	Description
<code>dataset_id</code>	path	string	yes	
<code>task_id</code>	path	string	yes	

Example — a representative request:

```
BASH
curl -X DELETE
"https://$ZONE.dp.progress.cloud/api/v1/dataset/{dataset_id}/task/{task_id}" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
200	<code>object</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

GET /dataset/{dataset_id}/task/{task_id}/inspect

Inspect Dataset Task

Get the status of an Ingestion Agent or a Task of a Dataset

Parameters

Name	In	Type	Req	Description
<code>dataset_id</code>	path	string	yes	

<code>task_id</code>	path	string	yes	
----------------------	------	--------	-----	--

Example — a representative request:

```
BASH
curl -X GET "https://$ZONE.dp.progress.cloud/api/v1/dataset/{dataset_id}/task/{task_id}/inspect" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
200	<code>PublicTaskRequest</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

POST /dataset/{dataset_id}/task/{task_id}/stop

Stop Dataset Task

Stop an Ingestion Agent or a Task of a Dataset

Parameters

Name	In	Type	Req	Description
<code>dataset_id</code>	path	string	yes	
<code>task_id</code>	path	string	yes	

Example — a representative request:

```
BASH
curl -X POST "https://$ZONE.dp.progress.cloud/api/v1/dataset/{dataset_id}/task/{task_id}/stop" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
200	<code>TaskResponse</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

GET /dataset/{dataset_id}/tasks

List Account Available Tasks

Get the available tasks for a Dataset

Parameters

Name	In	Type	Req	Description
<code>dataset_id</code>	path	string	yes	
<code>count</code>	query	integer		How many finished tasks to return

Example — a representative request:

BASH

```
curl -X GET "https://$ZONE.dp.progress.cloud/api/v1/dataset/{dataset_id}/tasks" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
200	<code>TaskList</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

POST /kb/{kb_uuid}/task/start

Start Knowledge Box Task

Start an Ingestion Agent or a Task on a Knowledge Box

Parameters

Name	In	Type	Req	Description
<code>kb_uuid</code>	path	string	yes	

Request body — `application/json` (required)

Field	Type	Req	Default	Description
-------	------	-----	---------	-------------

name	TaskName	yes		
parameters	DataAugmentation-Input SemanticModelMigrationParams null			Parameters to be passed to the task. These must match the validation field for the Task definition class
uuid_task	string null			UUID of an already configured task. This is used to start a task that was already configured
apply	ApplyOptions		ALL	Defines how the tasks should be applied to the existing data. - EXSITING: Only apply to existing data (starts a worker that executes the task) - NEW: Only apply to new data (enables the task at proces...
enabled	boolean		True	Whether the task should be enabled at the time of creation. This only applies to tasks involving 'NEW' apply options, as 'EXISTING' tasks start immediately

Example — a representative request:

BASH

```
curl -X POST "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/task/start" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY" \
-H "Content-Type: application/json" \
-d '{
  "name": "dummy",
  "apply": "EXISTING",
  "enabled": true
}'
```

Response codes

Status	Schema	Description
200	<code>TaskResponse</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

PATCH /kb/{kb_uuid}/task/{task_id}

Patch Knowledge Box Task

Modify the parameters of a task. This is only available for tasks that run on NEW resources, tasks that run on EXISTING resources cannot be modified.

Parameters

Name	In	Type	Req	Description
<code>kb_uuid</code>	path	string	yes	
<code>task_id</code>	path	string	yes	

Request body — `application/json` (required)

Field	Type	Req	Default	Description
<code>parameters</code>	DataAugmentation-Input SemanticModelMigrationParams null	yes		Parameters to be passed to the task. These must match the validation field for the Task definition class

Example — a representative request:

```
BASH
curl -X PATCH "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/task/{task_id}" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY" \
-H "Content-Type: application/json" \
-d '{
  "parameters": {
    "name": "",
    "on": 0
  }
}'
```

Response codes

Status	Schema	Description
200	<code>TaskRequest</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

DELETE /kb/{kb_uuid}/task/{task_id}

Delete Knowledge Box Task

Delete an Ingestion Agent or a Task of a Knowledge Box

Parameters

Name	In	Type	Req	Description
<code>kb_uuid</code>	path	string	yes	
<code>task_id</code>	path	string	yes	
<code>cleanup</code>	query	boolean		Cleanup all data associated with the task

Example — a representative request:

BASH

```
curl -X DELETE "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/task/{task_id}" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
200	<code>object</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

POST /kb/{kb_uuid}/task/{task_id}/cleanup

Cleanup Knowledge Box Task

Start a cleanup process for an Ingestion Agent or a Task of a Knowledge Box

Parameters

Name	In	Type	Req	Description
------	----	------	-----	-------------

<code>kb_uuid</code>	path	string	yes	
<code>task_id</code>	path	string	yes	

Example — a representative request:

```
BASH
curl -X POST
"https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/task/{task_id}/cleanup" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
200	<code>PublicTaskSet</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

POST /kb/{kb_uuid}/task/{task_id}/enable

Enable Knowledge Box Task

Enable or disable a task. This is only available for tasks that run on NEW resources. To disable a task that runs on EXISTING resources, use the stop endpoint.

Parameters

Name	In	Type	Req	Description
<code>task_id</code>	path	string	yes	
<code>kb_uuid</code>	path	string	yes	

Request body — `application/json` (required)

Field	Type	Req	Default	Description
<code>enabled</code>	boolean	yes		

Example — a representative request:

```
BASH
```

```
curl -X POST
"https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/task/{task_id}/enable" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY" \
-H "Content-Type: application/json" \
-d '{
  "enabled": true
}'
```

Response codes

Status	Schema	Description
200	<code>object</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

GET /kb/{kb_uuid}/task/{task_id}/inspect

Inspect Knowledge Box Task

Get the details of an Ingestion Agent or a Task of a Knowledge Box

Parameters

Name	In	Type	Req	Description
<code>kb_uuid</code>	path	string	yes	
<code>task_id</code>	path	string	yes	

Example — a representative request:

```
BASH
curl -X GET
"https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/task/{task_id}/inspect" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
200	<code>PublicTaskSet</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

POST /kb/{kb_uuid}/task/{task_id}/restart

Restart Knowledge Box Task

Restart an Ingestion Agent or a Task of a Knowledge Box. This only affects tasks that are configured to run on EXISTING resources, for tasks that run on NEW resources, use the enable endpoint.

Parameters

Name	In	Type	Req	Description
<code>kb_uuid</code>	path	string	yes	
<code>task_id</code>	path	string	yes	

Example — a representative request:

```
BASH
curl -X POST
"https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/task/{task_id}/restart" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
200	<code>TaskResponse</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

POST /kb/{kb_uuid}/task/{task_id}/stop

Stop Knowledge Box Task

Stop an Ingestion Agent or a Task of a Knowledge Box. This only affects tasks that are currently running on EXISTING resources, for disabling tasks that run on NEW resources, use the enable endpoint.

Parameters

Name	In	Type	Req	Description
<code>kb_uuid</code>	path	string	yes	
<code>task_id</code>	path	string	yes	

Example — a representative request:

BASH

```
curl -X POST "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/task/{task_id}/stop" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
200	TaskResponse	Successful Response
422	HTTPValidationError	Validation Error

GET /kb/{kb_uuid}/task/{task_id}/worker_config

Inspect Knowledge Box Task Worker

Get the worker config of a task on a knowledge box

Parameters

Name	In	Type	Req	Description
task_id	path	string	yes	

Example — a representative request:

BASH

```
curl -X GET
"https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/task/{task_id}/worker_config" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
200	TaskRequest	Successful Response
422	HTTPValidationError	Validation Error

GET /kb/{kb_uuid}/tasks

List Knowledge Box Available Tasks

Get the available Ingestion Agents and Tasks of a Knowledge Box

Parameters

Name	In	Type	Req	Description
<code>kb_uuid</code>	path	string	yes	
<code>count</code>	query	integer		How many finished tasks to return

Example — a representative request:

```
BASH
curl -X GET "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/tasks" \
  -H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
200	<code>TaskList</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

Training (1)

GET /download/{kbid}/model/{model_id}/{filename}

Download Account Predictive Model File

Download the trained model or any other generated file as a result of a training task on a Knowledge Box

Parameters

Name	In	Type	Req	Description
<code>kbid</code>	path	string	yes	
<code>model_id</code>	path	string	yes	
<code>filename</code>	path	string	yes	
<code>account_id</code>	query	string	yes	

Example — a representative request:

BASH

```
curl -X GET
"https://$ZONE.dp.progress.cloud/api/v1/download/$KB/model/{model_id}/{filename}"
\
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
200	<code>object</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

Extract Strategies (4)

GET /extract_strategies/{kbid}

Learning Extract Strategies

Get available extract strategies onprem

Parameters

Name	In	Type	Req	Description
<code>kbid</code>	path	string	yes	

Example — a representative request:

BASH

```
curl -X GET "https://$ZONE.dp.progress.cloud/api/v1/extract_strategies/$KB" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
200	<code>map<string, ExtractConfig-Output></code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

POST /extract_strategies/{kbid}

Add A Extract Strategy To A Kb Onprem

Add a extract strategy to a KB to be used

Parameters

Name	In	Type	Req	Description
kbid	path	string	yes	

Request body — application/json (required)

Field	Type	Req	Default	Description
name	string			
vllm_config	VLLMExtractionConfig-Input null			
ai_tables	AITables-Input null			
split	SplitConfig null			
max_parallel_llm_calls	integer		0	

Example — a representative request:

```
BASH
curl -X POST "https://$ZONE.dp.progress.cloud/api/v1/extract_strategies/$KB" \
  -H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY" \
  -H "Content-Type: application/json" \
  -d '{
    "name": "",
    "max_parallel_llm_calls": 0
  }'
```

Response codes

Status	Schema	Description
200	string	Successful Response
422	HTTPValidationError	Validation Error

GET /extract_strategies/{kbid}/strategies/{strategy_id}

Extract Strategy Configuration

Get extract strategy for a given id

Parameters

Name	In	Type	Req	Description
kbid	path	string	yes	
strategy_id	path	string	yes	

Example — a representative request:

```
BASH
curl -X GET
"https://$ZONE.dp.progress.cloud/api/v1/extract_strategies/$KB/strategies/
{strategy_id}" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
200	ExtractConfig-Output	Successful Response
422	HTTPValidationError	Validation Error

DELETE /extract_strategies/{kbid}/strategies/{strategy_id}

Disable Extract Strategy Onprem

Disable a extract strategy from a Knowledge Box

Parameters

Name	In	Type	Req	Description
kbid	path	string	yes	
strategy_id	path	string	yes	

Example — a representative request:

```
BASH
```

```
curl -X DELETE
"https://$ZONE.dp.progress.cloud/api/v1/extract_strategies/$KB/strategies/
{strategy_id}" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
200	<code>object</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

Task Assets (2)

GET /kb/{kb_uuid}/assets

List Kb Assets

List all assets for a Knowledge Box, ordered by creation date descending

Parameters

Name	In	Type	Req	Description
<code>kb_uuid</code>	path	string	yes	

Example — a representative request:

```
BASH
curl -X GET "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/assets" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
200	<code>AssetList</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

GET /kb/{kb_uuid}/assets/{asset_id}

Get Kb Asset

Get a signed download URL for a single task asset

Parameters

Name	In	Type	Req	Description
<code>kb_uuid</code>	path	string	yes	
<code>asset_id</code>	path	string	yes	

Example — a representative request:

BASH

```
curl -X GET "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/assets/{asset_id}" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
200	<code>AssetSignedURL</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

Search (3)

POST /kb/{kbid}/ask

Ask Knowledge Box

Ask questions on a Knowledge Box

Parameters

Name	In	Type	Req	Description
<code>kbid</code>	path	string	yes	
<code>x-ndb-client</code>	header	NucliaDBClientType		
<code>x-show-consumption</code>	header	boolean		
<code>x-nucliadb-user</code>	header	string		

x-forwarded-for	header	string		
x-synchronous	header	boolean		When set to true, outputs response as JSON in a non-streaming way. This is slower and requires waiting for entire answer to be ready.

Request body — `application/json` (required)

Field	Type	Req	Default	Description
audit_metadata	map<string, string> null			A dictionary containing optional audit-specific metadata, such as userid, environment, or other contextual information. This metadata can be leveraged for filtering and analyzing activity logs in futu...
query	string	yes		The query to get a generative answer for
top_k	integer		20	The top most relevant results to fetch at the retrieval step. The maximum number of results allowed is 200.
filter_expression	FilterExpression null			Returns only documents that match this filter expression. Filtering examples can be found here: (see docs) This allows building complex filtering expressions and replaces the following parameters:field...

<code>fields</code>	<code>array<string></code>		<code>[]</code>	The list of fields to search in. For instance: <code>a/title</code> to search only on title field. For more details on filtering by field, see: (see docs)
<code>filters</code>	<code>array<string> array<Filter></code>		<code>[]</code>	The list of filters to apply. Filtering examples can be found here: (see docs)
<code>keyword_filters</code>	<code>array<string> array<Filter></code>		<code>[]</code>	List of keyword filter expressions to apply to the retrieval step. The text block search will only be performed on the documents that contain the specified keywords. The filters are case-insensitive, ...
<code>vectorset</code>	<code>string null</code>			Vectors index to perform the search in. If not provided, NucliaDB will use the default one
<code>min_score</code>	<code>number MinScore null</code>			Minimum score to filter search results. Results with a lower score will be ignored. Accepts either a float or a dictionary with the minimum scores for the bm25 and vector indexes. If a float is provid...
<code>features</code>	<code>array<ChatOptions></code>		<code>['semantic', 'keyword']</code>	Features enabled for the chat endpoint. Semantic search is done if semantic is included. If keyword is included, the results will include matching paragraphs

				from the bm25 index. If relations is inclu...
<code>range_creation_start</code>	string null			Resources created before this date will be filtered out of search results. Datetime are represented as a str in ISO 8601 format, like: 2008-09-15T15:53:00+05:00.
<code>range_creation_end</code>	string null			Resources created after this date will be filtered out of search results. Datetime are represented as a str in ISO 8601 format, like: 2008-09-15T15:53:00+05:00.
<code>range_modification_start</code>	string null			Resources modified before this date will be filtered out of search results. Datetime are represented as a str in ISO 8601 format, like: 2008-09-15T15:53:00+05:00.
<code>range_modification_end</code>	string null			Resources modified after this date will be filtered out of search results. Datetime are represented as a str in ISO 8601 format, like: 2008-09-15T15:53:00+05:00.
<code>show</code>	array<ResourceProperties>		<code>['basic']</code>	Controls which types of metadata are serialized on resources of search results
<code>field_type_filter</code>	array<FieldTypeName>		<code>['text', 'file', 'link', 'conversation']</code>	Define which field types are serialized on resources of search results

			<code>', 'generic', 'key_value']</code>	
<code>extracted</code>	<code>array<ExtractedData TypeName></code>		<code>[]</code>	[Deprecated] Please use GET resource endpoint instead to get extracted metadata
<code>context</code>	<code>array<ChatContextMessage> null</code>			DEPRECATED! Please, use chathistory instead.
<code>chat_history</code>	<code>array<ChatContextMessage> null</code>			Use to rephrase the new LLM query by taking into account the chat conversation history. This will be passed to the LLM so that it is aware of the previous conversation.
<code>extra_context</code>	<code>array<string> null</code>			Additional context that is added to the retrieval context sent to the LLM. It allows extending the chat feature with content that may not be in the Knowledge Box.
<code>extra_context_images</code>	<code>array<Image> null</code>			Additional images added to the retrieval context sent to the LLM." It allows extending the chat feature with content that may not be in the Knowledge Box.
<code>query_image</code>	<code>Image null</code>			Image that will be used together with the query text for retrieval and then sent to the LLM as part of the context. If a query image is provided, the <code>extracontextimages</code> and

				ragimagesstrategies will be...
highlight	boolean		False	If set to true, the query terms will be highlighted in the results between <mark>...</mark> tags
resource_filters	array<string>		[]	List of resource ids to filter search results for. Only paragraphs from the specified resources will be returned.
prompt	string CustomPrompt null			Use to customize the prompts given to the generative model. Both system and user prompts can be customized. If a string is provided, it is interpreted as the user prompt.
rank_fusion	RankFusionName ReciprocalRankFusion		rrf	Rank fusion algorithm to use to merge results from multiple retrievers (keyword, semantic)
reranker	RerankerName PredictReranker		predict	Reranker let you specify which method you want to use to rerank your results at the end of retrieval
citations	boolean CitationsType null			Whether to include citations in the response. If set to None or False, no citations will be computed. If set to True or 'default', citations will be computed after answer generation and send as a sepa...

<code>citation_threshold</code>	number null			If citations is set to True or 'default', this will be the similarity threshold. Value between 0 and 1, lower values will produce more citations. If not set, it will be set to the optimized threshold ...
<code>security</code>	RequestSecurity null			Security metadata for the request. Please refer to the documentation for more details on how security works: (see docs)
<code>show_hidden</code>	boolean		<code>False</code>	If set to false (default), excludes hidden resources from search
<code>rag_strategies</code>	array<FieldExtensionStrategy FullResourceStrategy HierarchyResourceStrategy NeighbouringParagraphsStrategy MetadataExtensionStrategy ConversationalStrategy PreQueriesStrategy GraphStrategy>		<code>[]</code>	Options for tweaking how the context for the LLM model is crafted: - fullresource will add the full text of the matching resources to the context. This strategy cannot be combined with hierarchy, neig...
<code>rag_images_strategies</code>	array<PageImageStrategy ParagraphImageStrategy TableImageStrategy>		<code>[]</code>	Options for tweaking how the image based context for the LLM model is crafted: - pageimage will add the full page image of the matching resources to the context. - tables will send the table images

				fo...
<code>debug</code>	boolean		<code>False</code>	If set, the response will include some extra metadata for debugging purposes, like the list of queried nodes.
<code>generative_model</code>	string null			The generative model to use for the chat endpoint. If not provided, the model configured for the Knowledge Box is used.
<code>generative_model_seed</code>	integer null			The seed to use for the generative model for deterministic generation. Only supported by some models.
<code>max_tokens</code>	integer MaxTokens null			Use to limit the amount of tokens used in the LLM context and/or for generating the answer. If not provided, the default maximum tokens of the generative model will be used. If an integer is provided,...
<code>rephrase</code>	boolean		<code>False</code>	Rephrase the query for a more efficient retrieval. This will consume LLM tokens and make the request slower.
<code>chat_history_relevance_threshold</code>	number null			Threshold to determine if the past chat history is relevant to rephrase the user's question. 0 - Always treat previous messages as relevant (always

				rephrase).1 - Always treat previous messages as irre...
prefer_markdown	boolean		False	If set to true, the response will be in markdown format
answer_json_schema	object (free-form map) null			Desired JSON schema for the LLM answer. This schema is passed to the LLM so that it answers in a structured format following the schema. If not provided, textual response is returned. Note that when u...
generate_answer	boolean		True	Whether to generate an answer using the generative model. If set to false, the response will only contain the retrieval results.
search_configuration	string null			Load ask parameters from this configuration. Parameters in the request override parameters from the configuration.
reasoning	Reasoning boolean		False	Reasoning options for the generative model. Set to True to enable default reasoning, False to disable, or provide a Reasoning object for custom options.

Example — a real call captured against the live sandbox:

```
BASH
curl -X POST "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/ask" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY" \
```

```
-H "Content-Type: application/json" \  
-d '{  
  "query": "How much does a home battery store and when is it subsidised?",  
  "top_k": 5,  
  "features": [  
    "keyword",  
    "semantic"  
  ],  
  "citations": true  
'
```

Response:

```
JSON  
{  
  "answer": "A home battery stores 10kWh of usable capacity and is subsidised  
starting from July 2025 under the federal Cheaper Home Batteries program.",  
  "status": "success",  
  "retrieval_results": {  
    "resources": {  
      "86b5f7948b144b008f018cbd47a2bde4": {  
        "id": "86b5f7948b144b008f018cbd47a2bde4",  
        "slug": "battery-storage",  
        "title": "Home Battery Storage",  
        "summary": "",  
        "icon": "text/plain",  
        "thumbnail": "",  
        "metadata": {  
          "metadata": {},  
          "language": "en",  
          "languages": [  
            "en"  
          ],  
          "status": "PROCESSED"  
        },  
        "usermetadata": {  
          "classifications": [  
            {  
              "labelset": "topic",  
              "label": "energy",  
              "cancelled_by_user": false  
            }  
          ],  
          "relations": []  
        },  
        "fi  
... (truncated)
```

Response codes

Status	Schema	Description
--------	--------	-------------

200	<code>SyncAskResponse</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

POST `/kb/{kbid}/resource/{rid}/ask`

Ask a resource (by id)

Ask questions to a resource

Parameters

Name	In	Type	Req	Description
<code>kbid</code>	path	string	yes	
<code>rid</code>	path	string	yes	
<code>x-show-consumption</code>	header	boolean		
<code>x-ndb-client</code>	header	NucliaDBClientType		
<code>x-nucliadb-user</code>	header	string		
<code>x-forwarded-for</code>	header	string		
<code>x-synchronous</code>	header	boolean		When set to true, outputs response as JSON in a non-streaming way. This is slower and requires waiting for entire answer to be ready.

Request body — `application/json` (required)

Field	Type	Req	Default	Description
<code>audit_metadata</code>	<code>map<string, string> null</code>			A dictionary containing optional audit-specific metadata, such as <code>userid</code> , <code>environment</code> , or other contextual information. This metadata can be

				leveraged for filtering and analyzing activity logs in futu...
query	string	yes		The query to get a generative answer for
top_k	integer		20	The top most relevant results to fetch at the retrieval step. The maximum number of results allowed is 200.
filter_expression	FilterExpression null			Returns only documents that match this filter expression. Filtering examples can be found here: (see docs) This allows building complex filtering expressions and replaces the following parameters:field...
fields	array<string>		[]	The list of fields to search in. For instance: a/title to search only on title field. For more details on filtering by field, see: (see docs)
filters	array<string> array<Filter>		[]	The list of filters to apply. Filtering examples can be found here: (see docs)
keyword_filters	array<string> array<Filter>		[]	List of keyword filter expressions to apply to the retrieval step. The text block search will only be performed on the documents that contain the specified keywords. The filters are case-insensitive,

				...
<code>vectorset</code>	<code>string null</code>			Vectors index to perform the search in. If not provided, NucliaDB will use the default one
<code>min_score</code>	<code>number MinScore null</code>			Minimum score to filter search results. Results with a lower score will be ignored. Accepts either a float or a dictionary with the minimum scores for the bm25 and vector indexes. If a float is provid...
<code>features</code>	<code>array<ChatOptions></code>		<code>['semantic', 'keyword']</code>	Features enabled for the chat endpoint. Semantic search is done if semantic is included. If keyword is included, the results will include matching paragraphs from the bm25 index. If relations is inclu...
<code>range_creation_start</code>	<code>string null</code>			Resources created before this date will be filtered out of search results. Datetime are represented as a str in ISO 8601 format, like: 2008-09-15T15:53:00+05:00.
<code>range_creation_end</code>	<code>string null</code>			Resources created after this date will be filtered out of search results. Datetime are represented as a str in ISO 8601 format, like: 2008-09-15T15:53:00+05:00.
<code>range_modification_start</code>	<code>string null</code>			Resources modified before this date will

				be filtered out of search results. Datetime are represented as a str in ISO 8601 format, like: 2008-09-15T15:53:00+05:00.
<code>range_modification_end</code>	<code>string null</code>			Resources modified after this date will be filtered out of search results. Datetime are represented as a str in ISO 8601 format, like: 2008-09-15T15:53:00+05:00.
<code>show</code>	<code>array<ResourceProperties></code>		<code>['basic']</code>	Controls which types of metadata are serialized on resources of search results
<code>field_type_filter</code>	<code>array<FieldTypeName></code>		<code>['text', 'file', 'link', 'conversation', 'generic', 'key_value']</code>	Define which field types are serialized on resources of search results
<code>extracted</code>	<code>array<ExtractedDataTypeName></code>		<code>[]</code>	[Deprecated] Please use GET resource endpoint instead to get extracted metadata
<code>context</code>	<code>array<ChatContextMessage> null</code>			DEPRECATED! Please, use <code>chathistory</code> instead.
<code>chat_history</code>	<code>array<ChatContextMessage> null</code>			Use to rephrase the new LLM query by taking into account the chat conversation history. This will be passed to the LLM so that it is aware of the previous conversation.
<code>extra_context</code>	<code>array<string> null</code>			Additional context

				that is added to the retrieval context sent to the LLM. It allows extending the chat feature with content that may not be in the Knowledge Box.
<code>extra_context_images</code>	<code>array<Image> null</code>			Additional images added to the retrieval context sent to the LLM." It allows extending the chat feature with content that may not be in the Knowledge Box.
<code>query_image</code>	<code>Image null</code>			Image that will be used together with the query text for retrieval and then sent to the LLM as part of the context. If a query image is provided, the <code>extracontextimages</code> and <code>ragimagesstrategies</code> will be...
<code>highlight</code>	<code>boolean</code>		<code>False</code>	If set to true, the query terms will be highlighted in the results between <code><mark>...</mark></code> tags
<code>resource_filters</code>	<code>array<string></code>		<code>[]</code>	List of resource ids to filter search results for. Only paragraphs from the specified resources will be returned.
<code>prompt</code>	<code>string CustomPrompt null</code>			Use to customize the prompts given to the generative model. Both system and user prompts can be customized. If a string is provided, it

				is interpreted as the user prompt.
<code>rank_fusion</code>	RankFusionName ReciprocalRankFusion		<code>rrf</code>	Rank fusion algorithm to use to merge results from multiple retrievers (keyword, semantic)
<code>reranker</code>	RerankerName PredictReranker		<code>predict</code>	Reranker let you specify which method you want to use to rerank your results at the end of retrieval
<code>citations</code>	boolean CitationsType null			Whether to include citations in the response. If set to None or False, no citations will be computed. If set to True or 'default', citations will be computed after answer generation and send as a sepa...
<code>citation_threshold</code>	number null			If citations is set to True or 'default', this will be the similarity threshold. Value between 0 and 1, lower values will produce more citations. If not set, it will be set to the optimized threshold ...
<code>security</code>	RequestSecurity null			Security metadata for the request. Please refer to the documentation for more details on how security works: (see docs)
<code>show_hidden</code>	boolean		<code>False</code>	If set to false (default), excludes hidden resources from search

<code>rag_strategies</code>	array<FieldExtensionStrategy FullResourceStrategy HierarchyResourceStrategy NeighbouringParagraphsStrategy MetadataExtensionStrategy ConversationalStrategy PreQueriesStrategy GraphStrategy>		<code>[]</code>	Options for tweaking how the context for the LLM model is crafted: - fullresource will add the full text of the matching resources to the context. This strategy cannot be combined with hierarchy, neig...
<code>rag_images_strategies</code>	array<PageImageStrategy ParagraphImageStrategy TableImageStrategy>		<code>[]</code>	Options for tweaking how the image based context for the LLM model is crafted: - pageimage will add the full page image of the matching resources to the context. - tables will send the table images fo...
<code>debug</code>	boolean		<code>False</code>	If set, the response will include some extra metadata for debugging purposes, like the list of queried nodes.
<code>generative_model</code>	string null			The generative model to use for the chat endpoint. If not provided, the model configured for the Knowledge Box is used.
<code>generative_model_seed</code>	integer null			The seed to use for the generative model for deterministic generation. Only supported by some models.
<code>max_tokens</code>	integer MaxTokens null			Use to limit the amount of tokens

				used in the LLM context and/or for generating the answer. If not provided, the default maximum tokens of the generative model will be used. If an integer is provided,...
rephrase	boolean		False	Rephrase the query for a more efficient retrieval. This will consume LLM tokens and make the request slower.
chat_history_relevance_threshold	number null			Threshold to determine if the past chat history is relevant to rephrase the user's question. 0 - Always treat previous messages as relevant (always rephrase). 1 - Always treat previous messages as irre...
prefer_markdown	boolean		False	If set to true, the response will be in markdown format
answer_json_schema	object (free-form map) null			Desired JSON schema for the LLM answer. This schema is passed to the LLM so that it answers in a structured format following the schema. If not provided, textual response is returned. Note that when u...
generate_answer	boolean		True	Whether to generate an answer using the generative model. If set to false, the response will only contain the retrieval

				results.
<code>search_configuration</code>	string null			Load ask parameters from this configuration. Parameters in the request override parameters from the configuration.
<code>reasoning</code>	Reasoning boolean		<code>False</code>	Reasoning options for the generative model. Set to True to enable default reasoning, False to disable, or provide a Reasoning object for custom options.

Example — a representative request:

BASH

```
curl -X POST "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/resource/$RID/ask" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY" \
-H "Content-Type: application/json" \
-d '{
  "query": "your question here",
  "top_k": 20,
  "fields": [],
  "features": [
    "semantic",
    "keyword"
  ],
  "show": [
    "basic"
  ],
  "field_type_filter": [
    "text",
    "file",
    "link",
    "conversation",
    "generic",
    "key_value"
  ]
}'
```

Response codes

Status	Schema	Description
--------	--------	-------------

200	<code>SyncAskResponse</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

POST /kb/{kbid}/slug/{slug}/ask

Ask a resource (by slug)

Ask questions to a resource

Parameters

Name	In	Type	Req	Description
<code>kbid</code>	path	string	yes	
<code>slug</code>	path	string	yes	
<code>x-show-consumption</code>	header	boolean		
<code>x-ndb-client</code>	header	NucliaDBClientType		
<code>x-nucliadb-user</code>	header	string		
<code>x-forwarded-for</code>	header	string		
<code>x-synchronous</code>	header	boolean		When set to true, outputs response as JSON in a non-streaming way. This is slower and requires waiting for entire answer to be ready.

Request body — `application/json` (required)

Field	Type	Req	Default	Description
<code>audit_metadata</code>	map<string, string> null			A dictionary containing optional audit-specific metadata, such as userid, environment, or other contextual information. This metadata can be

				leveraged for filtering and analyzing activity logs in futu...
query	string	yes		The query to get a generative answer for
top_k	integer		20	The top most relevant results to fetch at the retrieval step. The maximum number of results allowed is 200.
filter_expression	FilterExpression null			Returns only documents that match this filter expression. Filtering examples can be found here: (see docs) This allows building complex filtering expressions and replaces the following parameters:field...
fields	array<string>		[]	The list of fields to search in. For instance: a/title to search only on title field. For more details on filtering by field, see: (see docs)
filters	array<string> array<Filter>		[]	The list of filters to apply. Filtering examples can be found here: (see docs)
keyword_filters	array<string> array<Filter>		[]	List of keyword filter expressions to apply to the retrieval step. The text block search will only be performed on the documents that contain the specified keywords. The filters are case-insensitive,

				...
<code>vectorset</code>	<code>string null</code>			Vectors index to perform the search in. If not provided, NucliaDB will use the default one
<code>min_score</code>	<code>number MinScore null</code>			Minimum score to filter search results. Results with a lower score will be ignored. Accepts either a float or a dictionary with the minimum scores for the bm25 and vector indexes. If a float is provid...
<code>features</code>	<code>array<ChatOptions></code>		<code>['semantic', 'keyword']</code>	Features enabled for the chat endpoint. Semantic search is done if semantic is included. If keyword is included, the results will include matching paragraphs from the bm25 index. If relations is inclu...
<code>range_creation_start</code>	<code>string null</code>			Resources created before this date will be filtered out of search results. Datetime are represented as a str in ISO 8601 format, like: 2008-09-15T15:53:00+05:00.
<code>range_creation_end</code>	<code>string null</code>			Resources created after this date will be filtered out of search results. Datetime are represented as a str in ISO 8601 format, like: 2008-09-15T15:53:00+05:00.
<code>range_modification_start</code>	<code>string null</code>			Resources modified before this date will

				be filtered out of search results. Datetime are represented as a str in ISO 8601 format, like: 2008-09-15T15:53:00+05:00.
<code>range_modification_end</code>	<code>string null</code>			Resources modified after this date will be filtered out of search results. Datetime are represented as a str in ISO 8601 format, like: 2008-09-15T15:53:00+05:00.
<code>show</code>	<code>array<ResourceProperties></code>		<code>['basic']</code>	Controls which types of metadata are serialized on resources of search results
<code>field_type_filter</code>	<code>array<FieldTypeName></code>		<code>['text', 'file', 'link', 'conversation', 'generic', 'key_value']</code>	Define which field types are serialized on resources of search results
<code>extracted</code>	<code>array<ExtractedDataTypeName></code>		<code>[]</code>	[Deprecated] Please use GET resource endpoint instead to get extracted metadata
<code>context</code>	<code>array<ChatContextMessage> null</code>			DEPRECATED! Please, use <code>chathistory</code> instead.
<code>chat_history</code>	<code>array<ChatContextMessage> null</code>			Use to rephrase the new LLM query by taking into account the chat conversation history. This will be passed to the LLM so that it is aware of the previous conversation.
<code>extra_context</code>	<code>array<string> null</code>			Additional context

				that is added to the retrieval context sent to the LLM. It allows extending the chat feature with content that may not be in the Knowledge Box.
<code>extra_context_images</code>	<code>array<Image> null</code>			Additional images added to the retrieval context sent to the LLM." It allows extending the chat feature with content that may not be in the Knowledge Box.
<code>query_image</code>	<code>Image null</code>			Image that will be used together with the query text for retrieval and then sent to the LLM as part of the context. If a query image is provided, the <code>extracontextimages</code> and <code>ragimagesstrategies</code> will be...
<code>highlight</code>	<code>boolean</code>		<code>False</code>	If set to true, the query terms will be highlighted in the results between <code><mark>...</mark></code> tags
<code>resource_filters</code>	<code>array<string></code>		<code>[]</code>	List of resource ids to filter search results for. Only paragraphs from the specified resources will be returned.
<code>prompt</code>	<code>string CustomPrompt null</code>			Use to customize the prompts given to the generative model. Both system and user prompts can be customized. If a string is provided, it

				is interpreted as the user prompt.
<code>rank_fusion</code>	RankFusionName ReciprocalRankFusion		<code>rrf</code>	Rank fusion algorithm to use to merge results from multiple retrievers (keyword, semantic)
<code>reranker</code>	RerankerName PredictReranker		<code>predict</code>	Reranker let you specify which method you want to use to rerank your results at the end of retrieval
<code>citations</code>	boolean CitationsType null			Whether to include citations in the response. If set to None or False, no citations will be computed. If set to True or 'default', citations will be computed after answer generation and send as a sepa...
<code>citation_threshold</code>	number null			If citations is set to True or 'default', this will be the similarity threshold. Value between 0 and 1, lower values will produce more citations. If not set, it will be set to the optimized threshold ...
<code>security</code>	RequestSecurity null			Security metadata for the request. Please refer to the documentation for more details on how security works: (see docs)
<code>show_hidden</code>	boolean		<code>False</code>	If set to false (default), excludes hidden resources from search

<code>rag_strategies</code>	array<FieldExtensionStrategy FullResourceStrategy HierarchyResourceStrategy NeighbouringParagraphsStrategy MetadataExtensionStrategy ConversationalStrategy PreQueriesStrategy GraphStrategy>		<code>[]</code>	Options for tweaking how the context for the LLM model is crafted: - fullresource will add the full text of the matching resources to the context. This strategy cannot be combined with hierarchy, neig...
<code>rag_images_strategies</code>	array<PageImageStrategy ParagraphImageStrategy TableImageStrategy>		<code>[]</code>	Options for tweaking how the image based context for the LLM model is crafted: - pageimage will add the full page image of the matching resources to the context. - tables will send the table images fo...
<code>debug</code>	boolean		<code>False</code>	If set, the response will include some extra metadata for debugging purposes, like the list of queried nodes.
<code>generative_model</code>	string null			The generative model to use for the chat endpoint. If not provided, the model configured for the Knowledge Box is used.
<code>generative_model_seed</code>	integer null			The seed to use for the generative model for deterministic generation. Only supported by some models.
<code>max_tokens</code>	integer MaxTokens null			Use to limit the amount of tokens

				used in the LLM context and/or for generating the answer. If not provided, the default maximum tokens of the generative model will be used. If an integer is provided,...
rephrase	boolean		False	Rephrase the query for a more efficient retrieval. This will consume LLM tokens and make the request slower.
chat_history_relevance_threshold	number null			Threshold to determine if the past chat history is relevant to rephrase the user's question. 0 - Always treat previous messages as relevant (always rephrase). 1 - Always treat previous messages as irre...
prefer_markdown	boolean		False	If set to true, the response will be in markdown format
answer_json_schema	object (free-form map) null			Desired JSON schema for the LLM answer. This schema is passed to the LLM so that it answers in a structured format following the schema. If not provided, textual response is returned. Note that when u...
generate_answer	boolean		True	Whether to generate an answer using the generative model. If set to false, the response will only contain the retrieval

				results.
<code>search_configuration</code>	string null			Load ask parameters from this configuration. Parameters in the request override parameters from the configuration.
<code>reasoning</code>	Reasoning boolean		<code>False</code>	Reasoning options for the generative model. Set to True to enable default reasoning, False to disable, or provide a Reasoning object for custom options.

Example — a representative request:

BASH

```
curl -X POST "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/slug/{slug}/ask" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY" \
-H "Content-Type: application/json" \
-d '{
  "query": "your question here",
  "top_k": 20,
  "fields": [],
  "features": [
    "semantic",
    "keyword"
  ],
  "show": [
    "basic"
  ],
  "field_type_filter": [
    "text",
    "file",
    "link",
    "conversation",
    "generic",
    "key_value"
  ]
}'
```

Response codes

Status	Schema	Description
--------	--------	-------------

200	<code>SyncAskResponse</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

Predict (23)

POST /predict/chat

Predict Chat

Chat interface for your NucliaDB

Parameters

Name	In	Type	Req	Description
<code>model</code>	query	string		
<code>x-show-consumption</code>	header	boolean		

Request body — `application/json` (required)

Field	Type	Req	Default	Description
<code>question</code>	string	yes		
<code>retrieval</code>	boolean		<code>True</code>	
<code>user_id</code>	string		<code>system</code>	
<code>system</code>	string null			
<code>chat_history</code>	array<AssistantMessage ToolMessage Message>			
<code>context</code>	array<Message>		<code>[]</code>	
<code>query_context</code>	array<string> map<string, string>		<code>{}</code>	
<code>query_context_order</code>	map<string, integer>		<code>{}</code>	
<code>truncate</code>	boolean null		<code>True</code>	
<code>user_prompt</code>	UserPrompt null			

<code>citations</code>	boolean CitationsType null			Whether to include citations in the response. If set to None or False, no citations will be computed. If set to True or 'default', citations will be computed after answer generation and send as a sepa...
<code>citation_threshold</code>	number null			If citations is set to True or 'default', this will be the similarity threshold. Value between 0 and 1, lower values will produce more citations. If not set, it will be set to the optimized threshold ...
<code>generative_model</code>	string null			
<code>max_tokens</code>	integer null			
<code>query_context_images</code>	array<Image> map<string, Image>		<code>{}</code>	
<code>prefer_markdown</code>	boolean null			
<code>json_schema</code>	object (free-form map) null			
<code>format_prompt</code>	boolean		<code>True</code>	
<code>rerank_context</code>	boolean		<code>False</code>	Whether to reorder the query context based on a reranker. This option will also make it so the first response will contain the scores given for each context piece.
<code>tools</code>	array<Tool>			List of tools to choose

<code>tool_choice</code>	ToolChoiceAuto ToolChoiceNone ToolChoiceRequired ToolChoiceForced		<code>{'type': 'required'}</code>	Tool choice strategy. auto: The model decides whether to use a tool or not based on the prompt and available tools. required (default): A tool must be used. none: Disables tool usage even if tools are ...
<code>seed</code>	integer null			Seed use for the generative model for a deterministic output.
<code>reasoning</code>	Reasoning boolean		<code>False</code>	Reasoning options for the generative model. Set to True to enable default reasoning, False to disable, or provide a Reasoning object for custom options.
<code>image_generation</code>	boolean		<code>False</code>	Whether to enable image generation in the response.

Example — a real call captured against the live sandbox:

BASH

```
curl -X POST "https://$ZONE.dp.progress.cloud/api/v1/predict/chat" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY" \
-H "Content-Type: application/json" \
-d '{
  "question": "What is 2+2? Answer with just the number.",
  "query_context": [],
  "user_id": "book-test"
}'
```

Response:

JSON

```
Not enough data to answer this.-2
```

Response codes

Status	Schema	Description
200	<code>object</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

POST /predict/chat/{kbid}

Predict Chat Kbid

Chat interface for your NucliaDB

Parameters

Name	In	Type	Req	Description
<code>kbid</code>	path	string	yes	
<code>x-show-consumption</code>	header	boolean		

Request body — `application/json` (required)

Field	Type	Req	Default	Description
<code>question</code>	string	yes		
<code>retrieval</code>	boolean		<code>True</code>	
<code>user_id</code>	string		<code>system</code>	
<code>system</code>	string null			
<code>chat_history</code>	array<AssistantMessage ToolMessage Message>			
<code>context</code>	array<Message>		<code>[]</code>	
<code>query_context</code>	array<string> map<string, string>		<code>{}</code>	
<code>query_context_order</code>	map<string, integer>		<code>{}</code>	
<code>truncate</code>	boolean null		<code>True</code>	
<code>user_prompt</code>	UserPrompt null			
<code>citations</code>	boolean CitationsType null			Whether to include citations in the

				response. If set to None or False, no citations will be computed. If set to True or 'default', citations will be computed after answer generation and send as a sepa...
<code>citation_threshold</code>	number null			If citations is set to True or 'default', this will be the similarity threshold. Value between 0 and 1, lower values will produce more citations. If not set, it will be set to the optimized threshold ...
<code>generative_model</code>	string null			
<code>max_tokens</code>	integer null			
<code>query_context_images</code>	array<Image> map<string, Image>		<code>{}</code>	
<code>prefer_markdown</code>	boolean null			
<code>json_schema</code>	object (free-form map) null			
<code>format_prompt</code>	boolean		<code>True</code>	
<code>rerank_context</code>	boolean		<code>False</code>	Whether to reorder the query context based on a reranker. This option will also make it so the first response will contain the scores given for each context piece.
<code>tools</code>	array<Tool>			List of tools to choose
<code>tool_choice</code>	ToolChoiceAuto ToolChoiceNone		<code>{'type': 'required'}</code>	Tool choice strategy. auto: The model

	ToolChoiceRequired ToolChoiceForced			decides whether to use a tool or not based on the prompt and available tools. required (default): A tool must be used.none: Disables tool usage even if tools are ...
seed	integer null			Seed use for the generative model for a deterministic output.
reasoning	Reasoning boolean		False	Reasoning options for the generative model. Set to True to enable default reasoning, False to disable, or provide a Reasoning object for custom options.
image_generation	boolean		False	Whether to enable image generation in the response.

Example — a representative request:

```
BASH
curl -X POST "https://$ZONE.dp.progress.cloud/api/v1/predict/chat/$KB" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY" \
-H "Content-Type: application/json" \
-d '{
  "question": "your question here",
  "retrieval": true,
  "user_id": "system",
  "format_prompt": true,
  "rerank_context": false,
  "image_generation": false
}'
```

Response codes

Status	Schema	Description
200	object	Successful Response
422	HTTPValidationError	Validation Error

--	--	--

POST /predict/compat/chat/completions

Chat Completions

OpenAI-compatible chat completions

Parameters

Name	In	Type	Req	Description
x-show-consumption	header	boolean		

Request body — `application/json` (required)

Field	Type	Req	Default	Description
messages	array<ChatCompletionDeveloperMessageParam ChatCompletionSystemMessageParam ChatCompletionUserMessageParam ChatCompletionAssistantMessageParam ChatCompletionToolMessageParam ChatCompletionFunctionMessageParam>	yes		
model	string null			
stream	boolean		False	
max_tokens	integer null			
seed	integer null			
temperature	number null			
top_p	number null			
user	string null			
tools	array<ChatCompletionFunctionToolParam			

	> null			
<code>tool_choice</code>	enum[none, auto, required] ChatCompletionAllowedToolChoiceParam ChatCompletionNamedToolChoiceParam ChatCompletionNamedToolChoiceCustomParam null			
<code>response_format</code>	ResponseFormatText ResponseFormatJSONSchema ResponseFormatJSONObject null			
<code>json_schema</code>	object (free-form map) null			
<code>reasoning_effort</code>	enum[none, minimal, low, medium, high, xhigh] null			
<code>image_generation</code>	boolean null			
<code>truncate</code>	boolean null			

Example — a representative request:

```
BASH
curl -X POST
"https://$ZONE.dp.progress.cloud/api/v1/predict/compat/chat/completions" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY" \
-H "Content-Type: application/json" \
-d '{
  "messages": [
    {
      "content": "string",
      "role": "string",
      "name": "string"
    }
  ],
  "stream": false
}'
```

Response codes

Status	Schema	Description
200	<code>object</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

POST /predict/compat/embeddings

Predict Sentences

Get the vector of a sentence

Parameters

Name	In	Type	Req	Description
<code>x-show-consumption</code>	header	boolean		

Request body — `application/json` (required)

Field	Type	Req	Default	Description
<code>input</code>	<code>array<string> string</code>	yes		
<code>model</code>	<code>string null</code>			

Example — a representative request:

```
BASH
curl -X POST "https://$ZONE.dp.progress.cloud/api/v1/predict/compat/embeddings" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY" \
-H "Content-Type: application/json" \
-d '{
  "input": [
    "string"
  ]
}'
```

Response codes

Status	Schema	Description
200	<code>EmbeddingsResponseModel</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

GET /predict/compat/models

List Models

OpenAI-compatible chat completions

Example — a representative request:

BASH

```
curl -X GET "https://$ZONE.dp.progress.cloud/api/v1/predict/compat/models" \
  -H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
200	ListModelsResponse	Successful Response
422	HTTPValidationError	Validation Error

GET /predict/compat/models/{model_id}

Get Model

OpenAI-compatible chat completions

Parameters

Name	In	Type	Req	Description
model_id	path	string	yes	

Example — a representative request:

BASH

```
curl -X GET
"https://$ZONE.dp.progress.cloud/api/v1/predict/compat/models/{model_id}" \
  -H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
200	ModelInfo	Successful Response

422	HTTPValidationError	Validation Error
-----	---------------------	------------------

GET /predict/query

Predict Query

Get the vector of a sentence. This endpoint is deprecated, please use the POST endpoint instead.

Parameters

Name	In	Type	Req	Description
text	query	string null		Text to process
semantic_model	query	string		
token_model	query	string		
generative_model	query	string		
rephrase	query	boolean		
rephrase_prompt	query	string null		Prompt to rephrase the sentence, if not provided, the default prompt will be used. It must include the {question} placeholder. The placeholder will be...
agentic_entities	query	boolean		If true, the model will return the entities detected in the sentence guided by an already defined Graph Extraction Agent in the Knowledge Box.
graph_nodes	query	array<string> null		List of graph node names to compute embeddings for. If not provided, no graph node embeddings will be computed.

<code>semantic_graph_h_node_models</code>	query	array<string> null		Semantic models to compute graph node embeddings for. If not provided, the default graph semantic model will be used.
<code>graph_edges</code>	query	array<string> null		List of graph edge labels to compute embeddings for. If not provided, no graph edge embeddings will be computed.
<code>semantic_graph_h_edge_models</code>	query	array<string> null		Semantic models to compute graph edge embeddings for. If not provided, the default graph semantic model will be used.
<code>x-show-consumption</code>	header	boolean		

Example — a real call captured against the live sandbox:

```
BASH
curl -X GET "https://$ZONE.dp.progress.cloud/api/v1/predict/query" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response:

```
JSON
{
  "language": "en",
  "stop_words": [],
  "semantic_threshold": 0.4,
  "semantic_thresholds": {
    "multilingual-2024-05-06": 0.4
  },
  "visual_llm": true,
  "max_context": 128000,
  "entities": {
    "tokens": [],
    "time": 0.005558967590332031,
    "input_tokens": 6
  }
}
```

```

},
"sentence": {
  "data": [
    -0.04191460460424423,
    0.0056065721437335014,
    -0.053347326815128326,
    -0.017205987125635147,
    -0.04948826879262924,
    -0.016326231881976128,
    0.005131538026034832,
    0.058367349207401276,
    -0.009004657156765461,
    0.01678045466542244,
    0.045586612075567245,
    0.04044400155544281,
    0.007959005422890186,
    0.02128644660115242,
    -0.0064488910138607025,
    -0.01187361590564251,
    0.026733797043561935,
    0.006086751818656921,
    0.0387454554438591,
    -0.005352390464395285,
    0.02820778824388981,
    ... (truncated)
  ]
}

```

Response codes

Status	Schema	Description
200	<code>QueryInfo</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

POST /predict/query

Predict Query Post

Preprocess a query.

Parameters

Name	In	Type	Req	Description
<code>x-show-consumption</code>	header	boolean		

Request body — `application/json` (required)

Field	Type	Req	Default	Description
<code>text</code>	string null			The query text to be processed
<code>query_image</code>	Image null			Image to be considered as part of the query. Even if the rephrase parameter is set to false, the rephrasing process will occur, combining the provided text with the image's visual features in the reph...
<code>rephrase</code>	boolean		<code>False</code>	If true, the model will rephrase the input text before processing
<code>rephrase_prompt</code>	string null			Custom prompt for rephrasing the input text
<code>generative_model</code>	string		<code>chatgpt-azure-4o</code>	The generative model to use for rephrasing.
<code>semantic_models</code>	array<string> null			Semantic models to compute the sentence vector for, if not provided, it will only compute the sentence vector for default semantic model in the Knowledge box's configuration.
<code>agentic_entities</code>	boolean		<code>False</code>	If true, the model will return the entities detected in the sentence guided by an already defined Graph Extraction Agent in the Knowledge Box.
<code>graph_nodes</code>	array<string> null			List of graph node names to compute embeddings for. If not provided, no graph node

				embeddings will be computed.
<code>semantic_graph_node_models</code>	<code>array<string> null</code>			Semantic models to compute graph node embeddings for. If not provided, it will only compute embeddings for the default graph semantic model in the Knowledge box's configuration.
<code>graph_edges</code>	<code>array<string> null</code>			List of graph edge labels to compute embeddings for. If not provided, no graph edge embeddings will be computed.
<code>semantic_graph_edge_models</code>	<code>array<string> null</code>			Semantic models to compute graph edge embeddings for. If not provided, it will only compute embeddings for the default graph edge semantic model in the Knowledge box's configuration.
<code>token_model</code>	<code>string</code>		<code>multilingual</code>	The NER model to use for extracting entities from the input text.
<code>semantic_model</code>	<code>string</code>		<code>multilingual-2024-05-06</code>	Semantic model to compute the sentence vector for.

Example — a representative request:

BASH

```
curl -X POST "https://$ZONE.dp.progress.cloud/api/v1/predict/query" \
  -H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY" \
  -H "Content-Type: application/json" \
  -d '{
    "rephrase": false,
```



```
"generative_model": "chatgpt-azure-4o",
"agentic_entities": false,
"token_model": "multilingual",
"semantic_model": "multilingual-2024-05-06"
}'
```

Response codes

Status	Schema	Description
200	QueryInfo	Successful Response
422	HTTPValidationError	Validation Error

GET /predict/query/{kbid}

Predict Query Kbid

Get the vector of a sentence

Parameters

Name	In	Type	Req	Description
kbid	path	string	yes	
text	query	string null		Text to process
semantic_model ls	query	array<string> null		Semantic models to compute the sentence vector for, if not provided, it will only compute the sentence vector for default semantic model in the Knowle...
graph_nodes	query	array<string> null		List of graph node names to compute embeddings for. If not provided, no graph node embeddings will be computed.
semantic_graph h_node_models	query	array<string> null		Semantic models to compute graph node embeddings for. If not provided, it will

				only compute embeddings for the default graph semantic model in the Kno...
graph_edges	query	array<string> null		List of graph edge labels to compute embeddings for. If not provided, no graph edge embeddings will be computed.
semantic_graph_edge_models	query	array<string> null		Semantic models to compute graph edge embeddings for. If not provided, it will only compute embeddings for the default graph edge semantic model in th...
generative_model	query	string		
rephrase	query	boolean		
rephrase_prompt	query	string null		Prompt to rephrase the sentence, if not provided, the default prompt will be used. It must include the {question} placeholder. The placeholder will be...
agentic_entities	query	boolean		If true, the model will return the entities detected in the sentence guided by an already defined Graph Extraction Agent in the Knowledge Box.
x-show-consumption	header	boolean		

Example — a representative request:

BASH

```
curl -X GET "https://$ZONE.dp.progress.cloud/api/v1/predict/query/$KB" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
200	QueryInfo	Successful Response
422	HTTPValidationError	Validation Error

POST /predict/query/{kbid}

Predict Query Post Kbid

Get the vector of a sentence

Parameters

Name	In	Type	Req	Description
kbid	path	string	yes	
x-show-consumption	header	boolean		

Request body — [application/json](#) (required)

Field	Type	Req	Default	Description
text	string null			The query text to be processed
query_image	Image null			Image to be considered as part of the query. Even if the rephrase parameter is set to false, the rephrasing process will occur, combining the provided text with the image's visual features in the reph...
rephrase	boolean		False	If true, the model will rephrase the input

				text before processing
rephrase_prompt	string null			Custom prompt for rephrasing the input text
generative_model	string null			The generative model to use for rephrasing
semantic_models	array<string> null			Semantic models to compute the sentence vector for, if not provided, it will only compute the sentence vector for default semantic model in the Knowledge box's configuration.
agentic_entities	boolean		False	If true, the model will return the entities detected in the sentence guided by an already defined Graph Extraction Agent in the Knowledge Box.
graph_nodes	array<string> null			List of graph node names to compute embeddings for. If not provided, no graph node embeddings will be computed.
semantic_graph_node_models	array<string> null			Semantic models to compute graph node embeddings for. If not provided, it will only compute embeddings for the default graph semantic model in the Knowledge box's configuration.
graph_edges	array<string> null			List of graph edge labels to compute embeddings for. If not provided, no

				graph edge embeddings will be computed.
<code>semantic_graph_h_edge_models</code>	<code>array<string> null</code>			Semantic models to compute graph edge embeddings for. If not provided, it will only compute embeddings for the default graph edge semantic model in the Knowledge box's configuration.

Example — a representative request:

```
BASH
curl -X POST "https://$ZONE.dp.progress.cloud/api/v1/predict/query/$KB" \
  -H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY" \
  -H "Content-Type: application/json" \
  -d '{
    "rephrase": false,
    "agentic_entities": false
  }'
```

Response codes

Status	Schema	Description
200	<code>QueryInfo</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

POST /predict/remi

Predict Remi

The REMi model is a model that computes metrics for a given question, answer and contexts. The metrics computed are: - Answer Relevance: Measures the relevance of the generated answer to the user query, in a scale of 0 to 5. - Context Relevance: Measures the relevance of the retrieved context to the user query, in a scale of 0 to 5. - Groundedness: Measures the degree to which the generated answer...

Parameters

Name	In	Type	Req	Description
------	----	------	-----	-------------

x-show-consumption	header	boolean		
--------------------	--------	---------	--	--

Request body — `application/json` (required)

Field	Type	Req	Default	Description
user_id	string	yes		
question	string null			
answer	string null			
contexts	array<string> null			

Example — a representative request:

```
BASH
curl -X POST "https://$ZONE.dp.progress.cloud/api/v1/predict/remi" \
  -H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY" \
  -H "Content-Type: application/json" \
  -d '{
    "user_id": "string"
  }'
```

Response codes

Status	Schema	Description
200	<code>RemiResponse</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

POST `/predict/remi/{kbid}`

Predict Remi Kbid

The REMi model is a model that computes metrics for a given question, answer and contexts. The metrics computed are: - Answer Relevance: Measures the relevance of the generated answer to the user query, in a scale of 0 to 5. - Context Relevance: Measures the relevance of the retrieved context to the user query, in a scale of 0 to 5. - Groundedness: Measures the degree to which the generated answer...

Parameters

Name	In	Type	Req	Description
------	----	------	-----	-------------

<code>kbid</code>	path	string	yes	
<code>x-show-consumption</code>	header	boolean		

Request body — `application/json` (required)

Field	Type	Req	Default	Description
<code>user_id</code>	string	yes		
<code>question</code>	string null			
<code>answer</code>	string null			
<code>contexts</code>	array<string> null			

Example — a representative request:

```
BASH
curl -X POST "https://$ZONE.dp.progress.cloud/api/v1/predict/remi/$KB" \
  -H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY" \
  -H "Content-Type: application/json" \
  -d '{
    "user_id": "string"
  }'
```

Response codes

Status	Schema	Description
200	<code>RemiResponse</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

POST `/predict/rephrase`

Predict Rephrase

Rephrase interface for your NucliaDB

Parameters

Name	In	Type	Req	Description
<code>model</code>	query	string		

Request body — `application/json` (required)

Field	Type	Req	Default	Description
<code>question</code>	string	yes		
<code>chat_history</code>	array<Message>			
<code>context</code>	array<Message>			
<code>user_context</code>	array<string> null			
<code>user_id</code>	string		<code>system</code>	
<code>generative_model</code>	string null			The generative model to use for the rephrase endpoint. If not provided, the model configured for the Knowledge Box is used.
<code>prompt</code>	string null			Prompt to send the model to rephrase the sentence, if not provided, the default prompt will be used. It must include the {question} placeholder. The placeholder will be replaced with the original ques...
<code>chat_history_relevance_threshold</code>	number null			

Example — a representative request:

BASH

```
curl -X POST "https://$ZONE.dp.progress.cloud/api/v1/predict/rephrase" \
  -H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY" \
  -H "Content-Type: application/json" \
  -d '{
    "question": "your question here",
    "user_id": "system"
  }'
```

Response codes

Status	Schema	Description
200	string	Successful Response
422	HTTPValidationError	Validation Error

POST /predict/rephrase/{kbid}

Predict Rephrase Kbid

Rephrase interface for your NucliaDB with Container

Parameters

Name	In	Type	Req	Description
kbid	path	string	yes	

Request body — application/json (required)

Field	Type	Req	Default	Description
question	string	yes		
chat_history	array<Message>			
context	array<Message>			
user_context	array<string> null			
user_id	string		system	
generative_model	string null			The generative model to use for the rephrase endpoint. If not provided, the model configured for the Knowledge Box is used.
prompt	string null			Prompt to send the model to rephrase the sentence, if not provided, the default prompt will be used. It must include the {question} placeholder. The placeholder will be

				replaced with the original ques...
chat_history_relevance_threshold	number null			

Example — a representative request:

```
BASH
curl -X POST "https://$ZONE.dp.progress.cloud/api/v1/predict/rephrase/$KB" \
  -H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY" \
  -H "Content-Type: application/json" \
  -d '{
    "question": "your question here",
    "user_id": "system"
  }'
```

Response codes

Status	Schema	Description
200	object	Successful Response
422	HTTPValidationError	Validation Error

POST /predict/rerank

Predict Rerank

Rerank your results

Parameters

Name	In	Type	Req	Description
x-show-consumption	header	boolean		

Request body — application/json (required)

Field	Type	Req	Default	Description
question	string	yes		
user_id	string	yes		

<code>context</code>	<code>map<string, string></code>	<code>{}</code>	
----------------------	--	-----------------	--

Example — a representative request:

```
BASH
curl -X POST "https://$ZONE.dp.progress.cloud/api/v1/predict/rerank" \
  -H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY" \
  -H "Content-Type: application/json" \
  -d '{
    "question": "your question here",
    "user_id": "string"
  }'
```

Response codes

Status	Schema	Description
200	<code>RerankResponse</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

POST /predict/rerank/{kbid}

Predict Rerank Kbid

Rerank your results

Parameters

Name	In	Type	Req	Description
<code>kbid</code>	path	string	yes	
<code>x-show-consumption</code>	header	boolean		

Request body — `application/json` (required)

Field	Type	Req	Default	Description
<code>question</code>	string	yes		
<code>user_id</code>	string	yes		
<code>context</code>	<code>map<string, string></code>		<code>{}</code>	

Example — a representative request:

```
BASH
curl -X POST "https://$ZONE.dp.progress.cloud/api/v1/predict/rerank/$KB" \
  -H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY" \
  -H "Content-Type: application/json" \
  -d '{
    "question": "your question here",
    "user_id": "string"
  }'
```

Response codes

Status	Schema	Description
200	RerankResponse	Successful Response
422	HTTPValidationError	Validation Error

GET /predict/sentence

Predict Sentence

Get the vector of a sentence

Parameters

Name	In	Type	Req	Description
text	query	string	yes	
model	query	string		
x-show-consumption	header	boolean		

Example — a real call captured against the live sandbox:

```
BASH
curl -X GET "https://$ZONE.dp.progress.cloud/api/v1/predict/sentence" \
  -H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response:

```
JSON
{
  "data": [
```

```

0.029169481247663498,
-0.018954990431666374,
-0.07776276022195816,
0.003512586699798703,
-0.0026530593167990446,
-0.04167490452528,
0.01399991661310196,
0.015164509415626526,
0.009860660880804062,
-0.01598522439599037,
0.01639706641435623,
0.021829016506671906,
-0.01719529926776886,
-0.0076917321421206,
-0.030735479667782784,
-0.034635890275239944,
-0.030568357557058334,
0.0011353847803547978,
-0.01868302747607231,
-0.023224780336022377,
0.009786454029381275,
-0.04288719967007637,
0.028424957767128944,
0.03647059574723244,
0.0408468134701252,
0.009546011686325073,
-0.03288205713033676,
0.018240636214613914,
-0.018396202474832535,
0.016910549253225327,
0.0065836310386657715,
0.002969736699014902,
-0.026616955175995827,
-0.01999562233686447,
... (truncated)

```

Response codes

Status	Schema	Description
200	SentenceResponse	Successful Response
422	HTTPValidationError	Validation Error

GET /predict/sentence/{kbid}

Predict Sentence Kbid

Get the vector of a sentence

Parameters

Name	In	Type	Req	Description
<code>kbid</code>	path	string	yes	
<code>text</code>	query	string	yes	
<code>models</code>	query	array<string> null		Semantic models to compute the sentence vector for, if not provided, it will only compute the sentence vector for default semantic model in the Knowle...
<code>x-show-consumption</code>	header	boolean		

Example — a representative request:

```
BASH
curl -X GET "https://$ZONE.dp.progress.cloud/api/v1/predict/sentence/$KB" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
200	<code>SentenceResponse</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

POST /predict/sentences

Predict Sentences

Generate the vectors of multiple texts in batch

Parameters

Name	In	Type	Req	Description
<code>x-show-consumption</code>	header	boolean		

Request body — `application/json` (required)

Field	Type	Req	Default	Description
<code>texts</code>	array<string>	yes		
<code>model</code>	string null			

Example — a representative request:

```
BASH
curl -X POST "https://$ZONE.dp.progress.cloud/api/v1/predict/sentences" \
  -H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY" \
  -H "Content-Type: application/json" \
  -d '{
    "texts": [
      "string"
    ]
  }'
```

Response codes

Status	Schema	Description
200	<code>SentencesResponse</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

POST /predict/summarize

Predict Summarize

Get summarized answers via NUA

Parameters

Name	In	Type	Req	Description
<code>model</code>	query	string		
<code>x-show-consumption</code>	header	boolean		

Request body — `application/json` (required)

Field	Type	Req	Default	Description
-------	------	-----	---------	-------------

<code>resources</code>	<code>map<string, SummarizeResource></code>	yes		
<code>summary_kind</code>	<code>SummaryKind</code>		<code>simple</code>	
<code>user_prompt</code>	<code>string null</code>			
<code>generative_model</code>	<code>string null</code>			

Example — a representative request:

```
BASH
curl -X POST "https://$ZONE.dp.progress.cloud/api/v1/predict/summarize" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY" \
-H "Content-Type: application/json" \
-d '{
  "resources": {},
  "summary_kind": "simple"
}'
```

Response codes

Status	Schema	Description
200	<code>SummarizedModelWithConsumption</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

POST /predict/summarize/{kbid}

Predict Summarize Kbid

Get summarized answers via NUA with Container

Parameters

Name	In	Type	Req	Description
<code>kbid</code>	path	string	yes	
<code>x-show-consumption</code>	header	boolean		

Request body — `application/json` (required)

Field	Type	Req	Default	Description
-------	------	-----	---------	-------------

resources	map<string, SummarizeResource>	yes		
summary_kind	SummaryKind		simple	
user_prompt	string null			
generative_model	string null			

Example — a representative request:

```
BASH
curl -X POST "https://$ZONE.dp.progress.cloud/api/v1/predict/summarize/$KB" \
  -H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY" \
  -H "Content-Type: application/json" \
  -d '{
    "resources": {},
    "summary_kind": "simple"
  }'
```

Response codes

Status	Schema	Description
200	SummarizedModelWithConsumption	Successful Response
422	HTTPValidationError	Validation Error

GET /predict/tokens

Predict Tokens

Get tokens of a sentence

Parameters

Name	In	Type	Req	Description
text	query	string	yes	
model	query	string		
x-show-consumption	header	boolean		

Example — a real call captured against the live sandbox:

```
BASH
curl -X GET "https://$ZONE.dp.progress.cloud/api/v1/predict/tokens" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response:

```
JSON
{
  "tokens": [
    {
      "text": "Sydney",
      "ner": "GPE",
      "start": 0,
      "end": 6
    },
    {
      "text": "Australia",
      "ner": "GPE",
      "start": 13,
      "end": 22
    },
    {
      "text": "Tesla",
      "ner": "ORG",
      "start": 27,
      "end": 32
    }
  ],
  "time": 0.006089448928833008,
  "input_tokens": 10
}
```

Response codes

Status	Schema	Description
200	TokenSearch	Successful Response
422	HTTPValidationError	Validation Error

GET /predict/tokens/{kbid}

Predict Tokens Kbid

Get tokens of a sentence with Container

Parameters

Name	In	Type	Req	Description
<code>kbid</code>	path	string	yes	
<code>text</code>	query	string	yes	
<code>x-show-consumption</code>	header	boolean		

Example — a representative request:

```
BASH
curl -X GET "https://$ZONE.dp.progress.cloud/api/v1/predict/tokens/$KB" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
200	<code>TokenSearch</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

Processing (6)

GET /processing/download

Download Binary File

Download a file referenced in an upload token

Parameters

Name	In	Type	Req	Description
<code>token</code>	query	string	yes	

Example — a representative request:

```
BASH
curl -X GET "https://$ZONE.dp.progress.cloud/api/v1/processing/download" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
200	<code>object</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

POST /processing/push

Send Data To Process

Push data to process

Request body — `application/json` (required)

Field	Type	Req	Default	Description
<code>kbid</code>	string null			
<code>uuid</code>	string null			
<code>title</code>	string null			
<code>labels</code>	array<string>		<code>[]</code>	
<code>slug</code>	string null			
<code>source</code>	Source null			
<code>userid</code>	string null			
<code>genericfield</code>	map<string, Text>		<code>{}</code>	
<code>filefield</code>	map<string, string>		<code>{}</code>	
<code>linkfield</code>	map<string, LinkUpload>		<code>{}</code>	
<code>textfield</code>	map<string, Text>		<code>{}</code>	
<code>layoutfield</code>	map<string, ProcessingLayoutDiff>		<code>{}</code>	
<code>conversationfield</code>	map<string, ProcessingConversation>		<code>{}</code>	
<code>generated_conversationfield</code>	map<string, ProcessingGenerated>		<code>{}</code>	

d	Conversation>			
processing_options	PushProcessingOptions null			
learning_config	LearningConfig null			
webhook_config	WebhookConfig null			

Example — a representative request:

```
BASH
curl -X POST "https://$ZONE.dp.progress.cloud/api/v1/processing/push" \
  -H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY" \
  -H "Content-Type: application/json" \
  -d '{
    "labels": []
  }'
```

Response codes

Status	Schema	Description
200	PushResponse	Successful Response
422	HTTPValidationError	Validation Error

GET /processing/requests

Get Processed Data

Retrieve status of processed and unprocessed payloads from a zone

Parameters

Name	In	Type	Req	Description
cursor	query	string null		
limit	query	integer		
scheduled	query	boolean null		
kbid	query	string null		
sort_on	query	SortFields		

sort_order	query	SortOrder		
------------	-------	-----------	--	--

Example — a representative request:

```
BASH
curl -X GET "https://$ZONE.dp.progress.cloud/api/v1/processing/requests" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
200	ProcessRequestStatusResults	Successful Response
422	HTTPValidationError	Validation Error

GET /processing/requests/{processing_id}

Get Processed Data

Retreive status of processed and unprocessed payloads from a zone

Parameters

Name	In	Type	Req	Description
processing_id	path	string	yes	

Example — a representative request:

```
BASH
curl -X GET
"https://$ZONE.dp.progress.cloud/api/v1/processing/requests/{processing_id}" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
200	ProcessRequestStatus	Successful Response
422	HTTPValidationError	Validation Error

GET /processing/requests/{processing_id}/results

Stream Processed Data

Retrieve SSE stream of processed data for a given processing ID.

Parameters

Name	In	Type	Req	Description
<code>processing_id</code>	path	string	yes	
<code>data_types</code>	query	array<StreamDataTypes> null		

Example — a representative request:

```
BASH
curl -X GET
"https://$ZONE.dp.progress.cloud/api/v1/processing/requests/{processing_id}/results" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
200	<code>object</code>	Server-Sent Events stream of processing data. Each event contains JSON object.
404		Processing ID not found or invalid
422	<code>HTTPValidationError</code>	Validation Error

POST /processing/upload

Upload Binary File

Upload a file using a multipart request

Parameters

Name	In	Type	Req	Description
<code>x-filename</code>	header	string null		
<code>x-password</code>	header	string null		

x-language	header	string null		
x-md5	header	string null		

Example — a representative request:

BASH

```
curl -X POST "https://$ZONE.dp.progress.cloud/api/v1/processing/upload" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
201	object	Successful Response
422	HTTPValidationError	Validation Error

Processing TUS uploads (4)

POST /processing/tusupload

Create New Upload

An empty POST request is used to create a new upload resource. The Upload-Length header indicates the size of the entire upload in bytes.

Example — a representative request:

BASH

```
curl -X POST "https://$ZONE.dp.progress.cloud/api/v1/processing/tusupload" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
201	object	Successful Response

PATCH /processing/tusupload/{upload_id}

Upload Data

Upload all bytes in the requests and append them in the specified offset

Parameters

Name	In	Type	Req	Description
<code>upload_id</code>	path	string	yes	

Example — a representative request:

```
BASH
curl -X PATCH
"https://$ZONE.dp.progress.cloud/api/v1/processing/tusupload/{upload_id}" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
200	<code>object</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

HEAD /processing/tusupload/{upload_id}

Upload Information

Get information about a current download (completed upload size)

Parameters

Name	In	Type	Req	Description
<code>upload_id</code>	path	string	yes	

Example — a representative request:

```
BASH
curl -X HEAD
"https://$ZONE.dp.progress.cloud/api/v1/processing/tusupload/{upload_id}" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
--------	--------	-------------

200	object	Successful Response
422	HTTPValidationError	Validation Error

OPTIONS /processing/tusupload/{upload_id}

Tus Server Information

Gather information about the Server's current configuration such as enabled extensions, version...

Parameters

Name	In	Type	Req	Description
upload_id	path	string null	yes	

Example — a representative request:

```
BASH
curl -X OPTIONS
"https://$ZONE.dp.progress.cloud/api/v1/processing/tusupload/{upload_id}" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
204		Successful Response
422	HTTPValidationError	Validation Error

split Strategies (4)

GET /split_strategies/{kbid}

Learning Split Strategies

Get available split strategies onprem

Parameters

Name	In	Type	Req	Description
------	----	------	-----	-------------

<code>kbid</code>	path	string	yes	

Example — a representative request:

BASH

```
curl -X GET "https://$ZONE.dp.progress.cloud/api/v1/split_strategies/$KB" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
200	<code>map<string, SplitConfiguration-Output></code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

POST /split_strategies/{kbid}

Add A Split Strategy To A Kb Onprem

Add a split strategy to a KB to be used

Parameters

Name	In	Type	Req	Description
<code>kbid</code>	path	string	yes	

Request body — `application/json` (required)

Field	Type	Req	Default	Description
<code>name</code>	string			
<code>max_paragraph</code>	integer		<code>0</code>	
<code>custom_split</code>	CustomSplitStrategy null		<code>0</code>	
<code>llm_split</code>	LLMSplitConfig-Input null			

<code>manual_split</code>	ManualSplitConfig null			
---------------------------	--------------------------	--	--	--

Example — a representative request:

```
BASH
curl -X POST "https://$ZONE.dp.progress.cloud/api/v1/split_strategies/$KB" \
  -H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY" \
  -H "Content-Type: application/json" \
  -d '{
    "name": "",
    "max_paragraph": 0
  }'
```

Response codes

Status	Schema	Description
200	<code>string</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

GET /split_strategies/{kbid}/strategies/{strategy_id}

Split Strategy Configuration

Get split strategy for a given id

Parameters

Name	In	Type	Req	Description
<code>kbid</code>	path	string	yes	
<code>strategy_id</code>	path	string	yes	

Example — a representative request:

```
BASH
curl -X GET
"https://$ZONE.dp.progress.cloud/api/v1/split_strategies/$KB/strategies/
{strategy_id}" \
  -H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
200	SplitConfiguration-Output	Successful Response
422	HTTPValidationError	Validation Error

DELETE /split_strategies/{kbid}/strategies/{strategy_id}

Disable Split Strategy Onprem

Disable a split strategy from a Knowledge Box

Parameters

Name	In	Type	Req	Description
kbid	path	string	yes	
strategy_id	path	string	yes	

Example — a representative request:

```
BASH
curl -X DELETE
"https://$ZONE.dp.progress.cloud/api/v1/split_strategies/$KB/strategies/
{strategy_id}" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
200	object	Successful Response
422	HTTPValidationError	Validation Error

PART 7 — REFERENCE APPENDIX

Appendix C — Zone API Reference

Managing Knowledge Boxes, keys, and connectors within a zone, fully expanded

The Regional Manager: KB lifecycle, service accounts/keys, connectors, backups, activity logs.

57 operations in **11 groups**, generated from the official [zone](#) v1 OpenAPI specification. Served from <https://<zone>.dp.progress.cloud>; authorized with a management credential (a NUA key with [allow_kb_management](#), or a user token / PAT). Every request-body field, parameter, and response is listed; object-typed fields reference named schemas documented in full in Appendix F (Schema Catalog).

Internal Sync Configs (1)

DELETE /api/internal/kb/{kb_id}/sync_configs

Delete All Configs

Parameters

Name	In	Type	Req	Description
kb_id	path	string	yes	

Example — a representative request:

```
BASH
curl -X DELETE "https://$ZONE.dp.progress.cloud/api/internal/kb/$KB/sync_configs" \
  -H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
204		Successful Response
422	HTTPValidationError	Validation Error

Backups (4)

DELETE /api/v1/account/{account_id}/backup/{backup_id}

Delete

Delete knowledgebox backup

Parameters

Name	In	Type	Req	Description
<code>account_id</code>	path	string	yes	
<code>backup_id</code>	path	string	yes	

Example — a representative request:

```
BASH
curl -X DELETE
"https://$ZONE.dp.progress.cloud/api/v1/account/$ACCOUNT/backup/{backup_id}" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
204		Successful Response
422	<code>HTTPValidationError</code>	Validation Error

POST /api/v1/account/{account_id}/backup/{backup_id}/restore

Restore

Restore a backup into a new knowledgebox

Parameters

Name	In	Type	Req	Description
<code>account_id</code>	path	string	yes	
<code>backup_id</code>	path	string	yes	

Request body — `application/json` (required)

Field	Type	Req	Default	Description
<code>slug</code>	string	yes		The slug of the new restored knowledgebox.
<code>title</code>	string	yes		The title of the new restored knowledgebox.

Example — a representative request:

```
BASH
curl -X POST
"https://$ZONE.dp.progress.cloud/api/v1/account/$ACCOUNT/backup/{backup_id}/
restore" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY" \
-H "Content-Type: application/json" \
-d '{
  "slug": "my-resource",
  "title": "My title"
}'
```

Response codes

Status	Schema	Description
201	<code>ItemCreated</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

GET /api/v1/account/{account_id}/backups

List

Get all backups from an account in a region

Parameters

Name	In	Type	Req	Description
<code>account_id</code>	path	string	yes	

Example — a representative request:

```
BASH
```



```
curl -X GET "https://$ZONE.dp.progress.cloud/api/v1/account/$ACCOUNT/backups" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
200	<code>array<BackupResponse></code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

POST /api/v1/account/{account_id}/backups

Create

Create knowledgebox backup

Parameters

Name	In	Type	Req	Description
<code>account_id</code>	path	string	yes	

Request body — `application/json` (required)

Field	Type	Req	Default	Description
<code>kb_id</code>	string	yes		The unique identifier of the knowledgebox to backup.

Example — a representative request:

```
BASH
curl -X POST "https://$ZONE.dp.progress.cloud/api/v1/account/$ACCOUNT/backups" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY" \
-H "Content-Type: application/json" \
-d '{
  "kb_id": "string"
}'
```

Response codes

Status	Schema	Description
--------	--------	-------------

201	<code>BackupCreateResponse</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

Knowledge Boxes (9)

GET /api/v1/account/{account_id}/kb/{kb_id}

Get Knowledge Box

Parameters

Name	In	Type	Req	Description
<code>account_id</code>	path	string	yes	
<code>kb_id</code>	path	string	yes	

Example — a representative request:

BASH

```
curl -X GET "https://$ZONE.dp.progress.cloud/api/v1/account/$ACCOUNT/kb/$KB" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
200	<code>KnowledgeBox</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

PATCH /api/v1/account/{account_id}/kb/{kb_id}

Modify Knowledge Box

Modifies an existing Knowledge Box.

Parameters

Name	In	Type	Req	Description
<code>account_id</code>	path	string	yes	

kb_id	path	string	yes	
-------	------	--------	-----	--

Request body — `application/json` (required)

Field	Type	Req	Default	Description
slug	string null			
title	string null			
state	KbState null			
description	string null			
allowed_origins	array<string> null			
allowed_ip_addresses	array<string> null			
search_configs	object (free-form map) null			
enforce_security	boolean null			Whether security is enforced by default on all requests to this Knowledge Box. Leaving this as None means no change, while setting it to True or False will update the setting accordingly.

Example — a representative request:

BASH

```
curl -X PATCH "https://$ZONE.dp.progress.cloud/api/v1/account/$ACCOUNT/kb/$KB" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY" \
-H "Content-Type: application/json" \
-d '{}'
```

Response codes

Status	Schema	Description
200	object (free-form map)	Successful Response
422	HTTPValidationError	Validation Error

--	--	--

DELETE /api/v1/account/{account_id}/kb/{kb_id}

Delete Knowledge Box

Parameters

Name	In	Type	Req	Description
account_id	path	string	yes	
kb_id	path	string	yes	

Example — a representative request:

BASH

```
curl -X DELETE "https://$ZONE.dp.progress.cloud/api/v1/account/$ACCOUNT/kb/$KB" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
204		Successful Response
422	HTTPValidationError	Validation Error

POST /api/v1/account/{account_id}/kb/{kb_id}/ephemeral_tokens

Create Ephemeral Token

Create an ephemeral token to access the Knowledge Box

Parameters

Name	In	Type	Req	Description
account_id	path	string	yes	
kb_id	path	string	yes	

Example — a representative request:

BASH

```
curl -X POST
"https://$ZONE.dp.progress.cloud/api/v1/account/$ACCOUNT/kb/$KB/ephemeral_tokens"
\
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY" \
-H "Content-Type: application/json" \
-d '{}'
```

Response codes

Status	Schema	Description
201	<code>EphemeralTokenResponse</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

GET /api/v1/account/{account_id}/kb/{kb_id}/logo.png

Download Logo

Knowledge Box logo

Parameters

Name	In	Type	Req	Description
<code>account_id</code>	path	string	yes	
<code>kb_id</code>	path	string	yes	

Example — a representative request:

```
BASH
curl -X GET
"https://$ZONE.dp.progress.cloud/api/v1/account/$ACCOUNT/kb/$KB/logo.png" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
200		Successful Response
422	<code>HTTPValidationError</code>	Validation Error

GET /api/v1/account/{account_id}/kb/{kb_id}/permissions

Get Knowledge Box Permissions

Parameters

Name	In	Type	Req	Description
<code>account_id</code>	path	string	yes	
<code>kb_id</code>	path	string	yes	

Example — a representative request:

```
BASH
curl -X GET
"https://$ZONE.dp.progress.cloud/api/v1/account/$ACCOUNT/kb/$KB/permissions" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
200	<code>PermissionsResponse</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

GET /api/v1/account/{account_id}/kbs

List Knowledge Boxes

List all region Knowledge Boxes that belongs to an specific account that the authenticated user can see

Parameters

Name	In	Type	Req	Description
<code>account_id</code>	path	string	yes	
<code>mode</code>	query	enum(kb, agent, agent_no_memory, agents)		Optional mode filter. 'kb' for standard KBs, 'agent' for agents with memory, 'agentnomemory' for agents without memory, 'agents' for all agent types (...)

<code>include_search_configs</code>	query	boolean		Whether to include searchconfigs in the response. Set to false to reduce payload size.
-------------------------------------	-------	---------	--	--

Example — a real call captured against the live sandbox:

BASH

```
curl -X GET "https://$ZONE.dp.progress.cloud/api/v1/account/$ACCOUNT/kbs" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response:

JSON

```
[
  {
    "id": "c610a68a-492a-499b-acc2-7ea27928f7ba",
    "slug": "ironbark-mutual-demo",
    "zone": "aws-ap-southeast-2-1",
    "title": "Ironbark Mutual - A Day on ARAG (synthetic insurance demo)",
    "state": "PRIVATE",
    "description": null,
    "role_on_kb": null,
    "allowed_origins": null,
    "allowed_ip_addresses": null,
    "search_configs": {},
    "external_index_provider": null,
    "hidden_resources_enabled": false,
    "hidden_resources_hide_on_creation": false,
    "prewarm_enabled": false,
    "enforce_security": false
  },
  {
    "id": "d6169de1-831b-4d1a-b0c7-1d924cc000a9",
    "slug": "factory-dam",
    "zone": "aws-ap-southeast-2-1",
    "title": "Factory Asset DAM",
    "state": "PRIVATE",
    "description": null,
    "role_on_kb": null,
    "allowed_origins": null,
    "allowed_ip_addresses": null,
    "search_configs": {},
    "external_index_provider": nu
    ... (truncated)
  ]
```

Response codes

Status	Schema	Description
--------	--------	-------------

200	<code>array<KnowledgeBoxSummary></code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

POST /api/v1/account/{account_id}/kbs

Create Knowledge Box

Creates a new Knowledge Box. The accepted values for the learningconfiguration are dynamic and may be dependant of the account used. Accepted field values needs to be checked on the 'Learning configuration schema' endpoint

Parameters

Name	In	Type	Req	Description
<code>account_id</code>	path	string	yes	

Request body — `application/json` (required)

Field	Type	Req	Default	Description
<code>slug</code>	string	yes		
<code>title</code>	string	yes		
<code>description</code>	string null			
<code>mode</code>	KBMode		<code>kb</code>	Mode of the Knowledge Box (immutable after creation). 'kb': Standard KB with NucliaDB storage. 'agent': Agent with memory backed by NucliaDB. 'agentnomemory': Agent without memory (no NucliaDB KB crea...
<code>learning_configuration</code>	object (free-form map)			

<code>allowed_origins</code>	<code>array<string> null</code>			
<code>allowed_ip_addresses</code>	<code>array<string> null</code>			
<code>external_index_provider</code>	<code>DummyKBIndexProvider null</code>			External index provider configuration for the Knowledge Box. If not set, the default NucliaDB's index will be used.
<code>search_configs</code>	<code>object (free-form map)</code>		<code>{}</code>	
<code>enforce_security</code>	<code>boolean</code>		<code>False</code>	Whether security is enforced by default on all requests to this Knowledge Box.

Example — a real call captured against the live sandbox:

BASH

```
curl -X POST "https://$ZONE.dp.progress.cloud/api/v1/account/$ACCOUNT/kbs" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY" \
-H "Content-Type: application/json" \
-d '{
  "slug": "arag-book-sandbox",
  "title": "ARAG Book Sandbox (safe to delete)",
  "description": "Throwaway KB for live-testing while writing the book."
}'
```

Response:

JSON

```
{
  "id": "$KB"
}
```

Response codes

Status	Schema	Description
201	<code>ItemCreated</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

POST /api/v1/ephemeral_token

Create Ephemeral Token From Sa

Create an ephemeral token scoped to a Knowledge Box (service account) or a path (OAuth user)

Request body — `application/json` (required)

Field	Type	Req	Default	Description
<code>agent_session</code>	string null			
<code>path</code>	string null			
<code>ttl</code>	integer null			

Example — a representative request:

BASH

```
curl -X POST "https://$ZONE.dp.progress.cloud/api/v1/ephemeral_token" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY" \
-H "Content-Type: application/json" \
-d '{}'
```

Response codes

Status	Schema	Description
201	<code>EphemeralTokenResponse</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

Invites (3)

POST /api/v1/account/{account_id}/kb/{kb_id}/invite

Invite To Knowledge Box

Send an invitation to join a Knowledge Box via email.

Parameters

Name	In	Type	Req	Description
------	----	------	-----	-------------

<code>account_id</code>	path	string	yes	
<code>kb_id</code>	path	string	yes	

Request body — `application/json` (required)

Field	Type	Req	Default	Description
<code>email</code>	string	yes		
<code>role</code>	KbRole	yes		
<code>came_from</code>	string null			

Example — a representative request:

```
BASH
curl -X POST
"https://$ZONE.dp.progress.cloud/api/v1/account/$ACCOUNT/kb/$KB/invite" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY" \
-H "Content-Type: application/json" \
-d '{
  "email": "user@example.com",
  "role": "SOWNER"
}'
```

Response codes

Status	Schema	Description
204		Successful Response
422	<code>HTTPValidationError</code>	Validation Error

DELETE /api/v1/account/{account_id}/kb/{kb_id}/invite

Delete Invite

Invalidate an invite to join a Knowledge Box

Parameters

Name	In	Type	Req	Description
<code>account_id</code>	path	string	yes	
<code>kb_id</code>	path	string	yes	

email	query	string	yes	
-------	-------	--------	-----	--

Example — a representative request:

```
BASH
curl -X DELETE
"https://$ZONE.dp.progress.cloud/api/v1/account/$ACCOUNT/kb/$KB/invite" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
204		Successful Response
422	HTTPValidationError	Validation Error

GET /api/v1/account/{account_id}/kb/{kb_id}/invites

Invited Users

List of users invited to a Knowledge Box, but not yet joined

Parameters

Name	In	Type	Req	Description
account_id	path	string	yes	
kb_id	path	string	yes	

Example — a representative request:

```
BASH
curl -X GET
"https://$ZONE.dp.progress.cloud/api/v1/account/$ACCOUNT/kb/$KB/invites" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
200	array<KbInvitedUser>	Successful Response
422	HTTPValidationError	Validation Error

API Keys (7)

DELETE /api/v1/account/{account_id}/kb/{kb_id}/service_account/{sa_id}

Delete Service Account

Delete a new Service Account for a Knowledge Box

Parameters

Name	In	Type	Req	Description
account_id	path	string	yes	
kb_id	path	string	yes	
sa_id	path	string	yes	

Example — a representative request:

```
BASH
curl -X DELETE
"https://$ZONE.dp.progress.cloud/api/v1/account/$ACCOUNT/kb/$KB/service_account/
{sa_id}" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
204		Successful Response
422	HTTPValidationError	Validation Error

DELETE /api/v1/account/{account_id}/kb/{kb_id}/service_account/{sa_id}/key/{sa_key_id}

Delete Api Key

Delete an API Key of a Knowledge Box

Parameters

Name	In	Type	Req	Description
------	----	------	-----	-------------

<code>account_id</code>	path	string	yes	
<code>kb_id</code>	path	string	yes	
<code>sa_id</code>	path	string	yes	
<code>sa_key_id</code>	path	string	yes	

Example — a representative request:

```
BASH
curl -X DELETE
"https://$ZONE.dp.progress.cloud/api/v1/account/$ACCOUNT/kb/$KB/service_account/
{sa_id}/key/{sa_key_id}" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
204		Successful Response
422	<code>HTTPValidationError</code>	Validation Error

POST /api/v1/account/{account_id}/kb/{kb_id}/service_account/{sa_id}/keys

Add Service Account Key

Create an API Key for a Service Account of a Knowledge Box

Parameters

Name	In	Type	Req	Description
<code>account_id</code>	path	string	yes	
<code>kb_id</code>	path	string	yes	
<code>sa_id</code>	path	string	yes	

Request body — `application/json` (required)

Field	Type	Req	Default	Description
<code>expires</code>	string	yes		

Example — a representative request:

BASH

```
curl -X POST
"https://$ZONE.dp.progress.cloud/api/v1/account/$ACCOUNT/kb/$KB/service_account/
{sa_id}/keys" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY" \
-H "Content-Type: application/json" \
-d '{
  "expires": "string"
}'
```

Response codes

Status	Schema	Description
201	ServiceAccountJWT	Successful Response
422	HTTPValidationError	Validation Error

GET /api/v1/account/{account_id}/kb/{kb_id}/service_accounts

List Service Accounts

List of Service Accounts of a Knowledge Box

Parameters

Name	In	Type	Req	Description
account_id	path	string	yes	
kb_id	path	string	yes	

Example — a representative request:**BASH**

```
curl -X GET
"https://$ZONE.dp.progress.cloud/api/v1/account/$ACCOUNT/kb/$KB/service_accounts"
\
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
200	array<ServiceAccount>	Successful Response
422	HTTPValidationError	Validation Error

POST /api/v1/account/{account_id}/kb/{kb_id}/service_accounts

Add Service Account

Create a new Service Account for a Knowledge Box

Parameters

Name	In	Type	Req	Description
account_id	path	string	yes	
kb_id	path	string	yes	

Request body — `application/json` (required)

Field	Type	Req	Default	Description
title	string	yes		
role	KbRole	yes		

Example — a representative request:

```
BASH
curl -X POST
"https://$ZONE.dp.progress.cloud/api/v1/account/$ACCOUNT/kb/$KB/service_accounts" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY" \
-H "Content-Type: application/json" \
-d '{
  "title": "My title",
  "role": "SOWNER"
}'
```

Response codes

Status	Schema	Description
201	ItemCreated	Successful Response
422	HTTPValidationError	Validation Error

POST /api/v1/service_account_agent_key

Create Temporal Agent Key

Creates a temporal key to use an agent session

Request body — `application/json` (required)

Field	Type	Req	Default	Description
<code>agent_session</code>	string	yes		
<code>ttl</code>	integer		10	

Example — a representative request:

```
BASH
curl -X POST "https://$ZONE.dp.progress.cloud/api/v1/service_account_agent_key" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY" \
-H "Content-Type: application/json" \
-d '{
  "agent_session": "string",
  "ttl": 10
}'
```

Response codes

Status	Schema	Description
201	<code>TemporalSecurityTokenResponse</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

POST /api/v1/service_account_temporal_key

Create Temporal Secure Token

Creates a temporal key for a service account. The key accepts

Example — a representative request:

```
BASH
curl -X POST
"https://$ZONE.dp.progress.cloud/api/v1/service_account_temporal_key" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY" \
-H "Content-Type: application/json" \
-d '{
  "ttl": 10
}'
```

Response codes

Status	Schema	Description
201	TemporalSecurityTokenResponse	Successful Response
422	HTTPValidationError	Validation Error

Users (2)

GET /api/v1/account/{account_id}/kb/{kb_id}/users

Users

List of users of a Knowledge Box

Parameters

Name	In	Type	Req	Description
account_id	path	string	yes	
kb_id	path	string	yes	
include_user_detail	query	boolean		

Example — a representative request:

```
BASH
curl -X GET
"https://$ZONE.dp.progress.cloud/api/v1/account/$ACCOUNT/kb/$KB/users" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
200	array<KbUser>	Successful Response
422	HTTPValidationError	Validation Error

PATCH /api/v1/account/{account_id}/kb/{kb_id}/users

Update Users

Assign or remove users from a Knowledge Box

Parameters

Name	In	Type	Req	Description
<code>account_id</code>	path	string	yes	
<code>kb_id</code>	path	string	yes	

Request body — `application/json` (required)

Field	Type	Req	Default	Description
<code>add</code>	array<AddKbUser> null			
<code>update</code>	array<UpdateKbUser> null			
<code>delete</code>	array<string> null			

Example — a representative request:

```
BASH
curl -X PATCH
"https://$ZONE.dp.progress.cloud/api/v1/account/$ACCOUNT/kb/$KB/users" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY" \
-H "Content-Type: application/json" \
-d '{}'
```

Response codes

Status	Schema	Description
204		Successful Response
422	<code>HTTPValidationError</code>	Validation Error

NUA Keys (6)

GET /api/v1/account/{account_id}/nua_client/{client_id}

Get Nua Client By Id

Gets the metadata of an existing Nuclia Understanding API

Parameters

Name	In	Type	Req	Description
<code>account_id</code>	path	string	yes	
<code>client_id</code>	path	string	yes	

Example — a representative request:

```
BASH
curl -X GET "https://$ZONE.dp.progress.cloud/api/v1/account/$ACCOUNT/nua_client/{client_id}" \
  -H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
200	<code>NUAClientMetadata</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

PATCH `/api/v1/account/{account_id}/nua_client/{client_id}`

Update Nua Client By Id

Updates an existing Nuclia Understanding API client

Parameters

Name	In	Type	Req	Description
<code>account_id</code>	path	string	yes	
<code>client_id</code>	path	string	yes	

Request body — `application/json` (required)

Field	Type	Req	Default	Description
<code>tokens_limit</code>	integer null			
<code>title</code>	string null			

<code>contact</code>	string null			
----------------------	---------------	--	--	--

Example — a representative request:

```
BASH
curl -X PATCH
"https://$ZONE.dp.progress.cloud/api/v1/account/$ACCOUNT/nua_client/{client_id}" \
  -H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY" \
  -H "Content-Type: application/json" \
  -d '{}'
```

Response codes

Status	Schema	Description
200	<code>NUAClientMetadata</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

DELETE /api/v1/account/{account_id}/nua_client/{client_id}

Delete Nua Client By Id

Deletes an existing Nuclia Understanding API client

Parameters

Name	In	Type	Req	Description
<code>account_id</code>	path	string	yes	
<code>client_id</code>	path	string	yes	

Example — a representative request:

```
BASH
curl -X DELETE
"https://$ZONE.dp.progress.cloud/api/v1/account/$ACCOUNT/nua_client/{client_id}" \
  -H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
--------	--------	-------------

204		Successful Response
422	<code>HTTPValidationError</code>	Validation Error

PUT /api/v1/account/{account_id}/nua_client/{client_id}/key

Set Account Nua Client Key

Recreates a new session key for an existing Nuclia Understanding API client

Parameters

Name	In	Type	Req	Description
<code>account_id</code>	path	string	yes	
<code>client_id</code>	path	string	yes	

Example — a representative request:

```
BASH
curl -X PUT "https://$ZONE.dp.progress.cloud/api/v1/account/$ACCOUNT/nua_client/{client_id}/key" \
  -H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
201	<code>NUAClient</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

GET /api/v1/account/{account_id}/nua_clients

Get Account Nua Clients

Gets the list of Nuclia Understanding API clients of an account

Parameters

Name	In	Type	Req	Description
<code>account_id</code>	path	string	yes	
<code>client_id</code>	query	string null		

Example — a representative request:

```
BASH
curl -X GET "https://$ZONE.dp.progress.cloud/api/v1/account/$ACCOUNT/nua_clients" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
200	<code>NUAClientsMetadata</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

POST /api/v1/account/{account_id}/nua_clients

Create Account Nua Client

Create a new Nuclia Understanding API client for an account

Parameters

Name	In	Type	Req	Description
<code>account_id</code>	path	string	yes	

Request body — `application/json` (required)

Field	Type	Req	Default	Description
<code>contact</code>	string	yes		
<code>title</code>	string	yes		
<code>description</code>	string null			
<code>client_id</code>	string null			
<code>processing_webhook</code>	ProcessingWebhook null			
<code>allow_knowledge_management</code>	boolean		<code>False</code>	
<code>tokens_limit</code>	integer null			

Example — a representative request:

BASH

```
curl -X POST
"https://$ZONE.dp.progress.cloud/api/v1/account/$ACCOUNT/nua_clients" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY" \
-H "Content-Type: application/json" \
-d '{
  "contact": "string",
  "title": "My title",
  "allow_kb_management": false
}'
```

Response codes

Status	Schema	Description
201	NUAClient	Successful Response
422	HTTPValidationError	Validation Error

Knowledgebox activity (8)

GET /api/v1/kb/{kb_id}/activity/download_request/{request_id}

Get Download Status

Get the status of a download request by Knowledgebox ID and request ID

Parameters

Name	In	Type	Req	Description
kb_id	path	string	yes	
request_id	path	string	yes	

Example — a representative request:

BASH

```
curl -X GET
"https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/activity/download_request/
{request_id}" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
200	<code>DownloadRequest</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

GET /api/v1/kb/{kb_id}/activity/metrics

Activity Logs Metrics

Get metrics for search and ask logs

Parameters

Name	In	Type	Req	Description
<code>kb_id</code>	path	string	yes	
<code>aggregation</code>	query	Aggregation		
<code>from</code>	query	string	yes	
<code>to</code>	query	string null		

Example — a representative request:

```
BASH
curl -X GET "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/activity/metrics" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
200	<code>array<ActivityUsageResponse></code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

GET /api/v1/kb/{kb_id}/activity/{event_type}/months

Get Knowledgebox Activity Downloads List

Get the list of monthly activity audit files available to download

Parameters

Name	In	Type	Req	Description
<code>kb_id</code>	path	string	yes	
<code>event_type</code>	path	EventType	yes	

Example — a representative request:

```
BASH
curl -X GET
"https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/activity/{event_type}/months" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
200	<code>EventDownloadsList</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

POST /api/v1/kb/{kb_id}/activity/{event_type}/query

Query Knowledgebox Activity Logs

Parameters

Name	In	Type	Req	Description
<code>kb_id</code>	path	string	yes	
<code>event_type</code>	path	EventType	yes	

Example — a representative request:

```
BASH
curl -X POST
"https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/activity/{event_type}/query" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY" \
-H "Content-Type: application/json" \
-d '{
  "year_month": "string",
  "filters": {}
}'
```

Response codes

Status	Schema	Description
200	<code>object</code>	Returns data in NDJSON format
422	<code>HTTPValidationError</code>	Validation Error

POST /api/v1/kb/{kb_id}/activity/{event_type}/query/download

Request The Preparation Of Knowledgebox Activity Logs Download

Request Knowledgebox Activity Logs download

Parameters

Name	In	Type	Req	Description
<code>kb_id</code>	path	string	yes	
<code>event_type</code>	path	EventType	yes	
<code>accept</code>	header	enum(application/x-ndjson, text/csv)	yes	

Example — a representative request:

```
BASH
curl -X POST
"https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/activity/{event_type}/query/download" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY" \
-H "Content-Type: application/json" \
-d '{
  "year_month": "string",
  "filters": {},
  "notify_via_email": false
}'
```

Response codes

Status	Schema	Description
202	<code>DownloadRequest</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

POST /api/v1/kb/{kb_id}/remi/query

Query Rag Requests By Scores

Get a list of rag request that matches a remi scores query

Parameters

Name	In	Type	Req	Description
<code>kb_id</code>	path	string	yes	

Request body — `application/json` (required)

Field	Type	Req	Default	Description
<code>context_relevance</code>	ContextRelevanceQuery null			
<code>month</code>	string	yes		
<code>from_date</code>	string null			
<code>to_date</code>	string null			
<code>feedback_good</code>	boolean null			
<code>status</code>	Status null			
<code>pagination</code>	Pagination		<code>{'limit': 10}</code>	

Example — a representative request:

```
BASH
curl -X POST "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/remi/query" \
  -H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY" \
  -H "Content-Type: application/json" \
  -d '{
    "month": "string"
  }'
```

Response codes

Status	Schema	Description
200	<code>RemiQueryResults</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

GET /api/v1/kb/{kb_uuid}/remi/events/{event_id}

Get A Rag Request With Related Remi Scores And Full Context

Get a rag request with full context and REMI scores. Intended for obtaining complete context for an item originating from a /remi/query

Parameters

Name	In	Type	Req	Description
kb_uuid	path	string	yes	
event_id	path	integer	yes	

Example — a representative request:

```
BASH
curl -X GET
"https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/remi/events/{event_id}" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
200	RemiQueryResultWithContext	Successful Response
422	HTTPValidationError	Validation Error

GET /api/v1/kb/{kb_uuid}/remi/scores

Aggregated Remi Scores Across Time

Get the evolution of remi scores of a kb on a period of time

Parameters

Name	In	Type	Req	Description
kb_uuid	path	string	yes	
aggregation	query	Aggregation		
from	query	string	yes	
to	query	string null		

Example — a representative request:

BASH

```
curl -X GET "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/remi/scores" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
200	<code>array<AggregatedRemiScoreMetric></code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

External Connections (7)

GET /api/v1/kb/{kb_id}/external_connection/{connection_id}

Get External Connection

Parameters

Name	In	Type	Req	Description
<code>kb_id</code>	path	string	yes	
<code>connection_id</code>	path	string	yes	

Example — a representative request:

BASH

```
curl -X GET "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/external_connection/{connection_id}" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
200	<code>ExternalConnectionOutput</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

DELETE /api/v1/kb/{kb_id}/external_connection/{connection_id}

Delete External Connection

Parameters

Name	In	Type	Req	Description
kb_id	path	string	yes	
connection_id	path	string	yes	

Example — a representative request:

```
BASH
curl -X DELETE
"https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/external_connection/
{connection_id}" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
204		Successful Response
422	HTTPValidationError	Validation Error

GET /api/v1/kb/{kb_id}/external_connection/{connection_id}/browse

Browse Storage

Browse cloud storage to discover sites, drives and folders. Use this endpoint to help users select a drive and folder path when creating a sync configuration. SharePoint/OneDrive flow (OAuth): 1. Resolve site URL: use the /resolvesite endpoint with a site URL → get siteid 2. List drives in a site: provide siteid → returns drives for that site - Use siteid=me for personal OneDrive (OAuth only) 3. L...

Parameters

Name	In	Type	Req	Description
kb_id	path	string	yes	
connection_id	path	string	yes	
site_search	query	string null		Search query for SharePoint sites. Only supported for

				certificate auth connections. For OAuth, use the /resolve_site endpoint instead.
site_id	query	string null		Microsoft Graph site ID to list drives for. Must be in the format 'hostname,site-collection-id,web-id' (e.g. 'contoso.sharepoint.com,<guid>,<guid>'). ...
drive_id	query	string null		Drive ID to browse folders in.
path	query	string null		Folder path to list within the drive. Cannot be combined with folder_id.
folder_id	query	string null		Folder ID to list contents of. Cannot be combined with path.
page_token	query	string null		Pagination token

Example — a representative request:

BASH

```
curl -X GET "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/external_connection/{connection_id}/browse" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
200	BrowseStorageOutput	Successful Response
422	HTTPValidationError	Validation Error

GET /api/v1/kb/{kb_id}/external_connection/{connection_id}/resolve_site

Resolve Site

Resolve a SharePoint site URL to its site ID, or return the OneDrive site. Given a full SharePoint site URL, this endpoint returns the site ID, name, and web URL. The returned site ID can then be used with the /browse endpoint to list drives for that site. This endpoint works with the Sites.Selected permission scope, as long as the specific site has been granted access to the application. For azur...

Parameters

Name	In	Type	Req	Description
<code>kb_id</code>	path	string	yes	
<code>connection_id</code>	path	string	yes	
<code>site_url</code>	query	string	yes	Full SharePoint site URL to resolve, or 'me'/'onedrive' for personal OneDrive. Examples: (see docs) me, onedrive. Passing 'me' or 'onedrive' is only s...

Example — a representative request:

BASH

```
curl -X GET "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/external_connection/{connection_id}/resolve_site" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
200	<code>StorageSiteOutput</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

GET /api/v1/kb/{kb_id}/external_connections

List External Connections

Parameters

Name	In	Type	Req	Description
<code>kb_id</code>	path	string	yes	

Example — a representative request:

```
BASH
curl -X GET "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/external_connections" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
200	<code>array<ExternalConnectionOutput></code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

POST /api/v1/kb/{kb_id}/external_connections

Create External Connection

Create an external connection for cloud storage sync. For OAuth providers (googleoauth, azureoauth, sharefileoauth): - Only provider is required in the request body - Returns an authorizeurl that the user should visit to complete the OAuth flow - After successful authorization, an ExternalConnection will be created - Returns HTTP 200 For awss3assumerole: - Both provider and credentials are require...

Parameters

Name	In	Type	Req	Description
<code>kb_id</code>	path	string	yes	

Request body — `application/json` (required)

Field	Type	Req	Default	Description
<code>provider</code>	Provider	yes		
<code>credentials</code>	AWSS3AssumeRoleCredentials AzureCertificateCredentials null			

Example — a representative request:

```
BASH
```

```
curl -X POST "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/external_connections" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY" \
-H "Content-Type: application/json" \
-d '{
  "provider": "google_oauth"
}'
```

Response codes

Status	Schema	Description
200	AuthorizeUrlOutput	OAuth authorization URL (for OAuth providers)
201	ExternalConnectionOutput	Created external connection (for access key providers)
422	HTTPValidationError	Validation Error

GET /api/v1/kb/{kb_id}/external_connections/s3/assume_role_info

Get S3 Assume Role Info

Returns the AWS assume-role information that a client needs to create an IAM role in their own AWS account for cross-account S3 access. The client should use the returned externalid, rolename, and awsaccountid to configure the IAM role trust policy.

Parameters

Name	In	Type	Req	Description
kb_id	path	string	yes	

Example — a representative request:

```
BASH
curl -X GET "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/external_connections/s3/assume_role_info" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
200	S3AssumeRoleInfoOutput	Successful Response

422	HTTPValidationError	Validation Error
-----	---------------------	------------------

Sync Configs (7)

GET /api/v1/kb/{kb_id}/sync_config/{config_id}

Get Config

Parameters

Name	In	Type	Req	Description
kb_id	path	string	yes	
config_id	path	string	yes	

Example — a representative request:

```
BASH
curl -X GET
"https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/sync_config/{config_id}" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
200	SyncConfigOutput	Successful Response
422	HTTPValidationError	Validation Error

PATCH /api/v1/kb/{kb_id}/sync_config/{config_id}

Update Config

Parameters

Name	In	Type	Req	Description
kb_id	path	string	yes	
config_id	path	string	yes	

Request body — `application/json` (required)

Field	Type	Req	Default	Description
<code>name</code>	string null			
<code>file_filter</code>	FileFilterInput null			
<code>labels</code>	array<LabelInput> null			
<code>modified_time_range</code>	ModifiedTimeRangeInput null			
<code>extract_strategy</code>	string null			

Example — a representative request:

```
BASH
curl -X PATCH
"https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/sync_config/{config_id}" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY" \
-H "Content-Type: application/json" \
-d '{}'
```

Response codes

Status	Schema	Description
200	<code>SyncConfigOutput</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

DELETE /api/v1/kb/{kb_id}/sync_config/{config_id}

Delete Config

Parameters

Name	In	Type	Req	Description
<code>kb_id</code>	path	string	yes	
<code>config_id</code>	path	string	yes	

Example — a representative request:

```
BASH
```

```
curl -X DELETE
"https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/sync_config/{config_id}" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
204		Successful Response
422	<code>HTTPValidationError</code>	Validation Error

POST /api/v1/kb/{kb_id}/sync_config/{config_id}/authorize

Get Authorize Url

Parameters

Name	In	Type	Req	Description
<code>kb_id</code>	path	string	yes	
<code>config_id</code>	path	string	yes	

Request body — `application/json` (required)

Field	Type	Req	Default	Description
<code>widget_redirect_url</code>	string null			
<code>rao_redirect_url</code>	string	yes		

Example — a representative request:

```
BASH
curl -X POST
"https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/sync_config/{config_id}/authorize" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY" \
-H "Content-Type: application/json" \
-d '{
  "rao_redirect_url": "https://example.com"
}'
```

Response codes

Status	Schema	Description
200	AuthorizeUrlOutput	Successful Response
422	HTTPValidationError	Validation Error

POST /api/v1/kb/{kb_id}/sync_config/{config_id}/validate_resources

Validate Resources

Validate that the user has read access to the specified resources. This endpoint checks if the provided credentials have access to the files/items specified in the request. The provider is determined from the sync config. Credentials must be encrypted using the same encryption key used for external connections. Args: kbid: The knowledge base ID. configid: The sync config ID (determines the provide...

Parameters

Name	In	Type	Req	Description
kb_id	path	string	yes	
config_id	path	string	yes	

Request body — [application/json](#) (required)

Field	Type	Req	Default	Description
credentials	string	yes		
resources	array<ResourceIdentifier>	yes		

Example — a representative request:

```
BASH
curl -X POST
"https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/sync_config/{config_id}/
validate_resources" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY" \
-H "Content-Type: application/json" \
-d '{
  "credentials": "string",
  "resources": [
    {
      "file_id": "string"
    }
  ]
}
```

```
]
}'
```

Response codes

Status	Schema	Description
200	<code>ResourceValidationOutput</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

GET /api/v1/kb/{kb_id}/sync_configs

List Configs

Parameters

Name	In	Type	Req	Description
<code>kb_id</code>	path	string	yes	

Example — a representative request:

BASH

```
curl -X GET "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/sync_configs" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
200	<code>array<SyncConfigOutput></code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

POST /api/v1/kb/{kb_id}/sync_configs

Add Config

Parameters

Name	In	Type	Req	Description
<code>kb_id</code>	path	string	yes	

Request body — `application/json` (required)

Field	Type	Req	Default	Description
<code>name</code>	string	yes		
<code>sync_root_path</code>	string null			
<code>folder_id</code>	string null			
<code>external_connection_id</code>	string	yes		
<code>drive_id</code>	string null			
<code>file_filter</code>	FileFilterInput null			
<code>labels</code>	array<LabelInput> null			
<code>modified_time_range</code>	ModifiedTimeRangeInput null			
<code>extract_strategy</code>	string null			

Example — a representative request:

BASH

```
curl -X POST "https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/sync_configs" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY" \
-H "Content-Type: application/json" \
-d '{
  "name": "string",
  "external_connection_id": "string"
}'
```

Response codes

Status	Schema	Description
201	<code>SyncConfigOutput</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

Sync Jobs (3)

GET /api/v1/kb/{kb_id}/sync_config/{config_id}/jobs

List Jobs

List sync jobs for a sync config with cursor-based pagination. Use the nextcursor from the response to fetch the next page of results.

Parameters

Name	In	Type	Req	Description
<code>kb_id</code>	path	string	yes	
<code>config_id</code>	path	string	yes	
<code>limit</code>	query	integer		Maximum number of jobs to return (1-100)
<code>cursor</code>	query	string null		Pagination cursor from previous response to fetch next page
<code>order</code>	query	SortOrder		Sort order by creation time: 'desc' (newest first) or 'asc' (oldest first)

Example — a representative request:

```
BASH
curl -X GET
"https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/sync_config/{config_id}/jobs" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
200	<code>PaginatedSyncJobsOutput</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

POST /api/v1/kb/{kb_id}/sync_config/{config_id}/sync

Trigger Sync

Trigger an immediate sync for a sync config. Creates a new pending sync job that will be processed by the sync job runner. Returns 409 Conflict if there's already an active (pending or in-progress) job for the specified config. Args: data: Optional input. Set options.fullsync=True to ignore incremental sync state and perform a full re-sync. Note: S3 and ShareFile providers always perform full sync...

Parameters

Name	In	Type	Req	Description
<code>kb_id</code>	path	string	yes	
<code>config_id</code>	path	string	yes	

Example — a representative request:

```
BASH
curl -X POST
"https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/sync_config/{config_id}/sync" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY" \
-H "Content-Type: application/json" \
-d '{
  "full_sync": false
}'
```

Response codes

Status	Schema	Description
201	<code>SyncJobOutput</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

GET /api/v1/kb/{kb_id}/sync_job/{job_id}/logs

Get Job Logs

Get logs for a sync job with cursor-based pagination. Use the nextcursor from the response to fetch the next page of results. Supports optional filtering by: - level: Log level (DEBUG, INFO, WARNING, ERROR, EXCEPTION, CRITICAL) - startdate / enddate: Timestamp range

Parameters

Name	In	Type	Req	Description
<code>kb_id</code>	path	string	yes	
<code>job_id</code>	path	string	yes	

<code>limit</code>	query	integer		Maximum number of log entries to return (1-500)
<code>cursor</code>	query	integer null		ID-based cursor from previous response to fetch next page (logs with id < cursor)
<code>level</code>	query	LogLevel null		Filter by log level
<code>start_date</code>	query	string null		Filter logs with timestamp >= startdate (ISO 8601 format)
<code>end_date</code>	query	string null		Filter logs with timestamp <= enddate (ISO 8601 format)
<code>order</code>	query	SortOrder		Sort order by ID: 'desc' (newest first) or 'asc' (oldest first)

Example — a representative request:

```
BASH
curl -X GET
"https://$ZONE.dp.progress.cloud/api/v1/kb/$KB/sync_job/{job_id}/logs" \
-H "X-NUCLIA-NUAKEY: Bearer $NUA_KEY"
```

Response codes

Status	Schema	Description
200	<code>PaginatedLogsOutput</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

PART 7 — REFERENCE APPENDIX

Appendix D — Global / Account API Reference

Accounts, users, and zones, fully expanded

Account-level concerns: accounts, users and invites, zones, usage.

27 operations in **7 groups**, generated from the official [global v1](#) OpenAPI specification. Served from <https://rag.progress.cloud/api/v1>; authorized with a user token or PAT (NUA keys are rejected here). Every request-body field, parameter, and response is listed; object-typed fields reference named schemas documented in full in Appendix F (Schema Catalog).

Invites (3)

DELETE /api/v1/account/{account_id}/invite/{email}

Delete Invite

Invalidate an invite to join an account

Parameters

Name	In	Type	Req	Description
<code>account_id</code>	path	string	yes	
<code>email</code>	path	string	yes	

Example — a representative request:

```
BASH
curl -X DELETE
"https://rag.progress.cloud/api/v1/account/$ACCOUNT/invite/{email}" \
-H "Authorization: Bearer $PAT"
```

Response codes

Status	Schema	Description
--------	--------	-------------

204		Successful Response
422	HTTPValidationError	Validation Error

GET /api/v1/account/{account_id}/invites

Invited Users

List of users invited to an account, but not yet joined

Parameters

Name	In	Type	Req	Description
account_id	path	string	yes	

Example — a representative request:

BASH

```
curl -X GET "https://rag.progress.cloud/api/v1/account/$ACCOUNT/invites" \
-H "Authorization: Bearer $PAT"
```

Response codes

Status	Schema	Description
200	array<InvitedUser>	Successful Response
422	HTTPValidationError	Validation Error

POST /api/v1/account/{account_slug}/invite

Invite To Account

Invite somebody by email

Parameters

Name	In	Type	Req	Description
account_slug	path	string	yes	

Request body — `application/json` (required)

Field	Type	Req	Default	Description
-------	------	-----	---------	-------------

email	string	yes		
role	AccountRole		AMEMBER	
came_from	string null			

Example — a representative request:

```
BASH
curl -X POST "https://rag.progress.cloud/api/v1/account/{account_slug}/invite" \
  -H "Authorization: Bearer $PAT" \
  -H "Content-Type: application/json" \
  -d '{
    "email": "user@example.com",
    "role": "AOWNER"
  }'
```

Response codes

Status	Schema	Description
200	object	Successful Response
422	HTTPValidationError	Validation Error

Accounts (9)

GET /api/v1/account/{account_id}/usage

Get Usage

Retrieve statistics for a knowledgebox or account

Parameters

Name	In	Type	Req	Description
account_id	path	string	yes	
aggregation	query	Aggregation		
knowledgebox	query	string null		
nua_key_id	query	string null		

from	query	string	yes	
to	query	string null		

Example — a representative request:

BASH

```
curl -X GET "https://rag.progress.cloud/api/v1/account/$ACCOUNT/usage" \
-H "Authorization: Bearer $PAT"
```

Response codes

Status	Schema	Description
200	<code>array<KbUsageDashboard></code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

GET /api/v1/account/{account_slug_or_id}

Get Account

Parameters

Name	In	Type	Req	Description
<code>account_slug_or_id</code>	path	string	yes	

Example — a representative request:

BASH

```
curl -X GET "https://rag.progress.cloud/api/v1/account/{account_slug_or_id}" \
-H "Authorization: Bearer $PAT"
```

Response codes

Status	Schema	Description
200	<code>Account</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

PATCH /api/v1/account/{account_slug}

Modify Account

Parameters

Name	In	Type	Req	Description
<code>account_slug</code>	path	string	yes	

Request body — `application/json` (required)

Field	Type	Req	Default	Description
<code>title</code>	string null			
<code>description</code>	string null			
<code>saml</code>	OldSAMLConfig null			
<code>saml_config</code>	SAMLConfig null			
<code>slug</code>	string null			
<code>workflow</code>	WorkflowMode null			
<code>eula_accepted</code>	boolean null			

Example — a representative request:

```
BASH
curl -X PATCH "https://rag.progress.cloud/api/v1/account/{account_slug}" \
  -H "Authorization: Bearer $PAT" \
  -H "Content-Type: application/json" \
  -d '{}'
```

Response codes

Status	Schema	Description
200	<code>object</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

DELETE /api/v1/account/{account_slug}

Delete Account

Parameters

Name	In	Type	Req	Description
<code>account_slug</code>	path	string	yes	

Example — a representative request:

BASH

```
curl -X DELETE "https://rag.progress.cloud/api/v1/account/{account_slug}" \
-H "Authorization: Bearer $PAT"
```

Response codes

Status	Schema	Description
204		Successful Response
422	<code>HTTPValidationError</code>	Validation Error

GET /api/v1/account/{account_slug}/logo.png

Download Account Logo

Parameters

Name	In	Type	Req	Description
<code>account_id</code>	query	string	yes	

Example — a representative request:

BASH

```
curl -X GET "https://rag.progress.cloud/api/v1/account/{account_slug}/logo.png" \
-H "Authorization: Bearer $PAT"
```

Response codes

Status	Schema	Description
200		Successful Response
422	<code>HTTPValidationError</code>	Validation Error

PATCH /api/v1/account/{account_slug}/logo.png

Upload Account Logo

Parameters

Name	In	Type	Req	Description
<code>account_id</code>	query	string	yes	

Request body — `multipart/form-data` (required)

Field	Type	Req	Default	Description
<code>file</code>	string	yes		

Example — a representative request:

```
BASH
curl -X PATCH "https://rag.progress.cloud/api/v1/account/{account_slug}/logo.png" \
-H "Authorization: Bearer $PAT"
```

Response codes

Status	Schema	Description
204		Successful Response
422	<code>HTTPValidationError</code>	Validation Error

GET /api/v1/account/{account_slug}/permissions

Get Account Permissions

Parameters

Name	In	Type	Req	Description
<code>account_id</code>	query	string	yes	

Example — a representative request:

```
BASH
curl -X GET
```

```
"https://rag.progress.cloud/api/v1/account/{account_slug}/permissions" \
-H "Authorization: Bearer $PAT"
```

Response codes

Status	Schema	Description
200	<code>object</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

GET /api/v1/accounts

List Accounts

Example — a representative request:

```
BASH
curl -X GET "https://rag.progress.cloud/api/v1/accounts" \
-H "Authorization: Bearer $PAT"
```

Response codes

Status	Schema	Description
200	<code>array<AccountListItem></code>	Successful Response

POST /api/v1/accounts

Create Account

Request body — `application/json` (required)

Field	Type	Req	Default	Description
<code>slug</code>	string null			
<code>title</code>	string	yes		
<code>description</code>	string null			
<code>email</code>	string null			
<code>zone</code>	string null			
<code>workflow</code>	WorkflowMode		<code>classic</code>	

eula_accepted	boolean		False	
---------------	---------	--	-------	--

Example — a representative request:

```
BASH
curl -X POST "https://rag.progress.cloud/api/v1/accounts" \
  -H "Authorization: Bearer $PAT" \
  -H "Content-Type: application/json" \
  -d '{
    "title": "My title",
    "workflow": "classic",
    "eula_accepted": false
  }'
```

Response codes

Status	Schema	Description
201	ItemCreated	Successful Response
422	HTTPValidationError	Validation Error

Zones (2)

GET /api/v1/account/{account_slug_or_id}/zones

List Account Zones

List zones visible to a specific account. Accepts either the account slug or UUID.

Parameters

Name	In	Type	Req	Description
account_slug_or_id	path	string	yes	

Example — a representative request:

```
BASH
curl -X GET
"https://rag.progress.cloud/api/v1/account/{account_slug_or_id}/zones" \
  -H "Authorization: Bearer $PAT"
```

Response codes

Status	Schema	Description
200	<code>array<Zone></code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

GET /api/v1/zones

List Zones

List available zones across all user accounts (union)

Example — a representative request:

```
BASH
curl -X GET "https://rag.progress.cloud/api/v1/zones" \
  -H "Authorization: Bearer $PAT"
```

Response codes

Status	Schema	Description
200	<code>array<Zone></code>	Successful Response

Users (4)

GET /api/v1/account/{account_slug}/user/{user_id}

Get Account User Details

Parameters

Name	In	Type	Req	Description
<code>account_slug</code>	path	string	yes	
<code>user_id</code>	path	string	yes	

Example — a representative request:

```
BASH
curl -X GET
```

```
"https://rag.progress.cloud/api/v1/account/{account_slug}/user/{user_id}" \
-H "Authorization: Bearer $PAT"
```

Response codes

Status	Schema	Description
200	<code>AccountUser</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

GET /api/v1/account/{account_slug}/users

List Account Users

Parameters

Name	In	Type	Req	Description
<code>account_slug</code>	path	string	yes	

Example — a representative request:

```
BASH
curl -X GET "https://rag.progress.cloud/api/v1/account/{account_slug}/users" \
-H "Authorization: Bearer $PAT"
```

Response codes

Status	Schema	Description
200	<code>array<AccountUser></code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

PATCH /api/v1/account/{account_slug}/users

Update Users

Account user modification

Parameters

Name	In	Type	Req	Description
------	----	------	-----	-------------

<code>account_slug</code>	path	string	yes	
---------------------------	------	--------	-----	--

Request body — `application/json` (required)

Field	Type	Req	Default	Description
<code>add</code>	<code>array<AddAccountUser> null</code>			
<code>delete</code>	<code>array<string> null</code>			

Example — a representative request:

```
BASH
curl -X PATCH "https://rag.progress.cloud/api/v1/account/{account_slug}/users" \
-H "Authorization: Bearer $PAT" \
-H "Content-Type: application/json" \
-d '{}'
```

Response codes

Status	Schema	Description
204		Successful Response
422	<code>HTTPValidationError</code>	Validation Error

GET /api/v1/account/{account_slug}/users/search

Search Account Users

Parameters

Name	In	Type	Req	Description
<code>account_slug</code>	path	string	yes	
<code>query</code>	query	string	yes	

Example — a representative request:

```
BASH
curl -X GET
"https://rag.progress.cloud/api/v1/account/{account_slug}/users/search" \
-H "Authorization: Bearer $PAT"
```


Response codes

Status	Schema	Description
200	<code>array<AccountUser></code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

Authenticated User (4)

GET /api/v1/user

Get Authenticated User

Information about user

Example — a representative request:

```
BASH
curl -X GET "https://rag.progress.cloud/api/v1/user" \
  -H "Authorization: Bearer $PAT"
```

Response codes

Status	Schema	Description
200	<code>stashify_idp__api__models__User</code>	Successful Response

PATCH /api/v1/user

Modify User Profile

Set information about user

Request body — `application/json` (required)

Field	Type	Req	Default	Description
<code>name</code>	string null			
<code>avatar</code>	string null			
<code>email</code>	string null			

language	Languages null			
----------	------------------	--	--	--

Example — a representative request:

```
BASH
curl -X PATCH "https://rag.progress.cloud/api/v1/user" \
  -H "Authorization: Bearer $PAT" \
  -H "Content-Type: application/json" \
  -d '{}'
```

Response codes

Status	Schema	Description
200	object	Successful Response
422	HTTPValidationError	Validation Error

DELETE /api/v1/user

Delete Authenticated User

Delete the logged in user

Example — a representative request:

```
BASH
curl -X DELETE "https://rag.progress.cloud/api/v1/user" \
  -H "Authorization: Bearer $PAT"
```

Response codes

Status	Schema	Description
204		Successful Response

GET /api/v1/user/welcome

Welcome

Welcome information for a new user

Example — a representative request:

```
BASH
```

```
curl -X GET "https://rag.progress.cloud/api/v1/user/welcome" \
-H "Authorization: Bearer $PAT"
```

Response codes

Status	Schema	Description
200	Welcome	Successful Response

Personal Access Tokens (3)

DELETE /api/v1/user/pa_token/{token_id}

Delete Token

Delete a personal access token by its ID

Parameters

Name	In	Type	Req	Description
token_id	path	string	yes	

Example — a representative request:

```
BASH
curl -X DELETE "https://rag.progress.cloud/api/v1/user/pa_token/{token_id}" \
-H "Authorization: Bearer $PAT"
```

Response codes

Status	Schema	Description
204		Successful Response
422	HTTPValidationError	Validation Error

GET /api/v1/user/pa_tokens

List Tokens

List all personal access tokens for the authenticated user

Example — a representative request:

```
BASH
curl -X GET "https://rag.progress.cloud/api/v1/user/pa_tokens" \
  -H "Authorization: Bearer $PAT"
```

Response codes

Status	Schema	Description
200	<code>array<TokenItem></code>	Successful Response

POST /api/v1/user/pa_tokens

Create Token

Create a new personal access token

Request body — `application/json` (required)

Field	Type	Req	Default	Description
<code>description</code>	string	yes		
<code>expiration_date</code>	string null			

Example — a representative request:

```
BASH
curl -X POST "https://rag.progress.cloud/api/v1/user/pa_tokens" \
  -H "Authorization: Bearer $PAT" \
  -H "Content-Type: application/json" \
  -d '{
    "description": "string"
  }'
```

Response codes

Status	Schema	Description
201	<code>TokenCreated</code>	Successful Response
422	<code>HTTPValidationError</code>	Validation Error

Authentication - OAuth (2)

POST /oauth2/register

Register Oauth Client

Dynamic Client Registration - create (RFC 7591)

Example — a representative request:

```
BASH
curl -X POST "https://rag.progress.cloud/oauth2/register" \
  -H "Authorization: Bearer $PAT"
```

Response codes

Status	Schema	Description
201	object (free-form map)	Successful Response

PUT /oauth2/register/{client_id}

Update Oauth Client

Dynamic Client Registration - update (RFC 7591)

Parameters

Name	In	Type	Req	Description
client_id	path	string	yes	

Example — a representative request:

```
BASH
curl -X PUT "https://rag.progress.cloud/oauth2/register/{client_id}" \
  -H "Authorization: Bearer $PAT"
```

Response codes

Status	Schema	Description
200	object (free-form map)	Successful Response
422	HTTPValidationError	Validation Error

PART 7 — REFERENCE APPENDIX

Appendix E — Glossary and Schema Reference

Terms, enumerations, and the identifiers you meet in responses

Glossary

Term	Definition
Account	The billing and ownership boundary for an organization. Holds users and Knowledge Boxes. Identified by a UUID and a slug.
Zone	A geographic region a Knowledge Box lives in (e.g. <code>aws-ap-southeast-2-1</code>). A KB belongs to exactly one zone.
Knowledge Box (KB)	The primary container of content, indexes, and configuration. Fully isolated from other KBs. Identified by a <code>kbid</code> .
Resource	One item of content in a KB — a document, page, image, audio, video, or text — plus its metadata. Has a <code>uuid</code> and optional <code>slug</code> .
Field	A typed value inside a resource: text, file, link, conversation, or key-value. A resource can hold several fields.
Paragraph	The retrievable unit — a text block a field is split into during processing. Indexed for keyword and semantic search.
Vector	An embedding of a paragraph or sentence, produced by the KB's semantic model, used for meaning-based retrieval.
Vectorset	A named vector index within a KB. A KB may hold several (e.g. per embedding model), selectable with the <code>vectorset</code> parameter.
Label / Labelset	A tag (label) organized into a named group (labelset), applied to resources or paragraphs and used for filtering and faceting.

Entity / Entity group	A named thing (person, org, place, product) detected by NER or defined by you; grouped into entity groups such as ORG , GPE , DATE .
Relation	A typed connection between entities and/or resources, forming the knowledge graph.
Processing	The asynchronous NUA pipeline that extracts text (OCR/STT), detects language, splits paragraphs, runs NER, and computes vectors.
NUA	The Nuclia Understanding API — processing and Predict services. Involves no Knowledge Box; nothing is stored.
Predict	NUA services that run models directly: embeddings, tokens/NER, LLM chat, query analysis, rerank, and REMi.
Ingestion agent	A data-augmentation agent that enriches resources at ingestion — labeler, generator, graph, Q&A, safety guards.
Retrieval Agent (RAO)	A configurable, multi-source, multi-step agent with drivers and a workflow; the 'agentic' layer above a KB's fixed pipeline.
Driver	A source made available to a Retrieval Agent — a KB, SQL/Snowflake DB, pandas, internet search, or MCP server.
RAG strategy	A method for assembling the context sent to the LLM in /ask — full_resource , neighbouring_paragraphs , hierarchy , graph , etc.
REMi	RAG Evaluation Metrics — a model scoring Answer Relevance, Context Relevance, and Groundedness (the RAG triad).
Service account	A non-human identity on a KB with a role; its API keys authorize content and search calls.
NUA key	A credential for the NUA API (and, with allow_kb_management , KB management). Presented in X-NUCLIA-NUAKEY .
PAT	Personal Access Token — a long-lived user token for admin automation; works on all four APIs.

Response identifiers

Search and citation responses address content with compound identifiers. Learn to read them and the rest of the API becomes legible.

Identifier form	Meaning	Example
<code><rid></code>	A resource UUID	<code>ca2a8f3bf186491facde30cd b2fbc2b6</code>
<code><rid>/<ft>/<field></code>	A field within a resource (ft = field-type letter)	<code>ca2a.../t/body</code>
<code><rid>/<ft>/<field>/ <start>-<end></code>	A paragraph — a character range within a field	<code>ca2a.../t/body/0-301</code>

Field-type letters seen in responses: **t** text, **f** file, **u** link, **a** generic/title, **c** conversation. Citations map one of these paragraph identifiers to the character span(s) of the answer it supports.

Key enumerations

Values verified from the `nucliadb` v1 specification. These are the enums you set most often.

Text field formats (``format``)

Value	Meaning
PLAIN	Plain text
HTML	HTML markup
MARKDOWN	Markdown, converted to plain text on ingest
KEEP_MARKDOWN	Markdown, preserving the markup
RST	reStructuredText
JSON	Structured JSON (understood for structured Q&A)
JSONL	JSON Lines
PLAIN_BLANKLINE_SPLIT	Plain text split on blank lines

Retrieval features (``features``)

<code>`find`</code> (FindOptions)	<code>`ask`</code> (ChatOptions)
keyword	keyword
semantic	semantic
relations	relations

graph

— (use the graph RAG strategy)

Resource detail (`show` — ResourceProperties)

`basic`, `values`, `extracted`, `origin`, `extra`, `relations`, `security`, `errors`. Default is `[basic]`.

Combine several to control how much of a resource is serialized.

Extracted data types (`extracted` — ExtractedDataTypeName)

`text`, `metadata`, `shortened_metadata`, `large_metadata`, `vectors`, `link`, `file`, `question_answers`, `relation_vectors`. Selects which processed artifacts to return when reading a resource with `show=extracted`.

Field types (FieldName)

`text`, `file`, `link`, `conversation`, `generic`, `key_value` — the `field_type_filter` values, controlling which field types appear in results.

Rank fusion and rerankers

`rank_fusion`: `rrf` (Reciprocal Rank Fusion — the default and current option). `reranker`: `predict` (the platform's reranking model — the default). Both can also take an object form for advanced tuning.

RAG strategy names (`rag\_strategies[].name`)

`full_resource`, `field_extension`, `hierarchy`, `neighbouring_paragraphs`, `metadata_extension`, `conversational`, `prequeries`, `graph`. Image strategies (`rag_images_strategies`): `page_image`, `paragraph_image`, `table_image`.

Service-account roles

`SREADER` (read/search), `SWRITER` (write content), `SCONTRIBUTOR` (write, not configure), `SOWNER` (full control). Knowledge Box user roles: `READER`, `WRITER`, `MANAGER`.

The four base URLs, one more time

API	Base URL	Auth header
Global	<code>https://rag.progress.cloud/api/v1</code>	<code>Authorization: Bearer (user/PAT)</code>
Zone	<code>https://<zone>.dp.progress.cloud</code>	<code>X-NUCLIA-NUAKEY</code> or

		Authorization: Bearer
NucliaDB	https:// <zone>.dp.progress.cloud /api/v1	Authorization: Bearer (KB key)
NUA	https:// <zone>.dp.progress.cloud /api/v1	X-NUCLIA-NUAKEY: Bearer

Key idea If a call returns 403, check this table first. The overwhelmingly common cause is a credential presented to an API it is not cut for — the error message will name the context it could not access.

PART 7 — REFERENCE APPENDIX

Appendix F — Complete Schema Catalog

Every data model in all four APIs, documented field by field

This catalog documents every named schema referenced anywhere in the four OpenAPI specifications — request bodies, response models, and every nested object type. When a field in Appendices A–D shows an object type by name (for example `FilterExpression`, `AskRequest`, or `SyncAskResponse`), its full definition is here. Schemas are grouped by API and listed alphabetically.

1140 schemas in total: 340 (NucliaDB), 481 (NUA), 124 (Zone), 195 (Global). Names may repeat across APIs where a model is shared; each API's copy is documented under its own section for accuracy.

nucliadb — 340 schemas

AgentsFilter

Field	Type	Req	Default	Description
<code>type</code>	AgentType	yes		
<code>task_names</code>	array<string>			list of task names. If None or empty, all tasks for that operation are applied.

AgentType

Enum: `graph`, `label`, `ask`, `qa`, `extract`, `prompt_guard`, `llama_guard`

AITables

Field	Type	Req	Default	Description
<code>llm</code>	LLMConfig null			

And_FieldFilterExpressionType_

Field	Type	Req	Default	Description
<code>operands</code>	array<And_FieldFilterExpressionType_ Or_FieldFilterExpressionType_ Not_FieldFilterExpressionType_ Resource-Input Field Keyword DateCreated DateModified Label ResourceMimetype FieldMimetype Entity-Input Language OriginTag OriginMetadata OriginPath OriginSource OriginCollaborator nucliadb_models__filters__Generated>	yes		

And_GraphNodesQuery_

Field	Type	Req	Default	Description
<code>operands</code>	array<And_GraphNodesQuery_ Or_GraphNodesQuery_ Not_GraphNodesQuery_ AnyNode nucliadb_models__graph__requests__Generated>	yes		

And_GraphPathQuery_

Field	Type	Req	Default	Description
<code>operands</code>	array<And_GraphPathQuery_ Or_GraphPathQuery_	yes		

	 Not_GraphPathQuery_ _ GraphPath-Input SourceNode DestinationNode AnyNode Relation- Input nucliadb_models__gr aph__requests__Gene rated>			
--	---	--	--	--

And_GraphRelationsQuery_

Field	Type	Req	Default	Description
<code>operands</code>	array<And_GraphRel ationsQuery_ Or_GraphRelationsQ uery_ Not_GraphRelations Query_ Relation- Input nucliadb_models__gr aph__requests__Gene rated>	yes		

And_KVFilterExpression_

Field	Type	Req	Default	Description
<code>operands</code>	array<And_KVFilter Expression_ Or_KVFilterExpressi on_ Not_KVFilterExpress ion_ Eq Inequalities Contains>	yes		

And_ParagraphFilterExpression_

Field	Type	Req	Default	Description
<code>operands</code>	array<And_ParagraphFilterExpression_	yes		

	Or_ParagraphFilterExpression_ Not_ParagraphFilterExpression_ Label Kind>			
--	--	--	--	--

And_ResourceFilterExpression_

Field	Type	Req	Default	Description
<code>operands</code>	array<And_ResourceFilterExpression_ Or_ResourceFilterExpression_ Not_ResourceFilterExpression_ Resource-Input DateCreated DateModified Label ResourceMimetype Language OriginTag OriginMetadata OriginPath OriginSource OriginCollaborator Status>	yes		

AnonimizationModel

An enumeration.

Enum: `disabled`, `multilingual`

Answer

Field	Type	Req	Default	Description
<code>text</code>	string	yes		
<code>language</code>	string null			
<code>ids_paragraphs</code>	array<string>	yes		

AnthropicKey

Field	Type	Req	Default	Description
key	string			

AnyNode

Field	Type	Req	Default	Description
prop	string		node	
value	string null			
match	NodeMatchKindName		exact	
type	RelationNodeType null		entity	
group	string null			

AppliedDataAugmentation

Field	Type	Req	Default	Description
qas	QuestionAnswers null			Question and answers generated by the Question Answers agent
new_text_fields	array<NewTextField>			New text fields. Only generated by the Generator agent as of now.
changed	boolean		True	Indicates if the FieldMetadata was changed by the agents

AskConfig

Field	Type	Req	Default	Description
audit_metadata	map<string, string> null			A dictionary containing optional

				audit-specific metadata, such as userid, environment, or other contextual information. This metadata can be leveraged for filtering and analyzing activity logs in futu...
top_k	integer		20	The top most relevant results to fetch at the retrieval step. The maximum number of results allowed is 200.
filter_expression	FilterExpression null			Returns only documents that match this filter expression. Filtering examples can be found here: (see docs) This allows building complex filtering expressions and replaces the following parameters: field...
fields	array<string>		[]	The list of fields to search in. For instance: a/title to search only on title field. For more details on filtering by field, see: (see docs)
filters	array<string> array<Filter>		[]	The list of filters to apply. Filtering examples can be found here: (see docs)
keyword_filters	array<string> array<Filter>		[]	List of keyword filter expressions to apply to the retrieval step. The text block search will only be performed on the documents that

				contain the specified keywords. The filters are case-insensitive, ...
<code>vectorset</code>	string null			Vectors index to perform the search in. If not provided, NucliaDB will use the default one
<code>min_score</code>	number MinScore null			Minimum score to filter search results. Results with a lower score will be ignored. Accepts either a float or a dictionary with the minimum scores for the bm25 and vector indexes. If a float is provid...
<code>features</code>	array<ChatOptions>		<code>['semantic', 'keyword']</code>	Features enabled for the chat endpoint. Semantic search is done if semantic is included. If keyword is included, the results will include matching paragraphs from the bm25 index. If relations is inclu...
<code>range_creation_start</code>	string null			Resources created before this date will be filtered out of search results. Datetime are represented as a str in ISO 8601 format, like: 2008-09-15T15:53:00+05:00.
<code>range_creation_end</code>	string null			Resources created after this date will be filtered out of search results. Datetime are represented as a str in ISO 8601 format, like: 2008-09-

				15T15:53:00+05:00.
<code>range_modification_start</code>	string null			Resources modified before this date will be filtered out of search results. Datetime are represented as a str in ISO 8601 format, like: 2008-09-15T15:53:00+05:00.
<code>range_modification_end</code>	string null			Resources modified after this date will be filtered out of search results. Datetime are represented as a str in ISO 8601 format, like: 2008-09-15T15:53:00+05:00.
<code>show</code>	array<ResourceProperties>		<code>['basic']</code>	Controls which types of metadata are serialized on resources of search results
<code>field_type_filter</code>	array<FieldTypeName>		<code>['text', 'file', 'link', 'conversation', 'generic', 'key_value']</code>	Define which field types are serialized on resources of search results
<code>extracted</code>	array<ExtractedData TypeName>		<code>[]</code>	[Deprecated] Please use GET resource endpoint instead to get extracted metadata
<code>context</code>	array<ChatContextMessage> null			DEPRECATED! Please, use chathistory instead.
<code>chat_history</code>	array<ChatContextMessage> null			Use to rephrase the new LLM query by taking into account the chat conversation history. This will be passed to the LLM so that it is aware of the

				previous conversation.
<code>extra_context</code>	<code>array<string> null</code>			Additional context that is added to the retrieval context sent to the LLM. It allows extending the chat feature with content that may not be in the Knowledge Box.
<code>extra_context_images</code>	<code>array<Image> null</code>			Additional images added to the retrieval context sent to the LLM." It allows extending the chat feature with content that may not be in the Knowledge Box.
<code>query_image</code>	<code>Image null</code>			Image that will be used together with the query text for retrieval and then sent to the LLM as part of the context. If a query image is provided, the <code>extracontextimages</code> and <code>ragimagesstrategies</code> will be...
<code>highlight</code>	<code>boolean</code>		<code>False</code>	If set to true, the query terms will be highlighted in the results between <code><mark>...</mark></code> tags
<code>resource_filters</code>	<code>array<string></code>		<code>[]</code>	List of resource ids to filter search results for. Only paragraphs from the specified resources will be returned.
<code>prompt</code>	<code>string CustomPrompt null</code>			Use to customize the prompts given to the generative model.

				Both system and user prompts can be customized. If a string is provided, it is interpreted as the user prompt.
<code>rank_fusion</code>	RankFusionName ReciprocalRankFusion		<code>rrf</code>	Rank fusion algorithm to use to merge results from multiple retrievers (keyword, semantic)
<code>reranker</code>	RerankerName PredictReranker		<code>predict</code>	Reranker let you specify which method you want to use to rerank your results at the end of retrieval
<code>citations</code>	boolean CitationsType null			Whether to include citations in the response. If set to None or False, no citations will be computed. If set to True or 'default', citations will be computed after answer generation and send as a sepa...
<code>citation_threshold</code>	number null			If citations is set to True or 'default', this will be the similarity threshold. Value between 0 and 1, lower values will produce more citations. If not set, it will be set to the optimized threshold ...
<code>security</code>	RequestSecurity null			Security metadata for the request. Please refer to the documentation for more details on how security works: (see docs)

<code>show_hidden</code>	boolean		<code>False</code>	If set to false (default), excludes hidden resources from search
<code>rag_strategies</code>	array<FieldExtensionStrategy FullResourceStrategy HierarchyResourceStrategy NeighbouringParagraphsStrategy MetadataExtensionStrategy ConversationalStrategy PreQueriesStrategy GraphStrategy>		<code>[]</code>	Options for tweaking how the context for the LLM model is crafted: - fullresource will add the full text of the matching resources to the context. This strategy cannot be combined with hierarchy, neig...
<code>rag_images_strategies</code>	array<PageImageStrategy ParagraphImageStrategy TableImageStrategy>		<code>[]</code>	Options for tweaking how the image based context for the LLM model is crafted: - pageimage will add the full page image of the matching resources to the context. - tables will send the table images fo...
<code>debug</code>	boolean		<code>False</code>	If set, the response will include some extra metadata for debugging purposes, like the list of queried nodes.
<code>generative_model</code>	string null			The generative model to use for the chat endpoint. If not provided, the model configured for the Knowledge Box is used.
<code>generative_model_seed</code>	integer null			The seed to use for the generative model for deterministic generation. Only

				supported by some models.
<code>max_tokens</code>	integer MaxTokens null			Use to limit the amount of tokens used in the LLM context and/or for generating the answer. If not provided, the default maximum tokens of the generative model will be used. If an integer is provided,...
<code>rephrase</code>	boolean		<code>False</code>	Rephrase the query for a more efficient retrieval. This will consume LLM tokens and make the request slower.
<code>chat_history_relevance_threshold</code>	number null			Threshold to determine if the past chat history is relevant to rephrase the user's question. 0 - Always treat previous messages as relevant (always rephrase). 1 - Always treat previous messages as irre...
<code>prefer_markdown</code>	boolean		<code>False</code>	If set to true, the response will be in markdown format
<code>answer_json_schema</code>	object (free-form map) null			Desired JSON schema for the LLM answer. This schema is passed to the LLM so that it answers in a structured format following the schema. If not provided, textual response is returned. Note that when u...
<code>generate_answer</code>	boolean		<code>True</code>	Whether to generate

er				an answer using the generative model. If set to false, the response will only contain the retrieval results.
reasoning	Reasoning boolean		False	Reasoning options for the generative model. Set to True to enable default reasoning, False to disable, or provide a Reasoning object for custom options.
query	string null			

AskRequest

Field	Type	Req	Default	Description
audit_metadata	map<string, string> null			A dictionary containing optional audit-specific metadata, such as userid, environment, or other contextual information. This metadata can be leveraged for filtering and analyzing activity logs in futu...
query	string	yes		The query to get a generative answer for
top_k	integer		20	The top most relevant results to fetch at the retrieval step. The maximum number of results allowed is 200.
filter_expression	FilterExpression null			Returns only documents that match this filter expression. Filtering examples can be

				found here: (see docs) This allows building complex filtering expressions and replaces the following parameters:field...
<code>fields</code>	<code>array<string></code>		<code>[]</code>	The list of fields to search in. For instance: a/title to search only on title field. For more details on filtering by field, see: (see docs)
<code>filters</code>	<code>array<string> array<Filter></code>		<code>[]</code>	The list of filters to apply. Filtering examples can be found here: (see docs)
<code>keyword_filters</code>	<code>array<string> array<Filter></code>		<code>[]</code>	List of keyword filter expressions to apply to the retrieval step. The text block search will only be performed on the documents that contain the specified keywords. The filters are case-insensitive, ...
<code>vectorset</code>	<code>string null</code>			Vectors index to perform the search in. If not provided, NucliaDB will use the default one
<code>min_score</code>	<code>number MinScore null</code>			Minimum score to filter search results. Results with a lower score will be ignored. Accepts either a float or a dictionary with the minimum scores for the bm25 and vector indexes. If a float is provid...

<code>features</code>	<code>array<ChatOptions></code>		<code>['semantic', 'keyword']</code>	Features enabled for the chat endpoint. Semantic search is done if semantic is included. If keyword is included, the results will include matching paragraphs from the bm25 index. If relations is inclu...
<code>range_creation_start</code>	<code>string null</code>			Resources created before this date will be filtered out of search results. Datetime are represented as a str in ISO 8601 format, like: 2008-09-15T15:53:00+05:00.
<code>range_creation_end</code>	<code>string null</code>			Resources created after this date will be filtered out of search results. Datetime are represented as a str in ISO 8601 format, like: 2008-09-15T15:53:00+05:00.
<code>range_modification_start</code>	<code>string null</code>			Resources modified before this date will be filtered out of search results. Datetime are represented as a str in ISO 8601 format, like: 2008-09-15T15:53:00+05:00.
<code>range_modification_end</code>	<code>string null</code>			Resources modified after this date will be filtered out of search results. Datetime are represented as a str in ISO 8601 format, like: 2008-09-15T15:53:00+05:00.
<code>show</code>	<code>array<ResourceProperties></code>		<code>['basic']</code>	Controls which types of metadata are

				serialized on resources of search results
<code>field_type_filter</code>	<code>array<FieldTypeNam e></code>		<code>['text', 'file', 'link', 'conversation', 'generic', 'key_value']</code>	Define which field types are serialized on resources of search results
<code>extracted</code>	<code>array<ExtractedData TypeName></code>		<code>[]</code>	[Deprecated] Please use GET resource endpoint instead to get extracted metadata
<code>context</code>	<code>array<ChatContextM essage> null</code>			DEPRECATED! Please, use chathistory instead.
<code>chat_history</code>	<code>array<ChatContextM essage> null</code>			Use to rephrase the new LLM query by taking into account the chat conversation history. This will be passed to the LLM so that it is aware of the previous conversation.
<code>extra_context</code>	<code>array<string> null</code>			Additional context that is added to the retrieval context sent to the LLM. It allows extending the chat feature with content that may not be in the Knowledge Box.
<code>extra_context_images</code>	<code>array<Image> null</code>			Additional images added to the retrieval context sent to the LLM." It allows extending the chat feature with content that may not be in the Knowledge Box.
<code>query_image</code>	<code>Image null</code>			Image that will be used together with

				the query text for retrieval and then sent to the LLM as part of the context. If a query image is provided, the <code>extracontextimages</code> and <code>ragimagesstrategies</code> will be...
<code>highlight</code>	boolean		<code>False</code>	If set to true, the query terms will be highlighted in the results between <code><mark>...</mark></code> tags
<code>resource_filters</code>	array<string>		<code>[]</code>	List of resource ids to filter search results for. Only paragraphs from the specified resources will be returned.
<code>prompt</code>	string CustomPrompt null			Use to customize the prompts given to the generative model. Both system and user prompts can be customized. If a string is provided, it is interpreted as the user prompt.
<code>rank_fusion</code>	RankFusionName ReciprocalRankFusion		<code>rrf</code>	Rank fusion algorithm to use to merge results from multiple retrievers (keyword, semantic)
<code>reranker</code>	RerankerName PredictReranker		<code>predict</code>	Reranker let you specify which method you want to use to rerank your results at the end of retrieval
<code>citations</code>	boolean CitationsType null			Whether to include citations in the response. If set to None or False, no

				citations will be computed. If set to True or 'default', citations will be computed after answer generation and send as a sepa...
<code>citation_threshold</code>	number null			If citations is set to True or 'default', this will be the similarity threshold. Value between 0 and 1, lower values will produce more citations. If not set, it will be set to the optimized threshold ...
<code>security</code>	RequestSecurity null			Security metadata for the request. Please refer to the documentation for more details on how security works: (see docs)
<code>show_hidden</code>	boolean		<code>False</code>	If set to false (default), excludes hidden resources from search
<code>rag_strategies</code>	array<FieldExtensionStrategy FullResourceStrategy HierarchyResourceStrategy NeighbouringParagraphsStrategy MetadataExtensionStrategy ConversationalStrategy PreQueriesStrategy GraphStrategy>		<code>[]</code>	Options for tweaking how the context for the LLM model is crafted: - fullresource will add the full text of the matching resources to the context. This strategy cannot be combined with hierarchy, neig...
<code>rag_images_strategies</code>	array<PageImageStrategy ParagraphImageStrat		<code>[]</code>	Options for tweaking how the image based context for the LLM

	egy TableImageStrategy>			model is crafted: - pageimage will add the full page image of the matching resources to the context. - tables will send the table images fo...
debug	boolean		False	If set, the response will include some extra metadata for debugging purposes, like the list of queried nodes.
generative_model	string null			The generative model to use for the chat endpoint. If not provided, the model configured for the Knowledge Box is used.
generative_model_seed	integer null			The seed to use for the generative model for deterministic generation. Only supported by some models.
max_tokens	integer MaxTokens null			Use to limit the amount of tokens used in the LLM context and/or for generating the answer. If not provided, the default maximum tokens of the generative model will be used. If an integer is provided,...
rephrase	boolean		False	Rephrase the query for a more efficient retrieval. This will consume LLM tokens and make the request slower.
chat_history	number null			Threshold to

relevance_threshold				determine if the past chat history is relevant to rephrase the user's question. 0 - Always treat previous messages as relevant (always rephrase).1 - Always treat previous messages as irre...
prefer_markdown	boolean		False	If set to true, the response will be in markdown format
answer_json_schema	object (free-form map) null			Desired JSON schema for the LLM answer. This schema is passed to the LLM so that it answers in a structured format following the schema. If not provided, textual response is returned. Note that when u...
generate_answer	boolean		True	Whether to generate an answer using the generative model. If set to false, the response will only contain the retrieval results.
search_configuration	string null			Load ask parameters from this configuration. Parameters in the request override parameters from the configuration.
reasoning	Reasoning boolean		False	Reasoning options for the generative model. Set to True to enable default reasoning, False to disable, or provide a Reasoning object for custom

				options.
--	--	--	--	----------

AskRetrievalMatch

Field	Type	Req	Default	Description
<code>id</code>	string	yes		Id of the matching text block

AskSearchConfiguration

Field	Type	Req	Default	Description
<code>kind</code>	string	yes		
<code>config</code>	AskConfig	yes		

AskTimings

Field	Type	Req	Default	Description
<code>generative_first_chunk</code>	number null			Time the LLM took to generate the first chunk of the answer
<code>generative_total</code>	number null			Total time the LLM took to generate the answer

AskTokens

Field	Type	Req	Default	Description
<code>input</code>	integer	yes		Number of LLM tokens used for the context in the query
<code>output</code>	integer	yes		Number of LLM tokens used for the answer
<code>input_nuclia</code>	number null			Number of Nuclia LLM tokens used for the context in the query

<code>output_nuclia</code>	number null			Number of Nuclia LLM tokens used for the answer
----------------------------	---------------	--	--	---

AugmentedContext

Field	Type	Req	Default	Description
<code>paragraphs</code>	map<string, AugmentedTextBlock>		<code>{}</code>	Paragraphs added to the context as a result of using the ragstrategies parameter, typically the neighbouring paragraphs or the conversation strategies
<code>fields</code>	map<string, AugmentedTextBlock>		<code>{}</code>	Field extracted texts added to the context as a result of using the ragstrategies parameter, typically the hierarchy or fullresource strategies.

AugmentedField

Field	Type	Req	Default	Description
<code>metadata</code>	FieldMetadata	yes		
<code>applied_data_augmentation</code>	AppliedDataAugmentation	yes		
<code>input_nuclia_tokens</code>	number	yes		
<code>output_nuclia_tokens</code>	number	yes		
<code>time</code>	number	yes		

AugmentedTextBlock

Field	Type	Req	Default	Description
-------	------	-----	---------	-------------

<code>id</code>	string	yes		The id of the augmented text bloc. It can be a paragraph id or a field id.
<code>text</code>	string	yes		The text of the augmented text block. It may include additional metadata to enrich the context
<code>position</code>	TextPosition null			Metadata about the position of the text block in the original document.
<code>parent</code>	string null			The parent text block that was augmented for.
<code>augmentation_type</code>	TextBlockAugmentationType	yes		Type of augmentation.

Author

Enum: `NUCLIA`, `USER`

AzureMistralKey

Field	Type	Req	Default	Description
<code>key</code>	string			
<code>url</code>	string			

AzureOpenAIKey

Field	Type	Req	Default	Description
<code>key</code>	string			
<code>url</code>	string			
<code>deployment</code>	string			
<code>model</code>	string			

CatalogFilterExpression

Returns only documents that match this filter expression. Filtering examples can be found here: (see docs)

This allows building complex filtering expressions and replaces the following parameters: `filters`, `range_*`, `with_status`.

Field	Type	Req	Default	Description
<code>resource</code>	And_ResourceFilterExpression_ Or_ResourceFilterExpression_ Not_ResourceFilterExpression_ Resource-Input DateCreated DateModified Label ResourceMimetype Language OriginTag OriginMetadata OriginPath OriginSource OriginCollaborator Status	yes		Filter to apply to resources

CatalogQuery

Field	Type	Req	Default	Description
<code>field</code>	CatalogQueryField		<code>title</code>	Field to search in
<code>match</code>	CatalogQueryMatch		<code>exact</code>	Operator to use for matching results
<code>query</code>	string	yes		Text to search for

CatalogQueryField

Enum: `title`, `slug`

CatalogQueryMatch

Enum: `exact`, `words`, `fuzzy`, `starts_with`, `ends_with`, `contains`

CatalogRequest

Field	Type	Req	Default	Description
-------	------	-----	---------	-------------

<code>query</code>	<code>string CatalogQuery</code>			The query to search for
<code>filter_expression</code>	<code>CatalogFilterExpression null</code>			Returns only documents that match this filter expression. Filtering examples can be found here: (see docs) This allows building complex filtering expressions and replaces the following parameters: <code>filter...</code>
<code>faceted</code>	<code>array<string></code>		<code>[]</code>	The list of facets to calculate. The facets follow the same syntax as filters: (see docs)
<code>sort</code>	<code>SortOptions null</code>			Options for results sorting
<code>page_number</code>	<code>integer</code>		<code>0</code>	The page number of the results to return
<code>page_size</code>	<code>integer</code>		<code>20</code>	The number of results to return per page. The maximum number of results per page allowed is 200.
<code>hidden</code>	<code>boolean null</code>			Set to filter only hidden or only non-hidden resources. Default is to return everything
<code>show</code>	<code>array<ResourceProperties></code>		<code>['basic', 'errors']</code>	Controls which types of metadata are serialized on resources of search results
<code>filters</code>	<code>array<string> array<Filter></code>		<code>[]</code>	The list of filters to apply. Filtering examples can be found here: (see docs)

<code>with_status</code>	ResourceProcessingStatus null			Filter results by resource processing status
<code>range_creation_start</code>	string null			Resources created before this date will be filtered out of search results. Datetime are represented as a str in ISO 8601 format, like: 2008-09-15T15:53:00+05:00.
<code>range_creation_end</code>	string null			Resources created after this date will be filtered out of search results. Datetime are represented as a str in ISO 8601 format, like: 2008-09-15T15:53:00+05:00.
<code>range_modification_start</code>	string null			Resources modified before this date will be filtered out of search results. Datetime are represented as a str in ISO 8601 format, like: 2008-09-15T15:53:00+05:00.
<code>range_modification_end</code>	string null			Resources modified after this date will be filtered out of search results. Datetime are represented as a str in ISO 8601 format, like: 2008-09-15T15:53:00+05:00.

ChatContextMessage

Field	Type	Req	Default	Description
<code>author</code>	Author	yes		
<code>text</code>	string	yes		

ChatOptions

Enum: `keyword`, `relations`, `semantic`

CitationsType

Enum: `none`, `default`, `llm_footnotes`

Classification

Field	Type	Req	Default	Description
<code>labelset</code>	string	yes		
<code>label</code>	string	yes		

CloudLink

Field	Type	Req	Default	Description
<code>uri</code>	string null			
<code>size</code>	integer null			
<code>content_type</code>	string null			
<code>filename</code>	string null			
<code>md5</code>	string null			

ComputedMetadata

The purpose of this field is to show a cherry-picked set of fields from computed metadata without having to load the whole computed metadata field.

Field	Type	Req	Default	Description
<code>field_classifications</code>	array<FieldClassification>		<code>[]</code>	

Consumption

Field	Type	Req	Default	Description
<code>normalized_to</code>	TokensDetail	yes		

kens				
customer_key_tokens	TokensDetail	yes		

Contains

Computes whether a value exists inside a range or a repeated field

Field	Type	Req	Default	Description
schema_id	string	yes		
key	string	yes		
contains	integer number string	yes		

ConversationalStrategy

Field	Type	Req	Default	Description
name	string		conversation	
attachments_text	boolean		False	Add attachments on context retrieved on conversation
attachments_images	boolean		False	Add attachments images on context retrieved on conversation if they are mime type image and using a visual LLM
full	boolean		False	Add all conversation fields on matched blocks
max_messages	integer		15	Max messages to append in case its not full field

ConversationFieldData

Field	Type	Req	Default	Description
-------	------	-----	---------	-------------

value	FieldConversation null			
extracted	ConversationFieldExtractedData null			
error	Error null			
status	string null			
errors	array<Error> null			

ConversationFieldExtractedData

Field	Type	Req	Default	Description
text	ExtractedText null			
metadata	FieldComputedMetadata null			
large_metadata	LargeComputedMetadata null			
vectors	VectorObject null			
question_answers	FieldQuestionAnswers null			
relation_node_vectors	map<string, array<RelationNodeVector>> null			
relation_edge_vectors	map<string, array<RelationEdgeVector>> null			

CreateResourcePayload

Field	Type	Req	Default	Description
title	string null			
summary	string null			
slug	string null			The slug is the user-defined id for the resource
icon	string null			The icon should be a media type string:

				(see docs)
thumbnail	string null			
metadata	InputMetadata null			Generic metadata for the resource. It can be used to store structured information about the resource that later is serialized on retrieval results, however this metadata can not be used for searching ...
usermetadata	UserMetadata null			
fieldmetadata	array<UserFieldMetadata> null			
origin	InputOrigin null			Origin metadata for the resource. Used to store information about the resource on the origin system. Most of its fields can later be used to filter at search time.
extra	Extra null			Extra metadata for the resource. It can be used to store structured information about the resource that can't be used to query at retrieval time.
hidden	boolean null			Set the hidden status of the resource. If not set, the default value for new resources in the KnowledgeBox will be used.
files	object		{ }	Dictionary of file fields to be added to the resource. The keys correspond to the field id, and must comply with the regex: <code>^[a-zA-Z0-9:-]</code>

				+\$
<code>links</code>	object		<code>{}</code>	Dictionary of link fields to be added to the resource. The keys correspond to the field id, and must comply with the regex: <code>^[a-zA-Z0-9:-]+\$</code>
<code>texts</code>	object		<code>{}</code>	Dictionary of text fields to be added to the resource. The keys correspond to the field id, and must comply with the regex: <code>^[a-zA-Z0-9:-]+\$</code>
<code>conversations</code>	object		<code>{}</code>	Dictionary of conversation fields to be added to the resource. The keys correspond to the field id, and must comply with the regex: <code>^[a-zA-Z0-9:-]+\$</code>
<code>key_values</code>	object		<code>{}</code>	Dictionary of key-value fields to be added to the resource. The key must be the schema name and must comply with the regex: <code>^[a-zA-Z0-9:-]+\$</code>
<code>processing_options</code>	PushProcessingOptions null		<code>{'ml_text': True}</code>	Options for processing the resource. If not set, the default options will be used.
<code>security</code>	ResourceSecurity null			Security metadata for the resource. It can be used to have fine-grained control over who can access the resource.

CustomPrompt

Field	Type	Req	Default	Description
<code>system</code>	string null			System prompt given to the generative model responsible of generating the answer. This can help customize the behavior of the model when generating the answer. If not specified, the default model prov...
<code>user</code>	string null			User prompt given to the generative model responsible of generating the answer. Use the words {context} and {question} in brackets where you want those fields to be placed, in case you want them in yo...
<code>rephrase</code>	string null			Rephrase prompt given to the generative model responsible for rephrasing the query for a more effective retrieval step. This is only used if the rephrase flag is set to true in the request. If not spe...

CustomSplitStrategy**Enum:** `0`, `1`, `2`**DateCreated**

Matches all fields created in a date range

Field	Type	Req	Default	Description
prop	string		created	
since	string null			Start of the date range. Leave blank for unbounded
until	string null			End of the date range. Leave blank for unbounded

DateModified

Matches all fields modified in a date range

Field	Type	Req	Default	Description
prop	string		modified	
since	string null			Start of the date range. Leave blank for unbounded
until	string null			End of the date range. Leave blank for unbounded

DestinationNode

Field	Type	Req	Default	Description
prop	string		destination_node	
value	string null			
match	NodeMatchKindName		exact	
type	RelationNodeType null		entity	
group	string null			

DirectionalRelation

Field	Type	Req	Default	Description
-------	------	-----	---------	-------------

entity	string	yes		
entity_type	RelationNodeType	yes		
entity_subtype	string	yes		
relation	RelationType	yes		
relation_label	string	yes		
direction	RelationDirection	yes		
metadata	RelationMetadata null			
resource_id	string	yes		

DummyIndexProvider

Field	Type	Req	Default	Description
type	ExternalIndexProviderType		unset	Enum for the different external index providers. For now none are supported, but we may add some in the future.

EntitiesGroup

Field	Type	Req	Default	Description
title	string null			Title of the entities group
color	string null			Color of the entities group. This is for display purposes only.
custom	boolean		False	Denotes if it has been created by the user
entities	map<string,		{}	

	nucliadb_models_entities__Entity>			
--	-----------------------------------	--	--	--

EntitiesGroupSummary

Field	Type	Req	Default	Description
title	string null			Title of the entities group
color	string null			Color of the entities group. This is for display purposes only.
custom	boolean		False	Denotes if it has been created by the user
entities	map<string, nucliadb_models_entities__Entity>		{}	This field is deprecated and will be removed in future versions. It will always be empty. Use the <code>/api/v1/kb/{kbid}/entitiesgroup/{group}</code> endpoint to get the entities of a group.

Entity

Matches fields that contains a detected entity

Field	Type	Req	Default	Description
prop	string		entity	
subtype	string	yes		Type of the entity. e.g: PERSON
value	string null			Value of the entity. e.g: Anna. If blank, matches any entity of the given type

Entity-Input

Matches fields that contains a detected entity

Field	Type	Req	Default	Description
<code>prop</code>	string		<code>entity</code>	
<code>subtype</code>	string	yes		Type of the entity. e.g: PERSON
<code>value</code>	string null			Value of the entity. e.g: Anna. If blank, matches any entity of the given type

Entity-Output

Field	Type	Req	Default	Description
<code>token</code>	string null			
<code>root</code>	string null			
<code>type</code>	string null			

EntitySubgraph

Field	Type	Req	Default	Description
<code>related_to</code>	array<DirectionalRelation>	yes		

Eq

Equal (==) operator

Field	Type	Req	Default	Description
<code>schema_id</code>	string	yes		
<code>key</code>	string	yes		
<code>eq</code>	boolean integer number string	yes		

Error

Field	Type	Req	Default	Description
-------	------	-----	---------	-------------

<code>body</code>	string	yes		
<code>code</code>	integer	yes		
<code>code_str</code>	string	yes		
<code>created</code>	string null	yes		
<code>severity</code>	string	yes		

ExternalIndexProviderType

Enum for the different external index providers. For now none are supported, but we may add some in the future.

Enum: `unset`

Extra

Field	Type	Req	Default	Description
<code>metadata</code>	object (free-form map)	yes		Arbitrary JSON metadata provided by the user that is not meant to be searchable, but can be serialized on results.

ExtractConfig

Field	Type	Req	Default	Description
<code>name</code>	string			
<code>vllm_config</code>	VLLMExtractionConfig null			
<code>ai_tables</code>	AITables null			
<code>split</code>	SplitConfig null			

ExtractedDataTypeName

Enum: `text`, `metadata`, `shortened_metadata`, `large_metadata`, `vectors`, `link`, `file`, `question_answers`, `relation_vectors`

ExtractedText

Field	Type	Req	Default	Description
<code>text</code>	string null			
<code>split_text</code>	map<string, string> null			
<code>deleted_splits</code>	array<string> null			

FeedbackRequest

Field	Type	Req	Default	Description
<code>ident</code>	string	yes		Id of the request to provide feedback for. This id is returned in the response header Nuclia-Learning-Id of the chat endpoint.
<code>good</code>	boolean	yes		Whether the result was good or not
<code>task</code>	FeedbackTasks	yes		The task the feedback is for. For now, only CHAT task is available
<code>feedback</code>	string null			Feedback text
<code>text_block_id</code>	string null			Text block id

FeedbackTasksEnum: `CHAT`**Field**

Matches a field or set of fields

Field	Type	Req	Default	Description
<code>prop</code>	string		<code>field</code>	
<code>type</code>	FieldTypeName	yes		Type of the field to match,

<code>name</code>	string null			Name of the field to match. If blank, matches all fields of the given type
-------------------	---------------	--	--	--

FieldClassification

Field	Type	Req	Default	Description
<code>field</code>	FieldID	yes		
<code>classifications</code>	array<Classification>		<code>[]</code>	

FieldComputedMetadata

Field	Type	Req	Default	Description
<code>metadata</code>	FieldMetadata	yes		
<code>split_metadata</code>	map<string, FieldMetadata> null			
<code>deleted_splits</code>	array<string> null			

FieldConversation

This is a metadata representation of a conversation about how many pages of messages and total of messages we have. This class is used mainly when exposing a conversation in the resource level

Field	Type	Req	Default	Description
<code>pages</code>	integer null			
<code>size</code>	integer null			
<code>total</code>	integer null			
<code>extract_strategy</code>	string null			
<code>split_strategy</code>	string null			

FieldEntities

Wrapper for the entities extracted from a field (required because protobuf doesn't support lists of lists)

Field	Type	Req	Default	Description
<code>entities</code>	array<FieldEntity>	yes		

FieldEntity

Field	Type	Req	Default	Description
<code>text</code>	string	yes		
<code>label</code>	string	yes		
<code>positions</code>	array<Position>	yes		

FieldExtensionStrategy

Field	Type	Req	Default	Description
<code>name</code>	string		<code>field_extension</code>	
<code>fields</code>	array<string>		<code>[]</code>	List of field ids to extend the context with. It will try to extend the retrieval context with the specified fields in the matching resources. The field ids have to be in the format {fieldtype}/{field...}
<code>data_augmentation_field_prefixes</code>	array<string>		<code>[]</code>	List of prefixes for data augmentation added fields to extend the context with. For example, if the prefix is 'simpson', all fields that are a result of data augmentation with that prefix will be used...

FieldFile

Field	Type	Req	Default	Description
added	string null			
file	CloudLink null			
language	string null			
password	string null			
external	boolean		False	
extract_strategy	string null			Id of the Nuclia extract strategy used at processing time. If not set, the default strategy was used. Extract strategies are defined at the learning configuration api.
split_strategy	string null			Id of the Nuclia split strategy used at processing time. If not set, the default strategy was used. Split strategies are defined at the learning configuration api.

FieldID

Field	Type	Req	Default	Description
field_type	FieldType	yes		
field	string	yes		

FieldLargeMetadata

Field	Type	Req	Default	Description
entities	array<Entity-Output> null			

tokens	map<string, integer> null			
--------	-----------------------------	--	--	--

FieldLink

Field	Type	Req	Default	Description
added	string null			
headers	map<string, string> null			
cookies	map<string, string> null			
uri	string null			
language	string null			
localStorage	map<string, string> null			
css_selector	string null			
xpath	string null			
extract_strategy	string null			Id of the Nuclia extract strategy used at processing time. If not set, the default strategy was used. Extract strategies are defined at the learning configuration api.
split_strategy	string null			Id of the Nuclia split strategy used at processing time. If not set, the default strategy was used. Split strategies are defined at the learning configuration api.

FieldMetadata

Field	Type	Req	Default	Description
-------	------	-----	---------	-------------

links	array<string>	yes		
paragraphs	array<nucliadb_model_common_Paragraph>	yes		
ner	map<string, string>	yes		
entities	map<string, FieldEntities>	yes		
classifications	array<Classification>	yes		
last_index	string null			
last_understanding	string null			
last_extract	string null			
last_summary	string null			
last_processing_start	string null			
thumbnail	CloudLink null			
language	string null			
summary	string null			
positions	map<string, Positions>	yes		
relations	array<Relation-Output> null			
mime_type	string null			

FieldMimetype

Matches fields with a mimetype

Field	Type	Req	Default	Description
prop	string		field_mimetype	
type	string	yes		Type of the mimetype to match. e.g: In image/jpeg, type is image

subtype	string null			Type of the mimetype to match. e.g: In image/jpeg, subtype is jpeg. Leave blank to match all mimetype of the type
---------	---------------	--	--	--

FieldQuestionAnswers

Field	Type	Req	Default	Description
question_answers	QuestionAnswers	yes		
split_question_answers	map<string, QuestionAnswers> null			
deleted_splits	array<string> null			

FieldRef

Field	Type	Req	Default	Description
field_type	FieldTypeName	yes		This map assumes that both values and extracted data field containers use the same names for its fields. See models.ResourceFieldValues and models.ResourceFieldExtractedData
field_id	string	yes		
split	string null			

FieldText

Field	Type	Req	Default	Description
body	string null			

<code>format</code>	<code>TextFormat null</code>			
<code>md5</code>	<code>string null</code>			
<code>extract_strategy</code>	<code>string null</code>			Id of the Nuclia extract strategy used at processing time. If not set, the default strategy was used. Extract strategies are defined at the learning configuration api.
<code>split_strategy</code>	<code>string null</code>			Id of the Nuclia split strategy used at processing time. If not set, the default strategy was used. Split strategies are defined at the learning configuration api.

FieldType

Enum: `file`, `link`, `text`, `generic`, `conversation`

FieldTypeName

This map assumes that both values and extracted data field containers use the same names for its fields. See `models.ResourceFieldValues` and `models.ResourceFieldExtractedData`

Enum: `text`, `file`, `link`, `conversation`, `generic`, `key_value`

File

Field	Type	Req	Default	Description
<code>filename</code>	<code>string null</code>			
<code>content_type</code>	<code>string</code>		<code>application/octet-stream</code>	
<code>payload</code>	<code>string null</code>			Base64 encoded file content
<code>md5</code>	<code>string null</code>			
<code>uri</code>	<code>string null</code>			

extra_headers	map<string, string>		{}	
---------------	---------------------	--	----	--

FileB64

Field	Type	Req	Default	Description
filename	string	yes		
content_type	string		application/ octet-stream	
payload	string	yes		
md5	string	yes		

FileExtractedData

Field	Type	Req	Default	Description
language	string null			
md5	string null			
metadata	map<string, string> null			
nested	map<string, string> null			
file_generated	map<string, CloudLink> null			
file_rows_previews	map<string, RowsPreview> null			
file_preview	CloudLink null			
file_pages_previews	FilePages null			
file_thumbnail	CloudLink null			
field	string null			
icon	string null			
nested_position	map<string, NestedPosition> null			

nested_list_position	map<string, NestedListPosition> null			
----------------------	--	--	--	--

FileField

Field	Type	Req	Default	Description
language	string null			
password	string null			
file	File	yes		
extract_strategy	string null			Id of the Nuclia extract strategy to use at processing time. If not set, the default strategy will be used. Extract strategies are defined at the learning configuration api.
split_strategy	string null			Id of the Nuclia split strategy used at processing time. If not set, the default strategy was used. Split strategies are defined at the learning configuration api.

FileFieldData

Field	Type	Req	Default	Description
value	FieldFile null			
extracted	FileFieldExtractedData null			
error	Error null			
status	string null			
errors	array<Error> null			

FileFieldExtractedData

Field	Type	Req	Default	Description
text	ExtractedText null			
metadata	FieldComputedMetadata null			
large_metadata	LargeComputedMetadata null			
vectors	VectorObject null			
question_answers	FieldQuestionAnswers null			
relation_node_vectors	map<string, array<RelationNodeVector>> null			
relation_edge_vectors	map<string, array<RelationEdgeVector>> null			
file	FileExtractedData null			

FilePages

Field	Type	Req	Default	Description
pages	array<CloudLink> null			
positions	array<PagePositions> null			
structures	array<PageStructure> null			

Filter

Field	Type	Req	Default	Description
all	array<string> null			
any	array<string> null			
none	array<string> null			

<code>not_all</code>	<code>array<string> null</code>			
----------------------	---	--	--	--

FilterExpression

Returns only documents that match this filter expression. Filtering examples can be found here: (see docs)

This allows building complex filtering expressions and replaces the following parameters: `fields`, `filters`, `range_*`, `resource_filters`, `keyword_filters`.

Field	Type	Req	Default	Description
<code>field</code>	And_FieldFilterExpressionType_ Or_FieldFilterExpressionType_ Not_FieldFilterExpressionType_ Resource-Input Field Keyword DateCreated DateModified Label ResourceMimetype FieldMimetype Entity-Input Language OriginTag OriginMetadata OriginPath OriginSource OriginCollaborator nucliadb_models__filters__Generated null			Filter to apply to fields
<code>paragraph</code>	And_ParagraphFilterExpression_ Or_ParagraphFilterExpression_ Not_ParagraphFilterExpression_ Label Kind null			Filter to apply to each text block
<code>key_value</code>	And_KVFilterExpression_ Or_KVFilterExpression_ Not_KVFilterExpression_ Eq Inequalities Contains null			Filter to apply to key-value fields (JSON index prefilter)

<code>operator</code>	Operator		<code>and</code>	How to combine field and paragraph filters (default is AND).AND returns text blocks that match both filters.OR returns textblocks that match one of the two filters
-----------------------	----------	--	------------------	---

FindConfig

Field	Type	Req	Default	Description
<code>audit_metadata</code>	map<string, string> null			A dictionary containing optional audit-specific metadata, such as userid, environment, or other contextual information. This metadata can be leveraged for filtering and analyzing activity logs in futu...
<code>query</code>	string			The query to search for
<code>filter_expression</code>	FilterExpression null			Returns only documents that match this filter expression.Filtering examples can be found here: (see docs) This allows building complex filtering expressions and replaces the following parameters:field...
<code>fields</code>	array<string>		<code>[]</code>	The list of fields to search in. For instance: a/title to search only on title field. For more details on filtering by

				field, see: (see docs)
<code>filters</code>	<code>array<string> array<Filter></code>		<code>[]</code>	The list of filters to apply. Filtering examples can be found here: (see docs)
<code>top_k</code>	<code>integer</code>		<code>20</code>	The number of results search should return. The maximum number of results allowed is 200.
<code>min_score</code>	<code>number MinScore null</code>			Minimum score to filter search results. Results with a lower score will be ignored. Accepts either a float or a dictionary with the minimum scores for the bm25 and vector indexes. If a float is provid...
<code>range_creation_start</code>	<code>string null</code>			Resources created before this date will be filtered out of search results. Datetime are represented as a str in ISO 8601 format, like: 2008-09-15T15:53:00+05:00.
<code>range_creation_end</code>	<code>string null</code>			Resources created after this date will be filtered out of search results. Datetime are represented as a str in ISO 8601 format, like: 2008-09-15T15:53:00+05:00.
<code>range_modification_start</code>	<code>string null</code>			Resources modified before this date will be filtered out of search results. Datetime are represented as a str in ISO 8601 format,

				like: 2008-09-15T15:53:00+05:00.
<code>range_modification_end</code>	string null			Resources modified after this date will be filtered out of search results. Datetime are represented as a str in ISO 8601 format, like: 2008-09-15T15:53:00+05:00.
<code>debug</code>	boolean		<code>False</code>	If set, the response will include some extra metadata for debugging purposes, like the list of queried nodes.
<code>highlight</code>	boolean		<code>False</code>	If set to true, the query terms will be highlighted in the results between <code><mark>...</mark></code> tags
<code>show</code>	array<ResourceProperties>		<code>['basic']</code>	Controls which types of metadata are serialized on resources of search results
<code>field_type_filter</code>	array<FieldTypeName>		<code>['text', 'file', 'link', 'conversation', 'generic', 'key_value']</code>	Define which field types are serialized on resources of search results
<code>extracted</code>	array<ExtractedData TypeName>		<code>[]</code>	[Deprecated] Please use GET resource endpoint instead to get extracted metadata
<code>vector</code>	array<number> null			The vector to perform the search with. If not provided, NucliaDB will use Nuclia Predict API to create the vector off from

				the query.
<code>vectorset</code>	<code>string null</code>			Vectors index to perform the search in. If not provided, NucliaDB will use the default one
<code>with_duplicates</code>	<code>boolean</code>		<code>False</code>	Whether to return duplicate paragraphs on the same document
<code>with_synonyms</code>	<code>boolean</code>		<code>False</code>	Whether to return matches for custom knowledge box synonyms of the query terms. Note: only supported for keyword and fulltext search options.
<code>resource_filters</code>	<code>array<string></code>		<code>[]</code>	List of resource ids to filter search results for. Only paragraphs from the specified resources will be returned.
<code>security</code>	<code>RequestSecurity null</code>			Security metadata for the request. Please refer to the documentation for more details on how security works: (see docs)
<code>show_hidden</code>	<code>boolean</code>		<code>False</code>	If set to false (default), excludes hidden resources from search
<code>rephrase</code>	<code>boolean</code>		<code>False</code>	Rephrase the query for a more efficient retrieval. This will consume LLM tokens and make the request slower.
<code>rephrase_prompt</code>	<code>string null</code>			Rephrase prompt given to the generative model

				responsible for rephrasing the query for a more effective retrieval step. This is only used if the rephrase flag is set to true in the request. If not spe...
query_image	Image null			Image that will be used together with the query text for retrieval.
graph_query	And_GraphPathQuery_ Or_GraphPathQuery_ Not_GraphPathQuery_ GraphPath SourceNode DestinationNode AnyNode nucliadb_models__graph__requests__Relation nucliadb_models__graph__requests__Generated null			Query for the knowledge graph. Paths (node-relation-node) extracted from a paragraphid will be used to extend the results
features	array<FindOptions>		['keyword', 'semantic']	List of search features to use. Each value corresponds to a lookup into one of the different indexes
rank_fusion	RankFusionName ReciprocalRankFusion		rrf	Rank fusion algorithm to use to merge results from multiple retrievers (keyword, semantic)
reranker	RerankerName PredictReranker		predict	Reranker let you specify which method you want to use to rerank your results at the end of retrieval
keyword_filters	array<string> array<Filter>		[]	List of keyword filter expressions to apply to the retrieval step.

				The text block search will only be performed on the documents that contain the specified keywords. The filters are case-insensitive, ...
<code>generative_model</code>	string null			The generative model used to rephrase the query. If not provided, the model configured for the Knowledge Box is used.

FindField

Field	Type	Req	Default	Description
<code>paragraphs</code>	map<string, FindParagraph>	yes		

FindOptions

Enum: `keyword`, `semantic`, `relations`, `graph`

FindParagraph

Field	Type	Req	Default	Description
<code>score</code>	number	yes		
<code>score_type</code>	SCORE_TYPE	yes		
<code>order</code>	integer		<code>0</code>	
<code>text</code>	string	yes		
<code>id</code>	string	yes		
<code>labels</code>	array<string> null		<code>[]</code>	
<code>position</code>	TextPosition null			
<code>fuzzy_result</code>	boolean		<code>False</code>	
<code>page_with_visual</code>	boolean		<code>False</code>	This flag informs if the page may have

				information that has not been extracted
reference	string null			Reference to the extracted image that represents this paragraph
is_a_table	boolean		False	The referenced image of the paragraph is a table
relevant_relations	Relations null			Relevant relations from which the paragraph was found, will only be filled if using the Graph RAG Strategy

FindRequest

Field	Type	Req	Default	Description
audit_metadata	map<string, string> null			A dictionary containing optional audit-specific metadata, such as userid, environment, or other contextual information. This metadata can be leveraged for filtering and analyzing activity logs in futu...
query	string			The query to search for
filter_expression	FilterExpression null			Returns only documents that match this filter expression. Filtering examples can be found here: (see docs) This allows building complex filtering expressions and replaces the following

				parameters:field...
fields	array<string>		[]	The list of fields to search in. For instance: a/title to search only on title field. For more details on filtering by field, see: (see docs)
filters	array<string> array<Filter>		[]	The list of filters to apply. Filtering examples can be found here: (see docs)
top_k	integer		20	The number of results search should return. The maximum number of results allowed is 200.
min_score	number MinScore null			Minimum score to filter search results. Results with a lower score will be ignored. Accepts either a float or a dictionary with the minimum scores for the bm25 and vector indexes. If a float is provid...
range_creation_start	string null			Resources created before this date will be filtered out of search results. Datetime are represented as a str in ISO 8601 format, like: 2008-09-15T15:53:00+05:00.
range_creation_end	string null			Resources created after this date will be filtered out of search results. Datetime are represented as a str in ISO 8601 format, like: 2008-09-15T15:53:00+05:00.

<code>range_modification_start</code>	string null			Resources modified before this date will be filtered out of search results. Datetime are represented as a str in ISO 8601 format, like: 2008-09-15T15:53:00+05:00.
<code>range_modification_end</code>	string null			Resources modified after this date will be filtered out of search results. Datetime are represented as a str in ISO 8601 format, like: 2008-09-15T15:53:00+05:00.
<code>debug</code>	boolean		<code>False</code>	If set, the response will include some extra metadata for debugging purposes, like the list of queried nodes.
<code>highlight</code>	boolean		<code>False</code>	If set to true, the query terms will be highlighted in the results between <code><mark>...</mark></code> tags
<code>show</code>	array<ResourceProperties>		<code>['basic']</code>	Controls which types of metadata are serialized on resources of search results
<code>field_type_filter</code>	array<FieldTypeName>		<code>['text', 'file', 'link', 'conversation', 'generic', 'key_value']</code>	Define which field types are serialized on resources of search results
<code>extracted</code>	array<ExtractedData TypeName>		<code>[]</code>	[Deprecated] Please use GET resource endpoint instead to get extracted

				metadata
vector	array<number> null			The vector to perform the search with. If not provided, NucliaDB will use Nuclia Predict API to create the vector off from the query.
vectorset	string null			Vectors index to perform the search in. If not provided, NucliaDB will use the default one
with_duplicates	boolean		False	Whether to return duplicate paragraphs on the same document
with_synonyms	boolean		False	Whether to return matches for custom knowledge box synonyms of the query terms. Note: only supported for keyword and fulltext search options.
resource_filters	array<string>		[]	List of resource ids to filter search results for. Only paragraphs from the specified resources will be returned.
security	RequestSecurity null			Security metadata for the request. Please refer to the documentation for more details on how security works: (see docs)
show_hidden	boolean		False	If set to false (default), excludes hidden resources from search
rephrase	boolean		False	Rephrase the query for a more efficient

				retrieval. This will consume LLM tokens and make the request slower.
rephrase_prompt	string null			Rephrase prompt given to the generative model responsible for rephrasing the query for a more effective retrieval step. This is only used if the rephrase flag is set to true in the request. If not spe...
query_image	Image null			Image that will be used together with the query text for retrieval.
graph_query	And_GraphPathQuery_ Or_GraphPathQuery_ Not_GraphPathQuery_ GraphPath-Input SourceNode DestinationNode AnyNode Relation-Input nucliadb_models__graph__requests__Generated null			Query for the knowledge graph. Paths (node-relation-node) extracted from a paragraphid will be used to extend the results
features	array<FindOptions>		['keyword', 'semantic']	List of search features to use. Each value corresponds to a lookup into on of the different indexes
rank_fusion	RankFusionName ReciprocalRankFusion		rrf	Rank fusion algorithm to use to merge results from multiple retrievers (keyword, semantic)
reranker	RerankerName PredictReranker		predict	Reranker let you specify which method you want to use to

				rerank your results at the end of retrieval
<code>keyword_filters</code>	<code>array<string> array<Filter></code>		<code>[]</code>	List of keyword filter expressions to apply to the retrieval step. The text block search will only be performed on the documents that contain the specified keywords. The filters are case-insensitive, ...
<code>search_configuration</code>	<code>string null</code>			Load find parameters from this configuration. Parameters in the request override parameters from the configuration.
<code>generative_model</code>	<code>string null</code>			The generative model used to rephrase the query. If not provided, the model configured for the Knowledge Box is used.

FindResource

Field	Type	Req	Default	Description
<code>id</code>	<code>string</code>	yes		
<code>slug</code>	<code>string null</code>			
<code>title</code>	<code>string null</code>			
<code>summary</code>	<code>string null</code>			
<code>icon</code>	<code>string null</code>			
<code>thumbnail</code>	<code>string null</code>			
<code>metadata</code>	<code>Metadata null</code>			
<code>usermetadata</code>	<code>UserMetadata null</code>			

fieldmetadata	array<UserFieldMetadata> null			
computedmetadata	ComputedMetadata null			
created	string null			
modified	string null			
last_seqid	integer null			
last_account_seq	integer null			
queue	QueueType null			
hidden	boolean null			
origin	Origin null			
extra	Extra null			
relations	array<Relation-Output> null			
data	ResourceData null			
security	ResourceSecurity null			Resource security metadata
fields	map<string, FindField>	yes		

FindSearchConfiguration

Field	Type	Req	Default	Description
kind	string	yes		
config	FindConfig	yes		

FullResourceApplyTo

Field	Type	Req	Default	Description
exclude	array<string>			Resources from matches containing any of these labels won't expand to the full resource. This

				may be useful to exclude long and not interesting resources and expend less tokens
--	--	--	--	--

FullResourceStrategy

Field	Type	Req	Default	Description
<code>name</code>	string		<code>full_resource</code>	
<code>count</code>	integer null			Maximum number of full documents to retrieve. If not specified, all matching documents are retrieved.
<code>include_remaining_text_blocks</code>	boolean		<code>False</code>	Whether to include the remaining text blocks after the maximum number of resources has been reached.
<code>apply_to</code>	FullResourceApplyTo null			Define which resources to exclude from serialization

GenerativeModel

An enumeration.

Enum: `generative-multilingual-2023`, `chatgpt`

Generator

Enum: `data-augmentation`, `processor`, `user`

GenericFieldData

Field	Type	Req	Default	Description
<code>value</code>	string null			
<code>extracted</code>	TextFieldExtractedData null			

<code>error</code>	Error null			
<code>status</code>	string null			
<code>errors</code>	array<Error> null			

GraphFilterExpression

Returns only relations from documents that match this filter expression. Filtering examples can be found here: (see docs)

Field	Type	Req	Default	Description
<code>field</code>	And_FieldFilterExpressionType_ Or_FieldFilterExpressionType_ Not_FieldFilterExpressionType_ Resource-Input Field Keyword DateCreated DateModified Label ResourceMimetype FieldMimetype Entity-Input Language OriginTag OriginMetadata OriginPath OriginSource OriginCollaborator nucliadb_models__filters__Generated	yes		Filter to apply to fields

GraphNode

Field	Type	Req	Default	Description
<code>value</code>	string null			
<code>match</code>	NodeMatchKindName		<code>exact</code>	
<code>type</code>	RelationNodeType null		<code>entity</code>	
<code>group</code>	string null			

GraphNode-Input

Field	Type	Req	Default	Description
value	string null			
match	NodeMatchKindName		exact	
type	RelationNodeType null		entity	
group	string null			

GraphNode-Output

Field	Type	Req	Default	Description
value	string	yes		
type	RelationNodeType	yes		
group	string	yes		
score	number null			

GraphNodesSearchRequest

Field	Type	Req	Default	Description
top_k	integer		50	
filter_expression	GraphFilterExpression null			Returns only relations from documents that match this filter expression. Filtering examples can be found here: (see docs)
security	RequestSecurity null			Security metadata for the request. Please refer to the documentation for more details on how security works: (see docs)

<code>show_hidden</code>	boolean		False	If set to false (default), excludes hidden resources from search
<code>query</code>	And_GraphNodesQuery_ Or_GraphNodesQuery_ Not_GraphNodesQuery_ AnyNode nucliadb_models__graph__requests__Generated	yes		

GraphNodesSearchResponse

Field	Type	Req	Default	Description
<code>nodes</code>	array<GraphNode-Output>	yes		

GraphPath

Field	Type	Req	Default	Description
<code>prop</code>	string		<code>path</code>	
<code>source</code>	GraphNode null			
<code>relation</code>	GraphRelation null			
<code>destination</code>	GraphNode null			
<code>undirected</code>	boolean		False	

GraphPath-Input

Field	Type	Req	Default	Description
<code>prop</code>	string		<code>path</code>	
<code>source</code>	GraphNode-Input null			
<code>relation</code>	GraphRelation-Input null			

<code>destination</code>	GraphNode-Input null			
<code>undirected</code>	boolean		<code>False</code>	

GraphPath-Output

Field	Type	Req	Default	Description
<code>source</code>	GraphNode-Output	yes		
<code>relation</code>	GraphRelation-Output	yes		
<code>destination</code>	GraphNode-Output	yes		
<code>metadata</code>	PathMetadata null	yes		
<code>score</code>	number null			

GraphRelation

Field	Type	Req	Default	Description
<code>label</code>	string null			
<code>type</code>	RelationType null			
<code>match</code>	RelationMatchKindName		<code>exact</code>	

GraphRelation-Input

Field	Type	Req	Default	Description
<code>label</code>	string null			
<code>type</code>	RelationType null			
<code>match</code>	RelationMatchKindName		<code>exact</code>	

GraphRelation-Output

Field	Type	Req	Default	Description
-------	------	-----	---------	-------------

label	string	yes		
type	RelationType	yes		
score	number null			

GraphRelationsSearchRequest

Field	Type	Req	Default	Description
top_k	integer		50	
filter_expression	GraphFilterExpression null			Returns only relations from documents that match this filter expression. Filtering examples can be found here: (see docs)
security	RequestSecurity null			Security metadata for the request. Please refer to the documentation for more details on how security works: (see docs)
show_hidden	boolean		False	If set to false (default), excludes hidden resources from search
query	And_GraphRelationsQuery_ Or_GraphRelationsQuery_ Not_GraphRelationsQuery_ Relation-Input nucliadb_models__graph__requests__Generated	yes		

GraphRelationsSearchResponse

Field	Type	Req	Default	Description
-------	------	-----	---------	-------------

<code>relations</code>	array<GraphRelation-Output>	yes		
------------------------	-----------------------------	-----	--	--

GraphSearchRequest

Field	Type	Req	Default	Description
<code>top_k</code>	integer		<code>50</code>	
<code>filter_expression</code>	GraphFilterExpression null			Returns only relations from documents that match this filter expression. Filtering examples can be found here: (see docs)
<code>security</code>	RequestSecurity null			Security metadata for the request. Please refer to the documentation for more details on how security works: (see docs)
<code>show_hidden</code>	boolean		<code>False</code>	If set to false (default), excludes hidden resources from search
<code>query</code>	And_GraphPathQuery_ Or_GraphPathQuery_ Not_GraphPathQuery_ GraphPath-Input SourceNode DestinationNode AnyNode Relation-Input nucliadb_models__graph__requests__Generated	yes		

GraphSearchResponse

Field	Type	Req	Default	Description
-------	------	-----	---------	-------------

<code>paths</code>	array<GraphPath-Output>	yes		
--------------------	-------------------------	-----	--	--

GraphStrategy

This strategy retrieves context pieces by exploring the Knowledge Graph, starting from the entities present in the query. It works best if the Knowledge Box has a user-defined Graph Extraction agent enabled.

Field	Type	Req	Default	Description
<code>name</code>	string		<code>graph_beta</code>	
<code>hops</code>	integer		<code>3</code>	Number of hops to take when exploring the graph for relevant context. For example, - hops=1 will explore the neighbors of the starting entities. - hops=2 will explore the neighbors of the neighbors of...
<code>top_k</code>	integer			Number of relationships to keep after each hop after ranking them by relevance to the query. This number correlates to more paragraphs being sent as context. If not set, this number will be set to 30 ...
<code>agentic_graph_only</code>	boolean		<code>True</code>	If set to true, only relationships extracted from a graph extraction agent are considered for context expansion.
<code>relation_text_as_paragraphs</code>	boolean		<code>False</code>	If set to true, the text of the relationships is to create context paragraphs, this enables to use bigger top K values without

				running into the generative model's context limits. If set to false, the ...
relation_ranking	RelationRanking		reranker	Method to rank relationships. - reranker uses the reranker model to rank relationships. - generative uses first the reranker to first lower the amount of relationships and then the generative model to...
query_entity_detection	QueryEntityDetection		predict	Method to detect entities in the query. - predict uses NUA to detect entities in the query, slower and more accurate but requires an exact text match between Knowledge Box entities and entities in the...
weight	number		3.0	Weight of the graph strategy in the context. The weight is used to scale the results of the strategy before adding them to the context. The weight should be a positive number.

HFEmbeddingKey

Some models require a specific template (including prefix) to work correctly in each task For example Snowflake's Arctic-embed requires a specific prefix to work correctly. In that case, the query prompt will be ` passage_prompt: "" query_prompt: "Represent this sentence for searching relevant pas...

Field	Type	Req	Default	Description
url	string			

key	string			
matryoshka	array<integer>			
similarity	string			
size	integer		0	
threshold	number		0.0	
passage_prompt	string			
query_prompt	string			

HFLLMKey

Field	Type	Req	Default	Description
key	string			
url	string			
model	ModelType		0	

HierarchyResourceStrategy

Field	Type	Req	Default	Description
name	string		hierarchy	
count	integer		0	Number of extra characters that are added to each matching paragraph when adding to the context.

HTTPValidationError

Field	Type	Req	Default	Description
detail	array<ValidationError>			

Image

Field	Type	Req	Default	Description
<code>content_type</code>	string	yes		
<code>b64encoded</code>	string	yes		

Inequalities

Inequality operators that can be grouped to perform range queries.

Field	Type	Req	Default	Description
<code>schema_id</code>	string	yes		
<code>key</code>	string	yes		
<code>gte</code>	integer number string null			
<code>lte</code>	integer number string null			

InputConversationField

Field	Type	Req	Default	Description
<code>messages</code>	array<InputMessage>			List of messages in the conversation field. Each message must have a unique ident. A single conversation can contain up to 51,200 messages. You can add up to 2,048 messages per request.
<code>extract_strategy</code>	string null			Id of the Nuclia extract strategy used at processing time. If not set, the default strategy was used. Extract strategies are defined at the learning configuration api.

<code>split_strategy</code>	string null			Id of the Nuclia split strategy used at processing time. If not set, the default strategy was used. Split strategies are defined at the learning configuration api.
-----------------------------	---------------	--	--	---

InputMessage

Field	Type	Req	Default	Description
<code>timestamp</code>	string null			Time at which the message was sent, in ISO 8601 format.
<code>who</code>	string null			Sender of the message, e.g. 'user' or 'assistant'
<code>to</code>	array<string>			List of recipients of the message, e.g. ['assistant'] or ['user']
<code>content</code>	InputMessageContent	yes		
<code>ident</code>	string	yes		Unique identifier for the message. Must be unique within the conversation.
<code>type</code>	MessageType null			

InputMessageContent

Field	Type	Req	Default	Description
<code>text</code>	string	yes		
<code>format</code>	MessageFormat		PLAIN	
<code>attachments</code>	array<FileB64>		[]	
<code>attachments_fields</code>	array<FieldRef>		[]	

InputMetadata

Field	Type	Req	Default	Description
metadata	map<string, string>		{}	
language	string null			
languages	array<string> null			

InputOrigin

Field	Type	Req	Default	Description
source_id	string null			
url	string null			
created	string null			Creation date of the resource at the origin system. This can be later used for date range filtering on search endpoints. Have a look at the advanced search documentation page: (see docs)
modified	string null			Modification date of the resource at the origin system. This can be later used for date range filtering on search endpoints. Have a look at the advanced search documentation page: (see docs)
metadata	map<string, string>		{}	Generic metadata from the resource at the origin system. It can later be used for filtering on search endpoints with '/origin.metadata/{key}/{value}'
tags	array<string>		[]	Resource tags about

				the origin system. It can later be used for filtering on search endpoints with <code>'/origin.tags/{tag}'</code>
<code>collaborators</code>	<code>array<string></code>		<code>[]</code>	
<code>filename</code>	<code>string null</code>			
<code>related</code>	<code>array<string></code>		<code>[]</code>	
<code>path</code>	<code>string null</code>			Path of the original resource. Typically used to store folder structure information of the resource at the origin system. It can be later used for filtering on search endpoints with <code>'/origin.path/{pat...'</code>
<code>sync_metadata</code>	<code>SyncMetadata null</code>			Metadata related to the resource from the origin system fetched by the Progress Agentic RAG's Cloud Storage Sync service.

KBKVSchemas

Field	Type	Req	Default	Description
<code>schemas</code>	<code>map<string, KVSchema></code>			

KeyValueField

A key-value field value. The field id (key in the resource's `key_values` dict) must equal the schema name — enforcing one KV field per schema per resource.

Field	Type	Req	Default	Description
<code>data</code>	<code>map<string, string array<string> integer number boolean Range></code>			Key-value pairs conforming to the schema.

KeyValueFieldData

Field	Type	Req	Default	Description
value	KeyValueField null			
error	Error null			
status	string null			
errors	array<Error> null			

Keyword

Matches all fields that contain a keyword

Field	Type	Req	Default	Description
prop	string		keyword	
word	string	yes		Keyword to find

Kind

Matches paragraphs of a certain kind

Field	Type	Req	Default	Description
prop	string		kind	
kind	TypeParagraph	yes		The kind of paragraph to match

KnowledgeBoxConfig

Field	Type	Req	Default	Description
uuid	string null			UUID for the Knowledge Box. If not provided, a new UUID will be generated.
slug	string null			Slug for the Knowledge Box.

<code>title</code>	string null			Title for the Knowledge Box.
<code>description</code>	string null			Description for the Knowledge Box.
<code>learning_configuration</code>	object (free-form map) null			Learning configuration for the Knowledge Box. If provided, NucliaDB will set the learning configuration for the Knowledge Box.
<code>external_index_provider</code>	DummyIndexProvider null			External index provider for the Knowledge Box.
<code>configured_external_index_provider</code>	object (free-form map) null			Metadata for the configured external index provider (if any)
<code>similarity</code>	VectorSimilarity null			This field is deprecated. Use 'learningconfiguration' instead.
<code>hidden_resources_enabled</code>	boolean		<code>False</code>	Allow hiding resources
<code>hidden_resources_hide_on_creation</code>	boolean		<code>False</code>	Hide newly created resources
<code>enforce_security</code>	boolean null			Whether to enforce security groups by default on all requests. If set to None, the value will not be updated. If set to a boolean value, it will update the setting for the Knowledge Box accordingly.

KnowledgeboxCounters

Field	Type	Req	Default	Description
-------	------	-----	---------	-------------

resources	integer	yes		
paragraphs	integer	yes		
fields	integer	yes		
sentences	integer	yes		
shards	array<string> null			
index_size	number		0.0	

KnowledgeBoxEntities

Field	Type	Req	Default	Description
uuid	string	yes		
groups	map<string, EntitiesGroupSummary>		{}	

KnowledgeboxFindResults

Find on knowledgebox results

Field	Type	Req	Default	Description
resources	map<string, FindResource>	yes		
relations	Relations null			
query	string null			
rephrased_query	string null			
total	integer		0	
page_number	integer		0	Pagination will be deprecated, please, refer to topk in the request
page_size	integer		20	Pagination will be deprecated, please, refer to topk in the request

<code>next_page</code>	boolean		<code>False</code>	Pagination will be deprecated, please, refer to <code>topk</code> in the request
<code>nodes</code>	array<map<string, string>> null			List of nodes queried in the search
<code>shards</code>	array<string> null			The list of shard replica ids used for the search.
<code>autofilters</code>	array<string>		<code>[]</code>	[deprecated] list of filters automatically applied to the search query
<code>min_score</code>	number MinScore null		<code>{'bm25': 0.0}</code>	The minimum scores that have been used for the search operation.
<code>best_matches</code>	array<string>		<code>[]</code>	List of ids of best matching paragraphs. The list is sorted by decreasing relevance (most relevant first).
<code>metrics</code>	object (free-form map) null			Metrics information about the search operation. The metadata included in this field is subject to change and should not be used in production. This is only available if the debug parameter is set to t...

KnowledgeBoxLabels

Field	Type	Req	Default	Description
<code>uuid</code>	string	yes		
<code>labelsets</code>	map<string, LabelSet>		<code>{}</code>	

KnowledgeBoxObj

The API representation of a Knowledge Box object.

Field	Type	Req	Default	Description
slug	string null			
uuid	string	yes		
config	KnowledgeBoxConfig null			
model	SemanticModelMetadata null			

KnowledgeBoxObjID

Field	Type	Req	Default	Description
uuid	string	yes		

KnowledgeboxSearchResults

Search on knowledgebox results

Field	Type	Req	Default	Description
resources	map<string, Resource-Output>		{}	
sentences	Sentences null			
paragraphs	Paragraphs null			
fulltext	Resources null			
relations	Relations null			
nodes	array<map<string, string>> null			
shards	array<string> null			
autofilters	array<string>		[]	[deprecated] list of filters automatically applied to the search query

KnowledgeboxSuggestResults

Suggest on resource results

Field	Type	Req	Default	Description
paragraphs	Paragraphs null			
entities	RelatedEntities null			
shards	array<string> null			

KnowledgeBoxSynonyms

Field	Type	Req	Default	Description
synonyms	map<string, array<string>>	yes		

KnowledgeGraphEntity

Field	Type	Req	Default	Description
name	string	yes		
type	RelationNodeType null			
subtype	string null			

KVFieldType**Enum:** text, integer, float, boolean, date**KVSchema**

Field	Type	Req	Default	Description
id	string	yes		
description	string			
fields	array<KVSchemaField>			

KVSchemaField

Field	Type	Req	Default	Description
key	string	yes		
type	KVFieldType	yes		
description	string			
required	boolean		True	
range	boolean		False	When enabled, this field stores a range instead of a single value and operations like contains are possible
repeated	boolean		False	When enabled, this field stores a list of elements instead of a single value and operations like contains are possible

Label

Matches fields/paragraphs with a label (or labelset)

Field	Type	Req	Default	Description
prop	string		label	
labelset	string	yes		The ID of the labelset to match
label	string null			The label to match. If blank, matches all labels in the given labelset

LabelSet

Field	Type	Req	Default	Description
title	string null			Title of the labelset. It is a prettier display name for the labelset

				shown in the UI but it is not intended to be used for searching.
color	string null		blue	
multiple	boolean		True	
kind	array<LabelSetKind>		[]	
labels	array<nucliadb_model_labels__Label>			List of labels in the labelset. The titles of the labels must be unique within the labelset.

LabelSetKind

Enum: RESOURCES, PARAGRAPHS, SENTENCES, SELECTIONS

Language

Matches the language of the field

Field	Type	Req	Default	Description
prop	string		language	
only_primary	boolean		False	Match only the primary language of the document. By default, matches any language that appears in the document
language	string	yes		The code of the language to match, e.g: en

LargeComputedMetadata

Field	Type	Req	Default	Description
metadata	FieldLargeMetadata null			
split_metadata	map<string, FieldLargeMetadata> null			

<code>deleted_splits</code>	<code>array<string> null</code>			
-----------------------------	---	--	--	--

LearningConfigurationUpdate

Field	Type	Req	Default	Description
<code>anonymization_model</code>	object		<code>disabled</code>	
<code>generative_model</code>	object		<code>chatgpt</code>	
<code>ner_model</code>	object		<code>multilingual</code>	

LinkExtractedData

Field	Type	Req	Default	Description
<code>date</code>	<code>string null</code>			
<code>language</code>	<code>string null</code>			
<code>title</code>	<code>string null</code>			
<code>metadata</code>	<code>map<string, string> null</code>			
<code>link_thumbnail</code>	<code>CloudLink null</code>			
<code>link_preview</code>	<code>CloudLink null</code>			
<code>field</code>	<code>string null</code>			
<code>link_image</code>	<code>CloudLink null</code>			
<code>description</code>	<code>string null</code>			
<code>type</code>	<code>string null</code>			
<code>embed</code>	<code>string null</code>			
<code>file_generated</code>	<code>map<string, CloudLink> null</code>			

LinkField

Field	Type	Req	Default	Description
-------	------	-----	---------	-------------

headers	map<string, string> null		{}	
cookies	map<string, string> null		{}	
uri	string	yes		
language	string null			
localstorage	map<string, string> null		{}	
css_selector	string null			
xpath	string null			
extract_strategy	string null			Id of the Nuclia extract strategy to use at processing time. If not set, the default strategy will be used. Extract strategies are defined at the learning configuration api.
split_strategy	string null			Id of the Nuclia split strategy used at processing time. If not set, the default strategy was used. Split strategies are defined at the learning configuration api.

LinkFieldData

Field	Type	Req	Default	Description
value	FieldLink null			
extracted	LinkFieldExtractedData null			
error	Error null			
status	string null			

<code>errors</code>	<code>array<Error> null</code>			
---------------------	--	--	--	--

LinkFieldExtractedData

Field	Type	Req	Default	Description
<code>text</code>	<code>ExtractedText null</code>			
<code>metadata</code>	<code>FieldComputedMetadata null</code>			
<code>large_metadata</code>	<code>LargeComputedMetadata null</code>			
<code>vectors</code>	<code>VectorObject null</code>			
<code>question_answers</code>	<code>FieldQuestionAnswers null</code>			
<code>relation_node_vectors</code>	<code>map<string, array<RelationNodeVector>> null</code>			
<code>relation_edge_vectors</code>	<code>map<string, array<RelationEdgeVector>> null</code>			
<code>link</code>	<code>LinkExtractedData null</code>			

LLMConfig

Field	Type	Req	Default	Description
<code>user_keys</code>	<code>UserLearningKeys null</code>			
<code>generative_model</code>	<code>string</code>			
<code>generative_provider</code>	<code>string</code>			
<code>generative_prompt_id</code>	<code>string</code>			

LLMSplitConfig

Field	Type	Req	Default	Description
-------	------	-----	---------	-------------

<code>rules</code>	array<string>			
<code>llm</code>	LLMConfig null			

ManualSplitConfig

Field	Type	Req	Default	Description
<code>splitter</code>	string			

MaxTokens

Field	Type	Req	Default	Description
<code>context</code>	integer null			Use to limit the amount of tokens used in the LLM context
<code>answer</code>	integer null			Use to limit the amount of tokens used in the LLM answer

MessageFormat

Enum: `PLAIN`, `HTML`, `RST`, `MARKDOWN`, `KEEP_MARKDOWN`, `JSON`

MessageType

Enum: `UNSET`, `QUESTION`, `ANSWER`

Metadata

Field	Type	Req	Default	Description
<code>metadata</code>	map<string, string>		<code>{}</code>	
<code>language</code>	string null			
<code>languages</code>	array<string> null			
<code>status</code>	ResourceProcessingStatus	yes		

MetadataExtensionStrategy

RAG strategy to enrich the context with metadata of the matching paragraphs or its resources. This strategy can be combined with any of the other strategies.

Field	Type	Req	Default	Description
<code>name</code>	string		<code>metadata_extension</code>	
<code>types</code>	array<MetadataExtensionType>	yes		List of resource metadata types to add to the context. - 'origin': origin metadata of the resource. - 'classificationlabels': classification labels of the resource. - 'ner': Named Entity Recognition e...

MetadataExtensionType

Enum: `origin`, `classification_labels`, `ners`, `extra_metadata`

MinScore

Field	Type	Req	Default	Description
<code>semantic</code>	number null			Minimum semantic similarity score used to filter vector index search. If not specified, the default minimum score of the semantic model associated to the Knowledge Box will be used. Check out the docu...
<code>bm25</code>	number		<code>0</code>	Minimum score used to filter bm25 index search. Check out the documentation for more information on how to use this parameter: (see docs)

MistralKey

Field	Type	Req	Default	Description
key	string			

ModelType

Enum: 0, 1

NeighbouringParagraphsStrategy

Field	Type	Req	Default	Description
name	string		neighbouring_paragraphs	
before	integer		2	Number of previous neighbouring paragraphs to add to the context, for each matching paragraph in the retrieval step.
after	integer		2	Number of following neighbouring paragraphs to add to the context, for each matching paragraph in the retrieval step.

NERModel

An enumeration.

Enum: multilingual

NestedListPosition

Field	Type	Req	Default	Description
positions	array<NestedPosition>	yes		

NestedPosition

Field	Type	Req	Default	Description
<code>start</code>	integer null			
<code>end</code>	integer null			
<code>page</code>	integer null			

NewTextField

Field	Type	Req	Default	Description
<code>text_field</code>	FieldText	yes		
<code>destination</code>	string	yes		

NodeMatchKindName

Enum: `exact`, `fuzzy`, `fuzzy_words`, `semantic`

Not_FieldFilterExpressionType_

Field	Type	Req	Default	Description
<code>operand</code>	And_FieldFilterExpressionType_ Or_FieldFilterExpressionType_ Not_FieldFilterExpressionType_ Resource-Input Field Keyword DateCreated DateModified Label ResourceMimetype FieldMimetype Entity-Input Language OriginTag OriginMetadata OriginPath OriginSource OriginCollaborator nucliadb_models__filters__Generated	yes		

Not_GraphNodesQuery_

Field	Type	Req	Default	Description
operand	And_GraphNodesQuery_ Or_GraphNodesQuery_ Not_GraphNodesQuery_ AnyNode nucliadb_models__graph__requests__Generated	yes		

Not_GraphPathQuery_

Field	Type	Req	Default	Description
operand	And_GraphPathQuery_ Or_GraphPathQuery_ Not_GraphPathQuery_ GraphPath-Input SourceNode DestinationNode AnyNode Relation-Input nucliadb_models__graph__requests__Generated	yes		

Not_GraphRelationsQuery_

Field	Type	Req	Default	Description
operand	And_GraphRelationsQuery_ Or_GraphRelationsQuery_ Not_GraphRelationsQuery_ Relation-Input nucliadb_models__graph__requests__Generated	yes		

Not_KVFilterExpression_

Field	Type	Req	Default	Description
operand	And_KVFilterExpression_ Or_KVFilterExpression_ Not_KVFilterExpression_ Eq Inequalities Contains	yes		

Not_ParagraphFilterExpression_

Field	Type	Req	Default	Description
operand	And_ParagraphFilterExpression_ Or_ParagraphFilterExpression_ Not_ParagraphFilterExpression_ Label Kind	yes		

Not_ResourceFilterExpression_

Field	Type	Req	Default	Description
operand	And_ResourceFilterExpression_ Or_ResourceFilterExpression_ Not_ResourceFilterExpression_ Resource-Input DateCreated DateModified Label ResourceMimetype Language OriginTag OriginMetadata OriginPath OriginSource OriginCollaborator	yes		

	Status			
--	--------	--	--	--

nucliadb_models__common__Paragraph

Field	Type	Req	Default	Description
start	integer null			
end	integer null			
start_seconds	array<integer> null			
end_seconds	array<integer> null			
kind	TypeParagraph null			
classifications	array<Classification> null			
sentences	array<nucliadb_models__common__Sentence> null			
key	string null			
page	PageInformation null			
representation	Representation null			
relations	ParagraphRelations null			

nucliadb_models__common__Sentence

Field	Type	Req	Default	Description
start	integer null			
end	integer null			
key	string null			

nucliadb_models__entities__Entity

Field	Type	Req	Default	Description
value	string	yes		

<code>merged</code>	boolean		<code>False</code>	
<code>represents</code>	array<string>		<code>[]</code>	

nucliadb_models__extracted__Entity

Field	Type	Req	Default	Description
<code>token</code>	string null			
<code>root</code>	string null			
<code>type</code>	string null			

nucliadb_models__filters__Entity

Matches fields that contains a detected entity

Field	Type	Req	Default	Description
<code>prop</code>	string		<code>entity</code>	
<code>subtype</code>	string	yes		Type of the entity. e.g: PERSON
<code>value</code>	string null			Value of the entity. e.g: Anna. If blank, matches any entity of the given type

nucliadb_models__filters__Generated

Matches if the field was generated by the given source

Field	Type	Req	Default	Description
<code>prop</code>	string		<code>generated</code>	
<code>by</code>	string	yes		Generator for this field. Currently, only data-augmentation is supported
<code>da_task</code>	string null			Matches field generated by an specific DA task, given its prefix

nucliadb_models__filters__Label

Matches fields/paragraphs with a label (or labelset)

Field	Type	Req	Default	Description
prop	string		label	
labelset	string	yes		The ID of the labelset to match
label	string null			The label to match. If blank, matches all labels in the given labelset

nucliadb_models__filters__Resource

Matches all fields of a resource given its id or slug

Field	Type	Req	Default	Description
prop	string		resource	
id	string null			UUID of the resource to match
slug	string null			Slug of the resource to match

nucliadb_models__graph__requests__Generated

Matches if the relation was generated by the given source

Field	Type	Req	Default	Description
prop	string		generated	
by	Generator	yes		Generator for this field.
da_task	string null			Matches relations generated by an specific DA task, given its prefix

nucliadb_models__graph__requests__Relation

Field	Type	Req	Default	Description
prop	string		relation	
label	string null			
type	RelationType null			
match	RelationMatchKindName		exact	

nucliadb_models__labels__Label

Field	Type	Req	Default	Description
title	string	yes		Title of the label. This is the display name for the label shown in the UI and also used for searching.
related	string null			
text	string null			
uri	string null			

nucliadb_models__metadata__Relation

Field	Type	Req	Default	Description
relation	RelationType	yes		
label	string null			
metadata	RelationMetadata null			
from	RelationEntity null			
to	RelationEntity	yes		

nucliadb_models__resource__Resource

Field	Type	Req	Default	Description
-------	------	-----	---------	-------------

id	string	yes		
slug	string null			
title	string null			
summary	string null			
icon	string null			
thumbnail	string null			
metadata	Metadata null			
usermetadata	UserMetadata null			
fieldmetadata	array<UserFieldMetadata> null			
computedmetadata	ComputedMetadata null			
created	string null			
modified	string null			
last_seqid	integer null			
last_account_seq	integer null			
queue	QueueType null			
hidden	boolean null			
origin	Origin null			
extra	Extra null			
relations	array<nucliadb_models__metadata__Relation> null			
data	ResourceData null			
security	ResourceSecurity null			Resource security metadata

nucliadb_models__search__Paragraph

Field	Type	Req	Default	Description
score	number	yes		

rid	string	yes		
field_type	string	yes		
field	string	yes		
text	string	yes		
labels	array<string>		[]	
start_seconds	array<integer> null			
end_seconds	array<integer> null			
position	TextPosition null			
fuzzy_result	boolean		False	

nucliadb_models__search__Sentence

Field	Type	Req	Default	Description
score	number	yes		
rid	string	yes		
text	string	yes		
field_type	string	yes		
field	string	yes		
index	string null			
position	TextPosition null			

NucliadbClientType

Enum: `api`, `widget`, `web`, `dashboard`, `desktop`, `chrome_extension`

OpenAIKey

Field	Type	Req	Default	Description
key	string			
org	string			

Operator

Enum: `and`, `or`

Or_FieldFilterExpressionType_

Field	Type	Req	Default	Description
operands	array<And_FieldFilterExpressionType_ Or_FieldFilterExpressionType_ Not_FieldFilterExpressionType_ Resource-Input Field Keyword DateCreated DateModified Label ResourceMimetype FieldMimetype Entity-Input Language OriginTag OriginMetadata OriginPath OriginSource OriginCollaborator nucliadb_models__filters__Generated>	yes		

Or_GraphNodesQuery_

Field	Type	Req	Default	Description
operands	array<And_GraphNodesQuery_ Or_GraphNodesQuery_ Not_GraphNodesQuery_ AnyNode nucliadb_models__graph__requests__Generated>	yes		

Or_GraphPathQuery_

Field	Type	Req	Default	Description
operands	array<And_GraphPathQuery_ Or_GraphPathQuery_	yes		

	 Not_GraphPathQuery_ _ GraphPath-Input SourceNode DestinationNode AnyNode Relation- Input nucliadb_models__gr aph__requests__Gene rated>			
--	---	--	--	--

Or_GraphRelationsQuery_

Field	Type	Req	Default	Description
<code>operands</code>	array<And_GraphRel ationsQuery_ Or_GraphRelationsQ uery_ Not_GraphRelations Query_ Relation- Input nucliadb_models__gr aph__requests__Gene rated>	yes		

Or_KVFilterExpression_

Field	Type	Req	Default	Description
<code>operands</code>	array<And_KVFilter Expression_ Or_KVFilterExpressi on_ Not_KVFilterExpress ion_ Eq Inequalities Contains>	yes		

Or_ParagraphFilterExpression_

Field	Type	Req	Default	Description
<code>operands</code>	array<And_Paragrap hFilterExpression_	yes		

	Or_ParagraphFilterExpression_ Not_ParagraphFilterExpression_ Label Kind>			
--	--	--	--	--

Or_ResourceFilterExpression_

Field	Type	Req	Default	Description
<code>operands</code>	array<And_ResourceFilterExpression_ Or_ResourceFilterExpression_ Not_ResourceFilterExpression_ Resource-Input DateCreated DateModified Label ResourceMimetype Language OriginTag OriginMetadata OriginPath OriginSource OriginCollaborator Status>	yes		

Origin

Field	Type	Req	Default	Description
<code>source_id</code>	string null			
<code>url</code>	string null			
<code>created</code>	string null			
<code>modified</code>	string null			
<code>metadata</code>	map<string, string>		<code>{}</code>	Generic metadata from the resource at the origin system. It can later be used for filtering on search endpoints with <code>'origin.metadata/{key}/{value}'</code>

tags	array<string>		[]	Resource tags about the origin system. It can later be used for filtering on search endpoints with '/origin.tags/{tag}'
collaborators	array<string>		[]	
filename	string null			
related	array<string>		[]	
path	string null			Path of the original resource. Typically used to store folder structure information of the resource at the origin system. It can be later used for filtering on search endpoints with '/origin.path/{pat...'
sync_metadata	SyncMetadata null			Metadata related to the resource from the origin system fetched by the Progress Agentic RAG's Cloud Storage Sync service.
source	Source null		API	

OriginCollaborator

Matches the origin collaborators

Field	Type	Req	Default	Description
prop	string		origin_collaborator	
collaborator	string	yes		Collaborator

OriginMetadata

Matches metadata from the origin

Field	Type	Req	Default	Description
-------	------	-----	---------	-------------

<code>prop</code>	string		<code>origin_metadata</code>	
<code>field</code>	string	yes		Metadata field
<code>value</code>	string null			Value of the metadata field. If blank, matches any document with the given metadata field set (to any value)

OriginPath

Matches the origin path

Field	Type	Req	Default	Description
<code>prop</code>	string		<code>origin_path</code>	
<code>prefix</code>	string null			Prefix of the path, matches all paths under this prefix.g: prefix=/dir/ matches /dir and /dir/a/b but not /dirrrr

OriginSource

Matches the origin source id

Field	Type	Req	Default	Description
<code>prop</code>	string		<code>origin_source</code>	
<code>id</code>	string null			Source ID

OriginTag

Matches all fields with a given origin tag

Field	Type	Req	Default	Description
<code>prop</code>	string		<code>origin_tag</code>	

tag	string	yes		The tag to match
-----	--------	-----	--	------------------

PageImageStrategy

Field	Type	Req	Default	Description
name	string		page_image	
count	integer null			Maximum number of page images to retrieve. By default, at most 5 images are retrieved.

PageInformation

Field	Type	Req	Default	Description
page	integer null			
page_with_visual	boolean null			

PagePositions

Field	Type	Req	Default	Description
start	integer null			
end	integer null			

PageStructure

Field	Type	Req	Default	Description
page	PageStructurePage	yes		
tokens	array<PageStructureToken>	yes		

PageStructurePage

Field	Type	Req	Default	Description
-------	------	-----	---------	-------------

<code>width</code>	integer	yes		
<code>height</code>	integer	yes		

PageStructureToken

Field	Type	Req	Default	Description
<code>x</code>	number	yes		
<code>y</code>	number	yes		
<code>width</code>	number	yes		
<code>height</code>	number	yes		
<code>text</code>	string	yes		
<code>line</code>	number	yes		

PalmKey

Field	Type	Req	Default	Description
<code>credentials</code>	string			
<code>location</code>	string			

Paragraph

Field	Type	Req	Default	Description
<code>start</code>	integer null			
<code>end</code>	integer null			
<code>start_seconds</code>	array<integer> null			
<code>end_seconds</code>	array<integer> null			
<code>kind</code>	TypeParagraph null			
<code>classifications</code>	array<Classification> null			
<code>sentences</code>	array<Sentence> null			
<code>key</code>	string null			

<code>page</code>	PageInformation null			
<code>representation</code>	Representation null			
<code>relations</code>	ParagraphRelations null			

ParagraphAnnotation

Field	Type	Req	Default	Description
<code>classifications</code>	array<UserClassification>		<code>[]</code>	
<code>key</code>	string	yes		

ParagraphImageStrategy

Field	Type	Req	Default	Description
<code>name</code>	string		<code>paragraph_image</code>	

ParagraphRelations

Field	Type	Req	Default	Description
<code>parents</code>	array<string>		<code>[]</code>	
<code>siblings</code>	array<string>		<code>[]</code>	
<code>replacements</code>	array<string>		<code>[]</code>	

Paragraphs

Field	Type	Req	Default	Description
<code>results</code>	array<nucliadb_model_search__Paragraph>		<code>[]</code>	
<code>facets</code>	object (free-form map) null			

query	string null			
total	integer		0	
page_number	integer		0	
page_size	integer		20	
next_page	boolean		False	
min_score	number	yes		Minimum bm25 score used to filter bm25 index search. Results with a lower score have been ignored.

PathMetadata

Field	Type	Req	Default	Description
field_id	string null			Field id where the relation has been extracted from
paragraph_id	string null			Paragraph id where the relation has been extracted from

Position

Field	Type	Req	Default	Description
start	integer	yes		
end	integer	yes		

Positions

Field	Type	Req	Default	Description
position	array<Position>	yes		
entity	string	yes		

PredictProxiedEndpoints

Enum for the different endpoints that are proxied to the Predict API

Enum: `tokens`, `chat`, `rephrase`, `run-agents-text`, `summarize`, `rerank`, `remi`

PredictReranker

Field	Type	Req	Default	Description
<code>name</code>	string		<code>predict</code>	
<code>window</code>	integer null			Number of elements reranker will use. Window must be greater or equal to topk. Greater values will improve results at cost of retrieval and reranking time. By default, this reranker uses a default of ...

PreQueriesStrategy

This strategy allows to run a set of queries before the main query and add the results to the context. It allows to give more importance to some queries over others by setting the weight of each query. The weight of the main query can also be set with the `main_query_weight` parameter.

Field	Type	Req	Default	Description
<code>name</code>	string		<code>prequeries</code>	
<code>queries</code>	array<PreQuery>	yes		List of queries to run before the main query. The results are added to the context with the specified weights for each query. There is a limit of 10 prequeries per request.
<code>main_query_weight</code>	number		<code>1.0</code>	Weight of the main query in the context. Use this to control the importance of the main query in the

				context.
--	--	--	--	----------

PreQuery

Field	Type	Req	Default	Description
<code>request</code>	FindRequest	yes		The request to be executed before the main query.
<code>weight</code>	number		<code>1.0</code>	Weight of the prequery in the context. The weight is used to scale the results of the prequery before adding them to the context. The weight should be a positive number, and they are normalized so that...
<code>id</code>	string null			Identifier of the prequery. If not specified, it is autogenerated based on the index of the prequery in the list (prequery0, prequery1, ...).
<code>prefilter</code>	boolean		<code>False</code>	If set to true, the prequery results are used to filter the scope of the remaining queries. The resources of the most relevant paragraphs of the prefilter queries are used as resource filters for the ...

PushProcessingOptions

Field	Type	Req	Default	Description
-------	------	-----	---------	-------------

<code>ml_text</code>	boolean null		<code>True</code>	
----------------------	----------------	--	-------------------	--

QueryEntityDetection

Enum: `predict`, `suggest`

Question

Field	Type	Req	Default	Description
<code>text</code>	string	yes		
<code>language</code>	string null			
<code>ids_paragraphs</code>	array<string>	yes		

QuestionAnswer

Field	Type	Req	Default	Description
<code>question</code>	Question	yes		
<code>answers</code>	array<Answer>	yes		

QuestionAnswerAnnotation

Field	Type	Req	Default	Description
<code>question_answer</code>	QuestionAnswer	yes		
<code>cancelled_by_user</code>	boolean		<code>False</code>	

QuestionAnswers

Field	Type	Req	Default	Description
<code>question_answer</code>	array<QuestionAnswer>	yes		

QueueTypeEnum: `private`, `shared`**Range**Type: `object` (free-form map)**RankFusionName**Enum: `rrf`**Reasoning**

Field	Type	Req	Default	Description
<code>display</code>	boolean		<code>True</code>	Whether to display the reasoning steps in the response.
<code>effort</code>	enum[none, minimal, low, medium, high, xhigh]		<code>medium</code>	Level of reasoning effort. Used by OpenAI models to control the depth of reasoning. This parameter will be automatically mapped to <code>budgettokens</code> if the chosen model does not support effort.
<code>budget_tokens</code>	integer		<code>15000</code>	Token budget for reasoning. Used by Anthropic or Google models to limit the number of tokens used for reasoning. This parameter will be automatically mapped to <code>effort</code> if the chosen model does not supp...

ReciprocalRankFusion

Field	Type	Req	Default	Description
<code>name</code>	string		<code>rrf</code>	

k	number		60.0	k parameter changes the influence top-ranked and lower-ranked elements have. Research has shown that 60 is a performant value across datasets
window	integer null			Number of elements for retrieval to do RRF. Window must be greater or equal to topk. Greater values will increase probability of multi match at cost of retrieval time
boosting	ReciprocalRankFusionWeights			Define different weights for each retriever. This allows to assign different priorities to different retrieval methods. RRF scores will be multiplied by this value. The default is 1 for each retriever...

ReciprocalRankFusionWeights

Field	Type	Req	Default	Description
keyword	number		1.0	
semantic	number		1.0	

RelatedEntities

Field	Type	Req	Default	Description
total	integer		0	
entities	array<RelatedEntity>		[]	

RelatedEntity

Field	Type	Req	Default	Description
family	string	yes		
value	string	yes		

Relation-Input

Field	Type	Req	Default	Description
prop	string		relation	
label	string null			
type	RelationType null			
match	RelationMatchKindName		exact	

Relation-Output

Field	Type	Req	Default	Description
relation	RelationType	yes		
label	string null			
metadata	RelationMetadata null			
from	RelationEntity null			
to	RelationEntity	yes		

RelationDirection

Enum: `in`, `out`

RelationEdgeVector

Field	Type	Req	Default	Description
relation_label	string	yes		
vector	array<number>	yes		

RelationEntity

Field	Type	Req	Default	Description
value	string	yes		
type	RelationNodeType	yes		
group	string null			

RelationMatchKindName

Enum: `exact`, `semantic`

RelationMetadata

Field	Type	Req	Default	Description
paragraph_id	string null			
source_start	integer null			
source_end	integer null			
to_start	integer null			
to_end	integer null			
data_augmentation_task_id	string null			

RelationNodeType

Enum: `entity`, `label`, `resource`, `user`

RelationNodeVector

Field	Type	Req	Default	Description
node_value	string	yes		
vector	array<number>	yes		

RelationRanking

Enum: `reranker`, `generative`

Relations

Field	Type	Req	Default	Description
<code>entities</code>	map<string, EntitySubgraph>	yes		

RelationType

Enum: `ABOUT`, `CHILD`, `COLAB`, `ENTITY`, `OTHER`, `SYNONYM`

Representation

Field	Type	Req	Default	Description
<code>is_a_table</code>	boolean null			
<code>reference_file</code>	string null			

RequestSecurity

Security metadata for the search request

Field	Type	Req	Default	Description
<code>access_groups</code>	array<string>		<code>[]</code>	List of group ids to do the request with.

RequestsResult

Field	Type	Req	Default	Description
<code>processing_id</code>	string	yes		Processing ID of the resource.
<code>resource_id</code>	string	yes		Resource ID.
<code>kbid</code>	string	yes		
<code>title</code>	string null			Title of the resource.
<code>labels</code>	array<string>		<code>[]</code>	Labels of the resource.
<code>completed</code>	boolean	yes		Whether the resource has been completed

<code>scheduled</code>	boolean	yes		Whether the resource has been scheduled
<code>timestamp</code>	string	yes		Timestamp of when the resource was first scheduled.
<code>completed_at</code>	string null			Timestamp of when the resource was completed
<code>scheduled_at</code>	string null			Timestamp of when the resource was first scheduled.
<code>failed</code>	boolean		<code>False</code>	Whether the resource has failed to process
<code>retries</code>	integer		<code>0</code>	Number of retries for the resource.
<code>schedule_eta</code>	number		<code>0.0</code>	Estimated time until the resource is scheduled.
<code>schedule_order</code>	integer		<code>0</code>	Order of the resource in the schedule queue.

RequestsResults

Field	Type	Req	Default	Description
<code>results</code>	array<RequestsResult>		<code>[]</code>	List of results.
<code>cursor</code>	string null			Cursor to use for the next page of results.

RerankerName

Rerankers - Predict reranker: after retrieval, send the results to Predict API to rerank it. This method uses a reranker model, so one can expect better results at the expense of more latency. This will be the new default - No-operation (noop) reranker: maintain order and do not rerank the results a...

Enum: `predict`, `noop`

Resource

Matches all fields of a resource given its id or slug

Field	Type	Req	Default	Description
<code>prop</code>	string		<code>resource</code>	
<code>id</code>	string null			UUID of the resource to match
<code>slug</code>	string null			Slug of the resource to match

Resource-Input

Matches all fields of a resource given its id or slug

Field	Type	Req	Default	Description
<code>prop</code>	string		<code>resource</code>	
<code>id</code>	string null			UUID of the resource to match
<code>slug</code>	string null			Slug of the resource to match

Resource-Output

Field	Type	Req	Default	Description
<code>id</code>	string	yes		
<code>slug</code>	string null			
<code>title</code>	string null			
<code>summary</code>	string null			
<code>icon</code>	string null			
<code>thumbnail</code>	string null			
<code>metadata</code>	Metadata null			
<code>usermetadata</code>	UserMetadata null			
<code>fieldmetadata</code>	array<UserFieldMetadata> null			
<code>computedmetadata</code>	ComputedMetadata null			
<code>created</code>	string null			

modified	string null			
last_seqid	integer null			
last_account_seq	integer null			
queue	QueueType null			
hidden	boolean null			
origin	Origin null			
extra	Extra null			
relations	array<Relation-Output> null			
data	ResourceData null			
security	ResourceSecurity null			Resource security metadata

ResourceAgentsRequest

Field	Type	Req	Default	Description
filters	array<AgentsFilter> null			Filters to apply to the agents. If None, all currently configured agents are applied.
agent_ids	array<string> null			

ResourceAgentsResponse

Field	Type	Req	Default	Description
results	map<string, AugmentedField>	yes		

ResourceCreated

Field	Type	Req	Default	Description
uuid	string	yes		
elapsed	number null			

<code>seqid</code>	integer null			
--------------------	----------------	--	--	--

ResourceData

Field	Type	Req	Default	Description
<code>texts</code>	map<string, TextFieldData> null			
<code>files</code>	map<string, FileFieldData> null			
<code>links</code>	map<string, LinkFieldData> null			
<code>conversations</code>	map<string, ConversationFieldData> null			
<code>generics</code>	map<string, GenericFieldData> null			
<code>key_values</code>	map<string, KeyValueFieldData> null			

ResourceField

Field	Type	Req	Default	Description
<code>field_type</code>	FieldTypeName	yes		This map assumes that both values and extracted data field containers use the same names for its fields. See <code>models.ResourceFieldValues</code> and <code>models.ResourceFieldExtractedData</code>
<code>field_id</code>	string	yes		
<code>value</code>	object			
<code>extracted</code>	TextFieldExtractedData FileFieldExtractedData			

	ta LinkFieldExtractedData ConversationFieldExtractedData null			
error	Error null			
status	string null			
errors	array<Error> null			

ResourceFieldAdded

Field	Type	Req	Default	Description
seqid	integer null			

ResourceFieldPrefix

Matches a field or set of fields. This filter is for internal use only and is not exposed in the public API schema.

Field	Type	Req	Default	Description
prop	string		resource_field_prefix	
resource_id	string null			ID of the resource containing the field(s) to match
resource_slug	string null			Slug of the resource containing the field(s) to match.
field_type	FieldTypeName	yes		Type of the fields to match
field_name_prefix	string	yes		Prefix of the name of the field to match. If blank, matches all fields of the given type in the given resource

ResourceFieldProperties

Enum: `value`, `extracted`, `error`

ResourceFileUploaded

Field	Type	Req	Default	Description
<code>seqid</code>	integer null			
<code>uuid</code>	string null			
<code>field_id</code>	string null			

ResourceList

Field	Type	Req	Default	Description
<code>resources</code>	array<nucliadb_model_resource>	yes		
<code>pagination</code>	ResourcePagination	yes		

ResourceMimetype

Matches resources with a mimetype. The mimetype of a resource can be assigned independently of the mimetype of its fields. In resources with multiple fields, you may prefer to use `field_mimetype`

Field	Type	Req	Default	Description
<code>prop</code>	string		<code>resource_mimetype</code>	
<code>type</code>	string	yes		Type of the mimetype to match. e.g: In image/jpeg, type is image
<code>subtype</code>	string null			Type of the mimetype to match. e.g: In image/jpeg, subtype is jpeg. Leave blank to match all mimetype of the type

ResourcePagination

Field	Type	Req	Default	Description
page	integer	yes		
size	integer	yes		
last	boolean	yes		

ResourceProcessingStatus

Enum: PENDING, PROCESSED, ERROR, EMPTY, BLOCKED, EXPIRED

ResourceProperties

Enum: basic, origin, extra, relations, values, extracted, errors, security

ResourceResult

Field	Type	Req	Default	Description
score	number integer	yes		
rid	string	yes		
field_type	string	yes		
field	string	yes		
labels	array<string> null			

Resources

Field	Type	Req	Default	Description
results	array<ResourceResult>	yes		
facets	object (free-form map) null			
query	string null			
total	integer		0	
page_number	integer		0	
page_size	integer		20	
next_page	boolean		False	

<code>min_score</code>	number	yes		Minimum bm25 score used to filter bm25 index search. Results with a lower score have been ignored.
------------------------	--------	-----	--	--

ResourceSearchResults

Search on resource results

Field	Type	Req	Default	Description
<code>sentences</code>	Sentences null			
<code>paragraphs</code>	Paragraphs null			
<code>relations</code>	Relations null			
<code>nodes</code>	array<map<string, string>> null			
<code>shards</code>	array<string> null			

ResourceSecurity

Security metadata for the resource

Field	Type	Req	Default	Description
<code>access_groups</code>	array<string>		<code>[]</code>	List of group ids that can access the resource.

ResourceUpdated

Field	Type	Req	Default	Description
<code>seqid</code>	integer null			

Row

Field	Type	Req	Default	Description
<code>cell</code>	array<string> null			

RowsPreview

Field	Type	Req	Default	Description
<code>sheets</code>	<code>map<string, Sheet> null</code>			

SCORE_TYPE

Enum: `VECTOR`, `BM25`, `BOTH`, `RERANKER`, `RELATION_RELEVANCE`

SearchOptions

Enum: `fulltext`, `keyword`, `relations`, `semantic`

SearchRequest

Field	Type	Req	Default	Description
<code>audit_metadata</code>	<code>map<string, string> null</code>			A dictionary containing optional audit-specific metadata, such as userid, environment, or other contextual information. This metadata can be leveraged for filtering and analyzing activity logs in futu...
<code>query</code>	<code>string</code>			The query to search for
<code>filter_expression</code>	<code>FilterExpression null</code>			Returns only documents that match this filter expression. Filtering examples can be found here: (see docs) This allows building complex filtering expressions and replaces the following parameters:field...
<code>fields</code>	<code>array<string></code>		<code>[]</code>	The list of fields to

				search in. For instance: a/title to search only on title field. For more details on filtering by field, see: (see docs)
filters	array<string> array<Filter>		[]	The list of filters to apply. Filtering examples can be found here: (see docs)
top_k	integer		20	The number of results search should return. The maximum number of results allowed is 200.
min_score	number MinScore null			Minimum score to filter search results. Results with a lower score will be ignored. Accepts either a float or a dictionary with the minimum scores for the bm25 and vector indexes. If a float is provid...
range_creation_start	string null			Resources created before this date will be filtered out of search results. Datetime are represented as a str in ISO 8601 format, like: 2008-09-15T15:53:00+05:00.
range_creation_end	string null			Resources created after this date will be filtered out of search results. Datetime are represented as a str in ISO 8601 format, like: 2008-09-15T15:53:00+05:00.
range_modification_start	string null			Resources modified before this date will

				be filtered out of search results. Datetime are represented as a str in ISO 8601 format, like: 2008-09-15T15:53:00+05:00.
<code>range_modification_end</code>	string null			Resources modified after this date will be filtered out of search results. Datetime are represented as a str in ISO 8601 format, like: 2008-09-15T15:53:00+05:00.
<code>debug</code>	boolean		<code>False</code>	If set, the response will include some extra metadata for debugging purposes, like the list of queried nodes.
<code>highlight</code>	boolean		<code>False</code>	If set to true, the query terms will be highlighted in the results between <code><mark>...</mark></code> tags
<code>show</code>	array<ResourceProperties>		<code>['basic']</code>	Controls which types of metadata are serialized on resources of search results
<code>field_type_filter</code>	array<FieldTypeName>		<code>['text', 'file', 'link', 'conversation', 'generic', 'key_value']</code>	Define which field types are serialized on resources of search results
<code>extracted</code>	array<ExtractedDataTypeName>		<code>[]</code>	[Deprecated] Please use GET resource endpoint instead to get extracted metadata
<code>vector</code>	array<number> null			The vector to perform

				the search with. If not provided, NucliaDB will use Nuclia Predict API to create the vector off from the query.
<code>vectorset</code>	string null			Vectors index to perform the search in. If not provided, NucliaDB will use the default one
<code>with_duplicates</code>	boolean		<code>False</code>	Whether to return duplicate paragraphs on the same document
<code>with_synonyms</code>	boolean		<code>False</code>	Whether to return matches for custom knowledge box synonyms of the query terms. Note: only supported for keyword and fulltext search options.
<code>resource_filters</code>	array<string>		<code>[]</code>	List of resource ids to filter search results for. Only paragraphs from the specified resources will be returned.
<code>security</code>	RequestSecurity null			Security metadata for the request. Please refer to the documentation for more details on how security works: (see docs)
<code>show_hidden</code>	boolean		<code>False</code>	If set to false (default), excludes hidden resources from search
<code>rephrase</code>	boolean		<code>False</code>	Rephrase the query for a more efficient retrieval. This will consume LLM tokens and make the request

				slower.
<code>rephrase_prompt</code>	string null			Rephrase prompt given to the generative model responsible for rephrasing the query for a more effective retrieval step. This is only used if the rephrase flag is set to true in the request. If not spe...
<code>query_image</code>	Image null			Image that will be used together with the query text for retrieval.
<code>features</code>	array<SearchOptions>		<code>['keyword', 'fulltext', 'semantic']</code>	List of search features to use. Each value corresponds to a lookup into one of the different indexes
<code>faceted</code>	array<string>		<code>[]</code>	The list of facets to calculate. The facets follow the same syntax as filters: (see docs)
<code>sort</code>	SortOptions null			Options for results sorting
<code>offset</code>	integer		<code>0</code>	The number of results to skip, starting from the beginning in sort order. Used for pagination. It can only be used with the keyword and fulltext indexes.

SemanticModelMetadata

Metadata of the semantic model associated to the KB

Field	Type	Req	Default	Description
<code>similarity_fu</code>	VectorSimilarity	yes		Vector similarity

<code>nction</code>				algorithm that is applied on search
<code>vector_dimension</code>	integer null			Dimension of the indexed vectors/embeddings
<code>default_min_score</code>	number null			Deprecated

Sentence

Field	Type	Req	Default	Description
<code>start</code>	integer null			
<code>end</code>	integer null			
<code>key</code>	string null			

Sentences

Field	Type	Req	Default	Description
<code>results</code>	array<nucliadb_model_search__Sentence>		<code>[]</code>	
<code>facets</code>	object (free-form map)	yes		
<code>page_number</code>	integer		<code>0</code>	
<code>page_size</code>	integer		<code>20</code>	
<code>min_score</code>	number	yes		Minimum similarity score used to filter vector index search. Results with a lower score have been ignored.

Sheet

Field	Type	Req	Default	Description
<code>rows</code>	array<Row> null			

SortField

Enum: `score`, `created`, `modified`, `title`

SortOptions

Field	Type	Req	Default	Description
<code>field</code>	SortField	yes		
<code>order</code>	SortOrder		<code>desc</code>	

SortOrder

Enum: `asc`, `desc`

Source

Enum: `WEB`, `DESKTOP`, `API`, `PYSDK`

SourceNode

Field	Type	Req	Default	Description
<code>prop</code>	string		<code>source_node</code>	
<code>value</code>	string null			
<code>match</code>	NodeMatchKindName		<code>exact</code>	
<code>type</code>	RelationNodeType null		<code>entity</code>	
<code>group</code>	string null			

SplitConfig

Field	Type	Req	Default	Description
<code>max_paragraph</code>	integer		<code>0</code>	

SplitConfiguration

Hey, developer! Keep this in sync with corresponding pydantic model in `learning_config.models`

Field	Type	Req	Default	Description
-------	------	-----	---------	-------------

name	string			
max_paragraph	integer		0	
custom_split	CustomSplitStrategy null			
llm_split	LLMSplitConfig null			
manual_split	ManualSplitConfig null			

Status

Matches resource in a certain processing status

Field	Type	Req	Default	Description
prop	string		status	
status	ResourceProcessingStatus	yes		The status of the resource

StatusResponse

Field	Type	Req	Default	Description
status	Status	yes		Matches resource in a certain processing status
total	integer		0	
processed	integer		0	
retries	integer		0	

StoredLearningConfiguration

Model to map in a generic way what we really store on the db, without validations. As enum values containing the versions change from time to time, and we don't keep historics, we cannot use the model enums here, as it will fail with older values

Field	Type	Req	Default	Description
-------	------	-----	---------	-------------

<code>semantic_model</code>	string	yes		
<code>anonymization_model</code>	string	yes		
<code>generative_model</code>	string	yes		
<code>ner_model</code>	string	yes		
<code>semantic_vector_similarity</code>	string	yes		
<code>semantic_vector_size</code>	integer			

SuggestOptions

Enum: `paragraph`, `entities`

SuggestRequest

Field	Type	Req	Default	Description
<code>query</code>	string	yes		The query to get suggestions for
<code>features</code>	array<SuggestOptions>		<code>['paragraph', 'entities']</code>	Features enabled for the suggest endpoint.
<code>filter_expression</code>	FilterExpression null			Returns only documents that match this filter expression. Filtering examples can be found here: (see docs) This allows building complex filtering expressions and replaces the following parameters: <code>field...</code>
<code>security</code>	RequestSecurity null			Security metadata for the request. Please refer to the documentation for more details on how security works: (see

				docs)
<code>show_hidden</code>	boolean		<code>False</code>	If set to false (default), excludes hidden resources from search
<code>highlight</code>	boolean		<code>False</code>	If set to true, the query terms will be highlighted in the results between <code><mark>...</mark></code> tags

SummarizedResource

Field	Type	Req	Default	Description
<code>summary</code>	string	yes		Summary of the resource
<code>tokens</code>	integer	yes		

SummarizedResponse

Field	Type	Req	Default	Description
<code>resources</code>	map<string, SummarizedResource>		<code>{}</code>	Individual resource summaries. The key is the resource id or slug.
<code>summary</code>	string			Global summary of all resources combined.
<code>consumption</code>	Consumption null			

SummarizeRequest

Model for the request payload of the summarize endpoint

Field	Type	Req	Default	Description
<code>generative_model</code>	string null			The generative model to use for the summarization. If not

				provided, the model configured for the Knowledge Box is used.
<code>user_prompt</code>	string null			Optional custom prompt input by the user
<code>resources</code>	array<string>	yes		Uids or slugs of the resources to summarize. If the resources are not found, they will be ignored.
<code>summary_kind</code>	SummaryKind		<code>simple</code>	Option to customize how the summary will be

SummaryKind

Enum: `simple`, `extended`

SyncAskMetadata

Field	Type	Req	Default	Description
<code>tokens</code>	AskTokens null			Number of tokens used in the LLM context and answer
<code>timings</code>	AskTimings null			Timings of the generative model

SyncAskResponse

Field	Type	Req	Default	Description
<code>answer</code>	string	yes		The generative answer to the query
<code>reasoning</code>	string null			The reasoning steps followed by the LLM to generate the answer. This is returned only if the reasoning feature is enabled in the

				request.
<code>answer_json</code>	object (free-form map) null			The generative JSON answer to the query. This is returned only if the <code>answerjsonschema</code> parameter is provided in the request.
<code>status</code>	string	yes		The status of the query execution. It can be 'success', 'error', 'nocontext' or 'noretrievaldata'
<code>retrieval_results</code>	KnowledgeboxFindResults	yes		The retrieval results of the query
<code>retrieval_best_matches</code>	array<AskRetrievalMatch>		<code>[]</code>	Sorted list of best matching text blocks in the retrieval step. This includes the main query and prequeries results, if any.
<code>prequeries</code>	map<string, KnowledgeboxFindResults> null			The retrieval results of the prequeries
<code>learning_id</code>	string			The id of the learning request. This id can be used to provide feedback on the learning process.
<code>relations</code>	Relations null			The detected relations of the answer
<code>citations</code>	object (free-form map)			The citations of the answer. List of references to the resources used to generate the answer.
<code>citation_footnote_to_context</code>	map<string, string>			Maps ids in the footnote citations to querycontext keys (normally paragraph ids)

<code>augmented_context</code>	AugmentedContext null			Augmented text blocks that were sent to the LLM as part of the RAG strategies applied on the retrieval results in the request.
<code>prompt_context</code>	array<string> null			The prompt context used to generate the answer. Returned only if the debug flag is set to true
<code>predict_request</code>	object (free-form map) null			The internal predict request used to generate the answer. Returned only if the debug flag is set to true
<code>metadata</code>	SyncAskMetadata null			Metadata of the query execution. This includes the number of tokens used in the LLM context and answer, and the timings of the generative model.
<code>consumption</code>	Consumption null			The consumption of the query execution. Return only if 'X-show-consumption' header is set to true in the request.
<code>error_details</code>	string null			Error details message in case there was an error
<code>debug</code>	object (free-form map) null			Debug information about the ask operation. The metadata included in this field is subject to change and should not be used in production. Note that it is only available if the debug parameter

				is set t...
--	--	--	--	-------------

SyncMetadata

Field	Type	Req	Default	Description
<code>file_id</code>	string	yes		Identifier of the file in the origin cloud storage system
<code>auth_provider</code>	string	yes		Authentication provider used to access the origin cloud storage system
<code>content_hash</code>	string	yes		Content hash of the file in the origin cloud storage system. The hash algorithm used depends on the origin system.

TableImageStrategy

Field	Type	Req	Default	Description
<code>name</code>	string		<code>tables</code>	

TextBlockAugmentationType

Enum: `neighbouring_paragraphs`, `conversation`, `hierarchy`, `full_resource`, `field_extension`, `metadata_extension`

TextField

Field	Type	Req	Default	Description
<code>body</code>	string	yes		The text body. The format of the text should be specified in the format field. The sum of all text fields in the request may not exceed 2MB. If you need to store more text, consider

				using a file field...
<code>format</code>	TextFormat		PLAIN	The format of the text.
<code>extract_strategy</code>	string null			Id of the Nuclia extract strategy to use at processing time. If not set, the default strategy will be used. Extract strategies are defined at the learning configuration api.
<code>split_strategy</code>	string null			Id of the Nuclia split strategy used at processing time. If not set, the default strategy was used. Split strategies are defined at the learning configuration api.

TextFieldData

Field	Type	Req	Default	Description
<code>value</code>	FieldText null			
<code>extracted</code>	TextFieldExtractedData null			
<code>error</code>	Error null			
<code>status</code>	string null			
<code>errors</code>	array<Error> null			

TextFieldExtractedData

Field	Type	Req	Default	Description
<code>text</code>	ExtractedText null			
<code>metadata</code>	FieldComputedMetadata null			

<code>large_metadata</code>	LargeComputedMetadata null			
<code>vectors</code>	VectorObject null			
<code>question_answers</code>	FieldQuestionAnswers null			
<code>relation_node_vectors</code>	map<string, array<RelationNodeVector>> null			
<code>relation_edge_vectors</code>	map<string, array<RelationEdgeVector>> null			

TextFormat

Enum: `PLAIN`, `HTML`, `RST`, `MARKDOWN`, `JSON`, `KEEP_MARKDOWN`, `JSONL`, `PLAIN_BLANKLINE_SPLIT`

TextGenerationKey

Field	Type	Req	Default	Description
<code>model</code>	string			

TextPosition

Field	Type	Req	Default	Description
<code>page_number</code>	integer null			
<code>index</code>	integer	yes		
<code>start</code>	integer	yes		
<code>end</code>	integer	yes		
<code>start_seconds</code>	array<integer> null			
<code>end_seconds</code>	array<integer> null			

TokensDetail

Field	Type	Req	Default	Description
<code>input</code>	number	yes		

<code>output</code>	number	yes		
<code>image</code>	number	yes		

TypeParagraph

Enum: `TEXT`, `OCR`, `INCEPTION`, `DESCRIPTION`, `TRANSCRIPT`, `TITLE`, `TABLE`

UpdateKVSchema

Field	Type	Req	Default	Description
<code>description</code>	string null			
<code>fields</code>	array<KVSchemaField> null			

UpdateResourcePayload

Field	Type	Req	Default	Description
<code>title</code>	string null			
<code>summary</code>	string null			
<code>slug</code>	string null			The slug is the user-defined id for the resource
<code>thumbnail</code>	string null			
<code>metadata</code>	InputMetadata null			
<code>usermetadata</code>	UserMetadata null			
<code>fieldmetadata</code>	array<UserFieldMetadata> null			
<code>origin</code>	InputOrigin null			
<code>extra</code>	Extra null			Extra metadata for the resource. It can be used to store structured information about the resource that can't be used to query at retrieval time. If not set, the existing extra metadata will not be

				mo...
files	object		{}	Dictionary of file fields to be added to the resource. The keys correspond to the field id, and must comply with the regex: <code>^[a-zA-Z0-9:-]+\$</code>
links	object		{}	Dictionary of link fields to be added to the resource. The keys correspond to the field id, and must comply with the regex: <code>^[a-zA-Z0-9:-]+\$</code>
texts	object		{}	Dictionary of text fields to be added to the resource. The keys correspond to the field id, and must comply with the regex: <code>^[a-zA-Z0-9:-]+\$</code>
conversations	object		{}	Dictionary of conversation fields to be added to the resource. The keys correspond to the field id, and must comply with the regex: <code>^[a-zA-Z0-9:-]+\$</code>
key_values	object		{}	Dictionary of key-value fields to be added to the resource. The key must be the schema name and must comply with the regex: <code>^[a-zA-Z0-9:-]+\$</code>
processing_options	PushProcessingOptions null		{'ml_text': True}	Options for processing the resource. If not set,

				the default options will be used.
<code>security</code>	ResourceSecurity null			Security metadata for the resource. It can be used to have fine-grained control over who can access the resource.
<code>hidden</code>	boolean null			Modify the hidden status of the resource. If not set, the hidden status will not be modified.

UserClassification

Field	Type	Req	Default	Description
<code>labelset</code>	string	yes		
<code>label</code>	string	yes		
<code>cancelled_by_user</code>	boolean		<code>False</code>	

UserFieldMetadata

Field-level metadata set by the user via the rest api

Field	Type	Req	Default	Description
<code>paragraphs</code>	array<ParagraphAnnotation>		<code>[]</code>	
<code>question_answers</code>	array<QuestionAnswerAnnotation>		<code>[]</code>	
<code>field</code>	FieldID	yes		

UserLearningKeys

Field	Type	Req	Default	Description
<code>openai</code>	OpenAIKey null			

azure_openai	AzureOpenAIKey null			
palm	PalmKey null			
anthropic	AnthropicKey null			
claude3	AnthropicKey null			
text_generation	TextGenerationKey null			
mistral	MistralKey null			
azure_mistral	AzureMistralKey null			
hf_llm	HFLLMKey null			
hf_embedding	HFEmbeddingKey null			

UserMetadata

Field	Type	Req	Default	Description
classifications	array<UserClassification>		[]	
relations	array<Relation-Output>		[]	

ValidationError

Field	Type	Req	Default	Description
loc	array<string integer>	yes		
msg	string	yes		
type	string	yes		
input	object			
ctx	object			

Vector

Field	Type	Req	Default	Description
-------	------	-----	---------	-------------

<code>start</code>	integer null			
<code>end</code>	integer null			
<code>start_paragraph</code>	integer null			
<code>end_paragraph</code>	integer null			
<code>vector</code>	array<number> null			

VectorObject

Field	Type	Req	Default	Description
<code>vectors</code>	Vectors null			
<code>split_vectors</code>	map<string, Vectors> null			
<code>deleted_splits</code>	array<string> null			

Vectors

Field	Type	Req	Default	Description
<code>vectors</code>	array<Vector> null			

VectorSimilarity

Enum: `cosine`, `dot`

VLLMExtractionConfig

Field	Type	Req	Default	Description
<code>rules</code>	array<string>			
<code>llm</code>	LLMConfig null			

nua — 481 schemas

AccountTypes

Enum: stash-trial, stash-starter, stash-growth, stash-startup, stash-enterprise, stash-developer, stash-business, v3starter, v3fly, v3growth, v3pro, v3enterprise, cowork

AgentID

Field	Type	Req	Default	Description
id	string	yes		

AITables-Input

Field	Type	Req	Default	Description
llm	LLMConfig-Input null			
rules	array<string>			
merge_pages	boolean null		False	
max_pages_to_merge	integer null		0	

AITables-Output

Field	Type	Req	Default	Description
llm	LLMConfig-Output null			
rules	array<string>			
merge_pages	boolean null		False	
max_pages_to_merge	integer null		0	

AITablesConfig

Field	Type	Req	Default	Description
-------	------	-----	---------	-------------

<code>generative_model</code>	string null			
<code>generative_provider</code>	string null			
<code>rules</code>	array<string>		<code>[]</code>	
<code>merge_pages</code>	boolean null			
<code>max_pages_to_merge</code>	integer null			

And_FieldFilterExpressionType_

Field	Type	Req	Default	Description
<code>operands</code>	array<And_FieldFilterExpressionType_ Or_FieldFilterExpressionType_ Not_FieldFilterExpressionType_ Resource-Input Field Keyword DateCreated DateModified Label ResourceMimetype FieldMimetype Entity-Input Language OriginTag OriginMetadata OriginPath OriginSource OriginCollaborator Generated>	yes		

And_GraphPathQuery_

Field	Type	Req	Default	Description
<code>operands</code>	array<And_GraphPathQuery_ Or_GraphPathQuery_ Not_GraphPathQuery_ GraphPath	yes		

	SourceNode DestinationNode AnyNode Relation- Input nucliadb_models__gr aph__requests__Gene rated>			
--	---	--	--	--

And_KVFilterExpression_

Field	Type	Req	Default	Description
operands	array<And_KVFilter Expression_ Or_KVFilterExpressi on_ Not_KVFilterExpress ion_ Eq Inequalities Contains>	yes		

And_ParagraphFilterExpression_

Field	Type	Req	Default	Description
operands	array<And_Paragraph FilterExpression_ Or_ParagraphFilterE xpression_ Not_ParagraphFilterE xpression_ Label Kind>	yes		

AnonymizationModel

Enum: disabled, multilingual

Answer

Field	Type	Req	Default	Description
text	string	yes		
language	string null			
ids_paragraph	array<string>	yes		

s				
---	--	--	--	--

AnswerRelevance

Field	Type	Req	Default	Description
score	integer	yes		
reason	string	yes		

AnthropicBedrockKey

Field	Type	Req	Default	Description
access_key	string	yes		
secret_key	string	yes		

AnthropicKey

Field	Type	Req	Default	Description
key	string			

AnthropicUserPrompt

Field	Type	Req	Default	Description
prompt	string			

AnthropicVertexKey

Field	Type	Req	Default	Description
credentials	string	yes		Service Account Key credentials in JSON format for Vertex AI access
location	string			GCP region where the model will be accessed from. E.g. us-central1

AnyNode

Field	Type	Req	Default	Description
prop	string		node	
value	string null			
match	NodeMatchKindName		exact	
type	RelationNodeType null		entity	
group	string null			

ApplyOptions

Defines how the tasks should be applied to the existing data. - EXSITING: Only apply to existing data (starts a worker that executes the task) - NEW: Only apply to new data (enables the task at processing time) - ALL: Apply to all data (both of the above)

Enum: `EXISTING`, `NEW`, `ALL`

ApplyTo

Enum: `0`, `1`

AskOperation

Field	Type	Req	Default	Description
question	string	yes		
destination	string	yes		
json	boolean		False	
triggers	array<Trigger>			
user_prompt	string			Prompt to use when executing the agent over a field, it must include the {context} placeholder and optionally the {question} placeholder if you wish. If not given, the

				default prompt will be used. Onl...
store_as_key_value	boolean		False	When set to true and json is also true, the generated JSON output is stored as a FieldKeyValue instead of a JSON TextField. Only applies to ApplyTo.FIELD mode.
kv_schema_id	string			The ID of the KV schema registered in the KnowledgeBox that the generated JSON output should conform to. Required when storeaskeyvalue is true. Falls back to destination if not set.

AskRequest

Field	Type	Req	Default	Description
audit_metadata	map<string, string> null			A dictionary containing optional audit-specific metadata, such as userid, environment, or other contextual information. This metadata can be leveraged for filtering and analyzing activity logs in futu...
query	string	yes		The query to get a generative answer for
top_k	integer		20	The top most relevant results to fetch at the retrieval step. The maximum number of results allowed is 200.

<code>filter_expression</code>	FilterExpression null			Returns only documents that match this filter expression. Filtering examples can be found here: (see docs) This allows building complex filtering expressions and replaces the following parameters: field...
<code>fields</code>	array<string>		<code>[]</code>	The list of fields to search in. For instance: a/title to search only on title field. For more details on filtering by field, see: (see docs)
<code>filters</code>	array<string> array<Filter>		<code>[]</code>	The list of filters to apply. Filtering examples can be found here: (see docs)
<code>keyword_filters</code>	array<string> array<Filter>		<code>[]</code>	List of keyword filter expressions to apply to the retrieval step. The text block search will only be performed on the documents that contain the specified keywords. The filters are case-insensitive, ...
<code>vectorset</code>	string null			Vectors index to perform the search in. If not provided, NucliaDB will use the default one
<code>min_score</code>	number MinScore null			Minimum score to filter search results. Results with a lower score will be ignored. Accepts either a float or a dictionary with

				the minimum scores for the bm25 and vector indexes. If a float is provid...
features	array<ChatOptions>		['semantic', 'keyword']	Features enabled for the chat endpoint. Semantic search is done if semantic is included. If keyword is included, the results will include matching paragraphs from the bm25 index. If relations is inclu...
range_creation_start	string null			Resources created before this date will be filtered out of search results. Datetime are represented as a str in ISO 8601 format, like: 2008-09-15T15:53:00+05:00.
range_creation_end	string null			Resources created after this date will be filtered out of search results. Datetime are represented as a str in ISO 8601 format, like: 2008-09-15T15:53:00+05:00.
range_modification_start	string null			Resources modified before this date will be filtered out of search results. Datetime are represented as a str in ISO 8601 format, like: 2008-09-15T15:53:00+05:00.
range_modification_end	string null			Resources modified after this date will be filtered out of search results. Datetime are represented as a str in ISO 8601 format,

				like: 2008-09-15T15:53:00+05:00.
<code>show</code>	<code>array<ResourceProperties></code>		<code>['basic']</code>	Controls which types of metadata are serialized on resources of search results
<code>field_type_filter</code>	<code>array<FieldTypeName></code>		<code>['text', 'file', 'link', 'conversation', 'generic', 'key_value']</code>	Define which field types are serialized on resources of search results
<code>extracted</code>	<code>array<ExtractedDataTypeName></code>		<code>[]</code>	[Deprecated] Please use GET resource endpoint instead to get extracted metadata
<code>context</code>	<code>array<ChatContextMessage> null</code>			DEPRECATED! Please, use <code>chathistory</code> instead.
<code>chat_history</code>	<code>array<ChatContextMessage> null</code>			Use to rephrase the new LLM query by taking into account the chat conversation history. This will be passed to the LLM so that it is aware of the previous conversation.
<code>extra_context</code>	<code>array<string> null</code>			Additional context that is added to the retrieval context sent to the LLM. It allows extending the chat feature with content that may not be in the Knowledge Box.
<code>extra_context_images</code>	<code>array<Image> null</code>			Additional images added to the retrieval context sent to the LLM." It allows extending the chat feature with content

				that may not be in the Knowledge Box.
query_image	Image null			Image that will be used together with the query text for retrieval and then sent to the LLM as part of the context. If a query image is provided, the extracontextimages and ragimagesstrategies will be...
highlight	boolean		False	If set to true, the query terms will be highlighted in the results between <mark>...</mark> tags
resource_filters	array<string>		[]	List of resource ids to filter search results for. Only paragraphs from the specified resources will be returned.
prompt	string CustomPrompt null			Use to customize the prompts given to the generative model. Both system and user prompts can be customized. If a string is provided, it is interpreted as the user prompt.
rank_fusion	RankFusionName ReciprocalRankFusion		rrf	Rank fusion algorithm to use to merge results from multiple retrievers (keyword, semantic)
reranker	RerankerName PredictReranker		predict	Reranker let you specify which method you want to use to rerank your results at the end of retrieval

<code>citations</code>	boolean CitationsType null			Whether to include citations in the response. If set to None or False, no citations will be computed. If set to True or 'default', citations will be computed after answer generation and send as a sepa...
<code>citation_threshold</code>	number null			If citations is set to True or 'default', this will be the similarity threshold. Value between 0 and 1, lower values will produce more citations. If not set, it will be set to the optimized threshold ...
<code>security</code>	RequestSecurity null			Security metadata for the request. Please refer to the documentation for more details on how security works: (see docs)
<code>show_hidden</code>	boolean		<code>False</code>	If set to false (default), excludes hidden resources from search
<code>rag_strategies</code>	array<FieldExtensionStrategy FullResourceStrategy HierarchyResourceStrategy NeighbouringParagraphsStrategy MetadataExtensionStrategy ConversationalStrategy PreQueriesStrategy		<code>[]</code>	Options for tweaking how the context for the LLM model is crafted: - fullresource will add the full text of the matching resources to the context. This strategy cannot be combined with hierarchy, neig...

	GraphStrategy>			
rag_images_strategies	array<PageImageStrategy ParagraphImageStrategy TableImageStrategy>		[]	Options for tweaking how the image based context for the LLM model is crafted: - pageimage will add the full page image of the matching resources to the context. - tables will send the table images fo...
debug	boolean		False	If set, the response will include some extra metadata for debugging purposes, like the list of queried nodes.
generative_model	string null			The generative model to use for the chat endpoint. If not provided, the model configured for the Knowledge Box is used.
generative_model_seed	integer null			The seed to use for the generative model for deterministic generation. Only supported by some models.
max_tokens	integer MaxTokens null			Use to limit the amount of tokens used in the LLM context and/or for generating the answer. If not provided, the default maximum tokens of the generative model will be used. If an integer is provided,...
rephrase	boolean		False	Rephrase the query for a more efficient retrieval. This will

				consume LLM tokens and make the request slower.
<code>chat_history_relevance_threshold</code>	number null			Threshold to determine if the past chat history is relevant to rephrase the user's question. 0 - Always treat previous messages as relevant (always rephrase).1 - Always treat previous messages as irre...
<code>prefer_markdown</code>	boolean		False	If set to true, the response will be in markdown format
<code>answer_json_schema</code>	object (free-form map) null			Desired JSON schema for the LLM answer. This schema is passed to the LLM so that it answers in a structured format following the schema. If not provided, textual response is returned. Note that when u...
<code>generate_answer</code>	boolean		True	Whether to generate an answer using the generative model. If set to false, the response will only contain the retrieval results.
<code>search_configuration</code>	string null			Load ask parameters from this configuration. Parameters in the request override parameters from the configuration.
<code>reasoning</code>	Reasoning boolean		False	Reasoning options for the generative model. Set to True to enable

				default reasoning, False to disable, or provide a Reasoning object for custom options.
--	--	--	--	--

AskRetrievalMatch

Field	Type	Req	Default	Description
<code>id</code>	string	yes		Id of the matching text block

AskTimings

Field	Type	Req	Default	Description
<code>generative_first_chunk</code>	number null			Time the LLM took to generate the first chunk of the answer
<code>generative_total</code>	number null			Total time the LLM took to generate the answer

AskTokens

Field	Type	Req	Default	Description
<code>input</code>	integer	yes		Number of LLM tokens used for the context in the query
<code>output</code>	integer	yes		Number of LLM tokens used for the answer
<code>input_nuclia</code>	number null			Number of Nuclia LLM tokens used for the context in the query
<code>output_nuclia</code>	number null			Number of Nuclia LLM tokens used for the answer

AssetList

Field	Type	Req	Default	Description
<code>assets</code>	array<AssetResponse>	yes		

AssetResponse

Returned to the worker after a successful upload.

Field	Type	Req	Default	Description
<code>id</code>	string	yes		
<code>task_id</code>	string	yes		
<code>task_name</code>	string	yes		
<code>kbid</code>	string	yes		
<code>account_id</code>	string	yes		
<code>created_at</code>	string	yes		
<code>content_type</code>	string	yes		

AssetSignedURL

Field	Type	Req	Default	Description
<code>url</code>	string	yes		

AssistantMessage

Field	Type	Req	Default	Description
<code>type</code>	string		<code>assistant</code>	
<code>author</code>	Author		NUCLIA	
<code>text</code>	string			
<code>content</code>	object			
<code>tool_calls</code>	array<MessageToolCall>			

AssumeRoleType

Enum: 0, 1

Audio

Data about a previous audio response from the model. [\[Learn more\]](#)((see docs))

Field	Type	Req	Default	Description
<code>id</code>	string	yes		

AugmentedContext

Field	Type	Req	Default	Description
<code>paragraphs</code>	map<string, AugmentedTextBlock>		<code>{}</code>	Paragraphs added to the context as a result of using the ragstrategies parameter, typically the neighbouringparagraphs or the conversation strategies
<code>fields</code>	map<string, AugmentedTextBlock>		<code>{}</code>	Field extracted texts added to the context as a result of using the ragstrategies parameter, typically the hierarchy or fullresource strategies.

AugmentedTextBlock

Field	Type	Req	Default	Description
<code>id</code>	string	yes		The id of the augmented text bloc. It can be a paragraph id or a field id.
<code>text</code>	string	yes		The text of the augmented text block. It may include additional metadata

				to enrich the context
<code>position</code>	TextPosition null			Metadata about the position of the text block in the original document.
<code>parent</code>	string null			The parent text block that was augmented for.
<code>augmentation_type</code>	TextBlockAugmentationType	yes		Type of augmentation.

Author

Enum: `NUCLIA`, `USER`

AzureAIKey

Field	Type	Req	Default	Description
<code>key</code>	string			
<code>url</code>	string			
<code>model</code>	string			

AzureMistralKey

Field	Type	Req	Default	Description
<code>key</code>	string			
<code>url</code>	string			

AzureMistralUserPrompt

Field	Type	Req	Default	Description
<code>prompt</code>	string			
<code>system</code>	string			

AzureOpenAIKey

Field	Type	Req	Default	Description
key	string			
url	string			
deployment	string			
model	string			

AzureUserPrompt

Field	Type	Req	Default	Description
system	string			
prompt	string			

BaseConfigType: `object`**BaseModel**Type: `object`**BedrockAssumeRoleFinishRequest**

Field	Type	Req	Default	Description
role_arn	string	yes		ARN of the role created in AWS IAM

BedrockAssumeRoleStartResponse

Field	Type	Req	Default	Description
external_id	string	yes		External ID for assuming role
aws_account_id	string	yes		AWS Account ID to assume role into
role_name	string	yes		Name of the role to be created in AWS IAM

Body_import_retrieval_agent_api_v1_agent__agent_id__import_post

Field	Type	Req	Default	Description
<code>file</code>	string	yes		
<code>passphrase</code>	string	yes		Passphrase for the encrypted export file
<code>overwrite</code>	boolean		<code>False</code>	Whether to overwrite existing agent configuration. If false and the agent configuration is not empty, the import will fail.

ChatCompletionAllowedToolChoiceParam

Constrains the tools available to the model to a pre-defined set.

Field	Type	Req	Default	Description
<code>allowed_tools</code>	ChatCompletionAllowedToolsParam	yes		Constrains the tools available to the model to a pre-defined set.
<code>type</code>	string	yes		

ChatCompletionAllowedToolsParam

Constrains the tools available to the model to a pre-defined set.

Field	Type	Req	Default	Description
<code>mode</code>	enum[auto, required]	yes		
<code>tools</code>	array<object (free-form map)>	yes		

ChatCompletionAssistantMessageParam

Messages sent by the model in response to user messages.

Field	Type	Req	Default	Description
<code>role</code>	string	yes		
<code>audio</code>	Audio null			
<code>content</code>	string array<ChatCompletionContentPartTextParam ChatCompletionContentPartRefusalParam> null			
<code>function_call</code>	FunctionCall null			
<code>name</code>	string			
<code>refusal</code>	string null			
<code>tool_calls</code>	array<ChatCompletionMessageFunctionToolCallParam ChatCompletionMessageCustomToolCallParam>			

ChatCompletionContentPartImageParam

Learn about [image inputs]((see docs))

Field	Type	Req	Default	Description
<code>image_url</code>	ImageURL	yes		
<code>type</code>	string	yes		

ChatCompletionContentPartInputAudioParam

Learn about [audio inputs]((see docs))

Field	Type	Req	Default	Description
<code>input_audio</code>	InputAudio	yes		
<code>type</code>	string	yes		

ChatCompletionContentPartRefusalParam

Field	Type	Req	Default	Description
<code>refusal</code>	string	yes		
<code>type</code>	string	yes		

ChatCompletionContentPartTextParam

Learn about [text inputs](see docs)

Field	Type	Req	Default	Description
<code>text</code>	string	yes		
<code>type</code>	string	yes		

ChatCompletionDeveloperMessageParam

Developer-provided instructions that the model should follow, regardless of messages sent by the user. With o1 models and newer, `developer` messages replace the previous `system` messages.

Field	Type	Req	Default	Description
<code>content</code>	string array<ChatCompletionContentPartTextParam>	yes		
<code>role</code>	string	yes		
<code>name</code>	string			

ChatCompletionFunctionMessageParam

Field	Type	Req	Default	Description
<code>content</code>	string null	yes		
<code>name</code>	string	yes		
<code>role</code>	string	yes		

ChatCompletionFunctionToolParam

A function tool that can be used to generate a response.

Field	Type	Req	Default	Description
<code>function</code>	FunctionDefinition	yes		
<code>type</code>	string	yes		

ChatCompletionMessageCustomToolCallParam

A call to a custom tool created by the model.

Field	Type	Req	Default	Description
<code>id</code>	string	yes		
<code>custom</code>	openai__types__chat__chat_completion_message_custom_tool_call_param__Custom	yes		The custom tool that the model called.
<code>type</code>	string	yes		

ChatCompletionMessageFunctionToolCallParam

A call to a function tool created by the model.

Field	Type	Req	Default	Description
<code>id</code>	string	yes		
<code>function</code>	openai__types__chat__chat_completion_message_function_tool_call_param__Function	yes		The function that the model called.
<code>type</code>	string	yes		

ChatCompletionNamedToolChoiceCustomParam

Specifies a tool the model should use. Use to force the model to call a specific custom tool.

Field	Type	Req	Default	Description
-------	------	-----	---------	-------------

<code>custom</code>	<code>openai__types__chat__chat_completion_named_tool_choice_custom_param__Custom</code>	yes		
<code>type</code>	string	yes		

ChatCompletionNamedToolChoiceParam

Specifies a tool the model should use. Use to force the model to call a specific function.

Field	Type	Req	Default	Description
<code>function</code>	<code>openai__types__chat__chat_completion_named_tool_choice_param__Function</code>	yes		
<code>type</code>	string	yes		

ChatCompletionSystemMessageParam

Developer-provided instructions that the model should follow, regardless of messages sent by the user. With o1 models and newer, use `developer` messages for this purpose instead.

Field	Type	Req	Default	Description
<code>content</code>	string array<ChatCompletionContentPartTextParam>	yes		
<code>role</code>	string	yes		
<code>name</code>	string			

ChatCompletionToolMessageParam

Field	Type	Req	Default	Description
<code>content</code>	string array<ChatCompletionContentPartTextParam>	yes		

<code>role</code>	string	yes		
<code>tool_call_id</code>	string	yes		

ChatCompletionUserMessageParam

Messages sent by an end user, containing prompts or additional context information.

Field	Type	Req	Default	Description
<code>content</code>	string array<ChatCompletionContentPartTextParam ChatCompletionContentPartImageParam ChatCompletionContentPartInputAudioParam File>	yes		
<code>role</code>	string	yes		
<code>name</code>	string			

ChatContextMessage

Field	Type	Req	Default	Description
<code>author</code>	Author	yes		
<code>text</code>	string	yes		

ChatModel

Field	Type	Req	Default	Description
<code>question</code>	string	yes		
<code>retrieval</code>	boolean		True	
<code>user_id</code>	string		system	
<code>system</code>	string null			
<code>chat_history</code>	array<AssistantMessage ToolMessage Message>			

context	array<Message>		[]	
query_context	array<string> map<string, string>		{}	
query_context_order	map<string, integer>		{}	
truncate	boolean null		True	
user_prompt	UserPrompt null			
citations	boolean CitationsType null			Whether to include citations in the response. If set to None or False, no citations will be computed. If set to True or 'default', citations will be computed after answer generation and send as a sepa...
citation_threshold	number null			If citations is set to True or 'default', this will be the similarity threshold. Value between 0 and 1, lower values will produce more citations. If not set, it will be set to the optimized threshold ...
generative_model	string null			
max_tokens	integer null			
query_context_images	array<Image> map<string, Image>		{}	
prefer_markdown	boolean null			
json_schema	object (free-form map) null			
format_prompt	boolean		True	
rerank_context	boolean		False	Whether to reorder

<code>t</code>				the query context based on a reranker. This option will also make it so the first response will contain the scores given for each context piece.
<code>tools</code>	<code>array<Tool></code>			List of tools to choose
<code>tool_choice</code>	<code>ToolChoiceAuto ToolChoiceNone ToolChoiceRequired ToolChoiceForced</code>		<code>{'type': 'required'}</code>	Tool choice strategy. auto: The model decides whether to use a tool or not based on the prompt and available tools. required (default): A tool must be used. none: Disables tool usage even if tools are ...
<code>seed</code>	<code>integer null</code>			Seed use for the generative model for a deterministic output.
<code>reasoning</code>	<code>Reasoning boolean</code>		<code>False</code>	Reasoning options for the generative model. Set to True to enable default reasoning, False to disable, or provide a Reasoning object for custom options.
<code>image_generation</code>	<code>boolean</code>		<code>False</code>	Whether to enable image generation in the response.

ChatOptions

Enum: `keyword`, `relations`, `semantic`

CitationsType

Enum: `none`, `default`, `llm_footnotes`

Classification

Field	Type	Req	Default	Description
labelset	string	yes		
label	string	yes		

ClassificationLabel

Field	Type	Req	Default	Description
label	string	yes		
labelset	string	yes		

Claude3UserPrompt

Field	Type	Req	Default	Description
system	string			
prompt	string			

CloudLink

Field	Type	Req	Default	Description
uri	string null			
size	integer null			
content_type	string null			
filename	string null			
md5	string null			

ComputedMetadata

The purpose of this field is to show a cherry-picked set of fields from computed metadata without having to load the whole computed metadata field.

Field	Type	Req	Default	Description
-------	------	-----	---------	-------------

<code>field_classifications</code>	<code>array<FieldClassification></code>		<code>[]</code>	
------------------------------------	---	--	-----------------	--

ConfigSchema

Field	Type	Req	Default	Description
<code>semantic_model</code>	<code>ConfigSchemaOptions null</code>			
<code>semantic_models</code>	<code>ConfigSchemaOptions null</code>			
<code>default_semantic_model</code>	<code>ConfigSchemaOptions null</code>			
<code>anonymization_model</code>	<code>ConfigSchemaOptions</code>	yes		
<code>visual_labeling</code>	<code>ConfigSchemaOptions null</code>			
<code>generative_model</code>	<code>ConfigSchemaOptions null</code>			
<code>ner_model</code>	<code>ConfigSchemaOptions null</code>			
<code>relation_model</code>	<code>ConfigSchemaOptions null</code>			
<code>summary_model</code>	<code>ConfigSchemaOptions null</code>			
<code>summary</code>	<code>ConfigSchemaOptions null</code>			
<code>user_keys</code>	<code>ConfigSchemaElements null</code>			
<code>user_prompts</code>	<code>ConfigSchemaElements null</code>			
<code>summary_prompt</code>	<code>ConfigSchemaElements null</code>			
<code>prefer_markdown_generative_response</code>	<code>ConfigSchemaElements null</code>			
<code>semantic_graph_node_models</code>	<code>ConfigSchemaOptions null</code>			

default_semantic_graph_node_model	ConfigSchemaOptions null			
semantic_graph_edge_models	ConfigSchemaOptions null			
default_semantic_graph_edge_model	ConfigSchemaOptions null			

ConfigSchemaElements

Field	Type	Req	Default	Description
schemas	object (free-form map)	yes		
create	boolean		False	
update	boolean		False	

ConfigSchemaOptions

Field	Type	Req	Default	Description
options	array<GenerativeOption> array<Option>	yes		
default	string null			
create	boolean		False	
update	boolean		False	
multiple	boolean		False	

Consumption

Field	Type	Req	Default	Description
normalized_tokens	TokensDetail	yes		
customer_key_tokens	TokensDetail	yes		

ConsumptionResponse

Field	Type	Req	Default	Description
normalized_tokens	TokensDetail	yes		
customer_key_tokens	TokensDetail	yes		

Contains

Computes whether a value exists inside a range

Field	Type	Req	Default	Description
schema_id	string	yes		
key	string	yes		
contains	integer number	yes		

ConversationalStrategy

Field	Type	Req	Default	Description
name	string		conversation	
attachments_text	boolean		False	Add attachments on context retrieved on conversation
attachments_images	boolean		False	Add attachments images on context retrieved on conversation if they are mime type image and using a visual LLM
full	boolean		False	Add all conversation fields on matched blocks
max_messages	integer		15	Max messages to append in case its not full field

ConversationFieldData

Field	Type	Req	Default	Description
value	FieldConversation null			
extracted	ConversationFieldExtractedData null			
error	Error null			
status	string null			
errors	array<Error> null			

ConversationFieldExtractedData

Field	Type	Req	Default	Description
text	ExtractedText null			
metadata	FieldComputedMetadata null			
large_metadata	LargeComputedMetadata null			
vectors	VectorObject null			
question_answers	FieldQuestionAnswers null			
relation_node_vectors	map<string, array<RelationNodeVector>> null			
relation_edge_vectors	map<string, array<RelationEdgeVector>> null			

CreateAccountCustomModel

Field	Type	Req	Default	Description
model_types	array<StringModelTypes>	yes		
location	string	yes		

description	string null			
trained_date	string		2026-07-22T10:44:48.440834	
trained_kbid	string			
openai_compat	OpenAICompatMode null			

CreateDefaultModelConfig

Field	Type	Req	Default	Description
default_model_id	string	yes		Identifier of the public model to add to the account. e.g: 'chatgpt-azure-4-turbo' or 'claude-3'
description	string null			Description for the model config
user_keys	UserLearningKeys null			Custom user keys for the model. If not set, Nuclia keys will be used.
user_prompts	UserPrompts null			Custom user prompts for the model. If not set, Nuclia prompts will be used.
assume_role	AssumeRoleType null			Assume role type for models that require it (e.g., AWS Bedrock)
kbid	array<string>			List of KB IDs where this default model config is allowed. Empty list means it is allowed in all KBs (default).

CustomAccountModel

This is the pydantic model for the model details in the account.

Field	Type	Req	Default	Description
-------	------	-----	---------	-------------

<code>model_id</code>	string	yes		
<code>account</code>	string	yes		
<code>model_type</code>	StringModelTypes null	yes		
<code>trained_date</code>	string null	yes		
<code>location</code>	string null	yes		
<code>trained_kbid</code>	string null	yes		
<code>model_types</code>	array<StringModelTypes>		<code>[]</code>	
<code>log</code>	string null	yes		
<code>kbid</code>	array<string>	yes		
<code>openai_compat</code>	OpenAICompatMode Output null			

CustomAccountModelListItem

This is the pydantic model for the listing of models in the account.

Field	Type	Req	Default	Description
<code>model_id</code>	string	yes		
<code>account</code>	string	yes		
<code>model_type</code>	StringModelTypes null	yes		
<code>trained_date</code>	string null	yes		
<code>location</code>	string null	yes		
<code>trained_kbid</code>	string null	yes		
<code>model_types</code>	array<StringModelTypes>		<code>[]</code>	
<code>title</code>	string null	yes		

CustomPrompt

Field	Type	Req	Default	Description
-------	------	-----	---------	-------------

<code>system</code>	string null			System prompt given to the generative model responsible of generating the answer. This can help customize the behavior of the model when generating the answer. If not specified, the default model prov...
<code>user</code>	string null			User prompt given to the generative model responsible of generating the answer. Use the words {context} and {question} in brackets where you want those fields to be placed, in case you want them in yo...
<code>rephrase</code>	string null			Rephrase prompt given to the generative model responsible for rephrasing the query for a more effective retrieval step. This is only used if the rephrase flag is set to true in the request. If not spe...

CustomSplitStrategy

Enum: `NONE`, `MANUAL`, `LLM`

DataAugmentation-Input

Field	Type	Req	Default	Description
<code>name</code>	string			

on	ApplyTo		0	Defines if the task should be applied to paragraphs (0) or whole fields (1)
filter	Filter null			Filter to apply the data augmentation
operations	array<Operation-Input>			
llm	LLMConfig-Input			
filter_expression_json	string null			JSON representation of the filter expression, used to filter the data augmentation in the UI. If set, it takes precedence over the filter field

DataAugmentation-Output

Field	Type	Req	Default	Description
name	string			
on	ApplyTo		0	Defines if the task should be applied to paragraphs (0) or whole fields (1)
filter	Filter null			Filter to apply the data augmentation
operations	array<Operation-Output>			
llm	LLMConfig-Output			
filter_expression_json	string null			JSON representation of the filter expression, used to filter the data augmentation in the UI. If set, it takes precedence over the filter field

DataResidencyStatusEnum: `unknown`, `guaranteed`, `not_guaranteed`**Dataset**

Field	Type	Req	Default	Description
<code>id</code>	string	yes		
<code>timestamp</code>	string	yes		
<code>account_id</code>	string	yes		
<code>nua_client_id</code>	string	yes		
<code>account_type</code>	string	yes		
<code>type</code>	string	yes		
<code>path</code>	string	yes		
<code>filter</code>	object (free-form map)	yes		
<code>name</code>	string null			

DatasetCreated

Field	Type	Req	Default	Description
<code>id</code>	string	yes		

DatasetFilter

Field	Type	Req	Default	Description
<code>labels</code>	array<string>	yes		

DatasetInfo

Field	Type	Req	Default	Description
<code>name</code>	string	yes		
<code>filter</code>	DatasetFilter	yes		
<code>type</code>	Task-Input	yes		

DatasetsList

Field	Type	Req	Default	Description
<code>datasets</code>	array<Dataset>	yes		

DateCreated

Matches all fields created in a date range

Field	Type	Req	Default	Description
<code>prop</code>	string		<code>created</code>	
<code>since</code>	string null			Start of the date range. Leave blank for unbounded
<code>until</code>	string null			End of the date range. Leave blank for unbounded

DateModified

Matches all fields modified in a date range

Field	Type	Req	Default	Description
<code>prop</code>	string		<code>modified</code>	
<code>since</code>	string null			Start of the date range. Leave blank for unbounded
<code>until</code>	string null			End of the date range. Leave blank for unbounded

DeepSeekUserPrompt

Field	Type	Req	Default	Description
<code>system</code>	string			
<code>prompt</code>	string			

DefaultModelConfig

Field	Type	Req	Default	Description
<code>id</code>	string	yes		Unique identifier of the default model configuration (UUID format). e.g: '3fa85f64-5717-4562-b3fc-2c963f66afa6'
<code>default_model_id</code>	string	yes		Identifier of the default model to add to the account. e.g: 'chatgpt-azure-4-turbo' or 'claude-3'
<code>description</code>	string null			Description for the model config
<code>user_keys</code>	UserLearningKeys null			Custom user keys for the model. If not set, Nuclia keys will be used.
<code>user_prompts</code>	UserPrompts null			Custom user prompts for the model. If not set, Nuclia prompts will be used.
<code>assume_role</code>	AssumeRoleType null			Assume role type for models that require it (e.g., AWS Bedrock)
<code>kbids</code>	array<string>			List of KB IDs where this default model config is allowed. Empty list means it is allowed in all KBs.

DefaultModelConfigListItem

Field	Type	Req	Default	Description
<code>id</code>	string	yes		Unique identifier of the default model configuration (UUID format). e.g:

				'3fa85f64-5717-4562-b3fc-2c963f66afa6'
default_model_id	string	yes		Identifier of the default model to add to the account. e.g: 'chatgpt-azure-4-turbo' or 'claude-3'
description	string null			Description for the model config
assume_role	AssumeRoleType null			Assume role type for models that require it (e.g., AWS Bedrock)
kbid	array<string>			List of KB IDs where this default model config is allowed. Empty list means it is allowed in all KBs.

DestinationNode

Field	Type	Req	Default	Description
prop	string		destination_node	
value	string null			
match	NodeMatchKindName		exact	
type	RelationNodeType null		entity	
group	string null			

DirectionalRelation

Field	Type	Req	Default	Description
entity	string	yes		
entity_type	RelationNodeType	yes		
entity_subtype	string	yes		

<code>relation</code>	RelationType	yes		
<code>relation_label</code>	string	yes		
<code>direction</code>	RelationDirection	yes		
<code>metadata</code>	RelationMetadata null			
<code>resource_id</code>	string	yes		

DownloadStatus

Field	Type	Req	Default	Description
<code>id</code>	string	yes		
<code>type</code>	string	yes		
<code>status</code>	enum[pending, ready]	yes		
<code>download_url</code>	string null	yes		
<code>query</code>	object (free-form map)	yes		

DriverConfig

Field	Type	Req	Default	Description
<code>id</code>	string null			
<code>identifier</code>	string	yes		
<code>name</code>	string	yes		
<code>provider</code>	object	yes		The type of driver, e.g., 'google', 'marklogic', etc.
<code>config</code>	EncryptedPayload	yes		The configuration specific to the driver.

DriverID

Field	Type	Req	Default	Description
<code>id</code>	string	yes		

EdgeEmbeddingModel**Enum:** `multilingual-graph-v1`**Effort****Enum:** `0, 1, 2, 3, 4, 5`**Effort-Input****Enum:** `0, 1, 2, 3, 4, 5`**EmbeddingModel**

Field	Type	Req	Default	Description
<code>object</code>	string		<code>embedding</code>	
<code>index</code>	integer		<code>0</code>	
<code>embedding</code>	array<number>	yes		

EmbeddingsRequestModel

Field	Type	Req	Default	Description
<code>input</code>	array<string> string	yes		
<code>model</code>	string null			

EmbeddingsResponseModel

Field	Type	Req	Default	Description
<code>object</code>	string		<code>list</code>	
<code>data</code>	array<EmbeddingModel>	yes		
<code>model</code>	string	yes		
<code>usage</code>	EmbeddingUsageData	yes		

EmbeddingUsageData

Field	Type	Req	Default	Description
-------	------	-----	---------	-------------

<code>prompt_tokens</code>	integer	yes		
<code>total_tokens</code>	integer	yes		

EncryptedPayload

Type: `object`

Entity-Input

Matches fields that contains a detected entity

Field	Type	Req	Default	Description
<code>prop</code>	string		<code>entity</code>	
<code>subtype</code>	string	yes		Type of the entity. e.g: PERSON
<code>value</code>	string null			Value of the entity. e.g: Anna. If blank, matches any entity of the given type

Entity-Output

Field	Type	Req	Default	Description
<code>token</code>	string null			
<code>root</code>	string null			
<code>type</code>	string null			

EntityDefinition

Field	Type	Req	Default	Description
<code>label</code>	string	yes		Entity type
<code>description</code>	string null			Description of the entity type

EntityExample

Field	Type	Req	Default	Description
<code>name</code>	string	yes		Name associated with the entity
<code>label</code>	string	yes		Type of entity

EntitySubgraph

Field	Type	Req	Default	Description
<code>related_to</code>	array<DirectionalRelation>	yes		

Eq

Equal (==) operator

Field	Type	Req	Default	Description
<code>schema_id</code>	string	yes		
<code>key</code>	string	yes		
<code>eq</code>	boolean integer number string	yes		

Error

Field	Type	Req	Default	Description
<code>body</code>	string	yes		
<code>code</code>	integer	yes		
<code>code_str</code>	string	yes		
<code>created</code>	string null	yes		
<code>severity</code>	string	yes		

Extra

Field	Type	Req	Default	Description
-------	------	-----	---------	-------------

<code>metadata</code>	object (free-form map)	yes		Arbitrary JSON metadata provided by the user that is not meant to be searchable, but can be serialized on results.
-----------------------	------------------------	-----	--	--

ExtractConfig

Field	Type	Req	Default	Description
<code>name</code>	string	yes		
<code>vllm_config</code>	VLLMExtractionConfig null			
<code>ai_tables</code>	AITablesConfig null			

ExtractConfig-Input

Hey, developer! Keep this in sync with corresponding pydantic model in `learning_config.models`

Field	Type	Req	Default	Description
<code>name</code>	string			
<code>vllm_config</code>	VLLMExtractionConfig-Input null			
<code>ai_tables</code>	AITables-Input null			
<code>split</code>	SplitConfig null			
<code>max_parallel_llm_calls</code>	integer		0	

ExtractConfig-Output

Hey, developer! Keep this in sync with corresponding pydantic model in `learning_config.models`

Field	Type	Req	Default	Description
<code>name</code>	string			
<code>vllm_config</code>	VLLMExtractionConfig-Output null			
<code>ai_tables</code>	AITables-Output			

	null			
<code>split</code>	SplitConfig null			
<code>max_parallel_llm_calls</code>	integer		0	

ExtractedDataTypeName

Enum: `text`, `metadata`, `shortened_metadata`, `large_metadata`, `vectors`, `link`, `file`, `question_answers`, `relation_vectors`

ExtractedText

Field	Type	Req	Default	Description
<code>text</code>	string null			
<code>split_text</code>	map<string, string> null			
<code>deleted_splits</code>	array<string> null			

ExtractOperation

Field	Type	Req	Default	Description
<code>model</code>	Model		0	
<code>triggers</code>	array<Trigger>			

Features

Field	Type	Req	Default	Description
<code>reasoning</code>	ReasoningConfig null			
<code>tools</code>	BaseConfig null			
<code>structured_output</code>	StructuredOutputConfig null			
<code>vision</code>	BaseConfig null			
<code>streaming</code>	BaseConfig null			

<code>system_message</code>	BaseConfig null			
<code>text_extraction</code>	BaseConfig null			
<code>custom_keys</code>	BaseConfig null			
<code>file_input</code>	BaseConfig null			
<code>image_generation</code>	BaseConfig null			

Field

Matches a field or set of fields

Field	Type	Req	Default	Description
<code>prop</code>	string		<code>field</code>	
<code>type</code>	FieldTypeName	yes		Type of the field to match,
<code>name</code>	string null			Name of the field to match. If blank, matches all fields of the given type

FieldClassification

Field	Type	Req	Default	Description
<code>field</code>	FieldID	yes		
<code>classifications</code>	array<Classification>		<code>[]</code>	

FieldComputedMetadata

Field	Type	Req	Default	Description
<code>metadata</code>	FieldMetadata	yes		
<code>split_metadata</code>	map<string, FieldMetadata> null			
<code>deleted_split</code>	array<string> null			

s				
---	--	--	--	--

FieldConversation

This is a metadata representation of a conversation about how many pages of messages and total of messages we have. This class is used mainly when exposing a conversation in the resource level

Field	Type	Req	Default	Description
pages	integer null			
size	integer null			
total	integer null			
extract_strategy	string null			
split_strategy	string null			

FieldEntities

Wrapper for the entities extracted from a field (required because protobuf doesn't support lists of lists)

Field	Type	Req	Default	Description
entities	array<FieldEntity>	yes		

FieldEntity

Field	Type	Req	Default	Description
text	string	yes		
label	string	yes		
positions	array<Position>	yes		

FieldExtensionStrategy

Field	Type	Req	Default	Description
name	string		field_extension	

<code>fields</code>	<code>array<string></code>		<code>[]</code>	List of field ids to extend the context with. It will try to extend the retrieval context with the specified fields in the matching resources. The field ids have to be in the format {fieldtype}/{field...}
<code>data_augmentation_field_prefixes</code>	<code>array<string></code>		<code>[]</code>	List of prefixes for data augmentation added fields to extend the context with. For example, if the prefix is 'simpson', all fields that are a result of data augmentation with that prefix will be used...

FieldFile

Field	Type	Req	Default	Description
<code>added</code>	<code>string null</code>			
<code>file</code>	<code>CloudLink null</code>			
<code>language</code>	<code>string null</code>			
<code>password</code>	<code>string null</code>			
<code>external</code>	<code>boolean</code>		<code>False</code>	
<code>extract_strategy</code>	<code>string null</code>			Id of the Nuclia extract strategy used at processing time. If not set, the default strategy was used. Extract strategies are defined at the learning configuration api.
<code>split_strategy</code>	<code>string null</code>			Id of the Nuclia split strategy used at processing time. If not set, the default

				strategy was used. Split strategies are defined at the learning configuration api.
--	--	--	--	--

FieldID

Field	Type	Req	Default	Description
<code>field_type</code>	FieldType	yes		
<code>field</code>	string	yes		

FieldLargeMetadata

Field	Type	Req	Default	Description
<code>entities</code>	array<Entity-Output> null			
<code>tokens</code>	map<string, integer> null			

FieldLink

Field	Type	Req	Default	Description
<code>added</code>	string null			
<code>headers</code>	map<string, string> null			
<code>cookies</code>	map<string, string> null			
<code>uri</code>	string null			
<code>language</code>	string null			
<code>localStorage</code>	map<string, string> null			
<code>css_selector</code>	string null			
<code>xpath</code>	string null			
<code>extract_strat</code>	string null			Id of the Nuclia

egy				extract strategy used at processing time. If not set, the default strategy was used. Extract strategies are defined at the learning configuration api.
split_strategy	string null			Id of the Nuclia split strategy used at processing time. If not set, the default strategy was used. Split strategies are defined at the learning configuration api.

FieldMetadata

Field	Type	Req	Default	Description
links	array<string>	yes		
paragraphs	array<Paragraph>	yes		
ner	map<string, string>	yes		
entities	map<string, FieldEntities>	yes		
classifications	array<Classification>	yes		
last_index	string null			
last_understanding	string null			
last_extract	string null			
last_summary	string null			
last_processing_start	string null			
thumbnail	CloudLink null			
language	string null			

summary	string null			
positions	map<string, Positions>	yes		
relations	array<Relation-Output> null			
mime_type	string null			

FieldMimetype

Matches fields with a mimetype

Field	Type	Req	Default	Description
prop	string		field_mimetype	
type	string	yes		Type of the mimetype to match. e.g: In image/jpeg, type is image
subtype	string null			Type of the mimetype to match. e.g: In image/jpeg, subtype is jpeg. Leave blank to match all mimetype of the type

FieldQuestionAnswers

Field	Type	Req	Default	Description
question_answers	QuestionAnswers	yes		
split_question_answers	map<string, QuestionAnswers> null			
deleted_splits	array<string> null			

FieldText

Field	Type	Req	Default	Description
-------	------	-----	---------	-------------

body	string null			
format	TextFormat null			
md5	string null			
extract_strategy	string null			Id of the Nuclia extract strategy used at processing time. If not set, the default strategy was used. Extract strategies are defined at the learning configuration api.
split_strategy	string null			Id of the Nuclia split strategy used at processing time. If not set, the default strategy was used. Split strategies are defined at the learning configuration api.

FieldType

Enum: file, link, text, generic, conversation

FieldTypeName

This map assumes that both values and extracted data field containers use the same names for its fields. See `models.ResourceFieldValues` and `models.ResourceFieldExtractedData`

Enum: text, file, link, conversation, generic, key_value

File

Learn about [file inputs]((see docs) for text generation.

Field	Type	Req	Default	Description
file	FileFile	yes		
type	string	yes		

FileExtractedData

Field	Type	Req	Default	Description
language	string null			
md5	string null			
metadata	map<string, string> null			
nested	map<string, string> null			
file_generated	map<string, CloudLink> null			
file_rows_previews	map<string, RowsPreview> null			
file_preview	CloudLink null			
file_pages_previews	FilePages null			
file_thumbnail	CloudLink null			
field	string null			
icon	string null			
nested_position	map<string, NestedPosition> null			
nested_list_position	map<string, NestedListPosition> null			

FileFieldData

Field	Type	Req	Default	Description
value	FieldFile null			
extracted	FileFieldExtractedData null			
error	Error null			
status	string null			

<code>errors</code>	<code>array<Error> null</code>			
---------------------	--	--	--	--

FileFieldExtractedData

Field	Type	Req	Default	Description
<code>text</code>	<code>ExtractedText null</code>			
<code>metadata</code>	<code>FieldComputedMetadata null</code>			
<code>large_metadata</code>	<code>LargeComputedMetadata null</code>			
<code>vectors</code>	<code>VectorObject null</code>			
<code>question_answers</code>	<code>FieldQuestionAnswers null</code>			
<code>relation_node_vectors</code>	<code>map<string, array<RelationNodeVector>> null</code>			
<code>relation_edge_vectors</code>	<code>map<string, array<RelationEdgeVector>> null</code>			
<code>file</code>	<code>FileExtractedData null</code>			

FileFile

Field	Type	Req	Default	Description
<code>file_data</code>	<code>string</code>			
<code>file_id</code>	<code>string</code>			
<code>filename</code>	<code>string</code>			

FilePages

Field	Type	Req	Default	Description
<code>pages</code>	<code>array<CloudLink> null</code>			
<code>positions</code>	<code>array<PagePositions></code>			

	null			
structures	array<PageStructure> null			

Filter

Field	Type	Req	Default	Description
contains	array<string>			Text that must be contained in the field in order to apply the data augmentation, if multiple values are provided, they will be combined with the specified operator
resource_type	array<string>			Data Augmentation will only be applied to resources of the specified types, if multiple values are provided, they will be combined with the specified operator
field_types	array<string>			Data Augmentation will only be applied to fields of the specified types
not_field_types	array<string>			Data Augmentation will not be applied to fields of the specified types
rids	array<string>			Data Augmentation will only be applied resources matching the specified resource ids
fields	array<string>			Data Augmentation will only be applied to fields matching the specified field ids
splits	array<string>			Data Augmentation will only be applied

				to fields matching the specified field splits
labels	array<string>			Data Augmentation will only be applied to fields matching the specified labels, if multiple values are provided, they will be combined with the specified operator
apply_to_agent_generated_fields	boolean		False	If enabled, the data augmentation will also be applied to fields generated by other data augmentation agents. This functionality is only supported for agents that do not generate fields.
contains_operator	FilterLogicalOperator		0	Way of combining the values in the contains field
labels_operator	FilterLogicalOperator		0	Way of combining the values in the labels field

FilterExpression

Returns only documents that match this filter expression. Filtering examples can be found here: (see docs)

This allows building complex filtering expressions and replaces the following parameters: `fields`, `filters`, `range_*`, `resource_filters`, `keyword_filters`.

Field	Type	Req	Default	Description
field	And_FieldFilterExpressionType_ Or_FieldFilterExpressionType_ Not_FieldFilterExpressionType_ Resource-Input Field Keyword DateCreated DateModified Label			Filter to apply to fields

	ResourceMimetype FieldMimetype Entity-Input Language OriginTag OriginMetadata OriginPath OriginSource OriginCollaborator Generated null			
paragraph	And_ParagraphFilter Expression_ Or_ParagraphFilterE xpression_ Not_ParagraphFilterE xpression_ Label Kind null			Filter to apply to each text block
operator	Operator		and	How to combine field and paragraph filters (default is AND).AND returns text blocks that match both filters.OR returns textblocks that match one of the two filters

FilterLogicalOperator

Enum: 0, 1

FindField

Field	Type	Req	Default	Description
paragraphs	map<string, FindParagraph>	yes		

FindOptions

Enum: keyword, semantic, relations, graph

FindParagraph

Field	Type	Req	Default	Description
score	number	yes		

score_type	SCORE_TYPE	yes		
order	integer		0	
text	string	yes		
id	string	yes		
labels	array<string> null		[]	
position	TextPosition null			
fuzzy_result	boolean		False	
page_with_visual	boolean		False	This flag informs if the page may have information that has not been extracted
reference	string null			Reference to the extracted image that represents this paragraph
is_a_table	boolean		False	The referenced image of the paragraph is a table
relevant_relations	Relations null			Relevant relations from which the paragraph was found, will only be filled if using the Graph RAG Strategy

FindRequest

Field	Type	Req	Default	Description
audit_metadata	map<string, string> null			A dictionary containing optional audit-specific metadata, such as userid, environment, or other contextual information. This metadata can be leveraged for filtering and analyzing activity logs in futu...

query	string			The query to search for
filter_expression	FilterExpression null			Returns only documents that match this filter expression. Filtering examples can be found here: (see docs) This allows building complex filtering expressions and replaces the following parameters: field...
fields	array<string>		[]	The list of fields to search in. For instance: a/title to search only on title field. For more details on filtering by field, see: (see docs)
filters	array<string> array<Filter>		[]	The list of filters to apply. Filtering examples can be found here: (see docs)
top_k	integer		20	The number of results search should return. The maximum number of results allowed is 200.
min_score	number MinScore null			Minimum score to filter search results. Results with a lower score will be ignored. Accepts either a float or a dictionary with the minimum scores for the bm25 and vector indexes. If a float is provid...
range_creation_start	string null			Resources created before this date will be filtered out of search results.

				Datetime are represented as a str in ISO 8601 format, like: 2008-09-15T15:53:00+05:00.
<code>range_creation_end</code>	string null			Resources created after this date will be filtered out of search results. Datetime are represented as a str in ISO 8601 format, like: 2008-09-15T15:53:00+05:00.
<code>range_modification_start</code>	string null			Resources modified before this date will be filtered out of search results. Datetime are represented as a str in ISO 8601 format, like: 2008-09-15T15:53:00+05:00.
<code>range_modification_end</code>	string null			Resources modified after this date will be filtered out of search results. Datetime are represented as a str in ISO 8601 format, like: 2008-09-15T15:53:00+05:00.
<code>debug</code>	boolean		<code>False</code>	If set, the response will include some extra metadata for debugging purposes, like the list of queried nodes.
<code>highlight</code>	boolean		<code>False</code>	If set to true, the query terms will be highlighted in the results between <code><mark>...</mark></code> tags
<code>show</code>	array<ResourceProperties>		<code>['basic']</code>	Controls which types of metadata are serialized on

				resources of search results
<code>field_type_filter</code>	<code>array<FieldTypeNam e></code>		<code>['text', 'file', 'link', 'conversation', 'generic', 'key_value']</code>	Define which field types are serialized on resources of search results
<code>extracted</code>	<code>array<ExtractedData TypeName></code>		<code>[]</code>	[Deprecated] Please use GET resource endpoint instead to get extracted metadata
<code>vector</code>	<code>array<number> null</code>			The vector to perform the search with. If not provided, NucliaDB will use Nuclia Predict API to create the vector off from the query.
<code>vectorset</code>	<code>string null</code>			Vectors index to perform the search in. If not provided, NucliaDB will use the default one
<code>with_duplicates</code>	<code>boolean</code>		<code>False</code>	Whether to return duplicate paragraphs on the same document
<code>with_synonyms</code>	<code>boolean</code>		<code>False</code>	Whether to return matches for custom knowledge box synonyms of the query terms. Note: only supported for keyword and fulltext search options.
<code>resource_filters</code>	<code>array<string></code>		<code>[]</code>	List of resource ids to filter search results for. Only paragraphs from the specified resources will be returned.

security	RequestSecurity null			Security metadata for the request. Please refer to the documentation for more details on how security works: (see docs)
show_hidden	boolean		False	If set to false (default), excludes hidden resources from search
rephrase	boolean		False	Rephrase the query for a more efficient retrieval. This will consume LLM tokens and make the request slower.
rephrase_prompt	string null			Rephrase prompt given to the generative model responsible for rephrasing the query for a more effective retrieval step. This is only used if the rephrase flag is set to true in the request. If not spe...
query_image	Image null			Image that will be used together with the query text for retrieval.
graph_query	And_GraphPathQuery_ Or_GraphPathQuery_ Not_GraphPathQuery_ GraphPath SourceNode DestinationNode AnyNode Relation-Input nucliadb_models__graph__requests__Generated null			Query for the knowledge graph. Paths (node-relation-node) extracted from a paragraphid will be used to extend the results

features	array<FindOptions>		['keyword', 'semantic']	List of search features to use. Each value corresponds to a lookup into one of the different indexes
rank_fusion	RankFusionName ReciprocalRankFusion		rrf	Rank fusion algorithm to use to merge results from multiple retrievers (keyword, semantic)
reranker	RerankerName PredictReranker		predict	Reranker let you specify which method you want to use to rerank your results at the end of retrieval
keyword_filters	array<string> array<Filter>		[]	List of keyword filter expressions to apply to the retrieval step. The text block search will only be performed on the documents that contain the specified keywords. The filters are case-insensitive, ...
search_configuration	string null			Load find parameters from this configuration. Parameters in the request override parameters from the configuration.
generative_model	string null			The generative model used to rephrase the query. If not provided, the model configured for the Knowledge Box is used.

FindResource

Field	Type	Req	Default	Description
-------	------	-----	---------	-------------

id	string	yes		
slug	string null			
title	string null			
summary	string null			
icon	string null			
thumbnail	string null			
metadata	Metadata null			
usermetadata	UserMetadata null			
fieldmetadata	array<UserFieldMetadata> null			
computedmetadata	ComputedMetadata null			
created	string null			
modified	string null			
last_seqid	integer null			
last_account_seq	integer null			
queue	QueueType null			
hidden	boolean null			
origin	Origin null			
extra	Extra null			
relations	array<Relation-Output> null			
data	ResourceData null			
security	ResourceSecurity null			Resource security metadata
fields	map<string, FindField>	yes		

Format

Enum: 0, 1, 2, 3, 4, 5, 6, 7

FullResourceApplyTo

Field	Type	Req	Default	Description
<code>exclude</code>	array<string>			Resources from matches containing any of these labels won't expand to the full resource. This may be useful to exclude long and not interesting resources and expend less tokens

FullResourceStrategy

Field	Type	Req	Default	Description
<code>name</code>	string		<code>full_resource</code>	
<code>count</code>	integer null			Maximum number of full documents to retrieve. If not specified, all matching documents are retrieved.
<code>include_remaining_text_blocks</code>	boolean		<code>False</code>	Whether to include the remaining text blocks after the maximum number of resources has been reached.
<code>apply_to</code>	FullResourceApplyTo null			Define which resources to exclude from serialization

FunctionCall

Deprecated and replaced by `tool_calls`. The name and arguments of a function that should be called, as generated by the model.

Field	Type	Req	Default	Description
<code>arguments</code>	string	yes		

name	string	yes		
------	--------	-----	--	--

FunctionDefinition

Field	Type	Req	Default	Description
name	string	yes		
description	string			
parameters	object (free-form map)			
strict	boolean null			

Generated

Matches if the field was generated by the given source

Field	Type	Req	Default	Description
prop	string		generated	
by	string	yes		Generator for this field. Currently, only data-augmentation is supported
da_task	string null			Matches field generated by an specific DA task, given its prefix

GenerationConfig

Stores parameters involved in generation that may vary from request to request

Field	Type	Req	Default	Description
temperature	number		0.0	Temperature to use when sampling
default_max_completion_tokens	integer		800	Maximum number of output tokens to request by default
max_input_tok	integer		64000	Maximum number of input tokens

<code>ens</code>				supported by the model.
<code>max_supported_completion_tokens</code>	integer null			Maximum number of output tokens supported by the model, if not set, no check will be done that the supplied maxtokens is less than this value

GenerativeDrivers

Enum: `microsoft`, `google`, `nuclia`, `hf`, `mistral`, `claude3`, `microsoft_mistral`, `openai`, `huggingface`, `vertex`, `azure_aii`, `openai_compat`, `bedrock`, `none`, `custom`, `dummy`

GenerativeOption

Field	Type	Req	Default	Description
<code>name</code>	string	yes		
<code>description</code>	string null			
<code>value</code>	string	yes		
<code>provider</code>	string	yes		
<code>user_prompt</code>	string null			
<code>user_key</code>	string null			

Generator

Enum: `data-augmentation`, `processor`, `user`

GenericFieldData

Field	Type	Req	Default	Description
<code>value</code>	string null			
<code>extracted</code>	TextFieldExtractedData null			
<code>error</code>	Error null			
<code>status</code>	string null			

<code>errors</code>	<code>array<Error> null</code>			
---------------------	--	--	--	--

GraphEdgeSearch

Field	Type	Req	Default	Description
<code>vectors</code>	<code>map<string, map<string, array<number>>>></code>	yes		Graph edge embeddings for each edge and semantic model. Structure: {model: {edge: vector}}
<code>timings</code>	<code>map<string, number></code>	yes		Time taken to compute the graph edge embedding for each semantic model
<code>consumption</code>	<code>ConsumptionResponse null</code>			Consumption details for the graph edge embeddings operation

GraphExtractionExample

Field	Type	Req	Default	Description
<code>entities</code>	<code>array<EntityExample></code>	yes		Examples of entities extracted from the example text
<code>relations</code>	<code>array<RelationExample></code>	yes		Examples of relations extracted from the example text, all entities must be included in the entities list
<code>text</code>	<code>string</code>	yes		Example text where entities and relations were extracted

GraphNode

Field	Type	Req	Default	Description
-------	------	-----	---------	-------------

value	string null			
match	NodeMatchKindName		exact	
type	RelationNodeType null		entity	
group	string null			

GraphNodeSearch

Field	Type	Req	Default	Description
vectors	map<string, map<string, array<number>>>	yes		Graph node embeddings for each node and semantic model. Structure: {model: {node: vector}}
timings	map<string, number>	yes		Time taken to compute the graph node embedding for each semantic model
consumption	ConsumptionResponse null			Consumption details for the graph node embeddings operation

GraphOperation-Input

Field	Type	Req	Default	Description
entity_defs	array<EntityDefinition>			Types of entities that will be considered for extraction
examples	array<GraphExtractionExample>			Examples to be sent to the model as few-shot learning
ident	string	yes		
triggers	array<Trigger>			

GraphOperation-Output

Field	Type	Req	Default	Description
<code>entity_defs</code>	array<EntityDefinition>			Types of entities that will be considered for extraction
<code>examples</code>	array<GraphExtractionExample>			Examples to be sent to the model as few-shot learning
<code>ident</code>	string	yes		
<code>triggers</code>	array<Trigger>			

GraphPath

Field	Type	Req	Default	Description
<code>prop</code>	string		<code>path</code>	
<code>source</code>	GraphNode null			
<code>relation</code>	GraphRelation null			
<code>destination</code>	GraphNode null			
<code>undirected</code>	boolean		<code>False</code>	

GraphRelation

Field	Type	Req	Default	Description
<code>label</code>	string null			
<code>type</code>	RelationType null			

GraphStrategy

This strategy retrieves context pieces by exploring the Knowledge Graph, starting from the entities present in the query. It works best if the Knowledge Box has a user-defined Graph Extraction agent enabled.

Field	Type	Req	Default	Description
<code>name</code>	string		<code>graph_beta</code>	
<code>hops</code>	integer		<code>3</code>	Number of hops to

				take when exploring the graph for relevant context. For example, - hops=1 will explore the neighbors of the starting entities. - hops=2 will explore the neighbors of the neighbors of...
top_k	integer			Number of relationships to keep after each hop after ranking them by relevance to the query. This number correlates to more paragraphs being sent as context. If not set, this number will be set to 30 ...
agentic_graph_only	boolean		True	If set to true, only relationships extracted from a graph extraction agent are considered for context expansion.
relation_text_as_paragraphs	boolean		False	If set to true, the text of the relationships is to create context paragraphs, this enables to use bigger top K values without running into the generative model's context limits. If set to false, the ...
relation_ranking	RelationRanking		reranker	Method to rank relationships. - reranker uses the reranker model to rank relationships. - generative uses first the reranker to first lower the amount of relationships and then

				the generative model to...
<code>query_entity_detection</code>	QueryEntityDetection		<code>predict</code>	Method to detect entities in the query. - predict uses NUA to detect entities in the query, slower and more accurate but requires an exact text match between Knowledge Box entities and entities in the...
<code>weight</code>	number		<code>3.0</code>	Weight of the graph strategy in the context. The weight is used to scale the results of the strategy before adding them to the context. The weight should be a positive number.

GuardOperation

Field	Type	Req	Default	Description
<code>enabled</code>	boolean		<code>False</code>	
<code>triggers</code>	array<Trigger>			

HFEembeddingKey

Field	Type	Req	Default	Description
<code>url</code>	string	yes		
<code>key</code>	string	yes		
<code>matryoshka</code>	array<integer>		<code>[]</code>	
<code>similarity</code>	string	yes		
<code>size</code>	integer	yes		
<code>threshold</code>	number	yes		

passage_prompt	string	yes		
query_prompt	string	yes		

HFLLMKey

Field	Type	Req	Default	Description
key	string			
url	string			
model	ModelType		0	

HFUserPrompt

Field	Type	Req	Default	Description
system	string			
prompt	string	yes		

HierarchyResourceStrategy

Field	Type	Req	Default	Description
name	string		hierarchy	
count	integer		0	Number of extra characters that are added to each matching paragraph when adding to the context.

HistoryQuestionAnswer

Field	Type	Req	Default	Description
question	string	yes		
answer	string	yes		

HTTPValidationError

Field	Type	Req	Default	Description
<code>detail</code>	array<ValidationError>			

Image

Field	Type	Req	Default	Description
<code>content_type</code>	string	yes		
<code>b64encoded</code>	string	yes		

ImageURL

Field	Type	Req	Default	Description
<code>url</code>	string	yes		
<code>detail</code>	enum[auto, low, high]			

Inequalities

Inequality operators that can be grouped to perform range queries.

Field	Type	Req	Default	Description
<code>schema_id</code>	string	yes		
<code>key</code>	string	yes		
<code>gte</code>	integer number string null			
<code>lte</code>	integer number string null			

InputAudio

Field	Type	Req	Default	Description
<code>data</code>	string	yes		

<code>format</code>	<code>enum[wav, mp3]</code>	<code>yes</code>		
---------------------	-----------------------------	------------------	--	--

InteractionOperation

Enum: `0`, `1`

InteractionRequest

Field	Type	Req	Default	Description
<code>question</code>	<code>string</code>	<code>yes</code>		
<code>headers</code>	<code>map<string, string></code>		<code>{}</code>	
<code>arguments</code>	<code>map<string, string></code>		<code>{}</code>	
<code>chat_history</code>	<code>array<HistoryQuestionAnswer> null</code>			Client-managed chat history. When set (even to an empty list), overrides any server-side session history for agents that use previous Q&A context (rephrase, summarize, smart, etc.). Omit the field ent...
<code>operation</code>	<code>InteractionOperation</code>		<code>0</code>	
<code>streaming</code>	<code>boolean</code>		<code>False</code>	

InteractionsAuditDownloadRequest

Field	Type	Req	Default	Description
<code>session_id</code>	<code>string null</code>			Filter by session ID
<code>year</code>	<code>integer null</code>			Filter by year (e.g., 2024). If not specified, defaults to the current year.
<code>month</code>	<code>integer null</code>			Filter by month (1-12). If not specified, defaults to the past month.

JobStatus

Enum: `not_running`, `finished`, `running`, `started`, `stopped`, `failed`, `terminating`, `starting`

JSONSchema

Structured Outputs configuration options, including a JSON Schema.

Field	Type	Req	Default	Description
<code>name</code>	string	yes		
<code>description</code>	string			
<code>schema</code>	object (free-form map)			
<code>strict</code>	boolean null			

KeyValueField

A key-value field value. The field id (key in the resource's `key_values` dict) must equal the schema name — enforcing one KV field per schema per resource.

Field	Type	Req	Default	Description
<code>data</code>	map<string, string integer number boolean Range>			Key-value pairs conforming to the schema.

KeyValueFieldData

Field	Type	Req	Default	Description
<code>value</code>	KeyValueField null			
<code>error</code>	Error null			
<code>status</code>	string null			
<code>errors</code>	array<Error> null			

Keyword

Matches all fields that contain a keyword

Field	Type	Req	Default	Description
-------	------	-----	---------	-------------

<code>prop</code>	string		<code>keyword</code>	
<code>word</code>	string	yes		Keyword to find

Kind

Matches paragraphs of a certain kind

Field	Type	Req	Default	Description
<code>prop</code>	string		<code>kind</code>	
<code>kind</code>	TypeParagraph	yes		The kind of paragraph to match

KnowledgeboxFindResults

Find on knowledgebox results

Field	Type	Req	Default	Description
<code>resources</code>	map<string, FindResource>	yes		
<code>relations</code>	Relations null			
<code>query</code>	string null			
<code>rephrased_query</code>	string null			
<code>total</code>	integer		<code>0</code>	
<code>page_number</code>	integer		<code>0</code>	Pagination will be deprecated, please, refer to topk in the request
<code>page_size</code>	integer		<code>20</code>	Pagination will be deprecated, please, refer to topk in the request
<code>next_page</code>	boolean		<code>False</code>	Pagination will be deprecated, please, refer to topk in the request
<code>nodes</code>	array<map<string, string>> null			List of nodes queried in the search

<code>shards</code>	<code>array<string> null</code>			The list of shard replica ids used for the search.
<code>autofilters</code>	<code>array<string></code>		<code>[]</code>	[deprecated] list of filters automatically applied to the search query
<code>min_score</code>	<code>number MinScore null</code>		<code>{'bm25': 0.0}</code>	The minimum scores that have been used for the search operation.
<code>best_matches</code>	<code>array<string></code>		<code>[]</code>	List of ids of best matching paragraphs. The list is sorted by decreasing relevance (most relevant first).
<code>metrics</code>	<code>object (free-form map) null</code>			Metrics information about the search operation. The metadata included in this field is subject to change and should not be used in production. This is only available if the debug parameter is set to t...

Label

Field	Type	Req	Default	Description
<code>label</code>	<code>string</code>	yes		
<code>description</code>	<code>string</code>			
<code>examples</code>	<code>array<string></code>			

LabelOperation

Field	Type	Req	Default	Description
<code>labels</code>	<code>array<Label></code>			

<code>ident</code>	string	yes		
<code>description</code>	string			
<code>multiple</code>	boolean		False	
<code>triggers</code>	array<Trigger>			

Language

Matches the language of the field

Field	Type	Req	Default	Description
<code>prop</code>	string		<code>language</code>	
<code>only_primary</code>	boolean		False	Match only the primary language of the document. By default, matches any language that appears in the document
<code>language</code>	string	yes		The code of the language to match, e.g: en

LargeComputedMetadata

Field	Type	Req	Default	Description
<code>metadata</code>	FieldLargeMetadata null			
<code>split_metadata</code>	map<string, FieldLargeMetadata> null			
<code>deleted_splits</code>	array<string> null			

LayoutFormat

Enum: 0

learning_config__models__AnthropicKey

Field	Type	Req	Default	Description
-------	------	-----	---------	-------------

key	string	yes		

learning_config_models__AzureOpenAIKey

Field	Type	Req	Default	Description
key	string	yes		
url	string	yes		
deployment	string	yes		
model	string	yes		

learning_config_models__HFEmbeddingKey

Field	Type	Req	Default	Description
url	string	yes		URL of the Huggingface embedding model endpoint
key	string			API Key with access the desired Huggingface endpoint
matryoshka	array<integer> null			List of matryoshka dimensions to use for the embedding, in descending order. If unknown, leave it empty
similarity	string		DOT	Similarity function to use. It can be DOT or COSINE, leave it as DOT if unknown
size	integer	yes		Size of the embeddings generated by the model. Must be set to the correct value for the specific model

				hosted in the endpoint
<code>threshold</code>	number		<code>0.4</code>	Threshold to use for semantic similarity
<code>passage_prompt</code>	string		<code>{}</code>	Prompt to use for passage embedding, only set this value if your embedding model requires it
<code>query_prompt</code>	string		<code>{}</code>	Prompt to use for query embedding, only set this value if your embedding model requires it

`learning_config__models__MistralKey`

Field	Type	Req	Default	Description
<code>key</code>	string	yes		

`learning_config__models__OpenAIKey`

Field	Type	Req	Default	Description
<code>key</code>	string	yes		
<code>org</code>	string	yes		

`learning_config__models__PalmKey`

Field	Type	Req	Default	Description
<code>credentials</code>	string null			If using Vertex AI, this should contain the Service Account Key credentials
<code>location</code>	string null			If using Vertex AI, this is the GCP region where the model will be accessed from. E.g. us-central1

<code>gemini_key</code>	string null			If using Gemini API, only an API key for the Gemini Developer API is required
-------------------------	---------------	--	--	---

`learning_models__providers__models__Effort`

Enum: `none`, `minimal`, `low`, `medium`, `high`, `xhigh`

`learning_protos__config_p2p__AnthropicVertexKey`

Field	Type	Req	Default	Description
<code>credentials</code>	string	yes		Service Account Key credentials in JSON format for Vertex AI access
<code>location</code>	string			GCP region where the model will be accessed from. E.g. us-central1

`learning_protos__config_p2p__HFEmbeddingKey`

Some models require a specific template (including prefix) to work correctly in each task For example Snowflake's Arctic-embed requires a specific prefix to work correctly. In that case, the query prompt will be `` passage_prompt: "" query_prompt: "Represent this sentence for searching relevant pas...`

Field	Type	Req	Default	Description
<code>url</code>	string			
<code>key</code>	string			
<code>matryoshka</code>	array<integer>			
<code>similarity</code>	string			
<code>size</code>	integer		<code>0</code>	
<code>threshold</code>	number		<code>0.0</code>	
<code>passage_prompt</code>	string			
<code>query_prompt</code>	string			

learning_protos__config_p2p__HFLLMKey

Field	Type	Req	Default	Description
key	string			
url	string			
model	ModelType		0	

learning_protos__config_p2p__UserLearningKeys-Input

Field	Type	Req	Default	Description
openai	OpenAIKey null			
azure_openai	AzureOpenAIKey null			
palm	PalmKey null			
anthropic	AnthropicKey null			
claude3	AnthropicKey null			
anthropic_vertex	learning_protos__config_p2p__AnthropicVertexKey null			
anthropic_bedrock	AnthropicBedrockKey null			
text_generation	TextGenerationKey null			
mistral	MistralKey null			
azure_mistral	AzureMistralKey null			
hf_llm	learning_protos__config_p2p__HFLLMKey null			
hf_embedding	HFEmbeddingKey null			
azure_aii	AzureAIIKey null			
openai_compat	OpenAICompatModel-Input null			

learning_protos__config_p2p__UserLearningKeys-Output

Field	Type	Req	Default	Description
openai	OpenAIKey null			
azure_openai	AzureOpenAIKey null			
palm	PalmKey null			
anthropic	AnthropicKey null			
claude3	AnthropicKey null			
anthropic_vertex	learning_protos__config_p2p__AnthropicVertexKey null			
anthropic_bedrock	AnthropicBedrockKey null			
text_generation	TextGenerationKey null			
mistral	MistralKey null			
azure_mistral	AzureMistralKey null			
hf_llm	learning_protos__config_p2p__HFLLMKey null			
hf_embedding	HFEmbeddingKey null			
azure_ait	AzureAITKey null			
openai_compat	OpenAICompatModel-Output null			

LearningConfig

Field	Type	Req	Default	Description
semantic_models	array<string> null			
generative_model	string null			
ner_model	string null			

<code>anonymization_model</code>	<code>string null</code>			
<code>visual_labeling</code>	<code>string null</code>			
<code>split_configs</code>	<code>map<string, SplitConfiguration> null</code>			
<code>extract_configs</code>	<code>map<string, ExtractConfig> null</code>			

LearningConfiguration

This model is to add the fields that are computed dynamically and not stored in the db, it also serves as a bridge between the stored config and what we want to send through the APIs

Field	Type	Req	Default	Description
<code>resource_labelers_models</code>	<code>array<string> null</code>			
<code>paragraph_labelers_models</code>	<code>array<string> null</code>			
<code>intent_models</code>	<code>array<string> null</code>			
<code>visual_labeling</code>	<code>string null</code>		<code>disabled</code>	
<code>ner_model</code>	<code>string</code>		<code>multilingual</code>	
<code>relation_model</code>	<code>string null</code>		<code>base</code>	
<code>anonymization_model</code>	<code>string</code>		<code>disabled</code>	
<code>semantic_model</code>	<code>string null</code>			
<code>semantic_models</code>	<code>array<string> null</code>			
<code>default_semantic_model</code>	<code>string null</code>			
<code>semantic_graph_node_models</code>	<code>array<string> null</code>			
<code>default_seman</code>	<code>string null</code>			

tic_graph_node_model				
semantic_graph_edge_models	array<string> null			
default_semantic_graph_edge_model	string null			
semantic_vector_similarity	string null		DOT	
semantic_vector_size	integer null		1024	
semantic_matryoshka_dims	array<integer> null		[]	
semantic_threshold	number null		0.4	
generative_model	string		chatgpt-azure-4o	
user_keys	UserLearningKeys null			
user_prompts	UserPrompts null			
summary	Summary		simple	
summary_model	string		chatgpt-azure-4o	
summary_prompt	SummaryPrompt null			
prefer_markdown_generative_response	boolean		False	
semantic_model_configs	map<string, SemanticConfig>		{}	
semantic_graph_node_model_configs	map<string, SemanticConfig>		{}	
semantic_graph_edge_model_configs	map<string, SemanticConfig>		{}	

LearningConfigurationCreation

Field	Type	Req	Default	Description
<code>anonymization_model</code>	AnonymizationModel string		<code>disabled</code>	
<code>visual_labeling</code>	VisualLabeling string		<code>disabled</code>	
<code>generative_model</code>	string		<code>chatgpt-azure-4o</code>	
<code>ner_model</code>	NERModel string		<code>multilingual</code>	
<code>relation_model</code>	RelationModel string		<code>base</code>	
<code>user_keys</code>	UserLearningKeys null			
<code>user_prompts</code>	UserPrompts null			
<code>summary</code>	Summary null		<code>simple</code>	
<code>summary_model</code>	string		<code>chatgpt-azure-4o</code>	
<code>summary_prompt</code>	SummaryPrompt null			
<code>prefer_markdown_generative_response</code>	boolean		<code>False</code>	
<code>default_semantic_model</code>	string null			The default vector set / semantic model to use. Must be one of the models defined in the semanticmodels field. If not provided, the first model in the list will be used.
<code>semantic_model</code>	ParagraphEmbeddingModel null			
<code>semantic_models</code>	array<ParagraphEmbeddingModel> null			
<code>semantic_graph_node_models</code>	array<NodeEmbeddingModel> null			
<code>default_seman</code>	string null			

tic_graph_node_model				
semantic_graph_edge_models	array<EdgeEmbeddingModel> null			
default_semantic_graph_edge_model	string null			

LearningConfigurationUpdate

Field	Type	Req	Default	Description
anonymization_model	AnonymizationModel string		disabled	
visual_labeling	VisualLabeling string		disabled	
generative_model	string		chatgpt-azure-4o	
ner_model	NERModel string		multilingual	
relation_model	RelationModel string		base	
user_keys	UserLearningKeys null			
user_prompts	UserPrompts null			
summary	Summary null		simple	
summary_model	string		chatgpt-azure-4o	
summary_prompt	SummaryPrompt null			
prefer_markdown_generative_response	boolean		False	
default_semantic_model	string null			The default vector set / semantic model to use. Must be one of the models defined in the semanticmodels field. If not provided,

				the first model in the list will be used.
--	--	--	--	---

LinkExtractedData

Field	Type	Req	Default	Description
date	string null			
language	string null			
title	string null			
metadata	map<string, string> null			
link_thumbnail	CloudLink null			
link_preview	CloudLink null			
field	string null			
link_image	CloudLink null			
description	string null			
type	string null			
embed	string null			
file_generated	map<string, CloudLink> null			

LinkFieldData

Field	Type	Req	Default	Description
value	FieldLink null			
extracted	LinkFieldExtractedData null			
error	Error null			
status	string null			
errors	array<Error> null			

LinkFieldExtractedData

Field	Type	Req	Default	Description
text	ExtractedText null			
metadata	FieldComputedMetadata null			
large_metadata	LargeComputedMetadata null			
vectors	VectorObject null			
question_answers	FieldQuestionAnswers null			
relation_node_vectors	map<string, array<RelationNodeVector>> null			
relation_edge_vectors	map<string, array<RelationEdgeVector>> null			
link	LinkExtractedData null			

LinkUpload

Field	Type	Req	Default	Description
link	string	yes		
headers	map<string, string>	yes		
cookies	map<string, string>	yes		
localstorage	map<string, string>	yes		
css_selector	string null			
xpath	string null			
extract_strategy	string null			
split_strategy	string null			
classification_labels	array<ClassificationLabel> null			

ListModelsResponse

Field	Type	Req	Default	Description
object	string		list	
data	array<ModelInfo>	yes		
has_more	boolean		False	

LLMConfig-Input

Field	Type	Req	Default	Description
model	string			
provider	string			
keys	UserLearningKeys-Input null			
prompts	UserPrompts null			
reasoning_config	LLMReasoningConfig null			

LLMConfig-Output

Field	Type	Req	Default	Description
model	string			
provider	string			
keys	UserLearningKeys-Output null			
prompts	UserPrompts null			
reasoning_config	LLMReasoningConfig null			

LLMReasoningConfig

Keep this in sync with learning_generative.models.Reasoning

Field	Type	Req	Default	Description
-------	------	-----	---------	-------------

<code>effort</code>	Effort null		<code>0</code>	
<code>budget_tokens</code>	integer null		<code>0</code>	

LLMReasoningConfig-Input

Keep this in sync with `learning_generative.models.Reasoning`

Field	Type	Req	Default	Description
<code>effort</code>	Effort-Input null		<code>0</code>	
<code>budget_tokens</code>	integer null		<code>0</code>	

LLMReasoningConfig-Output

Keep this in sync with `learning_generative.models.Reasoning`

Field	Type	Req	Default	Description
<code>effort</code>	Effort null		<code>0</code>	
<code>budget_tokens</code>	integer null		<code>0</code>	

LLMSplitConfig

Field	Type	Req	Default	Description
<code>generative_model</code>	string null			
<code>generative_provider</code>	string null			
<code>rules</code>	array<string>		<code>[]</code>	

LLMSplitConfig-Input

Field	Type	Req	Default	Description
<code>llm</code>	LLMConfig-Input null			
<code>rules</code>	array<string>			

LLMSplitConfig-Output

Field	Type	Req	Default	Description
llm	LLMConfig-Output null			
rules	array<string>			

ManualSplitConfig

Field	Type	Req	Default	Description
splitter	string null			

MaxTokens

Field	Type	Req	Default	Description
context	integer null			Use to limit the amount of tokens used in the LLM context
answer	integer null			Use to limit the amount of tokens used in the LLM answer

MemoryOperation-Input

Field	Type	Req	Default	Description
ident	string	yes		
prompt	string			
rules	array<string>			
triggers	array<Trigger>			
graph_extraction	boolean		False	When true, entities and relations are extracted from each fact using the same LLM. entitydefs and examples are only

				valid when this is true.
<code>entity_defs</code>	<code>array<EntityDefinition></code>			Restrict graph extraction to these entity types. Only valid when <code>graphextraction</code> is true.
<code>examples</code>	<code>array<GraphExtractionExample></code>			Few-shot examples for graph extraction. Only valid when <code>graphextraction</code> is true.

MemoryOperation-Output

Field	Type	Req	Default	Description
<code>ident</code>	string	yes		
<code>prompt</code>	string			
<code>rules</code>	<code>array<string></code>			
<code>triggers</code>	<code>array<Trigger></code>			
<code>graph_extraction</code>	boolean		<code>False</code>	When true, entities and relations are extracted from each fact using the same LLM. <code>entitydefs</code> and <code>examples</code> are only valid when this is true.
<code>entity_defs</code>	<code>array<EntityDefinition></code>			Restrict graph extraction to these entity types. Only valid when <code>graphextraction</code> is true.
<code>examples</code>	<code>array<GraphExtractionExample></code>			Few-shot examples for graph extraction. Only valid when <code>graphextraction</code> is true.

Message

Field	Type	Req	Default	Description
type	string		message	
author	Author	yes		
text	string	yes		

MessageToolCall

Field	Type	Req	Default	Description
id	string	yes		
type	string		function	
function	MessageToolFunction	yes		

MessageToolFunction

Field	Type	Req	Default	Description
name	string	yes		
arguments	object			

Metadata

Field	Type	Req	Default	Description
deployed_region	string null			
deprecated	boolean		False	
shut_down_date	string null			
superseded_by	string null			
stability	Stability		stable	
knowledge_cutoff_date	string null			

<code>release_date</code>	string null			
<code>data_residency_status</code>	DataResidencyStatus		<code>not_guaranteed</code>	

MetadataExtensionStrategy

RAG strategy to enrich the context with metadata of the matching paragraphs or its resources. This strategy can be combined with any of the other strategies.

Field	Type	Req	Default	Description
<code>name</code>	string		<code>metadata_extension</code>	
<code>types</code>	array<MetadataExtensionType>	yes		List of resource metadata types to add to the context. - 'origin': origin metadata of the resource. - 'classificationlabels': classification labels of the resource. - 'ner': Named Entity Recognition e...

MetadataExtensionType

Enum: `origin`, `classification_labels`, `ners`, `extra_metadata`

MinScore

Field	Type	Req	Default	Description
<code>semantic</code>	number null			Minimum semantic similarity score used to filter vector index search. If not specified, the default minimum score of the semantic model associated to the Knowledge Box will be used. Check out the docu...

bm25	number		0	Minimum score used to filter bm25 index search. Check out the documentation for more information on how to use this parameter: (see docs)
------	--------	--	---	---

MistralKey

Field	Type	Req	Default	Description
key	string			

MistralUserPrompt

Field	Type	Req	Default	Description
prompt	string			

Model

Enum: 0

ModelConfig

Field	Type	Req	Default	Description
input_tokens	TokenValues	yes		
output_tokens	OutputTokenValues	yes		
prompt_id	string	yes		
driver	GenerativeDrivers	yes		
max_images	integer null			
assume_role	integer null			

ModelFeatures

Stores the supported functionality of a model or API

Field	Type	Req	Default	Description
-------	------	-----	---------	-------------

logprobs	boolean		False	Whether the model supports outputting logprobs
streaming	boolean		True	Whether the model supports streaming responses
system_message	boolean		True	Whether the model supports system messages
tool_use	boolean		True	Whether the model supports tooluse/function calling
vision	boolean		True	Whether the model supports image understanding (vision)
file_input	boolean		False	Whether the API and model support file inputs
max_images	integer		5	Maximum number of images that the model supports in a single request
seed_support	boolean		False	Whether the model supports seeded generation, for deterministic outputs. If not enabled, using a seed will result in an error.
reasoning	ReasoningType		0	Type of reasoning supported by the model. - Unspecified: Reasoning will be automatically detected, but won't be customizable. - None: No reasoning capabilities are available. - Customizable Effort: Re...

<code>reasoning_chat_template_key</code>	string			The key used in the models chat template to enable or disable reasoning. Only used for models with Reasoning Chat Template reasoning type. Defaults to enablethinking if empty.
--	--------	--	--	---

ModelID

Field	Type	Req	Default	Description
<code>id</code>	string	yes		

ModelInfo

Field	Type	Req	Default	Description
<code>id</code>	string	yes		
<code>name</code>	string	yes		
<code>created</code>	integer	yes		
<code>title</code>	string null			
<code>description</code>	string null			
<code>model_type</code>	string null			
<code>max_input_tokens</code>	integer null			
<code>max_output_tokens</code>	integer null			

ModelType

Enum: 0, 1

NeighbouringParagraphsStrategy

Field	Type	Req	Default	Description
-------	------	-----	---------	-------------

<code>name</code>	string		<code>neighbouring_</code> <code>paragraphs</code>	
<code>before</code>	integer		2	Number of previous neighbouring paragraphs to add to the context, for each matching paragraph in the retrieval step.
<code>after</code>	integer		2	Number of following neighbouring paragraphs to add to the context, for each matching paragraph in the retrieval step.

Ner

Field	Type	Req	Default	Description
<code>text</code>	string	yes		
<code>ner</code>	string	yes		
<code>start</code>	integer	yes		
<code>end</code>	integer	yes		

NERModel

Enum: `multilingual`

NestedListPosition

Field	Type	Req	Default	Description
<code>positions</code>	array<NestedPosition >	yes		

NestedPosition

Field	Type	Req	Default	Description
<code>start</code>	integer null			
<code>end</code>	integer null			

<code>page</code>	integer null			
-------------------	----------------	--	--	--

NodeEmbeddingModel

Enum: `multilingual-graph-v1`

NodeMatchKindName

Enum: `exact`, `fuzzy`, `fuzzy_words`, `semantic`

Not_FieldFilterExpressionType_

Field	Type	Req	Default	Description
<code>operand</code>	And_FieldFilterExpressionType_ Or_FieldFilterExpressionType_ Not_FieldFilterExpressionType_ Resource-Input Field Keyword DateCreated DateModified Label ResourceMimetype FieldMimetype Entity-Input Language OriginTag OriginMetadata OriginPath OriginSource OriginCollaborator Generated	yes		

Not_GraphPathQuery_

Field	Type	Req	Default	Description
<code>operand</code>	And_GraphPathQuery_ Or_GraphPathQuery_ Not_GraphPathQuery_ GraphPath SourceNode DestinationNode AnyNode Relation-	yes		

	Input nucliadb_models__graph__requests__Generated			
--	--	--	--	--

Not_KVFilterExpression_

Field	Type	Req	Default	Description
<code>operand</code>	And_KVFilterExpression_ Or_KVFilterExpression_ Not_KVFilterExpression_ Eq Inequalities Contains	yes		

Not_ParagraphFilterExpression_

Field	Type	Req	Default	Description
<code>operand</code>	And_ParagraphFilterExpression_ Or_ParagraphFilterExpression_ Not_ParagraphFilterExpression_ Label Kind	yes		

NovaUserPrompt

Field	Type	Req	Default	Description
<code>system</code>	string			
<code>prompt</code>	string			

nucliadb_models__graph__requests__Generated

Matches if the relation was generated by the given source

Field	Type	Req	Default	Description
-------	------	-----	---------	-------------

<code>prop</code>	string		<code>generated</code>	
<code>by</code>	Generator	yes		Generator for this field.
<code>da_task</code>	string null			Matches relations generated by an specific DA task, given its prefix

NucliaDBClientType

Enum: `api`, `widget`, `web`, `dashboard`, `desktop`, `chrome_extension`

openai__types__chat__chat_completion_message_custom_tool_call_param__Custom

The custom tool that the model called.

Field	Type	Req	Default	Description
<code>input</code>	string	yes		
<code>name</code>	string	yes		

openai__types__chat__chat_completion_message_function_tool_call_param__Function

The function that the model called.

Field	Type	Req	Default	Description
<code>arguments</code>	string	yes		
<code>name</code>	string	yes		

openai__types__chat__chat_completion_named_tool_choice_custom_param__Custom

Field	Type	Req	Default	Description
<code>name</code>	string	yes		

openai__types__chat__chat_completion_named_tool_choice_param__Function

Field	Type	Req	Default	Description
<code>name</code>	string	yes		

OpenAIChatCompletionRequest

Field	Type	Req	Default	Description
<code>messages</code>	array<ChatCompletionDeveloperMessageParam ChatCompletionSystemMessageParam ChatCompletionUserMessageParam ChatCompletionAssistantMessageParam ChatCompletionToolMessageParam ChatCompletionFunctionMessageParam>	yes		
<code>model</code>	string null			
<code>stream</code>	boolean		<code>False</code>	
<code>max_tokens</code>	integer null			
<code>seed</code>	integer null			
<code>temperature</code>	number null			
<code>top_p</code>	number null			
<code>user</code>	string null			
<code>tools</code>	array<ChatCompletionFunctionToolParam> null			
<code>tool_choice</code>	enum[none, auto, required] ChatCompletionAllowedToolChoiceParam ChatCompletionNamedToolChoiceParam ChatCompletionNamedToolChoiceCustomParam null			
<code>response_format</code>	ResponseFormatText ResponseFormatJSON			

	NSchema ResponseFormatJSO NObject null			
json_schema	object (free-form map) null			
reasoning_effort	enum[none, minimal, low, medium, high, xhigh] null			
image_generation	boolean null			
truncate	boolean null			

OpenAICompatModel-Input

Field	Type	Req	Default	Description
key	string			The API key to use, can be left empty if not required
url	string	yes		The URL that the api is hosted on
model_id	string	yes		The model id to use as defined in the API
tokenizer	Tokenizer		0	The tokenizer to use for counting tokens, if you require a different tokenizer, please contact support
generation_config	GenerationConfig			The sampling parameters to use for the model
model_features	ModelFeatures			The model features that the model or API supports

OpenAICompatModel-Output

Field	Type	Req	Default	Description
-------	------	-----	---------	-------------

key	string			The API key to use, can be left empty if not required
url	string	yes		The URL that the api is hosted on
model_id	string	yes		The model id to use as defined in the API
tokenizer	Tokenizer		0	The tokenizer to use for counting tokens, if you require a different tokenizer, please contact support
generation_config	GenerationConfig			The sampling parameters to use for the model
model_features	ModelFeatures			The model features that the model or API supports

OpenAICompatUserPrompt

Field	Type	Req	Default	Description
system	string			
prompt	string	yes		

OpenAIKey

Field	Type	Req	Default	Description
key	string	yes		The OpenAI API key to use
org	string			The organization to use, can be left empty if not required

OpenAIUserPrompt

Field	Type	Req	Default	Description
-------	------	-----	---------	-------------

system	string			
prompt	string			

Operation-Input

XXX: Hey developer! If you modify this message, update the `OperationType`` enum above

Field	Type	Req	Default	Description
graph	GraphOperation-Input null			
label	LabelOperation null			
ask	AskOperation null			
qa	QAOperation null			
extract	ExtractOperation null			
prompt_guard	GuardOperation null			
llama_guard	GuardOperation null			
memory	MemoryOperation-Input null			

Operation-Output

XXX: Hey developer! If you modify this message, update the `OperationType`` enum above

Field	Type	Req	Default	Description
graph	GraphOperation-Output null			
label	LabelOperation null			
ask	AskOperation null			
qa	QAOperation null			
extract	ExtractOperation null			
prompt_guard	GuardOperation null			
llama_guard	GuardOperation null			

memory	MemoryOperation- Output null			
--------	-----------------------------------	--	--	--

Operator

Enum: `and`, `or`

Option

Field	Type	Req	Default	Description
name	string	yes		
value	string	yes		

Or_FieldFilterExpressionType_

Field	Type	Req	Default	Description
operands	array<And_FieldFilterExpressionType_ Or_FieldFilterExpressionType_ Not_FieldFilterExpressionType_ Resource-Input Field Keyword DateCreated DateModified Label ResourceMimetype FieldMimetype Entity-Input Language OriginTag OriginMetadata OriginPath OriginSource OriginCollaborator Generated>	yes		

Or_GraphPathQuery_

Field	Type	Req	Default	Description
operands	array<And_GraphPathQuery_	yes		

	Or_GraphPathQuery_ Not_GraphPathQuery_ _ GraphPath SourceNode DestinationNode AnyNode Relation- Input nucliadb_models__gr aph__requests__Gene rated>			
--	--	--	--	--

Or_KVFilterExpression_

Field	Type	Req	Default	Description
operands	array<And_KVFilter Expression_ Or_KVFilterExpressi on_ Not_KVFilterExpress ion_ Eq Inequalities Contains>	yes		

Or_ParagraphFilterExpression_

Field	Type	Req	Default	Description
operands	array<And_ParagraphFilterExpression_ Or_ParagraphFilterE xpression_ Not_ParagraphFilterE xpression_ Label Kind>	yes		

Origin

Field	Type	Req	Default	Description
source_id	string null			
url	string null			

created	string null			
modified	string null			
metadata	map<string, string>		{}	Generic metadata from the resource at the origin system. It can later be used for filtering on search endpoints with '/origin.metadata/{key}/{value}'
tags	array<string>		[]	Resource tags about the origin system. It can later be used for filtering on search endpoints with '/origin.tags/{tag}'
collaborators	array<string>		[]	
filename	string null			
related	array<string>		[]	
path	string null			Path of the original resource. Typically used to store folder structure information of the resource at the origin system. It can be later used for filtering on search endpoints with '/origin.path/{pat...'
sync_metadata	SyncMetadata null			Metadata related to the resource from the origin system fetched by the Progress Agentic RAG's Cloud Storage Sync service.
source	Source null		API	

OriginCollaborator

Matches the origin collaborators

Field	Type	Req	Default	Description
-------	------	-----	---------	-------------

<code>prop</code>	string		<code>origin_collaborator</code>	
<code>collaborator</code>	string	yes		Collaborator

OriginMetadata

Matches metadata from the origin

Field	Type	Req	Default	Description
<code>prop</code>	string		<code>origin_metadata</code>	
<code>field</code>	string	yes		Metadata field
<code>value</code>	string null			Value of the metadata field. If blank, matches any document with the given metadata field set (to any value)

OriginPath

Matches the origin path

Field	Type	Req	Default	Description
<code>prop</code>	string		<code>origin_path</code>	
<code>prefix</code>	string null			Prefix of the path, matches all paths under this prefix.g: prefix=/dir/ matches /dir and /dir/a/b but not /dirrrr

OriginSource

Matches the origin source id

Field	Type	Req	Default	Description
-------	------	-----	---------	-------------

<code>prop</code>	string		<code>origin_source</code>	
<code>id</code>	string null			Source ID

OriginTag

Matches all fields with a given origin tag

Field	Type	Req	Default	Description
<code>prop</code>	string		<code>origin_tag</code>	
<code>tag</code>	string	yes		The tag to match

OutputTokenValues

Field	Type	Req	Default	Description
<code>min</code>	integer		<code>0</code>	
<code>max</code>	integer	yes		
<code>default_max</code>	integer null			

PageImageStrategy

Field	Type	Req	Default	Description
<code>name</code>	string		<code>page_image</code>	
<code>count</code>	integer null			Maximum number of page images to retrieve. By default, at most 5 images are retrieved.

PageInformation

Field	Type	Req	Default	Description
<code>page</code>	integer null			
<code>page_with_visual</code>	boolean null			

PagePositions

Field	Type	Req	Default	Description
<code>start</code>	integer null			
<code>end</code>	integer null			

PageStructure

Field	Type	Req	Default	Description
<code>page</code>	PageStructurePage	yes		
<code>tokens</code>	array<PageStructureToken>	yes		

PageStructurePage

Field	Type	Req	Default	Description
<code>width</code>	integer	yes		
<code>height</code>	integer	yes		

PageStructureToken

Field	Type	Req	Default	Description
<code>x</code>	number	yes		
<code>y</code>	number	yes		
<code>width</code>	number	yes		
<code>height</code>	number	yes		
<code>text</code>	string	yes		
<code>line</code>	number	yes		

PalmKey

Field	Type	Req	Default	Description
<code>credentials</code>	string			If using Vertex AI,

				this should contain the Service Account Key credentials
location	string			If using Vertex AI, this is the GCP region where the model will be accessed from. E.g. us-central1
gemini_key	string			If using Gemini API, only an API key for the Gemini Developer API is required

PalmUserPrompt

Field	Type	Req	Default	Description
prompt	string			

Paragraph

Field	Type	Req	Default	Description
start	integer null			
end	integer null			
start_seconds	array<integer> null			
end_seconds	array<integer> null			
kind	TypeParagraph null			
classifications	array<Classification> null			
sentences	array<Sentence> null			
key	string null			
page	PageInformation null			
representation	Representation null			

<code>relations</code>	ParagraphRelations null			
------------------------	---------------------------	--	--	--

ParagraphAnnotation

Field	Type	Req	Default	Description
<code>classifications</code>	array<UserClassification>		<code>[]</code>	
<code>key</code>	string	yes		

ParagraphEmbeddingModel

Enum: `en-2024-04-24`, `multilingual-2024-05-06`, `multilingual-2023-08-16`, `multilingual-2024-10-07`, `text-embedding-3-small`, `text-embedding-3-large`, `gecko-embeddings-multi`, `gemini-embedding-2`, `hf_embedding`

ParagraphImageStrategy

Field	Type	Req	Default	Description
<code>name</code>	string		<code>paragraph_image</code>	

ParagraphRelations

Field	Type	Req	Default	Description
<code>parents</code>	array<string>		<code>[]</code>	
<code>siblings</code>	array<string>		<code>[]</code>	
<code>replacements</code>	array<string>		<code>[]</code>	

PatchAccountCustomModel

Field	Type	Req	Default	Description
<code>model_types</code>	array<StringModelTypes> null			
<code>location</code>	string null			
<code>description</code>	string null			

trained_date	string		2026-07-22T10:44:48.442327	
trained_kbid	string			
openai_compat	OpenAICompatModel-Input null			

PatchDefaultModelConfig

Field	Type	Req	Default	Description
description	string null			Description for the default model config
user_keys	UserLearningKeys null			Custom user keys for the default model. If not set, Nuclia keys will be used.
user_prompts	UserPrompts null			Custom user prompts for the default model. If not set, Nuclia prompts will be used.
assume_role	AssumeRoleType null			Assume role type for models that require it (e.g., AWS Bedrock)
kbid	array<string> null			List of KB IDs where this default model config is allowed. Empty list means it is allowed in all KBs. If not set, the previous value will be kept.

Position

Field	Type	Req	Default	Description
start	integer	yes		
end	integer	yes		

Positions

Field	Type	Req	Default	Description
<code>position</code>	array<Position>	yes		
<code>entity</code>	string	yes		

PredictOrigin

Enum: `PREDICT`, `RAO`

PredictReranker

Field	Type	Req	Default	Description
<code>name</code>	string		<code>predict</code>	
<code>window</code>	integer null			Number of elements reranker will use. Window must be greater or equal to topk. Greater values will improve results at cost of retrieval and reranking time. By default, this reranker uses a default of ...

PreQueriesStrategy

This strategy allows to run a set of queries before the main query and add the results to the context. It allows to give more importance to some queries over others by setting the weight of each query. The weight of the main query can also be set with the `main_query_weight` parameter.

Field	Type	Req	Default	Description
<code>name</code>	string		<code>prequeries</code>	
<code>queries</code>	array<PreQuery>	yes		List of queries to run before the main query. The results are added to the context with the specified weights for each query. There is a limit of 10 prequeries

				per request.
<code>main_query_weight</code>	number		<code>1.0</code>	Weight of the main query in the context. Use this to control the importance of the main query in the context.

PreQuery

Field	Type	Req	Default	Description
<code>request</code>	FindRequest	yes		The request to be executed before the main query.
<code>weight</code>	number		<code>1.0</code>	Weight of the prequery in the context. The weight is used to scale the results of the prequery before adding them to the context. The weight should be a positive number, and they are normalized so that...
<code>id</code>	string null			Identifier of the prequery. If not specified, it is autogenerated based on the index of the prequery in the list (prequery0, prequery1, ...).
<code>prefilter</code>	boolean		<code>False</code>	If set to true, the prequery results are used to filter the scope of the remaining queries. The resources of the most relevant paragraphs of the prefilter queries are used as resource filters for the ...

ProcessingBlock

Field	Type	Req	Default	Description
x	integer	yes		
y	integer	yes		
cols	integer	yes		
rows	integer	yes		
type	TypeBlock	yes		
ident	string	yes		
payload	string	yes		

ProcessingConversation

Field	Type	Req	Default	Description
messages	array<ProcessingMessage>	yes		
extract_strategy	string null			
split_strategy	string null			
classification_labels	array<ClassificationLabel> null			

ProcessingGeneratedConversation

Field	Type	Req	Default	Description
source_field_id	string	yes		
conversationfield	ProcessingConversation	yes		

ProcessingLayoutDiff

Field	Type	Req	Default	Description
<code>format</code>	LayoutFormat	yes		
<code>blocks</code>	map<string, ProcessingBlock>	yes		

ProcessingMessage

Field	Type	Req	Default	Description
<code>timestamp</code>	string null			
<code>who</code>	string null			
<code>to</code>	array<string>		<code>[]</code>	
<code>content</code>	ProcessingMessageContent	yes		
<code>ident</code>	string	yes		

ProcessingMessageContent

Field	Type	Req	Default	Description
<code>text</code>	string null			
<code>format</code>	Format	yes		

ProcessingQueueType

Enum: `shared`, `private`

ProcessRequestStatus

Field	Type	Req	Default	Description
<code>processing_id</code>	string	yes		
<code>resource_id</code>	string null	yes		
<code>kbid</code>	string null	yes		
<code>title</code>	string null	yes		

labels	array<string>	yes		
completed	boolean	yes		
scheduled	boolean	yes		
timestamp	string	yes		
completed_at	string null	yes		
scheduled_at	string null	yes		
failed	boolean		False	
retries	integer		0	
schedule_eta	number		0.0	
schedule_order	integer		0	
request	string null			
response	string null			

ProcessRequestStatusResults

Field	Type	Req	Default	Description
results	array<ProcessRequestStatus>	yes		
cursor	string null	yes		

PromptArgument

An argument for a prompt template.

Field	Type	Req	Default	Description
name	string	yes		
description	string null			
required	boolean null			

PromptConfig

Field	Type	Req	Default	Description
-------	------	-----	---------	-------------

<code>name</code>	string	yes		
<code>description</code>	string	yes		
<code>prompt</code>	string	yes		
<code>arguments</code>	array<PromptArgument> null			
<code>icons</code>	map<string, string> null			
<code>meta</code>	map<string, string> null			
<code>prompt_id</code>	string null			

PromptID

Field	Type	Req	Default	Description
<code>id</code>	string	yes		

Provider

Field	Type	Req	Default	Description
<code>title</code>	string	yes		
<code>enterprise_readiness</code>	boolean	yes		
<code>models</code>	map<string, ModelInfo>	yes		

ProviderName

Enum: `anthropic`, `anthropic-vertexai`, `anthropic-bedrock`, `microsoft`, `google`, `openai`, `openai_compat`, `huggingface`, `vertex`, `amazon`, `default`, `onprem`, `progress-agentic-rag`, `none`, `dummy`

ProvidersResponse

Field	Type	Req	Default	Description
<code>providers</code>	map<string, Provider>	yes		

PublicTask

Field	Type	Req	Default	Description
<code>name</code>	string	yes		
<code>data_augmentation</code>	boolean		False	
<code>description</code>	string null			
<code>can_cleanup</code>	boolean		False	Whether the task supports the cleanup mode at task deletion time

PublicTaskConfig

Field	Type	Req	Default	Description
<code>task</code>	PublicTask	yes		
<code>kbid</code>	string null			
<code>account_id</code>	string	yes		
<code>account_type</code>	string	yes		
<code>nua_client_id</code>	string null			
<code>user_id</code>	string	yes		
<code>parameters</code>	DataAugmentation-Output SemanticModelMigrationParams null			
<code>id</code>	string	yes		
<code>timestamp</code>	string	yes		
<code>defined_at</code>	string null			
<code>enabled</code>	boolean	yes		

PublicTaskRequest

Field	Type	Req	Default	Description
<code>task</code>	PublicTask	yes		

source	TrainingTaskDataset Source	yes		
kbid	string null			
dataset_id	string null			
account_id	string	yes		
nua_client_id	string null			
user_id	string	yes		
id	string	yes		
task_config_id	string null	yes		
timestamp	string	yes		
scheduled	boolean		False	
completed	boolean		False	
stopped	boolean		False	
scheduled_at	string null			
completed_at	string null			
stopped_at	string null			
failed	boolean		False	
retries	integer		0	
parameters	DataAugmentation- Output SemanticModelMigra- tionParams null			
log	string null			
cleanup_parent_task_id	string null			The ID of the task that this cleanup task is cleaning up after. This is only set to cleanup tasks.

PublicTaskSet

Field	Type	Req	Default	Description
-------	------	-----	---------	-------------

<code>request</code>	PublicTaskRequest null			
<code>config</code>	PublicTaskConfig null			

PushPayload

Field	Type	Req	Default	Description
<code>kbid</code>	string null			
<code>uuid</code>	string null			
<code>title</code>	string null			
<code>labels</code>	array<string>		<code>[]</code>	
<code>slug</code>	string null			
<code>source</code>	Source null			
<code>userid</code>	string null			
<code>genericfield</code>	map<string, Text>		<code>{}</code>	
<code>filefield</code>	map<string, string>		<code>{}</code>	
<code>linkfield</code>	map<string, LinkUpload>		<code>{}</code>	
<code>textfield</code>	map<string, Text>		<code>{}</code>	
<code>layoutfield</code>	map<string, ProcessingLayoutDiff>		<code>{}</code>	
<code>conversationfield</code>	map<string, ProcessingConversation>		<code>{}</code>	
<code>generated_conversationfield</code>	map<string, ProcessingGeneratedConversation>		<code>{}</code>	
<code>processing_options</code>	PushProcessingOptions null			
<code>learning_config</code>	LearningConfig null			
<code>webhook_config</code>	WebhookConfig null			

PushProcessingOptions

Field	Type	Req	Default	Description
ml_text	boolean		True	
extract_strategy	string null			
split_strategy	string null			

PushResponse

Field	Type	Req	Default	Description
seqid	integer null			
account_seq	integer null			
queue	ProcessingQueueType null			
uuid	string null			
processing_id	string	yes		

QAOperation

Field	Type	Req	Default	Description
question_generator_prompt	string			
system_question_generator_prompt	string			
summary_prompt	string			
generate_answers_prompt	string			
triggers	array<Trigger>			
max_questions	integer		20	Maximum number of question-answer pairs

				to generate per chunk.
--	--	--	--	------------------------

QueryEntityDetection

Enum: `predict`, `suggest`

QueryInfo

Field	Type	Req	Default	Description
<code>language</code>	string	yes		
<code>stop_words</code>	array<string>	yes		
<code>semantic_threshold</code>	number	yes		
<code>semantic_thresholds</code>	map<string, number>	yes		Semantic threshold for each semantic model
<code>visual_llm</code>	boolean	yes		
<code>max_context</code>	integer	yes		
<code>entities</code>	TokenSearch null	yes		
<code>sentence</code>	SentenceResponse null	yes		
<code>graph_nodes</code>	GraphNodeSearch null			Graph node embeddings if requested
<code>graph_edges</code>	GraphEdgeSearch null			Graph edge embeddings if requested
<code>query</code>	string null	yes		
<code>rephrased_query</code>	string null			

QueryModel

Model to represent a query request

Field	Type	Req	Default	Description
-------	------	-----	---------	-------------

<code>text</code>	string null			The query text to be processed
<code>query_image</code>	Image null			Image to be considered as part of the query. Even if the rephrase parameter is set to false, the rephrasing process will occur, combining the provided text with the image's visual features in the reph...
<code>rephrase</code>	boolean		<code>False</code>	If true, the model will rephrase the input text before processing
<code>rephrase_prompt</code>	string null			Custom prompt for rephrasing the input text
<code>generative_model</code>	string null			The generative model to use for rephrasing
<code>semantic_models</code>	array<string> null			Semantic models to compute the sentence vector for, if not provided, it will only compute the sentence vector for default semantic model in the Knowledge box's configuration.
<code>agentic_entities</code>	boolean		<code>False</code>	If true, the model will return the entities detected in the sentence guided by an already defined Graph Extraction Agent in the Knowledge Box.
<code>graph_nodes</code>	array<string> null			List of graph node names to compute embeddings for. If not provided, no graph node embeddings will be computed.

<code>semantic_graph_node_models</code>	<code>array<string> null</code>			Semantic models to compute graph node embeddings for. If not provided, it will only compute embeddings for the default graph semantic model in the Knowledge box's configuration.
<code>graph_edges</code>	<code>array<string> null</code>			List of graph edge labels to compute embeddings for. If not provided, no graph edge embeddings will be computed.
<code>semantic_graph_edge_models</code>	<code>array<string> null</code>			Semantic models to compute graph edge embeddings for. If not provided, it will only compute embeddings for the default graph edge semantic model in the Knowledge box's configuration.

QueryModelAccount

Model to represent a query request without a defined KB

Field	Type	Req	Default	Description
<code>text</code>	<code>string null</code>			The query text to be processed
<code>query_image</code>	<code>Image null</code>			Image to be considered as part of the query. Even if the rephrase parameter is set to false, the rephrasing process will occur, combining the provided text with the image's visual

				features in the reph...
rephrase	boolean		False	If true, the model will rephrase the input text before processing
rephrase_prompt	string null			Custom prompt for rephrasing the input text
generative_model	string		chatgpt-azure-4o	The generative model to use for rephrasing.
semantic_models	array<string> null			Semantic models to compute the sentence vector for, if not provided, it will only compute the sentence vector for default semantic model in the Knowledge box's configuration.
agentic_entities	boolean		False	If true, the model will return the entities detected in the sentence guided by an already defined Graph Extraction Agent in the Knowledge Box.
graph_nodes	array<string> null			List of graph node names to compute embeddings for. If not provided, no graph node embeddings will be computed.
semantic_graph_node_models	array<string> null			Semantic models to compute graph node embeddings for. If not provided, it will only compute embeddings for the default graph semantic model in the Knowledge box's configuration.
graph_edges	array<string> null			List of graph edge

				labels to compute embeddings for. If not provided, no graph edge embeddings will be computed.
<code>semantic_graph_edge_models</code>	<code>array<string> null</code>			Semantic models to compute graph edge embeddings for. If not provided, it will only compute embeddings for the default graph edge semantic model in the Knowledge box's configuration.
<code>token_model</code>	<code>string</code>		<code>multilingual</code>	The NER model to use for extracting entities from the input text.
<code>semantic_model</code>	<code>string</code>		<code>multilingual-2024-05-06</code>	Semantic model to compute the sentence vector for.

Question

Field	Type	Req	Default	Description
<code>text</code>	<code>string</code>	yes		
<code>language</code>	<code>string null</code>			
<code>ids_paragraphs</code>	<code>array<string></code>	yes		

QuestionAnswer

Field	Type	Req	Default	Description
<code>question</code>	<code>Question</code>	yes		
<code>answers</code>	<code>array<Answer></code>	yes		

QuestionAnswerAnnotation

Field	Type	Req	Default	Description
<code>question_answer</code>	QuestionAnswer	yes		
<code>cancelled_by_user</code>	boolean		False	

QuestionAnswers

Field	Type	Req	Default	Description
<code>question_answer</code>	array<QuestionAnswer>	yes		

QueueType

Enum: `private`, `shared`

Range

Type: `object` (free-form map)

RankFusionName

Enum: `rrf`

Reasoning

Field	Type	Req	Default	Description
<code>display</code>	boolean		True	Whether to display the reasoning steps in the response.
<code>effort</code>	enum[none, minimal, low, medium, high, xhigh]		medium	How much reasoning the model should use before answering. Higher values can improve complex answers but may take longer.
<code>budget_tokens</code>	integer		15000	Maximum number of tokens the model can spend on reasoning before answering.

				Lower values can make responses faster; higher values allow deeper reasoning.
--	--	--	--	---

ReasoningConfig

Field	Type	Req	Default	Description
<code>default</code>	integer learning_models__providers__models__Effort	yes		
<code>available_efforts</code>	array<learning_models__providers__models__Effort> null			
<code>supports_disabling</code>	boolean		<code>True</code>	
<code>chat_template_key</code>	string null			

ReasoningTags

For detecting reasoning blocks

Field	Type	Req	Default	Description
<code>start</code>	string	yes		
<code>end</code>	string	yes		

ReasoningType

Enum: `0`, `1`, `2`, `3`

ReciprocalRankFusion

Field	Type	Req	Default	Description
<code>name</code>	string		<code>rrf</code>	
<code>k</code>	number		<code>60.0</code>	k parameter changes the influence top-

				ranked and lower-ranked elements have. Research has shown that 60 is a performant value across datasets
<code>window</code>	integer null			Number of elements for retrieval to do RRF. Window must be greater or equal to topk. Greater values will increase probability of multi match at cost of retrieval time
<code>boosting</code>	ReciprocalRankFusionWeights			Define different weights for each retriever. This allows to assign different priorities to different retrieval methods. RRF scores will be multiplied by this value. The default is 1 for each retriever...

ReciprocalRankFusionWeights

Field	Type	Req	Default	Description
<code>keyword</code>	number		1.0	
<code>semantic</code>	number		1.0	

Relation-Input

Field	Type	Req	Default	Description
<code>prop</code>	string		<code>relation</code>	
<code>label</code>	string null			
<code>type</code>	RelationType null			

Relation-Output

Field	Type	Req	Default	Description
<code>relation</code>	RelationType	yes		
<code>label</code>	string null			
<code>metadata</code>	RelationMetadata null			
<code>from</code>	RelationEntity null			
<code>to</code>	RelationEntity	yes		

RelationDirection

Enum: `in`, `out`

RelationEdgeVector

Field	Type	Req	Default	Description
<code>relation_label</code>	string	yes		
<code>vector</code>	array<number>	yes		

RelationEntity

Field	Type	Req	Default	Description
<code>value</code>	string	yes		
<code>type</code>	RelationNodeType	yes		
<code>group</code>	string null			

RelationExample

Field	Type	Req	Default	Description
<code>source</code>	string	yes		Entity name from which the relation starts
<code>target</code>	string	yes		Entity name to which the relation ends

label	string	yes		Type of relation
-------	--------	-----	--	------------------

RelationMetadata

Field	Type	Req	Default	Description
paragraph_id	string null			
source_start	integer null			
source_end	integer null			
to_start	integer null			
to_end	integer null			
data_augmentation_task_id	string null			

RelationModel

Enum: `base`

RelationNodeType

Enum: `entity`, `label`, `resource`, `user`

RelationNodeVector

Field	Type	Req	Default	Description
node_value	string	yes		
vector	array<number>	yes		

RelationRanking

Enum: `reranker`, `generative`

Relations

Field	Type	Req	Default	Description
entities	map<string, EntitySubgraph>	yes		

RelationType

Enum: ABOUT, CHILD, COLAB, ENTITY, OTHER, SYNONYM

RemiRequest

Model to represent a request for the REMi model Metrics will be computed accordingly to the inputs given

- Answer relevance will be computed if `answer` and `question` are provided
- Context relevance will be computed if `question` and `contexts` are provided
- Groundedness will be computed if `answ...`

Field	Type	Req	Default	Description
<code>user_id</code>	string	yes		
<code>question</code>	string null			
<code>answer</code>	string null			
<code>contexts</code>	array<string> null			

RemiResponse

Field	Type	Req	Default	Description
<code>time</code>	number	yes		
<code>answer_relevance</code>	AnswerRelevance null			Answer Relevance measures the relevance of the generated answer to the user query, in a scale of 0 to 5.
<code>context_relevance</code>	array<integer null> null			Context Relevance measures the relevance of the retrieved context to the user query, in a scale of 0 to 5. The score will be None if there was an error computing the score for a specific context.
<code>groundedness</code>	array<integer null> null			Groundedness measures the degree to which the generated answer is grounded in the retrieved context, in a

				scale of 0 to 5. The score will be None if there was an error computing the score for a specif...
consumption	ConsumptionResponse null			

RephraseConfig

Field	Type	Req	Default	Description
max_output_tokens	integer		1025	
fallback_model	string null			

RephraseModel

Field	Type	Req	Default	Description
question	string	yes		
chat_history	array<Message>			
context	array<Message>			
user_context	array<string> null			
user_id	string		system	
generative_model	string null			The generative model to use for the rephrase endpoint. If not provided, the model configured for the Knowledge Box is used.
prompt	string null			Prompt to send the model to rephrase the sentence, if not provided, the default prompt will be used. It must include the {question} placeholder. The

				placeholder will be replaced with the original ques...
<code>chat_history_relevance_threshold</code>	number null			

Representation

Field	Type	Req	Default	Description
<code>is_a_table</code>	boolean null			
<code>reference_file</code>	string null			

RequestSecurity

Security metadata for the search request

Field	Type	Req	Default	Description
<code>access_groups</code>	array<string>		<code>[]</code>	List of group ids to do the request with.

RerankerName

Rerankers - Predict reranker: after retrieval, send the results to Predict API to rerank it. This method uses a reranker model, so one can expect better results at the expense of more latency. This will be the new default
 - No-operation (noop) reranker: maintain order and do not rerank the results a...

Enum: `predict`, `noop`

RerankModel

Field	Type	Req	Default	Description
<code>question</code>	string	yes		
<code>user_id</code>	string	yes		
<code>context</code>	map<string, string>		<code>{}</code>	

RerankResponse

Field	Type	Req	Default	Description
<code>context_scores</code>	map<string, number>	yes		Scores for each context given by the reranker
<code>consumption</code>	ConsumptionResponse null			Consumption details for the rerank operation

Resource-Input

Matches all fields of a resource given its id or slug

Field	Type	Req	Default	Description
<code>prop</code>	string		<code>resource</code>	
<code>id</code>	string null			UUID of the resource to match
<code>slug</code>	string null			Slug of the resource to match

Resource-Output

Field	Type	Req	Default	Description
<code>id</code>	string	yes		
<code>slug</code>	string null			
<code>title</code>	string null			
<code>summary</code>	string null			
<code>icon</code>	string null			
<code>thumbnail</code>	string null			
<code>metadata</code>	Metadata null			
<code>usermetadata</code>	UserMetadata null			
<code>fieldmetadata</code>	array<UserFieldMetadata> null			

computedmetadata	ComputedMetadata null			
created	string null			
modified	string null			
last_seqid	integer null			
last_account_seq	integer null			
queue	QueueType null			
hidden	boolean null			
origin	Origin null			
extra	Extra null			
relations	array<Relation-Output> null			
data	ResourceData null			
security	ResourceSecurity null			Resource security metadata

ResourceCreated

Field	Type	Req	Default	Description
uuid	string	yes		
elapsed	number null			
seqid	integer null			

ResourceData

Field	Type	Req	Default	Description
texts	map<string, TextFieldData> null			
files	map<string, FileFieldData> null			
links	map<string, LinkFieldData> null			

<code>conversations</code>	<code>map<string, ConversationFieldData> null</code>			
<code>generics</code>	<code>map<string, GenericFieldData> null</code>			
<code>key_values</code>	<code>map<string, KeyValueFieldData> null</code>			

ResourceFieldPrefix

Matches a field or set of fields. This filter is for internal use only and is not exposed in the public API schema.

Field	Type	Req	Default	Description
<code>prop</code>	<code>string</code>		<code>resource_field_prefix</code>	
<code>resource_id</code>	<code>string null</code>			ID of the resource containing the field(s) to match
<code>resource_slug</code>	<code>string null</code>			Slug of the resource containing the field(s) to match.
<code>field_type</code>	<code>FieldTypeName</code>	yes		Type of the fields to match
<code>field_name_prefix</code>	<code>string</code>	yes		Prefix of the name of the field to match. If blank, matches all fields of the given type in the given resource

ResourceList

Field	Type	Req	Default	Description
<code>resources</code>	<code>array<Resource-Output></code>	yes		
<code>pagination</code>	<code>ResourcePagination</code>	yes		

ResourceMimetype

Matches resources with a mimetype. The mimetype of a resource can be assigned independently of the mimetype of its fields. In resources with multiple fields, you may prefer to use `field_mimetype`

Field	Type	Req	Default	Description
<code>prop</code>	string		<code>resource_mime_type</code>	
<code>type</code>	string	yes		Type of the mimetype to match. e.g: In image/jpeg, type is image
<code>subtype</code>	string null			Type of the mimetype to match. e.g: In image/jpeg, subtype is jpeg. Leave blank to match all mimetype of the type

ResourcePagination

Field	Type	Req	Default	Description
<code>page</code>	integer	yes		
<code>size</code>	integer	yes		
<code>last</code>	boolean	yes		

ResourceProcessingStatus

Enum: `PENDING`, `PROCESSED`, `ERROR`, `EMPTY`, `BLOCKED`, `EXPIRED`

ResourceProperties

Enum: `basic`, `origin`, `extra`, `relations`, `values`, `extracted`, `errors`, `security`

ResourceSecurity

Security metadata for the resource

Field	Type	Req	Default	Description
<code>access_groups</code>	array<string>		<code>[]</code>	List of group ids that can access the resource.

ResponseFormatJSONObject

JSON object response format. An older method of generating JSON responses. Using `json_schema` is recommended for models that support it. Note that the model will not generate JSON without a system or user message instructing it to do so.

Field	Type	Req	Default	Description
<code>type</code>	string	yes		

ResponseFormatJSONSchema

JSON Schema response format. Used to generate structured JSON responses. Learn more about [Structured Outputs](see docs)

Field	Type	Req	Default	Description
<code>json_schema</code>	JSONSchema	yes		Structured Outputs configuration options, including a JSON Schema.
<code>type</code>	string	yes		

ResponseFormatText

Default response format. Used to generate text responses.

Field	Type	Req	Default	Description
<code>type</code>	string	yes		

RetrievalAgentExportRequest

Field	Type	Req	Default	Description
<code>passphrase</code>	string	yes		Passphrase to encrypt the exported configuration. Will be required for import.

Row

Field	Type	Req	Default	Description
<code>cell</code>	array<string> null			

RowsPreview

Field	Type	Req	Default	Description
<code>sheets</code>	map<string, Sheet> null			

Rule

Field	Type	Req	Default	Description
<code>prompt</code>	string null			

Rules

Field	Type	Req	Default	Description
<code>rules</code>	array<Rule string>			List of rules that the workflow should follow. Each rule can be a string or a Rule object with a prompt.

SCORE_TYPE

Enum: `VECTOR`, `BM25`, `BOTH`, `RERANKER`, `RELATION_RELEVANCE`

SeedConfig

Field	Type	Req	Default	Description
<code>min</code>	integer		<code>0</code>	
<code>max</code>	integer		<code>2147483647</code>	
<code>default</code>	integer null			

SemanticConfig

Field	Type	Req	Default	Description
similarity	SimilarityFunction	yes		
size	integer	yes		
threshold	number	yes		
max_tokens	integer null			
matryoshka_dims	array<integer>		[]	
external	boolean		False	

SemanticModelMigrationParams

Field	Type	Req	Default	Description
semantic_model_id	string	yes		The id of the semantic model to migrate to. This must be a valid semantic model id available for the account
hf_embedding_key	HFEmbeddingKey null			The HuggingFace embedding key to use. This is only required if the semantic model is hfembedding

Sentence

Field	Type	Req	Default	Description
start	integer null			
end	integer null			
key	string null			

SentenceRequestData

Field	Type	Req	Default	Description
-------	------	-----	---------	-------------

<code>texts</code>	<code>array<string></code>	yes		
<code>model</code>	<code>string null</code>			

SentenceResponse

Field	Type	Req	Default	Description
<code>data</code>	<code>array<number></code>			
<code>vectors</code>	<code>map<string, array<number>></code>	yes		Sentence vectors for each semantic model
<code>time</code>	<code>number</code>	yes		
<code>timings</code>	<code>map<string, number></code>	yes		Time taken to compute the sentence vector for each semantic model
<code>consumption</code>	<code>ConsumptionResponse null</code>			Consumption details for the sentence operation

SentencesResponse

Field	Type	Req	Default	Description
<code>vectors</code>	<code>array<array<number>></code>	yes		Sentence vectors for each semantic model
<code>time</code>	<code>number</code>	yes		
<code>timings</code>	<code>map<string, number></code>	yes		Time taken to compute the sentence vector for each semantic model
<code>consumption</code>	<code>ConsumptionResponse null</code>			Consumption details for the sentence operation

SessionData

Field	Type	Req	Default	Description
-------	------	-----	---------	-------------

slug	string	yes		
name	string	yes		
summary	string	yes		
data	string	yes		
format	TextFormat	yes		

Sheet

Field	Type	Req	Default	Description
rows	array<Row> null			

SimilarityFunction

Enum: 0, 1

SortFields

Enum: timestamp, scheduled_at, completed_at

SortOrder

Enum: asc, desc

Source

Enum: 0, 1

SourceNode

Field	Type	Req	Default	Description
prop	string		source_node	
value	string null			
match	NodeMatchKindName		exact	
type	RelationNodeType null		entity	
group	string null			

SplitConfig

Field	Type	Req	Default	Description
<code>max_paragraph</code>	integer		0	
<code>custom_split</code>	CustomSplitStrategy null		0	

SplitConfiguration

Field	Type	Req	Default	Description
<code>name</code>	string	yes		
<code>max_paragraph</code>	integer		0	
<code>custom_split</code>	CustomSplitStrategy		NONE	
<code>llm_split</code>	LLMSplitConfig null			
<code>manual_split</code>	ManualSplitConfig null			

SplitConfiguration-Input

Field	Type	Req	Default	Description
<code>name</code>	string			
<code>max_paragraph</code>	integer		0	
<code>custom_split</code>	CustomSplitStrategy null		0	
<code>llm_split</code>	LLMSplitConfig-Input null			
<code>manual_split</code>	ManualSplitConfig null			

SplitConfiguration-Output

Field	Type	Req	Default	Description
<code>name</code>	string			

<code>max_paragraph</code>	integer		0	
<code>custom_split</code>	CustomSplitStrategy null		0	
<code>llm_split</code>	LLMSplitConfig-Output null			
<code>manual_split</code>	ManualSplitConfig null			

Stability

Enum: `stable`, `preview`

StreamDataTypes

Enum: `FullText`, `Chunk`, `Metadata`, `FilePreview`, `FileThumbnail`, `GeneratedFile`, `Links`

StreamPostprocessor

Postprocessors applied to streaming generative responses.

Enum: `deduplicate_llm_citations`

StringModelTypes

Enum: `GENERATIVE`, `NER`, `RESOURCE_LABELER`, `CLASSIFIER`, `ANONYMIZER`, `VISUAL_LABELER`, `SUMMARY`, `DUMMY`, `PARAGRAPH_LABELER`, `EMBEDDINGS`, `RELATIONS`

StructuredOutputConfig

Field	Type	Req	Default	Description
<code>use_tool_calling</code>	boolean		True	

SummarizedModelWithConsumption

Field	Type	Req	Default	Description
<code>resources</code>	map<string, SummarizedResource>	yes		
<code>summary</code>	string			
<code>consumption</code>	ConsumptionResponse null			

SummarizedResource

Field	Type	Req	Default	Description
<code>summary</code>	string	yes		
<code>tokens</code>	integer	yes		
<code>input_tokens</code>	integer	yes		
<code>output_tokens</code>	integer	yes		

SummarizeModel

Field	Type	Req	Default	Description
<code>resources</code>	map<string, SummarizeResource>	yes		
<code>summary_kind</code>	SummaryKind		<code>simple</code>	
<code>user_prompt</code>	string null			
<code>generative_model</code>	string null			

SummarizeResource

Field	Type	Req	Default	Description
<code>fields</code>	map<string, string>	yes		

Summary

Enum: `extended`, `simple`

SummaryKind

Enum: `simple`, `extended`

SummaryPrompt

Field	Type	Req	Default	Description
<code>prompt</code>	string	yes		

SyncAskMetadata

Field	Type	Req	Default	Description
<code>tokens</code>	AskTokens null			Number of tokens used in the LLM context and answer
<code>timings</code>	AskTimings null			Timings of the generative model

SyncAskResponse

Field	Type	Req	Default	Description
<code>answer</code>	string	yes		The generative answer to the query
<code>reasoning</code>	string null			The reasoning steps followed by the LLM to generate the answer. This is returned only if the reasoning feature is enabled in the request.
<code>answer_json</code>	object (free-form map) null			The generative JSON answer to the query. This is returned only if the <code>answerjsonschema</code> parameter is provided in the request.
<code>status</code>	string	yes		The status of the query execution. It can be 'success', 'error', 'nocontext' or 'noretrievaldata'
<code>retrieval_results</code>	KnowledgeboxFindResults	yes		The retrieval results of the query
<code>retrieval_best_matches</code>	array<AskRetrievalMatch>		<code>[]</code>	Sorted list of best matching text blocks in the retrieval step. This includes the main query and prequeries results, if

				any.
<code>prequeries</code>	<code>map<string, KnowledgeboxFindResults> null</code>			The retrieval results of the prequeries
<code>learning_id</code>	<code>string</code>			The id of the learning request. This id can be used to provide feedback on the learning process.
<code>relations</code>	<code>Relations null</code>			The detected relations of the answer
<code>citations</code>	<code>object (free-form map)</code>			The citations of the answer. List of references to the resources used to generate the answer.
<code>citation_footnote_to_context</code>	<code>map<string, string></code>			Maps ids in the footnote citations to querycontext keys (normally paragraph ids)
<code>augmented_context</code>	<code>AugmentedContext null</code>			Augmented text blocks that were sent to the LLM as part of the RAG strategies applied on the retrieval results in the request.
<code>prompt_context</code>	<code>array<string> null</code>			The prompt context used to generate the answer. Returned only if the debug flag is set to true
<code>predict_request</code>	<code>object (free-form map) null</code>			The internal predict request used to generate the answer. Returned only if the debug flag is set to true
<code>metadata</code>	<code>SyncAskMetadata null</code>			Metadata of the query execution. This includes the number of tokens used in the LLM context and

				answer, and the timings of the generative model.
consumption	Consumption null			The consumption of the query execution. Return only if 'X-show-consumption' header is set to true in the request.
error_details	string null			Error details message in case there was an error
debug	object (free-form map) null			Debug information about the ask operation. The metadata included in this field is subject to change and should not be used in production. Note that it is only available if the debug parameter is set t...

SyncMetadata

Field	Type	Req	Default	Description
file_id	string	yes		Identifier of the file in the origin cloud storage system
auth_provider	string	yes		Authentication provider used to access the origin cloud storage system
content_hash	string	yes		Content hash of the file in the origin cloud storage system. The hash algorithm used depends on the origin system.

TableImageStrategy

Field	Type	Req	Default	Description
name	string		tables	

Task-Input

Enum: PARAGRAPH_CLASSIFICATION, FIELD_CLASSIFICATION, SENTENCE_CLASSIFICATION, TOKEN_CLASSIFICATION, IMAGE_CLASSIFICATION, PARAGRAPH_STREAMING, QUESTION_ANSWER_STREAMING, FIELD_STREAMING

Task-Output

Field	Type	Req	Default	Description
name	string	yes		
image_repository	string		europe-west4-docker.pkg.dev/nuclia-internal/nuclia	
image_name	string null			
image_version	string		latest	
entrypoint	array<string> null			
sentry_id	string null			
memory	string		1Gi	
cpu	string		1000m	
imagePullPolicy	string		Always	
use_gpu	boolean		False	
trainset_type	integer null			
command	array<string>	yes		
disallowed_account_types	array<string>		[]	
model_type	object			
listing	boolean		True	

<code>data_augmentation</code>	boolean		False	
<code>description</code>	string null			
<code>feedback</code>	boolean		False	
<code>cleanup_command</code>	array<string> null			

TaskDefinition

Field	Type	Req	Default	Description
<code>name</code>	TaskName	yes		
<code>description</code>	string null			
<code>validation</code>	object (free-form map) null			

TaskEnable

Field	Type	Req	Default	Description
<code>enabled</code>	boolean	yes		

TaskList

Field	Type	Req	Default	Description
<code>tasks</code>	array<TaskDefinition>	yes		
<code>running</code>	array<PublicTaskRequest>	yes		
<code>configs</code>	array<PublicTaskConfig>	yes		
<code>done</code>	array<PublicTaskRequest>	yes		

TaskName

Enum: `dummy`, `env`, `demo-dataset`, `labeler`, `llm-graph`, `synthetic-questions`, `ask`, `llm-align`, `semantic-model-migrator`, `llama-guard`, `prompt-guard`, `memory`

TaskPatch

Field	Type	Req	Default	Description
<code>parameters</code>	DataAugmentation-Input SemanticModelMigrationParams null	yes		Parameters to be passed to the task. These must match the validation field for the Task definition class

TaskRequest

Field	Type	Req	Default	Description
<code>task</code>	Task-Output	yes		
<code>source</code>	TrainingTaskDataset Source	yes		
<code>kbid</code>	string null			
<code>dataset_id</code>	string null			
<code>dataset_path</code>	string null			
<code>account_id</code>	string	yes		
<code>account_type</code>	string	yes		
<code>nua_client_id</code>	string null			
<code>user_id</code>	string	yes		
<code>parameters</code>	DataAugmentation-Output SemanticModelMigrationParams null			
<code>id</code>	string	yes		
<code>task_config_id</code>	string null			
<code>timestamp</code>	string	yes		
<code>scheduled</code>	boolean		False	
<code>completed</code>	boolean		False	
<code>stopped</code>	boolean		False	
<code>scheduled_at</code>	string null			

completed_at	string null			
stopped_at	string null			
failed	boolean		False	
retries	integer		0	
token	string null			
key_id	string null			
ndb_token	string null			
log	string null			
cleanup_parent_task_id	string null			The ID of the task that this cleanup task is cleaning up after. This is only set to cleanup tasks.

TaskResponse

Field	Type	Req	Default	Description
name	TaskName	yes		
status	JobStatus	yes		
id	string	yes		

TaskStart

Field	Type	Req	Default	Description
name	TaskName	yes		
parameters	DataAugmentation-Input SemanticModelMigrationParams null			Parameters to be passed to the task. These must match the validation field for the Task definition class
uuid_task	string null			UUID of an already configured task. This is used to start a task that was already configured

TaskStartKB

Field	Type	Req	Default	Description
name	TaskName	yes		
parameters	DataAugmentation-Input SemanticModelMigrationParams null			Parameters to be passed to the task. These must match the validation field for the Task definition class
uuid_task	string null			UUID of an already configured task. This is used to start a task that was already configured
apply	ApplyOptions		ALL	Defines how the tasks should be applied to the existing data. - EXSITING: Only apply to existing data (starts a worker that executes the task) - NEW: Only apply to new data (enables the task at proces...
enabled	boolean		True	Whether the task should be enabled at the time of creation. This only applies to tasks involving 'NEW' apply options, as 'EXISTING' tasks start immediately

TemperatureConfig

Field	Type	Req	Default	Description
min	number		0.0	
max	number	yes		
default	number null	yes		

Text

Field	Type	Req	Default	Description
body	string	yes		
format	Format	yes		
extract_strategy	string null			
split_strategy	string null			
classification_labels	array<ClassificationLabel> null			

TextBlockAugmentationType

Enum: neighbouring_paragraphs, conversation, hierarchy, full_resource, field_extension, metadata_extension

TextFieldData

Field	Type	Req	Default	Description
value	FieldText null			
extracted	TextFieldExtractedData null			
error	Error null			
status	string null			
errors	array<Error> null			

TextFieldExtractedData

Field	Type	Req	Default	Description
text	ExtractedText null			
metadata	FieldComputedMetadata null			
large_metadata	LargeComputedMetadata null			

<code>vectors</code>	VectorObject null			
<code>question_answers</code>	FieldQuestionAnswers null			
<code>relation_node_vectors</code>	map<string, array<RelationNodeVector>> null			
<code>relation_edge_vectors</code>	map<string, array<RelationEdgeVector>> null			

TextFormat

Enum: `PLAIN`, `HTML`, `RST`, `MARKDOWN`, `JSON`, `KEEP_MARKDOWN`, `JSONL`, `PLAIN_BLANKLINE_SPLIT`

TextGenerationKey

Field	Type	Req	Default	Description
<code>model</code>	string			

TextGenerationUserPrompt

Field	Type	Req	Default	Description
<code>prompt</code>	string			

TextPosition

Field	Type	Req	Default	Description
<code>page_number</code>	integer null			
<code>index</code>	integer	yes		
<code>start</code>	integer	yes		
<code>end</code>	integer	yes		
<code>start_seconds</code>	array<integer> null			
<code>end_seconds</code>	array<integer> null			

Tokenizer

Enum: 0, 1, 2

TokenizerConfig

Field	Type	Req	Default	Description
name	string null			
path	string null			
extra_kwargs	object (free-form map) null			

TokensDetail

Field	Type	Req	Default	Description
input	number	yes		
output	number	yes		
image	number	yes		

TokenSearch

Field	Type	Req	Default	Description
tokens	array<Ner>		[]	
time	number	yes		
input_tokens	integer		0	
consumption	ConsumptionResponse null			

TokenValues

Field	Type	Req	Default	Description
min	integer		0	
max	integer	yes		

Tool

Field	Type	Req	Default	Description
<code>name</code>	string	yes		
<code>description</code>	string	yes		
<code>parameters</code>	object (free-form map)			Schema of the tool

ToolChoiceAuto

Field	Type	Req	Default	Description
<code>type</code>	string		<code>auto</code>	

ToolChoiceForced

Field	Type	Req	Default	Description
<code>type</code>	string		<code>forced</code>	
<code>name</code>	string	yes		

ToolChoiceNone

Field	Type	Req	Default	Description
<code>type</code>	string		<code>none</code>	

ToolChoiceRequired

Field	Type	Req	Default	Description
<code>type</code>	string		<code>required</code>	

ToolMessage

Field	Type	Req	Default	Description
<code>type</code>	string		<code>tool</code>	

author	Author		USER	
text	string			
tool_call_id	string	yes		
name	string null			
content	object			

TrainingTaskDatasetSource

Enum: `nucliadb`, `dataset`

Trigger

Field	Type	Req	Default	Description
url	string			
headers	map<string, string>			
params	map<string, string>			

TypeBlock

Enum: `0`, `1`, `2`, `3`, `4`, `5`, `6`

TypeParagraph

Enum: `TEXT`, `OCR`, `INCEPTION`, `DESCRIPTION`, `TRANSCRIPT`, `TITLE`, `TABLE`

UserClassification

Field	Type	Req	Default	Description
labelset	string	yes		
label	string	yes		
cancelled_by_user	boolean		<code>False</code>	

UserFieldMetadata

Field-level metadata set by the user via the rest api

Field	Type	Req	Default	Description
-------	------	-----	---------	-------------

paragraphs	array<ParagraphAnnotation>		[]	
question_answers	array<QuestionAnswerAnnotation>		[]	
field	FieldID	yes		

UserLearningKeys

Field	Type	Req	Default	Description
openai	learning_config_models__OpenAIKey null			
azure_openai	learning_config_models__AzureOpenAIKey null			
mistral	learning_config_models__MistralKey null			
palm	learning_config_models__PalmKey null			
anthropic	learning_config_models__AnthropicKey null			
claude3	learning_config_models__AnthropicKey null			
anthropic_vertex	AnthropicVertexKey null			
azure_mistral	learning_config_models__MistralKey null			
hf_llm	HFLLMKey null			
hf_embedding	learning_config_models__HFEmbeddingKey null			
azure_aai	AzureAIIKey null			
openai_compat	OpenAICompatModel-Output null			

UserLearningKeys-Input

Field	Type	Req	Default	Description
<code>openai</code>	OpenAIKey null			
<code>azure_openai</code>	AzureOpenAIKey null			
<code>palm</code>	PalmKey null			
<code>anthropic</code>	AnthropicKey null			
<code>claude3</code>	AnthropicKey null			
<code>anthropic_vertex</code>	AnthropicVertexKey null			
<code>anthropic_bedrock</code>	AnthropicBedrockKey null			
<code>text_generation</code>	TextGenerationKey null			
<code>mistral</code>	MistralKey null			
<code>azure_mistral</code>	AzureMistralKey null			
<code>hf_llm</code>	HFLLMKey null			
<code>hf_embedding</code>	learning_protos__config_p2p__HFEmbeddingKey null			
<code>azure_aai</code>	AzureAIIKey null			
<code>openai_compat</code>	OpenAICompatModel-Input null			

UserLearningKeys-Output

Field	Type	Req	Default	Description
<code>openai</code>	OpenAIKey null			
<code>azure_openai</code>	AzureOpenAIKey null			
<code>palm</code>	PalmKey null			

anthropic	AnthropicKey null			
claude3	AnthropicKey null			
anthropic_vertex	AnthropicVertexKey null			
anthropic_bedrock	AnthropicBedrockKey null			
text_generation	TextGenerationKey null			
mistral	MistralKey null			
azure_mistral	AzureMistralKey null			
hf_llm	HFLLMKey null			
hf_embedding	learning_protos__config_p2p__HFEmbeddingKey null			
azure_aai	AzureAIIKey null			
openai_compat	OpenAICompatModel-Output null			

UserMetadata

Field	Type	Req	Default	Description
classifications	array<UserClassification>		[]	
relations	array<Relation-Output>		[]	

UserPrompt

Field	Type	Req	Default	Description
prompt	string	yes		

UserPrompts

Field	Type	Req	Default	Description
-------	------	-----	---------	-------------

<code>openai</code>	OpenAIUserPrompt null			
<code>azure_openai</code>	AzureUserPrompt null			
<code>palm</code>	PalmUserPrompt null			
<code>anthropic</code>	AnthropicUserPrompt null			
<code>text_generation</code>	TextGenerationUserPrompt null			
<code>mistral</code>	MistralUserPrompt null			
<code>azure_mistral</code>	AzureMistralUserPrompt null			
<code>claude3</code>	Claude3UserPrompt null			
<code>anthropic_vertex</code>	Claude3UserPrompt null			
<code>anthropic_bedrock</code>	Claude3UserPrompt null			
<code>deepseek</code>	DeepSeekUserPrompt null			
<code>nova</code>	NovaUserPrompt null			

ValidationError

Field	Type	Req	Default	Description
<code>loc</code>	array<string integer>	yes		
<code>msg</code>	string	yes		
<code>type</code>	string	yes		

Vector

Field	Type	Req	Default	Description
-------	------	-----	---------	-------------

<code>start</code>	integer null			
<code>end</code>	integer null			
<code>start_paragraph</code>	integer null			
<code>end_paragraph</code>	integer null			
<code>vector</code>	array<number> null			

VectorObject

Field	Type	Req	Default	Description
<code>vectors</code>	Vectors null			
<code>split_vectors</code>	map<string, Vectors> null			
<code>deleted_splits</code>	array<string> null			

Vectors

Field	Type	Req	Default	Description
<code>vectors</code>	array<Vector> null			

VisualLabeling

Enum: `disabled`, `enabled`

VLLMExtractionConfig

Field	Type	Req	Default	Description
<code>generative_model</code>	string null			
<code>generative_provider</code>	string null			
<code>rules</code>	array<string>		<code>[]</code>	

VLLMExtractionConfig-Input

Field	Type	Req	Default	Description
<code>rules</code>	array<string>			
<code>llm</code>	LLMConfig-Input null			

VLLMExtractionConfig-Output

Field	Type	Req	Default	Description
<code>rules</code>	array<string>			
<code>llm</code>	LLMConfig-Output null			

WebhookConfig

Field	Type	Req	Default	Description
<code>uri</code>	string null			
<code>headers</code>	map<string, string>		<code>{}</code>	

WorkflowData

Field	Type	Req	Default	Description
<code>id</code>	string	yes		
<code>name</code>	string	yes		
<code>description</code>	string null	yes		
<code>parameters</code>	object (free-form map) null	yes		
<code>rules</code>	Rules null			
<code>required</code>	array<string>			

WorkflowInput

Field	Type	Req	Default	Description
-------	------	-----	---------	-------------

<code>id</code>	string	yes		
<code>name</code>	string	yes		
<code>description</code>	string null			
<code>parameters</code>	object (free-form map) null			
<code>rules</code>	Rules			
<code>required</code>	array<string>			

WorkflowUpdate

Field	Type	Req	Default	Description
<code>name</code>	string	yes		
<code>description</code>	string	yes		
<code>parameters</code>	object (free-form map)	yes		
<code>required</code>	array<string>			
<code>rules</code>	Rules null			

zone — 124 schemas

AccountType

Enum: `stash-trial`, `v3starter`, `v3pro`, `v3enterprise`, `cowork`, `v3fly`, `v3growth`

ActivityLogsAskQuery

Field	Type	Req	Default	Description
<code>pagination</code>	Pagination		<code>{'limit': 10}</code>	
<code>year_month</code>	string	yes		
<code>show</code>	array<enum[id, date, user_id, user_type, client_type, total_duration, audit_metadata, resource_id,		<code>['resources_count', 'id', 'rephrased_question', 'answer', 'question',</code>	

	nuclia_tokens, token_details...]> string		'rag_strategi es_names', 'date']	
filters	QueryFiltersAsk	yes		

ActivityLogsChatQuery

Field	Type	Req	Default	Description
pagination	Pagination		{'limit': 10}	
year_month	string	yes		
show	array<enum[id, date, user_id, user_type, client_type, total_duration, audit_metadata, resource_id, nuclia_tokens, token_details...]> string		['question', 'rag_strategi es_names', 'id', 'date', 'rephrased_qu estion', 'answer']	
filters	QueryFiltersChat	yes		

ActivityLogsQuery

Field	Type	Req	Default	Description
pagination	Pagination		{'limit': 10}	
year_month	string	yes		
show	array<enum[id, date, user_id, user_type, client_type, total_duration, audit_metadata, resource_id, nuclia_tokens, token_details]> string		['date', 'id']	
filters	QueryFiltersCommon	yes		

ActivityLogsSearchQuery

Field	Type	Req	Default	Description
<code>pagination</code>	Pagination		<code>{'limit': 10}</code>	
<code>year_month</code>	string	yes		
<code>show</code>	array<enum[id, date, user_id, user_type, client_type, total_duration, audit_metadata, resource_id, nuclia_tokens, token_details...]> string		<code>['date', 'id', 'question', 'resources_count']</code>	
<code>filters</code>	QueryFiltersSearch	yes		

ActivityUsageResponse

Field	Type	Req	Default	Description
<code>timestamp</code>	string	yes		
<code>search</code>	integer	yes		
<code>chat</code>	integer	yes		
<code>ask</code>	integer	yes		

AddKbUser

Field	Type	Req	Default	Description
<code>id</code>	string	yes		
<code>role</code>	KbRole	yes		

AggregatedRemiScoreMetric

Field	Type	Req	Default	Description
<code>timestamp</code>	string	yes		
<code>metrics</code>	array<AggregatedRe	yes		

	miScoreValues>			
--	----------------	--	--	--

AggregatedRemiScoreValues

Field	Type	Req	Default	Description
name	string	yes		
min	number	yes		
max	number	yes		
average	number	yes		

Aggregation

Enum: hour, day, week, month, quarter, year, millennium

AuditMetadata

Field	Type	Req	Default	Description
eq	string null			
gt	string null			
ge	string null			
lt	string null			
le	string null			
ne	string null			
isnull	boolean null			
isin	array<string> null			
isnotin	array<string> null			
like	string null			
ilike	string null			
key	string	yes		

AuthorizeUrlOutput

Field	Type	Req	Default	Description
-------	------	-----	---------	-------------

<code>authorize_url</code>	string	yes		
----------------------------	--------	-----	--	--

AWSS3AssumeRoleCredentials

Field	Type	Req	Default	Description
<code>role_arn</code>	string	yes		
<code>external_id</code>	string	yes		

AzureCertificateCredentials

Credentials for Azure/SharePoint certificate-based authentication. This uses the client credentials flow with a certificate instead of a client secret. The certificate must be registered in the Azure App Registration. Users provide a .pfx file (base64-encoded) and its password. The private key and c...

Field	Type	Req	Default	Description
<code>tenant_id</code>	string	yes		
<code>client_id</code>	string	yes		
<code>pfx_base64</code>	string	yes		
<code>pfx_password</code>	string	yes		

BackupCreate

Field	Type	Req	Default	Description
<code>kb_id</code>	string	yes		The unique identifier of the knowledgebox to backup.

BackupCreateResponse

Field	Type	Req	Default	Description
<code>id</code>	string	yes		The unique identifier of the created backup.

BackupResponse

Field	Type	Req	Default	Description
<code>id</code>	string	yes		The unique identifier of the backup.
<code>account_id</code>	string	yes		The unique identifier of the account associated with the backup.
<code>started_at</code>	string	yes		The timestamp when the backup process started.
<code>kb_data</code>	KBDataResponse	yes		Metadata of the backed-up knowledgebox.
<code>finished_at</code>	string null			The timestamp when the backup process finished.
<code>size</code>	integer null			The size of the backup in bytes.

BackupRestore

Field	Type	Req	Default	Description
<code>slug</code>	string	yes		The slug of the new restored knowledgebox.
<code>title</code>	string	yes		The title of the new restored knowledgebox.

BatchUpdateKbUsers

Field	Type	Req	Default	Description
<code>add</code>	array<AddKbUser> null			
<code>update</code>	array<UpdateKbUser> null			
<code>delete</code>	array<string> null			

BrowseCapabilities

Capabilities for browsing storage in this external connection.

Field	Type	Req	Default	Description
<code>has_sites</code>	boolean		<code>False</code>	Whether the provider supports site browsing (SharePoint only).
<code>requires_site_search</code>	boolean		<code>False</code>	Whether sitesearch is required (True) or can list all sites (False). Only applicable for certificate auth connections.
<code>requires_site_url_resolution</code>	boolean		<code>False</code>	Whether the client must use the /resolvesite endpoint to obtain a siteid from a site URL before browsing drives. Applies to OAuth connections using Sites.Selected permission.

BrowseStorageOutput

Result of browsing storage, containing sites, drives or folders.

Field	Type	Req	Default	Description
<code>sites</code>	array<StorageSiteOutput> null			
<code>drives</code>	array<StorageDriveOutput> null			
<code>folders</code>	array<StorageFolderOutput> null			
<code>next_page_token</code>	string null			

ClientType**Enum:** `api`, `widget`, `web`, `dashboard`, `desktop`, `chrome_extension`**ContextRelevanceQuery**

Field	Type	Req	Default	Description
<code>value</code>	integer	yes		
<code>operation</code>	enum[gt, lt, eq]	yes		
<code>aggregation</code>	enum[average, max, min]	yes		

DownloadActivityLogsAskQuery

Field	Type	Req	Default	Description
<code>year_month</code>	string	yes		
<code>show</code>	array<enum[id, date, user_id, user_type, client_type, total_duration, audit_metadata, resource_id, nuclia_tokens, token_details...]> string		['resources_count', 'id', 'rephrased_question', 'answer', 'question', 'rag_strategies_names', 'date']	
<code>filters</code>	QueryFiltersAsk	yes		
<code>email_address</code>	string null			
<code>notify_via_email</code>	boolean		False	

DownloadActivityLogsChatQuery

Field	Type	Req	Default	Description
<code>year_month</code>	string	yes		
<code>show</code>	array<enum[id, date, user_id, user_type, client_type, total_duration, audit_metadata,		['question', 'rag_strategies_names', 'id', 'date',	

	resource_id, nuclia_tokens, token_details...]> string		'rephrased_qu estion', 'answer']	
filters	QueryFiltersChat	yes		
email_address	string null			
notify_via_email	boolean		False	

DownloadActivityLogsQuery

Field	Type	Req	Default	Description
year_month	string	yes		
show	array<enum[id, date, user_id, user_type, client_type, total_duration, audit_metadata, resource_id, nuclia_tokens, token_details]> string		['date', 'id']	
filters	QueryFiltersCommon	yes		
email_address	string null			
notify_via_email	boolean		False	

DownloadActivityLogsSearchQuery

Field	Type	Req	Default	Description
year_month	string	yes		
show	array<enum[id, date, user_id, user_type, client_type, total_duration, audit_metadata, resource_id, nuclia_tokens, token_details...]> string		['date', 'id', 'question', 'resources_co unt']	

	string			
filters	QueryFiltersSearch	yes		
email_address	string null			
notify_via_email	boolean		False	

DownloadFormat

Enum: ndjson, csv

DownloadRequest

Field	Type	Req	Default	Description
request_id	string	yes		
download_type	DownloadRequestType	yes		
download_format	DownloadFormat	yes		
event_type	EventType	yes		
requested_at	string	yes		
kb_id	string	yes		
download_url	string null	yes		

DownloadRequestType

Enum: query

DummyKBIndexProvider

Field	Type	Req	Default	Description
type	KBExternalIndexProviderType		unset	

EphemeralTokenProperties

Field	Type	Req	Default	Description
agent_session	string null			

<code>ttl</code>	integer null			
------------------	----------------	--	--	--

EphemeralTokenRequestPayload

Field	Type	Req	Default	Description
<code>agent_session</code>	string null			
<code>path</code>	string null			
<code>ttl</code>	integer null			

EphemeralTokenResponse

Field	Type	Req	Default	Description
<code>token</code>	string	yes		

EventDownloadsList

Field	Type	Req	Default	Description
<code>downloads</code>	array<string>	yes		

EventType

Enum: `visited`, `modified`, `deleted`, `new`, `search`, `suggest`, `indexed`, `chat`, `ask`, `retrieve`, `augment`, `started`, `stopped`, `processed`

ExternalAuthorizationInput

Input schema for external connection authorization with additional parameters.

Field	Type	Req	Default	Description
<code>widget_redirect_url</code>	string null			
<code>rao_redirect_url</code>	string	yes		

ExternalConnectionInput

Input schema for creating an external connection. For OAuth providers (google_oauth, azure_oauth, sharefile_oauth): only provider is required. The response will contain an authorize_url to complete the OAuth flow. For providers aws_s3_assume_role, azure_certificate_credentials: provider and credenti...

Field	Type	Req	Default	Description
provider	Provider	yes		
credentials	AWSS3AssumeRoleCredentials AzureCertificateCredentials null			

ExternalConnectionOutput

External connection output without sensitive credential data.

Field	Type	Req	Default	Description
id	string	yes		
kb_id	string	yes		
created_by	string	yes		
created_at	string	yes		
updated_at	string	yes		
provider	Provider	yes		
capabilities	BrowseCapabilities	yes		Capabilities for browsing storage in this external connection.

FileFilterInput

Field	Type	Req	Default	Description
mode	enum[include, exclude]	yes		
extensions	array<string> null			
glob_patterns	array<string> null			

FileFilterOutput

Field	Type	Req	Default	Description
<code>mode</code>	enum[include, exclude]	yes		
<code>extensions</code>	array<string> null			
<code>glob_patterns</code>	array<string> null			

GenericFilter_bool_

Field	Type	Req	Default	Description
<code>eq</code>	boolean null			
<code>gt</code>	boolean null			
<code>ge</code>	boolean null			
<code>lt</code>	boolean null			
<code>le</code>	boolean null			
<code>ne</code>	boolean null			
<code>isnull</code>	boolean null			
<code>isin</code>	array<boolean> null			
<code>isnotin</code>	array<boolean> null			

GenericFilter_ClientType_

Field	Type	Req	Default	Description
<code>eq</code>	ClientType null			
<code>gt</code>	ClientType null			
<code>ge</code>	ClientType null			
<code>lt</code>	ClientType null			
<code>le</code>	ClientType null			
<code>ne</code>	ClientType null			
<code>isnull</code>	boolean null			

<code>isin</code>	<code>array<ClientType> null</code>			
<code>isnotin</code>	<code>array<ClientType> null</code>			

GenericFilter_datetime_

Field	Type	Req	Default	Description
<code>eq</code>	<code>string null</code>			
<code>gt</code>	<code>string null</code>			
<code>ge</code>	<code>string null</code>			
<code>lt</code>	<code>string null</code>			
<code>le</code>	<code>string null</code>			
<code>ne</code>	<code>string null</code>			
<code>isnull</code>	<code>boolean null</code>			
<code>isin</code>	<code>array<string> null</code>			
<code>isnotin</code>	<code>array<string> null</code>			

GenericFilter_float_

Field	Type	Req	Default	Description
<code>eq</code>	<code>number null</code>			
<code>gt</code>	<code>number null</code>			
<code>ge</code>	<code>number null</code>			
<code>lt</code>	<code>number null</code>			
<code>le</code>	<code>number null</code>			
<code>ne</code>	<code>number null</code>			
<code>isnull</code>	<code>boolean null</code>			
<code>isin</code>	<code>array<number> null</code>			
<code>isnotin</code>	<code>array<number> null</code>			

GenericFilter_int_

Field	Type	Req	Default	Description
eq	integer null			
gt	integer null			
ge	integer null			
lt	integer null			
le	integer null			
ne	integer null			
isnull	boolean null			
isin	array<integer> null			
isnotin	array<integer> null			

GenericFilter_str_

Field	Type	Req	Default	Description
eq	string null			
gt	string null			
ge	string null			
lt	string null			
le	string null			
ne	string null			
isnull	boolean null			
isin	array<string> null			
isnotin	array<string> null			

GenericFilter_UserType_

Field	Type	Req	Default	Description
eq	UserType null			
gt	UserType null			

ge	UserType null			
lt	UserType null			
le	UserType null			
ne	UserType null			
isnull	boolean null			
isin	array<UserType> null			
isnotin	array<UserType> null			

HTTPValidationError

Field	Type	Req	Default	Description
detail	array<ValidationError>			

InvalidResource

Details about a resource that failed validation.

Field	Type	Req	Default	Description
resource	ResourceIdentifier	yes		Identifier for an external resource. The meaning of fileid depends on the provider (determined by the sync config): - Google Drive: the Drive file id - SharePoint / OneDrive: the item id - S3: the obj...
reason	string	yes		

InviteKnowledgeBoxPayload

Field	Type	Req	Default	Description
-------	------	-----	---------	-------------

email	string	yes		
role	KbRole	yes		
came_from	string null			

ItemCreated

Field	Type	Req	Default	Description
id	string	yes		

KBDataResponse

Field	Type	Req	Default	Description
id	string	yes		The unique identifier of the knowledgebox.
slug	string	yes		A human-readable, URL-friendly identifier for the knowledgebox.
title	string	yes		The title of the knowledgebox.
created	string	yes		The timestamp when the knowledgebox was created.

KBExternalIndexProviderType

Enum: `unset`

KbInvitedUser

Field	Type	Req	Default	Description
email	string	yes		
role	KbRole	yes		
expires	string	yes		

KBMode

Knowledge Box modes define the architecture and components of a KB. - kb: KB in idp_regional + NucliaDB KB (same ID), no agent - agent: KB in idp_regional + NucliaDB KB (same ID) + Agent with memory (agent_id = kb_id) - agent_no_memory: KB in idp_regional + Agent without memory (agent_id = kb_id), n...

Enum: kb, agent, agent_no_memory

KbRole

Enum: SOWNER, SMEMBER, SCONTRIBUTOR

KbState

Enum: PUBLISHED, PRIVATE

KbUser

Field	Type	Req	Default	Description
id	string	yes		
role	KbRole	yes		
email	string null			
name	string null			

KnowledgeBox

Field	Type	Req	Default	Description
id	string	yes		
slug	string	yes		
account_id	string	yes		
zone	string	yes		
title	string	yes		
description	string null			
state	KbState	yes		
created	string	yes		
admin	boolean null			
contrib	boolean null			

<code>allowed_origins</code>	array<string> null			
<code>allowed_ip_addresses</code>	array<string> null			
<code>search_configs</code>	object (free-form map)		<code>{}</code>	
<code>external_index_provider</code>	string null			
<code>enforce_security</code>	boolean	yes		Whether security is enforced by default on all requests to this Knowledge Box.

KnowledgeBoxCreation

Field	Type	Req	Default	Description
<code>slug</code>	string	yes		
<code>title</code>	string	yes		
<code>description</code>	string null			
<code>mode</code>	KBMode		<code>kb</code>	Mode of the Knowledge Box (immutable after creation). 'kb': Standard KB with NucliaDB storage. 'agent': Agent with memory backed by NucliaDB. 'agentnomemory': Agent without memory (no NucliaDB KB crea...
<code>learning_configuration</code>	object (free-form map)			
<code>allowed_origins</code>	array<string> null			
<code>allowed_ip_addresses</code>	array<string> null			
<code>external_index</code>	DummyKBIndexProv			External index

<code>x_provider</code>	ider null			provider configuration for the Knowledge Box. If not set, the default NucliaDB's index will be used.
<code>search_configs</code>	object (free-form map)		<code>{}</code>	
<code>enforce_security</code>	boolean		<code>False</code>	Whether security is enforced by default on all requests to this Knowledge Box.

KnowledgeBoxExport

Field	Type	Req	Default	Description
<code>id</code>	string	yes		
<code>title</code>	string	yes		
<code>description</code>	string	yes		

KnowledgeBoxSummary

Field	Type	Req	Default	Description
<code>id</code>	string	yes		
<code>slug</code>	string	yes		
<code>zone</code>	string	yes		
<code>title</code>	string	yes		
<code>state</code>	KbState	yes		
<code>description</code>	string null			
<code>role_on_kb</code>	KbRole null			
<code>allowed_origins</code>	array<string> null			
<code>allowed_ip_addresses</code>	array<string> null			
<code>search_config</code>	object (free-form			

s	map) null			
external_index_provider	string null			
enforce_security	boolean	yes		Whether security is enforced by default on all requests to this Knowledge Box.

KnowledgeBoxUpdate

Field	Type	Req	Default	Description
slug	string null			
title	string null			
state	KbState null			
description	string null			
allowed_origins	array<string> null			
allowed_ip_addresses	array<string> null			
search_configs	object (free-form map) null			
enforce_security	boolean null			Whether security is enforced by default on all requests to this Knowledge Box. Leaving this as None means no change, while setting it to True or False will update the setting accordingly.

LabelInput

Field	Type	Req	Default	Description
labelset	string	yes		
label	string	yes		

LabelOutput

Field	Type	Req	Default	Description
<code>labelset</code>	string	yes		
<code>label</code>	string	yes		

LogEntryOutput

A single log entry from a sync job.

Field	Type	Req	Default	Description
<code>id</code>	integer	yes		Unique identifier for the log entry
<code>timestamp</code>	string	yes		Timestamp of the log entry
<code>level</code>	string	yes		Log level (INFO, WARNING, ERROR, etc.)
<code>message</code>	string	yes		Log message
<code>extra</code>	object (free-form map) null			Additional context data for the log entry

LogLevel

Log level for filtering sync job logs.

Enum: `DEBUG`, `INFO`, `WARNING`, `ERROR`, `EXCEPTION`, `CRITICAL`

ModifiedTimeRangeInput

Field	Type	Req	Default	Description
<code>from</code>	string null			
<code>to</code>	string null			

ModifiedTimeRangeOutput

Field	Type	Req	Default	Description
-------	------	-----	---------	-------------

from	string null			
to	string null			

NUAClient

Field	Type	Req	Default	Description
token	string	yes		
client_id	string	yes		
internal_id	string	yes		

NUAClientKeyCreation

Field	Type	Req	Default	Description
contact	string	yes		
title	string	yes		
description	string null			
client_id	string null			
processing_webhook	ProcessingWebhook null			
allow_kb_management	boolean		False	
tokens_limit	integer null			

NUAClientMetadata

Field	Type	Req	Default	Description
title	string	yes		
description	string null			
contact	string	yes		
created	string	yes		
partitions	integer	yes		
client_id	string	yes		

<code>internal_id</code>	string	yes		
<code>account_id</code>	string	yes		
<code>processing_webhook</code>	ProcessingWebhook null			
<code>tokens_limit</code>	integer null			

NUAClientsMetadata

Field	Type	Req	Default	Description
<code>clients</code>	array<NUAClientMetadata>	yes		

NUAClientUpdate

Field	Type	Req	Default	Description
<code>tokens_limit</code>	integer null			
<code>title</code>	string null			
<code>contact</code>	string null			

NUATokenMetadataUpdate

Field	Type	Req	Default	Description
<code>nua_internal_id</code>	string	yes		
<code>issuer</code>	string null			
<code>seen_at</code>	string	yes		

PaginatedLogsOutput

Paginated list of log entries.

Field	Type	Req	Default	Description
<code>items</code>	array<LogEntryOutput>	yes		

<code>next_cursor</code>	integer null			Cursor (log ID) to fetch the next page of results. None if no more results.
--------------------------	----------------	--	--	---

PaginatedSyncJobsOutput

Paginated list of sync jobs.

Field	Type	Req	Default	Description
<code>items</code>	array<SyncJobOutput>	yes		
<code>next_cursor</code>	string null			Cursor to fetch the next page of results. None if no more results.

Pagination

Field	Type	Req	Default	Description
<code>limit</code>	integer		10	
<code>starting_after</code>	integer null			
<code>ending_before</code>	integer null			

PermissionsResponse

Field	Type	Req	Default	Description
<code>permissions</code>	array<string>	yes		

ProcessingWebhook

Field	Type	Req	Default	Description
<code>uri</code>	string	yes		
<code>headers</code>	map<string, string>		{}	

Provider

Enum: `google_oauth`, `azure_oauth`, `azure_certificate_credentials`, `aws_s3_assume_role`, `sharefile_oauth`, `dropbox_oauth`

QueryFiltersAsk

Field	Type	Req	Default	Description
<code>id</code>	GenericFilter_int_ null			
<code>date</code>	GenericFilter_datetime_ null			
<code>user_id</code>	GenericFilter_str_ null			
<code>user_type</code>	GenericFilter_UserType_ null			
<code>client_type</code>	GenericFilter_ClientType_ null			
<code>total_duration</code>	GenericFilter_int_ null			
<code>audit_metadata</code>	array<AuditMetadata> null			
<code>nuclia_tokens</code>	GenericFilter_float_ null			
<code>question</code>	StringFilter null			
<code>rephrased_question</code>	StringFilter null			
<code>answer</code>	StringFilter null			
<code>learning_id</code>	StringFilter null			
<code>feedback_good</code>	GenericFilter_bool_ null			True if the feedback provided for the main question is positive.
<code>feedback_comment</code>	StringFilter null			User-provided comment on the feedback for the question.
<code>feedback_good_all</code>	GenericFilter_bool_ null			True if all feedback, including that on the main question and each related text

				block, is positive.
feedback_good_any	GenericFilter_bool_ null			True if there is any positive feedback on the question itself or any related text block.
model	StringFilter null			
status	GenericFilter_int_ null			
generative_answer_first_chunk_time	GenericFilter_int_ null			Time in milliseconds from when the user made the request to when the first answer chunk of data was returned by the generative model.
generative_reasoning_first_chunk_time	GenericFilter_int_ null			Time in milliseconds from when the user made the request to when the first reasoning chunk of data was returned by the generative model.
generative_answer_time	GenericFilter_int_ null			Time in milliseconds elapsed between the answer streaming request is done to the generative model until the last chunk of the answer is returned.
reasoning	StringFilter null			
resources_count	GenericFilter_int_ null			
vectorset	StringFilter null			
min_score_bm25	GenericFilter_float_ null			
min_score_semantic	GenericFilter_float_ null			
result_per_page	GenericFilter_int_ null			

retrieval_time	GenericFilter_int_ null			Time in milliseconds spent on the NucliaDB side for retrieval, including retrieval rephrase, calculation of query embedding, index search and results hydration
----------------	---------------------------	--	--	---

QueryFiltersChat

Field	Type	Req	Default	Description
id	GenericFilter_int_ null			
date	GenericFilter_datetime_ null			
user_id	GenericFilter_str_ null			
user_type	GenericFilter_UserType_ null			
client_type	GenericFilter_ClientType_ null			
total_duration	GenericFilter_int_ null			
audit_metadata	array<AuditMetadata> null			
nuclia_tokens	GenericFilter_float_ null			
question	StringFilter null			
rephrased_question	StringFilter null			
answer	StringFilter null			
learning_id	StringFilter null			
feedback_good	GenericFilter_bool_ null			True if the feedback provided for the main question is positive.
feedback_comm	StringFilter null			User-provided

ent				comment on the feedback for the question.
feedback_good_all	GenericFilter_bool_ null			True if all feedback, including that on the main question and each related text block, is positive.
feedback_good_any	GenericFilter_bool_ null			True if there is any positive feedback on the question itself or any related text block.
model	StringFilter null			
status	GenericFilter_int_ null			
generative_answer_first_chunk_time	GenericFilter_int_ null			Time in milliseconds from when the user made the request to when the first answer chunk of data was returned by the generative model.
generative_reasoning_first_chunk_time	GenericFilter_int_ null			Time in milliseconds from when the user made the request to when the first reasoning chunk of data was returned by the generative model.
generative_answer_time	GenericFilter_int_ null			Time in milliseconds elapsed between the answer streaming request is done to the generative model until the last chunk of the answer is returned.
reasoning	StringFilter null			

QueryFiltersCommon

Field	Type	Req	Default	Description
-------	------	-----	---------	-------------

<code>id</code>	GenericFilter_int_ null			
<code>date</code>	GenericFilter_datetim e_ null			
<code>user_id</code>	GenericFilter_str_ null			
<code>user_type</code>	GenericFilter_UserTy pe_ null			
<code>client_type</code>	GenericFilter_Client Type_ null			
<code>total_duratio n</code>	GenericFilter_int_ null			
<code>audit_metadat a</code>	array<AuditMetadata > null			
<code>nuclia_tokens</code>	GenericFilter_float_ null			

QueryFiltersSearch

Field	Type	Req	Default	Description
<code>id</code>	GenericFilter_int_ null			
<code>date</code>	GenericFilter_datetim e_ null			
<code>user_id</code>	GenericFilter_str_ null			
<code>user_type</code>	GenericFilter_UserTy pe_ null			
<code>client_type</code>	GenericFilter_Client Type_ null			
<code>total_duratio n</code>	GenericFilter_int_ null			
<code>audit_metadat a</code>	array<AuditMetadata > null			
<code>nuclia_tokens</code>	GenericFilter_float_ null			

<code>question</code>	StringFilter null			
<code>resources_count</code>	GenericFilter_int_ null			
<code>vectorset</code>	StringFilter null			
<code>min_score_bm25</code>	GenericFilter_float_ null			
<code>min_score_semantic</code>	GenericFilter_float_ null			
<code>result_per_page</code>	GenericFilter_int_ null			
<code>retrieval_time</code>	GenericFilter_int_ null			Time in milliseconds spent on the NucliaDB side for retrieval, including retrieval rephrase, calculation of query embedding, index search and results hydration

RegionalTokenMetadataPayload

Field	Type	Req	Default	Description
<code>sa_updates</code>	array<SATokenMetadataUpdate>		<code>[]</code>	
<code>nua_updates</code>	array<NUATokenMetadataUpdate>		<code>[]</code>	

RemiAnswerRelevance

Field	Type	Req	Default	Description
<code>score</code>	integer	yes		
<code>reason</code>	string	yes		

RemiQuery

Field	Type	Req	Default	Description
-------	------	-----	---------	-------------

context_relevance	ContextRelevanceQuery null			
month	string	yes		
from_date	string null			
to_date	string null			
feedback_good	boolean null			
status	Status null			
pagination	Pagination		{'limit': 10}	

RemiQueryResult

Field	Type	Req	Default	Description
id	integer	yes		
question	string	yes		
date	string	yes		
answer	string null	yes		
remi	RemiScores null	yes		

RemiQueryResults

Field	Type	Req	Default	Description
data	array<RemiQueryResult>	yes		
has_more	boolean	yes		

RemiQueryResultWithContext

Field	Type	Req	Default	Description
id	integer	yes		
question	string	yes		
date	string	yes		
answer	string null	yes		

<code>remi</code>	RemiScores null	yes		
<code>context</code>	array<RetrievedCont ext>	yes		

RemiScores

Field	Type	Req	Default	Description
<code>answer_relevance</code>	RemiAnswerRelevance	yes		
<code>context_relevance</code>	array<integer>	yes		
<code>groundedness</code>	array<integer>	yes		

ResourceIdentifier

Identifier for an external resource. The meaning of `file_id` depends on the provider (determined by the sync config): - Google Drive: the Drive file id - SharePoint / OneDrive: the item id - S3: the object key - ShareFile: the item id - Dropbox: the full file path including filename (e.g. /Docu...

Field	Type	Req	Default	Description
<code>file_id</code>	string	yes		

ResourceValidationInput

Input schema for validating multiple resources with existing external connection. Credentials should be a Fernet-encrypted string containing the JSON-serialized credentials dict, encrypted with the same key used for external connections.

Field	Type	Req	Default	Description
<code>credentials</code>	string	yes		
<code>resources</code>	array<ResourceIdentifier>	yes		

ResourceValidationOutput

Output schema for validated resources.

Field	Type	Req	Default	Description
-------	------	-----	---------	-------------

<code>valid_resources</code>	array<ResourceIdentifier>	yes		
<code>invalid_resources</code>	array<InvalidResource>	yes		

RetrievedContext

Field	Type	Req	Default	Description
<code>text</code>	string	yes		
<code>text_block_id</code>	string null			

S3AssumeRoleInfoOutput

Information the client needs to configure an IAM role for Nuclia S3 sync.

Field	Type	Req	Default	Description
<code>external_id</code>	string	yes		The external ID to use in the IAM role trust policy condition.
<code>role_name</code>	string	yes		The role name that should be created in the client's AWS account.
<code>aws_account_id</code>	string	yes		Nuclia's AWS account ID that will assume the role.

SATokenMetadataUpdate

Field	Type	Req	Default	Description
<code>sa_key_id</code>	string	yes		
<code>issuer</code>	string null			
<code>seen_at</code>	string	yes		

ServiceAccount

Field	Type	Req	Default	Description
<code>id</code>	string	yes		
<code>title</code>	string	yes		
<code>role</code>	KbRole	yes		
<code>keys</code>	array<ServiceAccountKey> null			

ServiceAccountCreation

Field	Type	Req	Default	Description
<code>title</code>	string	yes		
<code>role</code>	KbRole	yes		

ServiceAccountJWT

Field	Type	Req	Default	Description
<code>id</code>	string	yes		
<code>token</code>	string	yes		

ServiceAccountKey

Field	Type	Req	Default	Description
<code>id</code>	string	yes		
<code>created</code>	string	yes		
<code>expires</code>	string	yes		

ServiceAccountKeyCreation

Field	Type	Req	Default	Description
<code>expires</code>	string	yes		

ShowOnly

Marker class to indicate that a field is for display purposes only. Fields wrapped in this class are excluded from OpenAPI documentation to prevent confusion, as they are not intended for filtering. Instead, they are only used in the "show" parameter. This behavior is enforced in the FastAPI applica...

Type: `object`

SortOrder

Sort order for paginated queries.

Enum: `asc`, `desc`

Status

Enum: `SUCCESS`, `ERROR`, `NO_CONTEXT`

StorageDriveOutput

Represents a drive/storage location that can be selected for syncing.

Field	Type	Req	Default	Description
<code>id</code>	string	yes		
<code>name</code>	string	yes		
<code>drive_type</code>	string null			
<code>web_url</code>	string null			

StorageFolderOutput

Represents a folder that can be selected for syncing.

Field	Type	Req	Default	Description
<code>id</code>	string	yes		
<code>name</code>	string	yes		
<code>path</code>	string	yes		
<code>web_url</code>	string null			

StorageSiteOutput

Represents a SharePoint site that can be selected for browsing drives.

Field	Type	Req	Default	Description
-------	------	-----	---------	-------------

<code>id</code>	string	yes		
<code>name</code>	string	yes		
<code>web_url</code>	string null			

StringFilter

Field	Type	Req	Default	Description
<code>eq</code>	string null			
<code>gt</code>	string null			
<code>ge</code>	string null			
<code>lt</code>	string null			
<code>le</code>	string null			
<code>ne</code>	string null			
<code>isnull</code>	boolean null			
<code>isin</code>	array<string> null			
<code>isnotin</code>	array<string> null			
<code>like</code>	string null			
<code>ilike</code>	string null			

SyncConfigInput

Field	Type	Req	Default	Description
<code>name</code>	string	yes		
<code>sync_root_path</code>	string null			
<code>folder_id</code>	string null			
<code>external_connection_id</code>	string	yes		
<code>drive_id</code>	string null			
<code>file_filter</code>	FileFilterInput null			

labels	array<LabelInput> null			
modified_time_range	ModifiedTimeRangeInput null			
extract_strategy	string null			

SyncConfigOutput

Field	Type	Req	Default	Description
id	string	yes		
name	string	yes		
kb_id	string	yes		
created_at	string	yes		
updated_at	string	yes		
sync_root_path	string	yes		
sync_interval_minutes	integer	yes		
created_by	string	yes		
external_connection	ExternalConnectionOutput	yes		External connection output without sensitive credential data.
drive_id	string null			
file_filter	FileFilterOutput null			
labels	array<LabelOutput> null			
modified_time_range	ModifiedTimeRangeOutput null			
extract_strategy	string null			
last_sync_run	string null			

SyncConfigUpdateInput

Partial update model for SyncConfig. Omitted fields are left unchanged. Fields explicitly set to null clear that entry from extra_configuration.

Field	Type	Req	Default	Description
<code>name</code>	string null			
<code>file_filter</code>	FileFilterInput null			
<code>labels</code>	array<LabelInput> null			
<code>modified_time_range</code>	ModifiedTimeRangeInput null			
<code>extract_strategy</code>	string null			

SyncJobOptionsOutput

Options for a sync job.

Field	Type	Req	Default	Description
<code>full_sync</code>	boolean		<code>False</code>	

SyncJobOutput

Field	Type	Req	Default	Description
<code>id</code>	string	yes		
<code>created_at</code>	string	yes		
<code>finished_at</code>	string null	yes		
<code>config_id</code>	string	yes		
<code>status</code>	SyncJobStatus	yes		
<code>options</code>	SyncJobOptionsOutput		<code>{'full_sync': False}</code>	Options for a sync job.

SyncJobStatus

Enum: `pending`, `in_progress`, `completed`, `failed`

TemporalAgentKeyCreation

Field	Type	Req	Default	Description
<code>agent_session</code>	string	yes		
<code>ttl</code>	integer		10	

TemporalSecurityTokenResponse

Field	Type	Req	Default	Description
<code>token</code>	string	yes		

TemporalServiceAccountKeyCreation

Field	Type	Req	Default	Description
<code>ttl</code>	integer		10	
<code>security_groups</code>	array<string> null			

TriggerSyncInput

Input for triggering a sync job.

Field	Type	Req	Default	Description
<code>full_sync</code>	boolean		False	If True, ignore incremental sync state and perform a full sync. Note: S3 and ShareFile providers always perform full syncs as they don't support delta tracking.

UpdateKbUser

Field	Type	Req	Default	Description
<code>id</code>	string	yes		

role	KbRole	yes		
------	--------	-----	--	--

UserType

Enum: `user`, `nuakey`

ValidationError

Field	Type	Req	Default	Description
loc	array<string integer>	yes		
msg	string	yes		
type	string	yes		
input	object			
ctx	object			

global — 195 schemas

Account

Field	Type	Req	Default	Description
id	string	yes		
slug	string	yes		
title	string	yes		
type	AccountType	yes		
zone_visibility	ZoneVisibility	yes		
description	string null			
can_manage_account	boolean	yes		
current_users	integer	yes		
max_users	integer null			
current_kbs	integer	yes		

max_kbs	integer null			
current_agents	integer	yes		
max_agents	integer null			
current_memories	integer	yes		
max_memories	integer null			
current_arags	integer	yes		
max_arags	integer null			
blocking_state	AccountBlockingState		unblocked	
blocked_features	array<BlockedFeature>		[]	
limits	AccountLimits null			
creation_date	string	yes		
trial_expiration_date	string null			
domain	string null			
saml_entity_id	string null			
saml_sso_url	string null			
saml_x509_certificate	string null			
saml_config	SAMLConfig null			
workflow	WorkflowMode		classic	
eula_accepted	boolean		False	

AccountAWSMarketplaceSubscriptionInfo

Field	Type	Req	Default	Description
status	AWSMarketplaceSubscriptionStatus	yes		
on_demand_bud	number null			

get				
action_on_budget_exhausted	ActionOnBudgetExhausted null			
aws_product_code	string	yes		
aws_product_id	string	yes		
free_tokens_per_billing_cycle	integer	yes		
free_storage_nuclia_tokens	integer null			
billing_formula_id	string	yes		

AccountBillingConfig

Field	Type	Req	Default	Description
account_id	string	yes		
billing_formula_id	string	yes		

AccountBlockingState

Enum: unblocked, quota_manager, trial_expired, budget, subscription_canceled, subscription_unpaid

AccountCreated

Field	Type	Req	Default	Description
id	string	yes		

AccountCreation

Field	Type	Req	Default	Description
slug	string null			

<code>title</code>	string	yes		
<code>description</code>	string null			
<code>email</code>	string null			
<code>zone</code>	string null			
<code>workflow</code>	WorkflowMode		<code>classic</code>	
<code>eula_accepted</code>	boolean		<code>False</code>	

AccountId

Field	Type	Req	Default	Description
<code>id</code>	string	yes		
<code>slug</code>	string	yes		

AccountingTarget

Accounting target identifies a specific method of calculating a value for a stripe price from accounting metrics, but uncoupled from both price and metrics concepts. Originally, the purpose of target was intended to be filled in by the value of a specific metric in accounting. Given that accounting ...

Enum: `media`, `training`, `paragraphs_processed`, `predict`, `generative`, `ai-tokens-used`, `nuclia-tokens`, `paragraphs`, `searches`, `qa`, `self-hosted-qa`, `self-hosted-predict`, `self-hosted-processed-paragraphs`, `self-hosted-processed-documents`

AccountLabels

Field	Type	Req	Default	Description
<code>progress_account</code>	boolean null			

AccountLimits

Field	Type	Req	Default	Description
<code>upload</code>	AccountUploadLimits null			
<code>usage</code>	AccountUsageLimits null			

processing	AccountUsageLimits null			
------------	------------------------------	--	--	--

AccountListItem

Field	Type	Req	Default	Description
id	string	yes		
slug	string	yes		
title	string	yes		
type	AccountType	yes		
zone_visibility	ZoneVisibility	yes		

AccountManager

Field	Type	Req	Default	Description
id	string	yes		
slug	string	yes		
title	string null			
type	AccountType	yes		
creator	string	yes		
email	string	yes		
created	string	yes		
modified	string null			
avatar	string null			
labels	AccountLabels null			
zone_visibility	ZoneVisibility	yes		
blocked_features	array<BlockedFeature>	yes		
blocking_state	AccountBlockingState	yes		

<code>users</code>	<code>array<stashify_idp__api__manage__models__User></code>	yes		
<code>members</code>	<code>array<string></code>	yes		
<code>managers</code>	<code>array<string></code>	yes		
<code>stashes</code>	<code>AccountManagerStashes</code>	yes		
<code>max_kbs</code>	<code>integer</code>	yes		
<code>max_agents</code>	<code>integer</code>	yes		
<code>max_memories</code>	<code>integer</code>	yes		
<code>max_arags</code>	<code>integer</code>	yes		
<code>arags</code>	<code>integer</code>	yes		
<code>limits</code>	<code>AccountManagerLimits null</code>			
<code>trial_expiration_date</code>	<code>string null</code>			
<code>allow_access_non_enterprise_models</code>	<code>boolean</code>		<code>True</code>	
<code>saml_config</code>	<code>SAMLConfig null</code>			
<code>workflow</code>	<code>WorkflowMode</code>		<code>classic</code>	

AccountManagerLimits

Field	Type	Req	Default	Description
<code>upload</code>	<code>AccountManagerUploadLimits null</code>			
<code>usage</code>	<code>AccountManagerUsageLimits null</code>			

AccountManagerPatch

Field	Type	Req	Default	Description
<code>slug</code>	<code>string null</code>			

email	string null			
creator	string null			
type	AccountType null			
kbs	integer null			
agents	integer null			
arags	integer null			
memories	integer null			
data	object (free-form map) null			
saml_config	SAMLConfig null			
limits	AccountManagerLimits null			
trial_expiration_date	string null			
allow_access_non_enterprise_models	boolean null			
workflow	WorkflowMode null			
labels	AccountLabels null			
zone_visibility	ZoneVisibility null			

AccountManagerSetBlocking

Field	Type	Req	Default	Description
blocking_state	AccountBlockingState	yes		
blocked_features	array<BlockedFeature> null			

AccountManagerStashes

Field	Type	Req	Default	Description
max_stashes	integer	yes		

items	array<Stash>	yes		
-------	--------------	-----	--	--

AccountManagerSummary

Field	Type	Req	Default	Description
id	string	yes		
slug	string	yes		
title	string null			
type	AccountType	yes		
creator	string	yes		
email	string	yes		
created	string	yes		
modified	string null			
avatar	string null			
labels	AccountLabels null			
zone_visibility	ZoneVisibility	yes		

AccountManagerUploadLimits

Field	Type	Req	Default	Description
upload_limit_max_media_file_size	integer null			
upload_limit_max_non_media_file_size	integer null			

AccountManagerUsageLimits

Field	Type	Req	Default	Description
monthly_limit_paragraphs_p	integer null			

processed				
monthly_limit_docs_no_media_processed	integer null			
monthly_limit_media_seconds_processed	integer null			
monthly_limit_paragraphs_stored	integer null			
monthly_limit_hosted_searches_performed	integer null			
monthly_limit_hosted_answers_generated	integer null			
monthly_limit_self_hosted_searches_performed	integer null			
monthly_limit_self_hosted_answers_generated	integer null			
storage_limit_max_bytes_per_kb	integer null			
storage_limit_max_resources_per_kb	integer null			

AccountMeteredUsage

Field	Type	Req	Default	Description
budget	number null			
currency	Currency	yes		
invoice_items	InvoiceItems	yes		

start_billing_date	string	yes		
end_billing_date	string	yes		

AccountModification

Field	Type	Req	Default	Description
title	string null			
description	string null			
saml	OldSAMLConfig null			
saml_config	SAMLConfig null			
slug	string null			
workflow	WorkflowMode null			
eula_accepted	boolean null			

AccountNucliaCustomerTableSubscriptionInfo

Field	Type	Req	Default	Description
status	NucliaCustomerTableSubscriptionStatus	yes		
on_demand_budget	number null			
action_on_budget_exhausted	ActionOnBudgetExhausted null			
free_tokens_per_billing_cycle	integer	yes		
free_storage_nuclia_tokens	integer null			
billing_formula_id	string	yes		
nuclia_tokens_price	number	yes		

AccountRole**Enum:** AOWNER, AMEMBER**AccountStatus**

Field	Type	Req	Default	Description
blocking_state	AccountBlockingState		unblocked	
blocked_features	array<BlockedFeature>		[]	
requested_account_type	AccountType null			

AccountStatusAvailable

Field	Type	Req	Default	Description
available	boolean	yes		

AccountStripeCustomer

Field	Type	Req	Default	Description
customer_id	string	yes		
billing_details	BillingDetails	yes		

AccountStripeSubscription

Field	Type	Req	Default	Description
subscription_id	string	yes		
payment_method_id	string	yes		
currency	Currency	yes		
status	string	yes		

<code>requires_action</code>	boolean	yes		
<code>client_secret</code>	string null			
<code>error</code>	PaymentError null			

AccountStripeSubscriptionInfo

Field	Type	Req	Default	Description
<code>status</code>	SubscriptionStatus	yes		
<code>on_demand_budget</code>	number null			
<code>action_on_budget_exhausted</code>	ActionOnBudgetExhausted null			
<code>billing_interval</code>	IntervalType null		month	
<code>start_billing_period</code>	string	yes		
<code>end_billing_period</code>	string	yes		
<code>free_tokens_per_billing_cycle</code>	integer	yes		
<code>free_storage_nuclia_tokens</code>	integer null			
<code>billing_formula_id</code>	string	yes		

AccountSubscription

Field	Type	Req	Default	Description
<code>provider</code>	SubscriptionProvider	yes		
<code>subscription</code>	AccountStripeSubscriptionInfo AccountAWSMarketplaceSubscriptionInfo	yes		

	AccountNucliaCusto merTableSubscriptio nInfo			
--	--	--	--	--

AccountType

Enum: `stash-trial`, `v3starter`, `v3pro`, `v3enterprise`, `cowork`, `v3fly`, `v3growth`

AccountTypeDefaults

Field	Type	Req	Default	Description
<code>max_kbs</code>	integer	yes		
<code>max_agents</code>	integer	yes		
<code>max_memories</code>	integer	yes		
<code>max_trial_days</code>	integer		-1	
<code>monthly_limit_paragraphs_processed</code>	integer		-1	
<code>monthly_limit_docs_no_media_processed</code>	integer		-1	
<code>monthly_limit_media_seconds_processed</code>	integer		-1	
<code>monthly_limit_paragraphs_stored</code>	integer		-1	
<code>monthly_limit_hosted_searches_performed</code>	integer		-1	
<code>monthly_limit_hosted_answers_generated</code>	integer		-1	
<code>monthly_limit_self_hosted_searches_performed</code>	integer		-1	
<code>monthly_limit</code>	integer		-1	

<code>_self_hosted_answers_generated</code>				
<code>storage_limit_max_bytes_per_kb</code>	integer		-1	
<code>storage_limit_max_resources_per_kb</code>	integer		-1	
<code>upload_limit_max_media_file_size</code>	integer	yes		
<code>upload_limit_max_non_media_file_size</code>	integer	yes		
<code>allow_access_non_enterprise_models</code>	boolean		True	

AccountUploadLimits

Field	Type	Req	Default	Description
<code>upload_limit_max_media_file_size</code>	integer null			
<code>upload_limit_max_non_media_file_size</code>	integer null			

AccountUsageLimits

Field	Type	Req	Default	Description
<code>monthly_limit_paragraphs_processed</code>	integer null			
<code>monthly_limit_docs_no_media_processed</code>	integer null			

<code>monthly_limit_media_seconds_processed</code>	integer null			
<code>monthly_limit_paragraphs_stored</code>	integer null			
<code>monthly_limit_hosted_searches_performed</code>	integer null			
<code>monthly_limit_hosted_answers_generated</code>	integer null			
<code>monthly_limit_self_hosted_searches_performed</code>	integer null			
<code>monthly_limit_self_hosted_answers_generated</code>	integer null			
<code>storage_limit_max_bytes_per_kb</code>	integer null			
<code>storage_limit_max_resources_per_kb</code>	integer null			
<code>monthly_limit_non_media_files_processed</code>	integer null			
<code>monthly_limit_chars_processed</code>	integer null			

AccountUser

Field	Type	Req	Default	Description
<code>id</code>	string	yes		
<code>name</code>	string	yes		

email	string	yes		
role	AccountRole	yes		

AccountUserCreation

Field	Type	Req	Default	Description
id	string	yes		

AccountUserPatch

Field	Type	Req	Default	Description
type	string	yes		

AccountZoneEntry

Field	Type	Req	Default	Description
account_id	string	yes		
zone_id	string	yes		

AccountZoneGrant

Field	Type	Req	Default	Description
account_id	string	yes		

AccountZoneInfo

Field	Type	Req	Default	Description
account_type	AccountType	yes		
account_email	string	yes		
account_slug	string	yes		
account_title	string	yes		

<code>blocking_state</code>	AccountBlockingState	yes		
<code>blocked_features</code>	array<BlockedFeature>	yes		
<code>upload_limits</code>	AccountUploadLimits	yes		
<code>allow_access_non_enterprise_models</code>	boolean		True	

ActionId

Enum: `join_account`, `join_kb`, `go_account`, `edit`, `create_template`, `create`, `edit_template`, `signup_create_user`

ActionOnBudgetExhausted

Enum: `WARN_ACCOUNT_OWNER`, `BLOCK_ACCOUNT`

AddAccountUser

Field	Type	Req	Default	Description
<code>id</code>	string	yes		
<code>role</code>	AccountRole	yes		

Aggregation

Enum: `hour`, `day`, `week`, `month`, `quarter`, `year`, `millennium`

Agreement-Input

Field	Type	Req	Default	Description
<code>id</code>	string	yes		
<code>status</code>	string	yes		
<code>start_time</code>	string	yes		
<code>expiration_date</code>	string	yes		
<code>offer</code>	Offer	yes		
<code>entitlement_t</code>	map<string,		<code>{}</code>	

erms	EntitlementTerm-Input>			
metered_terms	map<string, MeteredTerm-Input>		{}	

Agreement-Output

Field	Type	Req	Default	Description
id	string	yes		
status	string	yes		
start_time	string	yes		
expiration_date	string	yes		
offer	Offer	yes		
entitlement_terms	map<string, EntitlementTerm-Output>		{}	
metered_terms	map<string, MeteredTerm-Output>		{}	

AWSMarketplaceCustomerData

Field	Type	Req	Default	Description
first_name	string	yes		
last_name	string	yes		
owner_email_address	string	yes		
phone	string	yes		
company	string	yes		
country	string	yes		
use_case	string	yes		
organization_size	string	yes		

<code>role</code>	string	yes		
<code>receive_updates</code>	boolean	yes		

AWSMarketplaceLoginRequest

Field	Type	Req	Default	Description
<code>customer_token</code>	string	yes		

AWSMarketplaceLoginResponse

Field	Type	Req	Default	Description
<code>access_token</code>	string	yes		
<code>refresh_token</code>	string	yes		

AWSMarketplaceSetupAccountResponse

Field	Type	Req	Default	Description
<code>account_id</code>	string	yes		
<code>access_token</code>	string	yes		
<code>refresh_token</code>	string	yes		

AWSMarketplaceSetupRequest

Field	Type	Req	Default	Description
<code>customer_token</code>	string	yes		
<code>account_slug</code>	string null			
<code>customer_data</code>	AWSMarketplaceCustomerData null			

AWSMarketplaceSubscriptionStatus

Enum: `no_subscription`, `active`, `failed`, `cancel_scheduled`, `canceled`

BatchUpdateAccountUsers

Field	Type	Req	Default	Description
<code>add</code>	array<AddAccountUser> null			
<code>delete</code>	array<string> null			

BillCurrentPeriodResponse

Field	Type	Req	Default	Description
<code>budget</code>	number null			
<code>nuclia_tokens</code>	number	yes		
<code>start_billing_date</code>	string	yes		
<code>end_billing_date</code>	string	yes		
<code>free_storage_nuclia_tokens</code>	number null	yes		
<code>free_tokens_per_billing_cycle</code>	integer	yes		
<code>nuclia_tokens_price</code>	number null	yes		

BillingDetails

Field	Type	Req	Default	Description
<code>name</code>	string	yes		
<code>email</code>	string	yes		
<code>address</code>	string	yes		
<code>country</code>	Country	yes		
<code>state</code>	string null			

city	string	yes		
postal_code	string	yes		
phone	string	yes		
is_company	boolean	yes		
company	string null			
vat	string null			

BillingDetailsModified

Field	Type	Req	Default	Description
name	string null			
email	string null			
address	string null			
state	string null			
city	string null			
postal_code	string null			
phone	string null			
is_company	boolean null			
company	string null			
vat	string null			

BillingFormulaCreate

Field	Type	Req	Default	Description
title	string	yes		
description	string null			
formula	array<string>		[]	
absolute_formula	array<string>		[]	

BillingFormulaResponse

Field	Type	Req	Default	Description
id	string	yes		
created	string	yes		
title	string	yes		
description	string	yes		
formula	array<string>	yes		
absolute_formula	array<string>	yes		

BlockedFeature

Enum: upload, processing, search, generative, training, public_upload, public_processing, public_search, public_generative

Body_acs_api_auth_saml_acs_post

Field	Type	Req	Default	Description
SAMLResponse	string	yes		
RelayState	string	yes		

Body_aws_marketplace_fulfillment_callback_api_marketplace_aws_callback_post

Field	Type	Req	Default	Description
x-amzn-marketplace-token	string	yes		

Body_oauth_login_api_auth_oauth_login_post

Field	Type	Req	Default	Description
login_challenge	string	yes		
username	string null			

password	string null			
----------	---------------	--	--	--

Body_signup_start_api_auth_signup_start_post

Field	Type	Req	Default	Description
email	string	yes		
app	string	yes		
fullname	string null			
company	string null			
workflow	WorkflowMode null			

Body_submit_consent_api_auth_oauth_consent_post

Field	Type	Req	Default	Description
consent_challenge	string	yes		
grant_scope	array<string>		[]	
reject	boolean		False	

Body_upload_account_logo_api_v1_account__account_slug__logo_png_patch

Field	Type	Req	Default	Description
file	string	yes		

CalculatorConfig

Field	Type	Req	Default	Description
version	string	yes		
models	map<string, RawFactor ModelFactor>	yes		
storage	map<string, RawFactor>	yes		

CancelStripeSubscription

Field	Type	Req	Default	Description
<code>feedback</code>	StripeSubscriptionCancellationFeedback null			
<code>comment</code>	string null			

CloudProvider

Enum: `AWS`, `GCP`, `ON_PREMISES`

CloudZeroCustomerData

Field	Type	Req	Default	Description
<code>account_id</code>	string	yes		
<code>created</code>	string null			
<code>modified</code>	string null			
<code>billing_formula_id</code>	string		00000000-0000-4000-8000-000000000001	
<code>free_storage_nuclia_tokens</code>	integer null			
<code>free_tokens_per_billing_cycle</code>	integer		0	
<code>nuclia_tokens_price</code>	number	yes		
<code>action_on_budget_get_exhausted</code>	ActionOnBudgetExhausted null			
<code>on_demand_budget</code>	number null			
<code>billing_history</code>	object (free-form map)			
<code>last_notifications</code>	LastNotifications			
<code>product</code>	string	yes		

product_line	string	yes		
business_unit	string	yes		

CloudZeroSubscriptionCreate

Field	Type	Req	Default	Description
nuclia_tokens_price	number	yes		
billing_formula_id	string		00000000 - 0000-4000 - 8000 - 000000000002	
free_storage_nuclia_tokens	integer		0	
free_tokens_per_billing_cycle	integer		0	
on_demand_budget	number null			
action_on_budget_exhausted	ActionOnBudgetExhausted null			
product	string	yes		
product_line	string	yes		
business_unit	string	yes		

CloudZeroSubscriptionUpdate

Field	Type	Req	Default	Description
billing_formula_id	string null			
nuclia_tokens_price	number null			
free_storage_nuclia_tokens	integer null			
free_tokens_p	integer null			

er_billing_cycle				
on_demand_budget	number null			
action_on_budget_exhausted	ActionOnBudgetExhausted null			

Country

Enum: AF, AX, AL, DZ, AD, AO, AI, AQ, AG, AR, AM, AW, AC, AU, AT, AZ, BS, BH, BD, BB, BY, BE, BZ, BJ, BM, BT, BO, BA, BW, BV, BR, IO, VG, BN, BG, BF, BI, KH, CM, CA, CV, BQ, KY, CF, TD, CL, CN, CO, KM, CG, CD, CK, CR, CI, HR, CW, CY, CZ, DK, DJ, DM, DO, EC, EG, SV, GQ, ER, EE, SZ, ET, FK, FO, FJ, FI, FR, GF, PF, TF, GA, GM, GE, DE, GH, GI, GR, GL, GD, GP, GU, GT, GG, GN, GW, GY, HT, HN, HK, HU, IS, IN, ID, IQ, IE, IM, IL, IT, JM, JP, JE, JO, KZ, KE, KI, XK, KW, KG, LA, LV, LB, LS, LR, LY, LI, LT, LU, MO, MG, MW, MY, MV, ML, MT, MQ, MR, MU, YT, MX, MD, MC, MN, ME, MS, MA, MZ, MM, NA, NR, NP, NL, NC, NZ, NI, NE, NG, NU, MK, NO, OM, PK, PS, PA, PG, PY, PE, PH, PN, PL, PT, PR, QA, RE, RO, RU, RW, WS, SM, ST, SA, SN, RS, SC, SL, SG, SX, SK, SI, SB, SO, ZA, GS, KR, SS, ES, LK, BL, SH, KN, LC, MF, PM, VC, SR, SJ, SE, CH, TW, TJ, TZ, TH, TL, TG, TK, TO, TT, TA, TN, TR, TM, TC, TV, UG, UA, AE, GB, US, UY, UZ, VU, VA, VE, VN, WF, EH, YE, ZM, ZW

CountryCurrency

Field	Type	Req	Default	Description
currency	Currency	yes		

CreatePaymentMethod

Field	Type	Req	Default	Description
token	string	yes		

CreateStripeCustomer

Field	Type	Req	Default	Description
billing_details	BillingDetails	yes		

CreateStripePaymentLink

Field	Type	Req	Default	Description
-------	------	-----	---------	-------------

account_id	string	yes		
account_type	AccountType	yes		
price_ids	array<string>	yes		
billing_formula_id	string	yes		
allow_promotion_codes	boolean		False	
currency	Currency		usd	

CreateStripeSubscription

Field	Type	Req	Default	Description
payment_method_id	string	yes		
on_demand_budget	number null			
action_on_budget_exhausted	ActionOnBudgetExhausted null			
billing_interval	IntervalType null		month	
account_type	AccountType	yes		

CreateTokenRequest

Field	Type	Req	Default	Description
description	string	yes		
expiration_date	string null			

Currency

Enum: eur, usd

EntitlementTerm-Input

Field	Type	Req	Default	Description
price	number string null			
duration	string	yes		
trial	boolean		False	

EntitlementTerm-Output

Field	Type	Req	Default	Description
price	string null			
duration	string	yes		
trial	boolean		False	

FlatIdentifier

Field	Type	Req	Default	Description
type	string	yes		
service	string	yes		
model	string null			
version	string	yes		

GlobalAccountUser

Field	Type	Req	Default	Description
user_id	string	yes		
email	string null			
account_role	AccountRole null			
account_slug	string null			
name	string null			

GlobalMagicToken

Field	Type	Req	Default	Description
token	string	yes		

HTTPValidationError

Field	Type	Req	Default	Description
detail	array<ValidationError>			

IntervalType

Enum: month, year

InvitedUser

Field	Type	Req	Default	Description
email	string	yes		
role	AccountRole	yes		
expires	string	yes		

InvitePayload

Field	Type	Req	Default	Description
email	string	yes		
role	AccountRole		AMEMBER	
came_from	string null			

InvoiceItems

Field	Type	Req	Default	Description
media	nucia_accounting__models__models__InvoiceItem null			

training	nuclia_accounting__models__models__In voiceItem null			
paragraphs_processed	nuclia_accounting__models__models__In voiceItem null			
predict	nuclia_accounting__models__models__In voiceItem null			
generative	nuclia_accounting__models__models__In voiceItem null			
searches	nuclia_accounting__models__models__In voiceItem null			
paragraphs	nuclia_accounting__models__models__In voiceItem null			
ai-tokens-used	nuclia_accounting__models__models__In voiceItem null			
nuclia-tokens	nuclia_accounting__models__models__In voiceItem null			
qa	nuclia_accounting__models__models__In voiceItem null			
self-hosted-qa	nuclia_accounting__models__models__In voiceItem null			
self-hosted-predict	nuclia_accounting__models__models__In voiceItem null			
self-hosted-processed-paragraphs	nuclia_accounting__models__models__In voiceItem null			
self-hosted-processed-documents	nuclia_accounting__models__models__In voiceItem null			

InvoicesList

Field	Type	Req	Default	Description
<code>items</code>	array<nuclia_accounting__api__models__InvoiceItem>	yes		
<code>pagination</code>	InvoicesListPagination	yes		

InvoicesListPagination

Field	Type	Req	Default	Description
<code>starting_after</code>	string null			
<code>limit</code>	integer null			
<code>has_more</code>	boolean	yes		

ItemCreated

Field	Type	Req	Default	Description
<code>id</code>	string	yes		

KBMode

Knowledge Box modes define the architecture and components of a KB. - kb: KB in idp_regional + NucliaDB KB (same ID), no agent - agent: KB in idp_regional + NucliaDB KB (same ID) + Agent with memory (agent_id = kb_id) - agent_no_memory: KB in idp_regional + Agent without memory (agent_id = kb_id), n...

Enum: `kb`, `agent`, `agent_no_memory`

KbState

Enum: `PUBLISHED`, `PRIVATE`

KbUsageDashboard

Field	Type	Req	Default	Description
<code>timestamp</code>	string	yes		

<code>metrics</code>	<code>array<Metric></code>	yes		
----------------------	----------------------------------	-----	--	--

KnowledgeBoxIndex

Field	Type	Req	Default	Description
<code>kb_id</code>	string	yes		
<code>account_id</code>	string	yes		
<code>zone_id</code>	string	yes		
<code>kb_mode</code>	KBMode	yes		Knowledge Box modes define the architecture and components of a KB. - kb: KB in idpregional + NucliaDB KB (same ID), no agent - agent: KB in idpregional + NucliaDB KB (same ID) + Agent with memory (ag...

KnowledgeBoxIndexCreationPayload

Field	Type	Req	Default	Description
<code>kb_id</code>	string	yes		
<code>account_id</code>	string	yes		
<code>kb_mode</code>	KBMode	yes		Knowledge Box modes define the architecture and components of a KB. - kb: KB in idpregional + NucliaDB KB (same ID), no agent - agent: KB in idpregional + NucliaDB KB (same ID) + Agent with memory (ag...

Languages

Enum: CA, ES, EN, FR, DE, PT

LastNotifications

Field	Type	Req	Default	Description
notified_approaching_budget	string null			
notified_budget_exhausted	string null			

LoginPayload

Field	Type	Req	Default	Description
username	string	yes		
password	string	yes		

LoginProvider

Enum: GOOGLE, APPLE, LOCAL, SAML, GITHUB, AWS_MARKETPLACE, MICROSOFT

ManageAccountCreation

Field	Type	Req	Default	Description
slug	string	yes		
title	string	yes		
description	string null			
email	string	yes		
zone	string null			
workflow	WorkflowMode		classic	
eula_accepted	boolean		False	
type	AccountType	yes		
zone_visibility	ZoneVisibility null			

ManageUser

Field	Type	Req	Default	Description
<code>id</code>	string	yes		
<code>name</code>	string null			
<code>email</code>	string	yes		
<code>language</code>	string null			
<code>type</code>	UserType	yes		
<code>providers</code>	array<LoginProvider>		<code>[]</code>	
<code>data</code>	object	yes		
<code>last_login</code>	string null			

ManageUserPatch

Field	Type	Req	Default	Description
<code>type</code>	UserType null			
<code>email</code>	string null			
<code>name</code>	string null			
<code>language</code>	Languages null			

ManageZone

Field	Type	Req	Default	Description
<code>id</code>	string	yes		
<code>slug</code>	string	yes		
<code>title</code>	string null			
<code>creator</code>	string null			
<code>account</code>	string null			
<code>created</code>	string null			
<code>modified</code>	string null			

cloud_provider	CloudProvider	yes		
private	boolean	yes		
origin	string null			

ManageZoneAdd

Field	Type	Req	Default	Description
slug	string	yes		
title	string	yes		
creator	string	yes		
cloud_provider	CloudProvider	yes		
private	boolean		False	
origin	string null			

ManageZoneCreated

Field	Type	Req	Default	Description
id	string	yes		
slug	string	yes		
title	string null			
creator	string null			
account	string null			
created	string null			
modified	string null			
cloud_provider	CloudProvider	yes		
private	boolean	yes		
origin	string null			
token	string	yes		

ManageZonePatch

Field	Type	Req	Default	Description
slug	string null			
title	string null			
cloud_provider	CloudProvider null			
private	boolean null			
origin	string null			

ManageZoneSessionToken

Field	Type	Req	Default	Description
token	string	yes		

ManualAgreementRequest

Request body for setting manual agreement data for CPPO customers.

Field	Type	Req	Default	Description
agreement	Agreement-Input	yes		

ManualSubscriptionCreate

Field	Type	Req	Default	Description
nuclia_tokens_price	number	yes		
billing_formula_id	string		00000000-0000-4000-8000-000000000001	
free_storage_nuclia_tokens	integer		0	
free_tokens_per_billing_cycle	integer		0	

<code>on_demand_budget</code>	number null			
<code>action_on_budget_exhausted</code>	ActionOnBudgetExhausted null			

ManualSubscriptionUpdate

Field	Type	Req	Default	Description
<code>billing_formula_id</code>	string null			
<code>nuclia_tokens_price</code>	number null			
<code>free_storage_nuclia_tokens</code>	integer null			
<code>free_tokens_per_billing_cycle</code>	integer null			
<code>on_demand_budget</code>	number null			
<code>action_on_budget_exhausted</code>	ActionOnBudgetExhausted null			

MarketplaceSubscriptionState

Field	Type	Req	Default	Description
<code>account_type</code>	AccountType	yes		
<code>is_trial</code>	boolean	yes		
<code>trial_expiration_date</code>	string null			
<code>active</code>	boolean		True	
<code>blocking_state</code>	AccountBlockingState null			

MarketplaceSubscriptionUpdateAction

Enum: `change_account_type`, `remove_trial`, `unblock_account`, `activate_subscription`, `deactivate_subscription`, `block_account_canceled`

MeteredTerm-Input

Field	Type	Req	Default	Description
<code>price</code>	number string	yes		

MeteredTerm-Output

Field	Type	Req	Default	Description
<code>price</code>	string	yes		

Metric

Field	Type	Req	Default	Description
<code>name</code>	MetricNames	yes		
<code>value</code>	number	yes		
<code>details</code>	array<NucliaTokenDetails>		<code>[]</code>	

MetricNames

Enum: `pre_processing_time`, `slow_processing_time`, `resources_processed`, `bytes_processed`, `chars_processed`, `media_seconds_processed`, `pages_processed`, `paragraphs_processed`, `searches_performed`, `suggestions_performed`, `predictions_performed`, `docs_no_media`, `ai_tokens_used`, `nuclia_tokens`, `nuclia_tokens_billed`

ModelFactor

Field	Type	Req	Default	Description
<code>input</code>	number null			
<code>output</code>	number null			
<code>image</code>	number null			

model	string	yes		
-------	--------	-----	--	--

NuaKeyIndexCreationPayload

Field	Type	Req	Default	Description
client_id	string	yes		
account_id	string	yes		
internal_id	string	yes		
tokens_limit	integer null			

NuaKeyIndexUpdatePayload

Field	Type	Req	Default	Description
tokens_limit	integer null			

NuaKeyInfoResponse

Field	Type	Req	Default	Description
blocked	integer		0	
tokens_limit	integer null			

nuclia_accounting_api_models_InvoiceItem

Field	Type	Req	Default	Description
id	string	yes		
period_start	string	yes		
period_end	string	yes		
order_number	string	yes		
account_type	AccountType	yes		
amount	number	yes		
pdf	string	yes		

nuclia_accounting__models__models__InvoiceItem

Field	Type	Req	Default	Description
threshold	integer	yes		
current_usage	integer	yes		
price	number	yes		
over_usage	integer	yes		
over_cost	number	yes		

NucliaCustomerData

Field	Type	Req	Default	Description
account_id	string	yes		
created	string null			
modified	string null			
billing_formula_id	string		00000000 - 0000-4000 - 8000 - 000000000001	
free_storage_nuclia_tokens	integer null			
free_tokens_per_billing_cycle	integer		0	
nuclia_tokens_price	number	yes		
action_on_budget_exhausted	ActionOnBudgetExhausted null			
on_demand_budget	number null			
billing_history	object (free-form map)			
last_notifications	LastNotifications			

NucliaCustomerTableSubscriptionStatusEnum: `active`**NucliaTokenDetails**

Field	Type	Req	Default	Description
<code>identifier</code>	FlatIdentifier	yes		
<code>nuclia_tokens</code>	Tokens	yes		
<code>nuclia_tokens_billed</code>	Tokens	yes		
<code>raw_usage</code>	Tokens	yes		
<code>requests</code>	Requests	yes		

OAuthTokens

Field	Type	Req	Default	Description
<code>access_token</code>	string	yes		
<code>refresh_token</code>	string	yes		

Offer

Field	Type	Req	Default	Description
<code>id</code>	string	yes		
<code>name</code>	string	yes		
<code>is_trial</code>	boolean		<code>False</code>	
<code>trial_duration</code>	string null			

OldSAMLConfig

Field	Type	Req	Default	Description
<code>domain</code>	string	yes		
<code>sso_url</code>	string	yes		

entity_id	string	yes		
x509_cert	string	yes		

PATTokenMetadataUpdate

Field	Type	Req	Default	Description
token_jti	string	yes		
issuer	string null			
seen_at	string	yes		

PaymentError

Field	Type	Req	Default	Description
code	string	yes		
decline_code	string	yes		

PaymentMethod

Field	Type	Req	Default	Description
payment_method_id	string	yes		
payment_method_type	string	yes		

RawFactor

Field	Type	Req	Default	Description
factor	number	yes		
source	string	yes		

RecoverPayload

Field	Type	Req	Default	Description
-------	------	-----	---------	-------------

username	string	yes		
app	string null			
login_challenge	string null			
initial_setpassword	boolean		False	

RefreshRequest

Field	Type	Req	Default	Description
refresh_token	string	yes		

RequestGlobalMagicToken

Field	Type	Req	Default	Description
action_id	ActionId	yes		
token	string null			
email	string null			
account_id	string null			
came_from	string null			

Requests

Field	Type	Req	Default	Description
api	number null			
web	number null			
widget	number null			
desktop	number null			
dashboard	number null			
chrome_extension	number null			
internal	number null			

ResetPayload

Field	Type	Req	Default	Description
password	string	yes		
token	string	yes		
login_challenge	string null			

SAMLConfig

Field	Type	Req	Default	Description
domains	array<string>	yes		
sso_url	string	yes		
entity_id	string	yes		
x509_cert	string	yes		
authn_context	enum[exact, minimum, better, maximum] null			

SetPasswordPayload

Field	Type	Req	Default	Description
password	string	yes		

SetupUserPayload

Field	Type	Req	Default	Description
token	string	yes		
password	string	yes		
name	string	yes		

SignupAction

Enum: `check-mail`, `user-exists`

SignupRequest

Field	Type	Req	Default	Description
<code>password</code>	string	yes		
<code>login_challenge</code>	string	yes		

SignupResponse

Field	Type	Req	Default	Description
<code>action</code>	SignupAction	yes		

SocialLoginAuthorizationRequest

Field	Type	Req	Default	Description
<code>code</code>	string	yes		
<code>login_challenge</code>	string null			

Stash

Field	Type	Req	Default	Description
<code>id</code>	string	yes		
<code>slug</code>	string	yes		
<code>title</code>	string null			
<code>description</code>	string null			
<code>avatar</code>	string null			
<code>creator</code>	string	yes		
<code>state</code>	KbState	yes		
<code>account</code>	string	yes		

zone	string	yes		
created	string	yes		
modified	string null			
data	StashPreferences null			

stashify_idp__api__manage__models__User

Field	Type	Req	Default	Description
id	string null			
last_login	string null			
last_login_ip	string null			
avatar	string null			
email	string	yes		
name	string null			
type	UserType	yes		
language	Languages null			
created	string null			
modified	string null			
data	UserPreferences null			

stashify_idp__api__models__User

Field	Type	Req	Default	Description
id	string	yes		
name	string	yes		
email	string	yes		
language	Languages null			
type	UserType	yes		
last_login	string null			
avatar	string null			

StashPreferences**Type:** `object`**StripeCharge**

Field	Type	Req	Default	Description
<code>id</code>	string	yes		
<code>object</code>	string	yes		
<code>customer</code>	string	yes		
<code>status</code>	StripeChargeStatus	yes		
<code>amount</code>	integer	yes		
<code>currency</code>	string	yes		
<code>description</code>	string	yes		
<code>invoice</code>	string	yes		
<code>failure_code</code>	string null			
<code>failure_message</code>	string null			
<code>paid</code>	boolean	yes		
<code>created</code>	integer	yes		
<code>outcome</code>	StripeChargeOutcome null			
<code>metadata</code>	object (free-form map)		<code>{}</code>	

StripeChargeOutcome

Field	Type	Req	Default	Description
<code>network_statuses</code>	string null			
<code>reason</code>	string null			
<code>risk_level</code>	string null			
<code>risk_score</code>	integer null			

<code>rule</code>	string null			
<code>seller_message</code>	string null			
<code>type</code>	string null			

StripeChargesList

Field	Type	Req	Default	Description
<code>object</code>	string	yes		
<code>data</code>	array<StripeCharge>	yes		

StripeChargeStatus

Enum: `succeeded`, `pending`, `failed`

StripeCheckoutSession

Field	Type	Req	Default	Description
<code>id</code>	string	yes		
<code>object</code>	string	yes		
<code>customer</code>	string	yes		
<code>payment_link</code>	string null	yes		
<code>subscription</code>	string	yes		
<code>mode</code>	string	yes		
<code>status</code>	string	yes		

StripeEvent

Field	Type	Req	Default	Description
<code>object</code>	string	yes		
<code>data</code>	StripeObject	yes		
<code>type</code>	string	yes		

StripeObject

Field	Type	Req	Default	Description
<code>object</code>	StripeSubscription StripePaymentIntent StripeCheckoutSession	yes		
<code>previous_attributes</code>	object (free-form map)		<code>{}</code>	

StripePaymentIntent

Field	Type	Req	Default	Description
<code>id</code>	string	yes		
<code>object</code>	string	yes		
<code>customer</code>	string	yes		
<code>status</code>	StripePaymentIntentStatus	yes		
<code>amount</code>	integer	yes		
<code>currency</code>	string	yes		
<code>description</code>	string	yes		
<code>charges</code>	StripeChargesList null			
<code>metadata</code>	object (free-form map)		<code>{}</code>	

StripePaymentIntentStatus

Enum: `requires_payment_method`, `requires_confirmation`, `requires_action`, `processing`, `requires_capture`, `canceled`, `succeeded`

StripePaymentLinkInfo

Field	Type	Req	Default	Description
<code>id</code>	string	yes		
<code>url</code>	string	yes		

StripePriceInfo

Field	Type	Req	Default	Description
<code>id</code>	string	yes		
<code>nickname</code>	string	yes		
<code>product</code>	string	yes		

StripePublicKey

Field	Type	Req	Default	Description
<code>public_key</code>	string	yes		

StripeSubscription

Field	Type	Req	Default	Description
<code>id</code>	string	yes		
<code>object</code>	string	yes		
<code>billing_cycle_anchor</code>	integer	yes		
<code>customer</code>	string	yes		
<code>default_payment_method</code>	string null			
<code>status</code>	StripeSubscriptionStatus	yes		
<code>items</code>	StripeSubscriptionItemsList	yes		
<code>metadata</code>	object (free-form map)		<code>{}</code>	

StripeSubscriptionCancellationFeedback

Enum: `too_expensive`, `missing_features`, `switched_service`, `unused`, `customer_service`, `too_complex`, `low_quality`, `other`

StripeSubscriptionItemsList

Field	Type	Req	Default	Description
-------	------	-----	---------	-------------

<code>object</code>	string	yes		
<code>data</code>	array<object (free-form map)>	yes		

StripeSubscriptionStatus

Enum: `incomplete`, `incomplete_expired`, `trialing`, `active`, `past_due`, `canceled`, `unpaid`

StripeUsageType

Enum: `metered`, `licensed`

SubscriptionProvider

Enum: `NO_SUBSCRIPTION`, `STRIPE`, `AWS_MARKETPLACE`, `MANUAL`, `CLOUD_ZERO`

SubscriptionStatus

Enum: `no_subscription`, `pending`, `active`, `payment_issues`, `cancel_scheduled`, `canceled`

SubscriptionSyncExecutionResponse

Field	Type	Req	Default	Description
<code>previous_state</code>	MarketplaceSubscriptionState	yes		
<code>current_state</code>	MarketplaceSubscriptionState	yes		
<code>actions</code>	array<MarketplaceSubscriptionUpdateAction>	yes		

SubscriptionSyncStateResponse

Field	Type	Req	Default	Description
<code>current_state</code>	MarketplaceSubscriptionState	yes		
<code>desired_state</code>	MarketplaceSubscriptionState	yes		
<code>actions</code>	array<MarketplaceSubscriptionUpdateAction>	yes		

Tier

Field	Type	Req	Default	Description
<code>recurring</code>	<code>map<string, TierRecurringBillingItem> null</code>	yes		
<code>usage</code>	<code>map<string, TierUsageBillingItem></code>	yes		

TierRecurringBillingItem

Field	Type	Req	Default	Description
<code>price</code>	<code>number</code>	yes		

TierUsageBillingItem

Field	Type	Req	Default	Description
<code>threshold</code>	<code>integer</code>	yes		
<code>price</code>	<code>number</code>	yes		

TokenCreated

Field	Type	Req	Default	Description
<code>id</code>	<code>string</code>	yes		
<code>token</code>	<code>string</code>	yes		
<code>expires</code>	<code>string null</code>	yes		

TokenExchangeRequestInput

Field	Type	Req	Default	Description
<code>token</code>	<code>string</code>	yes		

TokenItem

Field	Type	Req	Default	Description
<code>id</code>	string	yes		
<code>description</code>	string	yes		
<code>expires</code>	string null	yes		

TokenMetadataPayload

Field	Type	Req	Default	Description
<code>pat_updates</code>	array<PATTokenMetadataUpdate>		<code>[]</code>	

Tokens

Field	Type	Req	Default	Description
<code>input</code>	number null			
<code>output</code>	number null			
<code>image</code>	number null			
<code>search</code>	number null			
<code>storage</code>	number null			

UpdateAWSMarketplaceSubscription

Field	Type	Req	Default	Description
<code>on_demand_budget</code>	number null			
<code>action_on_budget_exhausted</code>	ActionOnBudgetExhausted null			
<code>free_tokens_per_billing_cycle</code>	integer null			

UpdateStripeCustomer

Field	Type	Req	Default	Description
<code>billing_details</code>	BillingDetailsModified	yes		

UpdateStripeSubscription

Field	Type	Req	Default	Description
<code>on_demand_budget</code>	number null			
<code>action_on_budget_exhausted</code>	ActionOnBudgetExhausted null			
<code>free_tokens_per_billing_cycle</code>	integer null			

UpdateStripeSubscriptionMeteredPrices

Field	Type	Req	Default	Description
<code>price_nicknames</code>	array<string>	yes		

UpgradeStripeSubscription

Field	Type	Req	Default	Description
<code>billing_interval</code>	IntervalType null		<code>month</code>	
<code>account_type</code>	AccountType	yes		

UserCreated

Field	Type	Req	Default	Description
<code>id</code>	string	yes		

UserCreation

Field	Type	Req	Default	Description
name	string	yes		
email	string	yes		
language	Languages null			
type	UserType null		USER	
last_login	string null			

UserOnboardingInquiry

Field	Type	Req	Default	Description
use_case	string null			
organization_size	string null			
receive_updates	boolean null			
company	string null			
phone	string null			
role	string null			
country	string null			
accept_privacy_policy	boolean null		False	

UserPreferences**Type:** object**UserSetPreferences**

Field	Type	Req	Default	Description
name	string null			
avatar	string null			
email	string null			

language	Languages null			
----------	------------------	--	--	--

UserType

Enum: ROOT, DEALER, USER, READONLY, MANAGER, SALES

ValidationError

Field	Type	Req	Default	Description
loc	array<string integer>	yes		
msg	string	yes		
type	string	yes		
input	object			
ctx	object			

Welcome

Field	Type	Req	Default	Description
preferences	WelcomeUser	yes		
accounts	array<string>	yes		
create	boolean	yes		
can_be_deleted	boolean	yes		
dependant_accounts	array<AccountId>	yes		

WelcomeUser

Field	Type	Req	Default	Description
last_login	string null			
email	string	yes		
name	string null			
type	UserType	yes		

language	Languages null			
----------	------------------	--	--	--

WorkflowMode

Enum: classic, cowork

Zone

Field	Type	Req	Default	Description
id	string	yes		
slug	string	yes		
title	string null			
cloud_provider	CloudProvider	yes		
private	boolean	yes		
origin	string null			

ZoneVisibility

Enum: DEFAULT, RESTRICTED

PART 7 — REFERENCE APPENDIX

Appendix G — API Verification Matrix

Exactly how every endpoint was verified — evidence, not assertion

This matrix records the verification status of every documented operation, established by running a systematic sweep against a live sandbox Knowledge Box in the `aws-ap-southeast-2-1` zone. It exists so that 'how much of this is proven' is something you can read, not something you have to take on trust.

Methodology

- **Read endpoints** (GET/HEAD) were called directly against the sandbox.
- **Lifecycle endpoints** (create/update/delete of labelsets, search configurations, KV schemas, extract strategies, resources and fields, service accounts and keys, ephemeral tokens, tasks) were exercised with full create → verify → delete cycles on disposable objects.
- **Write probes:** remaining write endpoints were called with a minimal body; a 422 validation response confirms the endpoint is live and reveals its required fields without mutating data.
- **Not executed on purpose:** irreversible account-level operations, and anything that would delete real data or incur meaningful cost, were deliberately skipped and are marked as such.

Status legend

Status	Meaning
Verified (2xx)	Called live and returned a success response
Confirmed live	Called live; returned a real 4xx (validation/auth/not-found) — the endpoint exists and its behaviour was observed
Server 5xx	Called live; returned a server error (recorded and explained)
Not run (safety)	A destructive/irreversible write not executed against real data
Needs agent	A Retrieval Agent sub-resource requiring a pre-created agent (provisioned in the dashboard, outside these four specs)
Needs user/PAT	A Global API operation; the NUA key used for testing is rejected there by design

Summary

Of 356 documented operations, 235 (66%) were exercised live and returned a real response; 102 (28%) returned a full success. The remainder were not executed for the honest reasons below.

Status	Count	Share
Verified (2xx)	102	28%
Confirmed live	131	36%
Server 5xx	2	0%
Not run (safety)	54	15%
Needs agent	40	11%
Needs user/PAT	27	7%

Key idea The ~34% not fully verified is bounded by the test credentials, not by gaps in the documentation: destructive operations were skipped on purpose, the Global API requires a user token / PAT that a NUA key cannot substitute for, and the Retrieval Agent sub-resources need an agent provisioned in the dashboard. Every one of these is still documented in full in Appendices A–F.

Notable behaviours observed during verification

- **Tasks (data-augmentation agents) require the Knowledge Box key, not the NUA key.** `POST /kb/{kbid}/task/start` returned 403 with a NUA key but 200 with the KB API key.
- **Activity logs and backups reject the NUA key** (`NuaKeyUser cannot access context ...`); they need a user token / PAT.
- `GET /kb/{kbid}/notifications` is a streaming endpoint` — it holds the connection open (the sweep saw a read-timeout), so treat it as a long-lived stream, not a request/response call.
- `GET /processing/download` and the KB `logo.png` returned 500 when no payload/logo exists` — call them only when you know the artifact is present.
- **Several Predict endpoints are POST-only** (`/predict/rerank, /predict/remi`) — a GET returns 405.

The matrix

nucliadb (115 operations)

Method	Path	Verification
GET	/kb/s/{slug}	Confirmed live · 403
GET	/kb/{kbid}	Verified (2xx) · 200
POST	/kb/{kbid}/ask	Verified (2xx) · 200
GET	/kb/{kbid}/catalog	Verified (2xx) · 200
POST	/kb/{kbid}/catalog	Verified (2xx) · 200
GET	/kb/{kbid}/configuration	Verified (2xx) · 200
PATCH	/kb/{kbid}/configuration	Verified (2xx) · 204
POST	/kb/{kbid}/configuration	Confirmed live · 405
GET	/kb/{kbid}/counters	Verified (2xx) · 200
DELETE	/kb/{kbid}/custom-synonyms	Verified (2xx) · 204
GET	/kb/{kbid}/custom-synonyms	Verified (2xx) · 200
PUT	/kb/{kbid}/custom-synonyms	Verified (2xx) · 204
GET	/kb/{kbid}/entitiesgroup/{group}	Confirmed live · 404
GET	/kb/{kbid}/entitiesgroups	Verified (2xx) · 200
GET	/kb/{kbid}/export/{export_id}	Confirmed live · 404
GET	/kb/{kbid}/export/{export_id}/status	Confirmed live · 404
GET	/kb/{kbid}/extract_strategies	Verified (2xx) · 200
POST	/kb/{kbid}/extract_strategies	Verified (2xx) · 200
DELETE	/kb/{kbid}/extract_strategies/	Not run (safety)

	strategy/{strategy_id}	
GET	/kb/{kbid}/ extract_strategies/ strategy/{strategy_id}	Confirmed live · 404
POST	/kb/{kbid}/feedback	Verified (2xx) · 200
GET	/kb/{kbid}/find	Verified (2xx) · 200
POST	/kb/{kbid}/find	Verified (2xx) · 200
GET	/kb/{kbid}/ generative_providers	Verified (2xx) · 200
POST	/kb/{kbid}/graph	Verified (2xx) · 200
POST	/kb/{kbid}/graph/nodes	Verified (2xx) · 200
POST	/kb/{kbid}/graph/ relations	Confirmed live · 422
GET	/kb/{kbid}/import/ {import_id}/status	Confirmed live · 404
GET	/kb/{kbid}/kv-schemas	Verified (2xx) · 200
POST	/kb/{kbid}/kv-schemas	Verified (2xx) · 201
DELETE	/kb/{kbid}/kv-schemas/ {schema_id}	Verified (2xx) · 204
GET	/kb/{kbid}/kv-schemas/ {schema_id}	Confirmed live · 404
PUT	/kb/{kbid}/kv-schemas/ {schema_id}	Confirmed live · 404
DELETE	/kb/{kbid}/labelset/ {labelset}	Verified (2xx) · 200
GET	/kb/{kbid}/labelset/ {labelset}	Verified (2xx) · 200
POST	/kb/{kbid}/labelset/ {labelset}	Verified (2xx) · 200
GET	/kb/{kbid}/labelsets	Verified (2xx) · 200
GET	/kb/{kbid}/model/ {model_id}	Confirmed live · 404
GET	/kb/{kbid}/models	Verified (2xx) · 200
GET	/kb/{kbid}/models/ {model_id}/{filename}	Confirmed live · 404

GET	/kb/{kbid}/notifications	Confirmed live · stream
GET	/kb/{kbid}/predict/ {endpoint}	Confirmed live · 422
POST	/kb/{kbid}/predict/ {endpoint}	Confirmed live · 405
GET	/kb/{kbid}/processing- status	Verified (2xx) · 200
POST	/kb/{kbid}/resource/ {path_rid}/file/ {field}/tusupload	Confirmed live · 412
HEAD	/kb/{kbid}/resource/ {path_rid}/file/ {field}/tusupload/ {upload_id}	Verified (2xx) · 200
POST	/kb/{kbid}/resource/ {path_rid}/file/ {field}/upload	Verified (2xx) · 201
DELETE	/kb/{kbid}/resource/ {rid}	Verified (2xx) · 204
GET	/kb/{kbid}/resource/ {rid}	Verified (2xx) · 200
HEAD	/kb/{kbid}/resource/ {rid}	Verified (2xx) · 200
PATCH	/kb/{kbid}/resource/ {rid}	Verified (2xx) · 200
POST	/kb/{kbid}/resource/ {rid}/ask	Confirmed live · 422
PUT	/kb/{kbid}/resource/ {rid}/conversation/ {field_id}	Verified (2xx) · 201
GET	/kb/{kbid}/resource/ {rid}/conversation/ {field_id}/download/ field/{message_id}/ {file_num}	Confirmed live · 404
PUT	/kb/{kbid}/resource/ {rid}/conversation/ {field_id}/messages	Confirmed live · 422

DELETE	/kb/{kbid}/resource/ {rid}/conversation/ {field_id}/messages/ {message_id}	Not run (safety)
PUT	/kb/{kbid}/resource/ {rid}/file/{field_id}	Confirmed live · 422
GET	/kb/{kbid}/resource/ {rid}/file/{field_id}/ download/field	Verified (2xx) · 200
POST	/kb/{kbid}/resource/ {rid}/file/{field_id}/ reprocess	Not run (safety)
PATCH	/kb/{kbid}/resource/ {rid}/file/{field}/ tusupload/{upload_id}	Confirmed live · 404
PUT	/kb/{kbid}/resource/ {rid}/link/{field_id}	Confirmed live · 422
POST	/kb/{kbid}/resource/ {rid}/reindex	Verified (2xx) · 200
POST	/kb/{kbid}/resource/ {rid}/reprocess	Verified (2xx) · 202
POST	/kb/{kbid}/resource/ {rid}/run-agents	Verified (2xx) · 200
GET	/kb/{kbid}/resource/ {rid}/search	Verified (2xx) · 200
PUT	/kb/{kbid}/resource/ {rid}/text/{field_id}	Confirmed live · 422
DELETE	/kb/{kbid}/resource/ {rid}/{field_type}/ {field_id}	Verified (2xx) · 204
GET	/kb/{kbid}/resource/ {rid}/{field_type}/ {field_id}	Confirmed live · 404
GET	/kb/{kbid}/resource/ {rid}/{field_type}/ {field_id}/download/ extracted/ {download_field}	Confirmed live · 404

GET	/kb/{kbid}/resources	Verified (2xx) · 200
POST	/kb/{kbid}/resources	Verified (2xx) · 201
GET	/kb/{kbid}/schema	Verified (2xx) · 200
GET	/kb/{kbid}/search	Verified (2xx) · 200
POST	/kb/{kbid}/search	Verified (2xx) · 200
GET	/kb/{kbid}/ search_configurations	Verified (2xx) · 200
DELETE	/kb/{kbid}/ search_configurations/ {config_name}	Verified (2xx) · 204
GET	/kb/{kbid}/ search_configurations/ {config_name}	Verified (2xx) · 200
PATCH	/kb/{kbid}/ search_configurations/ {config_name}	Confirmed live · 422
POST	/kb/{kbid}/ search_configurations/ {config_name}	Verified (2xx) · 201
DELETE	/kb/{kbid}/slug/{rslug}	Not run (safety)
GET	/kb/{kbid}/slug/{rslug}	Verified (2xx) · 200
HEAD	/kb/{kbid}/slug/{rslug}	Verified (2xx) · 200
PATCH	/kb/{kbid}/slug/{rslug}	Verified (2xx) · 200
PUT	/kb/{kbid}/slug/ {rslug}/conversation/ {field_id}	Verified (2xx) · 201
GET	/kb/{kbid}/slug/ {rslug}/conversation/ {field_id}/download/ field/{message_id}/ {file_num}	Confirmed live · 404
PUT	/kb/{kbid}/slug/ {rslug}/conversation/ {field_id}/messages	Confirmed live · 422
DELETE	/kb/{kbid}/slug/ {rslug}/conversation/ {field_id}/messages/	Not run (safety)

	{message_ident}	
PUT	/kb/{kbid}/slug/ {rslug}/file/{field_id}	Confirmed live · 422
GET	/kb/{kbid}/slug/ {rslug}/file/ {field_id}/download/ field	Verified (2xx) · 200
POST	/kb/{kbid}/slug/ {rslug}/file/{field}/ tusupload	Confirmed live · 412
HEAD	/kb/{kbid}/slug/ {rslug}/file/{field}/ tusupload/{upload_id}	Verified (2xx) · 200
PATCH	/kb/{kbid}/slug/ {rslug}/file/{field}/ tusupload/{upload_id}	Confirmed live · 404
POST	/kb/{kbid}/slug/ {rslug}/file/{field}/ upload	Verified (2xx) · 201
PUT	/kb/{kbid}/slug/ {rslug}/link/{field_id}	Confirmed live · 422
POST	/kb/{kbid}/slug/ {rslug}/reindex	Not run (safety)
POST	/kb/{kbid}/slug/ {rslug}/reprocess	Not run (safety)
PUT	/kb/{kbid}/slug/ {rslug}/text/{field_id}	Confirmed live · 422
DELETE	/kb/{kbid}/slug/ {rslug}/{field_type}/ {field_id}	Not run (safety)
GET	/kb/{kbid}/slug/ {rslug}/{field_type}/ {field_id}	Confirmed live · 404
GET	/kb/{kbid}/slug/ {rslug}/{field_type}/ {field_id}/download/ extracted/ {download_field}	Confirmed live · 404

POST	/kb/{kbid}/slug/{slug}/ask	Confirmed live · 422
POST	/kb/{kbid}/slug/{slug}/run-agents	Verified (2xx) · 200
GET	/kb/{kbid}/split_strategies	Verified (2xx) · 200
POST	/kb/{kbid}/split_strategies	Verified (2xx) · 200
DELETE	/kb/{kbid}/split_strategies/strategy/{strategy_id}	Not run (safety)
GET	/kb/{kbid}/split_strategies/strategy/{strategy_id}	Confirmed live · 404
GET	/kb/{kbid}/suggest	Verified (2xx) · 200
POST	/kb/{kbid}/suggest	Confirmed live · 422
POST	/kb/{kbid}/summarize	Verified (2xx) · 200
OPTIONS	/kb/{kbid}/tusupload	Verified (2xx) · 204
POST	/kb/{kbid}/tusupload	Confirmed live · 412
HEAD	/kb/{kbid}/tusupload/{upload_id}	Verified (2xx) · 200
PATCH	/kb/{kbid}/tusupload/{upload_id}	Confirmed live · 404
POST	/kb/{kbid}/upload	Verified (2xx) · 201
GET	/learning/configuration/schema	Verified (2xx) · 200

nua (157 operations)

Method	Path	Verification
DELETE	/account/{account_id}/assume_role/bedrock	Not run (safety)
GET	/account/{account_id}/assume_role/bedrock	Confirmed live · 403
POST	/account/{account_id}/	Not run (safety)

	<code>assume_role/bedrock</code>	
GET	<code>/account/{account_id}/ assume_role/bedrock/ validate</code>	Confirmed live · 403
DELETE	<code>/account/{account_id}/ dataset/{dataset_id}</code>	Not run (safety)
GET	<code>/account/{account_id}/ dataset/{dataset_id}</code>	Confirmed live · 403
GET	<code>/account/{account_id}/ dataset/{dataset_id}/ models</code>	Confirmed live · 403
DELETE	<code>/account/{account_id}/ dataset/{dataset_id}/ partition/{partition_id}</code>	Not run (safety)
PUT	<code>/account/{account_id}/ dataset/{dataset_id}/ partition/{partition_id}</code>	Confirmed live · 403
DELETE	<code>/account/{account_id}/ dataset/{dataset_id}/ partitions</code>	Not run (safety)
GET	<code>/account/{account_id}/ datasets</code>	Confirmed live · 403
POST	<code>/account/{account_id}/ datasets</code>	Confirmed live · 403
DELETE	<code>/account/{account_id}/ default_model/{model_id}</code>	Not run (safety)
GET	<code>/account/{account_id}/ default_model/{model_id}</code>	Confirmed live · 422
GET	<code>/account/{account_id}/ default_models</code>	Confirmed live · 403
DELETE	<code>/account/{account_id}/ model/{model_id}</code>	Not run (safety)
GET	<code>/account/{account_id}/ model/{model_id}</code>	Confirmed live · 403
GET	<code>/account/{account_id}/ models</code>	Confirmed live · 403
DELETE	<code>/account/{account_id}/</code>	Not run (safety)

	models/{kbid}/{model_id}	
GET	/account/{account_id}/ schema	Confirmed live · 403
DELETE	/agent/{agent_id}/ audit/download_request/ {request_id}	Needs agent
GET	/agent/{agent_id}/ audit/download_request/ {request_id}/status	Confirmed live · 403
GET	/agent/{agent_id}/ audit/download_requests	Confirmed live · 403
POST	/agent/{agent_id}/ audit/interactions/ download	Needs agent
DELETE	/agent/{agent_id}/ driver/{driver}	Needs agent
PATCH	/agent/{agent_id}/ driver/{driver}	Needs agent
GET	/agent/{agent_id}/ drivers	Confirmed live · 403
POST	/agent/{agent_id}/ drivers	Needs agent
POST	/agent/{agent_id}/export	Needs agent
GET	/agent/{agent_id}/ generation	Confirmed live · 403
POST	/agent/{agent_id}/ generation	Needs agent
DELETE	/agent/{agent_id}/ generation/{generation}	Needs agent
PATCH	/agent/{agent_id}/ generation/{generation}	Needs agent
POST	/agent/{agent_id}/import	Needs agent
GET	/agent/{agent_id}/ postprocess	Confirmed live · 403
POST	/agent/{agent_id}/ postprocess	Needs agent

DELETE	/agent/{agent_id}/ postprocess/ {postprocess}	Needs agent
PATCH	/agent/{agent_id}/ postprocess/ {postprocess}	Needs agent
GET	/agent/{agent_id}/ preprocess	Confirmed live · 403
POST	/agent/{agent_id}/ preprocess	Needs agent
DELETE	/agent/{agent_id}/ preprocess/{preprocess}	Needs agent
PATCH	/agent/{agent_id}/ preprocess/{preprocess}	Needs agent
DELETE	/agent/{agent_id}/ prompt/{prompt_id}	Needs agent
GET	/agent/{agent_id}/ prompt/{prompt_id}	Confirmed live · 403
PATCH	/agent/{agent_id}/ prompt/{prompt_id}	Needs agent
GET	/agent/{agent_id}/ prompts	Confirmed live · 403
POST	/agent/{agent_id}/ prompts	Needs agent
GET	/agent/{agent_id}/rules	Confirmed live · 403
POST	/agent/{agent_id}/rules	Needs agent
DELETE	/agent/{agent_id}/ session/{session}	Needs agent
GET	/agent/{agent_id}/ session/{session}	Confirmed live · 403
PATCH	/agent/{agent_id}/ session/{session}	Needs agent
POST	/agent/{agent_id}/ session/{session}	Needs agent
DELETE	/agent/{agent_id}/ session/{session}/mcp	Needs agent

GET	/agent/{agent_id}/ session/{session}/mcp	Confirmed live · 403
POST	/agent/{agent_id}/ session/{session}/mcp	Needs agent
GET	/agent/{agent_id}/ sessions	Confirmed live · 403
POST	/agent/{agent_id}/ sessions	Needs agent
DELETE	/agent/{agent_id}/ workflow/{workflow_id}	Needs agent
PATCH	/agent/{agent_id}/ workflow/{workflow_id}	Needs agent
GET	/agent/{agent_id}/ workflow/{workflow_id}/ generation	Confirmed live · 403
POST	/agent/{agent_id}/ workflow/{workflow_id}/ generation	Needs agent
DELETE	/agent/{agent_id}/ workflow/{workflow_id}/ generation/{generation}	Needs agent
PATCH	/agent/{agent_id}/ workflow/{workflow_id}/ generation/{generation}	Needs agent
GET	/agent/{agent_id}/ workflow/{workflow_id}/ postprocess	Confirmed live · 403
POST	/agent/{agent_id}/ workflow/{workflow_id}/ postprocess	Needs agent
DELETE	/agent/{agent_id}/ workflow/{workflow_id}/ postprocess/ {postprocess}	Needs agent
PATCH	/agent/{agent_id}/ workflow/{workflow_id}/ postprocess/ {postprocess}	Needs agent

GET	/agent/{agent_id}/ workflow/{workflow_id}/ preprocess	Confirmed live · 403
POST	/agent/{agent_id}/ workflow/{workflow_id}/ preprocess	Needs agent
DELETE	/agent/{agent_id}/ workflow/{workflow_id}/ preprocess/{preprocess}	Needs agent
PATCH	/agent/{agent_id}/ workflow/{workflow_id}/ preprocess/{preprocess}	Needs agent
GET	/agent/{agent_id}/ workflow/{workflow_id}/ rules	Confirmed live · 403
POST	/agent/{agent_id}/ workflow/{workflow_id}/ rules	Needs agent
POST	/agent/{agent_id}/ workflow/{workflow_id}/ session/{session}	Needs agent
GET	/agent/{agent_id}/ workflows	Confirmed live · 403
POST	/agent/{agent_id}/ workflows	Needs agent
GET	/collect/feedback/{kbid}	Confirmed live · 404
GET	/collect/feedback/ {kbid}/{month}	Confirmed live · 404
DELETE	/config/{kbid}	Not run (safety)
GET	/config/{kbid}	Confirmed live · 404
PATCH	/config/{kbid}	Verified (2xx) · 200
POST	/config/{kbid}	Verified (2xx) · 200
DELETE	/dataset/{dataset_id}	Not run (safety)
PUT	/dataset/{dataset_id}/ partition/{partition_id}	Confirmed live · 404
POST	/dataset/{dataset_id}/ task/start	Not run (safety)

DELETE	/dataset/{dataset_id}/task/{task_id}	Not run (safety)
GET	/dataset/{dataset_id}/task/{task_id}/inspect	Confirmed live · 404
POST	/dataset/{dataset_id}/task/{task_id}/stop	Not run (safety)
GET	/dataset/{dataset_id}/tasks	Verified (2xx) · 200
POST	/datasets	Confirmed live · 422
GET	/download/{kbid}/model/{model_id}/{filename}	Confirmed live · 404
GET	/extract_strategies/{kbid}	Confirmed live · 404
POST	/extract_strategies/{kbid}	Confirmed live · 404
DELETE	/extract_strategies/{kbid}/strategies/{strategy_id}	Not run (safety)
GET	/extract_strategies/{kbid}/strategies/{strategy_id}	Confirmed live · 404
GET	/generative_providers/{kbid}	Verified (2xx) · 200
GET	/kb/{kb_uuid}/assets	Confirmed live · 404
GET	/kb/{kb_uuid}/assets/{asset_id}	Confirmed live · 404
POST	/kb/{kb_uuid}/task/start	Verified (2xx) · 200
DELETE	/kb/{kb_uuid}/task/{task_id}	Not run (safety)
PATCH	/kb/{kb_uuid}/task/{task_id}	Not run (safety)
POST	/kb/{kb_uuid}/task/{task_id}/cleanup	Not run (safety)
POST	/kb/{kb_uuid}/task/{task_id}/enable	Not run (safety)
GET	/kb/{kb_uuid}/task/	Confirmed live · 403

	<code>{task_id}/inspect</code>	
POST	<code>/kb/{kb_uuid}/task/{task_id}/restart</code>	Not run (safety)
POST	<code>/kb/{kb_uuid}/task/{task_id}/stop</code>	Not run (safety)
GET	<code>/kb/{kb_uuid}/task/{task_id}/worker_config</code>	Confirmed live · 403
GET	<code>/kb/{kb_uuid}/tasks</code>	Verified (2xx) · 200
POST	<code>/kb/{kbid}/ask</code>	Confirmed live · 403
GET	<code>/kb/{kbid}/mcp</code>	Confirmed live · 403
POST	<code>/kb/{kbid}/mcp</code>	Not run (safety)
POST	<code>/kb/{kbid}/resource/{rid}/ask</code>	Confirmed live · 403
POST	<code>/kb/{kbid}/slug/{slug}/ask</code>	Confirmed live · 403
GET	<code>/learning/configuration/schema</code>	Verified (2xx) · 200
GET	<code>/models/{kbid}</code>	Verified (2xx) · 200
POST	<code>/models/{kbid}</code>	Confirmed live · 422
GET	<code>/models/{kbid}/model/{model_id}</code>	Confirmed live · 404
POST	<code>/predict/chat</code>	Verified (2xx) · 200
POST	<code>/predict/chat/{kbid}</code>	Confirmed live · 422
POST	<code>/predict/compat/chat/completions</code>	Confirmed live · 422
POST	<code>/predict/compat/embeddings</code>	Confirmed live · 422
GET	<code>/predict/compat/models</code>	Verified (2xx) · 200
GET	<code>/predict/compat/models/{model_id}</code>	Confirmed live · 404
GET	<code>/predict/query</code>	Verified (2xx) · 200
POST	<code>/predict/query</code>	Verified (2xx) · 200
GET	<code>/predict/query/{kbid}</code>	Verified (2xx) · 200
POST	<code>/predict/query/{kbid}</code>	Verified (2xx) · 200

POST	/predict/remi	Confirmed live · 422
POST	/predict/remi/{kbid}	Confirmed live · 422
POST	/predict/rephrase	Verified (2xx) · 200
POST	/predict/rephrase/{kbid}	Confirmed live · 422
POST	/predict/rerank	Confirmed live · 422
POST	/predict/rerank/{kbid}	Confirmed live · 422
GET	/predict/sentence	Verified (2xx) · 200
GET	/predict/sentence/{kbid}	Verified (2xx) · 200
POST	/predict/sentences	Confirmed live · 422
POST	/predict/summarize	Verified (2xx) · 200
POST	/predict/summarize/{kbid}	Confirmed live · 422
GET	/predict/tokens	Verified (2xx) · 200
GET	/predict/tokens/{kbid}	Verified (2xx) · 200
GET	/processing/download	Server 5xx · 500
POST	/processing/push	Not run (safety)
GET	/processing/requests	Verified (2xx) · 200
GET	/processing/requests/{processing_id}	Confirmed live · 404
GET	/processing/requests/{processing_id}/results	Confirmed live · 404
POST	/processing/tusupload	Confirmed live · 412
HEAD	/processing/tusupload/{upload_id}	Verified (2xx) · 200
OPTIONS	/processing/tusupload/{upload_id}	Verified (2xx) · 200
PATCH	/processing/tusupload/{upload_id}	Confirmed live · 412
POST	/processing/upload	Verified (2xx) · 200
GET	/schema	Verified (2xx) · 200
GET	/schema/{kbid}	Verified (2xx) · 200
GET	/split_strategies/{kbid}	Confirmed live · 404
POST	/split_strategies/{kbid}	Confirmed live · 404

DELETE	/split_strategies/ {kbid}/strategies/ {strategy_id}	Not run (safety)
GET	/split_strategies/ {kbid}/strategies/ {strategy_id}	Confirmed live · 404

zone (57 operations)

Method	Path	Verification
DELETE	/api/internal/kb/ {kb_id}/sync_configs	Not run (safety)
DELETE	/api/v1/account/ {account_id}/backup/ {backup_id}	Not run (safety)
POST	/api/v1/account/ {account_id}/backup/ {backup_id}/restore	Not run (safety)
GET	/api/v1/account/ {account_id}/backups	Confirmed live · 403
POST	/api/v1/account/ {account_id}/backups	Confirmed live · 403
DELETE	/api/v1/account/ {account_id}/kb/{kb_id}	Not run (safety)
GET	/api/v1/account/ {account_id}/kb/{kb_id}	Verified (2xx) · 200
PATCH	/api/v1/account/ {account_id}/kb/{kb_id}	Verified (2xx) · 200
POST	/api/v1/account/ {account_id}/kb/ {kb_id}/ephemeral_tokens	Verified (2xx) · 201
DELETE	/api/v1/account/ {account_id}/kb/ {kb_id}/invite	Not run (safety)
POST	/api/v1/account/ {account_id}/kb/	Not run (safety)

	<code>{kb_id}/invite</code>	
GET	<code>/api/v1/account/ {account_id}/kb/ {kb_id}/invites</code>	Verified (2xx) · 200
GET	<code>/api/v1/account/ {account_id}/kb/ {kb_id}/logo.png</code>	Server 5xx · 500
GET	<code>/api/v1/account/ {account_id}/kb/ {kb_id}/permissions</code>	Verified (2xx) · 200
DELETE	<code>/api/v1/account/ {account_id}/kb/ {kb_id}/ service_account/{sa_id}</code>	Not run (safety)
DELETE	<code>/api/v1/account/ {account_id}/kb/ {kb_id}/ service_account/ {sa_id}/key/{sa_key_id}</code>	Not run (safety)
POST	<code>/api/v1/account/ {account_id}/kb/ {kb_id}/ service_account/ {sa_id}/keys</code>	Not run (safety)
GET	<code>/api/v1/account/ {account_id}/kb/ {kb_id}/service_accounts</code>	Verified (2xx) · 200
POST	<code>/api/v1/account/ {account_id}/kb/ {kb_id}/service_accounts</code>	Verified (2xx) · 201
GET	<code>/api/v1/account/ {account_id}/kb/ {kb_id}/users</code>	Verified (2xx) · 200
PATCH	<code>/api/v1/account/ {account_id}/kb/ {kb_id}/users</code>	Not run (safety)
GET	<code>/api/v1/account/ {account_id}/kbs</code>	Verified (2xx) · 200
POST	<code>/api/v1/account/</code>	Verified (2xx) · 201

	<code>{account_id}/kbs</code>	
DELETE	<code>/api/v1/account/ {account_id}/ nua_client/{client_id}</code>	Not run (safety)
GET	<code>/api/v1/account/ {account_id}/ nua_client/{client_id}</code>	Confirmed live · 403
PATCH	<code>/api/v1/account/ {account_id}/ nua_client/{client_id}</code>	Confirmed live · 403
PUT	<code>/api/v1/account/ {account_id}/ nua_client/{client_id}/ key</code>	Confirmed live · 403
GET	<code>/api/v1/account/ {account_id}/nua_clients</code>	Confirmed live · 403
POST	<code>/api/v1/account/ {account_id}/nua_clients</code>	Confirmed live · 403
POST	<code>/api/v1/ephemeral_token</code>	Not run (safety)
GET	<code>/api/v1/kb/{kb_id}/ activity/ download_request/ {request_id}</code>	Confirmed live · 403
GET	<code>/api/v1/kb/{kb_id}/ activity/metrics</code>	Confirmed live · 403
GET	<code>/api/v1/kb/{kb_id}/ activity/{event_type}/ months</code>	Confirmed live · 403
POST	<code>/api/v1/kb/{kb_id}/ activity/{event_type}/ query</code>	Confirmed live · 403
POST	<code>/api/v1/kb/{kb_id}/ activity/{event_type}/ query/download</code>	Confirmed live · 403
DELETE	<code>/api/v1/kb/{kb_id}/ external_connection/ {connection_id}</code>	Not run (safety)
GET	<code>/api/v1/kb/{kb_id}/</code>	Confirmed live · 403

	<code>external_connection/ {connection_id}</code>	
GET	<code>/api/v1/kb/{kb_id}/ external_connection/ {connection_id}/browse</code>	Confirmed live · 403
GET	<code>/api/v1/kb/{kb_id}/ external_connection/ {connection_id}/ resolve_site</code>	Confirmed live · 403
GET	<code>/api/v1/kb/{kb_id}/ external_connections</code>	Confirmed live · 403
POST	<code>/api/v1/kb/{kb_id}/ external_connections</code>	Not run (safety)
GET	<code>/api/v1/kb/{kb_id}/ external_connections/ s3/assume_role_info</code>	Confirmed live · 403
POST	<code>/api/v1/kb/{kb_id}/ remi/query</code>	Confirmed live · 403
DELETE	<code>/api/v1/kb/{kb_id}/ sync_config/{config_id}</code>	Not run (safety)
GET	<code>/api/v1/kb/{kb_id}/ sync_config/{config_id}</code>	Confirmed live · 403
PATCH	<code>/api/v1/kb/{kb_id}/ sync_config/{config_id}</code>	Not run (safety)
POST	<code>/api/v1/kb/{kb_id}/ sync_config/ {config_id}/authorize</code>	Not run (safety)
GET	<code>/api/v1/kb/{kb_id}/ sync_config/ {config_id}/jobs</code>	Confirmed live · 403
POST	<code>/api/v1/kb/{kb_id}/ sync_config/ {config_id}/sync</code>	Not run (safety)
POST	<code>/api/v1/kb/{kb_id}/ sync_config/ {config_id}/ validate_resources</code>	Not run (safety)
GET	<code>/api/v1/kb/{kb_id}/</code>	Confirmed live · 403

	sync_configs	
POST	/api/v1/kb/{kb_id}/ sync_configs	Not run (safety)
GET	/api/v1/kb/{kb_id}/ sync_job/{job_id}/logs	Confirmed live · 403
GET	/api/v1/kb/{kb_uuid}/ remi/events/{event_id}	Confirmed live · 403
GET	/api/v1/kb/{kb_uuid}/ remi/scores	Confirmed live · 403
POST	/api/v1/ service_account_agent_key	Not run (safety)
POST	/api/v1/ service_account_temporal_key	Not run (safety)

global (27 operations)

Method	Path	Verification
DELETE	/api/v1/account/ {account_id}/invite/ {email}	Needs user/PAT
GET	/api/v1/account/ {account_id}/invites	Needs user/PAT
GET	/api/v1/account/ {account_id}/usage	Needs user/PAT
GET	/api/v1/account/ {account_slug_or_id}	Needs user/PAT
GET	/api/v1/account/ {account_slug_or_id}/ zones	Needs user/PAT
DELETE	/api/v1/account/ {account_slug}	Needs user/PAT
PATCH	/api/v1/account/ {account_slug}	Needs user/PAT
POST	/api/v1/account/	Needs user/PAT

	<code>{account_slug}/invite</code>	
GET	<code>/api/v1/account/ {account_slug}/logo.png</code>	Needs user/PAT
PATCH	<code>/api/v1/account/ {account_slug}/logo.png</code>	Needs user/PAT
GET	<code>/api/v1/account/ {account_slug}/ permissions</code>	Needs user/PAT
GET	<code>/api/v1/account/ {account_slug}/user/ {user_id}</code>	Needs user/PAT
GET	<code>/api/v1/account/ {account_slug}/users</code>	Needs user/PAT
PATCH	<code>/api/v1/account/ {account_slug}/users</code>	Needs user/PAT
GET	<code>/api/v1/account/ {account_slug}/users/ search</code>	Needs user/PAT
GET	<code>/api/v1/accounts</code>	Needs user/PAT
POST	<code>/api/v1/accounts</code>	Needs user/PAT
DELETE	<code>/api/v1/user</code>	Needs user/PAT
GET	<code>/api/v1/user</code>	Needs user/PAT
PATCH	<code>/api/v1/user</code>	Needs user/PAT
DELETE	<code>/api/v1/user/pa_token/ {token_id}</code>	Needs user/PAT
GET	<code>/api/v1/user/pa_tokens</code>	Needs user/PAT
POST	<code>/api/v1/user/pa_tokens</code>	Needs user/PAT
GET	<code>/api/v1/user/welcome</code>	Needs user/PAT
GET	<code>/api/v1/zones</code>	Needs user/PAT
POST	<code>/oauth2/register</code>	Needs user/PAT
PUT	<code>/oauth2/register/ {client_id}</code>	Needs user/PAT